

On the design of Zeth transaction relays

Antoine Rondelet* and Duncan Tebbs*

*Clearmatics, UK

2020-12-21

Abstract

6 Privacy preserving protocols suffer from the need to pay transaction fees on blockchain
7 systems. While such fees constitute a sound economic barrier to a wide class of Denial
8 of Service attacks, they also represent impediments to the design of state transitions
9 with anonymous initiators. In this paper, we navigate the design space for “Zeth cryp-
10 tographic relays”. These enable Zeth users to carry out Zeth payments anonymously
11 by relying on an extra party – a relay – to settle and execute state transitions on the
12 blockchain.

13 **Keywords**— Privacy, Ethereum, Zeth, Relays, Sender Anonymity, Digital Cash

Contents

15	1 Preliminaries	3
16	1.1 Prerequisites	3
17	1.2 Notation	3
18	1.3 Terminology	4
19	1.4 Introduction	4
20	1.4.1 Turning “front-runners” into relays	5
21	1.4.2 Relay incentives and risks	5
22	1.4.3 Assumptions	6
23	2 Relays with Proof of Permission	8
24	2.1 Protocol Overview	9
25	2.2 Protocol	10
26	2.2.1 Relay Request and Permission Data	10
27	2.2.2 User Operations	11
28	2.2.3 Relay Operations	12
29	2.2.4 RelayExec Contract	15
30	3 Relays with Private Fees	18
31	3.1 Introduction	18
32	3.2 Protocol overview	18
33	3.2.1 Relay-originated mix transactions	19
34	3.3 Limitations and extensions to the protocol	19
35	3.3.1 Limitation of output notes	19
36	3.3.2 Increasing JSOUT	19
37	3.3.3 Support for Ether output	20
38	3.3.4 Fees as Ether or Zeth notes	21
39	A Relays with Stake	22
40	A.1 Introduction	22
41	A.2 Protocol overview	23
42	A.2.1 Stake contract	23
43	A.2.2 Relay	24
44	A.2.3 User	25

45	A.2.4 Reclaiming stake	25
46	A.3 Remarks	26
47	B Network structure	28
48	B.1 Binding requests to relays	28
49	B.1.1 Background	28
50	B.1.2 Emulating “free” relay requests	28
51	B.2 Unicast vs broadcast networks	29
52	B.3 Network anonymity	30

DRAFT

Chapter 1

Preliminaries

1.1 Prerequisites

This document assumes familiarity with **Ethereum** and **Zeth**. It does not, in any way, aim to replace the Ethereum yellow paper [Woo19] or the Zeth specifications [Cle20]. The reader is strongly advised to read about **Ethereum** and **Zeth** before delving into this document.

1.2 Notation

Unless stated otherwise, this document follows the notation of the **Zeth** protocol specifications [Cle20]. We note in particular the following notations (some of which originate from [Cle20]), used throughout the document.

$\widetilde{\text{Mixer}}$ A deployed instance of the Zeth contract.

\mathcal{U} A user, with identity $\mathcal{U}_{\mathcal{E}}$ on the **Ethereum** network, and/or Zeth identity $\mathcal{U}_{\mathcal{Z}}$.

\mathcal{R} An entity operating a relay, with identity $\mathcal{R}_{\mathcal{E}}$ on the **Ethereum** network, and/or Zeth identity $\mathcal{R}_{\mathcal{Z}}$.

$\widetilde{\text{RelayExec}}$ A contract deployed on the blockchain acting as a proxy to $\widetilde{\text{Mixer}}$ (removing the need for trust between users and relays).

$fee_{\mathcal{R}}$ The fee of the relay service \mathcal{R} .

genCallTx Algorithm which creates, signs and broadcasts a transaction to call a contract entry point with specific parameters. That is, given a contract method call of the form **Contract.method**(param_1, \dots), a secret key sk for some **Ethereum** address, a gas price $gasP$ and gas limit $gasL$, the algorithm **genCallTx**(**Contract.method**(param_1, \dots), $sk, gasP, gasL$) creates a signed transaction that performs the given contract call.

77 `broadcastTx` Algorithm that accepts a signed `Ethereum` transaction and broadcasts it
78 to the network, returning the transaction ID. The caller (in this document, the
79 relay) can use the transaction ID to monitor the asynchronous completion of the
80 transaction. The exact details will depend on the relay implementation, but once
81 the transaction is complete, the relay can retrieve its result and update any internal
82 state.

83 1.3 Terminology

84 The key words `MUST`, `MUST NOT`, `SHOULD`, `SHOULD NOT`, `MAY`, and `RECOMMENDED` in this docu-
85 ment are to be interpreted as described in [Bra97] when they appear in `ALL CAPS`. These
86 words may also appear in this document in lower case as plain English words, absent
87 their normative meanings.

88 1.4 Introduction

89 The `Zeth` protocol allows users to carry out privacy-preserving state transactions on
90 “smart-contract enabled blockchains” such as `Ethereum` or `Autonity`¹. Like all `Ethereum`
91 transactions, `Zeth` transactions require a fee to be paid. This is inherited from the base
92 protocol, which uses transaction fees as a security mechanism against Denial of Service
93 (DoS) attacks.

94 As pointed out in the `Zeth` paper [RZ19], the need to pay transaction fees represents
95 a challenge for designers of privacy preserving protocols, since transaction originators
96 must carry out `Zeth` contract calls from a funded `Ethereum` address (which in turn must
97 have been funded by other user(s) on the system², and for which the “controlling user
98 identity” must be known by at least one network member). As such, while `Zeth` provides
99 strong privacy guarantees (recipient anonymity, private payment amount etc.), sender
100 anonymity remains hard to achieve.

101 This document proposes some designs for “cryptographic relays” and investigates the
102 space of tradeoffs for both users and relay operators. The protocols enable `Ethereum`
103 peer-to-peer (P2P) nodes to act as *relays*, receiving requests (off-chain), and signing and
104 broadcasting transactions (which incorporate these requests) on behalf of `Zeth` users. In
105 exchange for this service, relays receive some fee from the original users.

106 As described below, relay fees are of paramount importance in establishing a sound
107 incentive structure, which in turn is necessary for the overall robustness of the system.
108 The primary goal of this study on “cryptographic relays” is to achieve `Zeth` sender
109 anonymity on blockchain systems, and the proposals in this document suggest multiple
110 ways in which relay fees can be paid while maintaining this anonymity. Furthermore,
111 under additional network assumptions (e.g. namely that `Zeth` users and relay nodes
112 communicate via an Anonymous Communication (AC) network, e.g. [PHE⁺17]), `Zeth`

¹<https://github.com/clearmatics/autonity>

²Unless the user is a miner.

113 users can make use of relay nodes without revealing any identifying information to the
114 relay. See Appendix B.3 for further discussion.

115 1.4.1 Turning “front-runners” into relays

116 As noted in [DGK⁺19] and [lsa20] (among others), on blockchains such as **Ethereum**,
117 so-called “front-runners” actively seek out transactions that are profitable for the sender
118 and attempt to replace them with modified versions, in order to steal the profit from the
119 original sender. Front-running strategies leverage the mempool ordering policy adopted
120 by miners. Namely, they set higher gas prices in order to overtake the targeted transac-
121 tions.

122 With the proliferation of “bots”³ inspecting the mempool and front-running prof-
123 itable transactions (see e.g. [RK20]), it becomes key for “layer 2”-protocol designers
124 to design state transitions that are secure against such replay attacks. As presented
125 in [Cle20, Section 2.3], the **Zeth** protocol prevents “front-running”/“replay” attacks by
126 design (see derivation of *hsig* and *dataToBeSigned*).

127 While “front-runners” present a threat to users of Decentralized Applications (DApps),
128 they can potentially be leveraged to act as transaction relays [lsa20]. Since front-runners
129 examine the mempool, looking for profitable transactions to overtake (by extracting the
130 transaction payload, creating a new transaction, signing it and broadcasting it on the
131 network with a higher gas price), a user may exploit this behavior by voluntarily broad-
132 casting a transaction with low gas price on the network, in the hope that a front-runner
133 will replay/overtake it. By doing so, the user may thus trigger a state transition on the
134 blockchain without paying the associated transaction fees. Nevertheless, “front-runners”
135 should be modelled as rational agents, meaning that such transactions must be profitable
136 to them. As such, for users to leverage “front-runners” as “relays” in practice, transac-
137 tions must be crafted such that “front runners” receive a fee in exchange for replacing
138 them.

139 Finally, in order for a user’s transactions to be added to a miner’s transaction pool,
140 it is necessary for the transaction to pass the “initial tests of intrinsic validity” [Woo19,
141 Section 6] (see also transaction pool implementation in Geth⁴). This means that users
142 who wish to have their transactions “front-run”/“relayed” must hold a funded **Ethereum**
143 account. This may not be desirable in all scenarios (especially in settings where sender
144 anonymity is a primary motivation).

145 1.4.2 Relay incentives and risks

146 Besides the potential profitability of “front-running”, mentioned above, relaying trans-
147 actions that haven’t been added to a miner’s mempool is inherently risky, and sound
148 incentive structures must reward such risk appropriately.

149 Firstly, it must be noted that relay nodes may be vulnerable to DoS attacks. In
150 such attacks, malicious clients “flood” targeted relays with an overwhelming stream of

³see, for instance, <https://github.com/Uniswap/uniswap-interface/issues/248>

⁴https://github.com/ethereum/go-ethereum/blob/master/core/tx_pool.go#L578-L583

151 transactions. While such attacks may be mitigated by relay operators using existing
 152 network monitoring techniques (e.g. packet filtering, rate limiting etc.), it is also impor-
 153 tant for relay operators to assess the profitability of the transactions that they relay to
 154 the blockchain. In fact, running a relay may quickly become a “money drain” if the cost
 155 of operating the relay service (i.e. infrastructure costs, transaction fees etc.) outweighs
 156 the relay fees received. As such, it is necessary for relays to have an efficient way to
 157 gauge the on-chain cost and profitability of a transaction. (Carrying out this operation
 158 may also exacerbate the DoS vector on relays, since a flood of maliciously crafted trans-
 159 actions – such as transactions that take a long time to execute, but fail to release any
 160 funds – may cause a relay node to spend all of its resources on transaction verification in
 161 return for no income⁵.) Additionally, it is worth remembering that “front runners” and
 162 “relayers” may in turn be front-run by other competitors. As a consequence, allocat-
 163 ing non-negligible computation resources to “assessing the profitability” of transactions
 164 represents a risk – other “front runners” may overtake the relay’s (verified) transaction
 165 to avoid carrying out this verification work locally.

166 1.4.3 Assumptions

167 Based on the remarks given in the previous sections, we make the following assumptions
 168 in the rest of the document.

169 **Transactions to be relayed are immune to front-running.** We assume that all re-
 170 layed transactions are inputs to the Zeth contract Mix function. As such, the inputs
 171 prevent front-running by construction.

172 **Transactions to be relayed target specific relays.** As mentioned above, Zeth is
 173 designed to avoid “front-running”/“malleability” attacks. To achieve this, several
 174 parameters (e.g. *dataToBeSigned*) are derived using the address of the **Ethereum**
 175 user that must send the transaction. As such, we assume that users *choose* a relay
 176 service to process their request. Note that, as discussed in Appendix B.1.2, users
 177 are free to target multiple relay services by creating multiple requests, but the
 178 underlying assumption is there exists a market of competing “relays”, from which
 179 users can select the service that best suits their needs. For example, different re-
 180 lays may offer different trade-offs between settlement latency and fees, while others
 181 may offer “aggregation services” (see, e.g. [Ron20])

182 **Relays are “discoverable”.** Discovery of relay nodes by users may be achieved through
 183 several possible mechanisms. For example, relay nodes may publish their IP ad-
 184 dress, current fees and **Ethereum** address (some of which may potentially be pub-
 185 lished on-chain), allowing users to discover their services. Overall, we assume that
 186 relay operators take the necessary steps to be “discoverable” by users.

⁵Additional security measures may alleviate such issues. For example, using properly crafted veri-
 fication thresholds, discarding transactions that take too long to be verified. Note however that such
 mechanisms “specialize” a relay into relaying only certain classes of transactions. Again, proper tradeoffs
 need to be adopted depending on the use-cases and threat model.

187 In the remainder of this document, we propose a set of protocols to relay Zeth
188 transactions, each with their own trade-offs and specific goals.

Note

Importantly, it is worth keeping in mind that in most scenarios, sound “relay economics” will imply that for a given state transition, relay fees are greater than the on-chain fees normally paid by blockchain users. Hence, we stress that using relays should not be seen as a “cheap way” to transact on a blockchain, but rather as a way to achieve otherwise impossible objectives on the system (e.g. to achieve sender anonymity).

189

DRAFT

Chapter 2

Relays with Proof of Permission

In this section, we propose a protocol with the aim of preserving the anonymity of Zeth transaction senders. This allows a user \mathcal{U}_Z to interact (anonymously) with a Zeth deployment to either “pour” the value of some of his *ZethNotes* into new ones (i.e. carry out a private payment), or withdraw some value *vout* from **Mixer** to a newly created **Ethereum** address $Addr_{new}$.

We assume that \mathcal{U}_Z is willing to pay a fee to achieve this anonymity, and that at least one party \mathcal{R}_E is willing to act as a “relay” in return for this fee.

By providing a mechanism to carry out private withdrawals of Zeth funds to a newly generated Ethereum address, we allow Ethereum users to manipulate Ether “privately” in future transactions - either “plain EOA-to-EOA transactions” or smart-contract calls.

We list below the set of characteristics that are desired in this setting.

From the user’s perspective

- The user must be able to leverage relays to anonymously carry out a Zeth state transition on-chain. This includes, anonymous private transfer (i.e. “pouring” the value of existing *ZethNotes* to new ones), and anonymous withdrawals to a newly created **Ethereum** account.
- Not only must the user be anonymous with regard to the blockchain network, but he must also remain anonymous to the relay¹ for the mechanism to be robust against malicious/compromised relays. The user must not be required to reveal any identifying information, including any pre-funded **Ethereum** addresses.
- The user must be guaranteed (up to some negligible probability) that the relay will only call the **Mix** function, on-chain, with the *correct* inputs (i.e. the relay may not execute any state transition on behalf of the user that the user did not request).

¹Further assumptions need to be made about the underlying network. The user must be able to communicate with the relay without revealing any network-layer identifying information.

215 From the relay's perspective

- 216 • The relay must be guaranteed (up to some negligible probability) that he will
217 receive the agreed upon fee in exchange for carrying out his role in the protocol
218 (and therefore he will not incur costs such as transaction fees for no reward)².

219 2.1 Protocol Overview

220 We start by assuming that a Zeth user \mathcal{U}_Z has chosen a relay service \mathcal{R} (with Ethereum ac-
221 count \mathcal{R}_E) which relays transactions for a fee $fee_{\mathcal{R}}$. We further assume that \mathcal{U}_Z knows
222 the address of the **RelayExec** contract.

223 The protocol consists of the following steps:

224 **Step 1 (User creates the Mix parameters).** \mathcal{U}_Z creates the Mix parameters $mixParams$
225 that spend her note(s), including the public output value $vout$. $mixParams$ are
226 generated such that only **RelayExec** can successfully use them. (This is achieved
227 by leveraging the property of Mix parameters, described in [Cle20, Sections 2.4,
228 2.5], which restricts the Ethereum address of the caller, possibly a contract³.)

229 **Step 2 (User generates a proof-of-relay-permission).** With $mixParams$ properly
230 created, \mathcal{U}_Z generates a proof-of-relay permission $\pi_{\mathcal{R}}^{(mixParams)}$ (described in further
231 detail below) for $mixParams$. This proves that the owner of the Zeth notes to be
232 spent by $mixParams$ agrees that \mathcal{R}_E may relay the Mix parameters via **RelayExec**
233 for a fee $fee_{\mathcal{R}}$. \mathcal{U}_Z also specifies the address $outAddr$ to which the remaining balance
234 $vout - fee_{\mathcal{R}}$ (if any) should be sent. (In general, $outAddr$ is expected to be a newly
235 generated address $Addr_{new}$).

236 **Step 3 (User sends parameters to the relay).** \mathcal{U}_Z sends a “relay request” req to
237 the chosen relay \mathcal{R} , containing $mixParams$, $\pi_{\mathcal{R}}^{(mixParams)}$ and other data such as
238 $outAddr$. Note that, as long as $outAddr$ is a newly generated address (with no
239 history on the blockchain), this request contains no information that identifies \mathcal{U}_Z
240 as the originator. \mathcal{U}_Z is also expected to leverage anonymising mechanisms to avoid
241 revealing any identifying information at the transport level.

242 **Step 4 (Relayer verifies and broadcasts the received request).** Upon receipt of
243 req , the relay performs a set of checks to gain confidence that $mixParams$ and
244 $\pi_{\mathcal{R}}^{(mixParams)}$ are valid, and indeed grant the relay fee to \mathcal{R}_E . (Note that the relay
245 has some scope to choose the extent of such checks, trading off the cost of checking
246 against the risk of losing money by broadcasting an invalid transaction.) If the
247 relay is satisfied that req is valid, he signs (using the \mathcal{R}_E identity) and broadcasts
248 a transaction that calls **RelayExec**.

²This would question the profitability of operating a relay node and would jeopardize the “relay network” as well as the viability of the “relaying activity”.

³This is currently the mechanism used to prevent front-runners from claiming $vout$.

249 **Step 5 (The intermediary contract checks all parameters and executes the Zeth mixer).**

250 The **RelayExec** contract acts as an intermediary, trusted by both users and re-
251 lays. It first checks $\pi_{\mathcal{R}}^{(mixParams)}$ to ensure that the caller ($\mathcal{R}_{\mathcal{E}}$) has indeed been
252 granted permission to use the $mixParams$ in exchange for $fee_{\mathcal{R}}$. **RelayExec** then
253 calls **Mixer** with parameters $mixParams$ and checks that the call succeeds. The
254 transaction is aborted if any of these checks fail.

255 **Step6 (The intermediary contract distributes the value $vout$).** If the Mix call is
256 successful, **RelayExec** now holds the $vout$ from **Mixer**. From this, it pays $fee_{\mathcal{R}}$
257 to $\mathcal{R}_{\mathcal{E}}.Addr$ and the remainder $vout - fee_{\mathcal{R}}$ to $outAddr$.

258 **Remark 1.** Using **RelayExec** to distribute the fee and fund the user-specified $outAddr$
259 address, gives $\mathcal{U}_{\mathcal{Z}}$ confidence that $outAddr$ will receive the correct output for the agreed
260 fee. Further, $\mathcal{R}_{\mathcal{E}}$ can be sure that no other relay can use the same set of Mix parameters
261 and forge a request from $\mathcal{U}_{\mathcal{Z}}$ to receive the relay fee.

262 2.2 Protocol

263 Below we give further details of the protocol outlined above, enabling anonymous **Zeth**
264 transfers via relayed transactions. The protocol leverages specific characteristics of the
265 **Zeth** protocol design (some of which were described above). In particular, it builds on
266 the fact that Mix parameters can be generated without owning an **Ethereum** account (see
267 derivation of $hsig$ and related discussion [Cle20, Remark A.2.2]), and that Mix parameters
268 are “bound” to the address of the **Ethereum** account that must call the **Mixer** contract
269 (see derivation of $dataToBeSigned$ [Cle20, Section 2.3]).

270 2.2.1 Relay Request and Permission Data

271 We use **MixInputDType** and related datatypes from [Cle20, Section 2.1], and define the
272 following new data type to represent a relay request with proof of relay permission.

273 **Definition 1.** The datatype **RelayRequest** is defined as:

Field	Description	Value
<i>mixParams</i>	Parameters to the Mix call	MixInputDType
<i>outAddr</i>	Ethereum address credited with the funds withdrawn from Mixer	$\mathbb{B}^{\text{ADDRLEN}}$
<i>relayAddr</i>	Ethereum address allowed to relay <i>mixParams</i> (and perceive <i>fee</i>)	$\mathbb{B}^{\text{ADDRLEN}}$
<i>fee</i>	Fee to pay (out of <i>vout</i>) to the authorized relay	$\mathbb{N}^{\text{ETHWORDLEN}}$
<i>permission</i>	Signature proving the authenticity of the request	SigOtsDType

Table 2.1: **RelayRequest** type

274 The *permission* attribute is used to indicate that the user has given permission for
275 the relay to forward the specific Mix call. We reuse the signature key *mixParams.otsvk*
276 for the scheme $\text{SigSch}_{\text{OT-SIG}}$, used to create *mixParams.otssig* (see [Cle20, Section 2.3]),
277 since only the author of the Mix call parameters can generate valid signatures.

TODO

Tighten the security requirements of the signature scheme used. For now, this proposal uses $\text{SigSch}_{\text{OT-SIG}}$ to create a second signature with the same private key. However, $\text{SigSch}_{\text{OT-SIG}}$ is “one-time”. Additionally, the security claim on the nested signature is only that it is UF-CMA (see specs), relying on the fact that the wrapping signature scheme (signing the transaction object) is SUF-CMA. Here however, since the relay request does not contain a signed blockchain transaction object, an adversary may be able to maul the signature and pass the UF-CMA game. More consideration of the threat model and the security requirements is required here.

278
279 As described below, the user must sign the attributes *relayAddr*, *fee* and *outAddr*
280 for the signature to be considered valid.

2.2.2 User Operations

282 We assume that user \mathcal{U}_Z has decided to use \mathcal{R} (controlling **Ethereum** account \mathcal{R}_E) to
283 relay her **Zeth** transaction in order to either withdraw some value *vout* to **Ethereum**
284 address *outAddr* and/or carry out a private transfer. We further assume that \mathcal{U}_Z agrees
285 to pay a fee $fee_{\mathcal{R}}$ to \mathcal{R}_E in order to achieve this. Here $fee_{\mathcal{R}}$ is the fee that \mathcal{R} is willing to
286 accept in exchange for relaying a single Mix call.⁴

287 \mathcal{U}_Z executes the following steps:

⁴This fee should be strictly greater than the gas cost of Mix in order for the relay to be profitable (more refined profitability forecasts would internalize the infrastructure operational costs to adjust the fee). Additionally, the relay may adjust and republish $fee_{\mathcal{R}}$ in light of gas price fluctuations on the blockchain, other changes to cost and risk, or to compete with other relays.

- 288 1. Create a valid $mixParams \in \text{MixInputDType}$, where:
- 289 • $mixParams$ spends previously unspent *ZethNotes* owned by \mathcal{U}_Z , with a public
 - 290 output of $vout$.
 - 291 • $mixParams.otssig$ is created using the address of $\widetilde{\text{RelayExec}}$, as described
 - 292 in [Cle20, Section 2.3]
 - 293 • The one-time signing key $sk_{\text{OT-SIG}}$, used to create $mixParams.otssig$, can be
 - 294 (securely) extracted for use in the following step.

2. Use the signing key $sk_{\text{OT-SIG}}$ to create a proof of relay permission:

$$\begin{aligned} data_{\mathcal{R}} &\leftarrow \text{encode}(\mathcal{R}_{\mathcal{E}}.Addr) \parallel \text{encode}(fee_{\mathcal{R}}) \parallel \text{encode}(outAddr) \\ \pi_{\mathcal{R}}^{(mixParams)} &\leftarrow \text{SigSch}_{\text{OT-SIG}}.\text{Sig}(sk_{\text{OT-SIG}}, \text{CRH}^{\text{ots}}(data_{\mathcal{R}})) \end{aligned}$$

3. Create a relay request $req \in \text{RelayRequest}$:

$$\begin{aligned} req &\leftarrow \{ \\ &\quad mixParams : mixParams, \\ &\quad outAddr : outAddr, \\ &\quad relayAddr : relayEthAccount.Addr, \\ &\quad fee : fee_{\mathcal{R}}, \\ &\quad permission : \pi_{\mathcal{R}}^{(mixParams)} \\ &\} \end{aligned}$$

295 **Remark 2.** If the user simply wants to leverage a relay to carry out a private *Zeth*

296 transfer (without withdrawing funds to a newly created address $outAddr$), \mathcal{U}_Z can set

297 $outAddr \leftarrow 0x0$ (see Fig. 2.3 for more details).

298 The user then sends this request to the relay \mathcal{R} , via a secure (anonymous) commu-

299 nication channel.

300 2.2.3 Relay Operations

301 Let $\mathcal{R}_{\mathcal{E}}.sk$ be the secret key corresponding to the address $\mathcal{R}_{\mathcal{E}}.Addr$ of *Ethereum* account

302 $\mathcal{R}_{\mathcal{E}}$. In what follows, we assume that $\mathcal{R}_{\mathcal{E}}.Addr$ is funded with enough *Ether* to pay the

303 gas required to call $\widetilde{\text{RelayExec}}$ on-chain.

304 Let $\mathcal{R}_{\mathcal{E}}.Addr$ be the *Ethereum* address of \mathcal{R} charging relay fee $fee_{\mathcal{R}}$. Further, as-

305 sume that $\widetilde{\text{RelayExec}}$ and $\widetilde{\text{Mixer}}$ are deployed with addresses $\widetilde{\text{RelayExec}}.Addr$ and

306 $\widetilde{\text{Mixer}}.Addr$ respectively.

307 Given the current *Ethereum* state ς (or a copy holding at least $\varsigma[\widetilde{\text{RelayExec}}.Addr]$

308 and $\varsigma[\widetilde{\text{Mixer}}.Addr]$) and a relay request $req \in \text{RelayRequest}$, the algorithm RelayCheckRequest

309 (see Fig. 2.1) succeeds if the request req is valid and will result in $\mathcal{R}_{\mathcal{E}}$ receiving $fee_{\mathcal{R}}$.

Note that we assume the existence of an algorithm `RelayCheckMixParams` which checks whether a given set of `Mix` parameters $\widetilde{mixParams}$ result in a successful `Mix` call in the context of the current blockchain state $\varsigma[\widetilde{\mathbf{Mixer}.Addr}]$. That is, given the state $\varsigma[\widetilde{\mathbf{Mixer}.Addr}]$ of the `Mixer` contract, an `Ethereum` address $Addr_{caller} \in \mathbb{B}^{\text{ADDRLEN}}$ and `Mix` parameters $\widetilde{mixParams} \in \text{MixInputDType}$,

`RelayCheckMixParams`($\varsigma[\widetilde{\mathbf{Mixer}.Addr}]$, $Addr_{caller}$, $\widetilde{mixParams}$)

310 returns the result of `ZethVerifyTx`(tx) (see [Cle20, Section 2.5]) where tx is a transaction
 311 that calls `Mix`($\widetilde{mixParams}$) from address $Addr_{caller}$, and `ZethVerifyTx` executes in the
 312 context of `Mixer` with state $\varsigma[\widetilde{\mathbf{Mixer}.Addr}]$.

`RelayCheckRequest`(ς , req)

```

1 : // Check the fee and relay address in the request
2 : if (req.fee ≠ feeR) ∨ (req.relayAddr ≠ RE.Addr) then
3 :   return false
4 : endif
5 : // Check that vout can pay the fee and reject deposits
6 : if (req.mixParams.vout < feeR) ∨ (req.mixParams.vin ≠ 0) then
7 :   return false
8 : endif
9 : // Check the proof-of-relay-permission
10 : dataR ← encode (RE.Addr) || encode(feeR) || encode(req.outAddr)
11 : if ¬SigSchOT-SIG.Vf(req.mixParams.otsvk, CRHots(dataR), req.permission) then
12 :   return false
13 : endif
14 : return RelayCheckMixParams(ϰ[Mixer.Addr], RE.Addr, req.mixParams)

```

Figure 2.1: `RelayCheckRequest` algorithm. The relay address $R_E.Addr$, desired relay fee fee_R and contract addresses `RelayExec.Addr` and `Mixer.Addr` are implicitly available as variables.

313 **Remark 3.** For this check to be meaningful, we assume here that R_E has access to
 314 $\varsigma[\widetilde{\mathbf{Mixer}.Addr}]$ for some recent block height⁵.

315 Further, we introduce the relay logic in Fig. 2.2 which illustrates the `RelayHandleRequest`
 316 algorithm that processes a received `RelayRequest` request object req and relays it
 317 on the blockchain network. Here again, we assume that the tuple $(R_E.Addr, fee_R,$

⁵Implementations are not necessarily expected to track $\varsigma[\widetilde{\mathbf{Mixer}.Addr}]$ by themselves, but rather to leverage an existing `Ethereum` full node implementation. `RelayCheckMixParams` can then be implemented as queries to the node (e.g. via RPC).

318 $\widetilde{\text{RelayExec.Addr}}, \widetilde{\text{Mixer.Addr}}$) is implicitly available to the algorithm. In `RelayHandleRequest`
 319 a gas price $gasP$ and gas limit $gasL$ are (possibly dynamically) set at the relay's discre-
 320 tion (based on its “relaying service and strategy”, defining the trade-off between relay
 321 fees, settlement latency, etc.) and passed as explicit parameters to `RelayHandleRequest`.

`RelayHandleRequest(ς, req)`

```

1 : // Check the request
2 : if  $\neg \text{RelayCheckRequest}(\varsigma, req)$  then
3 :   abort
4 : endif
5 : // Create and broadcast the relay Ethereum transaction
6 :  $tx_{\mathcal{R}} \leftarrow \text{genCallTx}(\widetilde{\text{RelayExec.ProcessRequest}}(req), \mathcal{R}_{\mathcal{E}}.sk, gasP, gasL)$ 
7 :  $txid \leftarrow \text{broadcastTx}(tx_{\mathcal{R}})$ 
8 : // Record the transaction id
9 : recordTransaction( $txid$ )
10 : return

```

Figure 2.2: Algorithm to process relay requests. `recordTransaction` represents the relay-specific handling of the transaction id.

322 **Remark 4.** *The checks in `RelayCheckRequest` provide some level of assurance that $tx_{\mathcal{R}}$*
 323 *will be successfully executed on chain. However, the relay cannot rule out the possibility*
 324 *that a conflicting transaction $tx_{\mathcal{R}}^*$ exists on the network, such that, if $tx_{\mathcal{R}}^*$ were mined*
 325 *first it would alter the blockchain state and affect the execution of $tx_{\mathcal{R}}$. For example,*
 326 *in the case of **Zeth**, if some transaction $tx_{\mathcal{R}}^*$ which spends the same notes as $tx_{\mathcal{R}}$ were*
 327 *mined first, $tx_{\mathcal{R}}$ would fail and no payment would be made to the relay.*

328 *Relays may perform further checks to increase their level of confidence that $tx_{\mathcal{R}}$ will*
 329 *execute as expected (such as examining their own mempool for conflicting transactions)*
 330 *but all such checks add to the cost of running a relay, which may be reflected in the relay*
 331 *fees. It is expected that relays will compete with each other on the basis of their fees.*
 332 *Consequently, relays are expected to adopt some strategy that trades off risk, validation*
 333 *costs and competitiveness, and thereby determine an appropriate price range for relaying.*

Note

Most relay implementations are likely to be able to receive and process multiple requests simultaneously, and at times may receive requests at a faster rate than they can be processed. In this case, the relay has scope to prioritise certain requests over others. In the simplest case, requests may be handled in the order in which they are received (for example via a FIFO queue). More sophisticated relays may employ a strategy allowing them to accept “relay bribes”, in which case $fee_{\mathcal{R}}$ is composed of a “base fee” covering the cost of relaying the request, complemented by a relaying premium/bribe to be processed ahead of other requests. Additionally, we note that, the strategy adopted by the relays for ordering and processing relay requests is likely to be impacted by the economics of the underlying platform (see e.g. [BCD⁺19, Rou20]), and so, can be adjusted at the relay’s discretion.

334

335 2.2.4 $\widetilde{\text{RelayExec}}$ Contract

336 $\widetilde{\text{RelayExec}}$ is a smart contract, deployed to the blockchain, with knowledge of the ad-
337 dress of the $\widetilde{\text{Mixer}}$ contract to which transactions are to be forwarded. Any participant
338 (relay or potential user) can verify its byte code, and be sure that it cannot be modified
339 by any other party. It executes the Mix call on behalf of the relay, distributing fees as
340 described in the user’s request, thereby allowing user and relay to interact in a trustless
341 way. That is, neither the user or the relay are required to trust the the other party to
342 behave in a certain way – $\widetilde{\text{RelayExec}}$ constrains how a relay request will be handled
343 once both parties have “agreed” to it.

344 Relays create transactions that call the ProcessRequest method on $\widetilde{\text{RelayExec}}$. This
345 method carries out processing of relay requests, and distribution of $vout$, as defined
346 in Fig. 2.3

RelayExec.ProcessRequest(*req*)

```

1 : // Check that the fee is redeemable
2 : if req.mixParams.vout < req.fee then
3 :   abort
4 : endif
5 : // Check the relay permission
6 : dataR ← encode(req.relayAddr) || encode(req.fee) || encode(req.outAddr)
7 : if ¬SigSchOT-SIG.Vf(req.mixParams.otsvk, CRHots(dataR), req.permission) then
8 :   abort
9 : endif
10 : // Cross-contract call to the zeth mixer
11 : mixSuccess ← Mixer.Mix(req.mixParams)
12 : if ¬mixSuccess then
13 :   abort
14 : endif
15 : // Distribute the fee and the withdrawn funds
16 : send(req.relayAddr, req.fee)
17 : funds ← safeSub(req.mixParams.vout, req.fee)
18 : if req.outAddr ≠ 0x0 ∧ funds ≠ 0 then
19 :   send(req.outAddr, funds)
20 : endif
21 : return

```

Figure 2.3: ProcessRequest function, where $\text{send}(\text{addr}, \text{amt})$ refers to the **Ethereum** operation that sends funds amt to some address addr , and where $\text{safeSub}: \mathbb{N}_{\text{ETHWORDLEN}} \times \mathbb{N}_{\text{ETHWORDLEN}} \rightarrow \mathbb{N}_{\text{ETHWORDLEN}}$, such that $\text{safeSub}(x, y) = x - y$ if $x > y$, 0 otherwise.

347 Note that the definition of **ProcessRequest** does not check that the transaction has
348 originated from req.relayAddr . Hence, it is technically possible for a third party to
349 “front-run” this transaction. However any value returned from the transaction is always
350 explicitly distributed to req.relayAddr and req.outAddr , and so any other party signing
351 this would essentially be paying the transaction fee on behalf of $\mathcal{R}_{\mathcal{E}}$, while still allowing
352 $\mathcal{R}_{\mathcal{E}}$ to keep the relay fee.

353 The system would still function if **ProcessRequest** did perform such a check (i.e. that
354 req.relayAddr signed the transaction). However, omitting such a check allows the relay
355 some flexibility when sending the transaction. For example, the relay may wish to pay
356 the transaction fee from another account, or may wish to call RelayExec from another
357 contract of some kind.

358 Finally, we note that users create relay requests such that only **RelayExec** may
359 successfully pass *mixParams* to **Mixer**, however there is no mechanism to ensure that
360 *mixParams* may only be used by a specific *method* of **RelayExec**. Thus, before entering
361 into the protocol, the user should convince himself that **RelayExec** cannot be called in
362 such a way as to violate the protocol (consider the case of a malicious deployer colluding
363 with a relay to provide a second method on **RelayExec** which returns all value to
364 *relayAddr*). For simplicity, we stipulate that **RelayExec** MUST NOT have methods other
365 than **ProcessRequest**.

DRAFT

Chapter 3

Relays with Private Fees

3.1 Introduction

Chapter 2 describes a protocol that allows users to carry out **Zeth** transactions while remaining anonymous (i.e. make **Mix** calls without controlling an **Ethereum** account in ς). Under this scheme, relays receive their fee as part of the public output *vout* withdrawn from the **Zeth** mixer contract. It is clear that, under the assumptions described in Section 1.4.3, fees paid publicly *in real time* in this way can be detected by observers of the **Ethereum** network. Such observers may then learn about the profits made by each relay service, which may not be desirable.

In this chapter we introduce a protocol in which relays can receive fees of hidden denominations. This setting is of particular interest in the context of a “relaying market” in which a set of competing relay services operate with the aim of capturing as much bandwidth¹ (on the “relay network”) as possible, in order to maximize their revenue.

3.2 Protocol overview

As in Chapter 2, we propose a protocol built on top of **Zeth** which allows holders of **Zeth** notes to securely spend their notes without needing to hold **Ether**. We assume that relays are willing to sign and broadcast **Mix** calls, and therefore pay for the gas, in exchange for fee payment in the form of **Zeth** notes. To do so, relays must publish their public **Zeth** address $\mathcal{R}_Z.pub$ and their fee $fee_{\mathcal{R}}$, as well as additional network information such as endpoints which accept relay requests.

Users create parameters *mixParams* to the **Mix** call, such that one of the newly created notes corresponds to payment of $fee_{\mathcal{R}}$ to $\mathcal{R}_Z.pub$. These parameters are then sent to the relay via an established communication channel. Users must create the signature *mixParams.otssig* using the relay’s **Ethereum** address $\mathcal{R}_{\mathcal{E}}.Addr$, to allow $\mathcal{R}_{\mathcal{E}}$ to use *mixParams*. (In some simple scenarios, $\mathcal{R}_{\mathcal{E}}.Addr$ may be made public alongside other relay information. Here, however, we assume that $\mathcal{R}_{\mathcal{E}}.Addr$ is passed securely from the

¹i.e. “market share”

393 relay to the user client as part of the protocol, as this gives the relay more flexibility –
394 see Section 3.3.3 for discussion of how this may enable further relay privacy.)

395 When the relay receives *mixParams* from the user, it checks to ensure that *mixParams*
396 is valid and indeed contains an output note that pays their fee. Relays can then sign and
397 broadcast a transaction directly calling $\text{Mix}(\text{mixParams})$ on the **Mixer** contract. Note
398 that *mixParams.otssig* ensures that the transaction cannot be front-run.

399 3.2.1 Relay-originated mix transactions

400 One potential advantage of relays receiving their fees in the form of **Zeth** notes is that
401 they maintain a level of privacy with respect to their fees. Observers that know the
402 relay’s **Ethereum** address can tell that a given transaction is likely to be on behalf of
403 some user, and therefore that one of the output notes is likely to be addressed to the
404 relay (although they will be unable to see any amounts). However, relays can generate
405 their own **Mix** transactions (which increase privacy by mixing their notes). These **Mix**
406 transactions are indistinguishable from regular relay transactions created on behalf of
407 other users.

408 3.3 Limitations and extensions to the protocol

409 3.3.1 Limitation of output notes

410 The **Mixer** contract is deployed with a hard-coded number of inputs and outputs (de-
411 noted JSIN and JSOUT respectively). In any **Mix** call that is anonymized using the relay
412 system described above, one of the outputs must be used to pay the relay. For the case
413 where $\text{JSOUT} = 2$ (a reasonable default value when relays are not considered), the utility
414 of the system is greatly reduced since users may only set the one remaining output freely.
415 In this case, users are able to combine multiple of their notes into another, but are unable
416 to “split” input notes into multiple output notes. In particular, they are unable to pay
417 a specific amount to another **Zeth** user and receive change.

418 3.3.2 Increasing JSOUT

419 To address the issues of limited output notes, **Zeth** could be instantiated with different
420 parameters (in particular $\text{JSOUT} \leftarrow 3$), in order to support “note-splitting” and relay fee
421 payment in a single transaction. However, such a change to the configuration may have
422 several consequences for the protocol.

423 To support more output notes, each transaction requires more data to be trans-
424 ferred and processed. This increases the storage and processing requirements of the
425 **Mixer** contract (increasing transaction gas costs), and in turn decreases the lifetime of
426 a **Zeth** deployment for a given Merkle tree size (note that the Merkle tree size is de-
427 fined when **Mixer** is deployed). Furthermore, the **Zeth** statement must be made more
428 complex (in order to handle more commitments, and possibly to accommodate a deeper
429 Merkle tree), increasing the cost of generating zero-knowledge proofs.

Note also that if JSOUT is increased, there may be a tendency for each user's funds to become distributed over more notes. If JSIN is not also increased, and the ability of users to recombine Zeth notes is not balanced with this, users may more frequently be required to issue multiple transactions when spending their funds (to "recombine" their funds spread across many notes).

Hence, adjusting JSOUT may have important consequences which should be considered very carefully, especially if the extra output notes are unlikely to be used outside of the relay system.

3.3.3 Support for Ether output

By default, the **Mixer** contract will return any Ether value *vout* to the calling address which, in the protocol described here, would be $\mathcal{R}_\mathcal{E}.Addr$, belonging to the relay. Thus, the protocol as presented does not allow users to withdraw value as Ether while using a relay, unless he is willing to trust the relay to forward the Ether in a later transaction. Our aim is to remove any need for trust within the protocol, and a trustless way to withdraw Ether could be valuable in several scenarios.

As in Chapter 2, users could withdraw Ether to previously unseen Ethereum addresses. Such anonymous addresses could then be used to pay for Zeth transactions, apparently disconnected from any other transactions in the blockchain history. Note that this provides a means for users to anonymously perform Zeth transactions that utilize all JSOUT output notes, without changing the Zeth configuration. Clearly, a user performing two transactions (one to withdraw and one to execute the original Zeth transaction) must pay the relay fee *and* the transaction fee for his subsequent transaction. This may have an impact on the economic model for relay fees.

It is clear that relays will be required to regularly withdraw Zeth notes as Ether, in order to continue to pay transaction fees. They can, of course, simply issue Zeth transactions to withdraw to $\mathcal{R}_\mathcal{E}.Addr$. However, if the relay protocol supports output to Ether, relays could also use this mechanism to withdraw to new Ethereum addresses. To an observer, such transactions would appear to be standard relay transactions on behalf of some user, but would provide the relay with anonymous Ether, further increasing their privacy.

We next identify two approaches to supporting withdrawals to Ether using the protocol given here.

Arbitrary *vout* address in Zeth protocol

The Zeth protocol could be slightly modified so that *mixParams* contains an explicit output address *mixParams.outAddr* to which *vout* should be sent by **Mixer**. This new field **Mixer.outAddr** must be included in the data signed by *mixParams.otssig*, ensuring that it cannot be altered by front-runners or malicious relays. This approach adds a small overhead to the generation of *mixParams*, and to the cost of Mix calls, since this output address must be passed as an extra parameter and used in signature validation.

469 However, supporting this would add versatility to the **Zeth** protocol and may allow other
470 applications to be built on top of it.

471 Note that this new address *outAddr* is distinct from the *sender's address* already in-
472 cluded in *mixParams.otssig*. *mixParams* must ensure that the encapsulating **Ethereum** trans-
473 action originates from $\mathcal{R}_{\mathcal{E}}.Addr$, and that *vout* **Ether** are paid to *mixParams.outAddr*.

474 **Relay via intermediary contract**

475 An alternative approach to support secure withdrawal of **Ether** via relays is to use an
476 intermediary contract, as described in Chapter 2. This change to the relay protocol
477 has the benefit that *vout* can be distributed to one or more parties in a trustless way.
478 However, it does have a potential down-side in terms of privacy – namely that it is
479 trivial for observers to distinguish between transactions issued by relays, and regular
480 transactions issued by users, even if the observer does not know any **Ethereum** addresses
481 owned by relays.

482 **3.3.4 Fees as Ether or Zeth notes**

483 In order to address the problem of limited output notes in **Zeth** (see Section 3.3.1), the
484 protocol could allow the user to choose between 2 fee payment methods: as a **Zeth** note
485 (as described here) or as **Ether** via *mixParams.vout*. In this case, users can use *all*
486 JSOUT outputs from the Mix call for their own purposes, potentially avoiding the need
487 to adjust JSOUT in the **Zeth** configuration, and all the associated problems (as described
488 in Section 3.3.2).

489 A simple way to pay fees as **Ether** is for users to set $vout = fee_{\mathcal{R}}$ when creating
490 *mixParams*. Upon receiving *mixParams*, relays then check for *either* a **Zeth** note *or*
491 *mixParams.vout* that pays their fees. The **Zeth** protocol extension described in Sec-
492 tion 3.3.3 to add *mixParams.outAddr* would then be desirable, to prevent front-runners
493 from claiming the relay fee.

494 An alternative approach would be to use an intermediary contract as described in
495 Chapter 2 (already partially mentioned in Section 3.3.3 to support **Ether** withdrawals).
496 The protocol would then require the extra complexity of a request structure and *proof-*
497 *of-relay-permission*, but would provide maximal flexibility for users. A single Mix call
498 could withdraw **Ether** to a new user address, use all output **Zeth** notes *and* pay the relay
499 fee (in **Ether**).

500 We expect that, given the choice, users would favour fee payment in **Ether** more often
501 than payment in **Zeth** notes, since fee payment in **Ether** allows them to control all JSOUT
502 output notes from the **Zeth** transaction. Further, it seems reasonable to assume that
503 there will always be some relay operators willing to accept relay fees in **Ether**, and thereby
504 users will have some element of choice in how fees are paid. Hence, we should expect
505 some divergence between the relays fees paid in **Ether** and those paid in **Zeth** notes –
506 namely that fees paid in **Zeth** notes will tend to be lower, in order to incentivise users
507 to adopt this protocol.

508 Appendix A

509 Relays with Stake

Note

This section is concerned with a high-level description of a work-in-progress proposal. There are several issues still to be addressed before it can be considered complete. It is given here as a first step towards the goal of addressing possible relay DoS vectors, with the hope that it can be iterated and eventually turned into a practical solution.

510

511 A.1 Introduction

512 In the above proposals (Chapters 2 and 3), relays receive requests and perform “offline”
513 checks to gain a high level of confidence that the transaction (which the relay must sign
514 and therefore pay for) will “succeed” and the relay will receive the designated fee. Under
515 these protocols, a relay is exposed to risk in two forms:

- 516 1. Users are free to make invalid relay requests with no consequences (we assume
517 that they connect via anonymising networks). At the same time, relays are highly
518 incentivised to filter out such invalid requests and avoid signing and broadcasting
519 the corresponding transaction (which would result in the relay paying the invalid
520 transaction’s gas, without receiving their relay fee). In order to detect invalid
521 relay requests, the relay must essentially simulate execution of the full transaction
522 against the current state of the blockchain. Although much less costly than paying
523 the corresponding gas, this may still require significant compute resources. If the
524 relay request is judged to be invalid, the relay will necessarily receive no relay fee
525 in exchange for these verification costs.
- 526 2. As mentioned in Remark 4, there is a chance that a transaction appears valid
527 (i.e. it passes all checks performed by the relay), but later fails due to a conflicting
528 transaction (unseen by the relay) being mined ahead of it. While this risk can be

529 reduced by more thorough checks, the relay can never rule out the possibility that
530 a conflicting transaction exists somewhere on the network.

531 Both of these risks represent possible DoS attack vectors, especially 1, since clients
532 can very easily craft invalid transactions with maximal cost of validation.

533 We consider a potential approach to address these problems. Specifically, we outline
534 a protocol involving collateral staked by users and associated with some specific relay
535 request. This supports a very lightweight check that relays can perform *before* they
536 commit either **Ether** as gas for transactions, or compute resources to fully verify the
537 validity of the request. If this fast upfront check passes, the relay is effectively guaranteed
538 income as a result of relaying the transaction - whether or not the transaction is valid
539 at the time it is mined. In this way, relays can avoid DoS attacks that force them to
540 waste resources for no return.

541 The user's stake is pre-deposited with a contract **RelayStake** in such a way that it is
542 bound to a specific relay request. On receipt of a relay request, relays need only confirm
543 that an associated stake exists before creating the relay transaction and executing it via
544 **RelayStake**. The relay can be sure that **RelayStake** will release the stake to the relay,
545 even if the associated relay transaction fails for any reason. In this way, relays can very
546 quickly gain assurance (up to some negligible probability) that they will not lose any
547 operating costs by proceeding with the request.

548 Since they are exposed to reduced risk, relays should be able to charge users lower
549 fees for their services, while still allowing users to perform **Zeth** operations that are
550 not directly connected to any of their **Ethereum** addresses. However, we note that this
551 reduction in risk for the relay comes with a trade-off for the user, who must deposit
552 **Ether** from some address in order to use the system. Although users do not reveal *which*
553 **Zeth** operation they are performing, they reveal that the owner of the **Ethereum** address
554 paying for the deposit *may* interact with **Zeth** at some future time.

555 While this does not provide as much anonymity as the systems described in Chap-
556 ters 2 and 3, this does represent an improvement in anonymity over the plain **Zeth** pro-
557 tocol. (Note that users interacting directly with **Zeth** can employ strategies to obfuscate
558 their actions, such as broadcasting “dummy” transactions which mix their commitments,
559 however this may incur a relatively high cost and always reveal to observers the exact
560 set of commitments they have created.) The nature of the improvement achieved by
561 the stake system will depend on the details of any final design. We discuss possible
562 trade-offs, and potential strategies to mitigate them, in Appendix A.3.

563 A.2 Protocol overview

564 We describe the components involved in the protocol, and their role in the full workflow.

565 A.2.1 Stake contract

566 We assume the existence of some contract **RelayStake** which performs multiple func-
567 tions:

- Accept and hold stake as collateral against a specific relay request (and in turn specific relay). Users depositing stake should be able to generate a proof π_{req} that **RelayStake** holds some stake for a specific request req . Further, verifiers of π_{req} should not learn which transaction caused the deposit (and thereby the **Ethereum** address of the user who deposited it).
- Act as a relay intermediary, accepting relay requests and proof-of-relay-stake objects. If a valid request and proof are received, the sender has permission to act as a relay for the given request, and the stake is still unspent, **RelayStake** executes the relay request (calls **Mixer.Mix**($req.mixParams$), using the notation of Chapter 2) and releases the stake to the relay, regardless of the outcome of the **Mix** call.

Note that this may be implemented as multiple interacting contracts.

A.2.2 Relay

Relays receive pairs (req, π_{req}) of relays requests and associated proof-of-relay-stake objects. When a relay \mathcal{R} (with **Ethereum** address $\mathcal{R}_{\mathcal{E}}.Addr$) receives such a pair, the operations he performs are relatively straightforward:

Step 1. Check the correctness of req and π_{req} , namely that:

1. req names the relay's **Ethereum** address as the recipient of the relay fee
2. a valid stake exists in **RelayStake**, corresponding to req and π_{req}

The relay should not learn which transaction deposited the stake that corresponds to req and π_{req} (which would reveal an **Ethereum** account of the user).

Step 2. Create a transaction $tx_{\mathcal{R}}$ which calls **RelayStake** with parameters req and π_{req} .

Step 3. Broadcast $tx_{\mathcal{R}}$ to the network and asynchronously wait to receive the relay fee.

As described above, after the initial check of stake corresponding to req , the relay can be sure he will receive his fee even if the **Mix** call fails. Note also that the relay's transaction cannot be front-run, since **RelayStake** will only release the stake to the relay mentioned in req . A user may be able to spend the notes in $req.mixParams$ via another transaction, but he will not be able to prevent the relay from claiming the stake he previously deposited.

For these reasons, the relay need only perform the checks listed above. There is little to be gained by checking any further details of req , including $req.mixParams$ or the state of the **Zeth** mixer contract **Mixer**.

A.2.3 User

A user who wants to make use of a specific relay \mathcal{R} with **Ethereum** address $\mathcal{R}_{\mathcal{E}}.Addr$, using Mix parameters $mixParams$, performs the following actions:

Step 1. Create the appropriate $mixParams$ and corresponding relay request req , bound to the relay's address $\mathcal{R}_{\mathcal{E}}.Addr$.

Step 2. Stake some **Ether** with **RelayStake** against req and generate a corresponding proof-of-relay-stake π_{req} . Note that this collateral is “bound” to req .

Step 3. Send req , π_{req} to the chosen relay (via an anonymous channel), and asynchronously wait for a corresponding transaction to be mined.

Note that the user must have a funded **Ethereum** account $\mathcal{U}_{\mathcal{E}}$ in order to stake collateral, impacting their anonymity to some extent.

A.2.4 Reclaiming stake

Under some circumstances, a user may wish to reclaim his stake after depositing it in **RelayStake**. In fact, if no mechanism were available to accomplish this, the protocol would be vulnerable to withholding attacks. That is, malicious relays could accept relay requests but not relay them within a reasonable time, essentially locking up the user's stake. Eventually, the victims of withholding attacks would be forced to forfeit their stake and find another means to carry out their **Zeth** operations (potentially via another relay). Once the operation has been completed via another transaction, the malicious relay (still holding a request with associated stake) can then claim the user's stake by broadcasting the relay transaction, which will now fail. The victim must pay for his transaction more than once, losing his stake (which, may be greater than the original relay fee - see Appendix A.3). This attack would also cause significant disruption to the relay network and associated market.

Note

Note that the affect of such an attack may be mitigated by some kind of reputation system alongside a relay market fee, so the threat posed by this may not be considered severe. In fact, a reputation system may be a vital part of any relay market.

If the user were allowed to reclaim their stake at any time, the transaction to reclaim it could be unseen by the relay, yet mined ahead of $tx_{\mathcal{R}}$. Thus, the relay would no longer have any guarantee that he would receive the stake in exchange for broadcasting the relay transaction.

It may be possible to facilitate reclaim of the stake by setting a “period of validity” for the stake, if this can be implemented such that:

- the user cannot reclaim the stake until the period of validity elapses,

- the relay can obtain some estimate regarding the period of validity (minimally, if the relay knows some lower bound on the remaining period of validity, it can gain a high degree of certainty that the relay transaction will be mined before the user is able to reclaim the stake),
- for a relay transaction $tx_{\mathcal{R}}$, observers (and relays) should not be able to infer significant information about the transaction that deposited the collateral for $tx_{\mathcal{R}}$,
- for a stake reclaim transaction, observers should not be able to infer significant information about the transaction that deposited the stake being reclaimed,
- the transaction to reclaim a stake cannot be front-run.

A.3 Remarks

In this section we give some remarks about the proposed stake system above. Note that details are highly dependent on the specific cryptographic primitives used. The steps above give an outline of how a staking system may work, with the caveat that they are very likely to be modified in any fully specified protocol.

Stake deposit leaks information. The transaction to deposit the user’s stake reveals an **Ethereum** address of the user, and the value of the stake. Similarly, the transaction to reclaim an unused stake is also associated with the user’s **Ethereum** address. In particular, observers may learn about **Ethereum** addresses of users likely to interact with the relay system. Note that this could potentially be mitigated if users are able, via some other mechanism such as Chapter 2, to anonymously receive funds at new “unused” **Ethereum** addresses. However, in isolation, the stake system always requires an initial transaction from a funded address.

Relays with stake as part of a wider market. Reliance on external mechanisms (such as the ability to withdraw to new **Ethereum** addresses, as mentioned above), may not necessarily be a problem for a stake-based relay protocol. We note that there may be significant benefit to an ecosystem of various relay types (or relays supporting a variety of request types) including those discussed in this document. In particular, we may imagine a market in which users can choose between a range of relay request types - those supporting more privacy and not requiring **Ether** (with higher fees, as a consequence of the extra risk assumed by the relay), and relays with lower fees which place more requirements on users (able to mitigate much of their risk). Concretely, such a system may allow users to withdraw **Ether** to an anonymous address via an expensive relay, and then use this **Ether** as collateral for further relay transactions with lower relay fees. Obviously, in such a scenario, users (and their wallet software) must take a great deal of care not to reveal relationships between transactions.

668 **Collateral value vs Relay Fee.** The above outline suggests a very simple scenario
669 in which the stake is unconditionally used as the relay fee (that is, the stake must
670 have the same value as the relay fee, which cannot be negotiated after the stake has
671 been deposited. We note that any concrete instantiation would require somewhat more
672 flexibility, and would likely extend fee payment mechanisms in one of several ways.

- 673 • The user could be required to stake some fixed value from which fees are paid,
674 with the “change” being paid either to the user, or as *vin* to the **Zeth** mixer. The
675 fixed stake value would then define an upper bound on relay fees. This allows a
676 lot more scope for relays to dynamically change their fees.
- 677 • We may prefer to distinguish between successful and failed **Mix** calls. On success,
678 **RelayStake** may refund the user with the difference between the stake and the
679 relay fee (as described above), while users could be punished for invalid relay
680 requests by forfeiting their entire stake to the relay.
- 681 • Similarly, it may be preferable for relay fees to be paid directly from the **Mix** call in
682 the case of successful transactions. That is, the **Mix** call parameters *req.mixParams*
683 must include an output paying the relay (either *vout* as in Chapter 2, or a **Zeth** note
684 as in Chapter 3), and the stake is refunded to the user (either to an address of their
685 choosing, or as *vin* to the **Mix** call. This allows relays to maintain privacy with
686 respect to their fee payments, while minimizing risk through the stake system.
- 687 • It may be possible to allow users to deposit a single stake, bound to a specific
688 relay, which can then be used for multiple relay requests. The relay could still
689 claim the stake by presenting an invalid request from the user, but the user could
690 make multiple relay requests (possibly within some time limit - see the following
691 paragraph) without the need to redeposit. User anonymity would be improved
692 because, while observers would still learn the **Ethereum** address depositing each
693 stake, they would not be able to determine how many relay requests were carried
694 out on behalf of each depositor.

695 **Period of validity.** The period of validity gives observers information about when
696 the users relay transaction will be carried out. Depending on the frequency of relay
697 transactions interacting with **RelayStake**, users should ensure that the period of validity
698 is sufficiently long, to avoid compromising their anonymity. Similarly, a predictable
699 interval between a stake being deposited and the corresponding relay transaction that
700 claims it, would also reveal information about the user originating each relay transaction.
701 To avoid this, the period of validity should be sufficiently long to allow some “noise”
702 in the interval between stake deposits and relay transactions. Note that this presents a
703 trade-off, since the user’s funds are potentially “locked up” for a longer period.

704 Appendix B

705 Network structure

706 B.1 Binding requests to relays

707 B.1.1 Background

708 The protocols presented in this document require users to create relay requests that
709 can only be successfully processed by a specific relay. This serves as a mechanism to
710 prevent other network participants from “stealing” the relay requests or front-running
711 relay transactions. The alternative to this would be to support “free” relay requests,
712 not bound to specific relays, which could therefore be processed by any participant. If
713 these “free” requests are made available (or “broadcast”) to multiple relays, those relays
714 must then “race” to process the request and broadcast a corresponding relay transaction.
715 When the first relay transaction is accepted by the blockchain, the “winning” relay will
716 receive the relay fee and later transactions from other relays will be rendered invalid (as
717 a consequence of the nullifiers declared in the Zeth relay request being marked as used).

718 While such an approach is entirely feasible, it increases the risk for relays, making
719 it much harder for them to hedge against lost fees and wasted compute resources. As
720 a consequence, it becomes much more difficult for relays to assess the risk associated
721 with a given request, which in turn is likely to result in an increase in relay fees. All
722 resources used by “losers” of the race are wasted. In contrast, in the case where requests
723 are bound to specific relays, these resources can be used to process multiple requests in
724 parallel, increasing the efficiency of the system.

725 B.1.2 Emulating “free” relay requests

726 Despite the mechanisms to prevent front-running, the protocols presented in this docu-
727 ment could be leveraged by some user (say \mathcal{U}) to force relays to “compete” for relay re-
728 quests. Specifically, in order to call the Zeth mixer $\widetilde{\text{Mixer}}$ with parameters mixParams ,
729 \mathcal{U} can run multiple instances of a protocol in parallel, generating N relay requests for
730 mixParams , each targeting a different relay. If the user then sends each of these N
731 requests to the targeted relay, the desired state transition will be carried out by the

first relay transaction to be accepted into the blockchain, rendering the remaining $N - 1$ requests invalid (by the nullifier mechanism cited previously).

As in the case of “free” requests, relays are exposed to extra risk for the reasons given above. However, this “emulating” approach does provide partial mitigation of this risk, due to the extra cost that the user \mathcal{U} must incur. Under the protocols given in this document, in order to create N requests targeting different relays, the user must generate N zk-SNARKs, which is computationally demanding and therefore represents a cost to the user. Therefore, request generation may act as a user-side proof-of-work, preventing malicious messages from flooding the network (as originally designed for by Dwork et al. [DN92, JJ99]). This naturally leads to the following process by which relays can partially protect against some DoS vectors, by performing the following checks on relay requests:

1. Verify that none of the nullifiers in the request has been seen in previously received requests (inspect the mempool and the blockchain state). If one or more nullifiers is a duplicate then reject the request, else proceed.
2. Verify the zk-SNARK proof in the request. If the proof is invalid, reject the request. Otherwise the request can be considered for processing.

In this way, the cost of generating N proofs imposes some upper bound on the message output rate of a potential attacker.

B.2 Unicast vs broadcast networks

The discussions in this document assume only that some transport mechanism exists for users to send relay requests to specific relays. As noted in Section 1.4, users can achieve further anonymity if this transport mechanism does not require the user to reveal any identifying information at the network level. We now discuss some specific implementations of the transport layer (namely “unicast” vs “broadcast” networks), and their respective properties.

By design, relay requests are bound to specific relays, which intuitively implies a “unicast” style transport mechanism. That is, relays publish a network address of some form, and users send requests to this address. Observers of the physical network may determine that a message has been sent from the user to the relay, but the content of the message (i.e. the details of the relay request) is not visible to other participants. This is a natural choice given that relay request data cannot be used by parties other than the targeted relay. While not a requirement of any of these protocols, unicast channels (in particular point-to-point communication channels, which we assume to be encrypted by default) only reveal the relay request content to the relay itself. Adversaries able to gain control a physical network node along the route between user and relay (in general a limited set of nodes, which varies depending on user and relay) may learn that a message was sent from the user to the relay, but they will not learn anything about the message content. Such communication channels also allow for interaction between relay and user

771 (for example, the relay could dynamically select an **Ethereum** or **Zeth** address to receive
772 payment, or privately negotiate fees with the user).

773 Clearly it would be entirely possible to implement these protocols using a “broadcast”
774 system, such as those employed by blockchains to propagate transactions and blocks.
775 Requests could be broadcast unencrypted without impacting the reliability of the system,
776 as long as messages were eventually seen by the target relay. Broadcast networks provide
777 some inherent receiver anonymity, in the sense that it is more difficult to identify which
778 *network node* is the recipient of a given message, however in this setting, all participants
779 in the system would be able to see the content of relay requests and potentially determine
780 the number of requests received by each relay identity (and, in turn, infer information
781 about their profit). Despite relay requests being visible to other participants, the protocol
782 would still prevent other relays from profiting from these requests, since they are bound
783 to the target relay. Further, the content of requests could be hidden by encryption so
784 that only the intended relay may read them. Instead of publishing a network “address”
785 of some form, relays could publish an encryption key, with which users must encrypt
786 requests before broadcasting them (although care must be taken to use a key-private
787 encryption scheme to avoid leaking information about the recipient).

788 While broadcast networks could theoretically be used in these relay protocols, uni-
789 cast networks are more bandwidth efficient (i.e. a given message needs only to find a
790 path through the network in order to flow from the sender to the recipient). In con-
791 trast, broadcast networks may provide a level of sender anonymity in the face of net-
792 work observers, although even in broadcast networks methods exist to infer the message
793 originator (e.g. nodes with high degree¹ – also referred to as “supernodes” – can be
794 used to infer the sender of a message on a broadcast channel by using timing informa-
795 tion [KKM14]). Broadcast communication channels are also of great interest to achieve
796 “recipient anonymity” (see [PW87] for more details on “anonymity”).

797 At first sight, the use of a unicast transport may appear to increase the centralization
798 of the system. However, this is demonstrably not the case for the protocols discussed
799 here, which can (as described above) be implemented using a transparent broadcast
800 network and do not inherently rely on any centralization.

801 Finally, we note that, although broadcast networks could theoretically be used, we
802 suggest that unicast networks are likely to be more suitable, given their lower bandwidth
803 and complexity.

804 B.3 Network anonymity

805 Relay protocol designers may choose to transmit requests via the method that best fits
806 their needs, taking into consideration the tradeoffs mentioned in Appendix B.2 above. As
807 well as overhead, network topology also has a strong influence on anonymity [DMT10]).
808 However, in order to achieve strong privacy guarantees, anonymisation techniques (e.g. cover
809 traffic, message padding etc.) must be used, to minimize communication leakages.

¹In the graph theoretical sense.

810 While protocols like Dandelion [VFV17, FVB⁺18] were initially introduced to im-
 811 prove *diffusion* mechanisms and improve network anonymity on Bitcoin, they could also
 812 be of interest in the context of relay request broadcasts (as alluded to in Appendix B.2).
 813 However, other techniques (providing different properties) may also be of interest for
 814 relay network. Some of these are given below:

- 815 • DC-nets [Cha88] provide strong guarantees with respect to the sender anonymity
 816 (but generally incur a big overhead and require large amounts of randomness).
- 817 • Crowds [RR98] follow a “blending into a crowd” approach (i.e. hiding one’s actions
 818 among the actions of many others), in which a user’s request is randomly circulated
 819 in a “crowd” (set of users) before being submitted – by a random member of the
 820 crowd – and sent to the destination². Note that such approaches generally do not
 821 provide strong guarantees with respect to recipient anonymity³.
- 822 • Mix networks (or *mixnets*) [Cha03], in which nodes (“mix nodes”) are routers that
 823 perform cryptographic operations (providing bit-wise unlinkability), and modify
 824 the order in which output messages are emitted. This hides any correspondence
 825 between input and output messages.
- 826 • Onion routing [GRS99] (which also underlies “garlic routing” [Din00]) consists of
 827 multiple layers of encryption (one per “hop” on the network). Requests are sent
 828 through a chosen set of routers (forming a “circuit”) in order to obfuscate the link
 829 between sender and recipient, as seen by non-global adversaries (i.e. those that do
 830 not control all nodes on the circuit⁴). This generally achieves low-latency relative
 831 to other approaches.

832 Importantly, modern protocols building on these techniques use additional mecha-
 833 nisms for enhanced robustness (e.g. “cover traffic” to prevent timing attacks etc.).

834 **Remark 5.** *We note that accountable anonymous communication networks [DP07] are*
 835 *also of great interest in the context of transaction relay protocols as a way to further*
 836 *prevent DoS attacks.*

²Crowd members *cannot* identify the initiator of the request. The initiator is indistinguishable from a member that forwards a request from another user.

³While relay anonymity is not our principal focus, it is worth keeping in mind the impact of side channel leakages which can be used to infer information about the sender. For instance, a powerful adversary – monitoring a big part of the Internet – may notice a client access the relay’s public information (such as the relay’s website) followed by a message to the relay from a crowd to which the client belongs. The adversary may then infer that the client was the relay user. Hence, additional care needs to be allocated to the relay discovery mechanism itself, and the right trade-offs must be made depending on the application and associated threat model.

⁴In some cases, controlling the “entry” and “exit” nodes (i.e. first and last nodes of the chain/circuit) is sufficient to carry out so-called “correlation attacks”. See <https://github.com/Attacks-on-Tor/Attacks-on-Tor> for a list of attacks on Tor [DMS04]

Bibliography

- [BCD⁺19] Vitalik Buterin, Eric Conner, Rick Dudley, Matthew Slipper, Ian Norden, and Abdelhamid Bakhta. Fee market change for eth 1.0 chain. <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-1559.md>, 2019.
- [Bra97] S. Bradner. Key words for use in rfcs to indicate requirement levels. RFC 2119, RFC Editor, March 1997.
- [Cha88] David Chaum. The dining cryptographers problem: Unconditional sender and recipient untraceability. *J. Cryptol.*, 1(1):65–75, 1988.
- [Cha03] David Chaum. Untraceable electronic mail, return addresses and digital pseudonyms. In Dimitris Gritzalis, editor, *Secure Electronic Voting*, volume 7 of *Advances in Information Security*, pages 211–219. Springer, 2003.
- [Cle20] Clearmatics. Zeth Protocol Specification, 2020.
- [DGK⁺19] Philip Daian, Steven Goldfeder, Tyler Kell, Yunqi Li, Xueyuan Zhao, Iddo Bentov, Lorenz Breidenbach, and Ari Juels. Flash boys 2.0: Frontrunning, transaction reordering, and consensus instability in decentralized exchanges, 2019.
- [Din00] Roger Dingledine. The Free Haven Project: Design and Deployment of an Anonymous Secure Data Haven. MIT Master’s Thesis. <https://www.freehaven.net/papers.html>, 06 2000.
- [DMS04] Roger Dingledine, Nick Mathewson, and Paul F. Syverson. Tor: The second-generation onion router. In Matt Blaze, editor, *Proceedings of the 13th USENIX Security Symposium, August 9-13, 2004, San Diego, CA, USA*, pages 303–320. USENIX, 2004.
- [DMT10] Claudia Díaz, Steven J. Murdoch, and Carmela Troncoso. Impact of network topology on anonymity and overhead in low-latency anonymity networks. In Mikhail J. Atallah and Nicholas J. Hopper, editors, *Privacy Enhancing Technologies, 10th International Symposium, PETS 2010, Berlin, Germany, July 21-23, 2010. Proceedings*, volume 6205 of *Lecture Notes in Computer Science*, pages 184–201. Springer, 2010.

- [DN92] Cynthia Dwork and Moni Naor. Pricing via processing or combatting junk mail. In Ernest F. Brickell, editor, *Advances in Cryptology - CRYPTO '92, 12th Annual International Cryptology Conference, Santa Barbara, California, USA, August 16-20, 1992, Proceedings*, volume 740 of *Lecture Notes in Computer Science*, pages 139–147. Springer, 1992.
- [DP07] Claudia Díaz and Bart Preneel. Accountable anonymous communication. In Milan Petkovic and Willem Jonker, editors, *Security, Privacy, and Trust in Modern Data Management, Data-Centric Systems and Applications*, pages 239–253. Springer, 2007.
- [FVB⁺18] Giulia C. Fanti, Shaileshh Bojja Venkatakrishnan, Surya Bakshi, Bradley Denby, Shruti Bhargava, Andrew Miller, and Pramod Viswanath. Dandelion++: Lightweight cryptocurrency networking with formal anonymity guarantees. *Proc. ACM Meas. Anal. Comput. Syst.*, 2(2):29:1–29:35, 2018.
- [GRS99] David M. Goldschlag, Michael G. Reed, and Paul F. Syverson. Onion routing. *Commun. ACM*, 42(2):39–41, 1999.
- [JJ99] M. Jakobsson and A. Juels. Proofs of work and bread pudding protocols. In *Communications and Multimedia Security*, 1999.
- [KKM14] Philip Koshy, Diana Koshy, and Patrick D. McDaniel. An analysis of anonymity in bitcoin using P2P network traffic. In Nicolas Christin and Reihaneh Safavi-Naini, editors, *Financial Cryptography and Data Security - 18th International Conference, FC 2014, Christ Church, Barbados, March 3-7, 2014, Revised Selected Papers*, volume 8437 of *Lecture Notes in Computer Science*, pages 469–485. Springer, 2014.
- [lsa20] lsankar4033. Surrogeth: Tricking frontrunners into being transaction relayers. <https://ethresear.ch/t/surrogeth-tricking-frontrunners-into-being-transaction-relayers/6937>, 2020.
- [PHE⁺17] Ania M. Piotrowska, Jamie Hayes, Tariq Elahi, Sebastian Meiser, and George Danezis. The loopix anonymity system. In Engin Kirda and Thomas Ristenpart, editors, *26th USENIX Security Symposium, USENIX Security 2017, Vancouver, BC, Canada, August 16-18, 2017*, pages 1199–1216. USENIX Association, 2017.
- [PW87] Andreas Pfitzmann and Michael Waidner. Networks without user observability. *Comput. Secur.*, 6(2):158–166, 1987.
- [RK20] Dan Robinson and Georgios Konstantopoulos. Ethereum is a dark forest. <https://medium.com/@danrobinson/ethereum-is-a-dark-forest-ecc5f0505dff>, 2020.

- 903 [Ron20] Antoine Rondelet. Zecale: Reconciling privacy and scalability on ethereum.
904 *CoRR*, abs/2008.05958, 2020.
- 905 [Rou20] Tim Roughgarden. Transaction Fee Mechanism Design for the Ethereum
906 Blockchain: An Economic Analysis of EIP-1559. <https://timroughgarden.org/papers/eip1559.pdf>, 2020.
- 907
- 908 [RR98] Michael K. Reiter and Aviel D. Rubin. Crowds: Anonymity for web trans-
909 actions. *ACM Trans. Inf. Syst. Secur.*, 1(1):66–92, 1998.
- 910 [RZ19] Antoine Rondelet and Michal Zajac. ZETH: On Integrating Zerocash on
911 Ethereum. [Online; released April-2019], 2019.
- 912 [VFV17] Shaileshh Bojja Venkatakrisnan, Giulia C. Fanti, and Pramod Viswanath.
913 Dandelion: Redesigning the bitcoin network for anonymity. *Proc. ACM*
914 *Meas. Anal. Comput. Syst.*, 1(1):22:1–22:34, 2017.
- 915 [Woo19] Dr Gavin Wood. ETHEREUM: A Secure Decentralised Generalised Trans-
916 action Ledger Byzantium. <https://ethereum.github.io/yellowpaper/paper.pdf>, 2019. [VERSION 7e819ec - 2019-10-20].
917