

1

# Zeth Protocol Specification

2

Clearmatics Cryptography R&D

3

November 23, 2020

## Abstract

5 This document specifies the Zeth protocol with various security fixes and performance  
6 improvements from the initial design [RZ19].

7 **Keywords**— Ethereum, Zerocash, Zcash, financial-privacy, zero-knowledge proofs,  
8 Zeth, privacy-preserving state transitions

# 9 Contents

10	<b>1 Preliminaries</b>	<b>11</b>
11	1.1 Data structures and representation . . . . .	11
12	1.1.1 Structured data . . . . .	11
13	1.1.2 Representations . . . . .	14
14	1.2 Ethereum . . . . .	15
15	1.2.1 Ethereum account . . . . .	16
16	1.2.2 Ethereum transaction . . . . .	17
17	1.2.3 Ethereum events and Bloom filters . . . . .	20
18	1.3 zk-SNARKs . . . . .	21
19	1.3.1 Preliminary definitions . . . . .	21
20	1.3.2 Computation representation – arithmetization . . . . .	23
21	1.4 Decentralized Anonymous Payment schemes (DAP) . . . . .	24
22	1.5 Definitions . . . . .	26
23	1.5.1 Negligible function . . . . .	26
24	1.5.2 Basic algebra notions . . . . .	27
25	1.5.3 Security assumptions . . . . .	27
26	1.5.4 Symmetric encryption . . . . .	28
27	1.5.5 Asymmetric encryption . . . . .	29
28	1.5.6 Block cipher-based compression functions . . . . .	30
29	1.5.7 Hash functions . . . . .	31
30	1.5.8 Pseudo Random Functions . . . . .	33
31	1.5.9 Commitment scheme . . . . .	33
32	1.5.10 Digital Signature . . . . .	34
33	1.5.11 Message Authentication Code . . . . .	35
34	<b>2 Zeth protocol</b>	<b>37</b>
35	2.1 Zeth Data Types . . . . .	37
36	2.2 Zeth statement . . . . .	40
37	2.3 Generating the inputs of the Mix function ( $\text{Mix}_{in}$ ) . . . . .	41
38	2.4 Creating an Ethereum transaction $tx_{\text{Mix}}$ to call <b>Mixer</b> . . . . .	43
39	2.5 Processing $tx_{\text{Mix}}$ . . . . .	43
40	2.6 Receiving <i>ZethNotes</i> . . . . .	45
41	2.7 Security requirements for the primitives . . . . .	46

42	2.7.1	Additional notes . . . . .	46
43	<b>3</b>	<b>Instantiation of the cryptographic primitives</b>	<b>48</b>
44	3.1	Instantiating the PRFs, ComSch and CRHs . . . . .	49
45	3.1.1	Blake2 primitive . . . . .	49
46	3.1.2	Commitment scheme . . . . .	50
47	3.1.3	PRFs . . . . .	50
48	3.1.4	Collision resistant hashes . . . . .	52
49	3.2	Instantiating MKHASH . . . . .	52
50	3.2.1	MIMC Encryption . . . . .	52
51	3.2.2	MIMC-based compression function . . . . .	53
52	3.2.3	An efficient instantiation of MIMC primitives . . . . .	54
53	3.2.4	Security requirements satisfaction . . . . .	55
54	3.3	Zeth statement after primitive instantiation . . . . .	58
55	3.3.1	Instantiating the packing functions . . . . .	59
56	3.4	Instantiate SigSch <sub>OT-SIG</sub> . . . . .	63
57	3.4.1	Security requirements satisfaction . . . . .	64
58	3.4.2	Data types . . . . .	65
59	3.5	Instantiate EncSch . . . . .	65
60	3.5.1	DHAES encryption scheme . . . . .	65
61	3.5.2	A DHAES instance . . . . .	66
62	3.5.3	EncSch instantiation . . . . .	69
63	3.5.4	Security requirements satisfaction . . . . .	72
64	3.5.5	Final notes and observations . . . . .	73
65	3.6	ZkSnarkSch instantiation . . . . .	76
66	<b>4</b>	<b>Implementation considerations and optimizations</b>	<b>79</b>
67	4.1	Client security considerations . . . . .	79
68	4.1.1	Syncing and waiting . . . . .	80
69	4.1.2	Note management . . . . .	81
70	4.1.3	Prepare arguments for Mix transaction . . . . .	82
71	4.1.4	Wallet backup and recovery . . . . .	82
72	4.2	Contract security considerations . . . . .	83
73	4.3	Efficiency and scalability . . . . .	83
74	4.3.1	Importance of performance . . . . .	83
75	4.3.2	Cost centers . . . . .	84
76	4.3.3	Client performance . . . . .	84
77	4.3.4	Contract performance . . . . .	85
78	4.3.5	Optimizing Blake2's circuit. . . . .	86
79	4.4	Encryption of the notes . . . . .	93

80	<b>A Transaction non malleability</b>	<b>95</b>
81	A.1 Transaction malleability attack on Zeth . . . . .	96
82	A.1.1 The attack . . . . .	97
83	A.2 Solutions to address the transaction malleability attack . . . . .	98
84	A.2.1 ZeroCash solution . . . . .	98
85	A.2.2 Zcash’s solution . . . . .	99
86	A.2.3 Solution on Ethereum . . . . .	99
87	<b>B Double spend attack on equivalent class</b>	<b>101</b>
88	<b>C Side-channel attacks and information leaks</b>	<b>102</b>
89	C.1 Counterfeit data . . . . .	102
90	C.2 Data leaked during synchronization . . . . .	103
91	C.3 Queries on successful decryption . . . . .	104
92	C.4 Invalid ciphertext . . . . .	104
93	C.5 Using (and retrieving) nullifiers . . . . .	105
94	C.6 Proof generation . . . . .	106
95	C.7 Simple mixer calls . . . . .	106
96	C.7.1 Small anonymity sets . . . . .	107
97	<b>D Security proofs of Blake2</b>	<b>109</b>
98	D.1 Security model of Blake2 . . . . .	109
99	D.1.1 Weakly Ideal Cipher Model . . . . .	109
100	D.2 Security proofs . . . . .	111
101	D.2.1 Blake2 is a PRF . . . . .	111
102	D.2.2 Proof of Blake2 collision resistance . . . . .	112
103	D.2.3 Blake2 as a commitment scheme . . . . .	113
104	D.2.4 Proof of commitment scheme security . . . . .	114
105		

# 106 Notation

## 107 Basic mathematical notation

108	$\emptyset$	The empty set, i.e. $\emptyset = \{\}$
109	$\#S$	The number of elements in the finite set $S$ (also referred to as “cardinality of the set $S$ ”). By convention, $\#\emptyset = 0$
111	$x \in S$	Represents that $x$ is an element of $S$ . If $x$ is a variable such that $x \in S$ , we will
112		say that “ $x$ has type $S$ ”, i.e. the unordered collection of objects $S$ represents all
113		the values that $x$ can take
114	$S \setminus T$	Set difference of sets $S$ and $T$ , i.e. $S \setminus T = \{x \in S : x \notin T\}$ (voiced “the set of
115		elements $x$ in $S$ such that $x$ is not in $T$ ”)
116	$S \subseteq T$	$S$ is a subset of $T$ , i.e. $x \in S \Rightarrow x \in T$
117	$S \subset T$	$S$ is a <i>proper</i> (or “strict”) subset of $T$ , i.e. $x \in S \Rightarrow x \in T \wedge \exists y \in T, y \notin S$
118	$S = T$	$S \subseteq T \wedge T \subseteq S$
119	$S \cup T$	Union of set $S$ and set $T$ , i.e. $\{x : x \in S \vee x \in T\}$
120	$S \cap T$	Intersection of set $S$ with set $T$ , i.e. $\{x : x \in S \wedge x \in T\}$
121	$f: S \rightarrow T$	Function $f$ that maps elements of the non-empty set $S$ , the “domain”, to the
122		non-empty set $T$ , the “codomain”
123	$\mathbb{N}$	Set of natural numbers. $\mathbb{N}^+$ represents $\mathbb{N} \setminus \{0\} = \{1, 2, \dots\}$ , where $\{n, \dots\}$ rep-
124		resents the application of the successor operator $\text{Succ}(n) = n + 1$ , defined by the
125		Peano axioms, infinitely many times
126	$\mathbb{Z}$	Set of integers, i.e. $\{\dots, -2, -1, 0, 1, 2, \dots\}$ , where $\{\dots, n\}$ represents the appli-
127		cation of the predecessor operator $\text{Pred}(n) = n - 1$ infinitely many times
128	$\mathbb{Q}, \mathbb{R}$	Set of rational, real numbers
129	$[n]$	Set $\{0, \dots, n - 1\}$ , where $n \in \mathbb{N}$
130	$\{a, \dots, b\}$	Set of integers from $a$ through $b$ inclusive, where $a \leq b$

131	$(a_0, a_1, \dots, a_{n-1})$	$n$ -tuple, i.e. ordered collection of items of length $n$ . If $n = 1$ , we call
132		it a “singleton”, if $n = 2$ , we call the tuple a “pair”. Finally, if $n = 3$ , we call it
133		a “triple”. We use the terms “tuples” and “lists” interchangeably.
134	$S \times T$	Cartesian product of sets $S$ and $T$ , i.e. set of all ordered pairs $\{(x, y) : x \in S \wedge y \in T\}$
135	$S^n$	$n$ -fold Cartesian product of $S$ with itself, i.e. $S^n = \{(x_0, \dots, x_{n-1}) : x_i \in S \forall i \in [n]\}$ , where $n \in \mathbb{N}$
136		
137	$\Lambda$	General notation for an alphabet, i.e. a <i>non-empty finite set</i> such that every string
138		(ordered collection of symbols, or letters, all in $\Lambda$ ) has a unique decomposition.
139		The number of symbols in a string is denoted the “length” of the string
140	$\varepsilon$	The empty string. $\varepsilon$ is a string over any alphabet.
141	$\Lambda^n$	Set of all strings, defined over alphabet $\Lambda$ , containing $n$ symbols (i.e. “of length
142		$n$ ”)
143	$\Lambda^*$	The Kleene star of $\Lambda$ represents the set of all strings of finite length, defined over
144		alphabet $\Lambda$ , including the empty string $\varepsilon$ , i.e. $\Lambda^* = \bigcup_{n \in \mathbb{N}} \Lambda^n$
145	$\text{length}(x)$	$\text{length} : \Lambda^* \rightarrow \mathbb{N}$ computes the length of a string $x$ defined over $\Lambda$ , i.e. $\text{length}(x)$
146		returns the number of symbols composing the string $x$ . By convention, $\text{length}(\varepsilon) =$
147		0
148	$x \parallel y$	Infix notation for the concatenation function, $\parallel : \Lambda^* \times \Lambda^* \rightarrow \Lambda^*$ . If $\text{length}(x) =$
149		$n, \text{length}(y) = m$ and $(n, m) \in \mathbb{N}^2$ , then for $z = x \parallel y$ holds $\text{length}(z) = n + m$
150	$\text{trunc}_x(k)$	$\text{trunc} : \Lambda^* \rightarrow \Lambda^k$ is the truncation function that returns the sequence formed
151		from the first $k$ elements of $x$ , where $x \in \Lambda^*$ . If $k > \text{length}(x)$ , then $\text{trunc}_x(k) = x$
152	$x[a:b]$	$[\cdot] : \Lambda^n \times \mathbb{N} \times \mathbb{N} \rightarrow \Lambda^{\leq b-a}$ is the slice function that, if $b \geq a$ , returns the string
153		starting at index $\min(n, a)$ of $x$ and finishing at index $\min(n, b)$ . The function
154		additionally interprets $x[:b]$ as $x[0:b]$ and $x[a:]$ as $x[a:n]$
155	$\text{pad}_n(x)$	$\text{pad} : \Lambda^{\leq n} \rightarrow \Lambda^n$ is the padding function which pads $x$ by 0’s to reach a size of
156		$n$ . The padding depends on the variable type and endianness.
157	$\text{append}(l, x)$	$\text{append} : D^n \times D \rightarrow D^{n+1}$ is the algorithm that appends $x$ to the list of $n$
158		element(s) $l$ , if all $x$ and $l$ share the same data type $D$
159	$\mathbb{B}$	Alphabet of binary symbols, i.e. $\{0, 1\}$
160	$\langle \mathbf{g}_1, \dots, \mathbf{g}_l \rangle$	Cyclic group generated by $\{\mathbf{g}_1, \dots, \mathbf{g}_l\}$
161	$(q, \mathbb{G}, \mathbf{g}, \otimes)$	Description of the cyclic group $\mathbb{G} = \langle \mathbf{g} \rangle$ of order $q$ , with operation $\otimes$

162	$\mathbb{Z}/r\mathbb{Z}$	Quotient group defined as the set of equivalence classes modulo $r$ . $\mathbb{Z}/r\mathbb{Z}$ , also written $\mathbb{Z}_r$ , is an additive group. If $r = p$ a prime number, then $\mathbb{Z}_p = \{0, \dots, p-1\} = \mathbb{Z}/p\mathbb{Z}$ is a finite field of elements modulo prime $p$ , also denoted $\mathbb{F}_p$ , where $0_{\mathbb{F}_p}$ and $1_{\mathbb{F}_p}$ respectively represent the additive and multiplicative identity
166	$\mathbb{F}_q$	Finite field of cardinality $q = p^m$ , where $p$ is prime, and $m \in \mathbb{N}$
167	$\llbracket x \rrbracket$	Represents the encoding of the scalar $x$ in a group $\mathbb{G}$ described as $(p, \mathbb{G}, \langle \mathbf{g} \rangle, \otimes)$ , i.e. $\llbracket x \rrbracket = x \cdot \llbracket 1 \rrbracket = \mathbf{g} \otimes \dots \otimes \mathbf{g}$ ( $x$ times). Thus, by convention, $\llbracket 1 \rrbracket = \mathbf{g}$
169	$\bullet$	Represents an inline operator for bilinear pairing. That is for a bilinear pairing from $\mathbb{G}_1 \times \mathbb{G}_2$ to $\mathbb{G}_T$ and elements $\llbracket a \rrbracket_1, \llbracket b \rrbracket_2$ we write $\llbracket ab \rrbracket_t = \llbracket a \rrbracket_1 \bullet \llbracket b \rrbracket_2$
171	$\lceil x \rceil$	Round $x \in \mathbb{R}$ to the next integer
172	$\lfloor x \rfloor$	Round $x \in \mathbb{R}$ to the previous integer
173	$\log_b(x)$	Logarithm with respect to base $b$ , i.e. $x = b^y, \log_b(x) = y$
174	<b>Algorithmic notation</b>	
175	$x \leftarrow \$ \mathcal{X}$	Element chosen uniformly at random from set $\mathcal{X}$
176	$x \leftarrow y$	The value $y$ is assigned to the variable $x$ (i.e. “ $x$ receives the value $y$ ”)
177	PPT	Probabilistic polynomial time. A polynomial time algorithm $A$ is one for which there exists a polynomial $f$ such that the running time of $A$ on input $x \in \{0, 1\}^*$ is $f( x )$ . A probabilistic algorithm has the ability to “flip” random coins and use the result of these coin tosses in its computation
181	NUPPT	Non-uniform probabilistic polynomial time
182	$\mathcal{O}(f)$	Big-O notation
183	il, kl, nl, rl, ol	The input il, key kl, nonce nl, randomness rl and output ol length
184	<b>Cryptography notation</b>	
185	$\mathcal{O}^X(n)$	Public oracle for algorithm $X$ which can be accessed at most $n$ times; $\mathcal{O}^X$ is an unrestricted oracle for algorithm $X$
187	$\lambda$	Security parameter ( $\lambda \in \mathbb{N}$ )
188	negl	Negligible function. In this document, negligible will usually mean $\mathcal{O}(2^{-\lambda})$
189	poly	Polynomial function
190	$\mathcal{A}$	Adversary algorithm



191	$\text{Adv}_{F,\mathcal{A}}^{\text{prop}}(\lambda)$	Advantage of the adversary $\mathcal{A}$ with regard to the attack game <b>prop</b> on $F$	
192		(e.g. $F$ can be a function, a family of functions or a group on which a given	
193		property represented by the game <b>prop</b> is supposed to hold)	
194	$\text{prop}^{\mathcal{A}}$	Adversary $\mathcal{A}$ running a security game <b>prop</b>	
195	<b>Zeth notation</b>		
196	$\pi$	Output of the proving algorithm of a zk-SNARK scheme. $\pi$ is also informally	
197		referred to as a “zk-SNARK proof”, “zk-proof”, or simply “proof”	
198	$\mathcal{P}_{\mathcal{Z}}$	Standard notation for a <b>Zeth</b> user	
199	$\widetilde{\text{Mixer}}$	The mixer smart-contract instance	
200	<b>EncSch</b>	In-band encryption scheme used to share <b>Zeth</b> notes	
201	<b>Ethereum notation</b>		
202	<b>Account</b>	Standard notation for an <b>Ethereum</b> account object	
203	$\widetilde{\text{Cntrct}}$	Standard notation for an <b>Ethereum</b> smart-contract instance	
204	$\mathcal{P}_{\mathcal{E}}$	Standard notation for an <b>Ethereum</b> user	
205	$\varsigma$	Mapping representing the <b>Ethereum</b> state (i.e. “World state”)	
206	$\varsigma[a]$	Account object stored at address $a$ in $\varsigma$ if it exists, $\perp$ is returned otherwise	
207	<b>Constants</b>		
208	<b>ADDRLEN</b>	The bit-length of an <b>Ethereum</b> address	160 <i>bits</i>
209	<b>BLAKE2sCLEN</b>	Output size of Blake2s compression function [ANWOW13]	256 <i>bits</i>
210	<b>FIELD<sub>CAP</sub><sub>BLS</sub></b>	Field capacity of $\mathbb{F}_{\mathbf{r}_{\text{BLS}}}$ .	$\lfloor \log_2 \mathbf{r}_{\text{BLS}} \rfloor = 252$ bits
211	<b>FIELD<sub>LEN</sub><sub>BLS</sub></b>	Bit-length of a field element $x \in \mathbb{F}_{\mathbf{r}_{\text{BLS}}}$	$\lceil \log_2 \mathbf{r}_{\text{BLS}} \rceil = 253$ bits
212	<b>FIELD<sub>CAP</sub><sub>BN</sub></b>	Field capacity of $\mathbb{F}_{\mathbf{r}_{\text{BN}}}$ .	$\lfloor \log_2 \mathbf{r}_{\text{BN}} \rfloor = 253$ bits
213	<b>FIELD<sub>LEN</sub><sub>BN</sub></b>	Bit-length of a field element $x \in \mathbb{F}_{\mathbf{r}_{\text{BN}}}$	$\lceil \log_2 \mathbf{r}_{\text{BN}} \rceil = 254$ bits
214	<b>BYTELEN</b>	Bit-length of a byte	8 <i>bits</i>
215	<b>ENCZETHNOTELEN</b>	Size of an encrypted note (see Section 3.5.3)	$\text{CTBYTELEN} * \text{BYTELEN}$ <i>bits</i>
216	<b>ETHWORDLEN</b>	Width of a storage cell on the Ethereum Virtual Machine stack, i.e. size of	
217		a word on the EVM	256 <i>bits</i>

218	<b>FIELD CAP</b>	Field capacity of $\mathbb{F}_r$ , defined as the maximum bit length $l$ such that all	
219		numbers $x$ encoded on $l$ bits are elements of $\mathbb{F}_r$ . In other words, $\text{FIELD CAP} =$	
220		$\max_{x \in \mathbb{F}_r} \{\lceil \log_2 x \rceil\}$ s.t. $\sum_{i \in [\text{FIELD CAP}]} 2^i \in \mathbb{F}_r$	
221	<b>FIELD LEN</b>	Bit-length of elements in field element $x \in \mathbb{F}_r$	$\lceil \log_2 r \rceil$ bits
222	<b>JSIN, JSOUT, JS MAX</b>	The number of inputs and outputs of a joinsplit and $\text{JS MAX} = \max \{\text{JSIN}, \text{JSOUT}\}$	
223	<b>KEK256 DLEN</b>	Message digest size of Keccak256 [GJMG11]	256 bits
224	<b>MKDEPTH</b>	The depth of the Merkle tree used to store commitments	
225	<b>p<sub>SECP</sub></b>	Prime defining the finite field of curve secp256k1 [wik]	
226	<b>r<sub>BLS</sub></b>	Characteristic of the scalar field of BLS12-377, $r_{\text{BLS}} = 84444617494283704242488$	
227		$24938781546531375899335154063827935233455917409239041$ [BCG <sup>+</sup> 20]	
228	<b>r<sub>BN</sub></b>	Characteristic of the scalar field of BN-254, $r_{\text{BN}} = 2188824287183927522246405$	
229		$745257275088548364400416034343698204186575808495617$ [Rk19]	
230	<b>r</b>	Characteristic of the scalar field of some chosen curve <b>Curve</b>	
231	<b>SECP FIELD LEN</b>	Bit-length of a field element $x \in \mathbb{F}_{\text{pSECP}}$	$\lceil \log_2 \text{pSECP} \rceil = 256$ bits
232	<b>SHA256 BLEN</b>	Block size of SHA256 [oST15, Figure 1]	512 bits
233	<b>SHA256 DLEN</b>	Message digest size of SHA256 [oST15, Figure 1]	256 bits
234	<b>SHA256 MLEN</b>	Message size of SHA256 [oST15, Figure 1]	$< 2^{64}$ bits
235	<b>DGAS</b>	The default/intrinsic gas cost of an <b>Ethereum</b> transaction	21000 Wei
236	<b>ZVALUE LEN</b>	Size in bits of the transferable maximal value	64 bits

# Change log

- 238 • **Version:** 0.0, **Date:** 04/12/2019, **Contributor:** Antoine Rondelet, **Descrip-**  
 239 **tion:** Creation of the document. Established initial table of content and started  
 240 to populate sections with bullet lists to develop in further versions of the document.
- 241 • **Version:** 0.1, **Date:** 20/12/2019, **Contributor:** Antoine Rondelet, **Descrip-**  
 242 **tion:** Refactored the structure of the document. Finalized the table of content,  
 243 wrote sections on notations and preliminaries, and introduced the content related  
 244 to the malleability fix.
- 245 • **Version:** 0.2, **Date:** 24/02/2020, **Contributor:** Clearmatics Cryptography R&D,  
 246 **Description:**
  - 247 – **Date:** 26/02/2020, **Contributor:** Duncan Tebbs, **Description:** Wrote sec-  
 248 tion on wallet implementation and side-channel attacks considerations.
  - 249 – **Date:** 02/03/2020, **Contributor:** Giuseppe Giffone, **Description:** Changed  
 250 Merkle tree hash function to MiMC compression function.
  - 251 – **Date:** 04/03/2020, **Contributor:** Duncan Tebbs, Michal Zajac, **Descrip-**  
 252 **tion:** Added background on Groth16 SNARK and SNARK scheme instanti-  
 253 ation in the protocol.
  - 254 – **Date:** 04/03/2020, **Contributor:** Raphael Toledo, **Description:** Wrote  
 255 section on the packing policy and corresponding attack.
  - 256 – **Date:** 04/03/2020, **Contributor:** Duncan Tebbs, **Description:** Refactored  
 257 the data structures preliminary section.
  - 258 – **Date:** 24/03/2020, **Contributor:** Raphael Toledo, **Description:** Changed  
 259 the PRF and commitment instantiation with Blake2s compression function.
  - 260 – **Date:** 17/04/2020, **Contributor:** Giuseppe Giffone, **Description:** Added  
 261 DHAES encryption scheme.
- 262 • **Version:** 0.3, **Date:** 09/06/2020, **Contributor:** Antoine Rondelet, **Descrip-**  
 263 **tion:** Fixed various inconsistencies throughout the document (notational mistakes  
 264 in document body and in proofs, latex macros, and typos).
- 265 • **Version:** 1.0, **Date:** 30/06/2020, **Contributor:** Antoine Rondelet, Duncan  
 266 Tebbs, Michal Zajac, **Description:** Global review of the document: fixed incon-  
 267 sistencies in definitions and notations, corrected grammatical mistakes and typos,

268 added examples, figures and merged sections 5 and 6 of Chapter 1 for clarity, added  
269 missing references.

270 • **Version:** 1.1, **Date:** 06/11/2020, **Contributor:** Duncan Tebbs, **Description:**  
271 Specification in terms of a generic curve, with constants provided for BN-254 and  
272 BLS12-377. Some clarification and grammar fixes.

# Chapter 1

## Preliminaries

**Zeth** is a protocol which enables private transactions on **Ethereum** [Woo19]. It is a modification of the Decentralized Anonymous Payment (DAP) system **ZeroCash** [BSCG<sup>+</sup>14]. The design described in [RZ19] presents the mechanisms by which **ZeroCash** can be used on **Ethereum**, and argues that the information leakages of the solution are well defined and controlled. This document, however, serves as a specification of the protocol and provides security fixes and optimizations from the first proof of concept release of the protocol [Cle19].

This document assumes familiarity with blockchain and **Ethereum** in particular. It does not, in any way, aim at replacing the Ethereum yellow paper [Woo19]. The reader is strongly advised to read about **Ethereum** before delving into this specification document.

The key words **MUST**, **MUST NOT**, **SHOULD**, **SHOULD NOT**, **MAY**, and **RECOMMENDED** in this document are to be interpreted as described in [Bra97] when they appear in **ALL CAPS**. These words may also appear in this document in lower case as plain English words, absent their normative meanings.

### 1.1 Data structures and representation

#### 1.1.1 Structured data

When describing the operations to be performed and the data to be manipulated as part of the protocol, we commonly employ tuples of related data where each element of the tuple has some associated semantic meaning and which must often satisfy some conditions. In this section, we develop a framework to reason about such *structured* data, where a single datum may consist of one or more logical parts (called *fields*). The framework is built on top of simple mathematical concepts such as sets, and mappings between them, ensuring that we can always reason about structured data in a rigorous way. We also define notation to aid the specification of structured data, and to refer to specific components of a datum. This will be used extensively in the specification of the protocol.

301 As a simple motivating example, consider a protocol that processes data relating to  
 302 individual people. This fictional system may send and receive data such as *name*, *age*  
 303 and *address* for a single person, grouping this data into a logical unit. Further, each  
 304 piece of data must satisfy specific conditions (*name* must be a series of characters from  
 305 some alphabet, *age* must be a positive integer, etc.) We shall make use of this example  
 306 several times during the formulation below.

307 In what follows, let  $\text{STR} = \{a, b, \dots, y, z\}^*$  (the Kleene star of the *Roman alphabet*).  
 308 In our formulation, field names  $f_i$  will be elements in this set.

309 **Remark 1.1.1.** Note that a similar formulation could be made using an arbitrary set,  
 310 such as the same alphabet augmented with specific symbols, or the alphabet of a different  
 311 language. Our choice of  $\text{STR}$  here is for simplicity.

312 We begin by defining a data type as a set of values called “fields”, each with a “name”  
 313 from  $\text{STR}$ . Abstract sets are used to constrain the values of each field.

**Definition 1.1.2** (Structured Data Type). Let  $f_0, \dots, f_{n-1}$  be  $n$  distinct elements of  $\text{STR}$  and let  $V_0, \dots, V_{n-1}$  be sets, for some  $n \in \mathbb{N}$ . We define the *structured data type*  $\mathbf{T}$  with fields  $\{(f_i, V_i)\}_{i \in [n]}$  to be a set of values:

$$\mathbf{T} = V_0 \times \dots \times V_{n-1}$$

with associated post-fix “dot” operators  $.f_i : \mathbf{T} \rightarrow V_i$  for  $i = 0, \dots, n-1$ , acting on values  $\mathbf{x} \in \mathbf{T}$  to extract the individual elements:

$$\mathbf{x}.f_i = v_i, \text{ where } \mathbf{x} = (v_0, \dots, v_{n-1}) \in \mathbf{T}$$

314 Here, we say that the  $i$ -th field has *field name*  $f_i$ , with *value set*  $V_i$ . Each “dot”  
 315 operator  $.f_i$  *extracts* the  $i$ -th component, or the *value with field name*  $f_i$ .

**Example 1.1.3.** Consider our example protocol that processes information about people. A potentially useful structured data type **Person** may be defined with fields:

$$\{(name, \text{STR}), (age, \mathbb{N}), (height, \mathbb{R}^+)\}$$

316 Values  $\mathbf{p}$  in **Person** are simply tuples in  $\text{STR} \times \mathbb{N} \times \mathbb{R}^+$ , with semantic meaning (name,  
 317 age, height) assigned to each component of  $\mathbf{p}$ .

Examples of valid elements in **Person** include  $\mathbf{a} = (alice, 28, 1.65)$  and  $\mathbf{b} = (bob, 31, 1.74)$ , where the following equalities hold:

$$\mathbf{a}.name = alice,$$

$$\mathbf{b}.age = 31,$$

$$\mathbf{b}.height = 1.74;$$

318 For clarity, structured data types may be specified using tables of names, descriptions  
 319 and value sets, rather than sets of the form  $\{(f_i, V_i)\}_{i \in [n]}$ . Similarly, it is frequently  
 320 convenient to include the *field names* alongside values when specifying structured data  
 321 values.

322 **Example 1.1.4.** Person from Example 1.1.3 might be described in table-form as follows:

Field	Description	Data type
<i>name</i>	Name of the person	STR
<i>age</i>	Age in years	$\mathbb{N}$
<i>height</i>	Height in meters	$\mathbb{R}^+$

**Example 1.1.5.** The values **a** and **b** in Example 1.1.3 might be written as follows:

$$\begin{aligned}\mathbf{a} &= \{name : \textit{alice}, age : 28, height : 1.65\} \\ \mathbf{b} &= \{name : \textit{bob}, age : 31, height : 1.74\}\end{aligned}$$

323 **Remark 1.1.6** (“dot” operators in assignment). The “dot” operators may be used in  
324 algorithm descriptions to indicate *assignment to a specific component*. For example  
325  $\mathbf{a}.age \leftarrow 29$  means that the value of the *age* field of **a** is replaced by the value 29.

Formally, for a structured data type **T** with fields  $\{(f_i, V_i)\}_{i \in [n]}$  where  $\mathbf{x} = (v_0, \dots, v_{n-1}) \in \mathbf{T}$  and  $v'_i \in V_i$ :

$$\mathbf{x}.f_i \leftarrow v'_i$$

is equivalent to:

$$\mathbf{x} \leftarrow (v_0, \dots, v_{i-1}, v'_i, v_{i+1}, \dots, v_{n-1})$$

326 We define one further operator and related assignment notation, convenient in cases  
327 where  $V_i = X^m$  for sets  $X$  and  $m \in \mathbb{N}$ .

**Definition 1.1.7** (Square bracket operator). For  $m \in \mathbb{N}$  and set  $X$ , define the operator  $[ ] : X^m \times [m] \rightarrow X$  as:

$$\mathbf{x}[i] = x_i \text{ where } \mathbf{x} = (x_0, \dots, x_m)$$

For the set  $X^*$ , the operator takes the form  $[ ] : X^* \times \mathbb{N} \rightarrow X$ , defined as:

$$\mathbf{x}[i] = \begin{cases} x_i & \text{if } \text{length}(\mathbf{x}) > i \text{ where } \mathbf{x} = (x_0, \dots) \\ \perp & \text{otherwise} \end{cases}$$

**Remark 1.1.8** (Square bracket operators in assignment). Similarly to Remark 1.1.6, we develop assignment notation for the square bracket operator  $[ ]$ . Let  $\mathbf{x} = (x_0, \dots, x_{m-1})$  be a member of  $X^m$ , and  $x'_i$  be some element in  $X$ . The statement:

$$\mathbf{x}[i] \leftarrow x'_i$$

is equivalent to:

$$\mathbf{x} \leftarrow (x_0, \dots, x_{i-1}, x'_i, x_{i+1}, \dots, x_{m-1})$$

328 Informally, this can be interpreted as replacing the  $i$ -th component of  $\mathbf{x}$  with  $x'_i$ .

329 **Remark 1.1.9** (Deep structures and chained “dot” operators). Consider the case of  
 330 structured data  $\mathbf{T}$  with fields  $\{(f_i, V_i)\}_{i \in [n]}$  for  $n \in \mathbb{N}$ . Let  $\mathbf{T}'$  be another structured data  
 331 type with fields  $\{(f'_i, V'_i)\}_{i \in [n']}$  for  $n' \in \mathbb{N}$ , and assume that  $V_j = \mathbf{T}'$  for some  $j \in [n]$ .  
 332 Informally, the values of the  $j$ -th field of elements of  $\mathbf{T}$  are themselves structured data  
 333 of type  $\mathbf{T}'$ .

334 In this case, “dot” operators may be *chained*, so that  $\mathbf{x}.f_j.f'_k$  refers to the  $k$ -th field  
 335  $v'_k$  of the  $j$ -th field  $v_j$  of  $\mathbf{x} \in \mathbf{T}$ .

336 **Example 1.1.10.** Define a structured data type **Address** with fields  $(country, \text{STR}), (zipcode, \text{STR})$ .  
 337 We redefine the structured data type **Person** from Example 1.1.3, with an extra field  
 338 *address* of type **Address**. That is, **Person** is the structured data type with fields:

Field	Description	Data type
<i>name</i>	Name of the person	STR
<i>age</i>	Age in years	$\mathbb{N}$
<i>height</i>	Height in meters	$\mathbb{R}^+$
<i>address</i>	Address of the person	<b>Address</b>

An example element  $\mathbf{a}$  in **Person** is:

$$\mathbf{a} = \{ \begin{array}{l} name : \textit{alice}, \\ age : 28, \\ height : 1.65, \\ address : (country : UK, zipcode : SW1A) \end{array} \}$$

In this case, the following equalities using the dot and square bracket operators all hold:

$$\begin{aligned} \mathbf{a}.name &= \textit{alice} \\ \mathbf{a}.height &= 1.65 \\ \mathbf{a}.address.country &= UK \\ \mathbf{a}.address.zipcode &= SW1A \\ \mathbf{a}.address.country[1] &= K \end{aligned}$$

## 339 1.1.2 Representations

340 The binary alphabet  $\{0, 1\}$ , denoted  $\mathbb{B}$ , is used to represent the presence or absence of an  
 341 electrical signal in a computer. In fact, every piece of information in a computer is rep-  
 342 resented as a string of bits. We assume the existence of an efficient binary representation  
 343 for some set of primitive datatypes (such as the natural numbers  $\mathbb{N}$ , or alphanumeric



344 characters). Structured data types built up from primitive types (as described above)  
 345 can then recursively be assigned similarly efficient representations. This is used to define  
 346 the following functions to *encode* data to its bit-string representation, and to *decode* such  
 347 bit-strings back to elements of the original type.

**Definition 1.1.11.** For a set  $X$  of values which are to be represented as bit strings, we define functions:

$$\begin{aligned}\text{encode}_X &: X \rightarrow \mathbb{B}^* \\ \text{decode}_X &: \mathbb{B}^* \rightarrow X \cup \perp\end{aligned}$$

satisfying

$$\text{decode}_X(\text{encode}_X(x)) = x \quad \forall x \in X$$

348 to be the functions which encode (resp. decode) elements of  $X$  into (resp. from) the  
 349 bit-string representations chosen above. We note that  $\text{decode}_X$  may return  $\perp$  in the case  
 350 that the input bit-string is malformed.

351 Without ambiguity, we overload the functions **encode** and **decode** to mean  $\text{encode}_X$   
 352 and  $\text{decode}_X$  where the set  $X$  is clear from context.

In the following sections, we assume that elements of  $\mathbb{N}$  are encoded as big-endian binary numbers in the natural way. We denote by  $\mathbb{N}_b$  the set of natural numbers that can be uniquely encoded in this way using  $b$  bits (possibly with padding). In other words,

$$\mathbb{N}_b = \left\{ x \in \mathbb{N} \text{ s.t. } \text{encode}_{\mathbb{N}}(x) \in \mathbb{B}^b \right\}$$

## 353 1.2 Ethereum

354 In a nutshell, **Ethereum** is a distributed deterministic state machine, consisting of a glob-  
 355 ally accessible singleton state (“the World state”) and a virtual machine that applies  
 356 changes to that state [AG18]. State transitions in the state machine are represented  
 357 by transactions on the system. As such, each transaction represents a change in the  
 358 global state represented as a Merkle Patricia Tree [wc] whose nodes are objects called  
 359 “accounts” (Section 1.2.1). The Ethereum Virtual Machine (EVM) allows state transi-  
 360 tions to be specified by creating a type of accounts which are associated with a piece  
 361 of code (smart-contracts). The code of such accounts, and so, the corresponding state  
 362 transitions, can be executed to transition to another state in the automata, by creating  
 363 a transaction that calls the given piece of code (Section 1.2.2).

364 To prevent unbounded state transitions in the state machine, each instruction exe-  
 365 cuted by the EVM is associated with a cost in **Wei**, referred to as “the gas necessary to  
 366 run the operation”. The “gas cost” of a transaction needs to be paid by the transaction  
 367 originator (deduced from their account balance), and is awarded to the miner (added  
 368 to their account balance) who successfully mines the block containing the transaction.  
 369 In addition to the cost of every instruction executed as part of a state transition, every  
 370 transaction has an intrinsic “gas cost” of  $\text{DGAS Wei}$  [Woo19, Appendix G]. Bounding

modifications to the **Ethereum** state by the amount of **Wei** held in the transaction originator’s account allows the system to avoid the Halting problem<sup>1</sup> and protects against a range of Denial of Service (DOS) attacks.

### 1.2.1 Ethereum account

An **Ethereum** account [Woo19, Section 4.1] is an object containing 4 attributes, as represented Table 1.1. We distinguish two types of accounts:

- “Externally Owned Accounts” (EOA), that are created by derivation of an ECDSA secret key; and
- Smart-contract accounts, that are derived from EVM code specifying a state transition on the state machine.

Each account object is accessible in the Merkle Patricia Tree representing the “World state” by a unique **ADDRLEN**-bit long identifier called the address. In the context of EOA, the address is obtained by generating a new ECDSA [JMV01] key pair  $(sk, vk)$  over curve **secp256k1** [Qu99] and taking the rightmost **ADDRLEN** bits of the **Keccak256** hash of the verification key  $vk$ .

Field	Description	Data type
<i>nce</i>	The nonce of an account is a scalar value representing the number of transactions that have originated from the account, starting at 0.	$\mathbb{N}^{\text{ETHWORDLEN}}$
<i>bal</i>	The balance of an account is a scalar value representing the amount of <b>Wei</b> in the account.	$\mathbb{N}^{\text{ETHWORDLEN}}$
<i>sRoot</i>	The storage root is the <b>Keccak256</b> hash representing the storage of the account.	$\mathbb{B}^{\text{KEK256DLEN}}$
<i>codeh</i>	The code hash is the hash of the EVM code governing the account. If this field is the <b>Keccak256</b> hash of the empty string, then the account is said to be an “Externally owned Account” (EOA), and is controlled by the corresponding ECDSA private key. If, however, this field is not the <b>Keccak256</b> hash of the empty string, the account represents a smart contract whose interactions are governed by its EVM code.	$\mathbb{B}^{\text{KEK256DLEN}}$

Table 1.1: Ethereum Account structure

<sup>1</sup>[https://en.wikipedia.org/wiki/Halting\\_problem](https://en.wikipedia.org/wiki/Halting_problem)

### Note

In the rest of this document, we will refer to an *Ethereum user*  $\mathcal{U}_{\mathcal{E}}$  as a person, modeled as an object, holding *one*<sup>a</sup> secret key,  $sk$  (object attribute), associated with an existing EOA in the “World state”. We denote by  $\mathcal{U}_{\mathcal{E}}.Addr$  the **Ethereum** address of  $\mathcal{U}_{\mathcal{E}}$  derived from  $\mathcal{U}_{\mathcal{E}}.sk$ , and which allows  $\mathcal{U}_{\mathcal{E}}$  to access the state of their account  $\varsigma[\mathcal{U}_{\mathcal{E}}.Addr]$ .

We denote by **SmartC** a smart-contract instance/object (i.e. deployed smart-contract with an address, Section 1.2.2), and denote by **SmartC.Addr** its address.

<sup>a</sup>The same physical person may correspond to multiple “Ethereum users” and thus control multiple accounts in the Merkle Patricia Tree.

## 1.2.2 Ethereum transaction

We now briefly mention what **Ethereum** transactions [Woo19, Section 4.2] are, and how they are created, signed and validated. Once more, the reader is highly encouraged to refer to [Woo19] for a detailed presentation. Informally, a transaction object ( $tx$ ) is a signed message originating from an **Ethereum** user  $\mathcal{U}_{\mathcal{E}}$  (the *transaction originator*, or simply *sender*) that represents a state transition on the distributed state machine (i.e. a change in the “World state”  $\varsigma$ ).

### Raw transaction

In the following, we define a raw transaction as an unsigned transaction (Table 1.2).

Field	Description	Data type
$nce$	Transaction nonce	$\mathbb{N}_{\text{ETHWORDLEN}}$
$gasP$	gasPrice	$\mathbb{N}_{\text{ETHWORDLEN}}$
$gasL$	gasLimit	$\mathbb{N}_{\text{ETHWORDLEN}}$
$to$	Recipient’s address	$\mathbb{B}^{\text{ADDRLEN}}$
$val$	Value of the transaction in Wei	$\mathbb{N}_{\text{ETHWORDLEN}}$
$init / data$	Contract Creation data $init$ Message call data $data$	$\mathbb{B}^*$

Table 1.2: Structure of a *raw transaction data type* TxRawDType

### Finalizing raw transactions

A raw transaction needs to be finalized to be accepted. In the context of this document, “finalizing a raw transaction” will be a synonym of “signing a raw transaction”. The transaction structure is represented in Table 1.3.

Field	Description	Data type
$tx_{raw}$	Raw transaction object	<b>TxRawDType</b>
$v$	Field $v$ of ECDSA signature used for public key recovery	$\mathbb{B}^{\text{BYTELEN}}$
$r$	Field $r$ of ECDSA signature [Por13]	$\mathbb{F}_{\text{PSECP}}$
$s$	Field $s$ of ECDSA signature [Por13]	$\mathbb{F}_{\text{PSECP}}$

Table 1.3: Structure of a (finalized) *transaction data type* TxDType

We define the transaction generation function, cf. Fig. 1.1, as the function taking the sender’s ECDSA signing key and the components of a raw transaction as arguments, and returning a signed (or finalized) transaction ( $tx_{final}$  or  $tx$  for short).

$$\begin{aligned}
tx_{final} &= \text{TxGen}(sk_{\text{ECDSA}}, nce_{in}, gasP_{in}, gasL_{in}, to_{in}, val_{in}, init_{in}, data_{in}) \\
tx_{final} &= \{ \\
&\quad \left. \begin{array}{ll} nce & : nce_{in}, \\ gasP & : gasP_{in}, \\ gasL & : gasL_{in}, \\ to & : to_{in}, \\ val & : val_{in}, \\ init/data & : init_{in}/data_{in}, \end{array} \right\} tx_{raw} \\
&\quad \left. \begin{array}{l} v : \sigma_{\text{ECDSA}}.v, \\ r : \sigma_{\text{ECDSA}}.r, \\ s : \sigma_{\text{ECDSA}}.s \end{array} \right\} \sigma_{\text{ECDSA}} \\
&\}
\end{aligned}$$

400 To sign a transaction, the sender first computes the hash of the raw transaction using  
401 Keccak256, cf. Eq. (1.1), and then uses their ECDSA signing key,  $sk_{\text{ECDSA}}$ , to sign the  
402 obtained digest. cf. Eq. (1.2). The signature is then appended to the raw transaction to  
403 obtain a finalized transaction, cf. Fig. 1.1.

$$digest_{\text{ECDSA}} = \text{Keccak256}(nce_{in}, gasP_{in}, gasL_{in}, to_{in}, val_{in}, init_{in}/data_{in}) \quad (1.1)$$

$$\sigma_{\text{ECDSA}} = \text{SigSch}_{\text{ECDSA}}.\text{Sig}(sk_{\text{ECDSA}}, digest_{\text{ECDSA}}) (= (v, r, s)) \quad (1.2)$$

---

```

TxGen( $sk_{\text{ECDSA}}, nce_{in}, gasP_{in}, gasL_{in}, to_{in}, val_{in}, init_{in}, data_{in}$ )
1 : if  $to_{in} = \emptyset$  do
2 :    $tx_{raw} \leftarrow \{nce : nce_{in}, gasP : gasP_{in}, gasL : gasL_{in}, to : to_{in}, val : val_{in}, init : init_{in}\};$ 
3 : else
4 :    $tx_{raw} \leftarrow \{nce : nce_{in}, gasP : gasP_{in}, gasL : gasL_{in}, to : to_{in}, val : val_{in}, data : data_{in}\};$ 
5 : endif
6 :  $\sigma_{\text{ECDSA}} \leftarrow \text{SigSch}_{\text{ECDSA}}.\text{Sig}(sk_{\text{ECDSA}}, \text{Keccak256}(tx_{raw}));$ 
7 :  $tx_{final} \leftarrow \{tx_{raw}, v : \sigma_{\text{ECDSA}}.v, r : \sigma_{\text{ECDSA}}.r, s : \sigma_{\text{ECDSA}}.s\};$ 
8 : return  $tx_{final};$ 

```

---

Figure 1.1: Transaction generation function TxGen

**Remark 1.2.1.** As one can see, there is no “from” attribute in a transaction. The sender’s Ethereum address can be recovered from the ECDSA signature. This method is defined in the Ethereum yellow paper as a “sender function”  $S$  [Woo19, Appendix F] which maps each transaction to its sender.

## Types of transactions

While only two types of transactions are described in [Woo19, Section 4.2]; namely those which result in message calls and those which result in the creation of new accounts with associated code, we will instead differentiate the types of transactions based on their purpose. The reader is encouraged to read [Woo19] for a formal discussion.

Informally, a transaction can be used to achieve three things: transferring Wei from an EOA to another EOA, creating a new account with associated code (i.e. “deploying a smart-contract”), and calling a function of a smart-contract. We will detail here the differences between these usages.

**Creating a contract** The  $tx.to$  address is set to  $\emptyset$  in the transaction. The contract creation data ( $tx.init$ ) includes the new contract’s code. The contract address is computed as the rightmost ADDRLEN bits of the Keccak256 hash of the RLP encoding [wc19] of the transaction originator’s address and account nonce [Woo19, Section 6].

**Calling a contract function** The  $tx.to$  address is set to the address of the contract. The message call data byte array ( $tx.data$ ) is set to the contract’s function address (or “Function Selector” [abi]) which are the first 4 bytes of the Keccak256 hash of the function signature, and the function input arguments (ETHWORDLEN bits per input) [Woo19, Section 8].

**Transferring Wei from an EOA to another EOA** This corresponds to a “plain transaction” spending Wei from an address to send them to another. In that case the  $tx.to$  address corresponds to the recipient’s address while the transaction data is left empty.

### Note

In order to keep notations simple, we assume, in the rest of the document, that smart-contract functions are uniquely determined by their name. As such, we denote by  $\text{FS}(\cdot): \mathbb{B}^* \rightarrow \mathbb{B}^{4 \cdot \text{BYTELEN}}$  the function that takes a function name as input and returns its function selector.

## Transaction validity

Importantly, not all finalized transactions constitute valid state transitions on the state machine [Woo19, Section 6]. We denote by `EthVerifyTx` the function that takes an `Ethereum` transaction object  $tx$  as input and return `true` (resp. `false`) if  $tx$  is valid (resp. invalid). To be deemed valid, a transaction **MUST** satisfy *all* the following conditions:

1. The transaction is correctly RLP encoded, with no additional trailing bytes;
2. the transaction signature  $(v, r, s)$  is valid;
3. the transaction nonce  $(tx.nce)$  is valid, i.e. it is equal to the account nonce of the transaction originator;
4. the gas limit is no smaller than the gas used by the transaction;
5. the transactor has enough funds on his account balance to cover at least the cost  $tx.val + tx.gasP \cdot tx.gasL$ .

## Lifecycle of a transaction, and miners' incentives

After the creation of an `Ethereum` transaction  $tx$  by a user from an `Ethereum` client (machine running a piece of software that enables to be connected to the `Ethereum` network), the transaction is broadcasted to the network and received by a set of peers/nodes.

The transaction is then stored in each node's transaction pool, which is a data structure containing all transactions that should be validated (pending transactions) by the node and mined. To maximize miners' returns, the transaction pools are ordered according to the gas price of the transactions. As such, transactions with the highest  $tx.gasP$  are subject to be validated and included into a block first. Once  $tx$  is selected from the transaction pool, it is validated (fed into `EthVerifyTx`), executed, and included into a block (i.e. "mined"). The block is then broadcasted to all the nodes of the network and is used as the predecessor for the next block to be mined on the network (i.e. "it is added to the chain").

### 1.2.3 Ethereum events and Bloom filters

The EVM contains the set of "LOGX" instructions enabling smart-contract functions to "emit events" (i.e. log data) when they are executed<sup>2</sup>

---

<sup>2</sup>see <https://ethgastable.info/>

As such, when a block is generated by a miner or verified by the rest of the network, the address of any logging contract, and all the indexed fields from the logs generated by executing those transactions are added to a Bloom filter [Blo70], which is included in the block header [Woo19, Section 4.3]. Importantly, the actual logs *are not included in the block data* in order to save space. As such, when an application wants to find (“consume”) all the log entries from a given contract, or with specific indexed fields (or both), the node can quickly scan over the header of each block, checking the Bloom filter to see if it may contain relevant logs. If it does, *the node re-executes the transactions from that block, regenerating the logs, and returning the relevant ones to the application* [Joh16].

#### Note

The ability for a smart-contract function to “emit” some pieces of data when executed, and for an application to “consume” such pieces of data, is used in Zeth in order to construct a *confidential receiver-anonymous channel* [KMO<sup>+</sup>13].

## 1.3 zk-SNARKs

In this section we introduce notions necessary to understand zero-knowledge proofs, define properties crucial for them, and introduce zk-SNARKs. We refer the reader to Section 3.6 in which we describe the zk-SNARK scheme used in Zeth.

### 1.3.1 Preliminary definitions

**NP class of languages.** Since the considered proof systems are designed to work with languages in NP we begin with defining this class. Intuitively, a language  $\mathbf{L}$  belongs to NP if for each element  $prim$  from the language there is a short witness  $aux$  that allows to efficiently<sup>3</sup> verify that in fact  $prim \in \mathbf{L}$ .

**Definition 1.3.1** (NP class of languages, cf. [Gol01]). We say that a language  $\mathbf{L}$  belongs to a class NP if there exist a polynomial  $p$  and a Turing machine  $M$  such that for every primary input  $prim \in \{0, 1\}^*$ ,  $prim \in \mathbf{L}$  iff there exists an auxiliary input  $aux$  such that  $M$  accepts the pair  $(prim, aux)$  in time at most  $p(\text{length}(prim))$ .

The set of all pairs  $(prim, aux)$  acceptable by  $M$  constitutes an NP *relation*  $\mathbf{R}$  corresponding to the language  $\mathbf{L}$ .

**Non-interactive zero knowledge.** A non-interactive zero-knowledge proof system NIZK for an NP language  $\mathbf{L}$  is a tuple of four algorithms  $\text{NIZK} = (\text{KGen}, P, V, \text{Sim})$ . NIZK for a language  $\mathbf{L}$  and instance  $prim \in \mathbf{L}$  allows a party, called prover and denoted by  $P$ , to convince another party, called verifier and denoted by  $V$ , that  $prim \in \mathbf{L}$  and nothing else.

Without loss of generality, we focus on zk-proof systems that are universal, that is, are able to work with any given NP relation  $\mathbf{R}$ . To that end, we define a *relation*

<sup>3</sup>Informally we say that an algorithm is efficient if it runs in time polynomial in the size of its inputs.

generator  $\mathcal{R}$  that on input  $1^\lambda$  (i.e. the security parameter represented in unary) outputs an NP relation  $\mathbf{R}$ . We assume that the security parameter  $\lambda$  can be easily deduced from  $\mathbf{R}$ .

We require from a NIZK to have three substantial properties, cf. [Gro06]:

**Completeness** that assures that an honest prover, who proves that  $\text{prim} \in \mathbf{L}$  succeeds, i.e. gets his proof accepted by the verifier  $\mathbf{V}$ . Formally we require that for any  $\lambda$ ,  $\mathbf{R} \leftarrow \mathcal{R}(1^\lambda)$ ,  $(\text{prim}, \text{aux}) \in \mathbf{R}$

$$\Pr \left[ \mathbf{V}(\mathbf{R}, \text{crs}, \text{prim}, \mathbf{P}(\mathbf{R}, \text{crs}, \text{prim}, \text{aux})) \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (\text{crs}, \text{td}) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda) \end{array} \right] = 1 .$$

**Computational soundness** which states that in case  $\text{prim} \notin \mathbf{L}$  the verifier accepts the proof for  $\text{prim}$  with negligible probability only. Formally we require that for any  $\mathbf{R} \leftarrow \mathcal{R}(1^\lambda)$  and PPT adversary  $\mathcal{A}$

$$\Pr \left[ \mathbf{V}(\mathbf{R}, \text{crs}, \text{prim}, \pi) \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (\text{crs}, \text{td}) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (\text{prim}, \pi) \leftarrow \mathcal{A}(\mathbf{R}, \text{crs}); \\ \text{prim} \notin \mathbf{L} \end{array} \right] \leq \text{negl}(\lambda).$$

**Zero knowledge** assures that the verifier learns from a proof nothing except the veracity of the proven statement. More precisely we require that there exist a PPT algorithm  $\text{Sim}$  and negligible function  $\eta(\lambda)$  such that for every adversary  $\mathcal{A}$  and security parameter  $\lambda$

$$\left| \Pr \left[ \mathcal{A}(\mathbf{R}, \text{crs}, \pi) = 1 \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (\text{crs}, \text{td}) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (\text{prim}, \text{aux}) \leftarrow \mathcal{A}(\mathbf{R}, \text{crs}); \\ (\text{prim}, \text{aux}) \in \mathbf{R}; \\ \pi \leftarrow \text{Sim}(\mathbf{R}, \text{crs}, \text{td}, \text{prim}) \end{array} \right] - \Pr \left[ \mathcal{A}(\mathbf{R}, \text{crs}, \pi) = 1 \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (\text{crs}, \text{td}) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (\text{prim}, \text{aux}) \leftarrow \mathcal{A}(\mathbf{R}, \text{crs}); \\ (\text{prim}, \text{aux}) \in \mathbf{R}; \\ \pi \leftarrow \mathbf{P}(\mathbf{R}, \text{crs}, \text{prim}, \text{aux}) \end{array} \right] \right| \leq \eta(\lambda).$$

We say that NIZK is *perfectly* zero-knowledge if  $\eta = 0$ .

We note that the existence of the simulator which by using the trapdoor is able to make a proof for a false statement (i.e. for  $\text{prim} \notin \mathbf{L}$ ) makes the whole zk-proof system



499 vulnerable to adversaries that also know the trapdoor. More precisely, an adversary  
 500 who knows a trapdoor  $td$  can break the soundness property. This vulnerability comes  
 501 with each CRS-based NIZK (for languages in NP). Thus in the real-life deployment of a  
 502 CRS-based NIZK it has to be enforced that nobody learns the trapdoor.

503 A zk-SNARK scheme, denoted  $\text{ZkSnarkSch}$ , is a special type of NIZK which is equipped  
 504 with two more properties. First, zk-SNARKs are arguments *of knowledge*, as such they  
 505 have to follow a stronger definition of soundness, called *knowledge soundness*.

**Knowledge soundness** assures that if a prover provided a proof  $\pi$  for a statement  
 $prim$  acceptable to a verifier, then she knows the corresponding auxiliary input  $aux$ .  
 More precisely, we require that for each  $\mathbf{R} \leftarrow \mathcal{R}(1^\lambda)$ , and malicious PPT prover  $\mathcal{A}$   
 there exists a machine  $\text{Ext}_{\mathcal{A}}$ , called extractor, that given access to randomness  $r$   
 used by  $\mathcal{A}$  and its inputs, *extracts* the auxiliary input  $aux$  from  $\mathcal{A}$ ; that is:

$$\Pr \left[ \begin{array}{c} \neg(\mathbf{R}(prim, aux)) \wedge \\ \mathbf{V}(\mathbf{R}, crs, prim, \pi) \end{array} \middle| \begin{array}{c} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (crs, td) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (prim, \pi) \leftarrow \mathcal{A}(\mathbf{R}, crs; r); \\ aux \leftarrow \text{Ext}_{\mathcal{A}}(\mathbf{R}, crs; r) \end{array} \right] \leq \text{negl}(\lambda).$$

506 Second, zk-SNARKs are *succinct*, and so we require that proofs produced by  $\text{ZkSnarkSch.P}$   
 507 are short, i.e. sublinear to the size of the primary and auxiliary inputs. Importantly, in  
 508 many modern zk-SNARKs, like [Gro16, MBKM19, Gab19, GWC19, CHM<sup>+</sup>20] the proof  
 509 size is constant regardless the size of the input.

### 510 1.3.2 Computation representation – arithmetization

511 In **Zeth** the sender shows that the transaction is correct by arguing (in zero knowledge,  
 512 i.e. hiding private inputs) about correctness of evaluation of some predefined predicate.  
 513 This predicate ensures that the soundness of the blockchain system is not violated, i.e. the  
 514 zk-proof is used to prove that a transaction follows the “rules of the system” without  
 515 disclosing its attributes. The proof system that **Zeth** uses operates on an algebraic  
 516 representation of the “predicate to prove”. Informally, representing the computation as  
 517 a set of algebraic constraints is called *arithmetization*. One of such representations is  
 518 Quadratic Arithmetic Programs (QAP) [GGPR13], which, following [Gro16], is used in  
 519 **Zeth**. This representation is considered one of the most efficient for general arithmetic  
 520 circuits.

521 **Remark 1.3.2.** Preprocessing SNARKs such as [Gro16] rely on common reference  
 522 strings with a specific structure. As such, we may use  $crs$  and  $srs$  (*structured refer-*  
 523 *ence string*) interchangeably in the rest of this document.

524 **QAP (R1CS).** Let  $C$  be an arithmetic circuit of fan-in 2 over  $\mathbb{F}_p$ . The number of  
 525 multiplication gates in  $C$  is denoted by  $constNo$ . Likewise, the number of all wires in  $C$   
 526 is denoted by  $inpNo$ .

Before we formally introduce the QAP relation  $\mathbf{R}_{\text{QAP}}$  we provide some intuitions behind it. First, we observe that the circuit  $\mathbf{C}$  can be represented by three matrices  $\vec{A}, \vec{B}, \vec{C}$  all in  $\mathbb{F}_p^{\text{constNo} \times \text{inpNo} + 1}$  such that the  $i$ -th row in matrix  $\vec{A}$  (and  $\vec{B}$ ) denotes left (and right) input to the  $i$ -th multiplication gate, which is also the  $k$ -th input to the circuit. That is for a circuit evaluation  $z \in \mathbb{F}_p^{\text{inpNo} + 1}$  the left input for the  $i$ -th gate is  $\sum_{j=0}^{\text{inpNo}} A_{ij} z_j$  and the right input is  $\sum_{j=0}^{\text{inpNo}} B_{ij} z_j$ . Furthermore, entry  $\vec{C}_{ik}$  contains the output of  $i$ -th multiplication gate that is  $k$ -th input to the circuit.

Second, for the sake of efficiency we represent each matrix as a sequence of polynomials. Each matrix's column is represented by a polynomial in  $\mathbb{F}_p[X]$  such that the column's  $i$ -th input equals polynomial's evaluation at  $\omega^i$  – the  $i$ -th primitive root of unity modulo  $p$ . More precisely, we define polynomials:

- $u_j(X)$ , for  $j \in \{0, \dots, \text{inpNo}\}$ , such that  $u_j(\omega^i) = \vec{A}_{ij}$ ;
- $v_j(X)$ , for  $j \in \{0, \dots, \text{inpNo}\}$ , such that  $v_j(\omega^i) = \vec{B}_{ij}$ ;
- $w_j(X)$ , for  $j \in \{0, \dots, \text{inpNo}\}$ , such that  $w_j(\omega^i) = \vec{C}_{ij}$ .

We consider inputs from 1 to  $\text{inpNoPrim}$  public (primary input), for some  $\text{inpNoPrim} \leq \text{inpNo}$ . The rest of the inputs is considered private (auxiliary input). The QAP relation  $\mathbf{R}_{\text{QAP}}$  is defined as follows:

$$\mathbf{R}_{\text{QAP}} = \left\{ (prim, aux) \left| \begin{array}{l} a_0 = 1; prim = (a_1, \dots, a_{\text{inpNoPrim}}) \in \mathbb{F}_p^{\text{inpNoPrim}}; \\ aux = (a_{\text{inpNoPrim}+1}, \dots, a_{\text{inpNo}}) \in \mathbb{F}_p^{\text{inpNo} - \text{inpNoPrim}}; \\ \sum_{j=0}^{\text{inpNo}} a_j u_j(X) \cdot \sum_{j=0}^{\text{inpNo}} a_j v_j(X) = \sum_{j=0}^{\text{inpNo}} a_j w_j(X) \end{array} \right. \right\}.$$

#### Note

Importantly, we note that efficient computation on standard hardware may not necessarily lead to an efficient QAP representation. As such, a function can be very efficient to evaluate on a standard computer, but very slow to evaluate in QAP form.

541

## 1.4 Decentralized Anonymous Payment schemes (DAP)

**Zeth** [RZ19] is a Decentralized Anonymous Payment scheme (DAP) [BSCG<sup>+</sup>14, Section 3] defined on top of an **Ethereum** ledger  $L$ . A DAP is a tuple of polynomial-time algorithms  $\text{DAP} = (\text{Setup}, \text{GenAddr}, \text{SendTx}, \text{VerifyTx}, \text{Receive})$  that manipulate (*create*, *spend*) data objects called *Notes*. These objects are bound to a given owner and have a value  $v$  attribute (see Section 2.1).

**System Setup** The algorithm **Setup** takes the security parameter  $\lambda$  as input and generates the public parameters  $pp$ . The algorithm **Setup** is executed by a trusted

548  
549

550 party. The resulting public parameters  $pp$  are published and made available to all  
 551 parties.

552 **Creating Zeth addresses** The algorithm `GenAddr` takes as input the public parame-  
 553 ters  $pp$  and generates a new DAP address object  $Addr = \{pub : Addr_{pk}, priv : Addr_{sk}\}$ . More precisely,  $Addr_{pk}$  is an object referred to as the “payment ad-  
 554 dress” (Table 1.4), and  $Addr_{sk}$  is an object referred to as the “private address”  
 555 (Table 1.5) [ZCa19].  
 556

557 **Transfer notes** The algorithm `SendTx` is used to transfer some public input  $vin$  as  
 558 well as the value of a set of input (“old”)  $Notes$  into a set of output (“new”)  $Notes$   
 559 as well as some public output value  $vout$ . The inputs  $Notes$  are marked as  
 560 “consumed” (alternatively, we say that the input  $Notes$  are “spent”). `SendTx` takes  
 561 as inputs the public parameters  $pp$ , the input value and the input (“old”)  $Notes$   
 562 to be transferred, as well as the Merkle root and the Merkle authentications paths  
 563 of the commitments to the input  $Notes$ , the “spending keys” related to the input  
 564  $Notes$ , the output value to create and the “payment addresses” for the output  
 565 (“new”)  $Notes$ . If the joinsplit equation is satisfied, the algorithm returns the new  
 566  $Notes$  and the corresponding **Ethereum** transaction  $tx$ , else it returns  $\perp$ .

567 **Verifying transactions** The algorithm `VerifyTx` checks the validity of a transaction.  
 568 It takes as inputs the public parameters  $pp$ , a transaction and the current ledger  
 569  $L$  and outputs a bit equal 1 iff the transaction is valid, 0 otherwise.

570 **Receiving notes** The algorithm `Receive` scans the ledger  $L$  and retrieves unspent  $Notes$   
 571 paid to a particular user address. It takes as input the recipient address key pair  
 572  $\{pub : Addr_{pk}, priv : Addr_{sk}\}$  and the current ledger  $L$  and outputs the set of  
 573 (unspent) received  $Notes$ .

#### Note

In the rest of this document, we will refer to a *Zeth user*  $\mathcal{U}_{\mathcal{Z}}$  as a person, modeled as an object, holding one **Zeth** address (object attribute), and thus holding a *private address*,  $Addr_{sk}$ . We denote by  $\mathcal{U}_{\mathcal{Z}}.Addr$  the **Zeth** address of  $\mathcal{U}_{\mathcal{Z}}$  derived from  $Addr_{sk}$ , and which allows  $\mathcal{U}_{\mathcal{Z}}$  to be the recipient of payments via **Zeth**, and to send funds via **Zeth**. Importantly, *not all Ethereum users are Zeth users, and vice-versa!*

574

Field	Description
$apk$	The <i>paying key</i>
$pkenc$	The <i>transmission key</i>

Table 1.4: “Payment address”,  $Addr_{pk}$ , of a DAP address

Field	Description
<i>ask</i>	The <i>spending key</i>
<i>skenc</i>	The <i>receiving key</i>

Table 1.5: “Private address”,  $Addr_{sk}$ , of a DAP address

575 **Zeth** leverages zk-SNARKs (Section 1.3) and the possibility to deploy smart-contracts  
576 to specify privacy-preserving state transitions altering the **Ethereum** state  $\varsigma$  (Section 1.2).  
577 As such, **Zeth** defines a smart-contract,  $\widetilde{\mathbf{Mixer}}$ , that keeps track of the set of *ZethNotes*  
578 (Section 2.1) in a committed form, stored in a Merkle tree; and which verifies the va-  
579 lidity of the state transitions generated by the **Zeth** users. As such a **Zeth** DAP is  
580 entirely determined by  $\widetilde{\mathbf{Mixer}}$ , the instance of the mixer smart-contract deployed on the  
581 **Ethereum** ledger. State transitions are executed on-chain by calling the **Mix** function of  
582  $\widetilde{\mathbf{Mixer}}$ , which implements the algorithm **VerifyTx** of DAP, and which modifies  $\varsigma$  iff the  
583 transaction is deemed valid.

#### Note

We denote by  $Mix_{in}$  the inputs taken by the **Mix** function defined on  $\widetilde{\mathbf{Mixer}}$ . Let  $zdata$  be the value of the *data* field of an **Ethereum** transaction such that:

$$zdata = FS(\mathbf{Mix}) \parallel Mix_{in}$$

Then, we define  $tx_{\mathbf{Mix}}$  as being the **Ethereum** transaction object returned by **SendTx** such that:

$$tx_{\mathbf{Mix}}.to = \widetilde{\mathbf{Mixer}}.Addr \wedge tx_{\mathbf{Mix}}.data = zdata$$

Importantly, when it is clear from context, we will omit the function selector from the definition of  $zdata$ , and only assume that  $zdata = Mix_{in}$ .

584

## 1.5 Definitions

585

### 1.5.1 Negligible function

586

587 **Definition 1.5.1** (Negligible function, [KL14, Definition 3.4]). A function  $f$  from  $\mathbb{N}$  to  
588  $\mathbb{R}^+$  (positive real numbers) is negligible if for every positive polynomial  $p$  there exists  $N$   
589 such that for all integers  $n > N$  it holds that  $f(n) < \frac{1}{p(n)}$ .

## 1.5.2 Basic algebra notions

**Definition 1.5.2** (Group, see [Bou03, Section I.4]). A group is given by a tuple  $(\mathbb{G}, \otimes)$ , where  $\mathbb{G}$  is a set and  $\otimes$  is a binary operation in  $\mathbb{G}$ , i.e.  $\otimes : \mathbb{G} \times \mathbb{G} \rightarrow \mathbb{G}$ , with the following properties:

- $(\mathbf{g} \otimes \mathbf{h}) \otimes \mathbf{k} = \mathbf{g} \otimes (\mathbf{h} \otimes \mathbf{k})$  (associativity)
- There exists an element  $\epsilon \in \mathbb{G}$  s.t. for each  $\mathbf{g} \in \mathbb{G}$ ,  $\mathbf{g} \otimes \epsilon = \epsilon \otimes \mathbf{g} = \mathbf{g}$  (identity element).
- For each  $\mathbf{g} \in \mathbb{G}$  there exist  $\mathbf{h} \in \mathbb{G}$  s.t.  $\mathbf{g} \otimes \mathbf{h} = \mathbf{h} \otimes \mathbf{g} = \epsilon$  (inverse element).

For simplicity, we may also use the additive notation for groups:  $\otimes$  is denoted as  $+$ , the identity element as  $\mathbf{o}$  and the inverse element of  $\mathbf{g}$  as  $-\mathbf{g}$ . Given  $\mathbf{g} \in \mathbb{G}$  and  $x \in \mathbb{Z}$ , we have that:

$$x \cdot \mathbf{g} = \begin{cases} \mathbf{o} & \text{if } x = 0 \\ \mathbf{g} + \dots + \mathbf{g}, (x \text{ times}) & \text{if } x > 0. \\ -\mathbf{g} + \dots + (-\mathbf{g}), (x \text{ times}) & \text{if } x < 0 \end{cases}$$

**Definition 1.5.3** (Finite Cyclic Group, adapted from [KL14, Sections 7.1.3, 7.3.2]). A finite cyclic group is given by a tuple  $(q, \mathbb{G}, \mathbf{g}, \otimes)$ , called the *group description*, where  $\mathbb{G}$  represents the set of group elements,  $\mathbf{g}$  is a generator and  $q$  is the order. The generator  $\mathbf{g}$  generates the group; namely, each  $\mathbf{h} \in \mathbb{G}$  can be expressed by the generator as  $\mathbf{h} = \mathbf{g} \otimes \dots \otimes \mathbf{g}$ . Given a scalar  $x$ , we denote by  $\llbracket x \rrbracket$  the *encoding* of  $x$  in  $\mathbb{G}$ : i.e.  $\llbracket x \rrbracket = \mathbf{g} \otimes \dots \otimes \mathbf{g}$  ( $x$  times). As consequence,  $\llbracket 1 \rrbracket = \mathbf{g}$ .

For theoretical purposes, we introduce the **SetupG** algorithm that for a given security parameter  $\lambda$  outputs a cyclic group, formally:

**Definition 1.5.4** (Group Setup Algorithm, taken from [KL14, Sections 7.1.3, 7.3.2]). A group setup algorithm **SetupG** is a PPT algorithm which takes as input a security parameter  $1^\lambda$  and outputs a group description  $(q, \mathbb{G}, \mathbf{g}, \otimes)$ , where the binary representation of  $q$  is given by  $\lambda$  bits and each group element can be represented by  $gLen(\lambda)$  bits. Note that  $gLen$  is  $\text{poly}(\lambda)$ .<sup>4</sup>

## 1.5.3 Security assumptions

**Definition 1.5.5** (Discrete Log Problem(DLog), cf. [BS07]). Let  $\mathbb{G}$  denote a group (Section 1.5.2) whose order  $p$  is prime and written over  $\lambda$  bits. We let  $\log_{\mathbf{g}}(h)$  denote the discrete logarithm of  $h$  in the basis  $\mathbf{g}$ . We assume  $\mathbb{G}$ ,  $p$  are fixed and known to all parties. We denote the advantage of a PPT adversary  $\mathcal{A}$  in attacking the discrete logarithm problem as

$$\text{Adv}_{\mathbb{G}, \mathcal{A}}^{\text{dlog}} = \Pr[\mathbf{g} \leftarrow \mathbb{G}^*, x \leftarrow \mathbb{F}_p, x' \leftarrow \mathcal{A}(\llbracket 1 \rrbracket, \llbracket x \rrbracket) : \llbracket x' \rrbracket = \llbracket x \rrbracket]$$

<sup>4</sup>For simplicity, we may denote  $gLen(\lambda)$  as  $gLen$ .

612 We say that the DLog is hard in  $\mathbb{G}$  if and only if  $\text{Adv}_{\mathbb{G}, \mathcal{A}}^{\text{dlog}}(\lambda)$  is negligible for any PPT  
 613 adversary  $\mathcal{A}$ .

**Definition 1.5.6** (One More Discrete Log Problem (om-DLog), cf. [PV05]). Let  $\mathbb{G}$  denote a group whose order  $p$  is prime and written over  $\lambda$  bits. We let  $\log_{\mathbf{g}}(h)$  denote the discrete logarithm of  $h$  in the basis  $\mathbf{g}$ . A PPT adversary  $\mathcal{A}$  solving the om-DLog is given  $q + 1$  random group elements as well as limited access to a discrete logarithm oracle  $\text{O}^{\text{DLog}_{\mathbf{g}}}(q)$ .  $\mathcal{A}$  is allowed to query this oracle at most  $q$  times, thus obtaining the discrete logarithm of  $q$  group elements of his choice with respect to a fixed base  $\mathbf{g}$ . Eventually,  $\mathcal{A}$  must output the  $q + 1$  discrete logarithms. We denote the advantage of a PPT adversary  $\mathcal{A}$  in attacking the one more discrete logarithm problem as

$$\text{Adv}_{\mathbb{G}, \mathcal{A}}^{\text{om-dlog}}(\lambda) = \Pr \left[ \begin{array}{l} \mathbf{g} \leftarrow \mathbb{G}^*, \{ \llbracket r_i \rrbracket \}_{i \in [q+1]} \leftarrow \mathbb{G}^{q+1}, \\ \{ r'_i \}_{i \in [q+1]} \leftarrow \mathcal{A}^{\text{O}^{\text{DLog}_{\mathbf{g}}}(q)}(\llbracket 1 \rrbracket, \{ \llbracket r_i \rrbracket \}_{i \in [q+1]}) : \\ \forall i \in [q+1], r'_i = \log_{\mathbf{g}}(\llbracket r_i \rrbracket) \end{array} \right]$$

614 We say that the om-DLog is hard in  $\mathbb{G}$  if and only if  $\text{Adv}_{\mathbb{G}, \mathcal{A}}^{\text{om-dlog}}(\lambda)$  is negligible for any  
 615 PPT adversary  $\mathcal{A}$ .

#### 616 1.5.4 Symmetric encryption

617 **Definition 1.5.7** (Symmetric Encryption, [KL14, Definition 3.8]). A symmetric encryption  
 618 scheme  $\text{Sym}$  is given by a tuple of PPT algorithms  $(\text{KGen}, \text{Enc}, \text{Dec})$  where:

- 619 • **KGen**, the key generation algorithm, takes a security parameter  $1^\lambda$  and outputs a  
 620 secret key  $ek$ ; we assume, without loss of generality, that  $kLen(\lambda) = \text{length}(ek) \geq \lambda$ .  
 621 Note that  $kLen(\lambda)$  is a polynomial function in  $\lambda$ .<sup>5</sup>
- 622 • **Enc**, the encryption algorithm, takes a key  $ek$ , a plaintext  $m \in \{0, 1\}^*$  and returns  
 623 a ciphertext  $ct$ .
- 624 • **Dec**, the decryption algorithm, takes a key  $ek$  and a ciphertext  $ct$ , and returns a  
 625 message  $m$ . We assume, without loss of generality, that **Dec** is deterministic.

626 For every security parameter  $\lambda$ , key  $ek$  output by  $\text{KGen}(1^\lambda)$ , and message  $m \in \{0, 1\}^*$ ,  
 627 it holds that  $\text{Dec}(ek, \text{Enc}(ek, m)) = m$  (*correctness property*).

628 Let  $(\text{KGen}, \text{Enc}, \text{Dec})$  be a symmetric encryption scheme. If there exists a polynomial  
 629  $l$  such that, for all  $\lambda > 0$  and key  $ek$  output by  $\text{KGen}(1^\lambda)$ ,  $\text{Enc}(ek, \cdot)$  is only defined for  
 630 messages  $m \in \{0, 1\}^{l(\lambda)}$ , then we say that  $(\text{KGen}, \text{Enc}, \text{Dec})$  is a *fixed-length symmetric*  
 631 *encryption scheme* with *length parameter*  $l(\lambda)$ . A security notion for  $\text{Sym}$  follows:

**Definition 1.5.8** (IND-CPA). Let  $\text{Sym}$  be a symmetric encryption scheme and let  $\mathcal{A}$  be an adversary. Consider the IND-CPA game described in Figure 1.2. We define the IND-CPA advantage of  $\mathcal{A}$  as follows:

$$\text{Adv}_{\text{Sym}, \mathcal{A}}^{\text{ind-cpa}}(\lambda) = |2 \cdot \Pr[\text{IND-CPA}(\lambda) = 1] - 1|.$$

---

<sup>5</sup>For simplicity, we may denote  $kLen(\lambda)$  as  $kLen$ .

---

```

IND-CPA( $\lambda$ )
 $ek \leftarrow \text{KGen}(1^\lambda)$ 
 $(m_0, m_1, \text{state}) \leftarrow \mathcal{A}^{\text{O}^{\text{Enc}_{ek}}} \text{ with } \text{length}(m_0) = \text{length}(m_1)$ 
 $b \leftarrow \$\{0, 1\}$ 
 $ct \leftarrow \text{Enc}(ek, m_b)$ 
 $\tilde{b} \leftarrow \mathcal{A}^{\text{O}^{\text{Enc}_{ek}}}(ct, \text{state})$ 
return  $\tilde{b} = b$ 

```

---

Figure 1.2: IND-CPA game for Sym.

632 Sym is said to be IND-CPA secure if, for every PPT adversary  $\mathcal{A}$ , the advantage  $\text{Adv}_{\text{Sym}, \mathcal{A}}^{\text{ind-cpa}}(\lambda)$   
633 is a negligible function.

### 634 1.5.5 Asymmetric encryption

635 **Definition 1.5.9** (Asymmetric encryption, [KL14, Definition 10.1]). An *asymmetric*  
636 *encryption scheme*  $\text{Asym}$  is given by a tuple of PPT algorithms  $(\text{KGen}, \text{Enc}, \text{Dec})$  where:

- 637 •  $\text{KGen}$ , the key generation algorithm, takes a security parameter  $1^\lambda$  and returns a  
638 pair of keys  $(sk, pk)$ . We refer to the first of these as the *private key* and the second  
639 as the *public key*. We assume for convenience that  $pk$  and  $sk$  each have length at  
640 least  $\lambda$ , and that  $\lambda$  can be determined from  $pk, sk$ ;
- 641 •  $\text{Enc}$ , the encryption algorithm, takes a public key  $pk$ , a plaintext  $m$ , from some  
642 underlying plaintext space (that may depend on  $pk$ ) and returns a ciphertext  $ct$ ;
- 643 •  $\text{Dec}$ , the decryption algorithm, takes a private key  $sk$  and a ciphertext  $ct$ , and  
644 returns a message  $m$  or a special symbol  $\perp$  denoting decryption failure. We assume,  
645 without loss of generality, that  $\text{Dec}$  is deterministic.

646 We require that for all  $(sk, pk)$  returned by  $\text{KGen}$ , and every message  $m$  in the appropriate  
647 underlying plaintext space, it holds that  $\text{Dec}(sk, \text{Enc}(pk, m)) = m$  (*correctness property*).

648 Secure communication usually requires ciphertext indistinguishability (e.g. IND-CCA2  
649 [ABR99, Definition 8]). In **Zeth**, however, the key privacy property IK-CCA [BBDP01]  
650 is also required – it ensures indistinguishability of the key under which an encryption is  
651 performed.

**Definition 1.5.10** (IK-CCA). Let  $\text{Asym} = (\text{KGen}, \text{Enc}, \text{Dec})$  be an asymmetric encryption scheme and let  $\mathcal{A}$  be an adversary. Given the IK-CCA game described in Figure 1.3, with the condition that  $\mathcal{A}$  cannot query  $\text{O}^{\text{Dec}_{sk_0}}$  or  $\text{O}^{\text{Dec}_{sk_1}}$  on the challenge ciphertext

---

IK-CCA( $\lambda$ )

$(sk_0, pk_0), (sk_1, pk_1) \leftarrow \text{KGen}(1^\lambda)$

$(m, state) \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(pk_0, pk_1)$

$b \leftarrow \$\{0, 1\}$

$ct \leftarrow \text{Enc}(pk_b, m)$

$\tilde{b} \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(ct, state)$

**return**  $\tilde{b} = b$

Figure 1.3: IK-CCA game.

$ct^6$ , we define the IK-CCA advantage of  $\mathcal{A}$  as follows:

$$\text{Adv}_{\text{Asym}, \mathcal{A}}^{\text{ik-cca}}(\lambda) = |2 \cdot \Pr[\text{IK-CCA}(\lambda) = 1] - 1|$$

We say that Asym is IK-CCA secure if for every PPT adversary  $\mathcal{A}$  the advantage  $\text{Adv}_{\text{Asym}, \mathcal{A}}^{\text{ik-cca}}(\lambda)$  is a negligible function.

### 1.5.6 Block cipher-based compression functions

**Definition 1.5.11.** Let  $kl, il > 1$ . A *block cipher* is a map  $E: \{0, 1\}^{kl} \times \{0, 1\}^{il} \rightarrow \{0, 1\}^{il}$  where, for each key  $k \in \{0, 1\}^{kl}$ , the function  $E_k(\cdot) = E(k, \cdot)$  is a permutation on  $\{0, 1\}^{il}$ . If  $E$  is a block cipher then  $E^{-1}$  is its inverse, that on input  $(k, y)$  returns  $m$  such that  $E_k(m) = y$ .

Let  $\mathcal{BCK}(kl, il)$  be the set of all block ciphers  $E: \{0, 1\}^{kl} \times \{0, 1\}^{il} \rightarrow \{0, 1\}^{il}$ . In order to analyse the security properties of block cipher-based cryptographic constructions it is common to use a security model denoted *the ideal cipher model (ICM)*. Informally speaking, in ICM attackers are allowed to query an oracle simulating a random block cipher, but have no information about the oracle's internal structure. We formalize this notion in the following definition:

**Definition 1.5.12** (Ideal Cipher Model [HKT11]). The Ideal Cipher Model (ICM), is a security model where all parties are granted access to an ideal cipher  $E: \{0, 1\}^{kl} \times \{0, 1\}^{il} \rightarrow \{0, 1\}^{il}$ , a random primitive such that  $E(k, \cdot)$  for  $k \in \{0, 1\}^{kl}$  are  $2^{kl}$  independent random permutations.

For fixed  $kl$  and  $il$ , each party is given access to the oracles  $\mathcal{O}^E$  and  $\mathcal{O}^{E^{-1}}$ , simulating  $E$  and  $E^{-1}$ , which can be queried for encryption and decryption a polynomial number of times. The encryption oracle takes as input a key,  $k \in \{0, 1\}^{kl}$ , and a preimage,  $m \in \{0, 1\}^{il}$ , and returns a tuple comprising the image,  $y \in \{0, 1\}^{il}$ , along with the inputs,  $k$  and  $m$ . If  $(k, m)$  is queried for the first time, the image  $y$  is taken uniformly

---

<sup>6</sup>*state* is some state information that the adversary outputs after the choice of the message to encrypt. It can be some preprocessed information that can be helpful to win the game



$O^E(k, m)$	$O^{E^{-1}}(k, y)$
<b>if</b> $(k, m, \cdot) \notin \text{Table}_O$ $y \leftarrow \$ \{0, 1\}^{\text{il}}$ $\text{Table}_O.\text{append}(k, m, y)$	<b>if</b> $(k, \cdot, y) \notin \text{Table}_O$ $m \leftarrow \$ \{0, 1\}^{\text{il}}$ $\text{Table}_O.\text{append}(k, m, y)$
<b>else</b> $y \leftarrow \text{Table}_O(k, m)$	<b>else</b> $m \leftarrow \text{Table}_O(k, y)$
<b>return</b> $(k, m, y)$	<b>return</b> $(k, m, y)$

Figure 1.4: Oracles of an ideal block cipher, with  $\text{Table}_O$  being a table of tuples (key, preimage, image) of queries already answered by the oracle.

at random from  $\{0, 1\}^{\text{il}}$  and added to the oracle's table. Otherwise, the oracle returns  $y$  associated with query  $(k, m)$  in its table. The decryption oracle is defined similarly with the image and key defined as inputs and the preimage chosen randomly, for details see Fig. 1.4.

**Definition 1.5.13** (Block cipher-based compression function [BRS02]). A *block cipher-based compression function* is a map  $f$  such that

$$f: \mathcal{BK}(\text{kl}, \text{il}) \times \{0, 1\}^a \times \{0, 1\}^b \rightarrow \{0, 1\}^c$$

where  $\text{kl}, \text{il}, a, b, c > 1$  and  $a + b > c$ . The function  $f$ , given  $m \in \{0, 1\}^a \times \{0, 1\}^b$ , computes  $f(E, m)$  using an  $E$ -oracle.

**Remark 1.5.14.** We use  $f_E$  to denote a block cipher-based compression function  $f$  restricted to a given block cipher  $E$ , i.e.  $f_E: \{0, 1\}^a \times \{0, 1\}^b \rightarrow \{0, 1\}^c$  and  $f_E = f(E, \cdot)$ , for  $a, b, c$  as given in the definition above.

Let  $f$  be a compression function based on a block cipher. Fix a constant  $h_0 \in \{0, 1\}^c$  and an adversary  $\mathcal{A}$ . We define the advantage in finding a collision in  $f$  as

$$\text{Adv}_{f, \mathcal{A}}^{\text{coll}} = \Pr \left[ E \leftarrow \$ \mathcal{BK}(\text{kl}, \text{il}); ((k, m), (k', m')) \leftarrow \mathcal{A}^{O^E, O^{E^{-1}}}(f_E, h_0) : \begin{aligned} & ((k, m) \neq (k', m') \wedge f_E(k, m) = f_E(k', m')) \vee f_E(k, m) = h_0 \end{aligned} \right].$$

The previous definition gives credit for finding an  $(k, m)$  such that  $f_E(k, m) = h_0$  for a fixed  $h_0 \in \{0, 1\}^c$ .

### 1.5.7 Hash functions

**Definition 1.5.15** (Hash function, [KL14, Definition 4.9]). A hash function  $\mathcal{H}$  is a pair of algorithms  $(\text{Setup}, H)$  fulfilling the following properties:

- **Setup** is a PPT algorithm which takes as input a security parameter  $1^\lambda$  and outputs a key  $hk$ . We assume that  $1^\lambda$  is included in  $hk$ .

690 •  $H$  is (deterministic) polynomial-time algorithm that takes as input a key  $hk$  and  
 691 any string  $x \in \{0, 1\}^*$ , and outputs a string  $H(hk, x) = H_{hk}(x) \in \{0, 1\}^{hLen}$ , where  
 692  $hLen$  is a polynomial in  $\lambda$ .<sup>7</sup>

693 If for every  $\lambda$  and  $hk$ ,  $H_{hk}$  is defined only over inputs of length  $hInpLen(\lambda)$  and  $hInpLen(\lambda) >$   
 694  $hLen(\lambda)$ , then we say that  $\mathcal{H}$  is a *fixed-length hash function* with length parameter  
 695  $hInpLen$ . Note that  $hInpLen(\lambda)$  is a polynomial in  $\lambda$ .

696 Informally, for a given function  $f$  we say that  $(x, y)$  is a *collision* if  $f(x) = f(y)$  and  
 697  $x \neq y$ . In the following, we formalize this notion for a hash function  $\mathcal{H}$ .

**Definition 1.5.16** (Collision Resistance [KL14, Definitions 4.10]). A hash function  $\mathcal{H} = (\text{Setup}, H)$  is collision resistant if for all PPT adversaries  $\mathcal{A}$  there exists a negligible function  $\text{negl}(\lambda)$  such that:

$$\text{Adv}_{\mathcal{H}, \mathcal{A}}^{\text{cr}}(\lambda) = \Pr \left[ hk \leftarrow \text{Setup}(1^\lambda), (x, y) \leftarrow \mathcal{A}(hk) : x \neq y \wedge H_{hk}(x) = H_{hk}(y) \right] \leq \text{negl}(\lambda).$$

## 698 HDHI and HDHI2 assumptions

699 The Hash Diffie-Hellmann Independence (HDHI) assumption states that, given  $H$  in  $\mathcal{H}$   
 700 and a group description  $(p, \mathbb{G}, \mathbf{g}, \otimes)$ , for  $\llbracket u \rrbracket$  and  $\llbracket v \rrbracket$ , with  $u, v$  sampled at random, it  
 701 is hard for an attacker to distinguish  $H(\llbracket u \rrbracket \parallel \llbracket uv \rrbracket)$  from a random string of the same  
 702 size.<sup>8</sup> This is formalized in Definition 1.5.17, where an attacker can also access an oracle  
 703  $\mathcal{O}^{\text{HDHI}_v}$  that on input  $\mathbf{r} \in \mathbb{G}$  returns  $H(\mathbf{r} \parallel v \cdot \mathbf{r})$  (queries on  $\llbracket u \rrbracket$  are forbidden).<sup>9</sup> In other  
 704 words, the HDHI assumption measures the sense in which  $H$  is “independent” of the  
 705 underlying Diffie-Hellman problem.

**Definition 1.5.17** (HDHI, [ABR99, Definition 7]). Let  $\mathcal{H}$  be a hash function,  $\text{SetupG}$  be a group generation algorithm and  $\mathcal{A}$  be an adversary. Consider the HDHI game described in Figure 1.5. We define the advantage of  $\mathcal{A}$  in violating the HDHI assumption as:

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi}}(\lambda) = |2 \cdot \Pr[\text{HDHI}(\lambda) = 1] - 1|.$$

706 Note that the above definition corresponds to [ABR99, Section 3.2.1, Definition 3].  
 707 In the following, we introduce a similar notion denoted as HDHI2 (this is an adapta-  
 708 tion of the ODH2 notion in [ABN10, Section 6]) which will be useful in the IK-CCA  
 709 proof Section 3.5.4.

**Definition 1.5.18** (HDHI2). Let  $\mathcal{H}$  be a hash function,  $\text{SetupG}$  a group generation algorithm and let  $\mathcal{A}$  be an adversary. Consider the HDHI2 game described in Figure 1.6. We define the advantage of  $\mathcal{A}$  in violating the HDHI2 assumption as:

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi2}}(\lambda) = |2 \cdot \Pr[\text{HDHI2}(\lambda) = 1] - 1|.$$

<sup>7</sup>For simplicity, we may denote  $hLen(\lambda)$  as  $hLen$ .

<sup>8</sup>Note that  $H$  takes as inputs bit strings, so technically we should make use of an encoding function from  $\mathbb{G}$  to  $\{0, 1\}^{gLen}$  but we may omit this step through the document to improve readability.

<sup>9</sup>In [ABR99, Section 3.2.1] this notion is denoted as adaptive HDH independence assumption. Since we only introduce the adaptive version we denote it as HDHI.

---

$\text{HDHI}(\lambda)$   
 $hk \leftarrow \mathcal{H}.\text{Setup}(1^\lambda)$   
 $(q, \mathbb{G}, \mathfrak{g}, \otimes) \leftarrow \text{SetupG}(1^\lambda)$   
 $u, v \leftarrow \$[q]$   
 $w_0 \leftarrow \mathcal{H}.H_{hk}(\llbracket u \rrbracket \parallel \llbracket uv \rrbracket)$   
 $w_1 \leftarrow \$\{0, 1\}^{hLen}$   
 $b \leftarrow \$\{0, 1\}$   
 $\tilde{b} \leftarrow \mathcal{A}^{\text{O}^{\text{HDHI}_v}}(\llbracket u \rrbracket, \llbracket v \rrbracket, w_b)$   
**return**  $\tilde{b} = b$

Figure 1.5: HDHI game.

---

$\text{HDHI2}(\lambda)$   
 $hk \leftarrow \mathcal{H}.\text{Setup}(1^\lambda)$   
 $(q, \mathbb{G}, \mathfrak{g}, \otimes) \leftarrow \text{SetupG}(1^\lambda)$   
 $u, v_0, v_1 \leftarrow \$[q]$   
 $w_{0,0} \leftarrow \mathcal{H}.H_{hk}(\llbracket u \rrbracket \parallel \llbracket uv_0 \rrbracket), w_{0,1} \leftarrow \mathcal{H}.H_{hk}(\llbracket u \rrbracket \parallel \llbracket uv_1 \rrbracket)$   
 $w_{1,0} \leftarrow \$\{0, 1\}^{hLen}, w_{1,1} \leftarrow \$\{0, 1\}^{hLen}$   
 $b \leftarrow \$\{0, 1\}$   
 $\tilde{b} \leftarrow \mathcal{A}^{\text{O}^{\text{HDHI}_{v_0}}, \text{O}^{\text{HDHI}_{v_1}}}}(\llbracket u \rrbracket, \llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket, w_{b,0}, w_{b,1})$   
**return**  $\tilde{b} = b$

Figure 1.6: HDHI2 game.

**Lemma 1.5.1.** *Let  $\mathcal{A}$  be an adversary with advantage  $\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi2}}$  in solving the HDHI2 problem. Then there exists an adversary  $\mathcal{B}$  such that*

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi2}}(\lambda) \leq 2 \cdot \text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{B}}^{\text{hdhi}}(\lambda).$$

*Proof.* We reuse the proof described in [ABN10, Lemma 6.1] by applying minor modifications. In fact, HDHI and HDHI2 are, respectively, slightly different from ODH and ODH2 notions: in the related security games, if  $b = 0$  the challenges are constructed as  $H(\llbracket u \rrbracket \parallel \llbracket uv \rrbracket)$  and  $\{H(\llbracket u \rrbracket \parallel \llbracket uv_0 \rrbracket), H(\llbracket u \rrbracket \parallel \llbracket uv_1 \rrbracket)\}$  instead of  $H(\llbracket uv \rrbracket)$  and  $\{H(\llbracket uv_0 \rrbracket), H(\llbracket uv_1 \rrbracket)\}$ . By accordingly changing the instances of  $H$  in the games  $G_0, G_1, G_2$  of [ABN10, Lemma 6.1] our lemma follows.  $\square$

## 1.5.8 Pseudo Random Functions

Informally, a pseudorandom function family  $\mathcal{PRF} = \{\text{PRF}_k : D \rightarrow C\}_{k \in \mathcal{K}}$  is a collection of functions such that for a randomly chosen  $k \in \mathcal{K}$ , the function  $\text{PRF}_k$  is indistinguishable from a random function that maps  $D$  to  $C$ .

**Definition 1.5.19** (PRF Family [KL14, Definition 3.23]). Let  $\mathcal{F} : \{0, 1\}^* \times \{0, 1\}^* \rightarrow \{0, 1\}^*$  be an efficient, length-preserving, keyed function. We say  $\mathcal{F}$  is a pseudo random function if for all probabilistic polynomial-time distinguishers  $\text{Dist}$ , there exists a negligible function  $\text{negl}$  such that:

$$\text{Adv}_{\mathcal{F}, \text{Dist}}^{\text{prf}}(\lambda) = \left| \Pr[\text{Dist}^{\mathcal{F}_k(\cdot)}(1^\lambda) = 1] - \Pr[\text{Dist}^{f_\lambda(\cdot)}(1^\lambda) = 1] \right| \leq \text{negl}(\lambda),$$

where  $k \leftarrow \$\mathcal{K} = \{0, 1\}^\lambda$  is chosen uniformly at random and  $f_\lambda$  is chosen uniformly at random from the set of functions mapping  $\lambda$ -bit strings to  $\lambda$ -bit strings.

## 1.5.9 Commitment scheme

**Definition 1.5.20** (Non-interactive commitment scheme [BCC<sup>+</sup>15, Section 2.1]). A non-interactive commitment scheme  $\text{ComSch}$  is defined by the following algorithms:

- 725 • **Setup**, is a PPT algorithm that takes a security parameter  $1^\lambda$  and outputs public  
726 parameters  $pp$ .
- 727 • **Com**, is a polynomial-time algorithm that takes a message  $m \in \mathbb{B}^l$ , a random coin  
728  $r \in \mathbb{B}^n$  and outputs a commitment  $cm \in \mathbb{B}^l$ .

729 We assume that  $pp$  is implicitly passed to **Com**.

**Definition 1.5.21** (Computational hiding). We say that a commitment scheme is computationally hiding if for all PPT adversary  $\mathcal{A}$ , the advantage:

$$\left| \Pr \left[ pp \leftarrow \text{Setup}(1^\lambda), (m_0, m_1) \leftarrow \mathcal{A}(pp), b \leftarrow \{0, 1\}, \right. \right. \\ \left. \left. r \leftarrow \mathbb{B}^n, cm \leftarrow \text{Com}(m_b; r), \tilde{b} \leftarrow \mathcal{A}(cm), b = \tilde{b} \right] - \frac{1}{2} \right|$$

730 is at most negligible in  $\lambda$ .

**Definition 1.5.22** (Computational binding). We say that a commitment scheme is computationally binding if for all PPT adversary  $\mathcal{A}$ , the advantage:

$$\Pr \left[ pp \leftarrow \text{Setup}(1^\lambda), (m_0, r_0, m_1, r_1) \leftarrow \mathcal{A}(pp) \right. \\ \left. m_0 \neq m_1 \wedge \text{Com}(m_0; r_0) = \text{Com}(m_1; r_1) \right]$$

731 is at most negligible in  $\lambda$ .

732 Note that the previous definitions can be made *statistical* if we consider unbounded  
733 attackers  $\mathcal{A}$ .

### 734 1.5.10 Digital Signature

735 **Definition 1.5.23** (Digital signature [KL14, Definition 12.1]). A digital signature scheme  
736 **SigSch** is defined by the tuple of functions **SigSch** = (**KGen**, **Sig**, **Vf**),

- 737 •  $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$ . Key Generation randomized algorithm takes as input the  
738 security parameter  $1^\lambda$  and returns a signing key  $sk$  and verifying key  $vk$ .
- 739 •  $\sigma \leftarrow \text{Sig}(sk, m)$ . Given a signing key  $sk$  and a message  $m$ , the **Sig** algorithm  
740 computes and outputs a signature  $\sigma$ .
- 741 •  $\{0, 1\} \leftarrow \text{Vf}(vk, m, \sigma)$ . Given a verification key  $vk$ , a message  $m$  and a signature  
742  $\sigma$ , the **Vf** algorithm returns 1 if  $\sigma$  is a valid signature else 0.

743 A signature scheme must satisfy the *correctness property* (i.e  $\text{Vf}(vk, m, \text{Sig}(sk, m)) =$   
744 **true**, where  $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$ ) and be *unforgeable* (i.e. it is intractable to produce a  
745 signature, without knowing the signing key  $sk$ , on a message that has not been signed  
746 yet). In addition to these properties, certain digital signature schemes have an additional  
747 property called *one-timeness*, also defined below.

### UF-CMA( $1^\lambda, t, q$ )

---

```

1 :   $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$ 
2 :   $state \leftarrow \mathcal{A}^{\text{OSig}_{sk}}(vk, \cdot)$ 
3 :  //  $state = \{(m_i, \sigma_i)\}_{i \in [q]}$  where  $m_i$  denotes
4 :  // the  $i$ th query made to  $\text{OSig}_{sk}$  and
5 :  //  $\sigma_i$  denotes the  $i$ th oracle answers
6 :   $(m^*, \sigma^*) \leftarrow \mathcal{A}(state)$ 
7 :  return  $\text{Vf}(vk, m^*, \sigma^*) = 1$ 
8 :   $\wedge m^* \notin \{m_i\}_{i \in [q]}$ 

```

Figure 1.7: UF-CMA game

### SUF-CMA( $1^\lambda, t, q$ )

---

```

1 :   $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$ 
2 :   $state \leftarrow \mathcal{A}^{\text{OSig}_{sk}}(vk, \cdot)$ 
3 :  //  $state = \{(m_i, \sigma_i)\}_{i \in [q]}$  where  $m_i$  denotes
4 :  // the  $i$ th query made to  $\text{OSig}_{sk}$  and
5 :  //  $\sigma_i$  denotes the  $i$ th oracle answers
6 :   $(m^*, \sigma^*) \leftarrow \mathcal{A}(state)$ 
7 :  return  $\text{Vf}(vk, m^*, \sigma^*) = 1$ 
8 :   $\wedge (m^*, \sigma^*) \notin \{(m_i, \sigma_i)\}_{i \in [q]}$ 

```

Figure 1.8: SUF-CMA game

**Definition 1.5.24** (Unforgeability (UF-CMA) [KL14, Definition 12.2]). A digital signature scheme  $\text{SigSch}$  is UF-CMA if for any PPT adversary  $\mathcal{A}$ , the probability that  $\mathcal{A}$  wins the UF-CMA game, depicted in Fig. 1.7, is negligible.

**Definition 1.5.25** (Strong Unforgeability (SUF-CMA)). A digital signature scheme  $\text{SigSch}$  is SUF-CMA if the probability that any PPT adversary  $\mathcal{A}$  wins the SUF-CMA game, depicted in Fig. 1.8, is negligible.

**Definition 1.5.26** (One-Time (OT) Signature [KL14, Definition 12.6]). A *one-time* signature scheme is a digital signature scheme that uses each key-pair at most once.

**Remark 1.5.27.** It is worth noting that users may use one-time signing keys to sign multiple messages. In this case no security claims can be made.

## 1.5.11 Message Authentication Code

A message authentication code is a scheme that enables users to tag data for the purpose of authenticity and integrity. Formally:

**Definition 1.5.28** (Message Authentication Code, [KL14, Definition 4.1]). A message authentication code MAC is given by a tuple of PPT algorithms  $(\text{KGen}, \text{Tag}, \text{Vf})$  where:

- **KGen**, the key generation algorithm, takes a security parameter  $1^\lambda$ , and returns a key  $mk \in \{0, 1\}^{mLen(\lambda)}$ .<sup>10</sup>
- **Tag**, the tag generation algorithm, takes a key  $mk$  and a message  $y \in \{0, 1\}^*$  and returns a string  $\tau \in \{0, 1\}^*$ , called *tag*.
- **Vf**, the tag verification algorithm, takes a key  $mk$ , a message  $y \in \{0, 1\}^*$  and a tag  $\tau \in \{0, 1\}^*$ . It returns a value in  $\{0, 1\}$  where: 0 denotes that the message was rejected (i.e. deemed unauthentic) and 1 denotes that the message was accepted (i.e. deemed authentic).

<sup>10</sup>For simplicity, we may denote  $mLen(\lambda)$  as  $mLen$ .

---

SUF-CMA ( $\lambda$ )

---

$mk \leftarrow \text{KGen}(1^\lambda)$   
 $(\bar{y}, \bar{\tau}) \leftarrow \mathcal{A}^{\text{O}^{\text{Tag}_{mk}}, \text{O}^{\text{Vf}_{mk}}}$   
**return**  $\text{Vf}(mk, \bar{y}, \bar{\tau}) = 1$

Figure 1.9: SUF-CMA game.

771 We require that for all  $mk \in \{0, 1\}^\lambda$  and  $y \in \{0, 1\}^*$  we have  $\text{Vf}(mk, y, \text{Tag}(mk, y)) = 1$ .  
772 If  $\text{Tag}(mk, \cdot)$  is defined only over messages of length  $l(\lambda)$  and  $\text{Vf}(mk, y, \tau)$  outputs 0 for  
773 every  $y$  that is not of length  $l(\lambda)$ , then we say that  $(\text{KGen}, \text{Tag}, \text{Vf})$  is a *fixed-length MAC*  
774 with length parameter  $l(\lambda)$ .

775 A security notion for MAC follows:

776 **Definition 1.5.29** (SUF-CMA, [ABR99, Section 3.2.3]). Let  $\text{MAC} = (\text{KGen}, \text{Tag}, \text{Vf})$  be  
777 a message authentication scheme and let  $\mathcal{A}$  be an adversary. Consider the SUF-CMA  
778 game described in Figure 1.9, with the condition that  $\text{Tag}(mk, \bar{y}) \neq \bar{\tau}$ . We say that an  
779 adversary  $\mathcal{A}$  has *forged* a tag when it outputs a pair  $(\bar{y}, \bar{\tau})$  such that  $\text{Vf}_k(\bar{y}, \bar{\tau}) = 1$ , where  
780  $(\bar{y}, \bar{\tau})$  was not previously obtained via a query to the tag oracle.

We define the SUF-CMA advantage of  $\mathcal{A}$  as follows:

$$\text{Adv}_{\text{MAC}, \mathcal{A}}^{\text{suf-cma}}(\lambda) = \Pr[\text{SUF-CMA}(\lambda) = 1]$$

781 We say that MAC is SUF-CMA secure if for every PPT adversary  $\mathcal{A}$  the advantage  
782  $\text{Adv}_{\text{MAC}, \mathcal{A}}^{\text{suf-cma}}(\lambda)$  is a negligible function.

## Chapter 2

# Zeth protocol

In this section, we detail the **Zeth** protocol and provide a set of requirements that need to be respected to guarantee the security of the protocol.

## 2.1 Zeth Data Types

We begin by describing, and giving intuition about, the data types (see Section 1.1) used in **Zeth**. We follow some design rationale from **ZeroCash** [BSCG<sup>+</sup>14], and **Zcash** [ZCa19] in order to prevent the transaction malleability attack, and the Faerie Gold attack [ZCa19, Section 8.4]. We refer the reader to Appendix A for more details.

In what follows **Curve** represents a curve with scalar field  $\mathbb{F}_r$ , satisfying the requirements of Section 3.6. The specification is described in terms of this generic curve, with examples and notes relating to specific instances of interest (namely BN-254 and BLS12-377, see Chapter 3).

**ZethNoteDType** Represents a note in **Zeth**. This data type consists of the note owner's public address  $apk$ , identifier  $\rho$ , randomness  $r$  and value  $v$ .

Field	Description	Data type
$apk$	Note owner's paying key	$\mathbb{B}^{\text{PRFADDRROUTLEN}}$
$r$	Note randomness	$\mathbb{B}^{\text{RTRAPLEN}}$
$v$	Note value	$\mathbb{B}^{\text{ZVALUELEN}}$
$\rho$	Note identifier	$\mathbb{B}^{\text{PRFRHOOUTLEN}}$

Table 2.1: **ZethNoteDType** data type

**JSInputDType** Denotes a joinsplit input. It comprises the opening of a commitment  $cm$  which is in the set of leaves in the Merkle tree of **Mixer** (i.e. a *ZethNote*), its address  $mkaddr$  and authentication path  $mkpath$  on the contract's Merkle tree as well as the spending key  $ask$  of the note holder and the note nullifier  $nf$ .

Field	Description	Data type
<i>mkpath</i>	Merkle authentication path to the commitment corresponding to the <i>ZethNote</i> to spend	$(\mathbb{F}_r)^{\text{MKDEPTH}}$
<i>mkaddr</i>	Commitment address in the Merkle tree	$\mathbb{B}^{\text{MKDEPTH}}$
<i>znote</i>	<b>Zeth</b> note object	<b>ZethNoteDType</b>
<i>cm</i>	<b>Zeth</b> note commitment	$\mathbb{F}_r$
<i>ask</i>	Note owner's spending key	$\mathbb{B}^{\text{ASKLEN}}$
<i>nf</i>	Note nullifier	$\mathbb{B}^{\text{PRNFOUTLEN}}$

Table 2.2: JSInputDType data type

802 **PrimInputDType** Represents the primary inputs used to generate the zk-SNARK proof  
803  $\pi$ . *prim* is a tuple defined as the current Merkle root *mkroot* of the Merkle tree  
804 maintained by **Mixer**, the input notes nullifiers *nfs* =  $(nf_0, \dots, nf_{\text{JSIN}-1})$ , the  
805 output notes commitments *cms* =  $(cm_0, \dots, cm_{\text{JSOUT}-1})$ , the signature hash *hsig*,  
806 the message authentication tags *htags* =  $(h_0, \dots, h_{\text{JSIN}-1})$  and the residual bits  
807 field *rsd*, which aggregates the former's fields bits which could not be contained in  
808 a field element.

Field	Description	Data type
<i>mkroot</i>	Merkle root of the Merkle tree	$\mathbb{F}_r$
<i>nfs</i>	Indexed set of nullifiers of the “old” notes to spend (see Section 3.3.1 for definition of <b>NFFLEN</b> )	$((\mathbb{F}_r)^{\text{NFFLEN}})^{\text{JSIN}}$
<i>cms</i>	Indexed set of commitments to the newly created notes	$(\mathbb{F}_r)^{\text{JSOUT}}$
<i>hsig</i>	Signature hash (non-malleability, see Appendix A and Section 3.3.1 for definition of <b>HSIGFLEN</b> )	$(\mathbb{F}_r)^{\text{HSIGFLEN}}$
<i>htags</i>	Indexed set of message authentication tags (non-malleability, see Appendix A and Section 3.3.1 for definition of <b>HFLEN</b> )	$((\mathbb{F}_r)^{\text{HFLEN}})^{\text{JSIN}}$
<i>rsd</i>	Residual bits corresponding to unpacked bits of former fields (see Section 3.3.1 for definition of <b>RSDFLEN</b> )	$(\mathbb{F}_r)^{\text{RSDFLEN}}$

Table 2.3: PrimInputDType data type



809 **AuxInputDType** Represents the auxiliary inputs used to generate the zk-SNARK proof  
810  $\pi$ . *aux* is a tuple defined as joinsplit inputs (i.e. “old outputs to be spent”), the new  
811 *ZethNotes*, the joinsplit’s randomness  $\phi$  as well the public values *vin* and *vout*, the  
812 signature hash *hsig* and the message authentication tags *htags* =  $(h_0, \dots, h_{\text{JSIN}-1})$ .

Field	Description	Data type
<i>jsins</i>	Indexed set of JSIN joinsplit inputs	<b>JSInputDType</b> <sup>JSIN</sup>
<i>znotes</i>	Indexed set of JSOUT newly created notes	<b>ZethNoteDType</b> <sup>JSOUT</sup>
$\phi$	The joinsplit randomness (non-malleability, see Appendix A)	$\mathbb{B}^{\text{PHILEN}}$
<i>vin</i>	Public input value to the joinsplit	$\mathbb{B}^{\text{ZVALUELEN}}$
<i>vout</i>	Public output value to the joinsplit	$\mathbb{B}^{\text{ZVALUELEN}}$
<i>hsig</i>	Signature hash (non-malleability, see Appendix A)	$\mathbb{B}^{\text{CRHHSIGOUTLEN}}$
<i>htags</i>	Indexed set of message authentication tags (non-malleability, see Appendix A)	$(\mathbb{B}^{\text{PRFPKOUTLEN}})^{\text{JSIN}}$

Table 2.4: **AuxInputDType** data type

813 **MixInputDType** Represents the set of inputs to the Mix function of **Mixer**. The input of  
814 the Mix function is a tuple defined as the primary inputs *prim*, the zk-proof  $\pi$ , the  
815 ciphertexts of the newly created notes *ciphers* =  $(ct_0, \dots, ct_{\text{JSOUT}-1})$ , a one-time  
816 signature  $\sigma$  and the associated verification key *vk*.

Field	Description	Data type
<i>primIn</i>	Primary input object associated with the zk-proof $\pi$	<b>PrimInputDType</b>
<i>proof</i>	The zk-SNARK associated to the <b>Zeth</b> statement (see Section 2.2)	<b>ZKPDType</b> (see Section 3.6)
<i>otssig</i>	The one-time signature used to prevent transaction malleability (see Appendix A)	<b>SigOtsDType</b> (see Section 3.4.2)
<i>otsvk</i>	The verification key associated with the signature <i>otssig</i> used to prevent transaction malleability (see Appendix A)	<b>VKOtsDType</b> (see Section 3.4.2)
<i>ciphers</i>	Indexed set of ciphertexts of the newly generated notes	$(\mathbb{B}^{\text{ENCZETHNOTELEN}})^{\text{JSOUT}}$ (see Section 3.5)

Table 2.5: **MixInputDType** data type

817 **MixEventDType** Represents the data emitted as an **Ethereum** event (Section 1.2.3) dur-  
 818 ing a successful execution of the Mix function of **Mixer**. Clients are required to  
 819 read this data and use it to update their representation of **Mixer**'s state.

Field	Description	Data type
<i>mkroot</i>	New root of Merkle tree of commitments	$\mathbb{F}_r$
<i>nfs</i>	Nullifiers for input notes consumed	$(\mathbb{B}^{\text{PRNFOUTLEN}})^{\text{JSIN}}$
<i>cms</i>	Commitments to the output notes	$(\mathbb{F}_r)^{\text{JSOUT}}$
<i>ciphers</i>	Ciphertexts for the output notes	$(\mathbb{B}^{\text{ENCZETHNOTELEN}})^{\text{JSOUT}}$

Table 2.6: **MixEventDType** data type

## 820 2.2 Zeth statement

821 As explained in [RZ19], the Mix function of **Mixer** verifies the validity of  $\pi$  on the  
 822 given primary inputs in order to determine whether the state transition is valid. As  
 823 such, **Mixer** verifies whether for  $\pi$ , and primary input *prim*, there exists an auxiliary  
 824 input *aux*, such that the tuple  $(\text{prim}, \text{aux})$  satisfies the NP-relation  $\mathbf{R}^z$ , consisting of the  
 825 following constraints:

- 826 • For each  $i \in [\text{JSIN}]$ :
  - 827 1.  $\text{aux.jsins}[i].\text{znote.apk} = \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{addr}}(0)$
  - 828 2.  $\text{aux.jsins}[i].\text{cm} = \text{ComSch.Com}(\text{aux.jsins}[i].\text{znote.apk}, \text{aux.jsins}[i].\text{znote}.\rho, \text{aux.jsins}[i].\text{znote}.v;$   
 829  $\text{aux.jsins}[i].\text{znote}.r)$
  - 830 3.  $\text{aux.jsins}[i].\text{nf} = \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{nf}}(\text{aux.jsins}[i].\text{znote}.\rho)$
  - 831 4.  $\text{aux.htags}[i] = \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{pk}}(i, \text{aux.hsigs})$  (non-malleability, see Appendix A)
  - 832 5.  $(\text{aux.jsins}[i].\text{znote}.v) \cdot (1 - e) = 0$  is satisfied for the boolean value  $e$  set such  
 833 that if  $\text{aux.jsins}[i].\text{znote}.v > 0$  then  $e = 1$ .
  - 834 6. The Merkle root  $\text{mkroot}'$  obtained after checking the Merkle authentica-  
 835 tion path  $\text{aux.jsins}[i].\text{mkpath}$  of commitment  $\text{aux.jsins}[i].\text{cm}$ , with MKHASH,  
 836 equals to  $\text{prim.mkroot}$  if  $e = 1$ .
  - 837 7.  $\text{prim.nfs}[i]$   
 838  $= \{\text{Pack}_{\mathbb{F}_r}(\text{aux.jsins}[i].\text{nf}[k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}])\}_{k \in [\lfloor \text{PRNFOUTLEN}/\text{FIELD CAP} \rfloor]}$   
 839 (see Section 3.3.1 for definition of Pack)
  - 840 8.  $\text{prim.htags}[i]$   
 841  $= \{\text{Pack}_{\mathbb{F}_r}(\text{aux.htags}[i][k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}])\}_{k \in [\lfloor \text{PRFPKOUTLEN}/\text{FIELD CAP} \rfloor]}$   
 842 (see Section 3.3.1 for definition of Pack)
- 843 • For each  $j \in [\text{JSOUT}]$ :

- 844 1.  $aux.znotes[j].\rho = \text{PRF}_{aux.\phi}^{\text{rho}}(j, aux.hsigs)$  (non-malleability, see Appendix A)
- 845 2.  $prim.cms[j] = \text{ComSch.Com}(aux.znotes[j].apk, aux.znotes[j].\rho, aux.znotes[j].v;$
- 846  $aux.znotes[j].r)$
- 847 •  $prim.hsigs = \{\text{Pack}_{\mathbb{F}_r}(aux.hsigs[k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}])\}_{k \in [\lfloor \text{CRHHSIGOUTLEN}/\text{FIELD CAP} \rfloor]}$
- 848 (see Section 3.3.1 for definition of  $\text{Pack}$ )
- 849 •  $prim.rsd = \text{Pack}_{rsd}(\{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}, aux.vin, aux.vout, aux.hsigs, \{aux.htags[i]\}_{i \in [\text{JSIN}]})$
- 850 (see Section 3.3.1 for definition of  $\text{Pack}_{rsd}$ )
- Check that the “joinsplit is balanced”, i.e. check that the joinsplit equation holds:<sup>1</sup>

$$\begin{aligned}
& \text{Pack}_{\mathbb{F}_r}(aux.vin) + \sum_{i \in [\text{JSIN}]} \text{Pack}_{\mathbb{F}_r}(aux.jsins[i].znote.v) \\
&= \sum_{j \in [\text{JSOUT}]} \text{Pack}_{\mathbb{F}_r}(aux.znotes[j].v) + \text{Pack}_{\mathbb{F}_r}(aux.vout)
\end{aligned}$$

## 851 2.3 Generating the inputs of the Mix function ( $\text{Mix}_{in}$ )

852 In the following section, we assume that the system is initialized. In other words, we  
853 assume that a ledger  $L$  is available (i.e. an **Ethereum** network is operated by a set of  
854 miners), the **Mixer** contract is deployed on  $L$ . Likewise, we assume that the public  
855 parameters  $pp_{\text{ZkSnarkSch}} \leftarrow \text{ZkSnarkSch.KGen}(1^\lambda, \mathbf{R}^z)$  are available to **Mixer** and to all  
856 parties willing to call the Mix function of **Mixer**. Furthermore, we assume that there  
857 exists a set of **Ethereum** and **Zeth** users, and that the *payment address* of each **Zeth** user  
858 is easily discoverable. In the rest of this section, the set of *payment addresses* discovered  
859 by a zeth user  $\mathcal{U}_Z$  is represented as a list attribute  $\mathcal{U}_Z.\text{keystore}$  indexed by usernames.

860 In order for  $\mathcal{U}_Z$  to transact via **Zeth**,  $\mathcal{U}_Z$  needs to create an object  $\text{Mix}_{in}$  of type  
861 **MixInputDType** to pass to the Mix function of **Mixer**:

- 862 1. Create an object  $prim$  of type **PrimInputDType** to represent the primary input,
- 863 and an object  $aux$  of type **AuxInputDType** to represent the auxiliary input, where:
  - 864 (a)  $prim.mkroot \in \text{Roots}$ , where  $\text{Roots}$  is the set of *all* Merkle roots corresponding  
865 to one of the state of the Merkle tree on **Mixer** containing *all* the commit-  
866 ments to the input notes, in  $aux.jsins$ , in its set of leaves.
  - 867 (b)  $aux.znotes[j].r \leftarrow \mathbb{B}^{\text{RTRAPLEN}}, \forall j \in [\text{JSOUT}]$ , and  $aux.\phi \leftarrow \mathbb{B}^{\text{PHILEN}}$
  - 868 (c) The public values  $(aux.vin, aux.vout) \in (\mathbb{B}^{\text{ZVALUELEN}})^2$ ,  $aux.znotes[j].v$  and  
869  $aux.znotes[j].apk \forall j \in [\text{JSOUT}]$  are all set by the sender,  $\mathcal{U}_Z$ , as desired as  
870 long as they satisfy the joinsplit equation.

---

<sup>1</sup>where  $\text{Pack}_{\mathbb{F}_r}(x)$  outputs the numerical value of  $x$  in  $\mathbb{F}_r$ . We rely on the fact that  $\text{ZVALUELEN} < \text{FIELD CAP}$  to perform this sum.

- (d) All attributes of the *prim* and *aux* objects should be derived as specified in the statement (see Section 2.2), alongside a signature hash (*aux.hsig*) that is generated as the hash of the nullifiers and a one-time signing verification key (non-malleability, see Appendix A), using the desired signature scheme  $\text{SigSch}_{\text{OT-SIG}}$  (see Section 3.4):

$$(sk_{\text{OT-SIG}}, vk_{\text{OT-SIG}}) = \text{SigSch}_{\text{OT-SIG}}.\text{KGen}(1^\lambda) \quad (2.1)$$

$$aux.hsig = \text{CRH}^{\text{hsig}}(\{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}, vk_{\text{OT-SIG}}) \quad (2.2)$$

- 871 (e)  $\text{Mix}_{in}.primIn \leftarrow prim$

#### Note

If one of the attributes of *prim* and *aux* is not correctly generated, then the proof of computational integrity generated in the next step will be rejected on **Mixer**, and the state of **Mixer** will not be modified.

872

- 873 2. Generate a zk-SNARK proof  $\pi$  to prove, in zero-knowledge, that the relation  $\mathbf{R}^z$   
874 (Section 2.2) holds on the primary and auxiliary inputs, using the desired zk-  
875 SNARK scheme  $\text{ZkSnarkSch}$  (see Section 3.6):

876 (a)  $\pi \leftarrow \text{ZkSnarkSch}.\text{P}(pp_{\text{ZkSnarkSch}}, prim, aux)$

877 (b)  $\text{Mix}_{in}.proof \leftarrow \pi$

- 878 3. Encrypt all the *aux.znotes* using the recipient's *payment address*, using the en-  
879 cryptation scheme  $\text{EncSch}$  (see Section 3.5).

- (a) For all  $j \in [\text{JSOUT}]$ , do:

$$ct_j \leftarrow \text{EncSch}.\text{Enc}(aux.znotes[j], \mathcal{U}_Z.\text{keystore}[recipient_j].pub.pk_{enc})$$

880 (b)  $\text{Mix}_{in}.ciphers \leftarrow \{ct_j\}_{j \in [\text{JSOUT}]}$

- 881 4. Generate a signature  $\sigma_{\text{OT-SIG}}$  on the inputs of the **Mix** function, in order to prevent  
882 any malleability attacks (c.f. Appendix A), using the desired signature scheme  
883  $\text{SigSch}_{\text{OT-SIG}}$  (see Section 3.4):

- (a) Using the one-time signature keypair generated in Eq. (2.1), do:

$$\begin{aligned} dataToBeSigned &= \mathcal{S}_E.Addr \parallel \text{Mix}_{in}.primIn \parallel \text{Mix}_{in}.\pi \parallel \text{Mix}_{in}.ciphers \\ \sigma_{\text{OT-SIG}} &= \text{SigSch}_{\text{OT-SIG}}.\text{Sig}(sk_{\text{OT-SIG}}, \text{CRH}^{\text{ots}}(dataToBeSigned)) \end{aligned}$$

884 (b)  $\text{Mix}_{in}.otssig \leftarrow \sigma_{\text{OT-SIG}}$

885 (c)  $\text{Mix}_{in}.otsvk \leftarrow vk_{\text{OT-SIG}}$

886 Here,  $\mathcal{S}_E.Addr$  represents the address of the **Ethereum** user  $\mathcal{S}_E$  who must sign the  
887 transaction (see Section 2.4). In general, this is likely to be owned by the holder  
888  $\mathcal{U}_Z$  of the **Zeth** notes to be spent, but this is not a requirement.

## 889 2.4 Creating an Ethereum transaction $tx_{\text{Mix}}$ to call $\widetilde{\text{Mixer}}$

After generating a  $\text{Mix}_{in}$  object,  $\mathcal{U}_Z$  can generate an object  $tx_{raw}$  of type  $\text{TxRawDType}$ , such that:

$$tx_{raw}.to = \widetilde{\text{Mixer}}.Addr \wedge tx_{raw}.data = zdata$$

890 Then, an **Ethereum** user  $\mathcal{S}_E$  can ECDSA sign  $tx_{raw}$ , under  $\mathcal{S}_E.sk$  in order to transform  
891 this object of type  $\text{TxRawDType}$  into an finalized transaction, i.e. an object  $tx_{\text{Mix}}$  of type  
892  $\text{TxDType}$ .

893 Finally, the transaction  $tx_{\text{Mix}}$  is broadcasted on the **Ethereum** network and eventually  
894 gets mined.

### Note

Here, the **Ethereum** user  $\mathcal{S}_E$  who sends the final transaction, and the **Zeth** user  $\mathcal{U}_Z$  may represent the same person or entity, but this is not necessarily the case. It is perfectly feasible (and in some cases may be desirable) for a **Zeth** user  $\mathcal{U}_Z$  to create a **Zeth** transaction which is later signed by a distinct party  $\mathcal{S}_E$ . In particular, the only identifying information that appears in plaintext on the ledger will be that of  $\mathcal{S}_E$ .

895

## 896 2.5 Processing $tx_{\text{Mix}}$

897 When a  $tx_{\text{Mix}}$  is mined (hence assuming that  $\text{EthVerifyTx}(tx_{\text{Mix}})$  returns **true**), the state  
898 transition specified by the **Mix** function of  $\widetilde{\text{Mixer}}$  is executed.

899 To preserve the soundness of **Zeth**, and make sure that no  $\mathcal{U}_Z$  is able to create  
900 value by double spending *ZethNotes*, various checks need to be satisfied. The function  
901 **ZethVerifyTx** is defined as the function that returns **true** if all the checks are satisfied,  
902 and **false** otherwise.

903 If **ZethVerifyTx**( $tx_{\text{Mix}}$ ) returns **true**, then **Mix** modifies the “World state”  $\varsigma$  to account  
904 for the spent *ZethNotes* and the newly generated ones. However, if **ZethVerifyTx**( $tx_{\text{Mix}}$ )  
905 returns **false**, then the state transition ends.

### Note

Even if **ZethVerifyTx**( $tx_{\text{Mix}}$ ) returns **false**,  $\varsigma$  is modified since the **Ethereum** balances of the transaction originator is decremented by the sum of **DGAS** and the gas consumed by the **ZethVerifyTx** function, and the balance of the **Ethereum** account of the miner gets incremented by the same amount.

906

907 Thus, **Mix** proceeds as follows:

1. Check that all the values of the primary inputs’ ( $\text{Mix}_{in}.primIn$ ) entries are elements of the scalar field over which the zk-proof is generated:

$$\text{Mix}_{in}.primIn \in \mathbb{F}_{\mathbf{r}}^*$$

2. Unpack the nullifiers, signature hash and public values (see Section 3.3.1 for the definitions of the `Unpack` functions):

$$\begin{aligned}
nf_i &= \text{Unpack}_{nf}(\text{Mix}_{in}.primIn.nfs[i], \text{Mix}_{in}.primIn.rsd) \quad \forall i \in [\text{JSIN}] \\
vin &= \text{decode}_{\mathbb{N}}(\text{Unpack}_{vin}(), \text{Mix}_{in}.primIn.rsd) \\
vout &= \text{decode}_{\mathbb{N}}(\text{Unpack}_{vout}(), \text{Mix}_{in}.primIn.rsd) \\
hsig &= \text{Unpack}_{hsig}(\text{Mix}_{in}.primIn.hsig, \text{Mix}_{in}.primIn.rsd)
\end{aligned}$$

- 908 3. Check the validity of the  $tx_{\text{Mix}}$  object (`ZethVerifyTx`):

- Check that  $\text{Mix}_{in}.primIn.hsig$  is correctly computed, i.e. check that the following equation holds (to prevent transaction malleability, see Appendix A):

$$hsig = \text{CRH}^{hsig}(\text{Mix}_{in}.primIn.nfs, \text{Mix}_{in}.otsvk)$$

- Check that  $\pi$  is a valid zk-SNARK proof for  $\text{Mix}_{in}.primIn$ , i.e. check that:

$$\text{ZkSnarkSch.V}(pp_{\text{ZkSnarkSch}}, \pi, \text{Mix}_{in}.primIn) = \text{true}$$

- Check that none of the nullifiers in  $\text{Mix}_{in}.primIn.nfs$  have already been used, i.e. check that:

$$nf_i \notin \text{Nulls}, \forall i \in [\text{JSIN}]$$

909 where  $\text{Nulls}$  is the set of all nullifiers that are “declared” on  $\widetilde{\text{Mixer}}$ .

- Check that  $\text{Mix}_{in}.otssig$  is a valid signature of the **Ethereum** sender’s address  $\text{Addr}$  (see Section 2.4) and the attributes of  $\text{Mix}_{in}$ , to prevent transaction malleability (see Appendix A), i.e. check that:

$$\begin{aligned}
&\text{SigSch}_{\text{OT-SIG}}.\text{Vf}(\text{Mix}_{in}.otsvk, m, \text{Mix}_{in}.otssig) = \text{true} \\
&\text{where } m = \text{CRH}^{\text{ots}}(\text{Addr} \parallel \text{Mix}_{in}.primIn \parallel \text{Mix}_{in}.\pi \parallel \text{Mix}_{in}.ciphers)
\end{aligned}$$

- Check that  $\text{Mix}_{in}.primIn.mkroot$  corresponds to a valid state of the Merkle tree held on  $\widetilde{\text{Mixer}}$ , i.e. check that:

$$\text{Mix}_{in}.primIn.mkroot \in \text{Roots}'$$

910 where  $\text{Roots}'$  is the set of all Merkle roots corresponding to one of the states  
911 of the Merkle tree.

- Check that  $vin$  corresponds to the value  $val$  of the transaction object, i.e. check that:

$$vin = tx_{\text{Mix}}.val$$

- 912 4. If all checks above pass, i.e. if `ZethVerifyTx(txMix)` returns `true`, then the following  
913 additional modifications are made in  $\varsigma$ :

- 914 • Add the commitments  $\text{Mix}_{in}.primIn.cms$  to the Merkle tree held on  $\widetilde{\text{Mixer}}$ .
- 915 •  $Roots' \leftarrow Roots' \cup \{mkroot'\}$ , where  $mkroot'$  is the Merkle root of the Merkle
- 916 tree after insertion of the commitments  $\text{Mix}_{in}.primIn.cms$  in the Merkle tree.
- 917 •  $Nulls \leftarrow Nulls \cup \{nf_i\}_{i \in [JSIN]}$ , i.e. the nullifiers  $nfs$  become “declared”.
- 918 • Modify the **Ethereum** balances according to the public values:
  - 919 –  $\varsigma[\mathcal{S}_{\mathcal{E}}.Addr].bal = \varsigma[\mathcal{S}_{\mathcal{E}}.Addr].bal - vin$
  - 920 –  $\varsigma[\mathcal{S}_{\mathcal{E}}.Addr].bal = \varsigma[\mathcal{S}_{\mathcal{E}}.Addr].bal + vout$
  - 921 –  $\widetilde{\text{Mixer}}.bal = \widetilde{\text{Mixer}}.bal + vin$
  - 922 –  $\widetilde{\text{Mixer}}.bal = \widetilde{\text{Mixer}}.bal - vout$
- 923 • Emit an event (Section 1.2.3)  $evMixOut$  of type **MixEventDType**, contain-
- 924 ing the new root  $mkroot'$  of the Merkle tree of commitments, the nullifiers
- 925  $\{nf_i\}_{i \in [JSIN]}$ , commitments to the newly created *ZethNotes*  $\text{Mix}_{in}.primIn.cms$ ,
- 926 and the corresponding ciphertexts  $\text{Mix}_{in}.primIn.ciphers$ .

## 927 2.6 Receiving *ZethNotes*

928 In order to confirm the reception of *ZethNotes*,  $\mathcal{R}_Z$  must listen to the events (Sec-  
 929 tion 1.2.3) of type **MixEventDType** emitted by the processing of  $tx_{\text{Mix}}$ , and try to decrypt  
 930 the ciphertexts using  $\mathcal{R}_Z.priv.skenc$  to see if he is the recipient of a **Zeth** payment. If  
 931 the decryption is successful ( $\mathcal{R}_Z$  is the recipient of a payment),  $\mathcal{R}_Z$  must verify that the  
 932 *ZethNote* recovered is the opening of a commitment in the Merkle tree of  $\widetilde{\text{Mixer}}$ . If not,  
 933  $\mathcal{R}_Z$  rejects the (invalid) payment.

934 We describe below the steps that  $\mathcal{R}_Z$  needs to carry out for all events  $evMixOut \in$   
 935 **MixEventDType** emitted by  $\widetilde{\text{Mixer}}$ , in order to receive payments:

- 936 1. Compute the new root  $mkroot'$  of the Merkle tree of commitments, after adding the
- 937 new values  $evMixOut.cms$ . If this value does not match the new root  $evMixOut.mkroot$
- 938 emitted by  $\widetilde{\text{Mixer}}$ , abort.

2. Try to decrypt the ciphertexts:

$$zn_j = \text{EncSch.Dec}(\mathcal{R}_Z.priv.skenc, evMixOut.ciphers[j])$$

- 939 3. For each successful decryption, let  $j$  be the index of the decrypted ciphertext:
  - 940 (a) Check whether the recovered plaintext  $zn_j$  is a well-formed *ZethNote*. Abort
  - 941 if it is not well-formed.
  - (b) Check that the recovered *ZethNote*  $zn_j$  is the opening of the corresponding
  - commitment  $evMixOut.cms[j]$ :

$$evMixOut.cms[j] = \text{ComSch.Com}(zn_j.apk, zn_j.\rho, zn_j.v; zn_j.r)$$

942 Abort if the note is not a valid opening.

943 (c) Additionally, if sender  $\mathcal{S}_{\mathcal{Z}}$ , and recipient  $\mathcal{R}_{\mathcal{Z}}$  had a contractual agreement,  
 944 then  $\mathcal{R}_{\mathcal{Z}}$  needs to check that the terms of this agreement are fulfilled by all  
 945 the recovered *ZethNotes*, abort otherwise.

946 Note that Steps 1 and 3b are required to ensure that data decrypted by  $\mathcal{R}_{\mathcal{Z}}$  ex-  
 947 actly matches the data committed to in **Mixer**. In particular, Step 1 requires  $\mathcal{R}_{\mathcal{Z}}$   
 948 to maintain or have access to some representation of the Merkle tree of commitments.  
 949 See Section 4.1.2 for further details.

## 950 2.7 Security requirements for the primitives

951 We list below the security requirements to instantiate the primitives of the **Zeth** protocol.

- 952 •  $\text{CRH}^{\text{hsig}}$  and  $\text{CRH}^{\text{ots}}$  MUST be collision resistant functions (see Definition 1.5.16).
- 953 •  $\text{PRF}^{\text{addr}}$ ,  $\text{PRF}^{\text{nf}}$ ,  $\text{PRF}^{\text{rho}}$  and  $\text{PRF}^{\text{pk}}$  MUST be PRF when keyed by  $ask$  and  $\phi$ , and be  
 954 collision resistant (see Definition 1.5.16, and Section 1.5.8).
- 955 •  $\text{SigSch}_{\text{OT-SIG}}$  MUST be UF-CMA (see Definition 1.5.24 and Appendix A.2.3).
- 956 •  $\text{ComSch}$  MUST be computationally hiding and binding (see Section 1.5.9).
- 957 •  $\text{MKHASH}$  MUST be collision resistant with  $h_0 = 0_{\mathbb{F}_r}$  (see Section 1.5.6).<sup>2</sup>
- 958 •  $\text{EncSch}$  MUST be IND-CCA2 and IK-CCA (see, respectively, [ABR99, Definition 8]  
 959 and Definition 1.5.10).
- 960 •  $\text{Unpack}(\text{Pack}(X)) = X$  and  $\text{Unpack}(\text{Pack}_{\text{rsd}}(X)) = X$  MUST hold.
- 961 •  $\text{decode}(\text{encode}(X)) = X$  MUST hold.

### 962 2.7.1 Additional notes

#### 963 Defining *hsig*

The signature hash *hsig* is a variable used to bind the signature keys to the primary inputs. We use the same definition of *hsig* as **Zcash** to prevent the Faerie Gold attack and thus

$$hsig = \text{CRH}^{\text{hsig}}(nfs, vk).$$

964 As a private transaction is uniquely determined by its nullifiers  $nfs = (nf_0, \dots, nf_{\text{JSIN}-1})$ ,  
 965 and because of the collision resistance of  $\text{CRH}^{\text{hsig}}$ , a transaction is uniquely determined  
 966 by *hsig* (with overwhelming probability). We did not use the *randomSeed* defined in  
 967 **Zcash** however, since this is only necessary to achieve uniqueness of *hsig* for transactions  
 968 *in transit* (i.e. not mined yet) [Hop16]. The uniqueness of *hsig* is a requirement to  
 969 prevent the Fairy Gold attack.

---

<sup>2</sup>This security requirement is equivalent to the one in [ZCa19, Section 5.4.1.3] where finding a preimage of  $0^{\text{SHA256DLEN}}$  must be hard.



970 **Security Requirement.**

- 971 • The variable *hsig* MUST be derived from the nullifiers  $\{nf_i\}_{i \in [JSIN]}$  and the signing  
 972 key *vk* using a collision resistant function. Doing so, makes sure that *hsig* is unique  
 973 for each  $tx_{\text{Mix}}$  with overwhelming probability.

974 **Defining  $\rho$**

We define  $\rho$  like in **Zcash** in order to prevent the Faerie Gold attack. A malicious sender could reuse the same  $\rho$  for a given recipient, hence correctly generating a *ZethNote* which could become unspendable by the recipient. Making  $\rho$  the output of a collision resistant PRF with random variable  $\phi$  as key and with  $tx_{\text{Mix}}$ 's *hsig* as input ensures, with overwhelming probability, the uniqueness of  $\rho$  and prevents this attack. Thus,

$$\rho_j = \text{PRF}_{\phi}^{\text{rho}}(j, \text{hsig}).$$

975 **Message authentication tags  $h_i$**

The message authentication tags are used to bind the signature hash to the input notes spending keys, to show ownership of the spent notes. Each tag derived from a note owner's spending key and the signature hash MUST be unique for each note with overwhelming probability. We define

$$h_i = \text{PRF}_{ask_i}^{\text{pk}}(i, \text{hsig}).$$

## Chapter 3

# Instantiation of the cryptographic primitives

In this chapter, we start by instantiating the cryptographic building blocks used in previous sections to describe the **Zeth** DAP design. Finally, we proceed by providing security proofs justifying that our instantiation complies with the security requirements listed in previous sections.

Note that, in several cases, it is necessary to specify details in terms of concrete properties of the curve **Curve** and associated scalar field  $\mathbb{F}_r$ . In these cases, we focus on two curves of interest: **BN-254** and **BLS12-377**. We note, however, that other suitable curves could be used.

**BN-254** [Rk19] has several properties that make it implementation-friendly. Elements of both the base field and scalar field can be represented in **ETHWORDLEN** bits (the native word size of the EVM), allowing efficient encoding and manipulation of such elements. Moreover, a subset of operations on **BN-254** are supported by the EVM through precompiled contracts. These precompiled contracts enable verification of signatures (Section 3.4) and zero-knowledge proofs (Section 3.6), required by this protocol, with minimal gas overhead.

**BLS12-377** [BCG<sup>+</sup>20], like **BN-254**, has the advantage that scalar field elements can be represented within **ETHWORDLEN**-bit words (although the same is not true of base field elements). However, the **EVM** provides no native support for **BLS12-377**, which increases the complexity of the **Mixer** implementation (see Section 2.5 for details of the operations to be performed). An advantage that **BLS12-377** does provide, is that it is the “inner” curve of a one-layer chain (as described in [BCG<sup>+</sup>20, HG20]). Therefore zero-knowledge proofs using **BLS12-377** can be efficiently verified by statements in other zero-knowledge proofs using an appropriate “outer” pairing. Support for **BLS12-377** in **Zeth** therefore admits several applications (not explicitly covered by this document), such as aggregation of proofs over multiple **Zeth** transactions (e.g. [Ron20]).

Further details related to implementation and optimization are given in Chapter 4.

### 3.1 Instantiating the PRFs, ComSch and CRHs

The functions  $\text{CRH}^{\text{hsig}}$  and  $\text{CRH}^{\text{ots}}$  are instantiated with SHA256 [oST15] which we assume to be collision resistant. Furthermore,  $\text{ComSch}$ ,  $\text{PRF}^{\text{pk}}(x)$ ,  $\text{PRF}^{\text{rho}}(x)$ ,  $\text{PRF}^{\text{addr}}(x)$ , and  $\text{PRF}^{\text{nf}}(x)$  are all instantiated with Blake2’s hash function optimized for 32-bit platforms, Blake2s, which we prove in the Weakly Ideal Cipher Model [LMN16] to be from a family of PRF and collision resistant functions. The Weakly Ideal Cipher model assumes that Blake2’s underlying block cipher is ideal and has no structural weaknesses (see Appendix D.2). In addition to that, and to ensure that the functions  $\text{PRF}^{\text{pk}}(x)$ ,  $\text{PRF}^{\text{rho}}(x)$ ,  $\text{PRF}^{\text{addr}}(x)$ , and  $\text{PRF}^{\text{nf}}(x)$  compute images lying in different domains, we use different message prefixes (or “domain separators”) for the PRFs inputs. This approach ensures that the  $\text{apk}_i$ ’s,  $\text{nf}_i$ ’s,  $\rho_i$ ’s, and  $h_i$ ’s have independent distributions from a PPT adversary point of view.

#### Note

It is important to note that, for this approach to be secure, the hash function used needs to be secure against *chosen-prefix collision attacks* [Ste15].

Furthermore, we take:

- $\text{RTRAPLEN}, \text{ASKLEN}, \text{PHILEN} = \text{BLAKE2sCLEN}$

#### 3.1.1 Blake2 primitive

Blake [AHMP08] is a hash family that was presented as a candidate at the SHA3 competition. Blake2 is the next iteration of the family which has been further optimized to achieve higher throughput thanks to some optimizations and by being less conservative on its security<sup>1</sup>. Blake and Blake2 are based on the ChaCha stream cipher [Ber08a] composed with the HAIFA framework [BD07]. ChaCha defined over 20 rounds, as used in Blake2, is deemed secure and a PRF based on today’s cryptanalysis [Pro14, CM16]. Blake2 is specified in RFC-7693 [MJS15] and licensed under CC0. Blake2s is an instantiation of Blake2 optimized for 32-bit platforms. As such, to reason about the security of Blake2s we prove the security of Blake2.

**Blake security** Blake security has been heavily scrutinized through the SHA3 competition [VNP10, MQZ10, AMP10, AAM12, AMPŠ12, ALM12, HMRS12]. Blake2 has also been thoroughly cryptanalyzed independently [GKN<sup>+</sup>14, Hao14, EFK15, NA19]. For  $n$ -bit long digests/outputs, the hash and compression functions present  $n/2$ -bit of collision resistance and  $n$ -bit of preimage resistance, immunity to length extension, and indistinguishability from a random oracle [ANWOW13]. They have furthermore been demonstrated secure in the Weakly Ideal Cipher Model [LMN16] (WICM, see Appendix D.1.1). More

<sup>1</sup>The authors increased the number of rounds of Blake for the SHA3 competition to be more conservative on security. They however showed afterwards that this change was not “meaningfully more secure” and thus reverted it for Blake2 (see [ANWOW13, Section 2.1]).

1037 particularly, Luykx et al. show that Blake2 is indifferentiable from a random oracle in  
 1038 this model and is a PRF. We use this result in Appendix D.2 to show the collision  
 1039 resistance of Blake2. We also prove that, given that Blake2 is collision resistant and a  
 1040 PRF, Blake2( $r||x$ ) is a computationally binding and computationally hiding commitment  
 1041 scheme for input  $x$  and randomness  $r$ .

#### Note

We assume that the encryption scheme used in the Blake2 underlying compression function – which is derived from ChaCha20 – has no exploitable structural behaviour. More precisely, that this encryption scheme behaves like a weak ideal cipher. We provide proofs in this model.

1042

### 1043 3.1.2 Commitment scheme

We define our commitment scheme as follows,

$$\begin{aligned} \text{ComSch.Setup} &: \{1^\lambda \text{ s.t. } \lambda \in \mathbb{N}\} \rightarrow \mathbb{B}^* \\ \text{ComSch.Com} &: (\mathbb{B}^{\text{PRFADDRROUTLEN}} \times \mathbb{B}^{\text{PRFRHOOUTLEN}} \times \mathbb{B}^{\text{ZVALUELEN}}) \times \mathbb{B}^{\text{RTRAPLEN}} \rightarrow \mathbb{F}_r \end{aligned}$$

We instantiate the commitment scheme with Blake2s as follows,

$$\begin{aligned} pp &= \text{ComSch.Setup}(1^\lambda) \text{ (corresponds to Blake2s's constant PB and } r) \\ cm &= \text{ComSch.Com}(m = (apk, \rho, v); r) \\ &= \text{decode}_{\mathbb{N}}(\text{Blake2s}(r||apk||\rho||v)) \pmod{r} \end{aligned}$$

1044 **Remark 3.1.1.** We set the commitment digest length in the parameter block PB [MJS15].

### 1045 Security proof

1046 The commitment scheme defined above is computationally hiding and binding in the  
 1047 WICM, see Appendix D.2.4. However, because of the modulo  $r$  operation, the scheme  
 1048 is only  $(\text{FIELDLEN}/2)$ -bit binding.

### 1049 3.1.3 PRFs

We show in this section how we instantiate the PRFs with Blake primitives. As a reminder the PRFs are defined as follows,

$$\begin{aligned} \text{PRF}^{\text{addr}} &: \mathbb{B}^{\text{ASKLEN}} \times \{0\} \rightarrow \mathbb{B}^{\text{PRFADDRROUTLEN}} \\ \text{PRF}^{\text{pk}} &: (\mathbb{B}^{\text{ASKLEN}} \times [\text{JSIN}]) \times \mathbb{B}^{\text{CRHHSIGOUTLEN}} \rightarrow \mathbb{B}^{\text{PRFPKOUTLEN}} \\ \text{PRF}^{\text{nf}} &: \mathbb{B}^{\text{ASKLEN}} \times \mathbb{B}^{\text{PRFRHOOUTLEN}} \rightarrow \mathbb{B}^{\text{PRFNFOUTLEN}} \\ \text{PRF}^{\text{rho}} &: (\mathbb{B}^{\text{PHILEN}} \times [\text{JSOUT}]) \times \mathbb{B}^{\text{CRHHSIGOUTLEN}} \rightarrow \mathbb{B}^{\text{PRFRHOOUTLEN}} \end{aligned}$$

As we instantiate the PRFs with Blake2s we have,

$$\text{PRFADDRROUTLEN, PRFNFOUTLEN, PRFPKOUTLEN, PRFRHOOUTLEN} = \text{BLAKE2sCLEN}$$

To ensure that the PRFs have independent distributions, we first introduce tagging functions  $\text{tag}^x$  which truncate and prepend with a distinct tag the PRFs key. We have,

$$\begin{aligned} \text{tag}^{\text{addr}} &: \mathbb{B}^{\text{ASKLEN}} \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \\ \text{tag}^{\text{pk}} &: \mathbb{B}^{\text{ASKLEN}} \times [\text{JSIN}] \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \\ \text{tag}^{\text{nf}} &: \mathbb{B}^{\text{ASKLEN}} \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \\ \text{tag}^{\text{rho}} &: \mathbb{B}^{\text{PHILEN}} \times [\text{JSOUT}] \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \end{aligned}$$

The tagging functions are instantiated as follows,

$$\begin{aligned} \text{tag}^{\text{addr}}(\text{aux.jsins}[i].\text{ask}) &= \text{tag}_{\text{ask}}^{\text{addr}} \\ &= (1) \parallel (1)^{\lceil \frac{\text{JSMAX}}{2} \rceil} \parallel (0, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN}-3-\lceil \frac{\text{JSMAX}}{2} \rceil}(\text{aux.jsins}[i].\text{ask}) \\ \text{tag}^{\text{nf}}(\text{aux.jsins}[i].\text{ask}) &= \text{tag}_{\text{ask}}^{\text{nf}} \\ &= (1) \parallel (1)^{\lceil \frac{\text{JSMAX}}{2} \rceil} \parallel (1, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN}-3-\lceil \frac{\text{JSMAX}}{2} \rceil}(\text{aux.jsins}[i].\text{ask}) \\ \text{tag}^{\text{pk}}(\text{aux.jsins}[i].\text{ask}, i) &= \text{tag}_{\text{ask}, i}^{\text{pk}} \\ &= (0) \parallel \text{pad}_{\lceil \frac{\text{JSMAX}}{2} \rceil}(\text{encode}_{\mathbb{N}}(i)) \parallel (0, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN}-3-\lceil \frac{\text{JSMAX}}{2} \rceil}(\text{aux.jsins}[i].\text{ask}) \\ \text{tag}^{\text{rho}}(\text{aux}.\phi, j) &= \text{tag}_{\text{ask}, j}^{\text{rho}} \\ &= (0) \parallel \text{pad}_{\lceil \frac{\text{JSMAX}}{2} \rceil}(\text{encode}_{\mathbb{N}}(j)) \parallel (1, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN}-3-\lceil \frac{\text{JSMAX}}{2} \rceil}(\text{aux}.\phi) \end{aligned}$$

1050 where  $\text{pad}_{\lceil \frac{\text{JSMAX}}{2} \rceil}(\text{encode}_{\mathbb{N}}(i))$  is the function that pads the binary representation of  $i$  by  
 1051 adding 0's before the most significant bit (e.g. assuming big endian encoding,  $\text{pad}_2(\text{encode}_{\mathbb{N}}(1)) =$   
 1052 01).

We now present how the PRFs are instantiated,

$$\begin{aligned} \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{addr}}(0) &= \text{aux.jsins}[i].\text{znote.apk} \\ &= \text{Blake2s}(\text{tag}^{\text{addr}}(\text{aux.jsins}[i].\text{ask}) \parallel \text{pad}_{\text{BLAKE2sCLEN}}(0)) \\ \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{nf}}(\text{aux.jsins}[i].\rho) &= \text{prim.nfs}[i] \\ &= \text{Blake2s}(\text{tag}^{\text{nf}}(\text{aux.jsins}[i].\text{ask}) \parallel \text{aux.jsins}[i].\text{znote}.\rho) \\ \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{pk}}(i, \text{prim.hsig}) &= \text{prim.htags}[i] \\ &= \text{Blake2s}(\text{tag}^{\text{pk}}(\text{aux.jsins}[i].\text{ask}, i) \parallel \text{prim.hsig}) \\ \text{PRF}_{\text{aux}.\phi}^{\text{rho}}(j, \text{prim.hsig}) &= \text{aux.znotes}[j].\rho \\ &= \text{Blake2s}(\text{tag}^{\text{rho}}(\text{aux}.\phi, j) \parallel \text{prim.hsig}) \end{aligned}$$

1053 **Remark 3.1.2.** We set the PRFs' output length in the Blake2s's parameter block PB.

## Security proof

The functions defined above are collision resistant and PRFs in the WICM, see Appendix D.2. Because of the tagging functions, the security parameter of the PRFs becomes  $\lambda = \text{BLAKE2sCLEN}/2 - \text{JSMAX}/4 - 3/2$ .

### 3.1.4 Collision resistant hashes

We instantiate in this section the collision resistant hash functions  $\text{CRH}^{\text{hsig}}$  and  $\text{CRH}^{\text{ots}}$  with SHA256. As a consequence, we have,

$$\text{CRHHSIGOUTLEN} = \text{CRHOTSOUTLEN} = \text{SHA256DLEN}$$

**SHA256 Security** SHA-256 (Secure Hash Algorithm 256) is a hash function designed by the National Security Agency (NSA) in 2001. It is based on the Merkle–Damgård structure, the Davies–Meyer compression function construct [BRS02, Function  $f_5$  in Figure 3] and the classified SHACAL-2 block cipher.

Collision attacks have been thoroughly studied by the research community [SS08, MNS11]. The best attacks at this day, are second-order differential attack by Lamberger et al. [LM11] on the SHA-256 compression function reduced to 46 out of 64 rounds.

Many researchers [IS09, AGM<sup>+</sup>09] have also studied preimage attacks on SHA-256 with reduced rounds. Guo et al. [GLRW10] in particular were among the first to use the meet in the middle strategy [AS09] and achieved more efficient ones on 42-step SHA-256. Khovratovich et al. in 2012 [KRS12] have so far presented the best preimage attacks, on 45-round and 52-round SHA-256 as well as a 52-round attack on the SHA-256 compression function.

Li et al. have published in 2012 [LIS12] a noteworthy paper on converting meet in the middle preimage attack into pseudo collision attack. Using preimage attacks by bicliques, they found pseudo collisions attacks on 52 steps of SHA-256.

**Claim 1.** SHA256 is 128-bit collision resistant.

## 3.2 Instantiating MKHASH

In this section we describe the instantiation of MKHASH with a compression function based on MIMC [AGR<sup>+</sup>16]. We firstly show how the compression function is constructed, and prove that this instantiation complies with the security requirements mentioned in Section 2.7

### 3.2.1 MIMC Encryption

MIMC is a block cipher with a simple design, consisting of a number of rounds (denoted *rounds*). During the  $i$ -th round, the message  $m$  is mixed with the encryption key  $k$  and a randomly chosen constant  $c[i]$ , and a permutation function is applied to generate a new value of  $m$ . The permutation function consists of exponentiation with a carefully chosen

1086 exponent  $e$  (see Section 3.2.1). Note that  $rounds$  depends on the desired security level  
 1087  $\lambda$ . We denote the encryption function by **MIMC-Encrypt** and illustrate it in Fig. 3.1.

```

MIMC-Encrypt( $k, m, c, e, rounds$ )
1 : foreach  $i \in [rounds]$  :
2 :    $m \leftarrow (k \text{ OP } c[i] \text{ OP } m)^e$ 
3 : return  $(m \text{ OP } k)$ 

```

Figure 3.1: MIMC Encryption function.

1088 **MIMC-Encrypt** can be defined on both binary and prime fields, and as such the **OP**  
 1089 operation corresponds to either  $\oplus$  or  $+$  (mod  $p$ ) [AGR<sup>+</sup>16, GRR<sup>+</sup>16]. For general prime  
 1090  $p$  (resp. positive integer  $n$ ), we denote by **MIMC<sub>p</sub>** (resp. **MIMC<sub>2<sup>n</sup></sub>**) the **MIMC-Encrypt**  
 1091 function defined over  $\mathbb{F}_p$  (resp.  $\mathbb{F}_{2^n}$ ). In this document, we only consider **MIMC** defined  
 1092 over prime fields (in particular, the field  $\mathbb{F}_r$  with elements written over  $\lambda$  bits, over which  
 1093 **ZkSnarkSch** operates).

Since block ciphers are usually defined over the product space of keys and messages,  
 we consider the variables  $c$ ,  $rounds$  and  $e$  as fixed. We thereby consider an instantiation  
 of **MIMC** with signature

$$\text{MIMCr} : \mathbb{F}_r \times \mathbb{F}_r \rightarrow \mathbb{F}_r$$

## 1094 Security parameters and analysis

1095 To ensure that the exponentiation leads to a permutation in  $\mathbb{F}_r$ , we consider  $e$  of the  
 1096 form  $e = 2^t - 1$  and  $e = 2^t + 1$  such that  $\gcd(e, r - 1) = 1$ . To achieve a security of  $\lambda$ ,  
 1097 we require that  $rounds = \left\lceil \frac{\log_2 r}{\log_2 e} \right\rceil$ .

1098 We refer to the **MIMC** paper [AGR<sup>+</sup>16, Section 4.2 and 5.1] for more details on the  
 1099 security analysis and attacks on the scheme. Note that **MIMCr** does not suffer from  
 1100 *inversion subfield attacks* as there are no proper subfields of  $\mathbb{F}_r$ .

## 1101 3.2.2 MIMC-based compression function

1102 There exist two main techniques to construct a hash function from a block-cipher (or  
 1103 permutation): sponge functions [BDPVA07] and iterated compression functions [BRS02].

1104 A Merkle tree is a binary tree of values of fixed size, where the values in each “layer”  
 1105 are generated by hashing pairs of values from the previous “layer”. That is, we require a  
 1106 compression function **MKHASH**, which we construct via the Miyaguchi-Preneel scheme.  
 1107 (Miyaguchi-Preneel is more secure [BRS02,  $f_5$  function] than the more flexible Davies-  
 1108 Meyer construct [GFBR06, Section 3], but this flexibility is not required in our case).

## 1109 Miyaguchi-Preneel compression construct

1110 Miyaguchi-Preneel (MP) [BRS02,  $f_3$  function] is a general scheme for constructing com-  
 1111 pression functions from block ciphers (see Section 1.5.6). Given a block cipher  $E$ , the

1112 corresponding compression function by  $f_E^{\text{MP}}$  is given in Fig. 3.2. The original construc-  
 1113 tion is defined over binary fields, however **Zeth** operates over prime fields. Hence, in the  
 1114 general discussion here we replace the bitwise addition operator  $\oplus$  by modular addition  
 1115 in  $\mathbb{F}_r$  (see [Har19]).

1116 We denote by **MIMC-MP** the compression function defined by the application of  
 1117 the Miyaguchi-Preneel construct over **MIMC**. Similarly, for general prime  $p$  we denote  
 1118 by **MIMC-MP<sub>p</sub>** (see Fig. 3.3) the compression function defined by application of the  
 1119 Miyaguchi-Preneel construct over **MIMC<sub>p</sub>**.

$f_E^{\text{MP}}(k, m)$
1 : $res \leftarrow E_k(m)$ 2 : <b>return</b> $(res + m + k) \pmod{r}$

Figure 3.2: MP construct in  $\mathbb{F}_r$ .

<b>MIMC-MP<sub>r</sub></b> ( $k, m$ )
1 : $res \leftarrow \text{MIMCr}(k, m)$ 2 : <b>return</b> $(res + k + m) \pmod{r}$

Figure 3.3: **MIMC-MP<sub>r</sub>** construction.

### 1120 3.2.3 An efficient instantiation of MIMC primitives

1121 To select appropriate instances of **MIMCr** and **MIMC-MP<sub>r</sub>**, we consider the cost (in terms  
 1122 of gas consumption and prover efficiency). For given  $e$  and *rounds*, the final definition  
 1123 of **MIMC-MP<sub>r</sub>** is given in Fig. 3.4 and Fig. 3.5.

<b>MIMCr</b> ( $k, m$ )	<b>InitRoundConstants</b> ()
1 : $c \leftarrow \text{InitRoundConstants}()$ 2 : <b>foreach</b> $i \in [\text{rounds}]$ : 3 : $m \leftarrow (k + c[i] + m)^e \pmod{r}$ 4 : <b>return</b> $(m + k) \pmod{r}$	$iv \leftarrow \text{Keccak256}(\text{"clearmatics\_mt\_seed"})$ $c[0] \leftarrow 0$ $c[1] \leftarrow \text{Keccak256}(iv)$ <b>foreach</b> $i \in \{2, \dots, \text{rounds}\}$ : $c[i] \leftarrow \text{Keccak256}(c[i-1])$ <b>return</b> $c = (c[0], \dots, c[\text{rounds} - 1])$

Figure 3.4: **MIMCr** full construction

<b>MIMC-MP<sub>r</sub></b> ( $k, m$ )
<b>return</b> <b>MIMCr</b> ( $k, m$ ) + $m + k \pmod{r}$

Figure 3.5: **MIMC-MP<sub>r</sub>** full construction

1124 **Remark 3.2.1.** Note that Keccak256 is the 256-bit digest instance of the Keccak family  
 1125 that won the NIST SHA-3 competition [GJMG11]. It is supported by the EVMvia an  
 1126 opcode (see [W<sup>+</sup>, Appendix G]), making it convenient for use in smart contracts.

1127 **Remark 3.2.2.** To increase the security of the MKHASH, different round constants for  
 1128 each level of the Merkle tree could be used.



We define MKHASH to be MIMC-MP over  $\mathbb{F}_r$ . Thereby, for input values  $m_0$  and  $m_1$ ,  $\text{MKHASH} : \mathbb{F}_r \times \mathbb{F}_r \rightarrow \mathbb{F}_r$  is defined by

$$\text{MKHASH}(m_0, m_1) = \text{MIMC-MP}_r(m_0, m_1) \quad (3.1)$$

For specific values of  $r$  (such as  $r_{\text{BN}}$  for BN-254 or  $r_{\text{BLS}}$  for BLS12-377), it remains to choose concrete values of  $e$  and *rounds*.

Note that small exponents  $e$  result in fewer constraints in the arithmetic circuit (see Section 2.2), while larger exponents can reduce the cost of Merkle tree operations on the contract (see Section 2.5). This is due to two factors, namely that exponentiation is cheaper to execute on a contract than in an arithmetic circuit, and that the number of rounds decreases with higher  $e$ . For instance, choosing  $e = 7$  results in 365 constraints and  $\approx 20k$  gas while  $e = 31$  corresponds to 417 constraints (+15%) and  $\approx 17k$  (−10%) in gas consumption. Repeating the same process for different exponents, we observe roughly the same order of magnitude gain on the gas consumption and loss on the number of constraints.

The number of constraints of MIMC-MP for several exponents  $e$  is given by the formula

$$\text{constraints} = \text{rounds} \cdot \text{mults} + 1$$

where  $\text{rounds} = \lceil \frac{\log_2 r}{\log_2 e} \rceil$ , *mults* is the number of multiplications required for exponentiation and the additional constraint (corresponding to +1 in the above formula) is a result of the final message and key addition. Note that for  $e = 2^t - 1$  we have  $\text{mults} = 2 \cdot t - 2$ , using the *square-and-multiply* algorithm [MVOV96], and for  $e = 2^t + 1$  we have  $\text{mults} = t + 1$ .

For several concrete values of  $e$ , the number of *rounds* required to attain the desired security level, along with the number of constraints, are shown in Table 3.1.

For the case of BN-254 we set  $e = 7$  and *rounds* = 91, targetting a 254-bit security level. For BLS12-377 we set  $e = 31$  and *rounds* = 51, targetting a 253-bit security level. These values are chosen such that they satisfy the requirement that  $\gcd(e, r - 1) = 1$  and give a good balance between the number of constraints in the arithmetic circuit and the gas cost of hashing on the contract.

### 3.2.4 Security requirements satisfaction

After presenting the state of the art of MiMC cryptanalysis, we present the security proof of MIMC-MP collision resistance.

#### Cryptanalysis of MIMC block cipher and primitives

MIMC's security is increasingly being analysed since the primitive has gained traction in zero-knowledge and cryptocurrency communities for its succinct algebraic constraint representation. As of today, we do not know of any attacks breaking MIMC on prime fields on full rounds.

$e$	BN-254		BLS12-377	
	<i>rounds</i>	<i>constraints</i>	<i>rounds</i>	<i>constraints</i>
5	110	331		
7	91	365		
17	65	316	62	311
31	52	417	51	409
127	37	445	37	445
257	32	289	32	289
511	29	465		
2047	24	481	23	461
8191	20	481	20	481
32676	17	477		
65537	16	273	16	273
131071	15	481	15	481
524287	14	505	14	505
1048577	13	274	13	274
2097151	13	521		

Table 3.1: Arithmetic constraints required to represent MIMC-MP as an R1CS program, for different exponents  $e$  and curves. Missing entries where  $\gcd(e, r - 1) \neq 1$

The first attack on MIMC was an interpolation attack [LP19] which targets a reduced-round version for a scenario in which the attacker has only limited memory. An attack on Feistel-based MIMC [Bon19] was discovered shortly after, by using generic properties of the used Feistel construction (instead of exploiting properties of the primitive itself). Additionally, [ACG<sup>+</sup>19] proposes an attack based on Gröbner basis. The authors state that by introducing a new intermediate variable in each round, the resulting multivariate system of equations is a Gröbner basis. As such, the first step of a Gröbner basis attack can be obtained for free. However, the following steps of the attack are so computationally demanding that the attack becomes infeasible in practice. A recent work [EGL<sup>+</sup>20] targets MIMC on binary fields, and achieves a full-round break of the scheme. While, the attack presented does not apply to prime fields, the authors note that it “can be generalized to include ciphers over  $\mathbb{F}_p$ ”, and that only the lack of efficient distinguishers over prime fields precludes this. Another attack from Beyne et al [BCD<sup>+</sup>20] uses a low complexity distinguisher against full MIMC permutation leading to a practical collision attack on reduced round sponge-based MIMC hash defined with security of 128 bits.

### Security proof of MIMC-MP collision resistance

We now prove that this compression scheme satisfies all the security requirements listed in Section 2.7. To do so, we first assume that the round constants are pseudo-random, i.e. that Keccak256 is a PRF.

**Lemma 3.2.1.** *Keccak256 is a PRF with  $\lambda = 128$ .*

The security of MIMC-MP derives from a more general result, i.e. from modelling MIMC as an ideal cipher (see Definition 1.5.12). More specifically, we show a security result for the MP construction on  $\mathbb{F}_r$  by proving that, in the Ideal Cipher Model, the collision resistance advantage of any adversary is bounded by  $\frac{q(q+1)}{r}$ , where  $q$  is the number of different queries that the attacker makes to the oracle. This means that, assuming a maximum  $q$  number of possible encryption/decryption queries, parameter  $r$  can be chosen to make the advantage small as needed and  $f_E^{\text{MP}}$  considered collision resistant. Similar result applies to the  $2^n$  case.

The instance of MIMC we use is modelled as an ideal cipher defined on field elements, for this reason we consider a variant of the ICM model where the keys, inputs and outputs are field elements in  $\mathbb{F}_r$  and the block cipher scheme, with key  $k$ , correspond to a family of  $r$  independent random permutations  $f_k : \mathbb{F}_r \times \mathbb{F}_r \rightarrow \mathbb{F}_r$ .

In the proof, without loss of generality, we assume the following conventions for an adversary  $\mathcal{A}$ :

- the adversary asks distinct queries: i.e. if  $\mathcal{A}$  asks a query  $O^E(k, m)$  and this returns  $y$ , then  $\mathcal{A}$  does not ask a subsequent query of  $O^E(k, m)$  or  $O^{E^{-1}}(k, y)$ , and inversely;
- the adversary necessarily obtained the candidate collision from the oracle. This property follows suite from modelling MIMC as an ideal cipher.

**Lemma 3.2.2.** *Let  $f_E^{\text{MP}}$  be the MP compression function built on an ideal block-cipher  $E$  on  $\mathbb{F}_r$ , the probability for an adversary  $\mathcal{A}$  to find a collision is not greater than  $q(q+1)/r$  where  $q$  is a (positive) number of distinct oracle queries.*

The following proof has been adapted from [BRS02, Lemma 3.3]<sup>2</sup>.

*Proof.* Fix  $h_0 \in \mathbb{F}_r$ . Let  $\mathcal{A}$  be an adversary attacking the compression function  $f_E^{\text{MP}}$ . Assume that  $\mathcal{A}$  asks the oracles  $O^E$  and  $O^{E^{-1}}$  a total of *distinct*  $q$  queries. Let us denote the result of the  $q$  queries and output of the attacker (candidate collision) as  $((k_1, m_1, y_1), \dots, (k_q, m_q, y_q), \text{out})$ . If  $\mathcal{A}$  is successful it means that it outputs  $(k, m)$ ,  $(k', m')$  such that either  $(k, m) \neq (k', m')$  and  $f_E^{\text{MP}}(k, m) = f_E^{\text{MP}}(k', m')$  or  $f_E^{\text{MP}}(k, m) = h_0$ . By the definition of  $f_E^{\text{MP}}$ , we have that  $E_k(m) + m + k = E_{k'}(m') + m' + k'$  for the first case, or  $E_k(m) + m + k = h_0$  for the second. So either there are distinct  $r, s \in [1, \dots, q]$  such that  $(k_r, m_r, y_r) = (k, m, E_k(m))$  and  $(k_s, m_s, y_s) = (k', m', E_{k'}(m'))$  and  $E_{k_r}(m_r) + m_r + k_r = E_{k_s}(m_s) + m_s + k_s$  or else there is an  $r \in [1, \dots, q]$  s.t.  $(k_r, m_r, y_r) = (k, m, h_0)$  and  $E_{k_r}(m_r) + m_r + k_r = h_0$ . We show that this event is unlikely.

In fact, for each  $i \in [1, \dots, q]$ , let  $C_i$  be the event that either  $y_i + m_i + k_i = h_0$  or does exist  $j \in [1, \dots, i-1]$  s.t.  $y_i + m_i + k_i = y_j + m_j + k_j$ . When carrying out the simulation  $y_i$  or  $m_i$  was randomly selected from a set of at least  $r - (i-1)$  elements, so

<sup>2</sup>It states the collision resistance of a set of compression functions  $f_1, \dots, f_{12}$ , denoted as *group-1 compression functions* and showed in [BRS02, Figure 3]. As mentioned above, Miyaguchi-Preneel corresponds to  $f_3$  of that group. Since the proof of [BRS02, Lemma 3.3] shows collision resistance of  $f_1$ , we slightly modified it to work for  $f_3$ .

1215  $\Pr[C_i] \leq i/(\mathbf{r} - i)$ . This means that for the collision advantage of  $\mathcal{A}$ ,  $\text{Adv}_{\text{f}_{\text{MP}}^{\text{coll}}, \mathcal{A}}^{\text{coll}}$  it holds  
 1216 that  $\text{Adv}_{\text{f}_{\text{MP}}^{\text{coll}}, \mathcal{A}}^{\text{coll}} \leq \Pr[C_1 \vee \dots \vee C_q] \leq \sum_{i=1}^q \Pr[C_i]$ . For  $q \leq \frac{\mathbf{r}}{2}$  this probability is bounded  
 1217 by  $l \cdot \frac{q(q+1)}{\mathbf{r}}$ . However, we allow only a polynomial number of queries, thus for  $q = \text{poly}(\lambda)$   
 1218 this probability becomes  $\frac{\text{poly}(\lambda)}{\mathbf{r}}$ , where  $\mathbf{r} \approx 2^\lambda$ .  $\square$

#### Note

Lemma 3.2.2 is applicable to our case by the strong assumption of MIMCr being an ideal cipher. In other words, the proof does not take into account any structural weakness or knowledge that an attacker is aware of. Any such additional information could make Lemma 3.2.2 invalid, and consequently could be used to break the collision resistance.

1219

1220 **Remark 3.2.3.** Note that from Lemma 3.2.2 follows that the collision resistance security  
 1221 of the Zeth Merkle tree is  $\log_2(\mathbf{r}/2)$  (around 127 bits for  $\mathbf{r} = \mathbf{r}_{\text{BN}}$  or  $\mathbf{r}_{\text{BLS}}$ ).

#### Note

MIMC has *not* received as much cryptanalytic scrutiny as other “older” and more established hash functions. This is important to note since, for these type of primitives which are not provably secure, the amount of attacks received by a scheme is a great indicator of its security and robustness. A natural alternative to MIMC here consists in using Pedersen hash which is provably collision resistant under the discrete-logarithm assumption.

1222

### 1223 3.3 Zeth statement after primitive instantiation

1224 After instantiating the various primitives and providing security proofs to justify that  
 1225 they comply with the security requirements listed in previous sections,  $\mathbf{R}^{\mathbf{Z}}$  now becomes:

- 1226 • For each  $i \in [\text{JSIN}]$ :
  - 1227 1.  $\text{aux.jsins}[i].\text{znote.apk} = \text{Blake2s}(\text{tag}_{\text{ask}}^{\text{addr}} \parallel \text{pad}_{\text{BLAKE2sCLEN}}(0))$   
 1228 with  $\text{tag}_{\text{ask}}^{\text{addr}}$  defined in Section 3.1.3
  - 1229 2.  $\text{aux.jsins}[i].\text{nf} = \text{Blake2s}(\text{tag}_{\text{ask}}^{\text{nf}} \parallel \text{aux.jsins}[i].\text{znote}.\rho)$   
 1230 with  $\text{tag}_{\text{ask}}^{\text{nf}}$  defined in Section 3.1.3
  - 1231 3.  $\text{aux.jsins}[i].\text{cm} = \text{Blake2s}(\text{aux.jsins}[i].\text{znote}.\text{r} \parallel m)$   
 1232 with  $m = \text{aux.jsins}[i].\text{znote.apk} \parallel \text{aux.jsins}[i].\text{znote}.\rho \parallel \text{aux.jsins}[i].\text{znote}.\text{v}$
  - 1233 4.  $\text{aux.htags}[i] = \text{Blake2s}(\text{tag}_{\text{ask},i}^{\text{pk}} \parallel \text{prim.hsigs})$  (malleability fix, see Appendix A)  
 1234 with  $\text{tag}_{\text{ask},i}^{\text{pk}}$  defined in Section 3.1.3
  - 1235 5.  $(\text{aux.jsins}[i].\text{znote}.\text{v}) \cdot (1 - e) = 0$  is satisfied for the boolean value  $e$  set such  
 1236 that if  $\text{aux.jsins}[i].\text{znote}.\text{v} > 0$  then  $e = 1$ .

- 1237 6. The Merkle root  $mkroot'$  used to check the Merkle authentication path  $aux.jsins[i].mkpath$   
 1238 of commitment  $aux.jsins[i].cm$ , with  $MIMC-MP_r$ , equals  $prim.mkroot$  if  $e = 1$ .
- 1239 7.  $prim.nfs[i]$   
 1240  $= \{\text{Pack}_{\mathbb{F}_r}(aux.jsins[i].nf[k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}])\}_{k \in [\lfloor \text{PRFNFOUTLEN}/\text{FIELD CAP} \rfloor]}$
- 1241 8.  $prim.htags[i]$   
 1242  $= \{\text{Pack}_{\mathbb{F}_r}(aux.htags[i][k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}])\}_{k \in [\lfloor \text{PRFPKOUTLEN}/\text{FIELD CAP} \rfloor]}$
- 1243 • For each  $j \in [\text{JSOUT}]$ :
- 1244 1.  $aux.znotes[j].\rho = \text{Blake2s}(tag_{ask,j}^\rho || prim.hsigs)$  (malleability fix, see Appendix A)  
 1245 with  $tag_{ask,j}^\rho$  defined in Section 3.1.3
- 1246 2.  $prim.cms[j] = \text{Blake2s}(aux.znotes[j].r || m)$   
 1247 with  $m = aux.znotes[j].apk || aux.znotes[j].\rho || aux.znotes[j].v$
- 1248 •  $prim.hsigs = \{\text{Pack}_{\mathbb{F}_r}(aux.hsigs[k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}])\}_{k \in [\lfloor \text{CRHHSIGOUTLEN}/\text{FIELD CAP} \rfloor]}$
- 1249 •  $prim.rsd = \text{Pack}_{rsd}(\{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}, aux.vin, aux.vout, aux.hsigs, \{aux.htags[i]\}_{i \in [\text{JSIN}]})$
- Check that the “joinsplit is balanced”, i.e. check that the joinsplit equation holds:

$$\begin{aligned} & \text{Pack}_{\mathbb{F}_r}(aux.vin) + \sum_{i \in [\text{JSIN}]} \text{Pack}_{\mathbb{F}_r}(aux.jsins[i].znote.v) \\ &= \sum_{j \in [\text{JSOUT}]} \text{Pack}_{\mathbb{F}_r}(aux.znotes[j].v) + \text{Pack}_{\mathbb{F}_r}(aux.vout) \end{aligned}$$

1250 **Remark 3.3.1.** For higher security, we could use Blake2b with 32-byte output instead  
 1251 of SHA256. In fact, since a precompiled contract computing the Blake2 compression  
 1252 function [MJS15] has been added to the Istanbul release of **Ethereum** (EIP 152 [THH15]),  
 1253 it could be possible to write a small wrapper on the smart contracts, in order to hash  
 1254 with Blake2b with any parameter.

### 1255 3.3.1 Instantiating the packing functions

1256 As we consider SNARKs based on arithmetic circuits defined over a prime field, all  
 1257 variables in the constraint system are interpreted as field elements. Nevertheless, as  
 1258 illustrated in Section 2.2, part of the statement consists of functions whose co-domains  
 1259 are sets of binary strings (which may be longer than the bit representation of elements of  
 1260 the finite field). While a bit (i.e.  $\{0, 1\}$ ) is an element of  $\mathbb{F}_p$  ( $p$  prime), it is important to  
 1261 minimize the number of gates in the arithmetic circuit (for proof generation efficiency),  
 1262 and to minimize the number of input wires (to improve verification time). This can be  
 1263 done by representing fragments of binary strings as the base 2 decomposition of field  
 1264 elements, thereby “packing” binary strings into multiple elements. Converting binary  
 1265 strings into field elements requires the addition of some arithmetic gates (extending  
 1266 the statement to be proven), but reduces the number of primary inputs (reducing the

complexity of the SNARK verification carried out on-chain). The cost of Groth16 zk-SNARK [Gro16] proof verification is linear in the number of primary inputs, since each input acts as a scalar in a costly scalar multiplication of a curve point in  $\mathbb{G}_1$ . Hence, while packing slightly increases the prover cost – by adding constraints to the circuit – it simplifies the verifier’s work.

In this section, we detail the method by which we encode (resp. decode) a set of binary strings to (resp. from) sets of field elements. In the rest of this section, the notion of *packing policy* refers to the set of *packing* and *unpacking* functions.

The set of primary inputs is composed of the input nullifiers, the output commitments, the public values (see [RZ19, Section 3.4.3]) along with the signature hash and the authentication tags for security (malleability fix, see Appendix A). The complete description of the public inputs is represented in Eq. (3.2).

$$(\{prim.nf_i\}_{i \in [JSIN]}, \{prim.cms[j]\}_{j \in [JSOUT]}, vin, vout, hsig, \{prim.htags[i]\}_{i \in [JSIN]}) \quad (3.2)$$

The primary inputs that consist of binary strings are: the nullifiers *nfs*, the public values *vin* and *vout*, the signature hash *hsig* and the authentication tags *htags*.

For a binary string  $x$ , let  $\alpha_x = \lceil \text{length}(x) / \text{FIELD CAP} \rceil$  be the number of field elements required to completely encode  $x$  and let  $\beta_x = \lfloor \text{length}(x) / \text{FIELD CAP} \rfloor$  be the number of field elements whose capacity is fully used. Let  $\gamma_x = \text{length}(x) \pmod{\text{FIELD CAP}}$  be the number of “residual” bits remaining after fully using  $\beta_x$  field elements.

**Example 3.3.2.** Consider binary strings  $A \in \{0, 1\}^7$  of length 7, to be encoded over the field  $\mathbb{F}_{41}$ . This field has a capacity of 5 bits, and therefore  $\alpha_A = 2$ ,  $\beta_A = 1$ , and  $\gamma_A = 2$ . That is,  $A$  can be represented as 2 field elements, or as 1 field element with 2 “residual” bits.

Consider  $A = (1111011)$ . Fig. 3.6 illustrates how  $A$  can be packed as field elements. Note that the 2 residual bits are taken from the “beginning” of the bit string, that is, the highest order bits.




Figure 3.6: Packing of string  $A$  (see Example 3.3.2)

We now consider strategies to pack all primary inputs that are binary strings. A naive approach is to encode each binary string  $x$  as  $\alpha_x$  field elements. In general, this results in significant waste (and consequently more field elements than necessary), especially when the number of residual bits is small compared to **FIELD CAP** (see Fig. 3.7). An alternative strategy could be to concatenate all binary strings into a single string  $y$  and

1293 pack this string into  $\alpha_y$  field elements. While this approach minimizes the set of unused  
 1294 bits, each unpack operation would require different shift and mask operations over 2 or  
 1295 3 field elements. This significantly increases the complexity of the unpacking operation  
 1296 that must be performed on-chain, resulting in a higher gas cost (due to extra logic) or  
 1297 more contract code (if each unpack operation is hard-coded).

The **Zeth** protocol requires that each binary string variable  $x$  is packed into  $\beta_x$  field elements, and the residual bits from all binary strings, along with the public values  $vin$  and  $vout$ , are aggregated into a variable  $rsd$ . Let **RSDBLEN** be the total number of residual bits, and **RSDFLEN** be the number of field elements required to represent  $rsd$ . We assume that **ZVALUELEN** < **FIELD CAP**, and define the notation  $\gamma_v = \text{ZVALUELEN}$  for the bit lengths of public values  $vin$  and  $vout$ . Thus **RSDBLEN** is given by

$$\text{RSDBLEN} = \gamma_{hsig} + 2 \cdot \gamma_v + \text{JSIN} \cdot (\gamma_{nf} + \gamma_h)$$

and the lengths, in field elements, of each of the corresponding public inputs are

$$\begin{aligned} \text{NFFLEN} &= \beta_{nf} \\ \text{HSIGFLEN} &= \beta_{hsig} \\ \text{HFLEN} &= \beta_h \\ \text{RSDFLEN} &= \lceil \text{RSDBLEN} / \text{FIELD CAP} \rceil \end{aligned}$$

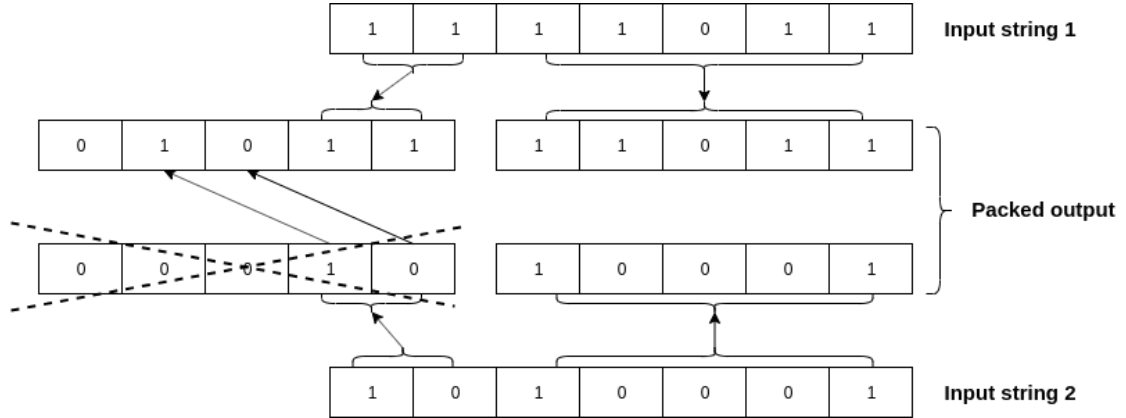


Figure 3.7: Packing of multiple strings. Observe that, by carefully arranging the bits of the input strings, it is possible to output fewer field elements

The residual bits  $rsd$  are formatted as follows:

$$\widetilde{hsig} \parallel \widetilde{nfs} \parallel \widetilde{htags} \parallel vin \parallel vout$$

1298 where  $\widetilde{hsig}, \widetilde{nfs}, \widetilde{htags}$  are, respectively, the  $\gamma_{hsig}, \gamma_{nf}, \gamma_h$  bits.

1299 Note that the public values are packed into the “last”, or lowest order,  $2 \cdot \gamma_v$  bits of  
 1300 the resulting field element(s). In this way, their unpack functions are independent of the

---

**Pack** <sub>$\mathbb{F}_r$</sub> ( $x$ )

---

$out \leftarrow 0_{\mathbb{F}_r};$   
**for**  $i \in [\text{length}(x)]$  **do** :  
    **if**  $x[i] = 1$  **do** :  
         $out \leftarrow out +_{\mathbb{F}_r} 2^{\text{length}(x)-1-i}$   
**return**  $out;$

Figure 3.8: Algorithm to pack bits into a field element.

---

**Pack** <sub>$r_{sd}$</sub> ( $nfs, vin, vout, hsig, htags$ )

---

$out \leftarrow []; r \leftarrow \epsilon;$   
 $r \leftarrow vout;$   
 $r \leftarrow vin || r;$   
**for**  $i \in [\text{JSIN}]$  **do** :  
     $r \leftarrow htags[i][\beta_{htags[i]} \cdot \text{FIELD CAP} : ] || r;$   
**for**  $i \in [\text{JSIN}]$  **do** :  
     $r \leftarrow nfs[i][\beta_{nfs[i]} \cdot \text{FIELD CAP} : ] || r;$   
 $r \leftarrow hsig[\beta_{hsig} \cdot \text{FIELD CAP} : ] || r;$   
**for**  $i \in [\lceil \text{length}(r) / \text{FIELD CAP} \rceil]$  **do** :  
     $out[i] \leftarrow \text{Pack}_{\mathbb{F}_{r_{BN}}}(r[i \cdot \text{FIELD CAP} : (i+1) \cdot \text{FIELD CAP}]);$   
**return**  $out;$

Figure 3.9: Algorithm to pack residual bits.

1301 values JSIN and JSOUT and of the number of residual bits required for each bit string  
1302 (and consequently, independent of the finite field used).

To format the unpacked primary inputs into field elements, we define the following functions. Given a bit string of length less than FIELD CAP, the algorithm Pack (see Fig. 3.8) returns a field element. Given the nullifiers, public values and authentication tags, the algorithm Pack <sub>$r_{sd}$</sub>  (see Fig. 3.9) outputs the residual bits. Given a set of packed field elements and the residual bits, the algorithm Unpack returns the variables reassembled as binary strings. In particular, we have that  $\text{Unpack}_{nf}(\text{prim}.nfs, rsd) = \{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}$ .

$\text{Pack} : \mathbb{B}^{\leq \text{FIELD CAP}} \rightarrow \mathbb{F}_r$

$\text{Pack}_{rsd} : (\mathbb{B}^{\text{PRNFOUTLEN}})^{\text{JSIN}} \times (\mathbb{B}^{\text{ZVALUELEN}})^2 \times \mathbb{B}^{\text{CRHHSIGOUTLEN}} \times (\mathbb{B}^{\text{PRFPKOUTLEN}})^{\text{JSIN}} \rightarrow (\mathbb{F}_r)^{\text{RSDFLEN}}$

$\text{Unpack} : \mathbb{F}_r^* \times (\mathbb{F}_r)^{\text{RSDFLEN}} \rightarrow \mathbb{B}^*$

The Unpack functions for nullifiers, public values and signature hash are represented as follows.

$\text{Unpack}_{hsig} : (\mathbb{F}_r)^{\text{HSIGFLEN}} \times (\mathbb{F}_r)^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{CRHHSIGOUTLEN}}$

$\text{Unpack}_{nf} : (\mathbb{F}_r)^{\text{NFFLEN}} \times (\mathbb{F}_r)^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{PRNFOUTLEN}}$

$\text{Unpack}_{vin} : \mathbb{F}_r^0 \times (\mathbb{F}_r)^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{ZVALUELEN}}$

$\text{Unpack}_{vout} : \mathbb{F}_r^0 \times (\mathbb{F}_r)^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{ZVALUELEN}}$



### 1303 Packing Policy Security

1304 **Proposition 3.3.1** (Packing security). *For a binary string  $x$ , it holds that  $\text{Unpack}(\text{Pack}(x)) =$*   
 1305  *$x$  and  $\text{Unpack}(\text{Pack}_{rsd}(x)) = x$ .*

### 1306 Packing Policy Example

In the case where  $\text{JSIN} = \text{JSOUT} = 2$ , the BN-254 is being used (in which field elements hold  $\text{FIELD CAP}_{\text{BN}}$  bits) and all PRFs and  $\text{CRH}^{\text{hsig}}$  output bit-strings of length 256, the unpacked primary inputs are 2167-bit long. The packing parameters are therefore:

$$\begin{aligned}\text{RSDBLEN} &= 5 \times 3 + 64 + 64 = 143 \\ \text{NFFLEN} &= \text{HSIGFLEN} = \text{HFLEN} = \text{RSDFLN} = 1\end{aligned}$$

The packed primary inputs are 2277 bits long, corresponding to a small space overhead of  $\approx 5\%$  unused bits. Moreover, the 143-bit residual bits can be packed into a single field element. As such, the primary inputs are encoded as 9 field elements. Finally, the residual bits are formatted as follows,

$$\underbrace{\text{padding}}_{113 \text{ bits}} \parallel \underbrace{\text{hsig}}_{3 \text{ bits}} \parallel \underbrace{\text{nf}_1}_{3 \text{ bits}} \parallel \underbrace{\text{nf}_0}_{3 \text{ bits}} \parallel \underbrace{h_1}_{3 \text{ bits}} \parallel \underbrace{h_0}_{3 \text{ bits}} \parallel \underbrace{\text{vin}}_{64 \text{ bits}} \parallel \underbrace{\text{vout}}_{64 \text{ bits}}$$

For the analogous case using  $\text{BLS12-377}$  (in which field elements hold  $\text{FIELD CAP}_{\text{BLS}}$  bits), the packing parameters are:

$$\begin{aligned}\text{RSDBLEN} &= 5 \times 4 + 64 + 64 = 148 \\ \text{NFFLEN} &= \text{HSIGFLEN} = \text{HFLEN} = \text{RSDFLN} = 1\end{aligned}$$

The residual bits can be packed into a single field element of the form

$$\underbrace{\text{padding}}_{108 \text{ bits}} \parallel \underbrace{\text{hsig}}_{4 \text{ bits}} \parallel \underbrace{\text{nf}_0}_{4 \text{ bits}} \parallel \underbrace{\text{nf}_1}_{4 \text{ bits}} \parallel \underbrace{h_0}_{4 \text{ bits}} \parallel \underbrace{h_1}_{4 \text{ bits}} \parallel \underbrace{\text{vin}}_{64 \text{ bits}} \parallel \underbrace{\text{vout}}_{64 \text{ bits}}$$

1307 and the primary inputs are again encoded as 9 field elements.

## 1308 3.4 Instantiate $\text{SigSch}_{\text{OT-SIG}}$

1309 **Zeth** uses the one-time Schnorr-based signature scheme introduced by Bellare and Shoup [BS07]  
 1310 for its long proven security, simplicity, speed and size. Its security relies on the one-more  
 1311 discrete log problem (see Definition 1.5.6) and the collision resistance of the underlying  
 1312 hash function  $\text{CRH}$  (see Definition 1.5.16) that we instantiate with  $\text{SHA256}$ .

1313 Note that no signature operations or data are used in the arithmetic circuit describing  
 1314 the **Zeth** statement. Hence the curve used for the signature scheme can be chosen  
 1315 independently of **Curve** (the scalar field of which is used for the arithmetic circuit, and  
 1316 consequently for commitments and bit string encodings described in Section 3.1 and

1317 Section 3.2). BN-254 is used since it is supported by the EVM, in the form of precompiled  
 1318 contracts. This allows a gas-efficient implementation in the **Mixer** contract.

1319 This one-time signature scheme (see Definition 1.5.26) is defined by the two-tier  
 1320 signature scheme over a cyclic group  $(p, \mathbb{G}, \langle \mathbf{g} \rangle, \otimes)$ . In the two-tier signature scheme, the  
 1321 hash function CRH only needs to be collision resistant (the random oracle model is not  
 1322 used). Similarly, the variable  $hk$  represents the key of the hash function (a particular  
 1323 instance).

1324 To turn this two-tier signature scheme into a one-time signature scheme, one simply  
 1325 has to define the one-time signature key generation KGen as the combination of both  
 1326 primary and secondary key generations of the two-tier (see [BS07, Section 6]). The  
 1327 one-time signing key (respectively verification key) of the one time signature scheme is  
 1328 defined as both the primary and secondary signing key (respectively verification key) of  
 1329 the two-tier scheme, Fig. 3.10

<u>KGen(<math>1^\lambda</math>) :</u>	<u>Sig(<math>sk, m</math>) :</u>	<u>Vf(<math>pk, m, \sigma</math>) :</u>
$hk \leftarrow \$ \mathbb{B}^{kl}$	$hk, \mathbf{g}, x = sk.sk1$	$hk, \mathbf{g}, \llbracket x \rrbracket = pk.pk1$
$\mathbf{g} \leftarrow \$ \mathbb{G}^*$	$y, \llbracket y \rrbracket = sk.sk2$	$\llbracket y \rrbracket = pk.pk2$
$x \leftarrow \$ \mathbb{F}_p$	$c = \text{CRH}(hk, \llbracket y \rrbracket \  m)$	$c = \text{CRH}(hk, \llbracket y \rrbracket \  m)$
$pk1 = (hk, \mathbf{g}, \llbracket x \rrbracket)$	$\sigma = y \bmod p$	<b>if</b> $\sigma = \llbracket y \rrbracket \otimes c \cdot \llbracket x \rrbracket$ <b>then</b>
$sk1 = (hk, \mathbf{g}, x)$	$\sigma += c \cdot x \bmod p$	<b>return</b> 1
$y \leftarrow \$ \mathbb{F}_p$	<b>return</b> $\sigma$	<b>else</b>
$pk2 = \llbracket y \rrbracket$		<b>return</b> 0
$sk2 = (y, \llbracket y \rrbracket)$		<b>endif</b>
$pk = (pk1, pk2)$		
$sk = (sk1, sk2)$		

Figure 3.10: One-time signature scheme from two tier Schnorr based signature scheme by Bellare and Shoup [BS07]

### 1330 3.4.1 Security requirements satisfaction

1331 We now prove that this signature scheme satisfies all the security requirements listed  
 1332 in Section 2.7.

1333 **Theorem 3.4.1.** *The One-Time Schnorr signature is strongly unforgeable under chosen-*  
 1334 *message attacks (SUF-CMA) assuming that the om-DLog problem is hard in  $\mathbb{G}$  and that*  
 1335 *the hash function CRH is collision resistant.*

1336 *Proof.* See [BS07, Theorems 5.1, 5.2 and 6.1]. □

### 1337 3.4.2 Data types

1338 We now describe the data types and operations associated with this signature scheme.

1339 **VK0tsDType** Denotes the verification key associated with the one-time signature scheme.

Field	Description	Data type
$pk1$	Encoding of the scalar $x$ in the group	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^2$
$pk2$	Encoding of the scalar $y$ in the group	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^2$

Table 3.2: **VK0tsDType** data type

1340 **SK0tsDType** Denotes the signing key associated with the one-time signature scheme.

Field	Description	Data type
$sk1$	Scalar element $x$	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}$
$sk21$	Scalar element $y$	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}$
$sk22$	Encoding of the scalar $y$ in the group	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^2$

Table 3.3: **SK0tsDType** data type

1341 **Sig0tsDType** Denotes the signature data type associated with the one-time signature  
 1342 scheme. **Sig0tsDType** is an alias for  $(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^2$ .

## 1343 3.5 Instantiate EncSch

1344 In this section we describe the instantiation of **EncSch** primitive introduced in Section 2.3.  
 1345 First, we present a general asymmetric encryption scheme called **DHAES** (Diffie-Hellman  
 1346 Asymmetric Encryption Scheme [ABR99]), which satisfies all the required security prop-  
 1347 erties for the in-band encryption scheme **EncSch** (see Section 1.5.3). Then, we give details  
 1348 of the concrete algorithms used for the implementation.

### 1349 3.5.1 DHAES encryption scheme

1350 Given a symmetric encryption scheme **Sym**, a group defined by **SetupG**, a family of hash  
 1351 function  $\mathcal{H}^3$  and a message authentication scheme **MAC** as defined in Section 1.5, we  
 1352 define a **DHAES** scheme as the following public-key encryption scheme:

- 1353 • **Setup**, setup algorithm, takes as input a security parameter  $1^\lambda$ . It runs  $\mathcal{H}.\text{Setup}$ ,  
 1354 **SetupG** and returns public parameters  $pp = (hk, (q, \mathbb{G}, \mathbf{g}, +))$ .

<sup>3</sup>Here, we only consider fixed-length hash functions with  $h\text{InpLen}(\lambda) = 2g\text{Len}$  and  $h\text{Len}(\lambda) = k\text{Len}(\lambda) + m\text{Len}(\lambda)$  (see Section 1.5).

- 1355 • KGen, key generation algorithm, takes as input public parameters  $pp$ . It samples  
1356 at random  $v \leftarrow \mathbb{S}[q]$  and returns a keypair  $(sk, pk) = (v, \llbracket v \rrbracket)$ .
- 1357 • Enc, encryption algorithm, takes as input public parameters  $pp$ , a message  $m$  and  
1358 a public key  $pk$ . It runs KGen that returns an ephemeral keypair  $(esk, epk) =$   
1359  $(u, \llbracket u \rrbracket)$ . Then, it computes a shared secret  $ss = H_{hk}(epk \parallel esk \cdot pk) = H_{hk}(epk \parallel sk \cdot$   
1360  $epk)$ , parsed as  $ek \parallel mk$ <sup>4</sup>. It computes  $ct_{\text{Sym}} = \text{Sym.Enc}(ek, m)$  and  $\tau = \text{MAC.Tag}(mk, ct_{\text{Sym}})$   
1361 and finally outputs the ciphertext  $epk \parallel ct_{\text{Sym}} \parallel \tau$ .
- 1362 • Dec, decryption algorithm, takes as input public parameters  $pp$ , a private key  $sk$   
1363 a ciphertext  $epk \parallel ct_{\text{Sym}} \parallel \tau$ . It computes  $ss = H_{hk}(epk \parallel sk \cdot epk)$  and parses it, as  
1364 above, as  $ek \parallel mk$ . If MAC verification passes, i.e.  $\text{MAC.Vf}(mk, \tau) = 1$ , the algorithm  
1365 returns  $\text{Sym.Dec}(ek, ct_{\text{Sym}})$  and  $\perp$  otherwise.

1366 The DHAES definition given above is an asymptotic adaptation of [ABR99, Section  
1367 1.3].

### 1368 Inclusion of ephemeral key in hash input

1369 Given an ephemeral keypair  $(u_0, \llbracket u_0 \rrbracket)$ , If the group  $\langle \mathbf{g} \rangle$ , generated by SetupG, has com-  
1370 posite order, then  $\llbracket u_0 \rrbracket$  is required to be part of the hash input because  $\llbracket u_0 v \rrbracket$  and  $\llbracket v \rrbracket$   
1371 together may not uniquely determine  $\llbracket u_0 \rrbracket$ . Equivalently, there may exist two values  
1372  $u_0$  and  $u_1$  such that  $u_0 \neq u_1$  and  $\llbracket u_0 v \rrbracket = \llbracket u_1 v \rrbracket$ . As a result, both  $u_0$  and  $u_1$  can be  
1373 used to produce two different *valid* ciphertexts of the same plaintext  $m$ , under different  
1374 ephemeral keys  $(\llbracket u_0 \rrbracket, \llbracket u_1 \rrbracket)$ . It is easy to show this, for example, in the multiplicative  
1375 group  $\mathbb{Z}_p \setminus \{0\}$ , where  $p$  is a prime (see [ABR99, Section 3.1]). A scheme having such  
1376 malleability property clearly cannot be proven IND-CCA2 secure: an attacker could eas-  
1377 ily win the related security game by altering the challenged ciphertext and query the  
1378 decryption oracle that would not recognize that as a not allowed query. If the group  
1379 has prime order this problem does not arise so only  $\llbracket u_0 v \rrbracket$  is required as input of the  $H$   
1380 function [ABR01, Section 3].

### 1381 3.5.2 A DHAES instance

#### 1382 Curve25519

1383 For a cyclic group we propose the use of a subgroup of Curve25519 described in [Ber06]  
1384 and in [LHT16]. Curve25519 is a Montgomery elliptic curve [Mon87] defined by the  
1385 equation  $y^2 = x^3 + 486662x^2 + x$  and coordinates on  $\mathbb{F}_p$ , where  $p$  is the prime number  
1386  $2^{255} - 19$ . It has a prime order subgroup of order  $2^{252} + 27742317777372353535851937$   
1387  $790883648493$  and cofactor 8. Curve25519 comes with an efficient scalar multiplication  
1388 denoted as X25519<sup>5</sup>. In a Diffie-Hellman-based scheme it allows to have 32-byte long

<sup>4</sup>Note that  $ek$  and  $mk$  must have the same length.

<sup>5</sup>X25519 is actually introduced in [LHT16] in order to avoid notation issues due to the use Curve25519 to indicate both curve and scalar multiplication as done in [Ber06]

1389 public and private keys (given a point  $P = (x, y)$  only the  $x$  coordinate is actually used)  
1390 and the 32-byte sequence representing 9 is specified as base point.

## 1391 Efficiency and security of Curve25519

1392 High-speed and timing-attack resistant implementations of X25519 are available and  
1393 its security level is conjectured to be 128 bits [Ber06, Section 1]. However, combined  
1394 attacks can lead to 124 bits of security (see [BL, Section “Twist Security”]). By design,  
1395 Curve25519 is resistant to state-of-the-art attacks and satisfies all security criteria and  
1396 principles listed in *Safecurves* [BL]<sup>6</sup>.

1397 Interestingly, Curve25519 does not require *public key validation*<sup>7</sup>, while we know that,  
1398 on other curves, active attacks – consisting of sending malformed public keys – could be  
1399 carried out by adversaries, to violate the confidentiality of private keys, e.g. [ABM<sup>+</sup>03].  
1400 However, Curve25519 specification mandates the *clamping* of private keys: that is, after  
1401 the random sampling of 32 bytes, the user clears bits 0, 1 and 2 of the first byte, clears  
1402 bit 7 and sets bit 6 of the last byte. The resulting 32 bytes are then used as private key.  
1403 This particular structure for private keys prevents various types of attacks (see [Ber06,  
1404 Section 3] for more details).

### Note

Note that the *clamping* procedure is vital to ensure the security guarantees of the Curve25519 specification, and implementations **MUST** perform this exactly as described.

1405

## 1406 ChaCha20

1407 ChaCha20 is an ARX-based<sup>8</sup> stream cipher introduced in [Ber08a]. It is an improved  
1408 version of Salsa20 [Ber08b] that won the *eSTREAM* challenge [est]. Compared with  
1409 Salsa20, it has been designed to improve diffusion per round, conjecturally increasing  
1410 resistance to cryptanalysis, while preserving time efficiency per round. It is considerably  
1411 faster than AES in software-only implementations and can be easily implemented to be  
1412 timing-attacks resistant. Several versions of the cipher can be used. The original paper  
1413 presents ChaCha20 with a 128-bit key and 64-bit nonce/block count. However, the length  
1414 of the key, nonce and block count – which indicates how many chunks can be processed  
1415 by using the same key and nonce – can be modified depending on the application.  
1416 In [LN18][Section 2.3], for instance, the key is a 256-bit string, the nonce is a string of  
1417 96 bits and the block count is encoded on a 32-bit word. This configuration allows to

<sup>6</sup>In this work, the authors take into account both Elliptic Curve Discrete Logarithm Problem (ECDLP) and Elliptic Curve Cryptosystems (ECC) security, that allows to have an overall evaluation of the security guarantees.

<sup>7</sup>Informally, it is a set of security checks that a user performs before using a not trusted public key (e.g. see [BCK<sup>+</sup>18])

<sup>8</sup>Addition-Rotation-XOR

1418 process around  $2^{32}$  blocks, corresponding to roughly 256 GB of data. We propose to use  
 1419 the same parameters in Zeth.

$$\text{ChaCha20} : \mathbb{B}^{256} \times \mathbb{B}^{32} \times \mathbb{B}^{96} \times \mathbb{B}^* \rightarrow \mathbb{B}^*$$

## 1420 Security of Chacha

1421 Recent cryptanalysis results for ChaCha are available in [AFK<sup>+</sup>08, Ish12, SZFW12,  
 1422 Mai16, CM16, CM17]: all of them make use of advanced cryptanalysis techniques able  
 1423 to perform key-recovery attacks only on reduced versions (6 and 7 rounds) of ChaCha.

### Note

Importantly, the security properties of ChaCha rely on the fact that, for a given key, all blocks are processed with distinct values in the state words 12 to 15 (storing the counter and the nonce) [LN18, Section 2.3].

1424

## 1425 Poly1305

1426 Poly1305 [Ber05] is a high-speed message authentication code, easy to implement and  
 1427 make side-channel attack resistant. It takes a 32-byte one-time key  $mk$  and a message  $m$   
 1428 and produces a 16-byte tag  $\tau$  that authenticates the message.  $mk$  must be unpredictable  
 1429 and it is represented as a couple  $(r, s)$ , where both components are given as a sequence  
 1430 of 16 bytes each. It can be generated by using pseudorandom algorithms: in [Ber05,  
 1431 Section 2], for example, AES and a nonce are used to generate  $s$ . The second part of  
 1432 the key,  $r$ , is expected to have a given form [Ber05, Section 2], and must be “clamped”  
 1433 as follows: top four bits of  $r[3]$ ,  $r[7]$ ,  $r[11]$ ,  $r[15]$  and bottom two bits of  $r[4]$ ,  $r[8]$ ,  $r[12]$   
 1434 are cleared (see also Section 3.5.3).

### Note

Similarly to Curve25519, the *clamping* procedure here is essential to the security of the Poly1305 scheme. Implementations MUST ensure that this is performed correctly in order for all security guarantees to hold.

1435

1436 We refer to [LN18, Section 2.5, Section 3] for Tag and Vf implementations of Poly1305.

$$\text{Poly1305.Tag} : \mathbb{B}_Y^{32} \times \mathbb{B}_Y^* \rightarrow \mathbb{B}_Y^{16}$$

$$\text{Poly1305.Vf} : \mathbb{B}_Y^{32} \times \mathbb{B}_Y^{16} \times \mathbb{B}_Y^* \rightarrow \mathbb{B}$$

## 1437 Security of Poly1305

1438 Citing Poly1305 [LN18, Section 4], “the Poly1305 authenticator is designed to ensure that  
 1439 forged messages are rejected with a probability of  $1 - (n/(2^{102}))$  for a  $16n$ -byte message,

1440 even after sending  $2^{64}$  legitimate messages, so it is SUF-CMA (strong unforgeability  
1441 against chosen-message attacks)."

### 1442 **Blake2b-512**

1443 Since we need a total of 64 bytes for the key material (32 for ChaCha20 and 32 for  
1444 Poly1305) Blake2b512 can be used. ZCash protocol [ZCa19, Section 5.4.3], instead, makes  
1445 use of Blake2b256 since a DHAES variant, denoted as ChaCha20-Poly1305, is adopted  
1446 (see [LN18, Section 2.8]).

$$\text{Blake2b512} : \mathbb{B}^* \rightarrow \mathbb{B}_Y^{32}$$

### 1447 **3.5.3 EncSch instantiation**

In the following we instantiate EncSch as a DHAES scheme, detailing the KGen, Enc and Dec components. First, we introduce some required constant values:

$$\begin{aligned} \text{ESKBYTELEN} &= 32 \\ \text{EPKBYTELEN} &= 32 \\ \text{NOTEBYTELEN} &= (\text{PRFADDRROUTLEN} + \text{RTRAPLEN} + \text{ZVALUELEN} + \text{PRFRHOOUTLEN})/\text{BYTELEN} \\ \text{SYMKEYBYTELEN} &= 32 \\ \text{MACKEYBYTELEN} &= 32 \\ \text{KDFDIGESTBYTELEN} &= \text{SYMKEYBYTELEN} + \text{MACKEYBYTELEN} \\ \text{CTBYTELEN} &= \text{EPKBYTELEN} + \text{NOTEBYTELEN} + \text{TAGBYTELEN} \\ \text{TAGBYTELEN} &= 16 \\ \text{CHACHANONCEVALUE} &= 0^{32} \\ \text{CHACHABLOCKCOUNTERVALUE} &= 0^{96} \end{aligned}$$

### 1448 **EncSch.KGen**

1449 The keypair  $(sk, pk)$  generation is defined as:

- 1450 • Randomly sample a sequence of ESKBYTELEN bytes and assign to  $sk$ .
- Clamp  $sk$  as follows:

$$\begin{aligned} sk[0] &\leftarrow sk[0] \& 0xF8 \\ sk[31] &\leftarrow sk[31] \& 0x7F \\ sk[31] &\leftarrow sk[31] \mid 0x40 \end{aligned}$$

1451 where  $|$  and  $\&$  denotes, respectively, OR and AND binary operators between bit  
 1452 strings of same the length.<sup>9</sup>

- 1453 • Compute  $pk = \text{X25519}(sk, 0x09)$ .
- 1454 • Return  $(sk, pk) \in \mathbb{B}_Y^{\text{ESKBYTELEN}} \times \mathbb{B}_Y^{\text{EPKBYTELEN}}$

1455 **EncSch.Enc**

1456 The encryption, on inputs  $(pk, m) \in \mathbb{B}_Y^{\text{EPKBYTELEN}} \times \mathbb{B}_Y^{\text{NOTEBYTELEN}}$ , is defined as follows:

- 1457 1. Generate an ephemeral Curve25519 keypair  $(esk, epk) \in \mathbb{B}_Y^{\text{ESKBYTELEN}} \times \mathbb{B}_Y^{\text{EPKBYTELEN}}$   
 1458 (as above).
2. Compute the shared secret<sup>10</sup>  $ss \in \mathbb{B}_Y^{\text{EPKBYTELEN}}$ :

$$ss = \text{X25519}(esk, pk) \in \mathbb{B}_Y^{\text{EPKBYTELEN}}$$

3. Generate a session key:

$$\text{Blake2b512}(\text{encTag} \| epk \| ss) \in \mathbb{B}_Y^{\text{KDFDIGESTBYTELEN}}$$

where  $\text{encTag} = 0x5A \| 0x65 \| 0x74 \| 0x68 \| 0x45 \| 0x6E \| 0x63$ , that is the UTF-8 encoding of “*ZethEnc*” string (used for domain separation purposes). The result, then, is parsed as follows:

$$\begin{aligned} ek &= \text{Blake2b512}(\text{encTag} \| epk \| ss)[\text{SYMKEYBYTELEN} - 1] \\ mk &= \text{Blake2b512}(\text{encTag} \| epk \| ss)[\text{SYMKEYBYTELEN} : \text{SYMKEYBYTELEN} + \text{MACKEYBYTELEN} - 1]. \end{aligned}$$

4. Encrypt the confidential data:

$$ct_{\text{Sym}} = \text{ChaCha20}(ek, \text{CHACHABLOCKCOUNTERVALUE}, \text{CHACHANONCEVALUE}, m) \in \mathbb{B}^{\text{NOTEBYTELEN} * \text{BYTELEN}}$$

1459 **Remark 3.5.1.** Formally speaking we should have written  $ct_{\text{Sym}} \in \mathbb{B}^n$ , where  
 1460  $n$  is the length of binary representation of the encrypted message  $m$ . In **Zeth**  
 1461 however, the only data encrypted are the notes. As such, the size of the plaintexts  
 1462 is  $\text{NOTEBYTELEN} * \text{BYTELEN}$  bits.

1463 **Remark 3.5.2.** In the following, we omit the explicit conversion from  $\mathbb{B}^n$  to  
 1464  $\mathbb{B}_Y^{[n/\text{BYTELEN}]}$  when passing the output of ChaCha20 to the Poly1305 algorithms.

5. Randomly generate  $(r, s) \in \mathbb{B}_Y^{\text{MACKEYBYTELEN}/2} \times \mathbb{B}_Y^{\text{MACKEYBYTELEN}/2}$  and clamp it:

$$r[3] \leftarrow r[3] \& 0x0F$$

<sup>9</sup>E.g Given two bytes  $0x15$  and  $0x03$  then  $0x15|0x03 = 0x17$  and  $0x15\&0x03 = 0x01$ .

<sup>10</sup>We assume here that  $esk$  has been clamped as discussed in Section 3.5.2



$$\begin{aligned}
r[7] &\leftarrow r[7] \& 0x0F \\
r[11] &\leftarrow r[11] \& 0x0F \\
r[15] &\leftarrow r[15] \& 0x0F \\
r[4] &\leftarrow r[4] \& 0xFC \\
r[8] &\leftarrow r[8] \& 0xFC \\
r[12] &\leftarrow r[12] \& 0xFC
\end{aligned}$$

6. Generate the related tag:

$$\tau = \text{Poly1305.Tag}(mk, ct_{\text{Sym}}) \in \mathbb{B}_{\mathbb{Y}}^{\text{TAGBYTELEN}}.$$

7. Create the asymmetric ciphertext as:

$$ct = epk \| ct_{\text{Sym}} \| \tau \in \mathbb{B}_{\mathbb{Y}}^{\text{CTBYTELEN}}.$$

1465 8. Return  $ct$ . As consequence  $\text{ENCZETHNOTELEN} = \text{CTBYTELEN} * \text{BYTELEN}$  bits.

1466 **EncSch.Dec**

1467 The decryption, on inputs  $(sk, ct) \in \mathbb{B}_{\mathbb{Y}}^{\text{ESKBYTELEN}} \times \mathbb{B}_{\mathbb{Y}}^{\text{CTBYTELEN}}$ , is defined as follows:

1. Parse the ciphertext  $ct$  as:

$$\begin{aligned}
epk &\leftarrow ct[: \text{EPKBYTELEN} - 1] \\
ct_{\text{Sym}} &\leftarrow ct[\text{EPKBYTELEN} : \text{EPKBYTELEN} + \text{NOTEBYTELEN} - 1] \\
\tau &\leftarrow ct[\text{EPKBYTELEN} + \text{NOTEBYTELEN} : \text{EPKBYTELEN} + \text{NOTEBYTELEN} + \text{TAGBYTELEN} - 1]
\end{aligned}$$

2. Recover the shared secret

$$ss = \text{X25519}(sk, epk).$$

3. Compute the  $ek \| mk$

$$\begin{aligned}
ek &= \text{Blake2b512}(\text{encTag} \| epk \| ss)[: \text{SYMKEYBYTELEN} - 1] \\
mk &= \text{Blake2b512}(\text{encTag} \| epk \| ss)[\text{SYMKEYBYTELEN} : \text{SYMKEYBYTELEN} + \text{MACKEYBYTELEN} - 1].
\end{aligned}$$

4. Verify that the ciphertext has not been forged:

$$\text{Poly1305.Vf}(mk, \tau, ct_{\text{Sym}})$$

5. (If the MAC verifies) decrypt:

$$m = \text{ChaCha20.Dec}(ek, \text{CHACHABLOCKCOUNTERVALUE}, \text{CHACHANONCEVALUE}, ct_{\text{Sym}})$$

1468 6. Return  $m$ .

### 3.5.4 Security requirements satisfaction

DHAES has already been proved to be IND-CCA2 secure (see [ABR99, Section 3.5, Theorem 6])<sup>11</sup> and to the best of our knowledge there is no paper showing IK-CCA security. The only proof we have found is related to DHIES scheme [ABN10], that is a prime order group version of DHAES. In the following, we provide a proof for IK-CCA security of DHAES by adapting that proof to our case.

**Theorem 3.5.1** (IK-CCA of DHAES). *Let DHAES be the asymmetric encryption scheme as defined above. Let  $\mathcal{A}$  be an adversary for the IK-CCA game, then there exists a HDHI adversary  $\mathcal{B}$  of  $(\mathcal{H}, \text{SetupG})$  and a SUF-CMA adversary  $\mathcal{C}$  of MAC such that*

$$\text{Adv}_{\text{DHAES}, \mathcal{A}}^{\text{ik-cca}}(\lambda) \leq 2 \cdot \text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{B}}^{\text{hdhi}}(\lambda) + \text{Adv}_{\text{MAC}, \mathcal{C}}^{\text{suf-cma}}(\lambda).$$

The adversaries  $\mathcal{B}$  and  $\mathcal{C}$  have the same running time as  $\mathcal{A}$ <sup>12</sup>.

*Informal proof.* As already mentioned, DHAES is similar to DHIES scheme, except for the underlying group and the way the symmetric keys are constructed. As consequence, IK-CCA property for DHAES can be shown similarly to the approach in [ABN10, Theorem 6.2]. More precisely, they show that one can construct from an attacker  $\mathcal{A}$  for the IK-CCA game two attackers  $\mathcal{B}$  and  $\mathcal{C}$  for the ODH and SUF-CMA games. Actually, they make use of a  $\bar{\mathcal{B}}$  attacker for the ODH2 game [ABN10, Figure 20] and then apply [ABN10, Lemma 6.1] to obtain an attacker  $\mathcal{B}$ <sup>13</sup> in the ODH game. We adopt a similar strategy, working with HDHI, HDHI2 and Lemma 1.5.1.

Let  $\mathcal{A}$  be an attacker for the IK-CCA game, and let  $\bar{\mathcal{B}}$  be an attacker for the HDHI2 game described in Fig. 3.11. We show that,

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \bar{\mathcal{B}}}^{\text{hdhi2}}(\lambda) = |\Pr[\text{IK-CCA}^{\mathcal{A}}(\lambda) = 1] + \Pr[\text{G}_0^{\mathcal{A}}(\lambda) = 1] - 1|$$

where  $\text{G}_0$  is the security game described in Fig. 3.12.

Given an HDHI2 challenge  $(\llbracket u \rrbracket, \llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket, w_{b_2,0}, w_{b_2,1})$ , an adversary  $\bar{\mathcal{B}}$  samples  $b \leftarrow \{0, 1\}$  and runs  $\mathcal{A}$  on  $\llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket$  (note that  $b_2$  is the random bit chosen by the  $\bar{\mathcal{B}}$  challenger in the HDHI2 game).  $\bar{\mathcal{B}}$  constructs oracles  $\text{O}^{\text{Dec}_{sk_i}}$  where the queries  $(\tau \| ct_{\text{Sym}} \| \tau)$  are processed as follows: if  $\tau \neq \llbracket u \rrbracket$ , then  $\bar{\mathcal{B}}$  queries related HDHI2 oracle to obtain  $ek \| mk \leftarrow \text{O}^{\text{HDHI}_{v_i}}(\tau)$  (see Fig. 3.11). If  $\tau = \llbracket u \rrbracket$ ,  $w_{b_2,i}$  is parsed as  $ek \| mk$ . In both cases, it checks that  $\text{MAC.Vf}(mk, ct_{\text{Sym}}, \tau) = 1$  and, if so, returns  $m \leftarrow \text{Sym.Dec}(ek, ct_{\text{Sym}})$ . We note that  $\mathcal{A}$  cannot query the challenged ciphertext.  $\bar{\mathcal{B}}$  returns 0 if and only if  $b = \tilde{b}$ . It easy to see that if  $b_2$  is equal to 0, then all symmetric encryption and MAC keys used for the challenge ciphertext  $(\tau^* \| ct_{\text{Sym}}^* \| \tau^*)$  and decryption responses are exactly as in a DHAES game.

<sup>11</sup>Specifically, if Sym is IND-CPA secure, it holds that H is HDHI secure and MAC is SUF-CMA secure.

<sup>12</sup>In order to give an asymptotic version of the theorem, the number of queries  $q$  has been substituted by the fact of considering PPT adversaries.

<sup>13</sup>Note that in [ABN10] the IK-CCA game is a particular case of the AI-CCA game that requires two input messages in the LR query. In order to reason only about the key-privacy, the two messages  $m_0$  and  $m_1$  are constrained to be equal.

Adversary $\bar{\mathcal{B}}(\llbracket u \rrbracket, \llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket, w_{b_2,0}, w_{b_2,1})$	$\bar{\mathcal{B}}$ simulation of $\mathcal{O}^{\text{Dec}_{sk_i}}(\tau \parallel ct_{\text{Sym}} \parallel \tau)$
$b \leftarrow \$\{0, 1\}$	<b>if</b> $\tau \neq \llbracket u \rrbracket$
$(m, state) \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(\llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket)$	$ek \parallel mk \leftarrow \mathcal{O}^{\text{HDH}_{v_i}}(\tau)$
$ek \parallel mk \leftarrow w_{b_2,b}$	<b>else</b>
$\tau^* \leftarrow u$	$ek \parallel mk \leftarrow w_{b_2,i}$
$ct_{\text{Sym}}^* \leftarrow \text{Sym.Enc}(ek, m)$	<b>fi</b>
$\tau^* \leftarrow \text{MAC.Tag}(mk, ct_{\text{Sym}}^*)$	<b>if</b> $\text{MAC.Vf}(mk, ct_{\text{Sym}}, \tau) = 1$
$\tilde{b} \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(\tau^* \parallel ct_{\text{Sym}}^* \parallel \tau^*, state)$	<b>return</b> $\text{Sym.Dec}(ek, ct_{\text{Sym}})$
<b>return</b> $\tilde{b} = b$	<b>else</b>
	<b>return</b> $\perp$
	<b>fi</b>

Figure 3.11: Description of the adversary  $\bar{\mathcal{B}}$  for HDH12, simulating DHAES game for  $\mathcal{A}$ .

If  $b_2 = 1$ , then  $w_{1,0}$  and  $w_{1,1}$  are random strings and the challenge ciphertext and decryption responses are given as in the  $\mathbf{G}_0$  game described in Fig. 3.12. So we get,

$$\Pr[\text{HDH12}^{\bar{\mathcal{B}}}(\lambda) = 1] = \frac{1}{2} \cdot \Pr[\text{IK-CCA}^{\mathcal{A}}(\lambda) = 1] + \frac{1}{2} \cdot \Pr[\mathbf{G}_0^{\mathcal{A}}(\lambda) = 1].$$

1495 And from the definition of HDH12 advantage we have

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \bar{\mathcal{B}}}^{\text{hdhi2}}(\lambda) = |\Pr[\text{IK-CCA}^{\mathcal{A}}(\lambda) = 1] + \Pr[\mathbf{G}_0^{\mathcal{A}}(\lambda) = 1] - 1|.$$

1496 At this point, we can conclude as in [ABN10, Theorem 6.2], with the only difference  
 1497 of applying Lemma 1.5.1 instead of [ABN10, Lemma 6.1] and by defining a game  $\mathbf{G}_1$   
 1498 that is *identical until bad*<sup>14</sup>  $\mathbf{G}_0$  defined in Fig. 3.12.  $\square$

### 1499 3.5.5 Final notes and observations

1500 In this section we list some notes regarding the approach taken in **Zcash** (see [ZCa19,  
 1501 Section 8.7]), and other observations:

- 1502 • *Key derivation parameters:* in DHAES construction, the only required input vari-  
 1503 ables are the shared secret  $ss$  and  $epk$ . In the Sprout release of **Zcash**, additional  
 1504 parameters were added (i.e.  $h_{sig}$ ,  $pk_{enc}$  and a counter  $i$ ) (see [ZCa19, 5.4.4.2]):  
 1505 they state that  $h_{sig}$  was used in order to get a different randomness extractor for  
 1506 each joinsplit transfer in order to limit the degradation of the security and weaken  
 1507 assumption on the hash. The authors believed, about the use of long-standing  
 1508 public key  $pk_{enc}$ , that it might be necessary for IND-CCA2 security and for post-  
 1509 quantum privacy (in the case where the quantum attacker does not have the public

<sup>14</sup>Games  $\mathbf{G}_i$  and  $\mathbf{G}_j$  are said to be *identical until bad* if they differ only in statements that follow the setting of the **bad** variable to *True*. **bad** is initialized with *False*

$G_0(\lambda)$	Oracle $O^{\overline{\text{Dec}}_{sk_i}}(\tau \  ct_{\text{Sym}} \  \tau)$
$(q, \mathbb{G}, g, +) \leftarrow \text{SetupG}(1^\lambda)$	<b>if</b> $\tau = \tau^*$
$(sk_0, pk_0), (sk_1, pk_1) \leftarrow \$ \text{KGen}(1^\lambda)$	$m \leftarrow \perp$
$\tau^* \leftarrow \$ \mathbb{G}$	<b>if</b> $\text{MAC.Vf}(mk^*, ct_{\text{Sym}}, \tau) = 1$
$ek^* \leftarrow \$ \{0, 1\}^{kLen}$	$\text{bad} \leftarrow \text{true}$
$mk^* \leftarrow \$ \{0, 1\}^{mLen}$	$m \leftarrow \text{Sym.Dec}(ek^*, ct_{\text{Sym}})$
$(m, state) \leftarrow \mathcal{A}^{O^{\overline{\text{Dec}}_{sk_0}}, O^{\overline{\text{Dec}}_{sk_1}}}(pk_0, pk_1)$	<b>fi</b>
$b \leftarrow \$ \{0, 1\}$	<b>else</b>
$ct_{\text{Sym}}^* \leftarrow \text{Sym.Enc}(ek^*, m)$	$m \leftarrow \text{Dec}(sk_i, \tau \  ct_{\text{Sym}} \  \tau)$
$\tau^* \leftarrow \text{MAC.Tag}(mk^*, ct_{\text{Sym}}^*)$	<b>fi</b>
$\tilde{b} \leftarrow \mathcal{A}^{O^{\overline{\text{Dec}}_{sk_0}}, O^{\overline{\text{Dec}}_{sk_1}}}(\tau^* \  ct_{\text{Sym}}^* \  \tau^*, state)$	<b>return</b> $m$
<b>return</b> $\tilde{b} = b$	

Figure 3.12:  $G_0$  game and related decryption oracles for Theorem 3.5.1.

key) [zcaa]. None of these additional components are used any longer starting from the Sapling release (see [ZCa19, 5.4.4.4]). To the best of our knowledge there is no formal reason to use the note counter  $i$  as an input to the KDF: an explanation could be to avoid the same session key being reused for multiple notes, but this should not be a problem since a different nonce or block counter is used for the symmetric cipher (actually this is already mandated in the case where  $epk$  is reused, as described below).

- *Reuse of ephemeral keys  $epk$ :* **Zcash** reuses the same ephemeral keys  $epk$  (and different nonces) for two ciphertexts in a joinsplit description, claiming that this does not affect the security of the scheme as soon as the HDHI assumption of the DHAES security proof is adapted. Note that the proof they refer to is related to the IND-CCA2 notion.

- Note that in **Zcash** Sprout and Sapling, being able to break the Elliptic Curve Diffie-Hellman Problem on Curve25519 or Jubjub would not help to decrypt the transmitted notes ciphertext unless the receiver  $pk_{enc}$  is known or guessed. On the other hand, having  $pk_{enc}$  into the hash (as used in Sprout) may violate in principle the key-privacy of the encryption scheme. For these reasons, we underline that the protocol should enforce a mechanism that does not reveal users public keys to increase the security.

- In [ABN10], the concept of *robustness* for an asymmetric encryption scheme is introduced: it formalizes the infeasibility of producing a ciphertext valid under two different public encryption keys. We note that this is particularly useful for **Zeth** since only the intended receiver will be able to decrypt the encrypted note. In fact, the definition is more general since it also covers the case in which a decryption

1534 is successful but returns an incorrect plaintext. This prevents situations where  
1535 a user, scanning the **Mixer** logs for incoming transactions, gets a false positive  
1536 decryption and stores garbage notes.

#### Note

We note however, that the “false-positive” situation above can be prevented by relying on a weaker notion of robustness called *collision-freeness* [Moh10]. In fact, as described in Section 2.6, the procedure to receive a *ZethNote* requires to decrypt the ciphertext emitted by the **Mixer**, and then to verify that the recovered plaintext is the opening of a commitment in the Merkle tree. As such, since the *collision-freeness* of the encryption ensures that plaintexts recovered under different keys are different (i.e. “do not produce a collision”), then we know that plaintexts recovered by parties who are not the intended recipient will fail the “commitment opening verification”, leading the payment to be rejected, and solving the aforementioned false-positive issue.

1537

1538 In [ABN10, Section 6], the authors prove that DHIES can be made strongly robust.  
1539 The proof can be easily adapted to work with DHAES.

- 1540 • *No public key validation for X25519*: cryptographers have been discussing the ab-  
1541 sence of any mandated public key validation or checks on the result of X25519.  
1542 For example, in [LHT16, Section 6.1], an optional zero check is introduced in order  
1543 to assure that the result of X25519 is not 0: this avoids a situation in which one  
1544 of the two parties can force the result of the key-exchange by using a small order  
1545 point as public key. This property is generally defined as *contributory behaviour*,  
1546 that is, none of the parties is able to force the output of a key exchange. However,  
1547 protocols do not have all the same security requirements and adding default checks  
1548 in the Curve25519 specifications would be superfluous in most cases and would add  
1549 complexities that Bernstein has deliberately chosen to avoid (*simple implementa-*  
1550 *tion principle*). More importantly, Diffie-Hellman does not require *contributory*  
1551 *behaviour* property [Per17]: modern view is that the only requirements are key  
1552 indistinguishability and, in case of an active attacker, that the output of the key  
1553 exchange should not produce a low-entropy function of the honest party’s private  
1554 key (e.g. small-subgroup and invalid-curve attacks). Since these two properties are  
1555 considered satisfied by Curve25519, there is no need to add extra checks to the  
1556 Curve25519 specification. We conclude by observing that in the Sprout release, the  
1557 Zcash protocol does not specify any point validation and makes use only of the  
1558 private key clamping to keep Diffie-Hellman key exchange secure.

### 3.6 ZkSnarkSch instantiation

Groth's proof system Groth16 [Gro16] is the most efficient known zk-SNARK (in terms of the proof size and proof and verification cost) for QAPs, and thus one of the most efficient NIZK for proving statements on arithmetic circuits. Below we present Groth16's key generation, prover, verifier, and simulator algorithms, adjusted as described in [BGM17] to further reduce the size of *srs* and proofs, and to make the KGen algorithm more amenable to implementation as a multi-party computation.

In what follows, let the number *constNo* of constraints in the relation **R** be fixed. Without loss of generality we consider *constNo* to be an *upper bound* on the number of constraints in the **R** parameter, and assume that there exists some *constNo*-th root of unity  $\omega \in \mathbb{F}_r$ . Define  $\ell_i(X)$  to be the *i*-th Lagrange polynomial of degree  $(\text{constNo} - 1)$  over the set  $\{\omega^i\}_{i \in [\text{constNo}]}$ , and let  $\ell(X)$  be the unique non-zero polynomial of degree *constNo* that satisfies  $\ell(\omega^i) = 0$  for all  $i \in [\text{constNo}]$ .

We note that the requirement that there exists a *constNo*-th root of unity  $\omega$  imposes a restriction on the maximum number of constraints in **R** that the scheme can support. In the particular case of  $\omega \in \mathbb{F}_{r_{\text{BN}}}$ , the restriction becomes  $\text{constNo} \leq 2^{28}$ . For  $\mathbb{F}_{r_{\text{BLS}}}$  this becomes  $\text{constNo} \leq 2^{47}$ .

KGen(**R**,  $1^\lambda$ ):

- i. Pick trapdoor  $td = (\tau, \alpha, \beta, \delta) \leftarrow (\mathbb{Z}_p^* \setminus \{\omega^{i-1}\}_{i=1}^{\text{constNo}}) \times (\mathbb{Z}_p^*)^3$ ;
- ii. For  $j \in \{1, \dots, \text{inpNo}\}$ , let

$$\begin{aligned} u_j(\tau) &= \sum_{i=1}^{\text{constNo}} U_{ij} \ell_i(\tau), \\ v_j(\tau) &= \sum_{i=1}^{\text{constNo}} V_{ij} \ell_i(\tau), \\ w_j(\tau) &= \sum_{i=1}^{\text{constNo}} W_{ij} \ell_i(\tau); \end{aligned}$$

- iii. Set

$$\begin{aligned} \text{srs}_{\mathbf{P}} &\leftarrow \left( [\alpha]_1, [\beta], [\delta], \{[u_j(\tau)]_1\}_{j=1}^{\text{inpNo}}, \{[v_j(\tau)]_1\}_{j=0}^{\text{inpNo}}, \right. \\ &\quad \left. \{[(u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau))/\delta]_1\}_{j=\text{inpNoPrim}+1}^{\text{inpNo}}, \right. \\ &\quad \left. \{[\tau^i \ell(\tau)/\delta]_1\}_{i=0}^{\text{constNo}-2} \right) \\ \text{srs}_{\mathbf{V}} &\leftarrow \left( [\alpha]_1, [\beta]_2, [\delta]_2, \{[\beta u_j(\tau) + \alpha v_j(\tau) + w_j]_1\}_{j=0}^{\text{inpNoPrim}} \right) \\ \text{srs} &\leftarrow (\text{srs}_{\mathbf{P}}, \text{srs}_{\mathbf{V}}) \end{aligned}$$

**return** *srs*, *td*

**P**(**R**, *srs*<sub>**P**</sub>, *prim* =  $(\text{inp}_j)_{j=1}^{\text{inpNoPrim}}$ , *aux* =  $(\text{inp}_j)_{j=\text{inpNoPrim}+1}^{\text{inpNo}}$ ):

i. Define

$$a^\dagger(X) = \sum_{j=1}^{inpNo} inp_j u_j(X), \quad b^\dagger(X) = \sum_{j=1}^{inpNo} inp_j v_j(X), \quad c^\dagger(X) = \sum_{j=1}^{inpNo} inp_j w_j(X);$$

1580 ii. Define the polynomial  $h(X) = (a^\dagger(X)b^\dagger(X) - c^\dagger(X))/\ell(X)$  and compute the  
 1581 coefficients  $\{h_i\}_{i=0}^{constNo-2}$  of  $h$ , such that  $h(X) = \sum_{i=0}^{constNo-2} h_i X^i$ .

1582 iii.  $r_a \leftarrow \mathbb{Z}_p$ ;

1583 iv.  $r_b \leftarrow \mathbb{Z}_p$ ;

v. Compute proof elements:

$$\begin{aligned} \mathbf{a} &\leftarrow \sum_{j=1}^{inpNo} inp_j \llbracket u_j(\tau) \rrbracket_1 + \llbracket \alpha \rrbracket_1 + r_a \llbracket \delta \rrbracket_1 \\ \mathbf{b} &\leftarrow \sum_{j=1}^{inpNo} inp_j \llbracket v_j(\tau) \rrbracket_2 + \llbracket \beta \rrbracket_2 + r_b \llbracket \delta \rrbracket_2 \\ \mathbf{c} &\leftarrow r_b \mathbf{a} + r_a \left( \sum_{j=1}^{inpNo} inp_j \llbracket v_j(\tau) \rrbracket_1 + \llbracket \beta \rrbracket_1 \right) + \\ &\quad \sum_{j=inpNoPrim+1}^{inpNo} inp_j \left\llbracket \frac{u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau)}{\delta} \right\rrbracket_1 + \\ &\quad \sum_{i=0}^{constNo-2} h_i \llbracket \tau^i \ell(\tau) / \delta \rrbracket_1 \end{aligned}$$

1584 **return**  $\pi \leftarrow (\mathbf{a}, \mathbf{b}, \mathbf{c})$ ;

1585  $\mathbf{V}(\mathbf{R}, srs_V, prim = (inp_j)_{j=1}^{inpNoPrim}, \pi)$ :

i. Check that:

$$\begin{aligned} \mathbf{a} \bullet \mathbf{b} &= \mathbf{c} \bullet \llbracket \delta \rrbracket_2 \\ &\quad + \left( \sum_{j=1}^{inpNoPrim} inp_j \llbracket u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau) \rrbracket_1 \right) \bullet \llbracket 1 \rrbracket_2 \\ &\quad + \llbracket \alpha \rrbracket_1 \bullet \llbracket \beta \rrbracket_2 \end{aligned}$$

1586 Note that  $\llbracket \alpha \rrbracket_1$  and  $\llbracket \beta \rrbracket_2$  are stored individually and used by the prover to re-  
 1587 compute  $\llbracket \alpha\beta \rrbracket_T$  seemingly redundantly. This is required in order to leverage the  
 1588 pairing check functionality built in to **Ethereum**, which accepts a sequence of tuples  
 1589 in  $\mathbb{G}_1 \times \mathbb{G}_2$  and returns **true** if and only if the product of the resulting pairings  
 1590 equals  $\llbracket 1 \rrbracket_T$ .

1591 **Sim**( $\mathbf{R}, srs, td, prim$ ):

1592

- i. Sample  $\mathbf{a}^* \leftarrow \$\mathbb{Z}_p$ ;  $\mathbf{b}^* \leftarrow \$\mathbb{Z}_p$ ;
- ii. Compute proof elements:

$$\begin{aligned}
 \mathbf{a} &\leftarrow \llbracket \mathbf{a}^* \rrbracket_1 + \llbracket \alpha \rrbracket_1 \\
 \mathbf{b} &\leftarrow \llbracket \mathbf{b}^* \rrbracket_1 + \llbracket \beta \rrbracket_2 \\
 \mathbf{c} &\leftarrow \frac{1}{\delta} \cdot \left[ \mathbf{a}^* \mathbf{b}^* \llbracket 1 \rrbracket_1 + \mathbf{a}^* \llbracket \beta \rrbracket_1 + \mathbf{b}^* \llbracket \alpha \rrbracket_1 \right. \\
 &\quad \left. - \sum_{j=1}^{inpNoPrim} inp_j \llbracket u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau) \rrbracket_1 \right]
 \end{aligned}$$

1593

**return**  $\pi \leftarrow (\mathbf{a}, \mathbf{b}, \mathbf{c})$ ;



## Chapter 4

# Implementation considerations and optimizations

### 4.1 Client security considerations

In this section we consider some details of client *wallet* software that manages user's private and public keys, **Zeth** notes, and interacts with the **Mixer** contract.

Due to the processing and storage requirements involved, we consider it impractical for all **Zeth** client implementations to assume that a dedicated **Ethereum** node (miner node or archive node) is run on the same host as the wallet. Therefore, in order to interact with the **Ethereum** network, wallet software must communicate with external **Ethereum** P2P nodes via their RPC channel, and must assume that these nodes are completely outside the wallet's control. *From a security standpoint, connected **Ethereum** nodes should therefore be considered untrusted, and in particular the details of all RPC calls and responses should be considered publicly visible.* Note that even if the connected **Ethereum** node itself is not malicious, 3rd parties able to see network traffic may also be able to gain an insight into the RPC communication of a specific **Zeth** client.

#### Note

Note that there are several possible models besides the fully untrusted **Ethereum** node. Organizations or individuals could host one or more “trusted” **Ethereum** nodes, which clients can securely connect to (if they trust the host). This centralization would represent a security trade-off. From the point of view of clients it would create a single point of trust, and for potential malicious observers or attackers it would represent a valuable target.

In what follows we focus on preventing data leaks through network traffic. We do not consider adversaries with physical access to the machine running the wallet (see Appendix C).

## Note

Importantly, we focus here on information leakages intrinsic to network communication patterns of the **Zeth** protocol. However, in order to protect against sophisticated adversaries, it is necessary to use network-level anonymity solutions to protect the source of messages emitted on the network. While this is outside of the scope of the **Zeth** protocol, we highly encourage implementers to establish threat models and consider using technologies like *mixnets* to protect against network analysis (see e.g. [PHE<sup>+</sup>17, DG09]).

### 4.1.1 Syncing and waiting

**Zeth** clients must periodically synchronize with the latest state of the blockchain. This is necessary to keep track of the data held by the **Mixer** contract, and to detect notes received by the user of the wallet, storing them for future transactions.

Clients should synchronize with **Ethereum** nodes in such a way that information is not leaked. As such:

1. Clients **MUST** use consensus evidence and block headers to verify all data they receive from **Ethereum** nodes.
2. Clients **MUST** locally store all parts of the **Mixer** state they require in order to function.
3. Clients **MUST** obtain all such information by “synchronizing” with the **Ethereum** blockchain and parsing relevant events emitted by **Mixer**. Clients **MUST NOT** query the **Mixer** state via RPC.
4. Clients **SHOULD** take steps to avoid being identified while synchronizing (see Appendix C.2. For example, clients **SHOULD** vary the set of **Ethereum** nodes that they connect to, and **SHOULD NOT** always sync from the block following the last one that they processed.
5. Clients **SHOULD NOT** re-request blocks or transaction receipts that are of particular interest to them. They **SHOULD** process all events, emitted by **Mixer**, in the same way.
6. Clients **SHOULD NOT** make any RPC calls or change their externally visible behaviour in response to blocks or transaction receipts that are of interest to them.

### Use of contract queries

We suggest that clients **SHOULD NOT** directly query the contract state, for the reasons discussed in Appendix C.2 and Appendix C.3 (and consequently, Section 4.2 suggests that the **Mixer** contract should, as far as possible, not expose public methods). The

1641 Zeth protocol prohibits direct queries of the state of  $\widetilde{\text{Mixer}}$  (via *public* smart-contract  
1642 functions) because they introduce a risk that client implementations will leak information  
1643 by using them.

1644 If implementers choose to add public methods to the  $\widetilde{\text{Mixer}}$  contract (for application-  
1645 specific reasons), they should consider carefully the security issues raised in Appendix C.  
1646 This specification assumes that `Mix` is the only public method of the  $\widetilde{\text{Mixer}}$  contract.

#### 1647 4.1.2 Note management

1648 `Mix` calls on the  $\widetilde{\text{Mixer}}$  contract emit log events containing new commitment values,  
1649 nullifiers, the new Merkle root and the secret data for new notes (encrypted using a key  
1650 derived from the recipients public key). As clients synchronize with the latest state of  
1651 the blockchain, they **MUST** read these events and correctly process the data they contain.

- 1652 1. Clients **MUST** process the `MixEventDType` event for every `Mix` transaction, in the  
1653 order in which they appear in the blockchain.
- 1654 2. Clients implementing spending functionality **MUST** use the commitment values in  
1655 events to track the state of the Merkle tree. The Merkle tree state will be used  
1656 to generate Merkle paths for future transactions, and **MUST** be made available to  
1657 the client without the need to query the contract. (Note that not all commitments  
1658 must necessarily be persisted – see Section 4.3).
- 1659 3. Clients that can receive notes **MUST** attempt to decrypt the ciphertexts for every  
1660 transaction (see Item 2 in Section 2.6).
- 1661 4. Clients **MUST NOT** perform any network-related action, including closing the RPC  
1662 connection, dependent on successful/unsuccessful decryption of ciphertexts (see  
1663 Appendix C.3).
- 1664 5. Clients that can receive notes **MUST** attempt to parse any successfully decrypted  
1665 plaintext (that is, ensure it is well-formed as in Item 3a in Section 2.6).
- 1666 6. Clients **MUST NOT** perform any network-related action, including closing the con-  
1667 nection, dependent on successful / unsuccessful parsing (see Appendix C.4).
- 1668 7. Clients that can receive notes **MUST** verify that successfully parsed plaintext data  
1669 is the opening of the corresponding commitment in the transaction (see Item 3b  
1670 in Section 2.6).
- 1671 8. Clients **MUST NOT** perform any network-related action, including closing the con-  
1672 nection, dependent on whether the parsed note data is the opening of the corre-  
1673 sponding commitment (see Appendix C.4).
- 1674 9. Clients **MUST** confirm that, after adding the new commitments, the local repre-  
1675 sentation of the Merkle tree of commitments has a root consistent with the event  
1676 data.

- 1677 10. Clients **SHOULD** keep a *local* record of the data related to valid decrypted notes.  
1678 This will be required in order to spend the notes in a future transaction.
- 1679 11. Clients implementing spending functionality **SHOULD** process all nullifiers in **Mix**  
1680 transaction events, checking for any corresponding notes previously recorded. Any  
1681 such note should be marked as spent in the client's local record.

#### 1682 4.1.3 Prepare arguments for Mix transaction

1683 Clients **MUST NOT** query **Ethereum** nodes while generating any arguments to a **Mix** call.  
1684 In particular, Merkle paths **MUST** be calculated using the client's local representation of  
1685 the Merkle tree of commitments that was constructed by parsing events.

1686 Where the zero-knowledge proof is generated by some external process, clients **MUST**  
1687 put in place sufficient security schemes to ensure that:

- 1688 • they are communicating with an authentic proof generation process (not a man-  
1689 in-the-middle), and
- 1690 • data sent to and from the proving process cannot be observed in transit and tam-  
1691 pered with by a third party, and
- 1692 • the proof has been generated for the correct instance-witness pair<sup>1</sup>

1693 Without these safe-guards, the operation of the system and the secret data required  
1694 to spend the input notes may be compromised. See Appendix C.6.

#### 1695 4.1.4 Wallet backup and recovery

1696 Given the restrictions placed on clients and their interaction with the **Mixer** contract,  
1697 it follows that clients must store all data required to spend notes owned by their users'  
1698 addresses, and to verify the validity of incoming notes. If this local data is lost, it must  
1699 be reconstructed before client operations can resume.

1700 **Zeth** private keys (see Table 1.5) can be used to fully restore client state. In this  
1701 case, clients **MUST** retrieve all events from the beginning of the **Mixer** contract's his-  
1702 tory, decrypting notes and tracking nullifiers, as described in the previous sections, to  
1703 reconstruct the set of unspent notes that they own.

1704 Without a backup of the private keys it is not possible to restore wallet state. As  
1705 such, private keys are the minimal set of data that must be securely stored and backed  
1706 up, and clients **SHOULD** provide support for this mode of recovery. However, to avoid the  
1707 need to scan all events emitted by **Mixer** (a very expensive operation) implementations  
1708 **SHOULD** also support back ups of further state data (such as the representation of the  
1709 Merkle tree of note commitments, the set of unspent notes, etc) to allow more efficient  
1710 modes of recovery.

---

<sup>1</sup>Although given an acceptable zk-proof  $\pi$  for an instance *prim* it is infeasible to check which witness has been used – which comes directly from the zero-knowledge property – we need to assure security measures that prevents any third party from mauling and tampering with the proof generation process.

## 1711 4.2 Contract security considerations

1712 Section 4.1 mentions several considerations for client implementations, concerning how  
1713 they interact with the contract. These must be taken into account when authoring the  
1714 contract code, to ensure that clients can securely retrieve the information they need to  
1715 operate without encouraging them to perform insecure operations.

- 1716 1. **Mixer** MUST validate inputs, the contract needs to ensure that the primary inputs  
1717 are elements of the scalar field  $\mathbb{F}_r$  (that is, they are in the range  $\{0, \dots, r - 1\}$ ).
- 1718 2. **Mixer** MUST output events for valid Mix calls, including:
  - 1719 (a) commitment for each new note;
  - 1720 (b) nullifier for each spent note;
  - 1721 (c) value of new Merkle root of commitments;
  - 1722 (d) ciphertexts for each new note;
  - 1723 (e) implementation-specific data (such as the one-time sender public key specified  
1724 in Section 3.5, required to decrypt the ciphertexts).
- 1725 3. The Mix function MUST be *payable*<sup>2</sup>, to support non-zero *vin*.
- 1726 4. **Mixer** MUST NOT expose any public methods except for Mix.

## 1727 4.3 Efficiency and scalability

### 1728 4.3.1 Importance of performance

1729 Poor performance and scalability has several impacts on the viability of the system.

1730 Efficiency and performance are arguably most important for the **Mixer** contract,  
1731 where gas usage directly affects the monetary cost of using **Zeth** to transfer value. That  
1732 is, high gas costs could make transactions very expensive, and therefore not practical for  
1733 many use-cases, undermining the utility and viability.

1734 High storage or compute requirements on the client would severely restrict the set  
1735 of devices on which **Zeth** client software can run, and long delays when sending or  
1736 receiving transactions can adversely affect the user-experience, discouraging some users  
1737 and undermining the privacy promises of the system.

1738 Although the proof-of-concept implementation of **Zeth** is not intended to be used in a  
1739 production environment, one of its aims is to demonstrate the practicality of the protocol  
1740 in terms of transaction costs. Therefore, some of the techniques described here have been  
1741 included in the proof-of-concept implementation, while in some cases implementers of  
1742 production software may wish to make different trade-offs.

---

<sup>2</sup>see <https://solidity.readthedocs.io/en/v0.6.2/types.html?highlight=payable#function-types>

### 4.3.2 Cost centers

One important factor, primarily affecting client performance, is the cost of zero-knowledge proof generation. This is directly related to the number of constraints used to represent the statement in Section 2.2, which in turn depends on the specific cryptographic primitives used (see Chapter 3).

Note that cryptographic primitives which are “snark-friendly” (i.e. can be implemented using fewer gates in an arithmetic circuit) may not necessarily run efficiently on the EVM or on standard hardware. As such, trade-offs must be made between proof generation cost and the gas costs of state transitions. An example of this is the hash function used in the Merkle tree of commitments. This is not only used in the statement of Section 2.2 (to verify Merkle proofs, see Section 2.2), but also on the client (to create Merkle proofs, see Section 2.3) and in the **Mixer** contract (to compute the Merkle root, see Section 2.5).

Aside from the specific hash function used, implementers have some freedom in the data structures and algorithms used to maintain the Merkle tree and generate proofs. Because of this freedom, and the importance of the chosen algorithms on performance across all components of the system, the majority of this section focuses on the details of the Merkle tree.

As described in Chapter 2, **Zeth** notes are maintained and secured by the Merkle tree, whose depth `MKDEPTH` must be fixed when the contract is deployed. Therefore, `MKDEPTH` determines the maximum number of notes ( $2^{\text{MKDEPTH}}$ ) that may be created over the lifetime of the deployment. To ensure the utility of **Zeth**, `MKDEPTH` must be sufficiently large, and therefore the following includes a discussion of *scalability* with respect to `MKDEPTH`.

Also, due to the fact that `MKDEPTH` is fixed, we assume that Merkle proofs are computed as `MKDEPTH`-tuples, no matter how many leaves have been populated. Unpopulated leaves are assumed to take some default value (usually a string of zero bits).

### 4.3.3 Client performance

#### Commitment Merkle tree

The simplest possible implementation, which stores only the data items at the leaves of the tree, requires  $2^{\text{MKDEPTH}} - 1$  hash invocations to compute the Merkle root or to generate a Merkle proof. The cost of this is too high to be practical for non-trivial values of `MKDEPTH`.

An immediate improvement in compute costs can be achieved by simply storing all nodes (or all nodes whose value is not the default value) and updating only those necessary as new commitments are added. When adding `JSOUT` consecutive leaves to the tree, after  $\mathcal{O}(\log_2(\text{JSOUT}))$  layers (requiring  $\mathcal{O}(\text{JSOUT})$  hashes) we reach the common ancestor of all new leaves and can update the Merkle tree by proceeding along a single branch (of approximately `MKDEPTH` -  $\log_2(\text{JSOUT})$  layers). Thus, the cost of updating the Merkle tree for a single transaction has a fixed bound which is linear in `JSOUT` and

1782 MKDEPTH. However, this doubles the storage cost of the tree since non-leaf nodes must  
1783 also be persisted.

1784 In the case of the client, the Merkle tree will only be used to generate proofs for notes  
1785 owned by the user of the client. Thereby **Zeth** clients need only store nodes of the Merkle  
1786 tree that are required for this purpose, and may discard all others. In particular, any  
1787 full sub-tree need only contain nodes that are part of Merkle paths associated with the  
1788 user's notes. Implementations that discard unnecessary nodes can achieve vast savings  
1789 in storage space.

#### 1790 **Zero-knowledge proof generation**

1791 As well as keeping the number of constraints as low as possible, it is important to ensure  
1792 that the prover implementation is optimal and thereby that proving times are as short as  
1793 possible. Proof generation should also exploit any available parallelism, to help reduce  
1794 the time taken. This may require specific programming languages or frameworks to be  
1795 used, necessitating that proof generation be performed by some external process (as is  
1796 the case in the proof-of-concept implementation).

1797 The proof generation process can also be very memory intensive (in part due to the  
1798 FFT calculations required), and so ensuring that enough RAM is present in the system  
1799 is important to avoid long proof times.

1800 See Appendix C.6 for a discussion of related security concerns.

#### 1801 **4.3.4 Contract performance**

1802 For most components of the contract, the set of operations to be performed is strictly  
1803 defined and the set of possible algorithmic optimizations that can be made is limited.  
1804 In these cases, it is important to ensure that code is benchmarked and optimized to  
1805 a reasonable degree, to minimize gas costs. We note that apart from the number and  
1806 type of compute instructions executed, store and lookup operations have a significant  
1807 impact on the gas used. In particular, storing new values is more expensive than over-  
1808 writing existing values, and a gas rebate is made when contracts release stored values.  
1809 See [Woo19, Appendix H.1] for further details.

1810 The primary component in which algorithmic optimizations can be made is the  
1811 Merkle tree of note commitments. The **Mixer** contract must compute (and store) the  
1812 new Merkle root after adding the **JSOUT** new commitments as leaves.

1813 As in Section 4.3.3, the simplest possible implementation which stores only the data  
1814 items at the leaves of the tree, requires the full root to be recomputed, involving  $2^{\text{MKDEPTH}} -$   
1815 1 hash invocations. This quickly becomes impractical for non-trivial values of MKDEPTH.

1816 The first-pass optimization (also described in Section 4.3.3) can be used to ensure  
1817 that the cost of updating the Merkle tree (number of hash computations, stores and  
1818 loads) is bounded by a constant that is linear in the Merkle tree depth. This is the  
1819 strategy used in the proof-of-concept implementation of **Mixer**.

1820 It may be possible to gain further improvements in gas costs by discarding nodes  
1821 from the Merkle tree that are not required. Unlike clients, **Mixer** is only required to

1822 compute the new Merkle root, and does not need to create or validate Merkle proofs  
 1823 (as these are checked as part of the zero-knowledge proof). Consequently, *all* nodes in a  
 1824 sub-tree can be discarded when the sub-tree is full, and the optimization is much simpler  
 1825 to implement than on the client.

1826 Another possible strategy for decreasing the gas costs associated with Merkle trees  
 1827 is *Merkle Shrubs*, described in [Lab19, Section 2.2]. Under this scheme, the contract  
 1828 maintains a “frontier” of roots of sub-trees and Merkle proofs provided by clients (as  
 1829 auxiliary inputs to the  $\mathbf{R}^Z$  circuit) contain a path from the leaf to one of the nodes in the  
 1830 frontier. The gas savings in this scheme are due to the fact that, for new commitments,  
 1831 the contract need only recompute the value of nodes from the leaf to the “frontier” (not  
 1832 all the way to the root of the tree). However this comes at the cost of complexity in the  
 1833 arithmetic circuit, which must verify a Merkle path to one of several frontier nodes.

1834 When choosing cryptographic primitives to be used on the EVM (and considering  
 1835 the trade-off with other platforms, described in Section 4.3.1) it may be valuable to note  
 1836 that the EVM supports so-called “pre-compiled contracts”. These behave like built-  
 1837 in functions providing very gas-efficient access to certain algorithms, such as Keccak.  
 1838 However, pre-compiled contracts exist only for a limited set of algorithms. Others must  
 1839 be implemented using EVM instructions.

#### 1840 4.3.5 Optimizing Blake2’s circuit.

1841 After presenting Blake2s circuit and its components working on little endian variables,  
 1842 we show a few optimizations.

#### 1843 Helper circuits

1844 We first define the following helper circuits needed in the Blake2s routine, operating on  
 1845  $w$ -bit long words.

1846 **XOR circuits** The following XOR circuits on  $w$ -bit long variables have been imple-  
 1847 mented, we assume the inputs are boolean (this is not checked in these circuits),

- 1848 • “Classic” XOR circuit, which xors 2 variables,  
 1849  $a \oplus b = c$ ;
- 1850 • XOR with constant, which xors two variables and a constant,  
 1851  $a \oplus b \oplus c = d$ , with  $c$  constant;
- 1852 • XOR with rotation, which xors two variables and rotates the result.  
 1853  $a \oplus b \ggg r = c$ , with  $r$  constant, and  $\ggg$  the rightward rotation [MJS15, Section  
 1854 2.3]; i.e. for and constant  $r < w$  we have  $a_i \oplus b_i = c_{i+r \pmod w}$ , for  $i = 0, \dots, w$ ,



Each of these circuits presents  $w$  constraints. Assuming that the inputs are boolean, the output is automatically boolean. To ascertain that both inputs are boolean ( $a$  and  $b$ ), we would need  $2 \cdot w$  more gates per circuit.<sup>3</sup>

**Modular addition** We present here two circuits to verify modular arithmetic.

**Double modular addition:  $a + b = c \pmod{2^w}$ .** This circuit checks that the sum of two  $w$ -bit long variables in little endian format modulo  $2^w$  is equal to a  $w$ -bit long variable. More precisely, it checks the equality of the modular addition of  $a + b \pmod{2^w}$  and  $c$  and the booleanness of the later. We assume the inputs are boolean (this is not checked in this circuit).

As the addition of two  $w$ -bit long integers results in at most an  $(w + 1)$ -bit integer, we consider  $c$  to be  $(w + 1)$ -bit long. We do not care about the last bit value,  $c_w$ , but have to ensure its booleanness.

The circuit presents the following  $w + 2$  constraints, for  $a$  and  $b$  of size  $w$ , where  $w = 32$  in practice, and variable  $c$  of size  $w + 1$ , that:

$$\sum_{i=0}^{w-1} (a_i + b_i) \cdot 2^i = \sum_{j=0}^w c_j \cdot 2^j \quad (4.1)$$

$$\forall j \in \{0, \dots, w\}, (c_j - 0) \cdot (c_j - 1) = 0 \quad (4.2)$$

**Triple modular addition:  $a + b + c = d \pmod{2^w}$ .** This circuit checks the equality of a  $w$ -bit long variable  $d$  with the sum of three  $w$ -bit long variables in little endian format modulo  $2^w$ . More precisely, it checks the equality of the modular addition of  $a + b + c \pmod{2^w}$  and  $d$  and the booleanness of the latter. We assume the inputs are boolean (this is not checked in this circuit).

As the addition of three  $w$ -bit long integers results in at most an  $(w + 2)$ -bit integer, we consider  $d$  to be  $(w + 2)$ -bit long. We do not care about the values of the last two bits ( $d_w$  and  $d_{w+1}$ ), but have to ensure their booleanness.

The circuit presents the following  $w + 3$  constraints, for  $a$ ,  $b$  and  $c$  of size  $w$ , where  $w = 32$  in practice, and variable  $d$  of size  $w + 2$ , that:

$$\sum_{i=0}^{w-1} (a_i + b_i + c_i) \cdot 2^i = \sum_{j=0}^{w+1} d_j \cdot 2^j \quad (4.3)$$

$$\forall j \in \{0, \dots, w + 1\}, (d_j - 0) \cdot (d_j - 1) = 0 \quad (4.4)$$

---

<sup>3</sup>Making sure that no gates are duplicated in the circuit is very important to keep the proving time as small as possible. One challenge of writing RICS programs is to make sure that the statement is correctly represented, without redundancy, in order to keep the constraint system as small as possible.

$G(a, b, c, d; x, y) \mapsto (a_2, b_2, c_2, d_2)$	$\text{getSigma}()$
1 : $a_1 \leftarrow a + b + x \pmod{2^w}$	1 : $\Sigma \in (\mathbb{N}^{16})^{10}$
2 : $d_1 \leftarrow d \oplus a_1 \ggg r_1$	2 : $\Sigma[0] \leftarrow (0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15)$
3 : $c_1 \leftarrow c + d_1 \pmod{2^w}$	3 : $\Sigma[1] \leftarrow (14, 10, 4, 8, 9, 15, 13, 6, 1, 12, 0, 2, 11, 7, 5, 3)$
4 : $b_1 \leftarrow b \oplus c_1 \ggg r_2$	4 : $\Sigma[2] \leftarrow (11, 8, 12, 0, 5, 2, 15, 13, 10, 14, 3, 6, 7, 1, 9, 4)$
5 : $a_2 \leftarrow a_1 + b_1 + y \pmod{2^w}$	5 : $\Sigma[3] \leftarrow (7, 9, 3, 1, 13, 12, 11, 14, 2, 6, 5, 10, 4, 0, 15, 8)$
6 : $d_2 \leftarrow d_1 \oplus a_2 \ggg r_3$	6 : $\Sigma[4] \leftarrow (9, 0, 5, 7, 2, 4, 10, 15, 14, 1, 11, 12, 6, 8, 3, 13)$
7 : $c_2 \leftarrow c_1 + d_2 \pmod{2^w}$	7 : $\Sigma[5] \leftarrow (2, 12, 6, 10, 0, 11, 8, 3, 4, 13, 7, 5, 15, 14, 1, 9)$
8 : $b_2 \leftarrow b_1 \oplus c_2 \ggg r_4$	8 : $\Sigma[6] \leftarrow (12, 5, 1, 15, 14, 13, 4, 10, 0, 7, 6, 3, 9, 2, 8, 11)$
9 : <b>return</b> $a_2, b_2, c_2, d_2$	9 : $\Sigma[7] \leftarrow (13, 11, 7, 14, 12, 1, 3, 9, 5, 0, 15, 4, 8, 6, 2, 10)$
	10 : $\Sigma[8] \leftarrow (6, 15, 14, 9, 11, 3, 0, 8, 12, 2, 13, 7, 1, 4, 10, 5)$
	11 : $\Sigma[9] \leftarrow (10, 2, 8, 4, 7, 6, 1, 5, 15, 11, 9, 14, 3, 12, 13, 0)$
	12 : <b>return</b> $\Sigma$

Figure 4.1:  $G$  primitive [MJS15, Section 3.1]

Figure 4.2: Blake2 permutation table [MJS15, Section 2.7]

#### 1877 Blake2s routine circuit

1878 We define in this section the circuit of the Blake2 routine (see [MJS15, Section 3.1]  
1879 and Fig. 4.1) known as “ $G$  function” [ANWOW13, Section 2.4].  $G$  is based on ChaCha  
1880 encryption [Ber08a]. It works on  $w$ -bit long words, and presents  $8 \cdot w + 10$  constraints.  
1881 The function mixes a state  $(a, b, c$  and  $d)$  with the inputs  $(x$  and  $y)$  and returns the  
1882 updated state.

1883 This circuit does not check the booleanness of the inputs or state. However, given that  
1884 the state is boolean, the output is automatically boolean due to the use of the modular  
1885 addition circuits.

1886 For Blake2s, we have  $w = 32$ ,  $r_1 = 16$ ,  $r_2 = 12$ ,  $r_3 = 3$  and  $r_4 = 7$ .

#### 1887 Blake2s compression function circuit

The compression function is defined as follows, for more details see Fig. 4.3,

$$\text{Blake2sC} : \mathbb{B}^n \times \mathbb{B}^{2n} \times \mathbb{B}^{n/4} \times \mathbb{B}^{n/4} \rightarrow \mathbb{B}^n.$$

1888 Blake2C takes as input a state  $h \in \mathbb{B}^n$  which is used as chaining value when hashing,  
1889 a message to compress  $x \in \mathbb{B}^{2n}$ , a message length written in binary  $t \in \mathbb{B}^{n/4}$  which is  
1890 incremented when hashing and a binary flag  $f \in \mathbb{B}^{n/4}$  to tell whether the current block  
1891 is the last to be compressed to prevent length extension attacks.

1892 Blake2C uses the  $G$  function iteratively over **rounds** number of rounds on a state  
1893 and message. The constant initialization vector  $IV$  and the permutation table  $\Sigma$  are  
1894 hard-coded. Blake2sC works in little endian (see [MJS15, Section 2.4]) on  $n$ -bit long  
1895 variables ( $n = 256$ ),  $w$ -bit long words ( $w = 32$ ), and the rotation constants specified

in Section 4.3.5 (see [MJS15, Section 2.1]). We have the following constants (see specifications [ANWOW13] and [MJS15, Section 2.2]),

- **IV** is the  $(8 \cdot w)$ -bit long initialization vector; it corresponds to the first  $w$  bits of the fractional parts of the square roots of the first eight prime numbers  $(2, 3, 5, 7, \dots)$  (see [MJS15, Section 2.6]);
- $\Sigma$  are the  $10 \cdot 16$  permutation constants of Blake2 (see Fig. 4.2 and [MJS15, Section 2.7]);
- **rounds**, the number of rounds: 10 for Blake2sC, 12 for Blake2bC.

We have the following variables (see specifications [ANWOW13] and [MJS15, Section 2.2]),

- **H** is the  $(8 \cdot w)$ -bit long initial state while  $v$  is the  $(16 \cdot w)$ -bit long final state;
- $T[i]$  are two  $w$ -bit long counters encoding the block length;
- $F[i]$  are two  $w$ -bit long finalization flags. We set the first one  $F[0]$  to  $2^w - 1$  to state when the input block is the last one to be hashed. The second,  $F[1] = 0$  is only set for tree hashing mode (which is not our case) and is therefore unused.

We introduce the following functions to write Blake2C (see specifications [ANWOW13] and [MJS15, Section 2.6]):

- The function **prime** takes a positive integer  $i$  as input and outputs the  $i$ -th prime number;
- The function **dec** takes a real number  $x$  as input outputs its positive decimal part.

This circuit presents  $((64 \cdot \text{rounds} + 8) \cdot w + 8 \cdot \text{rounds} + 10)$  constraints. For Blake2sC, as  $w = 32$  and **rounds** = 10, we have 21536 constraints.

We do not check the input block booleaness in this circuit. Given that the initial state is boolean, the output is automatically boolean. This can be proved iteratively by the booleaness of **G** primitive's output.

**Security requirement.** The inputs to Blake2sC **MUST** be boolean.

## Blake2s hash function

The hash function is defined as follows, for more details see Fig. 4.3,

$$\text{Blake2s} : \mathbb{B}^{\leq 2n} \times \mathbb{B}^* \rightarrow \mathbb{B}^n$$

Blake2 takes as input a hash key  $k \in \mathbb{B}^n$  and the message to hash  $x \in \mathbb{B}^{2n}$ . Blake2 uses the Blake2C function iteratively over each  $2n$ -bit long chunk of the padded message. If the key is non null, it is used as the first block to be hashed. The constant initialization vector **IV** and part of the parameter block **PB** are hard-coded. We have the following constants (see specifications [ANWOW13] and [MJS15, Section 2.2]),

### Blake2C( $h, m, t, f$ )

---

```

1 :   $\mathbf{T}, \mathbf{F}, \mathbf{H}, \mathbf{IV}, v \in (\mathbb{B}^w)^2 \times (\mathbb{B}^w)^2 \times (\mathbb{B}^w)^8 \times (\mathbb{B}^w)^8 \times (\mathbb{B}^w)^{16}$ 
2 :   $\{\mathbf{IV}[i]\}_{i \in [8]} \leftarrow \left\{ \left\lfloor 2^w \cdot \text{dec}(\sqrt{\text{prime}(i+1)}) \right\rfloor \right\}_{i \in [8]}$ 
3 :   $\Sigma \leftarrow \text{getSigma}()$ 
4 :   $\{\mathbf{H}[i]\}_{i \in [8]} \leftarrow \{h[i \cdot w : (i+1) \cdot w]\}_{i \in [8]}$ 
5 :   $\{m[i]\}_{i \in [8]} \leftarrow \{x[i \cdot w : (i+1) \cdot w]\}_{i \in [8]}$ 
6 :   $\mathbf{T}[0], \mathbf{T}[1] \leftarrow t[w:2w], t[0:w]$ 
7 :   $\mathbf{F}[0], \mathbf{F}[1] \leftarrow f[w:2w], f[0:w]$ 
8 :   $\{v[i]\}_{i \in [8]} \leftarrow \{\mathbf{H}[i]\}_{i \in [8]}$ 
9 :   $\{v[i+8]\}_{i \in [8]} \leftarrow \{\mathbf{IV}[i]\}_{i \in [8]}$ 
10 :  $v[12], v[13] \leftarrow v[12] \oplus \mathbf{T}[0], v[13] \oplus \mathbf{T}[1]$ 
11 :  $v[14], v[15] \leftarrow v[14] \oplus \mathbf{F}[0], v[15] \oplus \mathbf{F}[1]$ 
12 : foreach  $r \in [\text{rounds}]$  do
13 :    $\tau \leftarrow \Sigma[r \pmod{15}]$ 
14 :    $v[0], v[4], v[8], v[12] \leftarrow \mathbf{G}(v[0], v[4], v[8], v[12], m[\tau[0]], m[\tau[1]])$ 
15 :    $v[1], v[5], v[9], v[13] \leftarrow \mathbf{G}(v[1], v[5], v[9], v[13], m[\tau[2]], m[\tau[3]])$ 
16 :    $v[2], v[6], v[10], v[14] \leftarrow \mathbf{G}(v[2], v[6], v[10], v[14], m[\tau[4]], m[\tau[5]])$ 
17 :    $v[3], v[7], v[11], v[15] \leftarrow \mathbf{G}(v[3], v[7], v[11], v[15], m[\tau[6]], m[\tau[7]])$ 
18 :    $v[0], v[5], v[10], v[15] \leftarrow \mathbf{G}(v[0], v[5], v[10], v[15], m[\tau[8]], m[\tau[9]])$ 
19 :    $v[1], v[6], v[11], v[12] \leftarrow \mathbf{G}(v[1], v[6], v[11], v[12], m[\tau[10]], m[\tau[11]])$ 
20 :    $v[2], v[7], v[8], v[13] \leftarrow \mathbf{G}(v[2], v[7], v[8], v[13], m[\tau[12]], m[\tau[13]])$ 
21 :    $v[3], v[4], v[9], v[14] \leftarrow \mathbf{G}(v[3], v[4], v[9], v[14], m[\tau[14]], m[\tau[15]])$ 
22 : return  $\|_{i=0}^8 \mathbf{H}[i] \oplus v[i] \oplus v[i+8]$ 

```

Figure 4.3: Blake2 compression function [MJS15, Section 3.2]. Set  $n$ ,  $w$  and  $\mathbf{G}$ 's constants to obtain Blake2sC.

- 1928 •  $\mathbf{IV}$  is the  $(8 \cdot w)$ -bit long Initialization Vector; it corresponds to the first  $w$  bits of the  
1929 fractional parts of the square roots of the first eight prime numbers  $(2, 3, 5, 7, \dots)$   
1930 (see [MJS15, Section 2.6]).
- 1931 We have the following variables (see specifications [ANWOW13] and [MJS15, Section  
1932 2.2]),
- 1933 •  $\mathbf{PB}$  is the  $(16 \cdot w)$ -bit long parameter block used to initialize the state (see [MJS15,  
1934 Section 2.5]). In big endian encoding, the first byte corresponds to the digest  
1935 length (fixed to 32 bytes), the second byte to the key length, the third and fourth  
1936 bytes correspond to the use of the serial mode;
- 1937 •  $\mathbf{H} \in \mathbb{B}^{\text{BLAKE2sCLEN}}$ , the chaining value.

## Blake2( $k, x$ )

---

```

1 : H, IV, PB  $\in \mathbb{B}^{8w} \times \mathbb{B}^{8w} \times \mathbb{B}^{8w}$ 
2 : PB  $\leftarrow \text{pad}_{8 \cdot w}(\text{encode}_{\mathbb{N}}(0x0101) \parallel \text{pad}_w(\text{encode}_{\mathbb{N}}(\lceil \text{length}(k)/\text{BYTELEN} \rceil)) \parallel \text{encode}_{\mathbb{N}}(0x20))$ 
3 : IV  $\leftarrow \parallel_{i=0}^8 \left[ 2^w \cdot \text{dec}(\sqrt{\text{prime}(i+1)}) \right]$ 
4 : H  $\leftarrow \text{PB} \oplus \text{IV}$ 
5 :  $y \leftarrow x$ 
6 : if  $\text{length}(k) \neq 0$  do
7 :    $y \leftarrow \text{pad}_{2n}(k) \parallel y$ 
8 :    $z \leftarrow \text{pad}_{2n \cdot \lceil \text{length}(y)/2n \rceil}(y)$ 
9 :   for  $i \in [\lceil \text{length}(z)/2n \rceil]$  do
10 :    if  $i = \lceil \text{length}(z)/2n \rceil - 1$  do
11 :      H  $\leftarrow \text{Blake2C}(H, z[i \cdot 2n:(i+1) \cdot 2n], \text{pad}_{2w}(\text{encode}_{\mathbb{N}}(\lceil \text{length}(y)/\text{BYTELEN} \rceil)), \text{pad}_{2w}(\text{encode}_{\mathbb{N}}(2^w - 1)))$ 
12 :    else
13 :      H  $\leftarrow \text{Blake2C}(H, z[i \cdot 2n:(i+1) \cdot 2n], \text{pad}_{2w}(\text{encode}_{\mathbb{N}}((i+1) \cdot 2n/\text{BYTELEN})), \text{pad}_{2w}(0))$ 
14 :   return H

```

Figure 4.4: Blake2 hash function [MJS15, Section 3.3]. Set  $n = 16w$  and G's constants accordingly to obtain Blake2s.

1938 We do not check the input block booleaness in this circuit. Given that the initial  
1939 state is boolean, the output is automatically boolean. This can be proved iteratively by  
1940 the booleaness of Blake2C primitive's output.

1941 **Security requirement** To ensure the correct use of Blake2s, Blake2s's inputs **MUST** be  
1942 boolean.

## 1943 Optimizing the circuits

1944 The above helper circuits form the building blocks of the Blake2s compression function.  
1945 We show here two exclusive methods to optimize these circuits.

## 1946 Optimizing the Modular additions

**Double modular addition:**  $a + b = c \pmod{2^w}$ . We present here an optimization on the circuit to save one constraint by merging the modular constraint with a boolean constraint. The optimized circuit presents the following constraints:

$$\left( \sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i \right) \cdot \left( \sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i - 2^w \right) = 0 \quad (4.5)$$

$$\forall j \in \{0, \dots, w-1\}, (c_j - 0) \cdot (c_j - 1) = 0 \quad (4.6)$$

1947 with  $\sum_{i=0}^{w-1} x_i \cdot 2^i$  a binary encoding of  $x$  ( $x_i$  is the  $i^{th}$  bit of  $x$ ).

1948 These equations describe  $w + 1$  constraints to prove the bit equality  $a + b = c$  (note  
1949 that an additional  $2 \cdot w$  constraints would be required to prove the booleanness of input  
1950 variables  $a$  and  $b$ ). We now explain how we obtained them.

1951 *Proof.* The most straightforward way to prove that  $a + b = c \pmod{2^w}$  and  $c$  booleanness  
1952 is with the set of constraints illustrated in Eq. (4.1) and in Eq. (4.2).

As we perform arithmetic modulo  $2^w$ , we do not care about the value of  $c_w$  but would like to ensure its booleanness. As one may notice, the summing constraint Eq. (4.1) is an equality of two linear combinations with no multiplication by a variable. Hence, we can combine it with the boolean constraint of  $c_w$  to remove any reference to  $c_w$  and still have a bilinear gate. To do so, we first rewrite Eq. (4.1) as an equality check over  $c_w \cdot 2^w$  and multiply Eq. (4.2) for  $j = n$  by  $2^{2 \cdot w}$ .

$$\sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i = c_w \cdot 2^w \quad (4.7)$$

$$2^w \cdot (c_w - 0) \cdot 2^w \cdot (c_w - 1) = 0 \quad (4.8)$$

We finally replace  $c_w \cdot 2^w$  in Eq. (4.8) by the value from Eq. (4.7).

$$\begin{aligned} 0 &= 2^w \cdot (c_w - 0) \cdot 2^w \cdot (c_w - 1) = 2^w \cdot c_w \cdot (2^w \cdot c_w - 2^w) \\ &= \left( \sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i \right) \cdot \left( \left( \sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i \right) - 2^w \right) \end{aligned}$$

1953 This results in Eq. (4.5) and Eq. (4.6). All references to  $c_w$  have disappeared and, with  
1954 a single multiplication by a variable, we still have bilinear gates.  $\square$

1955 **Triple modular addition:  $a + b + c = d \pmod{2^w}$ .** To optimize, we use the  
1956 above circuit twice. We define a temporary variable  $d'$  such that  $a + b = d' \pmod{2^w}$ .  
1957 As such, we have  $c + d' = d \pmod{2^w}$ . As  $d'$  is the addition of two  $w$ -bit long variables,  
1958 it is  $(w + 1)$ -bit long. However as we evaluate the sum modulo  $2^w$ , we discard the last bit  
1959 of  $d'$ . We proceed similarly for  $d$ . To ensure that  $d$  is boolean, we check the booleanness  
1960 of the  $w + 1$  bits of  $d$  as well as the booleanness of the last bit of  $d'$  (to account for  $d$ 's  
1961  $w + 2^{th}$  bit in the original expression  $(a + b + c = d \pmod{2^w})$ ).

We thus obtain the following circuit with  $w + 2$  constraints,

$$\begin{aligned} \left( \sum_{i=0}^{w-1} (a_i + b_i - d'_i) \cdot 2^i \right) \cdot \left( \sum_{i=0}^{w-1} (a_i + b_i - d'_i) \cdot 2^i - 2^w \right) &= 0 \\ \left( \sum_{i=0}^{w-1} (c_i + d'_i - d_i) \cdot 2^i \right) \cdot \left( \sum_{i=0}^{w-1} (c_i + d'_i - d_i) \cdot 2^i - 2^w \right) &= 0 \\ \forall j \in \{0, \dots, w - 1\}, (d_j - 0) \cdot (d_j - 1) &= 0 \end{aligned}$$

1962 These optimizations lead to a gain of 320 constraints ( $= 4 \cdot 8 \cdot \text{rounds}$ ).

1963 **Optimizing Blake2s routine’s circuit** As seen in Fig. 4.1, our routine presents 2  
1964 double and 2 triple modular additions. Each of these circuits comprises at least one  
1965 modular constraint which pack several  $w$ -bit long variables. The circuit is however  
1966 processed in  $\mathbb{F}_r$ , that is to say most integers can be written over **FIELD**CAP bits. We can  
1967 thus batch the modular constraints. As the **G** primitive performs 2 double modular and  
1968 2 triple modular, we have in total 6 modular checks per iteration. We can batch up to  
1969 **FIELD**CAP/ $w$  constraints together. For  $w = 32$  and **FIELD**CAP  $\geq 224$  (which holds for  
1970 **BN-254** and **BLS12-377**), we can encode up to 7 words per field element, that is to say  
1971 we can include all the modular constraints into a single one.

1972 This optimization leads to a gain of 274 constraints ( $= 4 \cdot 8 \cdot 10 - \lceil \frac{4 \cdot 8 \cdot 10}{7} \rceil$ ).

1973 **Optimization conclusion** Using the more efficient optimization on the modular ad-  
1974 ditions, the Blake2s compression function comprises 21216 constraints.

#### 1975 **Increasing the PRF security with Blake**

1976 As Blake2 comprises a personalization tag in its parameter block **PB**, one could ensure the  
1977 independence of the PRFs by writing different tags for each of them (we would be able to  
1978 consider up to  $2^{30}$  inputs and outputs). We did not choose to write this enhancement in  
1979 the instantiation to keep a general tagging method in case of a change of hash function.

## 1980 **4.4 Encryption of the notes**

1981 In this section we give some remarks concerning the implementation of the **Zeth** en-  
1982 cryption scheme, described in Section 3.5. As noted, there are several details in the  
1983 specification of the underlying primitives which can impact security if not carefully im-  
1984 plemented. The following list is by no means exhaustive but includes several details  
1985 noted during development of the proof-of-concept system.

- 1986 • Private keys for **Curve25519** **MUST** be randomly generated as 32 bytes where the  
1987 first byte is a multiple of 8, and the last byte takes a value between 64 and 127  
1988 (inclusive). Further details are given in [Ber06], including an example algorithm  
1989 for generation. Implementations **MUST** take care to ensure that their code, or any  
1990 external libraries they rely upon, correctly perform this step.
- 1991 • A similar observation holds for **Poly1305** in which the  $r$  component of the **MAC**  
1992 key  $(r, s)$  **MUST** be *clamped* in a specific way (see Section 3.5.3). This step is also  
1993 essential and **MUST** be performed.
- 1994 • In the implementation of the **ChaCha** stream cipher, correct use of the *key*, *counter*  
1995 and *nonce* **MUST** be ensured in order to adhere to the standard and guarantee the  
1996 appropriate security properties.

1997        During the proof-of-concept implementation it was not obvious that the cryptogra-  
1998 phy library<sup>4</sup> adhered to the specifications with respect to the above points. In particular,  
1999 it was not clear whether key clamping was performed at generation time and/or when  
2000 performing operations. Moreover, the interface to the **ChaCha** cipher accepted a differ-  
2001 ent set of input parameters (namely *key* and *nonce* with no *counter*). This left some  
2002 ambiguity about the responsibility for clamping, and whether the **ChaCha** block data  
2003 would be updated as described in the specification. Details of how this was resolved are  
2004 given in the proof-of-concept encryption code, which may prove a useful reference for  
2005 implementers<sup>5</sup>.

---

<sup>4</sup><https://cryptography.io/en/latest/>

<sup>5</sup>see <https://github.com/clearmatics/zeth/blob/v0.4/client/zeth/encryption.py>



## Appendix A

# Transaction non malleability

The transaction malleability problem for a DAP (Section 1.4) is characterized by a game TR-NM involving a polynomial-time adversary  $\mathcal{A}$  as described below.

**Definition A.0.1.** Let DAP be a (candidate) Decentralized Anonymous Payment scheme.

$$\text{DAP} = (\text{Setup}, \text{GenAddr}, \text{SendTx}, \text{VerifyTx}, \text{Receive})$$

We say that DAP is TR-NM secure if, for every  $\text{poly}(\lambda)$ -time adversary  $\mathcal{A}$

$$\text{Adv}_{\text{DAP}, \mathcal{A}}^{\text{tr-nm}}(\lambda) < \text{negl}(\lambda),$$

where  $\text{Adv}_{\text{DAP}, \mathcal{A}}^{\text{tr-nm}}(\lambda) = \Pr[\text{TR-NM}(\text{DAP}, \mathcal{A}, \lambda) = 1]$  is  $\mathcal{A}$ 's advantage in the TR-NM experiment.

Below, we adapt [BSCG<sup>+</sup>14, Appendix C.2] to our specific DAP—Zeth.

We start by describing the TR-NM experiment. Given a (candidate) Zeth DAP, adversary  $\mathcal{A}$ , and security parameter  $\lambda$ , the (probabilistic) game  $\text{TR-NM}(\text{DAP}, \mathcal{A}, \lambda)$  consists of an interaction between  $\mathcal{A}$  and a challenger  $\mathcal{C}$ , terminating with a binary output by  $\mathcal{C}$ .

At the beginning of the game,  $\mathcal{C}$  samples  $pp \leftarrow \text{Setup}(\lambda)$  and sends  $pp$  to  $\mathcal{A}$ . Next,  $\mathcal{C}$  initializes a DAP oracle  $\text{O}^{\text{DAP}}$  with  $pp$  and allows  $\mathcal{A}$  to issue queries to it [RZ19, Appendix B].

At the end of the experiment,  $\mathcal{A}$  sends to  $\mathcal{C}$  a **Mixer** contract call transaction  $tx_{\text{Mix}}^*$ , and  $\mathcal{C}$  outputs 1 iff the following conditions hold. Letting  $T$  be the set of transactions generated by  $\text{O}^{\text{DAP}}$  in response to **SendTx** queries, there exists  $tx_{\text{Mix}} \in T$  such that:

1.  $tx_{\text{Mix}}$  was not inserted in  $L$  by  $\mathcal{A}$ ;
2.  $tx_{\text{Mix}}^*.data \neq tx_{\text{Mix}}.data$ ;
3.  $\text{VerifyTx}(pp, tx_{\text{Mix}}^*, L') = 1$  where  $L'$  is the portion of the ledger  $L$  preceding  $tx_{\text{Mix}}$ ;
4. a serial number revealed in  $tx_{\text{Mix}}^*$  is also revealed in  $tx_{\text{Mix}}$ .

## A.1 Transaction malleability attack on Zeth

In this section we present the threat related to the transaction malleability attack on Zeth and expose the solutions by ZeroCash [BSCG<sup>+</sup>14] and Zcash [ZCa19] that we adapted.

First, we start by assuming that none of the checks related to transaction malleability attack have been added in the protocol Chapter 2. As such, we assume that *hsig* and *htags* are not attributes of `PrimInputDType`,  $\phi$  is not an attribute of `AuxInputDType`, and *otssig* and *otsvk* are not attributes of the `MixInputDType` data type anymore. As a consequence, all checks related to these attributes are removed from the protocol. Moreover, if *zn* is an object of type `ZethNoteDType`, then *zn*. $\rho$  is chosen at random. Finally, the NP-relation used in Zeth, now denoted  $\mathbf{R}^{\text{mal}}$ , becomes the following:

- For each  $i \in [\text{JSIN}]$ :

1.  $\text{aux.jsins}[i].\text{znote.apk} = \text{Blake2s}(\text{tag}_{\text{ask}}^{\text{addr}} \parallel \text{pad}_{\text{BLAKE2SCLEN}}(0))$   
with  $\text{tag}_{\text{ask}}^{\text{addr}}$  defined in Section 3.1.3
2.  $\text{aux.jsins}[i].\text{nf} = \text{Blake2s}(\text{tag}_{\text{ask}}^{\text{nf}} \parallel \text{aux.jsins}[i].\text{znote}.\rho)$   
with  $\text{tag}_{\text{ask}}^{\text{nf}}$  defined in Section 3.1.3
3.  $\text{aux.jsins}[i].\text{cm} = \text{Blake2s}(\text{aux.jsins}[i].\text{znote}.r \parallel m)$   
with  $m = \text{aux.jsins}[i].\text{znote.apk} \parallel \text{aux.jsins}[i].\text{znote}.\rho \parallel \text{aux.jsins}[i].\text{znote}.v$
4.  $(\text{aux.jsins}[i].\text{znote}.v) \cdot (1 - e) = 0$  is satisfied for the boolean value  $e$  set such that if  $\text{aux.jsins}[i].\text{znote}.v > 0$  then  $e = 1$ .
5. The Merkle root  $\text{mkroot}'$  obtained after checking the Merkle authentication path  $\text{aux.jsins}[i].\text{mkpath}$  of commitment  $\text{aux.jsins}[i].\text{cm}$ , with MKHASH, is equal to  $\text{prim.mkroot}$  if  $e = 1$ .
6.  $\text{prim.nfs}[i]$   
 $= \{\text{Pack}_{\mathbb{F}_r}(\text{aux.jsins}[i].\text{nf}[k \cdot \text{FIELD CAP}:(k + 1) \cdot \text{FIELD CAP}])\}_{k \in [\lfloor \text{PRNFOUTLEN} / \text{FIELD CAP} \rfloor]}$

- For each  $j \in [\text{JSOUT}]$ :

1.  $\text{prim.cms}[j] = \text{Blake2s}(\text{aux.znotes}[j].r \parallel m)$   
with  $m = \text{aux.znotes}[j].\text{apk} \parallel \text{aux.znotes}[j].\rho \parallel \text{aux.znotes}[j].v$

- $\text{prim.rsd} = \text{Pack}_{\text{rsd}}(\{\text{aux.jsins}[i].\text{nf}\}_{i \in [\text{JSIN}]}, \text{aux.vin}, \text{aux.vout})$
- Check that the “joinsplit is balanced”, i.e. check that the joinsplit equation holds:

$$\begin{aligned} & \text{Pack}_{\mathbb{F}_r}(\text{aux.vin}) + \sum_{i \in [\text{JSIN}]} \text{Pack}_{\mathbb{F}_r}(\text{aux.jsins}[i].\text{znote}.v) \\ &= \sum_{j \in [\text{JSOUT}]} \text{Pack}_{\mathbb{F}_r}(\text{aux.znotes}[j].v) + \text{Pack}_{\mathbb{F}_r}(\text{aux.vout}) \end{aligned}$$

### A.1.1 The attack

In order to win the game TR-NM on the weak Zeth DAP above, an adversary  $\mathcal{A}$  intercepts a target transaction  $tx_{\text{Mix}}$  by passively listening to the network (remember that transactions are broadcasted to the Ethereum network in order to be mined, see Section 1.2.2), extracts the zk-proof and primary inputs from  $tx_{\text{Mix}}.data$  and uses these extracted pieces of information in order to create a malicious transaction  $tx_{\text{Mix}}'$ , where the ciphertexts are replaced by arbitrary data. The adversary can then broadcast  $tx_{\text{Mix}}'$  to the network in order for it to be mined. If the malicious transaction gets mined before the legitimate one, the input notes become spent and the ciphertexts are undecryptable making the new notes unredeemable (by any Zeth user!), since all attempts to decrypt the ciphertexts will fail (see Section 2.6).

---

```

TxMalGen( $sk'_{\text{ECDSA}}, nce_{in}, tx_{\text{Mix}}$ )
1 :  $p \leftarrow tx_{\text{Mix}}.gasP + 1$ 
2 :  $l \leftarrow tx_{\text{Mix}}.gasL + 1$ 
3 :  $zdata' \leftarrow tx_{\text{Mix}}.data$ 
4 :  $zdata'.ciphers \leftarrow \mathbb{B}^*$ 
5 :  $tx_{raw} \leftarrow \{nce : nce_{in}, gasP : p, gasL : l, to : tx_{\text{Mix}}.to, val : tx_{\text{Mix}}.val, data : zdata'\};$ 
6 :  $\sigma_{\text{ECDSA}} \leftarrow \text{SigSch}_{\text{ECDSA}}.\text{Sig}(sk'_{\text{ECDSA}}, \text{Keccak256}(tx_{raw}));$ 
7 :  $tx_{final} \leftarrow \{tx_{raw}, v : \sigma_{\text{ECDSA}}.v', r : \sigma_{\text{ECDSA}}.r', s : \sigma_{\text{ECDSA}}.s'\};$ 
8 : return  $tx_{final}$ ;

```

---

Figure A.1: Transaction malleability attack function TxMalGen

As shown on Fig. A.1, during the attack, the adversary extracts the proof and primary inputs from the honest transaction, and replaces the ciphertexts by some arbitrary information. The attacker then formats this data into a transaction that calls the Mix function of **Mixer**, and submits it to the network. While the data fields ( $tx_{\text{Mix}}.data$  and  $tx_{\text{Mix}}'.data$ ) are different, the nullifiers revealed by both transactions are the same (i.e.  $tx_{\text{Mix}}.data.proof = tx_{\text{Mix}}'.data.proof$ , and  $tx_{\text{Mix}}.data.prim = tx_{\text{Mix}}'.data.prim$ ). As a consequence, if the adversary makes sure that  $tx_{\text{Mix}}'$  satisfies all the checks of EthVerifyTx (Section 1.2.2), he can ensure that ZethVerifyTx( $tx_{\text{Mix}}'$ ) will return the same value as ZethVerifyTx( $tx_{\text{Mix}}$ ). Furthermore, if  $tx_{\text{Mix}}'.gasP > tx_{\text{Mix}}.gasP$ , then the adversary maximizes his chances of having his transaction mined first (Section 1.2.2), and so maximizes the chances for the malleability attack to be successful; leading to lost funds on **Mixer**.

**Remark A.1.1.** Note that, although not directly contained within the *data* field of a **Mixer** call transaction, the Ethereum address  $\mathcal{S}_{\mathcal{E}}.Addr$  of the transaction sender is also used by the **Mixer** call (this is either the calling contract's address, or the transaction signer's address recovered as described in Remark 1.2.1). In particular, the balance of this Ethereum address is incremented by the value *vout* by successful Mix calls. If

we again assume the absence of the malleability checks, an attacker could re-sign any **Mixer** call transaction with a key under his control, rebroadcast it as described above, and (with some reasonable probability) become the recipient of any public output value *vout*.

**Remark A.1.2.** We note that the attack described above cannot be prevented by merely substituting a malleable Groth16 zk-SNARK by a simulation-extractable one like e.g. [GM17]. This comes since the attack does not utilise malleability of the proof system, but malleability of data that are broadcasted along with the zk-proof.

## A.2 Solutions to address the transaction malleability attack

### A.2.1 ZeroCash solution

The idea of the solution presented in [BSCG<sup>+</sup>14] is to use a one-time SUF-CMA digital signature and bind its verification key with the zk-proof primary inputs to prevent an adversary from corrupting part of a transaction's data.

Specifically, to transact via **Zeth**, the user first samples a key pair  $(sk, vk)$  for a one-time signature scheme. He then computes the hash  $hsig = CRH(vk)$ , where  $CRH$  is a collision resistant hash function, see [BSCG<sup>+</sup>14], and derives a value  $h_i = PRF_{ask_i}^{pk}(hsig)$ , for each input note (i.e.  $i \in [JSIN]$ ), which acts as a MAC binding  $hsig$  to the address spending key of a note  $(ask_i)$ .

The user then generates the zk-proof with the additional statement that the values  $\{h_i\}_{i \in [JSIN]}$  are computed correctly. He finally uses  $sk$  to sign every value associated with the operation, thus obtaining a signature, which is included, along with the signature verification key  $vk$ , in the transaction. To verify a transaction on the DAP, it is necessary to verify that

- the primary inputs are correctly formatted,
- the Merkle root corresponds to one of the previous states of the Merkle tree,
- the nullifiers have not been declared in a previous transaction,
- the  $hsig$  is correctly computed from  $vk$ , and
- both the zk-proof and the one-time signature verifications pass successfully.

Now, an adversary trying to carry out the aforementioned attack has to either change the ciphertexts or the encryption key. Nevertheless, doing so should lead to the one-time signature verification to fail or should yield an attack that breaks the UF-CMA property of the one-time signature (as this corresponds to creating a forgery on a different message, not changing the signature). Thereby, the adversary also has to modify the signature, however he does not know the one-time signing key used by the creator of the targeted transaction. As such, the adversary needs to use another signing key pair, however

2120 this leads to the check verifying that  $hsig$  is correctly computed to fail. If the adversary  
 2121 attempts to change  $hsig$ , the zk-proof verification fails as the NP-statement has changed.  
 2122 Hence, any attempt to carry out a malleability attack results in the violation of at least  
 2123 one check in the verification of the transaction on the DAP. The solution presented  
 2124 effectively solves the transaction-malleability attack initially described.

2125 **Remark A.2.1.** The one-timeness property of the signature scheme was required in  
 2126 **ZeroCash** to retain anonymity. It also makes analysing non-adaptive adversary sufficient.  
 2127 As **Ethereum** transaction senders need to pay the gas cost associated with their trans-  
 2128 actions, the senders are not anonymous. This said, making sure that **Zeth** is designed  
 2129 with anonymity in mind is worth the effort in order to minimize information leakages  
 2130 and be ready if/when **Ethereum** incorporates protocol changes that enable anonymous  
 2131 transactions.

## 2132 A.2.2 Zcash’s solution

2133 In addition to the changes aforementioned, **Zcash**’s solution [ZCa19] also consists of:

- Redefining the variable  $hsig$  as,

$$hsig = \text{CRH}(\text{randomSeed}, \{nf_i\}_{i \in [\text{JSIN}]}, vk)$$

2134 for some random seed  $\text{randomSeed}$ .

- Defining a new random variable  $\phi$  and using it with  $hsig$ , as key and input of a  
 2135 PRF respectively, to compute the identifier of each output notes  $\rho_j$  ( $j \in [\text{JSOUT}]$ )  
 2136 and ensure their uniqueness (with overwhelming probability).  
 2137

2138 These changes were made to prevent the Faerie Gold attack [ZCa19, Section 8.4], as well  
 2139 as to prevent linkability: if  $hsig$  were repeated in two transactions, the circuit would  
 2140 leak, via  $\{h_i\}_{i \in [\text{JSIN}]}$ , the fact that the input notes in both transactions were spent with  
 2141 the same  $ask_i$  (if that were the case).

2142 More particularly, using the input notes’ nullifiers to derive  $hsig$  ensures that  $hsig$  is  
 2143 unique with overwhelming probability for all *accepted* transaction. Furthermore, us-  
 2144 ing  $\text{randomSeed}$  ensures the uniqueness of  $hsig$  for transactions *in transit* (as before  
 2145 validation there may be several in transit transactions with the same set of nullifiers).

## 2146 A.2.3 Solution on Ethereum

2147 As described in the **Ethereum** yellow paper [Woo19, Appendix F], **Ethereum** transactions  
 2148 are ECDSA signed. Further, as described in Section 2.3, the one-time signature used to  
 2149 sign the Mix data also signs the **Ethereum** address used to sign the transaction. As such,  
 2150 any modification to the transaction object will result in a new transaction hash, and  
 2151 any attempt to sign the transaction with a different ECDSA key will be rejected by the  
 2152 **Mixer** contract (see Section 2.5). We thereby conclude that the one-time signature used

2153 to sign the transaction data does not need to be **SUF-CMA**, but *only needs to achieve*  
 2154 **UF-CMA**.

2155 Specifically, carrying out any change on the one-time signature will change the  
 2156 **Ethereum** transaction data and result in a failure to verify the **ECDSA** signature of  
 2157 the **Ethereum** transaction. To obtain a new valid signature on this transaction, the  
 2158 adversary needs to break the **UF-CMA** property of the **ECDSA** signature scheme or use  
 2159 another **ECDSA** keypair to sign the transaction. In the last case, the one-time signature  
 2160 will no longer be valid.

2161 Note that including the **Ethereum** transaction sender in the data to be signed by the  
 2162 one-time signature scheme also addresses the possible attack described in Remark A.1.1.  
 2163 An attacker trying to resign the same **Ethereum** transaction with a different key will  
 2164 cause **Mixer** to reject the transaction when the one-time signature is checked.

**Remark A.2.2.** We note that the transaction malleability issue can also be addressed  
 in another way. In fact, one could use the **ECDSA** signatures on **Ethereum** transactions  
 to fix all inputs and ciphertexts, and then tie the sender of the **Ethereum** transaction to  
 the zk-snark by putting the sender address  $\mathcal{S}_{\mathcal{E}}.Addr$  in  $hsig$ . In other words, it is also  
 possible to define  $hsig$  as:

$$hsig = \text{CRH}(\{nf_i\}_{i \in [\text{JSIN}]}, \mathcal{S}_{\mathcal{E}}.Addr)$$

2165 As such, if an attacker extracts the ciphertexts of a  $tx_{\text{Mix}}$  transaction in order to craft  
 2166 another malicious transaction  $tx_{\text{Mix}}'$ , the key-pair used to sign  $tx_{\text{Mix}}'$  differs from the one  
 2167 used to sign  $tx_{\text{Mix}}$ , which changes the transaction sender address recovered on **Mixer**. As  
 2168 a consequence, the check on  $hsig$  would fail on the **Mixer**, invalidating the transaction,  
 2169 and preventing the attack.

2170 While such a solution would avoid the need to generate one-time signing keys and  
 2171 could avoid a signature check in the **Mixer**, it would also require every **Zeth** user to  
 2172 have an **Ethereum** account. Doing so, would be a major hindrance toward the design of  
 2173 mechanisms aiming to provide anonymity to **Zeth** transactions initiators. In fact, the  
 2174 addressing scheme used in **Zeth** along with the solution to the malleability introduced in  
 2175 **Zcash** makes it possible to generate raw **Zeth** transactions without having an **Ethereum**  
 2176 account. These raw transactions could then be broadcasted – to a set of **Ethereum** user  
 2177 nodes – on an anonymous p2p network, before being finalized and submitted to the  
 2178 **Ethereum** network by **Ethereum** users who would be rewarded according to an incentive  
 2179 structure. While such a protocol is outside of the scope of this document, it shows that  
 2180 defining  $hsig$  using the senders address alters the flexibility of **Zeth**; hence this solution  
 2181 has not been favoured.

## Appendix B

### Double spend attack on equivalent class

The primary inputs of our zk-SNARK are elements of  $\mathbb{F}_r$  and they can be written over  $\text{FIELDLEN}$  bits. Note that the projection of  $\mathbb{B}^{\text{FIELDLEN}}$  onto  $\mathbb{F}_r$  formed by interpreting elements in  $\mathbb{B}^{\text{FIELDLEN}}$  as  $\text{FIELDLEN}$ -bit numbers and reducing modulo  $r$ , is surjective.

When we pass the primary inputs to the **Mixer** contract, they are interpreted as elements of  $\mathbb{B}^{\text{ETHWORDLEN}}$ , and  $\mathbb{B}^{\text{FIELDLEN}} \subset \mathbb{B}^{\text{ETHWORDLEN}}$ . As previously noted, this means that there exist pairs of elements in  $\mathbb{B}^{\text{ETHWORDLEN}}$  with the same projection in  $\mathbb{F}_r$ . An adversary could make use of this to perform a double spend attack.

Indeed, to check that a note is not double spent, the contract stores the nullifiers of spent notes (as elements of  $\mathbb{B}^{\text{ETHWORDLEN}}$ ) and verifies that the nullifier of the note to be spent is not stored. The adversary could thus modify the nullifier to a different value with the same projection. As the SNARK verification operates in  $\mathbb{F}_r$ , the proof would still be valid. However, the value stored for this nullifier would be different from the adversarial one. Hence, the nullifier would be validated, the transaction would succeed and the note would be double spent. In practice, the adversary can perform the attack by simply adding  $r$  to one of the elements representing the nullifier.

To prevent this attack, the contract checks that all primary inputs are elements of  $\mathbb{F}_r$ , that is to say that they are smaller than  $r$ . As one may see, the attack described above is not due to the packing of hash digests into field elements but to the contract storage of field elements as **Ethereum** words.

## 2204 Appendix C

# 2205 Side-channel attacks and 2206 information leaks

2207 The following subsections describe several side-channel attacks and possible weaknesses  
2208 that implementers should be aware of and attempt to mitigate.

2209 We consider cases in which the attacker is able to observe the RPC communications  
2210 between **Zeth** client software, and Ethereum P2P nodes. This situation may occur if an  
2211 observer is able to monitor the network traffic between the Ethereum node and the **Zeth**  
2212 client software, or if the Ethereum node itself is compromised.

### Note

In this discussion, we do not consider adversaries with physical access to the machine running the client software. Such adversaries could make precise measurements of timing, power-consumption or other physical quantities that could reveal fine-grained details of the operations being carried out by the software, or the data it is operating on. Protecting against attacks of this kind often involves implementation techniques such as: avoiding branches based on private data, being careful with memory access patterns, and making all operations constant time, to only name a few. We leave consideration of these attacks and prevention methods for a future discussion.

2213

## 2214 C.1 Counterfeit data

2215 Malicious Ethereum nodes or attackers able to compromise the network have the oppor-  
2216 tunity to send invalid data to RPC clients. This could be used to inject invalid data  
2217 into the client's record of state, which could prevent it from generating valid **Mix** calls  
2218 or allow it to be identified in the future. In general, data from any remote host should  
2219 be treated as malicious, unless accompanied by evidence that convinces the client of its  
2220 authenticity.



2221 In the case of Ethereum event logs (the main source of data used to track the on-  
2222 chain state – see Section 4.1.1 for details), clients **MUST** leverage the consensus evidence  
2223 and block headers to verify that log data is genuine and has been committed to the  
2224 blockchain. See Section 1.2.3 for further information about how such data is secured.

## 2225 C.2 Data leaked during synchronization

2226 In order to receive private payments and keep their local data up-to-date, **Zeth** client  
2227 software **MUST** scan the blockchain and process *all* the event data emitted by **Mixer**  
2228 during Mix calls (as described in Section 4.1.1). There are several issues to consider  
2229 when determining exactly how and when this “synchronization” takes place.

2230 Client implementations that only connect to the RPC endpoint in response to user  
2231 input, or in preparation for performing a Mix call, may leak information. Observers may  
2232 deduce that such client are likely to be the recipient of a recent or upcoming transaction,  
2233 or that they may be about to perform a Mix call.

2234 Similarly, payment provider software that only listens for events when awaiting a  
2235 transaction, and remains disconnected otherwise, may reveal that it is the recipient of  
2236 an upcoming transaction, and possibly *which* transaction or block it was paid by (based  
2237 on when it stops listening).

2238 Further, consider wallet software that performs RPC operations to explicitly wait  
2239 for the Ethereum transaction corresponding to a specific Mix call. This would most  
2240 likely be for transactions emitted by the **Zeth** client, in order to inform the user and  
2241 update the wallet state once the payment is complete (but could possibly happen on  
2242 the receiver side, if he somehow knows the ID of the transaction of interest – e.g. via  
2243 off-chain communication with the sender). If such a *wait* procedure is implemented by  
2244 querying the status of a specific transaction by its ID, or by listening for blocks *until*  
2245 the transaction of interest is received, the connected Ethereum node may infer that this  
2246 client is interested in the transaction, and likely to be the sender or recipient.

2247 Consider a client which periodically connects to some Ethereum node and requests  
2248 all relevant data from the last block it saw, up to the latest block available. Each client  
2249 will have information up to some block  $n$  (where  $n$  varies per client), and  $n$  is known to  
2250 the Ethereum node that served the client. The client could then potentially be identified  
2251 by  $n$  (even if it hides its IP for each connection) since a client that connects and queries  
2252 **Zeth** transactions from block  $n + 1$  reveals that it is one of the clients who synced up to  
2253 block  $n$  when it last connected.

2254 Note that, if the client always broadcasts the Mix transaction via this same Ethereum  
2255 node, then the Ethereum node may already deduce that the client is the sender. However,  
2256 implementations may wish to use techniques (such as sending transactions from other  
2257 nodes or hiding their IP address in other way) to obfuscate any relationship between  
2258 transactions and the clients that originated them.

### C.3 Queries on successful decryption

The event data emitted by  $\widetilde{\text{Mixer}}$  contains the note data for new commitments, encrypted using a key derived from the recipients' public key. As described in Section 2.6, clients scan the blockchain for these events and attempt to decrypt the ciphertext using their secret decryption keys. If they are successful, they are the recipient of the note and can try to parse the plaintext to extract the secret note data.

When decryption is successful and the note data has been extracted from the plaintext (we discuss parsing failure in Appendix C.4), clients **MUST** check that this note data does indeed open the commitment for the note.

A naive implementation of this check could query the state of  $\widetilde{\text{Mixer}}$  via RPC to check the relevant entry in the set of commitments. However, this would reveal to an observer that the client had successfully decrypted and parsed the corresponding ciphertext, and was therefore the recipient of that note.

For this reason, the protocol specifies that  $\widetilde{\text{Mixer}}$  **MUST** emit events informing clients of new commitment values and locations in the Merkle tree. Clients **MUST** consume *all* such data to maintain their view of contract state (as described in Appendix C.2). Further, clients **MUST** attempt to decrypt *all* ciphertexts and, for successful decryptions, **MUST** verify that the plaintext opens the note's commitment. This avoids the need for any extra RPC queries that would reveal which ciphertexts were successfully decrypted.

#### Note

Emitting events containing all data necessary to carry out the local checks implemented in the wallet is a way to enforce that all wallets behave exactly the same to the eyes of network (passive) adversaries (regardless whether the user is the recipient of a note or not).

### C.4 Invalid ciphertext

The attack described in [TBP20, Section 4.2.1] illustrates the importance of correctly handling invalid data in client software. A so-called “REJECT Attack” is described whereby an attacker creates a  $\text{Mix}$  call with specially crafted ciphertext. The ciphertext can be successfully decrypted by the correct recipient – that is, the plaintext note is encrypted with an encryption key derived from the recipients public key – but the corresponding plaintext is invalid and cannot be parsed correctly by the recipient.

#### Note

Note that the above is possible because the plaintext is neither verified by the circuit encoding  $\mathbf{R}^Z$ , nor by the contract (which is unable to decrypt it). Hence,  $\widetilde{\text{Zeth}}$  allows such transactions with malicious ciphertexts to be accepted by the  $\widetilde{\text{Mixer}}$  contract, and clients must handle this case with care.

2287 In the case described in [TBP20], there is no distinction between “client” or “wallet”  
2288 software, and the underlying P2P nodes. Before a fix was applied (see [zcab]), nodes  
2289 explicitly rejected transactions of the above form, proving to their peers that they were  
2290 able to decrypt the ciphertext and were therefore the intended recipient.

2291 In **Zeth**, P2P nodes and wallet software are separated, so there will be no explicit  
2292 rejection of the transaction. However, careless error handling (such as exceptions which  
2293 causes the RPC connection to be closed) could potentially be detected by the connected  
2294 Ethereum node. As in the “REJECT Attack”, this reveals that the connected RPC  
2295 client is the intended recipient of a transaction, and the owner of the corresponding  
2296 encryption key.

## 2297 C.5 Using (and retrieving) nullifiers

2298 Any non-trivial wallet implementation will need to track which of the user’s **Zeth** notes  
2299 have been spent, and which are still available. Naturally, the wallet software could mark  
2300 the notes as it broadcasts transactions that spend them. However, this approach is  
2301 subject to several problems.

2302 Firstly, for each note spent, the client software must record the ID of the spending  
2303 transaction, in order to track it and confirm that it is accepted into a block. Once each  
2304 spending transaction is accepted the client can finally mark the appropriate **Zeth** notes  
2305 as “spent”. This requires significant complexity in order to asynchronously mark the  
2306 notes, and to deal with the issues described in Appendix C.2.

2307 Secondly, this approach does not support multiple wallets using the same key, or  
2308 wallets being restored from **Zeth** addresses. A user that wishes to rebuild his wallet  
2309 (see the discussion in Section 4.1.4), or check for any spending activity by other wallets,  
2310 would not be able to do so by simply scanning the blockchain.

2311 By using the nullifiers passed to **Mix** calls, clients can determine the availability of  
2312 notes in a more robust way. That is, to determine whether a note is spent or available,  
2313 the client can compute the nullifier and check whether that nullifier has been seen by  
2314 the **Mixer** contract.

2315 In a similar way to Appendix C.3, queries to **Mixer** for specific nullifiers reveals  
2316 to observers that the client was the sender of any previous or future transaction that  
2317 generates such a nullifier. To mitigate this, **Mixer** MUST include nullifier values in the  
2318 event data it emits, and clients SHOULD use this to track which of their notes are spent.  
2319 This MUST happen as part of the regular sync operation, so that no extra RPC traffic is  
2320 generated and observers cannot distinguish between clients that do and do not recognize  
2321 any given nullifier. Note that this approach also supports tracking spent notes from  
2322 multiple wallets, and rebuilding wallets by re-syncing the blockchain.

## C.6 Proof generation

Generation of the zero-knowledge proofs, required for valid Mix calls, is a very computationally intensive process. The proof generation itself does not require any communication with external parties, and so may not directly leak information about the client, but implementers should consider some indirect ways in which information may be leaked.

Implementers may also wish to consider the possible indirect impact of proof generation on the RPC channel. For example, a client that “waits” for proof generation without servicing the RPC connection may fail to respond to (or take significantly longer to respond to) new log events. The connected Ethereum node might then deduce that its peer is generating a proof and therefore likely to be the sender of an upcoming transaction.

### Note

As stated in the introduction to this chapter, this discussion does not consider general timing attacks. We mention this extreme case of a client that completely stalls during proof generation only to illustrate how a poor implementation may leak information to its RPC peer.

In the case where proof generation is carried out on some external host, or by an external process on the same host, there may be a risk of network traffic or other IPC traffic being observed. If an observer can detect that a given client is communicating with a prover process, it can reliably deduce that the client will be the sender of an upcoming transaction.

An observer able to see the content of the communication between the wallet and prover process will also gain knowledge of the auxiliary inputs to the proof (including the data required to spend the input notes and secret attributes of the output notes). It is therefore important to secure any such connection, protect any prover process from being maliciously modified or observed, and to ensure that wallets only communicate with trusted processes.

## C.7 Simple mixer calls

The public parameters to a Mix call can reveal information about the nature of a transaction, even though they do not reveal recipient details or note amounts. For example, a Mix call in which  $\text{Mix}_{in}.primIn.vout = 0$  and  $\text{Mix}_{in}.primIn.vin \neq 0$  may indicate a simple “deposit” of funds into the mixer. Similarly, if both  $\text{Mix}_{in}.primIn.vout$  and  $\text{Mix}_{in}.primIn.vin$  are zero, the transaction must be spending only notes already within **Mixer**, into new notes. Finally, if  $\text{Mix}_{in}.primIn.vin = 0$  and  $\text{Mix}_{in}.primIn.vout \neq 0$ , the sender may be performing a simple “withdrawal” of funds from some existing notes.

A Mix call can combine all of the above logical operations in a single transaction. That is, it can deposit value into the mixer, spend existing notes, create new notes, and withdraw value from **Mixer** *at the same time*. Combining logical operations in this way

2356 makes it much more difficult for an observer to attribute a specific purpose to the `Mix`  
2357 call.

2358 Clients can also perform `Mix` calls in which  $vin = vout = 0$  and 0-valued notes are  
2359 created from other 0-valued notes. Such “dummy” self-payments can further obfuscate  
2360 the activity of a wallet, by adding “noise” to the system. Note, however, that the gas  
2361 cost for such transactions must still be paid.

2362 Wallet implementations **SHOULD** encourage the use of these complex calls where pos-  
2363 sible, either via the user interface or by automatically adding complexity to transactions,  
2364 and **SHOULD** support features such as adding “noise”<sup>1</sup> if the user wishes to pay for extra  
2365 protection of this kind.

### 2366 C.7.1 Small anonymity sets

2367 Until there is a large number of commitments and users of the mixer, it may be easy  
2368 for an observer to infer some of the private data that is intended to be hidden by mixer  
2369 calls.

2370 In the simple case, if there are very few commitments in the `Mixer`’s Merkle tree, an  
2371 attacker has a small list of candidate commitments that are being spent by subsequent  
2372 `Mix` calls. Similarly, if the number of distinct Ethereum addresses that have been used  
2373 to call `Mixer` is very small, observers can trace the original source of funds subsequently  
2374 withdrawn to a small set of original depositors.

2375 Client software may wish to track metrics about the `Mixer` state, and either prevent  
2376 certain actions or design the user interface to discourage users<sup>2</sup> from creating trans-  
2377 actions whose features can be identified with high probability. We provide below a  
2378 non-exhaustive list of metrics of interest:

- 2379 • **Number of commitments.** If there is a low absolute number of commitments,  
2380 clearly any non-zero output must spend one of these (although we note that only  
2381  $vout$  can be publicly known to be non-zero).
- 2382 • **Number of unspent commitments.** If  $\#Comms - \#Nulls$  is small and a new  
2383 commitment is created and then spent, observers can deduce that there is a high  
2384 chance that the spend operation targeted the new commitment.
- 2385 • **Number of Ethereum addresses.** While very few distinct addresses (or groups  
2386 of addresses that are not associated) have used the contract, observers can de-  
2387 duce that subsequent `Mix` calls are likely to spend commitments created by clients  
2388 associated with one of this small set of addresses.

2389 The set of Ethereum addresses that have interacted with the contract can leak data  
2390 in other ways. An Ethereum address that withdraws value from the contract, but has not  
2391 previously been used to make a `Mix` call (or a `Mix` call that deposits value into `Mixer`),

---

<sup>1</sup>By randomly scheduling dummy payments, for instance

<sup>2</sup>By, for example, displaying warning messages and/or asking the user for confirmation

2392 must have been the recipient of zeth notes created by a previous depositor. The details  
2393 may not be directly available to an observer, but this is another example of information  
2394 which could be combined with other leaked data to infer connections between entities  
2395 and transactions.

## Appendix D

# Security proofs of Blake2

This appendix proves the collision resistance, PRF-ness, binding and hiding properties of the Blake2 hash function in the Weakly Ideal Cipher model (WICM, see [LMN16]). The proofs use definitions and results of Luykx et al. [LMN16], regarding the indistinguishability of Blake2 and a random oracle in the Weakly Ideal Cipher Model (WICM). In the following, we assume that the optimization of Blake2 for 8- to 32-bit platforms is as secure as Blake2 as described in [LMN16].

### D.1 Security model of Blake2

The security analysis treats Blake2 as hash function built on top of a block-cipher-based compression function in the WICM (which derives from the Ideal Cipher Model). In this section, we present the WICM and prove that Blake2 is a collision resistant PRF, and thus a commitment scheme.

#### D.1.1 Weakly Ideal Cipher Model

The research community believes that Blake’s underlying block cipher has no known weaknesses and could reasonably be modeled as an ideal cipher [LMN16, Section 2.1]. However, Blake2 admits weak keys with a specific structure [LMN16, Section 2.1]. Blake2 is therefore more appropriately analysed in the WICM, which is an extension of the Ideal Cipher Model that represents a block cipher as a set of independent random permutations [HKT11]. The WICM may also be viewed as a specialization for Blake2 of the Weak Cipher Model [MP15], which aims to be realistic by modeling particular characteristics, invariants or properties a block cipher may have.

A number of definitions in what follows are quoted directly from Luykx et al. [LMN16].

**The Weakly Ideal Cipher Model.** Let  $\mathcal{W}$  and  $\mathcal{S}$  be the following partition of  $\mathbb{B}^{2 \cdot \text{ol}}$  into weak and strong sets, where  $w$  is the word length ( $16 \cdot w = 2 \cdot \text{ol}$ ):

$$\mathcal{W} = \left\{ \text{aaaabbbbccccdddd} \in \mathbb{B}^{2 \cdot \text{ol}} \mid a, b, c, d \in \mathbb{B}^w \right\}$$

$$\mathcal{S} = \mathbb{B}^{2 \cdot \text{ol}} \setminus \mathcal{W}$$

Let  $\mathcal{BLC}(2 \cdot \text{ol}, 2 \cdot \text{ol})$  denote the set of all block ciphers  $E : \mathbb{B}^{2 \cdot \text{ol}} \times \mathbb{B}^{2 \cdot \text{ol}} \rightarrow \mathbb{B}^{2 \cdot \text{ol}}$ . Define  $\mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  as the set of all block ciphers  $E \in \mathcal{BLC}(2 \cdot \text{ol}, 2 \cdot \text{ol})$  with the additional restriction that  $E(k_w, \cdot)$  is  $\mathcal{W}$ - and  $\mathcal{S}$ -subspace invariant for all keys  $k_w \in \mathcal{K}_{\text{weak}}$ . That is, inputs in  $\mathcal{W}$  map to  $\mathcal{W}$ , and likewise for  $\mathcal{S}$ . Here,  $\mathcal{K}_{\text{weak}}$  is the set of weak keys, defined as

$$\mathcal{K}_{\text{weak}} = \left\{ k = \text{kkkkkkkkkkkkkkkk} \in \mathbb{B}^{2 \cdot \text{ol}} \mid k \in \mathbb{B}^w \right\}.$$

2419 A random  $E \leftarrow \$\mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  can now be modeled as follows:

- 2420 • on input of  $(k, x) \in \mathcal{K}_{\text{weak}} \times \mathcal{W}$ ,  $E$  generates its response  $y$  randomly from  $\mathcal{W}$  up  
2421 to repetition;
- 2422 • on input of  $(k, x) \in \mathcal{K}_{\text{weak}} \times \mathcal{S}$ ,  $E$  generates its response  $y$  randomly from  $\mathcal{S}$  up to  
2423 repetition.

2424 For key values  $k \in \mathbb{B}^{2 \cdot \text{ol}} \setminus \mathcal{K}_{\text{weak}}$ ,  $E$  behaves like an ideal cipher: it either outputs a  
2425 new random value or if the key-message-image tuple has already been queried the tuple's  
2426 image. The case of inverse queries is analogous.

Blake2C is defined over the following domains and codomain:

$$\text{Blake2C} : \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol}) \times \mathbb{B}^{\text{ol}} \times \mathbb{B}^{2 \cdot \text{ol}} \times \mathbb{B}^{\text{ol}/4} \times \mathbb{B}^{\text{ol}/4} \rightarrow \mathbb{B}^{\text{ol}}$$

2427 We write  $\text{Blake2C}_E(h, m, t, f)$  for the output of the Blake2 compression function, defined  
2428 over encryption scheme  $E$  on inputs  $h$ ,  $m$ ,  $t$  and  $f$ . The compression function, in par-  
2429 ticular, computes the state  $x = (h \parallel \text{pad}_{\text{ol}/2}(0) \parallel t \parallel f) \oplus (\text{pad}_{\text{ol}}(0) \parallel \text{IV})$  for some  $\text{IV}$ . It then  
2430 encrypts  $x$  under  $m$  (where  $m$  is treated as a key for the encryption) and splits  $E(m, x)$   
2431 in two same size variables, the left part  $l_E$  and right part  $r_E$ . It finally outputs  $l_E \oplus r_E \oplus h$ .  
2432

2433 Zeth uses the Blake2 compression function with a fixed encryption scheme  $E^*$  based on  
2434 ChaCha stream cipher [Ber08a]. Thus, we write  $\text{Blake2C}(h, m, t, f) = \text{Blake2C}_{E^*}(h, m, t, f)$ .

2435 **Indifferentiability.** One way to measure the extent to which a certain cryptographic  
2436 function behaves like a random function is via the indistinguishability framework where  
2437 a distinguisher is given oracle access to either the cryptographic function or the random  
2438 function with the goal of determining which one it has access to.

**Definition D.1.1.** Let  $\mathcal{C}$  be a construction with oracle access to an ideal primitive  $\mathcal{P}$ . Let  $\mathcal{R}$  be an ideal primitive with the same domain and codomain as  $\mathcal{C}$ . Let  $\text{Sim}$  be a simulator with the same domain and codomain as  $\mathcal{P}$  with oracle access to  $\mathcal{R}$ , and let  $\text{Dist}$  be a PPT distinguisher. The indifferentiability advantage of  $\text{Dist}$  is defined as:

$$\text{Indiff}_{\mathcal{C}^{\mathcal{P}}, \text{Sim}}(\text{Dist}) = \left| \Pr \left[ \text{Dist}^{\mathcal{C}^{\mathcal{P}}, \mathcal{P}} = 1 \right] - \Pr \left[ \text{Dist}^{\mathcal{R}, \text{Sim}^{\mathcal{R}}} = 1 \right] \right|$$



2439 The distinguisher  $\text{Dist}$  can query both its left oracle (either  $\mathcal{C}$  or  $\mathcal{R}$ ) and its right  
 2440 oracle (either  $\mathcal{P}$  or  $\text{Sim}$ ). We refer to  $\mathcal{C}^{\mathcal{P}}$ ,  $\mathcal{P}$  as the real world, and to  $\mathcal{R}$ ,  $\text{Sim}^{\mathcal{R}}$  as the  
 2441 simulated world; the distinguisher  $\text{Dist}$  converses with either of these worlds and its goal  
 2442 is to tell both worlds apart.

**Theorem D.1.1** (Indifferentiability of Blake2 [LMN16]). *Let an encryption scheme  $\mathbf{E} \leftarrow \$ \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  be a weakly ideal cipher, and consider the hash function  $\text{Blake2}_{\mathbf{E}}$  that internally uses  $\mathbf{E}$ . There exists a simulator  $\text{Sim}$  such that for any distinguisher  $\text{Dist}$  with total complexity  $q$ , we have:*

$$\text{Indiff}_{\text{Blake2}_{\mathbf{E}}, \text{Sim}}(\text{Dist}) \leq \frac{\binom{q}{2}}{2^{2\text{ol}}} + \frac{2\binom{q}{2}}{2^{\text{ol}}} + \frac{q}{2^{\text{ol}/2}}$$

2443 where  $\text{Sim}$  makes at most  $O(q^3)$  queries to a random function  $\mathcal{R}$ .

2444 *Proof.* See [LMN16, Corollary 1]. □

2445 For asymptotic security, we assume the distinguisher to be PPT and that the number  
 2446 of queries made is polynomial  $q \leq \text{poly}(\text{ol})$ .

2447 **Additional remarks.** Luykx et al. [LMN16] remark that, by resorting to the WICM,  
 2448 they do not make stronger assumptions than those used in previous results (ICM), and,  
 2449 despite the fact that they give distinguishers more power (by weakening the cipher),  
 2450 they are able to get similar results.

## 2451 D.2 Security proofs

### 2452 D.2.1 Blake2 is a PRF

2453 Luykx et al. already prove the PRFness of Blake2 *keyed* hash function in the multi-key  
 2454 setting.

**Definition D.2.1** (PRF in multi-key setting [ML15]). Let  $\mu \geq 1$  and  $k \leftarrow \$ \mathcal{K}^{\mu}$ . Let  $\mathcal{C}$  be a keyed construction with key space  $\mathcal{K}$  and with oracle access to an ideal primitive  $\mathcal{P}$ . Let  $\mathcal{R}_1, \dots, \mathcal{R}_{\mu}$  be random functions with the same domains and ranges as  $\mathcal{C}_{k_1}, \dots, \mathcal{C}_{k_{\mu}}$ . Let  $\text{D}$  be a distinguisher. The PRF distinguishing advantage of  $\text{D}$  is defined as,

$$\text{PRF}_{\mathcal{C}^{\mathcal{P}}}(\text{D}) = \left| \Pr[\text{Dist}_{\mathcal{C}_{k_1}^{\mathcal{P}}, \dots, \mathcal{C}_{k_{\mu}}^{\mathcal{P}}, \mathcal{P}} = 1] - \Pr[\text{Dist}_{\mathcal{R}_1, \dots, \mathcal{R}_{\mu}, \mathcal{P}} = 1] \right|$$

Blake2 supports keyed hashing by simply prepending the key to the message:

$$\text{Blake2}_{\mathbf{E}, k}(m) = \text{Blake2}_{\mathbf{E}}(k \| 0^{2\text{ol}-\text{kl}} \| m)$$

2455 where  $\text{kl} \leq 2\text{ol}$  denotes the key size. In other words, the key gets processed as other data  
 2456 and the HAIFA counter and flags are designated to the key in a similar fashion as if they  
 2457 were for normal data blocks.

**Theorem D.2.1** (PRF-security of Blake2 keyed mode [LMN16]). *Let  $\mu \geq 1$  and let  $k \leftarrow \$ (\mathbb{B}^{\text{kl}})^\mu$ . Let an encryption scheme  $E \leftarrow \$ \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  be a weakly ideal cipher, and consider the keyed hash function  $\text{Blake2}_{E,k}$  that internally uses  $\text{Blake2C}_E$  that internally uses  $E$ . For any distinguisher  $\text{Dist}$  with total complexity  $q$ :*

$$\text{PRF}_{\text{Blake2}_{E,k}}(\text{Dist}) \leq \frac{\binom{q}{2}}{2^{2\text{ol}}} + \frac{2\binom{q}{2}}{2^{\text{ol}}} + \frac{q}{2^{\text{ol}/2}} + \frac{\mu q}{2^{\text{kl}}} + \frac{\binom{\mu}{2}}{2^{\text{kl}}}$$

2458 *Proof.* See [LMN16, Corollary 3]. □

2459 **Remark D.2.2.** We can note that in the case of keyed hashing, the key is padded only  
 2460 to be processed in a single block to differentiate the key from the message. The security  
 2461 proof of Theorem D.2.1 does not rely on this padding and as such also works with no  
 2462 padding.

**Theorem D.2.2** (PRF-security of Blake2 with a single key [LMN16]). *Let  $k \leftarrow \$ \mathbb{B}^{\text{kl}}$ . Let an encryption scheme  $E \leftarrow \$ \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  be a weakly ideal cipher, and consider the keyed hash function  $\text{Blake2}_E(k, \cdot) = \text{Blake2}_E(k \| \cdot)$  that internally uses  $\text{Blake2C}_E$  that internally uses  $E$ . For any distinguisher  $\text{Dist}$  with total complexity  $q$ :*

$$\text{PRF}_{\text{Blake2}_E}(\text{Dist}) \leq \frac{\binom{q}{2}}{2^{2\text{ol}}} + \frac{2\binom{q}{2}}{2^{\text{ol}}} + \frac{q}{2^{\text{ol}/2}} + \frac{q}{2^{\text{kl}}}$$

2463 *Proof.* See Remark D.2.2 and Theorem D.2.1 with  $\mu = 1$ . □

2464 **Remark D.2.3.** Since we analyse the security of Blake2 asymptotically, we assume that  
 2465 for a security parameter  $\lambda$  holds  $\text{ol} = \mathcal{O}(\lambda)$ ,  $\text{kl} = \mathcal{O}(\lambda)$ , and  $q = \text{poly}(\lambda)$ .

## 2466 D.2.2 Proof of Blake2 collision resistance

2467 We want to prove here the collision resistance of Blake2. To do so, we are going to  
 2468 prove by contradiction that if Blake2 is not collision resistant, it is not indifferentiable  
 2469 according to Definition D.1.1.

2470 **Theorem D.2.3.** *Blake2 is collision resistant.*

2471 *Informal proof.* Let us assume that there exists a PPT adversary  $\mathcal{B}$  which breaks the  
 2472 collision resistance of Blake2. We build an adversary  $\mathcal{A}$  that uses this adversary to  
 2473 differentiate between the real and simulated worlds. More particularly,  $\mathcal{A}$  gets left and  
 2474 right oracles (see [LMN16, Figure 3]), which are either an oracle for a hash function and  
 2475 for a weakly ideal block cipher or a random oracle and an encryption simulator with  
 2476 oracle access to the random oracle.

2477 On each  $\mathcal{B}$ 's query  $m_i$ ,  $i \in \{1, \dots, q\}$ ,  $\mathcal{A}$  passes them to his left oracle and returns  
 2478 the answer  $h_i$  to  $\mathcal{B}$ . Eventually, if  $\mathcal{B}$  finds a collision, that is a pair  $(m_i, m_j)$  such that  
 2479  $m_i \neq m_j$  and  $h_i = h_j$ ,  $\mathcal{A}$  guesses that his oracles were real; else  $\mathcal{A}$  returns a random

2480 guess. Otherwise  $\mathcal{A}$  guesses his oracles were simulated – if the left oracle was a random  
 2481 oracle, the probability of finding a collision would be negligible for  $q \leq \text{poly}(\lambda)^1$ .  
 2482 On the other hand,  $\mathcal{B}$  finds a collision with non-negligible probability if the oracles  
 2483 were real. Hence,  $\mathcal{A}$  wins the indistinguishability game with non-negligible advantage,  
 2484 which is a contradiction.  $\square$

### 2485 D.2.3 Blake2 as a commitment scheme

2486 We prove here that Blake2 is a commitment scheme, i.e. is binding and hiding. To do so  
 2487 we rely on the previous results that Blake2 is collision resistant and a PRF.

2488 **Theorem D.2.4.** *Let  $E \leftarrow \$\mathcal{BLC}(2\text{ol}, 2\text{ol})$  and for a message  $x \in \mathbb{B}^*$  and randomness*  
 2489  *$r \in \mathbb{B}^l$  commitment to  $x$  using  $r$  be  $\text{ComSch.Com}(x; r) = \text{Blake2}_E(r \| x)$ . Then  $\text{ComSch}$*   
 2490 *is hiding and binding.*

2491 *Informal proof. Hiding.* A commitment scheme  $\text{ComSch}$  is computationally hiding if,  
 2492 knowing two potential openings, a PPT adversary cannot distinguish which was com-  
 2493 mitted. Let us assume that there exists a PPT adversary  $\mathcal{B}$  which breaks the hiding  
 2494 property of Blake2 with a non-negligible advantage  $\eta$ . We build an adversary  $\mathcal{A}$  that  
 2495 uses  $\mathcal{B}$  to break the PRF property of Blake2 with advantage  $\eta/2$ .

2496 First, the PRF game is initiated, that is, the challenger chooses a random encryption  
 2497 scheme  $E$  and key  $k \in \mathbb{B}^l$  and instantiates two oracles  $O^{\text{Blake2}_k} = \text{Blake2}_E(k, \cdot)$  and  $O^R$   
 2498 a random function. The challenger picks an oracle randomly and gives  $\mathcal{A}$  access to it.  
 2499  $\mathcal{B}$  sends  $q$  oracle queries  $m_1, \dots, m_q$  to  $\mathcal{A}$  (adaptively) who extends them with random  
 2500  $r_1, \dots, r_q$  and sends  $r_i \| m_i$  to his left oracle. Given the answer from the oracle,  $\mathcal{A}$  returns  
 2501 them to  $\mathcal{B}$ . Eventually,  $\mathcal{B}$  then outputs two challenge messages  $(\tilde{m}_0, \tilde{m}_1)$  and sends them  
 2502 to  $\mathcal{A}$  who randomly selects message  $\tilde{m}_b$ , extends it with  $r$  and sends  $r \| \tilde{m}_b$  to his left oracle.  
 2503 The oracle answers with  $y_b$  which is also sent to  $\mathcal{B}$ . Finally,  $\mathcal{B}$  returns the decision bit  $\tilde{b}$   
 2504 to  $\mathcal{A}$ . If  $b = \tilde{b}$ ,  $\mathcal{A}$  answers to the challenger that the oracle was instantiating the PRF.  
 2505 Otherwise,  $\mathcal{A}$  answers with a random guess. The advantage of  $\mathcal{A}$  equals advantage of  $\mathcal{B}$   
 2506 if it interacts with a real hash function. The advantage of  $\mathcal{A}$  equals half the advantage  
 2507 of  $\mathcal{B}$  when interacting with a random oracle and simulator.

2508 *Binding.* A commitment scheme  $\text{ComSch}$  is said to be computationally binding if  
 2509 it is infeasible to find  $x, x'$  and  $r, r'$  such that  $x \neq x'$  and  $\text{Com}(x; r) = \text{Com}(x'; r')$ .  
 2510 This is implied by collision resistance of Blake2. Thus if  $\mathcal{B}$  is an algorithm that breaks  
 2511 the biding property with advantage  $\eta$ , there is another algorithm  $\mathcal{A}$  that breaks Blake2  
 2512 collision resistance with the same advantage.  $\square$

---

<sup>1</sup>The probability would be  $\frac{q^2}{2^{\text{ol}}}$  which is negligible for a polynomial number of queries  $q$ . This is the sum of the probabilities of finding a collision when doing the  $i^{\text{th}}$  query. Indeed, let us suppose the adversary has done  $i - 1$ ,  $i > 2$ , queries without finding a collision, i.e. he knows  $i - 1$  distinct tuples of input-output. When receiving the  $i^{\text{th}}$  value, the adversary has thus  $i - 1$  chance to find a collision. The probability for the new output to be equal to any of the previous outputs is thus  $(i - 1) \cdot \frac{1}{2^{\text{ol}}}$  (as we are in the random oracle model). Summing this probability over all queries, we find the probability of finding a collision after doing  $q$  queries.

2513 Assuming that Blake2s is as secure as Blake2, a commitment scheme based on a  
 2514 Blake2s, i.e.  $\text{Com}(x; r) = \text{Blake2s}_E(r \| x)$  is hiding and binding.

#### 2515 D.2.4 Proof of commitment scheme security

To prove the binding and hiding property of  $\text{ComSch}$  (see Section 3.1.2), we introduce the following commitment scheme  $\text{ComSch}^*$ ,

$$\begin{aligned} \text{ComSch}^*.\text{Setup} : \{1^\lambda \text{ s.t. } \lambda \in \mathbb{N}\} &\rightarrow \mathbb{B}^* \\ \text{ComSch}^*.\text{Com} : \mathcal{BK}^*(2 \cdot \text{BLAKE2sCLEN}, 2 \cdot \text{BLAKE2sCLEN}) &\times \mathbb{B}^{2 \cdot \text{BLAKE2sCLEN}} \\ &\times (\mathbb{B}^{\text{PRFADDRROUTLEN}} \times \mathbb{B}^{\text{PRFRHOOUTLEN}} \times \mathbb{B}^{\text{ZVALUELEN}}) \times \mathbb{B}^{\text{RTRAPLEN}} \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \end{aligned}$$

The commitment scheme is defined as follows,

$$\begin{aligned} \text{ComSch}^*.\text{Setup}(1^\lambda) &= pp^* = \epsilon \\ \text{ComSch}^*.\text{Com}(m = (apk, \rho, v); r) &= cm \\ &= \text{Blake2E}^*(r \| apk \| \rho \| v) \end{aligned}$$

Given a commitment scheme  $\text{ComSch}^*$ , the bijective function  $\text{decode}_{\mathbb{N}}(\cdot)$  and  $p_\lambda \in \mathbb{N}$ , a prime which can be represented using  $\lambda$  bits, we define the commitment scheme  $\text{ComSch}'$  as follows:

$$\begin{aligned} \text{ComSch}'.\text{Setup}(1^\lambda) &= (\text{ComSch}^*.\text{Setup}(1^\lambda), p_\lambda) \\ \text{ComSch}'.\text{Com}(m; r) &= \text{decode}_{\mathbb{N}}(\text{ComSch}^*.\text{Com}(m; r)) \pmod{p_\lambda} \text{ for } m = (apk \| \rho \| v) \end{aligned}$$

2516 Note that  $\text{ComSch}$  (see Section 3.1.2) is a particular instantiation of  $\text{ComSch}'$  where  $E^*$   
 2517 is set as ChaCha encryption scheme [Ber08a],  $k^*$  is a random key, and  $p_\lambda$  is  $r$ .

2518 **Theorem D.2.5** (Hiding). *If  $\text{ComSch}^*$  is hiding then  $\text{ComSch}'$  is hiding.*

2519 *Proof.* We prove the theorem by contradiction i.e. we assume that there exists an adver-  
 2520 sary  $\mathcal{B}$  that breaks  $\text{ComSch}'$ 's hiding property and construct an adversary  $\mathcal{A}$  that uses  
 2521  $\mathcal{B}$  to break  $\text{ComSch}^*$ 's hiding property with non-negligible probability.

2522 Let  $\mathcal{C}$  be a challenger that sets up the hiding game for  $\text{ComSch}^*$  and  $\mathcal{A}$ . The adversary  
 2523  $\mathcal{A}$ , given public parameters  $pp^*$  of  $\text{ComSch}^*$  and access to an oracle that runs the  $\text{Com}$   
 2524 algorithm of  $\text{ComSch}^*$  scheme, simulates a hiding game for  $\text{ComSch}'$  for  $\mathcal{B}$ . The adversary  
 2525  $\mathcal{A}$  starts by setting public parameters  $pp'$  for  $\text{ComSch}'$  using public parameters  $pp^*$   
 2526 given by  $\mathcal{C}$ . Parameters  $pp'$  are passed to  $\mathcal{B}$  who outputs a pair of messages  $m_0, m_1$ .  
 2527 The adversary  $\mathcal{A}$  forwards them to the challenger who samples a bit  $b$  at random and  
 2528 generates  $cm^* = \text{ComSch}^*.\text{Com}(m_b; r)$  for some randomness  $r$ . The result is returned  
 2529 to  $\mathcal{A}$  (see Definition 1.5.21). Then  $\mathcal{A}$  passes  $cm = \text{decode}_{\mathbb{N}}(cm^*) \pmod{p_\lambda}$  to  $\mathcal{B}$  who  
 2530 returns his guess  $b'$ . The adversary  $\mathcal{A}$  returns the same  $b'$  to the challenger.

2531 By construction, it is clear that  $\mathcal{A}$  wins the hiding game with the same probability  
 2532 that  $\mathcal{B}$  wins the simulated hiding game. Since  $\mathcal{B}$ 's advantage is non-negligible, this means  
 2533 that  $\mathcal{A}$  wins the  $\text{ComSch}^*$  hiding game with non-negligible probability as well.  $\square$

2534 **Theorem D.2.6** (Binding). *Let  $\text{ComSch}^*$  be a computationally binding commitment*  
 2535 *scheme and  $\text{ComSch}^*.\text{Com}$  indifferentiable from a random oracle. Then  $\text{ComSch}'$  is also*  
 2536 *computationally binding if  $l = \lceil 2^\lambda/p_\lambda \rceil$  is at most  $\text{poly}(\lambda)$ .*

2537 *Proof.* Assume that  $\mathcal{A}$  asks the  $\text{ComSch}'$  commit and open oracles a total of  $q_\lambda$  distinct  
 2538 queries. Let us denote the result of the  $q_\lambda$  queries and output of the attacker (the  
 2539 candidate collision) as  $((m_1, r_1, y_1), \dots, (m_{q_\lambda}, r_{q_\lambda}, y_{q_\lambda}), \text{out})$ . If  $\mathcal{A}$  is successful it means  
 2540 that it outputs  $(m, r)$ ,  $(m', r')$  such that  $(m, r) \neq (m', r')$  and  $\text{ComSch}'.\text{Com}(m; r) =$   
 2541  $\text{ComSch}'.\text{Com}(m'; r')$ .

By the definition of  $\text{ComSch}'$ , we have that,

$$\text{ComSch}'.\text{Com}(m; r) = \text{decode}_\mathbb{N}(\text{ComSch}^*.\text{Com}(m; r)) \pmod{p_\lambda}$$

Hence, we have a collision in  $\text{ComSch}'$  if there exists  $k \in [l]$ ,  $l$  being the ratio of the  
 codomains of  $\text{ComSch}^*.\text{Com}$  and  $\text{ComSch}'.\text{Com}$ , such that,

$$|\text{decode}_\mathbb{N}(\text{ComSch}^*.\text{Com}(m; r)) - \text{decode}_\mathbb{N}(\text{ComSch}^*.\text{Com}(m'; r'))| = k \cdot p_\lambda.$$

2542 We show that this event is unlikely.

2543 In fact, for each  $i \in [q_\lambda]$ , let  $C_i$  be the event that the adversary wins at the  $i$ -th  
 2544 query. That is, the last commitment  $y_i$  is a  $\text{ComSch}'$  collision with one of the previous  
 2545  $y_j$ . More precisely there exists  $j \leq i$  and  $k < l$  such that  $y_i = y_j + k \cdot p_\lambda$ .

2546 Since  $\text{ComSch}^*$  is a random oracle,  $y_i$  is randomly selected from a set of at least  $p_\lambda$   
 2547 elements. As such, we have  $\Pr[C_i] \leq i \cdot l/p_\lambda$ .

Thus the probability of finding a collision after  $q_\lambda$  queries is  $\Pr[C_1 \vee \dots \vee C_{q_\lambda}] \leq$   
 $\sum_{i=1}^{q_\lambda} \Pr[C_i] = l/p_\lambda \cdot \sum_{i=1}^{q_\lambda} i$ . This probability is bounded by  $l \cdot \frac{q_\lambda(q_\lambda+1)}{p_\lambda}$ . However,  
 we allow only polynomial number of queries. Thus for  $q_\lambda = \text{poly}(\lambda)$  this probability  
 becomes,

$$\frac{2^\lambda \cdot \text{poly}(\lambda)}{p_\lambda^2},$$

2548 what is negligible for  $2^\lambda/p_\lambda \leq \text{poly}(\lambda)$ . □

2549 **Remark D.2.4.** Note that in Zeth's commitment scheme, we set  $p_\lambda = \mathbf{r}$  and  $2^\lambda =$   
 2550  $2^{\text{BLAKE2sCLEN}}$ . Thus, for BN-254 and BLS12-377 have  $l = 6$  and  $l = 14$ , respectively.  
 2551 Therefore, the probability of an attacker breaking the binding property due to reduction  
 2552 modulo  $\mathbf{r}$  increases approximately by these factors. This is still negligible.

2553 **Corollary.** *Assume that Blake2 is indifferentiable from a random oracle and a PRF,*  
 2554 *then  $\text{ComSch}^*$  is computationally binding and computationally hiding. Furthermore, the*  
 2555 *reduction is tight. That is, the advantage of any PPT adversary against the binding*  
 2556 *(resp. hiding) property is the same as the advantage of an adversary against collision*  
 2557 *resistance and binding (resp. hiding).*

2558

## 2559 Glossary

- 2560 **joinsplit** Set of JSIN input *ZethNotes*, and JSOUT output *ZethNotes* as well as the  
2561 public values *vin* and *vout* used in a  $tx_{\text{Mix}}$  transaction. 37, 39, 41, 59, 96, 115, 116
- 2562 **joinsplit equation** Equation that checks that the sum of the values of the SendTx  
2563 algorithm of DAP is equal to the sum of the values of its outputs. This equations  
2564 checks that the joinsplit is “balanced” and thus, that no value is created while  
2565 creating new *ZethNotes*. 25, 41, 59, 96, 115

## 2566 **Acronyms**

2567 **DOS** Denial of Service (Attack). 16, 115

2568 **EOA** Externally Owned Account. 16, 17, 19, 115

2569 **EVM** Ethereum Virtual Machine. 15, 16, 20, 48, 54, 64, 84, 86, 115

2570 **FFT** Fast Fourier Transform. 85, 115

2571 **MAC** Message Authentication Code. 98, 115

2572 **PoC** Proof of Concept. 115

2573 **RAM** Random-access Memory. 85, 115

2574 **RLP** Recursive Length Prefix. 19, 20, 115

# Bibliography

- 2576 [AAM12] Imad Fakhri Alshaikhli, Mohammad A Alahmad, and Khanssaa Munthir.  
2577 Comparison and analysis study of sha-3 finalists. In *2012 International*  
2578 *Conference on Advanced Computer Science Applications and Technologies*  
2579 *(ACSAT)*, pages 366–371. IEEE, 2012.
- 2580 [abi] Contract abi specification, section "function selector". [https:](https://solidity.readthedocs.io/en/develop/abi-spec.html#function-selector)  
2581 [//solidity.readthedocs.io/en/develop/abi-spec.html#](https://solidity.readthedocs.io/en/develop/abi-spec.html#function-selector)  
2582 [function-selector](https://solidity.readthedocs.io/en/develop/abi-spec.html#function-selector).
- 2583 [ABM<sup>+</sup>03] Adrian Antipa, Daniel Brown, Alfred Menezes, René Struik, and Scott  
2584 Vanstone. Validation of elliptic curve public keys. In *International Work-*  
2585 *shop on Public Key Cryptography*, pages 211–223. Springer, 2003.
- 2586 [ABN10] Michel Abdalla, Mihir Bellare, and Gregory Neven. Robust encryption.  
2587 In *Theory of Cryptography Conference*, pages 480–497. Springer, 2010.  
2588 <https://eprint.iacr.org/2008/440.pdf>.
- 2589 [ABR99] Michel Abdalla, Mihir Bellare, and Phillip Rogaway. Dhaes:  
2590 An encryption scheme based on the diffie-hellman prob-  
2591 lem. 1999. [https://pdfs.semanticscholar.org/95f4/](https://pdfs.semanticscholar.org/95f4/63d097086fba325086a4cf88706648dafd09.pdf)  
2592 [63d097086fba325086a4cf88706648dafd09.pdf](https://pdfs.semanticscholar.org/95f4/63d097086fba325086a4cf88706648dafd09.pdf).
- 2593 [ABR01] Michel Abdalla, Mihir Bellare, and Phillip Rogaway. Dhies: An en-  
2594 cryptation scheme based on the diffie-hellman problem., 2001. [https:](https://web.cs.ucdavis.edu/~rogaway/papers/dhies.pdf)  
2595 [//web.cs.ucdavis.edu/~rogaway/papers/dhies.pdf](https://web.cs.ucdavis.edu/~rogaway/papers/dhies.pdf).
- 2596 [ACG<sup>+</sup>19] Martin R Albrecht, Carlos Cid, Lorenzo Grassi, Dmitry Khovratovich,  
2597 Reinhard Lüftenegger, Christian Rechberger, and Markus Schofnegger.  
2598 Algebraic cryptanalysis of stark-friendly designs: application to marvel-  
2599 lous and mimc. In *International Conference on the Theory and Appli-*  
2600 *cation of Cryptology and Information Security*, pages 371–397. Springer,  
2601 2019.
- 2602 [AFK<sup>+</sup>08] Jean-Philippe Aumasson, Simon Fischer, Shahram Khazaei, Willi Meier,  
2603 and Christian Rechberger. New features of latin dances: analysis of salsa,  
2604 chacha, and rumba. In *International Workshop on Fast Software Encryp-*  
2605 *tion*, pages 470–488. Springer, 2008.



- 2606 [AG18] Andreas M. Antonopoulos and Wood Gavin. *Mastering Ethereum*.  
2607 O'Reilly Media, 2018.
- 2608 [AGM<sup>+</sup>09] Kazumaro Aoki, Jian Guo, Krystian Matusiewicz, Yu Sasaki, and Lei  
2609 Wang. Preimages for step-reduced sha-2. In *International Conference*  
2610 *on the Theory and Application of Cryptology and Information Security*,  
2611 pages 578–597. Springer, 2009. [https://link.springer.com/content/  
2612 pdf/10.1007/978-3-642-10366-7\\_34.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-10366-7_34.pdf).
- 2613 [AGR<sup>+</sup>16] Martin Albrecht, Lorenzo Grassi, Christian Rechberger, Arnab Roy, and  
2614 Tyge Tiessen. Mimc: Efficient encryption and cryptographic hashing  
2615 with minimal multiplicative complexity. In *International Conference on*  
2616 *the Theory and Application of Cryptology and Information Security*, pages  
2617 191–219. Springer, 2016.
- 2618 [AHMP08] Jean-Philippe Aumasson, Luca Henzen, Willi Meier, and Raphael C-W  
2619 Phan. Sha-3 proposal blake. *Submission to NIST*, 229:230, 2008.
- 2620 [ALM12] Elena Andreeva, Atul Luykx, and Bart Mennink. Provable security of  
2621 blake with non-ideal compression function. In *International Conference*  
2622 *on Selected Areas in Cryptography*, pages 321–338. Springer, 2012.
- 2623 [AMP10] Elena Andreeva, Bart Mennink, and Bart Preneel. Security reductions  
2624 of the second round sha-3 candidates. In *International Conference on*  
2625 *Information Security*, pages 39–53. Springer, 2010.
- 2626 [AMPŠ12] Elena Andreeva, Bart Mennink, Bart Preneel, and Marjan Škrobot. Se-  
2627 curity analysis and comparison of the sha-3 finalists blake, grøstl, jh,  
2628 keccak, and skein. In *International Conference on Cryptology in Africa*,  
2629 pages 287–305. Springer, 2012.
- 2630 [ANWOW13] Jean-Philippe Aumasson, Samuel Neves, Zooko Wilcox-O’Hearn, and  
2631 Christian Winnerlein. Blake2: simpler, smaller, fast as md5. In *Interna-*  
2632 *tional Conference on Applied Cryptography and Network Security*, pages  
2633 119–135. Springer, 2013. <https://eprint.iacr.org/2013/322.pdf>.
- 2634 [AS09] Kazumaro Aoki and Yu Sasaki. Meet-in-the-middle preimage attacks  
2635 against reduced sha-0 and sha-1. In *Annual International Cryptology Con-*  
2636 *ference*, pages 70–89. Springer, 2009.
- 2637 [BBDP01] Mihir Bellare, Alexandra Boldyreva, Anand Desai, and David Pointcheval.  
2638 Key-privacy in public-key encryption. In *International Conference on the*  
2639 *Theory and Application of Cryptology and Information Security*, pages  
2640 566–582. Springer, 2001. [https://iacr.org/archive/asiacrypt2001/  
2641 22480568.pdf](https://iacr.org/archive/asiacrypt2001/22480568.pdf).

- [BCC<sup>+</sup>15] Jonathan Bootle, Andrea Cerulli, Pyrros Chaidos, Essam Ghadafi, Jens Groth, and Christophe Petit. Short accountable ring signatures based on ddh. In *European Symposium on Research in Computer Security*, pages 243–265. Springer, 2015.
- [BCD<sup>+</sup>20] Tim Beyne, Anne Canteaut, Itai Dinur, Maria Eichlseder, Gregor Leander, Gaëtan Leurent, María Naya-Plasencia, Léo Perrin, Yu Sasaki, Yosuke Todo, and Friedrich Wiemer. Out of oddity – new cryptanalytic techniques against symmetric primitives optimized for integrity proof systems. Cryptology ePrint Archive, Report 2020/188, 2020. <https://eprint.iacr.org/2020/188>.
- [BCG<sup>+</sup>20] Sean Bowe, Alessandro Chiesa, Matthew Green, Ian Miers, Pratyush Mishra, and Howard Wu. ZEXE: enabling decentralized private computation. In *2020 IEEE Symposium on Security and Privacy, SP 2020, San Francisco, CA, USA, May 18-21, 2020*, pages 947–964. IEEE, 2020.
- [BCK<sup>+</sup>18] Elaine Barker, Lily Chen, Sharon Keller, Allen Roginsky, Apostol Vassilev, and Richard Davis. Recommendation for pair-wise key-establishment schemes using discrete logarithm cryptography. Technical report, National Institute of Standards and Technology, 2018. [Online; last accessed 10-January-2020].
- [BD07] Eli Biham and Orr Dunkelman. A framework for iterative hash functions—haifa. Technical report, Computer Science Department, Technion, 2007. <https://eprint.iacr.org/2007/278.pdf>.
- [BDPVA07] Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche. Sponge functions. In *ECRYPT hash workshop*, volume 2007. Citeseer, 2007.
- [Ber05] Daniel J Bernstein. The poly1305-aes message-authentication code. In *International Workshop on Fast Software Encryption*, pages 32–49. Springer, 2005. <https://cr.yp.to/mac/poly1305-20050329.pdf>.
- [Ber06] Daniel J Bernstein. Curve25519: new diffie-hellman speed records. In *International Workshop on Public Key Cryptography*, pages 207–228. Springer, 2006. <https://cr.yp.to/ecdh/curve25519-20060209.pdf>.
- [Ber08a] Daniel J Bernstein. Chacha, a variant of salsa20. In *Workshop Record of SASC*, volume 8, pages 3–5, 2008. <https://cr.yp.to/chacha/chacha-20080120.pdf>.
- [Ber08b] Daniel J. Bernstein. New stream cipher designs. chapter The Salsa20 Family of Stream Ciphers, pages 84–97. Springer-Verlag, Berlin, Heidelberg, 2008.

- 2679 [BGM17] Sean Bowe, Ariel Gabizon, and Ian Miers. Scalable multi-party compu-  
2680 tation for zk-snark parameters in the random beacon model. *Cryptology*  
2681 *ePrint Archive*, Report 2017/1050, 2017. [https://eprint.iacr.org/](https://eprint.iacr.org/2017/1050)  
2682 [2017/1050](https://eprint.iacr.org/2017/1050).
- 2683 [BL] Daniel J. Bernstein and Tanja Lange. Safecurves: choosing safe curves  
2684 for elliptic-curve cryptography. <https://safecurves.cr.yp.to>. [Online;  
2685 last accessed 09-December-2019].
- 2686 [Blo70] Burton H. Bloom. Space/time trade-offs in hash coding with allowable  
2687 errors. *Commun. ACM*, 13(7):422–426, 1970.
- 2688 [Bon19] Xavier Bonnetain. Collisions on feistel-mimc and univariate gmimc. 2019.
- 2689 [Bou03] Nicolas Bourbaki. *Elements of mathematics: Algebra*. Springer, 2003.
- 2690 [Bra97] S. Bradner. Key words for use in rfcs to indicate requirement levels. RFC  
2691 2119, RFC Editor, March 1997.
- 2692 [BRS02] John Black, Phillip Rogaway, and Thomas Shrimpton. Black-box  
2693 analysis of the block-cipher-based hash-function constructions from  
2694 pgv. In *Annual International Cryptology Conference*, pages 320–335.  
2695 Springer, 2002. [https://link.springer.com/content/pdf/10.1007/](https://link.springer.com/content/pdf/10.1007/3-540-45708-9_21.pdf)  
2696 [3-540-45708-9\\_21.pdf](https://link.springer.com/content/pdf/10.1007/3-540-45708-9_21.pdf).
- 2697 [BS07] Mihir Bellare and Sarah Shoup. Two-tier signatures, strongly unforge-  
2698 able signatures, and fiat-shamir without random oracles. In *International*  
2699 *Workshop on Public Key Cryptography*, pages 201–216. Springer, 2007.
- 2700 [BSCG<sup>+</sup>14] Eli Ben Sasson, Alessandro Chiesa, Christina Garman, Matthew Green,  
2701 Ian Miers, Eran Tromer, and Madars Virza. Zerocash: Decentralized  
2702 anonymous payments from bitcoin. In *2014 IEEE Symposium on Security*  
2703 *and Privacy*, pages 459–474. IEEE, 2014.
- 2704 [CHM<sup>+</sup>20] Alessandro Chiesa, Yuncong Hu, Mary Maller, Pratyush Mishra, Noah  
2705 Vesely, and Nicholas P. Ward. Marlin: Preprocessing zkSNARKs with  
2706 universal and updatable SRS. pages 738–768, 2020.
- 2707 [Cle19] Clearmatics. Zeth Sproken release. [https://github.com/clearmatics/](https://github.com/clearmatics/zeth/releases/tag/v0.2)  
2708 [zeth/releases/tag/v0.2](https://github.com/clearmatics/zeth/releases/tag/v0.2), 2019. [Online; released 04-April-2019].
- 2709 [CM16] Arka Rai Choudhuri and Subhamoy Maitra. Differential cryptanalysis of  
2710 salsa and chacha-an evaluation with a hybrid model. *IACR Cryptology*  
2711 *ePrint Archive*, 2016:377, 2016. [https://eprint.iacr.org/2016/377.](https://eprint.iacr.org/2016/377.pdf)  
2712 [pdf](https://eprint.iacr.org/2016/377.pdf).
- 2713 [CM17] Arka Rai Choudhuri and Subhamoy Maitra. Significantly improved multi-  
2714 bit differentials for reduced round salsa and chacha. *IACR Transactions*  
2715 *on Symmetric Cryptology*, 2016(2):261–287, Feb. 2017.

- 2716 [DG09] George Danezis and Ian Goldberg. Sphinx: A compact and provably  
2717 secure mix format. In *30th IEEE Symposium on Security and Privacy*  
2718 *(S&P 2009), 17-20 May 2009, Oakland, California, USA*, pages 269–282.  
2719 IEEE Computer Society, 2009.
- 2720 [EFK15] Thomas Espitau, Pierre-Alain Fouque, and Pierre Karpman. Higher-order  
2721 differential meet-in-the-middle preimage attacks on sha-1 and blake. In  
2722 *Annual Cryptology Conference*, pages 683–701. Springer, 2015. <https://eprint.iacr.org/2015/515>.  
2723
- 2724 [EGL<sup>+</sup>20] Maria Eichlseder, Lorenzo Grassi, Reinhard Lüftenegger, Morten Øygarden,  
2725 Christian Rechberger, Markus Schofnegger, and Qingju Wang. An  
2726 algebraic attack on ciphers with low-degree round functions: Application  
2727 to full mimc. Cryptology ePrint Archive, Report 2020/182, 2020.  
2728 <https://eprint.iacr.org/2020/182>.
- 2729 [est] Estream project. <https://en.wikipedia.org/wiki/ESTREAM>.
- 2730 [Gab19] Ariel Gabizon. AuroraLight: Improved prover efficiency and SRS size in  
2731 a sonic-like system. Cryptology ePrint Archive, Report 2019/601, 2019.  
2732 <https://eprint.iacr.org/2019/601>.
- 2733 [GFBR06] Decio Gazzoni Filho, Paulo SLM Barreto, and Vincent Rijmen. The  
2734 maelstrom-0 hash function. In *Brazilian Symposium on Information and*  
2735 *Computer System Security*. , 2006.
- 2736 [GGPR13] Rosario Gennaro, Craig Gentry, Bryan Parno, and Mariana Raykova.  
2737 Quadratic span programs and succinct nizks without peps. In Thomas  
2738 Johansson and Phong Q. Nguyen, editors, *Advances in Cryptology - EU-*  
2739 *ROCRYPT 2013, 32nd Annual International Conference on the Theory*  
2740 *and Applications of Cryptographic Techniques, Athens, Greece, May 26-*  
2741 *30, 2013. Proceedings*, volume 7881 of *Lecture Notes in Computer Science*,  
2742 pages 626–645. Springer, 2013.
- 2743 [GJMG11] B Guido, D Joan, P Michaël, and VA Gilles. The keccak sha-3 submission.  
2744 2011.
- 2745 [GKN<sup>+</sup>14] Jian Guo, Pierre Karpman, Ivica Nikolić, Lei Wang, and Shuang Wu.  
2746 Analysis of blake2. In *Cryptographers’ Track at the RSA Conference*,  
2747 pages 402–423. Springer, 2014. [https://eprint.iacr.org/2013/467](https://eprint.iacr.org/2013/467.pdf).  
2748 [pdf](https://eprint.iacr.org/2013/467.pdf).
- 2749 [GLRW10] Jian Guo, San Ling, Christian Rechberger, and Huaxiong Wang. Ad-  
2750 vanced meet-in-the-middle preimage attacks: First results on full tiger,  
2751 and improved results on md4 and sha-2. In *International Conference on*  
2752 *the Theory and Application of Cryptology and Information Security*, pages  
2753 56–75. Springer, 2010. <https://eprint.iacr.org/2010/016.pdf>.

- 2754 [GM17] Jens Groth and Mary Maller. Snarky signatures: Minimal signatures  
2755 of knowledge from simulation-extractable snarks. In Jonathan Katz and  
2756 Hovav Shacham, editors, *Advances in Cryptology - CRYPTO 2017 - 37th  
2757 Annual International Cryptology Conference, Santa Barbara, CA, USA,  
2758 August 20-24, 2017, Proceedings, Part II*, volume 10402 of *Lecture Notes  
2759 in Computer Science*, pages 581–612. Springer, 2017.
- 2760 [Gol01] Oded Goldreich. *Foundations of Cryptography: Basic Tools*, volume 1.  
2761 Cambridge University Press, Cambridge, UK, 2001.
- 2762 [Gro06] Jens Groth. Simulation-sound NIZK proofs for a practical language and  
2763 constant size group signatures. pages 444–459, 2006.
- 2764 [Gro16] Jens Groth. On the size of pairing-based non-interactive arguments. In  
2765 *Annual International Conference on the Theory and Applications of Cryptographic  
2766 Techniques*, pages 305–326. Springer, 2016.
- 2767 [GRR<sup>+</sup>16] Lorenzo Grassi, Christian Rechberger, Dragos Rotaru, Peter Scholl, and  
2768 Nigel P Smart. Mpc-friendly symmetric key primitives. In *Proceedings of  
2769 the 2016 ACM SIGSAC Conference on Computer and Communications  
2770 Security*, pages 430–443. ACM, 2016. [https://eprint.iacr.org/2016/  
2771 542](https://eprint.iacr.org/2016/542).
- 2772 [GWC19] Ariel Gabizon, Zachary J. Williamson, and Oana Ciobotaru. PLONK:  
2773 Permutations over lagrange-bases for oecumenical noninteractive argu-  
2774 ments of knowledge. Cryptology ePrint Archive, Report 2019/953, 2019.  
2775 <https://eprint.iacr.org/2019/953>.
- 2776 [Hao14] Yonglin Hao. The boomerang attacks on blake and blake2. In *International  
2777 Conference on Information Security and Cryptology*, pages 286–310.  
2778 Springer, 2014. <https://eprint.iacr.org/2014/1012.pdf>.
- 2779 [Har19] HarryR. Conversation about Miyaguchi-Preneel security. [https://  
2780 github.com/HarryR/ethsnarks/issues/119](https://github.com/HarryR/ethsnarks/issues/119), 2019. Online; accessed  
2781 June-2019.
- 2782 [HG20] Youssef El Housni and Aurore Guillevic. Optimized and secure pairing-  
2783 friendly elliptic curves suitable for one layer proof composition. Cryptol-  
2784 ogy ePrint Archive, Report 2020/351, 2020. [https://eprint.iacr.org/  
2785 2020/351](https://eprint.iacr.org/2020/351).
- 2786 [HKT11] Thomas Holenstein, Robin Künzler, and Stefano Tessaro. The equiva-  
2787 lence of the random oracle model and the ideal cipher model, revisited.  
2788 In *Proceedings of the forty-third annual ACM symposium on Theory of  
2789 computing*, pages 89–98. ACM, 2011.

- 2790 [HMRS12] Ekawat Homsirikamol, Paweł Morawiecki, Marcin Rogawski, and Marian  
2791 Srebrny. Security margin evaluation of sha-3 contest finalists through  
2792 sat-based attacks. In *IFIP International Conference on Computer Infor-*  
2793 *mation Systems and Industrial Management*, pages 56–67. Springer, 2012.
- 2794 [Hop16] Daira Hopwood. Daira’s comment on: ”ensure spec retains distinctness  
2795 assumption in hsig”, 2016.
- 2796 [IS09] Takanori Isobe and Kyoji Shibutani. Preimage attacks on reduced tiger  
2797 and sha-2. In *International Workshop on Fast Software Encryption*, pages  
2798 139–155. Springer, 2009. [https://link.springer.com/content/pdf/](https://link.springer.com/content/pdf/10.1007/978-3-642-03317-9_9.pdf)  
2799 [10.1007/978-3-642-03317-9\\_9.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-03317-9_9.pdf).
- 2800 [Ish12] Tsukasa Ishiguro. Modified version of” latin dances revisited: New ana-  
2801 lytic results of salsa20 and chacha”. 2012. [https://eprint.iacr.org/](https://eprint.iacr.org/2012/065.pdf)  
2802 [2012/065.pdf](https://eprint.iacr.org/2012/065.pdf).
- 2803 [JMV01] Don Johnson, Alfred Menezes, and Scott Vanstone. The elliptic curve  
2804 digital signature algorithm (ecdsa). *International journal of information*  
2805 *security*, 1(1):36–63, 2001.
- 2806 [Joh16] Nick Johnson. Response to ”how does ethereum make use of  
2807 bloom filters?”. [https://ethereum.stackexchange.com/questions/](https://ethereum.stackexchange.com/questions/3418/how-does-ethereum-make-use-of-bloom-filters)  
2808 [3418/how-does-ethereum-make-use-of-bloom-filters](https://ethereum.stackexchange.com/questions/3418/how-does-ethereum-make-use-of-bloom-filters), 2016. [On-  
2809 line; last accessed 10-January-2020].
- 2810 [KL14] Jonathan Katz and Yehuda Lindell. *Introduction to modern cryptogra-*  
2811 *phy*. Chapman and Hall/CRC, 2014. [https://repo.zenk-security.](https://repo.zenk-security.com/Cryptographie%20.%20Algorithmes%20.%20Steganographie/Introduction%20to%20Modern%20Cryptography.pdf)  
2812 [com/Cryptographie%20.%20Algorithmes%20.%20Steganographie/](https://repo.zenk-security.com/Cryptographie%20.%20Algorithmes%20.%20Steganographie/Introduction%20to%20Modern%20Cryptography.pdf)  
2813 [Introduction%20to%20Modern%20Cryptography.pdf](https://repo.zenk-security.com/Cryptographie%20.%20Algorithmes%20.%20Steganographie/Introduction%20to%20Modern%20Cryptography.pdf).
- 2814 [KMO<sup>+</sup>13] Markulf Kohlweiss, Ueli Maurer, Cristina Onete, Björn Tackmann, and  
2815 Daniele Venturi. Anonymity-preserving public-key encryption: A con-  
2816 structive approach. In Emiliano De Cristofaro and Matthew K. Wright,  
2817 editors, *Privacy Enhancing Technologies - 13th International Symposium,*  
2818 *PETS 2013, Bloomington, IN, USA, July 10-12, 2013. Proceedings*, vol-  
2819 *ume 7981 of Lecture Notes in Computer Science*, pages 19–39. Springer,  
2820 2013.
- 2821 [KRS12] Dmitry Khovratovich, Christian Rechberger, and Alexandra Savelieva.  
2822 Bicliques for preimages: attacks on skein-512 and the sha-2 family. In  
2823 *International Workshop on Fast Software Encryption*, pages 244–263.  
2824 Springer, 2012. [https://link.springer.com/content/pdf/10.1007/](https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5_15.pdf)  
2825 [978-3-642-34047-5\\_15.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5_15.pdf).
- 2826 [Lab19] Matter Labs. Merkle shrubs. [https://github.com/matter-labs/](https://github.com/matter-labs/MerkleShrubs)  
2827 [MerkleShrubs](https://github.com/matter-labs/MerkleShrubs), 2019.

- 2828 [LHT16] Adam Langley, Mike Hamburg, and Sean Turner. Elliptic curves for security. RFC 7748, <https://tools.ietf.org/pdf/rfc7748.pdf>, 2016.  
2829
- 2830 [LIS12] Ji Li, Takanori Isobe, and Kyoji Shibutani. Converting meet-in-the-middle preimage attack into pseudo collision attack: Application to sha-2. In *International Workshop on Fast Software Encryption*, pages 264–286. Springer, 2012. [https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5\\_16.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5_16.pdf).  
2831  
2832  
2833  
2834
- 2835 [LM11] Mario Lamberger and Florian Mendel. Higher-order differential attack on reduced sha-256. *IACR Cryptology ePrint Archive*, 2011:37, 2011. <https://eprint.iacr.org/2011/037.pdf>.  
2836  
2837
- 2838 [LMN16] Atul Luykx, Bart Mennink, and Samuel Neves. Security analysis of blake2’s modes of operation. *IACR Transactions on Symmetric Cryptology*, pages 158–176, 2016. <https://www.esat.kuleuven.be/cosic/publications/article-2705.pdf>.  
2839  
2840  
2841
- 2842 [LN18] Adam Langley and Yoav Nir. Chacha20 and poly1305 for ietf protocols. *RFC 8439*, 2018. <https://tools.ietf.org/html/rfc8439>.  
2843
- 2844 [LP19] Chaoyun Li and Bart Preneel. Improved interpolation attacks on cryptographic primitives of low algebraic degree. In *International Conference on Selected Areas in Cryptography*, pages 171–193. Springer, 2019.  
2845  
2846
- 2847 [Mai16] Subhamoy Maitra. Chosen iv cryptanalysis on reduced round chacha and salsa. *Discrete Applied Mathematics*, 208:88–97, 2016.  
2848
- 2849 [MBKM19] Mary Maller, Sean Bowe, Markulf Kohlweiss, and Sarah Meiklejohn. Sonic: Zero-knowledge SNARKs from linear-size universal and updatable structured reference strings. pages 2111–2128, 2019.  
2850  
2851
- 2852 [MJS15] Ed. M-J. Saarinen. Blake Compression Function F. <https://tools.ietf.org/html/rfc7693#section-3.2>, 2015. [Online; accessed November-2019].  
2853  
2854
- 2855 [ML15] Nicky Mouha and Atul Luykx. Multi-key security: The even-mansour construction revisited. In *Annual Cryptology Conference*, pages 209–223. Springer, 2015. <https://hal.inria.fr/hal-01240988/document>.  
2856  
2857
- 2858 [MNS11] Florian Mendel, Tomislav Nad, and Martin Schl  ffer. Finding sha-2 characteristics: searching through a minefield of contradictions. In *International Conference on the Theory and Application of Cryptology and Information Security*, pages 288–307. Springer, 2011. [https://link.springer.com/content/pdf/10.1007/978-3-642-25385-0\\_16.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-25385-0_16.pdf).  
2859  
2860  
2861  
2862

- 2863 [Moh10] Payman Mohassel. A closer look at anonymity and robustness in encryption schemes. In *Advances in Cryptology - ASIACRYPT 2010 - 16th International Conference on the Theory and Application of Cryptology and Information Security*, volume 6477 of *Lecture Notes in Computer Science*, pages 501–518. Springer, 2010.
- 2864
- 2865
- 2866
- 2867
- 2868 [Mon87] Peter L. Montgomery. Speeding the Pollard and elliptic curve methods of factorization. *Math. Comp.*, 48(177):243–264, 1987.
- 2869
- 2870 [MP15] Bart Mennink and Bart Preneel. On the impact of known-key attacks on hash functions. In *International Conference on the Theory and Application of Cryptology and Information Security*, pages 59–84. Springer, 2015.
- 2871
- 2872
- 2873 <https://eprint.iacr.org/2015/909.pdf>.
- 2874 [MQZ10] Mao Ming, He Qiang, and Shaokun Zeng. Security analysis of blake-32 based on differential properties. In *2010 International Conference on Computational and Information Sciences*, pages 783–786. IEEE, 2010.
- 2875
- 2876
- 2877 [MVOV96] Alfred J Menezes, Paul C Van Oorschot, and Scott A Vanstone. Handbook of applied cryptography. <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.99.2838&rep=rep1&type=pdf>, 1996.
- 2878
- 2879
- 2880 [NA19] Samuel Neves and Filipe Araujo. An observation on norx, blake2, and chacha. *Information Processing Letters*, 149:1–5, 2019.
- 2881
- 2882 [oST15] National Institute of Standards and Technology. Secure Hash Standard (SHS). <https://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.180-4.pdf>, 2015.
- 2883
- 2884
- 2885 [Per17] Trevor Perrin. X25519 and zero outputs. <https://moderncrypto.org/mail-archive/curves/2017/000896.html>, 2017. [Online; last accessed 08-January-2020].
- 2886
- 2887
- 2888 [PHE<sup>+</sup>17] Ania M. Piotrowska, Jamie Hayes, Tariq Elahi, Sebastian Meiser, and George Danezis. The loopix anonymity system. In Engin Kirda and Thomas Ristenpart, editors, *26th USENIX Security Symposium, USENIX Security 2017, Vancouver, BC, Canada, August 16-18, 2017*, pages 1199–1216. USENIX Association, 2017.
- 2889
- 2890
- 2891
- 2892
- 2893 [Por13] Thomas Pornin. Deterministic Usage of the Digital Signature Algorithm (DSA) and Elliptic Curve Digital Signature Algorithm (ECDSA). <https://tools.ietf.org/html/rfc6979>, 2013.
- 2894
- 2895
- 2896 [Pro14] Gordon Procter. A security analysis of the composition of chacha20 and poly1305. *IACR Cryptology ePrint Archive*, 2014:613, 2014.
- 2897



2898 [PV05] Pascal Paillier and Damien Vergnaud. Discrete-log-based signatures may  
2899 not be equivalent to discrete log. In *International Conference on the*  
2900 *Theory and Application of Cryptology and Information Security*, pages  
2901 1–20. Springer, 2005.

2902 [Qu99] Minghua Qu. Sec 2: Recommended elliptic curve domain parameters. *Cer-*  
2903 *ticom Res., Mississauga, ON, Canada, Tech. Rep. SEC2-Ver-0.6*, 1999.

2904 [Rk19] Antoine Rondelet and @karalabe. Go-ethereum BN256 package.  
2905 [https://github.com/ethereum/go-ethereum/blob/master/crypto/](https://github.com/ethereum/go-ethereum/blob/master/crypto/bn256/cloudflare/constants.go)  
2906 [bn256/cloudflare/constants.go](https://github.com/ethereum/go-ethereum/blob/master/crypto/bn256/cloudflare/constants.go), 2019. [Online; released 28-May-2019].

2907 [Ron20] Antoine Rondelet. Zecale: Reconciling privacy and scalability on  
2908 ethereum, 2020.

2909 [RZ19] Antoine Rondelet and Michal Zajac. ZETH: On Integrating Zerocash on  
2910 Ethereum. [Online; released April-2019], 2019.

2911 [SS08] Somitra Kumar Sanadhya and Palash Sarkar. New collision attacks  
2912 against up to 24-step sha-2. In *International conference on cryptology in*  
2913 *India*, pages 91–103. Springer, 2008. [https://eprint.iacr.org/2008/](https://eprint.iacr.org/2008/270.pdf)  
2914 [270.pdf](https://eprint.iacr.org/2008/270.pdf).

2915 [Ste15] Stevens, Marc. On collisions for md5. [https://www.win.tue.](https://www.win.tue.nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.%20Stevens.pdf)  
2916 [nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.](https://www.win.tue.nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.%20Stevens.pdf)  
2917 [%20Stevens.pdf](https://www.win.tue.nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.%20Stevens.pdf), 2015.

2918 [SZFW12] Zhenqing Shi, Bin Zhang, Dengguo Feng, and Wenling Wu. Improved  
2919 key recovery attacks on reduced-round salsa20 and chacha. In *Interna-*  
2920 *tional Conference on Information Security and Cryptology*, pages 337–351.  
2921 Springer, 2012.

2922 [TBP20] Florian Tramèr, Dan Boneh, and Kenneth G. Paterson. Remote side-  
2923 channel attacks on anonymous transactions. Cryptology ePrint Archive,  
2924 Report 2020/220, 2020. <https://eprint.iacr.org/2020/220>.

2925 [THH15] Piotr Dyrąga Tjaden Hess, Matt Luongo and James Hancock. EIP 152:  
2926 Add BLAKE2 compression function ‘F’. [https://eips.ethereum.org/](https://eips.ethereum.org/EIPS/eip-152)  
2927 [EIPS/eip-152](https://eips.ethereum.org/EIPS/eip-152), 2015. [Online; accessed November-2019].

2928 [VNP10] Janoš Vidali, Peter Nose, and Enes Pašalić. Collisions for variants of the  
2929 blake hash function. *Information processing letters*, 110(14-15):585–590,  
2930 2010.

2931 [W<sup>+</sup>] Gavin Wood et al. Ethereum: A secure decentralised generalised trans-  
2932 action ledger.

2933 [wc] Ethereum wiki contributors. Patricia tree. [https://github.com/](https://github.com/ethereum/wiki/wiki/Patricia-Tree)  
2934 [ethereum/wiki/wiki/Patricia-Tree](https://github.com/ethereum/wiki/wiki/Patricia-Tree).

2935 [wc19] Ethereum wiki contributors. RLP. [https://github.com/ethereum/](https://github.com/ethereum/wiki/wiki/RLP)  
2936 [wiki/wiki/RLP](https://github.com/ethereum/wiki/wiki/RLP), 2019. [Online; accessed December-2019].

2937 [wik] Bitcoin wiki. Secp256k1. <https://en.bitcoin.it/wiki/Secp256k1>.  
2938 [Online; last accessed 04-January-2020].

2939 [Woo19] Dr Gavin Wood. ETHEREUM: A Secure Decentralised Gener-  
2940 alised Transaction Ledger Byzantium. [https://ethereum.github.io/](https://ethereum.github.io/yellowpaper/paper.pdf)  
2941 [yellowpaper/paper.pdf](https://ethereum.github.io/yellowpaper/paper.pdf), 2019. [VERSION 7e819ec - 2019-10-20].

2942 [zcaa] On the security of sprout/sapling in-band en-  
2943 cryption. [https://forum.zcashcommunity.com/t/](https://forum.zcashcommunity.com/t/on-the-security-of-sprout-sapling-in-band-encryption/34986)  
2944 [on-the-security-of-sprout-sapling-in-band-encryption/34986](https://forum.zcashcommunity.com/t/on-the-security-of-sprout-sapling-in-band-encryption/34986).

2945 [zcab] "zcash alerts - security announcement 2019-09-24".  
2946 [https://z.cash/support/security/announcements/](https://z.cash/support/security/announcements/security-announcement-2019-09-24/a)  
2947 [security-announcement-2019-09-24/a](https://z.cash/support/security/announcements/security-announcement-2019-09-24/a).

2948 [ZCa19] ZCash. ZCash protocol specification. [https://github.com/zcash/zips/](https://github.com/zcash/zips/blob/master/protocol/protocol.pdf)  
2949 [blob/master/protocol/protocol.pdf](https://github.com/zcash/zips/blob/master/protocol/protocol.pdf), 2019. [Online; initially released  
2950 14-December-2015].