

1

# Zeth Protocol Specification

2

Clearmatics Cryptography R&D

3

2020-06-11

## Abstract

5 This document specifies the Zeth protocol with various security fixes and performance  
6 improvements from the initial design [RZ19].

7 **Keywords**— Ethereum, Zerocash, Zcash, financial-privacy, zero-knowledge proofs,  
8 Zeth, privacy-preserving state transitions

# 9 Contents

10	<b>1 Preliminaries</b>	<b>10</b>
11	1.1 Data Structures and Representation . . . . .	10
12	1.1.1 Structured Data . . . . .	10
13	1.1.2 Representations . . . . .	13
14	1.2 Ethereum . . . . .	14
15	1.2.1 Ethereum Account . . . . .	15
16	1.2.2 Ethereum Transaction . . . . .	16
17	1.2.3 Ethereum events and Bloom filters . . . . .	19
18	1.3 zk-SNARKs . . . . .	20
19	1.3.1 Preliminary definitions . . . . .	20
20	1.3.2 Computation representation — Arithmetization . . . . .	22
21	1.4 Decentralized Anonymous Payment schemes (DAP) . . . . .	23
22	1.5 Security Assumptions . . . . .	25
23	1.6 Definitions . . . . .	26
24	1.6.1 Negligible function . . . . .	26
25	1.6.2 Basic algebra notions . . . . .	26
26	1.6.3 Symmetric Encryption . . . . .	27
27	1.6.4 Asymmetric Encryption . . . . .	28
28	1.6.5 (Block-cipher-based) Compression functions . . . . .	29
29	1.6.6 Hash functions . . . . .	30
30	1.6.7 Pseudo Random Functions . . . . .	32
31	1.6.8 Commitment scheme . . . . .	32
32	1.6.9 Digital Signature . . . . .	33
33	1.6.10 Message Authentication Code . . . . .	34
34	<b>2 Zeth protocol</b>	<b>36</b>
35	2.1 Zeth Data Types . . . . .	36
36	2.2 Zeth statement . . . . .	39
37	2.3 Generating the inputs of the Mix function ( $\text{Mix}_{in}$ ) . . . . .	40
38	2.4 Creating an Ethereum transaction $tx_{\text{Mix}}$ to call <b>Mixer</b> . . . . .	42
39	2.5 Processing $tx_{\text{Mix}}$ . . . . .	42
40	2.6 Receiving <i>ZethNotes</i> . . . . .	44
41	2.7 Security requirements for the primitives . . . . .	45

42	2.7.1	Additional notes . . . . .	45
43	<b>3</b>	<b>Instantiation of the cryptographic primitives</b>	<b>47</b>
44	3.1	Instantiating the PRFs, ComSch and CRHs . . . . .	47
45	3.1.1	Blake2 primitive . . . . .	48
46	3.1.2	Commitment scheme . . . . .	48
47	3.1.3	PRFs . . . . .	49
48	3.1.4	Collision Resistant Hashes . . . . .	50
49	3.2	Instantiating MKHASH . . . . .	51
50	3.2.1	MIMC Encryption . . . . .	51
51	3.2.2	MIMC-based compression function . . . . .	52
52	3.2.3	An efficient instantiation of MIMC primitives . . . . .	53
53	3.2.4	MKHASH instantiation . . . . .	54
54	3.2.5	Security requirements satisfaction . . . . .	54
55	3.3	Zeth statement after primitive instantiation . . . . .	56
56	3.3.1	Instantiating the Packing functions . . . . .	58
57	3.4	Instantiate SigSch <sub>OT-SIG</sub> . . . . .	60
58	3.4.1	Security requirements satisfaction . . . . .	60
59	3.4.2	Data types . . . . .	61
60	3.5	Instantiate EncSch . . . . .	62
61	3.5.1	DHAES encryption scheme . . . . .	62
62	3.5.2	A DHAES instance . . . . .	63
63	3.5.3	EncSch instantiation . . . . .	66
64	3.5.4	Security requirements satisfaction . . . . .	68
65	3.5.5	Final notes and observations . . . . .	71
66	3.6	Instantiate ZkSnarkSch . . . . .	72
67	<b>4</b>	<b>Implementation considerations and optimizations</b>	<b>76</b>
68	4.1	Client Security Considerations . . . . .	76
69	4.1.1	Syncing and Waiting . . . . .	77
70	4.1.2	Note Management . . . . .	77
71	4.1.3	Prepare Arguments for Mix Transaction . . . . .	78
72	4.1.4	Wallet Backup and Recovery . . . . .	79
73	4.2	Contract Security Considerations . . . . .	80
74	4.3	Efficiency and Scalability . . . . .	80
75	4.3.1	Importance of Performance . . . . .	80
76	4.3.2	Cost Centers . . . . .	81
77	4.3.3	Client Performance . . . . .	81
78	4.3.4	Contract Performance . . . . .	82
79	4.3.5	Optimizing Blake2's circuit. . . . .	83
80	4.4	Encryption of the notes . . . . .	90

81	<b>A Transaction Non Malleability</b>	<b>91</b>
82	A.1 Transaction malleability attack on Zeth . . . . .	92
83	A.1.1 The transaction malleability attack . . . . .	93
84	A.2 Solutions to address the transaction malleability attack . . . . .	94
85	A.2.1 ZeroCash solution . . . . .	94
86	A.2.2 Zcash’s solution . . . . .	95
87	A.2.3 Solution on Ethereum . . . . .	95
88	<b>B Double Spend Attack on Equivalent class</b>	<b>97</b>
89	<b>C Side-Channel Attacks and Information Leaks</b>	<b>98</b>
90	C.1 Counterfeit Data . . . . .	98
91	C.2 Data Leaked during Synchronization . . . . .	99
92	C.3 Queries on Successful Decryption . . . . .	100
93	C.4 Invalid Ciphertext . . . . .	100
94	C.5 Using (and Retrieving) Nullifiers . . . . .	101
95	C.6 Proof Generation . . . . .	101
96	C.7 Simple Mixer Calls . . . . .	102
97	C.7.1 Small Anonymity Sets . . . . .	103
98	<b>D Security proofs of Blake2</b>	<b>104</b>
99	D.1 Security model of Blake2 . . . . .	104
100	D.1.1 Weakly Ideal Cipher Model . . . . .	104
101	D.2 Security proofs . . . . .	106
102	D.2.1 Proof of Blake2 function’s PRFness . . . . .	106
103	D.2.2 Proof of Blake2 collision resistance . . . . .	107
104	D.2.3 Blake2 as a commitment scheme . . . . .	108
105	D.2.4 Proof of commitment scheme security . . . . .	109
106		

# 107 Notations

## 108 Basic mathematical notations

- 109  $\emptyset$  The empty set, i.e.  $\emptyset = \{\}$
- 110  $\#S$  The number of elements in the finite set  $S$  (also referred to as “cardinality of the  
111 set  $S$ ”). By convention,  $\#\emptyset = 0$
- 112  $x \in S$  Represents that  $x$  is an element of  $S$ . If  $x$  is a variable such that  $x \in S$ , we will  
113 say that “ $x$  has type  $S$ ”, i.e. the unordered collection of objects  $S$  represents all  
114 the values that  $x$  can take
- 115  $S \setminus T$  Set difference of sets  $S$  and  $T$ , i.e.  $S \setminus T = \{x \in S : x \notin T\}$  (voiced “the set of  
116 elements  $x$  in  $S$  such that  $x$  is not in  $T$ ”)
- 117  $S \subseteq T$   $S$  is a subset of  $T$ , i.e.  $\forall x \in S \Rightarrow x \in T$
- 118  $S \subset T$   $S$  is a *proper* (or “strict”) subset of  $T$ , i.e.  $\forall x \in S \Rightarrow x \in T \wedge \exists y \in T, y \notin S$
- 119  $S = T$   $S \subseteq T \wedge T \subseteq S$
- 120  $S \cup T$  Union of set  $S$  and set  $T$ , i.e.  $\{x : x \in S \vee x \in T\}$
- 121  $S \cap T$  Intersection of set  $S$  with set  $T$ , i.e.  $\{x : x \in S \wedge x \in T\}$
- 122  $f: S \rightarrow T$  Function  $f$  that maps items from the non-empty set  $S$ , called “the domain”,  
123 to the non-empty set  $T$ , called “the co-domain”
- 124  $\mathbb{N}$  Set of natural numbers.  $\mathbb{N}^+$  represents  $\mathbb{N} \setminus \{0\} = \{1, 2, \dots\}$ , where  $\{n, \dots\}$  rep-  
125 represents the application of the successor operator  $\text{Succ}(n) = n + 1$ , defined by the  
126 Peano axioms, infinitely many times
- 127  $\mathbb{Z}$  Set of integers, i.e.  $\{\dots, -2, -1, 0, 1, 2, \dots\}$ , where  $\{\dots, n\}$  represents the appli-  
128 cation of the predecessor operator  $\text{Pred}(n) = n - 1$  infinitely many times
- 129  $\mathbb{Q}, \mathbb{R}$  Set of rational, real numbers
- 130  $[n]$  Set  $\{0, \dots, n - 1\}$ , where  $n \in \mathbb{N}$
- 131  $\{a, \dots, b\}$  Set of integers from  $a$  through  $b$  inclusive, where  $a \leq b$

132	$(a_0, a_1, \dots, a_{n-1})$	$n$ -tuple, i.e. ordered collection of items of length $n$ . If $n = 1$ , we call
133		it a “singleton”, if $n = 2$ , we call the tuple a “pair”. Finally, if $n = 3$ , we call it
134		a “triple”. We use the terms “tuples” and “lists” interchangeably.
135	$S \times T$	Cartesian product of sets $S$ and $T$ , i.e. set of all ordered pairs $\{(x, y) : x \in S \wedge y \in T\}$
136	$S^n$	$n$ -fold Cartesian product of $S$ with itself, i.e. $S^n = \{(x_1, \dots, x_n) : \forall i \in \{1, \dots, n\}, x_i \in S\}$ , where $n \in \mathbb{N}$
137		
138	$\Lambda$	General notation for an alphabet, i.e. a <i>non-empty finite set</i> such that every string
139		(ordered collection of symbols, or letters, all in $\Lambda$ ) has a unique decomposition.
140		The number of symbols in a string is denoted the “length” of the string
141	$\varepsilon$	The empty string. $\varepsilon$ is a string over any alphabet.
142	$\Lambda^n$	Set of all strings, defined over alphabet $\Lambda$ , containing $n$ symbols (i.e. “of length
143		$n$ ”)
144	$\Lambda^*$	The Kleene star of $\Lambda$ represents the set of all strings of finite length, defined over
145		alphabet $\Lambda$ , including the empty string $\varepsilon$ , i.e. $\Lambda^* = \bigcup_{n \in \mathbb{N}} \Lambda^n$
146	$\text{length}(x)$	$\text{length} : \Lambda^* \rightarrow \mathbb{N}$ computes the length of a string $x$ defined over $\Lambda$ , i.e. $\text{length}(x)$
147		returns the number of symbols composing the string $x$ . By convention, $\text{length}(\varepsilon) =$
148		0
149	$x \parallel y$	Infix notation for the concatenation function, $\parallel : \Lambda^* \times \Lambda^* \rightarrow \Lambda^*$ . If $\text{length}(x) =$
150		$n, \text{length}(y) = m : (n, m) \in \mathbb{N}^2$ , then $\text{length}(z) = n + m, z = x \parallel y$
151	$\text{trunc}_x(k)$	$\text{trunc} : \Lambda^* \rightarrow \Lambda^k$ is the truncation function that returns the sequence formed
152		from the the first $k$ elements of $x$ , where $x \in \Lambda^*$ . If $k > \text{length}(x)$ , then
153		$\text{trunc}_x(k) = x$
154	$x[a:b]$	$[\cdot] : \Lambda^n \times \mathbb{N} \times \mathbb{N} \rightarrow \Lambda^{\leq b-a}$ is the slice function that, if $b \geq a$ , returns the string
155		starting at index $\min(n, a)$ of $x$ and finishing at index $\min(n, b)$ . The function
156		additionally interprets $x[:b]$ as $x[0:b]$ and $x[a:]$ as $x[a:n]$
157	$\text{pad}_n(x)$	$\text{pad} : \Lambda^{\leq n} \rightarrow \Lambda^n$ is the padding function which pads $x$ by 0’s to reach a size of
158		$n$ . The padding depends on the variable type and endianness
159	$\text{append}(l, x)$	$\text{append} : D^n \times D \rightarrow D^{n+1}$ is the algorithm that appends $x$ to the list of $n$
160		element(s) $l$ , if all $x$ and $l$ share the same datatype $D$
161	$\mathbb{B}$	Alphabet of binary symbols, i.e. $\{0, 1\}$
162	$\langle \mathbf{g}_1, \dots, \mathbf{g}_l \rangle$	Cyclic group generated by $\{\mathbf{g}_1, \dots, \mathbf{g}_l\}$
163	$(q, \mathbb{G}, \mathbf{g}, \otimes)$	Description of the cyclic group $\mathbb{G} = \langle \mathbf{g} \rangle$ of order $q$ , with operation $\otimes$

164	$\mathbb{Z}/r\mathbb{Z}$	Quotient group defined as the set of equivalence classes modulo $r$ . $\mathbb{Z}/r\mathbb{Z}$ , also written $\mathbb{Z}_r$ , is an additive group. If $r = p$ a prime number, then $\mathbb{Z}_p = \{0, \dots, p-1\} = \mathbb{Z}/p\mathbb{Z}$ is a finite field of elements modulo prime $p$ , also denoted $\mathbb{F}_p$ , where $0_{\mathbb{F}_p}$ and $1_{\mathbb{F}_p}$ respectively represent the additive and multiplicative identity
165		
166		
167		
168	$\mathbb{F}_q$	Finite field of cardinality $q = p^m$ , where $p$ is prime
169	$\llbracket x \rrbracket$	Represents the encoding of the scalar $x$ in a group $\mathbb{G}$ described as $(p, \mathbb{G}, \langle \mathbf{g} \rangle, \otimes)$ , i.e. $\llbracket x \rrbracket = x \cdot \llbracket 1 \rrbracket = \mathbf{g} \otimes \dots \otimes \mathbf{g}$ ( $x$ times). Thus, by convention, $\llbracket 1 \rrbracket = \mathbf{g}$
170		
171	$\bullet$	Represents an inline operator for bilinear pairing. That is for a bilinear pairing from $\mathbb{G}_1 \times \mathbb{G}_2$ to $\mathbb{G}_T$ and elements $\llbracket a \rrbracket_1, \llbracket b \rrbracket_2$ we write $\llbracket ab \rrbracket_t = \llbracket a \rrbracket_1 \bullet \llbracket b \rrbracket_2$
172		
173	$\lceil x \rceil$	Round $x \in \mathbb{R}$ to the next integer
174	$\lfloor x \rfloor$	Round $x \in \mathbb{R}$ to the previous integer
175	$\log_b(x)$	Logarithm with respect to base $b$ , i.e. $x = b^y, \log_b(x) = y$
176	<b>Algorithmic notations</b>	
177	$x \leftarrow \$_{\mathcal{X}}$	Element chosen uniformly at random from set $\mathcal{X}$
178	$x \leftarrow y$	The value $y$ is assigned to the variable $x$ (i.e. “ $x$ receives the value $y$ ”)
179	PPT	Probabilistic polynomial time. A polynomial time algorithm $A$ is one for which there exists a polynomial $f$ such that the running time of $A$ on input $x \in \{0, 1\}^*$ is $f( x )$ . A probabilistic algorithm has the ability to “flip” random coins and use the result of these coin tosses in its computation
180		
181		
182		
183	NUPPT	Non-uniform probabilistic polynomial time
184	$\mathcal{O}(f)$	Big-O notation
185	il, kl, nl, rl, ol	The input il, key kl, nonce nl, randomness rl and output ol length
186	<b>Cryptography notations</b>	
187	$\mathcal{O}^X(n)$	Public oracle for algorithm $X$ which can be accessed at most $n$ times; $\mathcal{O}^X$ is an unrestricted oracle for algorithm $X$
188		
189	$\lambda$	Security parameter ( $\lambda \in \mathbb{N}$ )
190	negl	Negligible function. In this document, negligible will usually mean $\mathcal{O}(2^{-\lambda})$
191	poly	Polynomial function
192	$\mathcal{A}$	Adversary algorithm



193	$\text{Adv}_{F,\mathcal{A}}^{\text{prop}}(\lambda)$	Advantage of the adversary $\mathcal{A}$ with regard to the attack game <b>prop</b> on $F$	
194		(e.g. $F$ can be a function, a family of functions or a group on which a given	
195		property represented by the game <b>prop</b> is supposed to hold)	
196	$\text{prop}^{\mathcal{A}}$	Adversary $\mathcal{A}$ running a security game <b>prop</b>	
197	<b>Zeth notations</b>		
198	$\pi$	Zero-knowledge non-interactive argument of knowledge (zk-snark for short). $\pi$ is	
199		also informally referred to as a “zk-proof”, or simply “proof” when clear from	
200		context	
201	$\mathcal{P}_{\mathcal{Z}}$	Standard notation for a <b>Zeth</b> user	
202	$\widetilde{\text{Mixer}}$	The mixer smart-contract instance	
203	<b>EncSch</b>	In-band encryption scheme used to share <b>Zeth</b> notes	
204	<b>Ethereum notations</b>		
205	<b>Account</b>	Standard notation for an <b>Ethereum</b> account object	
206	$\widetilde{\text{Cntrct}}$	Standard notation for an <b>Ethereum</b> smart-contract instance	
207	$\mathcal{P}_{\mathcal{E}}$	Standard notation for an <b>Ethereum</b> user	
208	$\varsigma$	Mapping representing the <b>Ethereum</b> state (i.e. “World state”)	
209	$\varsigma[a]$	Account object stored at address $a$ in $\varsigma$ if it exists, $\perp$ is returned otherwise	
210	<b>Constants</b>		
211	<b>ADDRLEN</b>	The bit-length of an <b>Ethereum</b> address	160 <i>bits</i>
212	<b>BLAKE2sCLEN</b>	Output size of Blake2s compression function [ANWOW13]	256 <i>bits</i>
213	<b>BNFIELDCAP</b>	Field capacity of $\mathbb{F}_{\mathbf{r}_{\text{BN}}}$ . It is defined as the maximum bit length $l$ such that	
214		all numbers $x \in \mathbb{B}^l$ are field elements $x \in \mathbb{F}_{\mathbf{r}_{\text{BN}}}$ : $\max_n \{\forall x \in \{0, 1\}^n, x \in \mathbb{F}_{\mathbf{r}_{\text{BN}}}\}$	
215			$\lceil \log_2 \mathbf{r}_{\text{BN}} \rceil = 253 \text{ bits}$
216	<b>BNFIELDLEN</b>	Bit-length of a field element $x \in \mathbb{F}_{\mathbf{r}_{\text{BN}}}$	$\lceil \log_2 \mathbf{r}_{\text{BN}} \rceil = 254 \text{ bits}$
217	<b>BYTELEN</b>	Bit-length of a byte	8 <i>bits</i>
218	<b>ENCZETHNOTELEN</b>	Size of an encrypted note	1216 <i>bits</i>
219	<b>WORDLEN</b>	Width of a storage cell on the Ethereum Virtual Machine stack, i.e. size of a	
220		word on the EVM	256 <i>bits</i>
221	<b>JSIN, JSOUT, JSMAX</b>	The number of inputs and outputs of a joinsplit and $\text{JSMAX} = \max(\text{JSIN}, \text{JSOUT})$	

222	<b>KEK256DLEN</b>	Message digest size of Keccak256 [GJMG11]	256 <i>bits</i>
223	<b>MKDEPTH</b>	The depth of the Merkle tree used to store commitments	
224	<b>p<sub>SECP</sub></b>	Prime defining the finite field of curve secp256k1 [?]	
225	<b>r<sub>BN</sub></b>	Characteristic of the scalar field of BN-254, $r_{BN} = 21888242871839275222246405$	
226		745257275088548364400416034343698204186575808495617 [?]	
227	<b>SECPFIELDLEN</b>	Bit-length of a field element $x \in \mathbb{F}_{p_{SECP}}$	$\lceil \log_2 p_{SECP} \rceil = 256$ <i>bits</i>
228	<b>SHA256BLEN</b>	Block size of SHA256 [oST15, Figure 1]	512 <i>bits</i>
229	<b>SHA256DLEN</b>	Message digest size of SHA256 [oST15, Figure 1]	256 <i>bits</i>
230	<b>SHA256MLEN</b>	Message size of SHA256 [oST15, Figure 1]	$< 2^{64}$ <i>bits</i>
231	<b>DGAS</b>	The intrinsic gas cost of an Ethereum transaction	21000 Wei
232	<b>ZVALUELEN</b>	Size in bits of the transferable maximal value	64 <i>bits</i>

233

# Change Log

234

- **Version:** 0.0, **Date:** 04/12/2019, **Contributor:** Antoine Rondelet, **Description:** Creation of the document. Established initial table of content and started to populate sections with bullet lists to develop in further versions of the document.

235

236

237

- **Version:** 0.1, **Date:** 20/12/2019, **Contributor:** Antoine Rondelet, **Description:** Refactored the structure of the document. Finalized the table of content, wrote sections on notations and preliminaries, and introduced the content related to the malleability fix.

238

239

240

241

- **Version:** 0.2, **Date:** 24/02/2020, **Contributor:** Clearmatics Cryptography R&D, **Description:**

242

243

- **Date:** 26/02/2020, **Contributor:** Duncan Tebbs, **Description:** Wrote section on wallet implementation and side-channel attacks considerations.

244

245

- **Date:** 02/03/2020, **Contributor:** Giuseppe Giffone, **Description:** Changed merkle tree hash function to MiMC compression function.

246

247

- **Date:** 04/03/2020, **Contributor:** Duncan Tebbs, Michal Zajac, **Description:** Added background on Groth16 SNARK and SNARK scheme instantiation in the protocol.

248

249

250

- **Date:** 04/03/2020, **Contributor:** Raphael Toledo, **Description:** Wrote section on the packing policy and corresponding attack.

251

252

- **Date:** 04/03/2020, **contributor:** Duncan Tebbs, **Description:** Refactored the data structures preliminary section.

253

254

- **Date:** 24/03/2020, **Contributor:** Raphael Toledo, **Description:** Changed the PRF and commitment instantiation with Blake2s compression function.

255

256

- **Date:** 17/04/2020, **Contributor:** Giuseppe Giffone, **Description:** Added DHAES encryption scheme.

257

258

- **Version:** 0.3, **Date:** 09/06/2020, **Contributor:** Antoine Rondelet, **Description:** Fixed various inconsistencies throughout the document (notational mistakes in document body and in proofs, latex macros, and typos).

259

260

# Chapter 1

## Preliminaries

**Zeth** is a protocol which enables private transactions on **Ethereum** [Woo19]. It is a modification of the Decentralized Anonymous Payment (DAP) system **ZeroCash** [BSCG<sup>+</sup>14]. The design described in [RZ19] presents the mechanisms by which **ZeroCash** can be used on **Ethereum**, and argues that the information leakages of the solution are well defined and controlled. This document, however, serves as a specification of the protocol and provides security fixes and optimizations from the first proof of concept release of the protocol [Cle19].

This document assumes familiarity with blockchain and **Ethereum** in particular. It does not, in any way, aim at replacing the Ethereum yellow paper [Woo19]. The reader is strongly advised to read about **Ethereum** before delving into this specification document.

The key words **MUST**, **MUST NOT**, **SHOULD**, **SHOULD NOT**, **MAY**, and **RECOMMENDED** in this document are to be interpreted as described in [Bra97] when they appear in **ALL CAPS**. These words may also appear in this document in lower case as plain English words, absent their normative meanings.

### 1.1 Data Structures and Representation

#### 1.1.1 Structured Data

When describing the operations to be performed and the data to be manipulated as part of the protocol, we commonly employ tuples of related data where each element of the tuple has some associated semantic meaning often must satisfy some conditions. In this section, we develop a framework to reason about such *structured* data, where a single datum may consist of one or more logical parts (called *fields*). The framework is built on top of simple mathematical concepts such as sets, and mappings between them, ensuring that we can always reason about structured data in a rigorous way. We also define notation to aid the specification of structured data, and to refer to specific components of a datum. This will be used extensively in the specification of the protocol.

As a simple motivating example, consider a protocol that processes data relating to individual people. This fictional system may send and receive data such as *name*, *age*

or *address* for a single person, grouping this data into a logical unit. Further, each piece of data must satisfy specific conditions (*name* must be a series of characters from some alphabet, *age* must be a positive integer, etc.) We shall make use of this example several times during the formulation below.

In what follows, let  $\text{STR} = \{a, b, \dots, y, z\}^*$  (the Kleene star of the *Roman alphabet*). In our formulation, field names  $f_i$  will be elements in this set. (Note that a similar formulation could be made using an arbitrary set, such as the same alphabet augmented with specific symbols, or the alphabet of a different language. Our choice of  $\text{STR}$  here is for simplicity).

We begin by defining a datatype as a set of values called “fields”, each with a “name” from  $\text{STR}$ . Abstract sets are used constrain the values of each field.

**Definition 1** (Structured Data Type). *Let  $f_0, \dots, f_{n-1}$  be  $n$  distinct elements of  $\text{STR}$  and let  $V_0, \dots, V_{n-1}$  be sets, for some  $n \in \mathbb{N}$ . We define the structured datatype  $\mathbf{T}$  with fields  $\{(f_i, V_i)\}_{i \in [n]}$  to be a set of values:*

$$\mathbf{T} = V_0 \times \dots \times V_{n-1}$$

with associated post-fix “dot” operators  $.f_i : \mathbf{T} \rightarrow V_i$  for  $i = 0, \dots, n-1$ , acting on values  $\mathbf{x} \in \mathbf{T}$  to extract the individual elements:

$$\mathbf{x}.f_i = v_i, \text{ where } \mathbf{x} = (v_0, \dots, v_{n-1}) \in \mathbf{T}$$

Here, we say that the  $i$ -th field has field name  $f_i$ , with value set  $V_i$ . Each “dot” operator  $.f_i$  extracts the  $i$ -th component, or the value with field name  $f_i$ .

**Example 1.** Consider our example protocol that processes information about people. A potentially useful structured data type **Person** may be defined with fields:

$$\{(name, \text{STR}), (age, \mathbb{N}), (height, \mathbb{R}^+)\}$$

Values  $\mathbf{p}$  in **Person** are simply tuples in  $\text{STR} \times \mathbb{N} \times \mathbb{R}^+$ , with semantic meaning (*name*, *age*, *height*) assigned to each component of  $\mathbf{p}$ .

Examples of valid elements in **Person** include  $\mathbf{a} = (\text{alice}, 28, 1.65)$  and  $\mathbf{b} = (\text{bob}, 31, 1.74)$ , where the following equalities hold:

$$\mathbf{a}.name = \text{alice},$$

$$\mathbf{b}.age = 31,$$

$$\mathbf{b}.height = 1.74;$$

For clarity, structured data types may be specified using tables of names, descriptions and value sets, rather than sets of the form  $\{(f_i, V_i)\}_{i \in [n]}$ . Similarly, it is frequently convenient to specify the *field names* alongside values in tuples when defining structured data values.

309 **Example 2.** Person from Example 1 might be described in table-form as follows:

Field	Description	Values
<i>name</i>	Name of the person	STR
<i>age</i>	Age in years	$\mathbb{N}$
<i>height</i>	Height in meters	$\mathbb{R}^+$

**Example 3.** The values **a** and **b** in Example 1 might be written as follows:

$$\mathbf{a} = (\text{name} : \text{alice}, \text{age} : 28, \text{height} : 1.65)$$

$$\mathbf{b} = (\text{name} : \text{bob}, \text{age} : 31, \text{height} : 1.74)$$

310 **Remark 1** (“dot” operators in assignment). The “dot” operators may be used in al-  
 311 gorithm descriptions to indicate assignment to a specific component. For example  
 312  $\mathbf{a}.\text{age} \leftarrow 29$  means that the value of the age field of **a** is replaced by the value 29.

Formally, for a structured data type **T** with fields  $\{(f_i, V_i)\}_{i \in [n]}$  where  $\mathbf{x} = (v_0, \dots, v_{n-1}) \in \mathbf{T}$  and  $v'_i \in V_i$ :

$$\mathbf{x}.f_i \leftarrow v'_i$$

is equivalent to:

$$\mathbf{x} \leftarrow (v_0, \dots, v_{i-1}, v'_i, v_{i+1}, \dots, v_{n-1})$$

313 We define one further operator and related assignment notation, convenient in cases  
 314 where  $V_i = X^m$  for sets  $X$  and  $m \in \mathbb{N}$ .

**Definition 2** (Square bracket operator). For  $m \in \mathbb{N}$  and set  $X$ , define the operator  $[ ] : X^m \times [m] \rightarrow X$  as:

$$\mathbf{x}[i] = x_i \text{ where } \mathbf{x} = (x_0, \dots, x_m)$$

For the set  $X^*$ , the operator takes the form  $[ ] : X^* \times \mathbb{N} \rightarrow X$ , defined as:

$$\mathbf{x}[i] = \begin{cases} x_i & \text{if } \text{length}(\mathbf{x}) > i \text{ where } \mathbf{x} = (x_0, \dots) \\ \perp & \text{otherwise} \end{cases}$$

**Remark 2** (Square bracket operators in assignment). Similarly to the notation in Remark 1, we develop notation for square bracket operator  $[ ]$  in assignment statements. Let  $\mathbf{x} = (x_0, \dots, x_{m-1})$  be a member of  $X^m$ , and  $x'_i$  be some element in  $X$ . The statement:

$$\mathbf{x}[i] \leftarrow x'_i$$

is equivalent to:

$$\mathbf{x} \leftarrow (x_0, \dots, x_{i-1}, x'_i, x_{i+1}, \dots, x_{m-1})$$

315 Informally, this can be interpreted as replacing the  $i$ -th component of  $\mathbf{x}$  with  $x'_i$ .

316 **Remark 3** (Deep structures and chained “dot” operators). Consider the case of struc-  
 317 tured data  $\mathbf{T}$  with fields  $\{(f_i, V_i)\}_{i \in [n]}$  for  $n \in \mathbb{N}$ . Let  $\mathbf{T}'$  be another structured data type  
 318 with fields  $\{(f'_i, V'_i)\}_{i \in [n']}$  for  $n' \in \mathbb{N}$ , and assume that  $V_j = \mathbf{T}'$  for some  $j \in [n]$ . Infor-  
 319 mally, the values of the  $j$ -th field of elements of  $\mathbf{T}$  are themselves structured data of type  
 320  $\mathbf{T}'$ .

321 In this case, “dot” operators may be chained, so that  $\mathbf{x}.f_j.f'_k$  refers to the  $k$ -th field  
 322  $v'_k$  of the  $j$ -th field  $v_j$  of  $\mathbf{x} \in \mathbf{T}$ .

323 **Example 4.** Define a structured data type **Address** with fields  $(\text{country}, \text{STR}), (\text{zipcode}, \text{STR})$ .  
 324 We redefine the structured data type **Person** from Example 1, with an extra field *address*  
 325 of type **Address**. That is, **Person** is the structured data type with fields:

Field	Description	Values
<i>name</i>	Name of the person	STR
<i>age</i>	Age in years	$\mathbb{N}$
<i>height</i>	Height in meters	$\mathbb{R}^+$
<i>address</i>	Address of the person	<b>Address</b>

An example element **a** in **Person** is:

$$\begin{aligned} \mathbf{a} = & (\text{name} : \text{alice}, \\ & \text{age} : 28, \\ & \text{height} : 1.65, \\ & \text{address} : (\text{country} : \text{UK}, \text{zipcode} : \text{SW1A})) \end{aligned}$$

In this case, the following equalities using the dot and square bracket operators all hold:

$$\begin{aligned} \mathbf{a}.\text{name} &= \text{alice} \\ \mathbf{a}.\text{height} &= 1.65 \\ \mathbf{a}.\text{address}.\text{country} &= \text{UK} \\ \mathbf{a}.\text{address}.\text{zipcode} &= \text{SW1A} \\ \mathbf{a}.\text{address}.\text{country}[1] &= K \end{aligned}$$

### 326 1.1.2 Representations

327 The binary alphabet  $\{0, 1\}$ , denoted  $\mathbb{B}$ , is used to represent the presence or absence of an  
 328 electrical signal in a computer. In fact, every piece of information in a computer is rep-  
 329 resented as a string of bits. We assume the existence of an efficient binary representation  
 330 for some set of primitive datatypes (such as the natural numbers  $\mathbb{N}$ , or alphanumeric  
 331 characters). Structured data types built up from primitive types (as described above)  
 332 can then recursively be assigned similarly efficient representations. This is used to define  
 333 the following functions to *encode* data to its bit-string representation, and to *decode* such  
 334 bit-strings back to elements of the original type.

**Definition 3.** For a set  $X$  of values which are to be represented as bit strings, we define functions:

$$\begin{aligned}\text{encode}_X &: X \rightarrow \mathbb{B}^* \\ \text{decode}_X &: \mathbb{B}^* \rightarrow X \cup \perp\end{aligned}$$

satisfying

$$\text{decode}_X(\text{encode}_X(x)) = x \quad \forall x \in X$$

335 to be the functions which encode (resp. decode) elements of  $X$  into (resp. from) the bit-  
336 string representations chosen above. We note that  $\text{decode}_X$  may return  $\perp$  in the case  
337 that the input bit-string is malformed.

338 Without ambiguity, we overload the functions `encode` and `decode` to mean  $\text{encode}_X$   
339 and  $\text{decode}_X$  where the set  $X$  is clear from context.

In the following sections, we assume that elements of  $\mathbb{N}$  are encoded as big-endian binary numbers in the natural way. We denote by  $\mathbb{N}_b$  the set of natural numbers that can be uniquely encoded in this way using  $b$  bits. In other words,

$$\mathbb{N}_b = \{x \in \mathbb{N} \text{ s.t. } \text{encode}_{\mathbb{N}}(x) \in \mathbb{B}^b\}$$

## 340 1.2 Ethereum

341 In a nutshell, **Ethereum** is a distributed deterministic state machine, consisting of a glob-  
342 ally accessible singleton state (“the World state”) and a virtual machine that applies  
343 changes to that state [AG18]. State transitions in the state machine are represented  
344 by transactions on the system. As such, each transaction represents a change in the  
345 global state represented as a Merkle Patricia Tree [wc] whose nodes are objects called  
346 “accounts” (Section 1.2.1). The Ethereum Virtual Machine (EVM) allows state transi-  
347 tions to be specified by creating a type of accounts which are associated with a piece  
348 of code (smart-contracts). The code of such accounts, and so, the corresponding state  
349 transitions, can be executed to transition to another state in the automata, by creating  
350 a transaction that calls the given piece of code (Section 1.2.2).

351 To prevent unbounded state transitions in the state machine, each instruction exe-  
352 cuted by the EVM is associated with a cost in **Wei**, referred to as “the gas necessary to  
353 run the operation”. The “gas cost” of a transaction needs to be paid by the transaction  
354 originator (deduced from their account balance), and is awarded to the miner (added  
355 to their account balance) who successfully mines the block containing the transaction.  
356 In addition to the cost of every instruction executed as part of a state transition, every  
357 transaction has an intrinsic “gas cost” of  $\text{DGAS Wei}$  [Woo19, Appendix G]. Bounding  
358 modifications to the Ethereum state by the amount of **Wei** held in the transaction origi-  
359 nator’s account allows the system to avoid the Halting problem<sup>1</sup> and protect against a  
360 range of Denial of Service (DOS) attacks.

---

<sup>1</sup>[https://en.wikipedia.org/wiki/Halting\\_problem](https://en.wikipedia.org/wiki/Halting_problem)



### 361 1.2.1 Ethereum Account

362 An **Ethereum** account [Woo19, Section 4.1] is an object containing 4 attributes, as rep-  
 363 resented Table 1.1. We distinguish two types of accounts:

- 364 • “Externally Owned Accounts” (EOA), that are created by derivation of an ECDSA se-  
 365 cret key; and
- 366 • Smart-contract accounts, that are derived from EVM code specifying a state tran-  
 367 sition on the state machine.

368 Each account object is accessible in the Merkle Patricia Tree representing the “World  
 369 state”, by a unique **ADDRLEN**-bit long identifier called the address. In the context of EOA,  
 370 the address is obtained by generating a new ECDSA [JMV01] key pair  $(sk, vk)$  over curve  
 371 **secp256k1** [Qu99] and taking the the rightmost **ADDRLEN** bits of the **Keccak256** hash of  
 372 the verification key  $vk$ .

Field	Description	Value
<i>nce</i>	The nonce of an account is a scalar value representing the number of transactions that have originated from the account, starting at 0.	$\mathbb{N}_{\text{WORDLEN}}$
<i>bal</i>	The balance of an account is a scalar value representing the amount of Wei in the account.	$\mathbb{N}_{\text{WORDLEN}}$
<i>sRoot</i>	The storage root is the <b>Keccak256</b> hash representing the storage of the account.	$\mathbb{B}^{\text{KEK256DLEN}}$
<i>codeh</i>	The code hash is the hash of the EVM code governing the account. If this field is the <b>Keccak256</b> hash of the empty string, then the account is said to be an “Externally owned Account” (EOA), and is controlled by the corresponding ECDSA private key. If, however, this field is not the <b>Keccak256</b> hash of the empty string, the account represents a smart contract whose interactions are governed by its EVM code.	$\mathbb{B}^{\text{KEK256DLEN}}$

Table 1.1: Ethereum Account structure

### Note

In the rest of this document, we will refer to an *Ethereum user*  $\mathcal{U}_{\mathcal{E}}$  as a person, modeled as an object, holding *one*<sup>a</sup> secret key,  $sk$  (object attribute), associated with an existing EOA in the “World state”. We denote by  $\mathcal{U}_{\mathcal{E}}.Addr$  the **Ethereum** address of  $\mathcal{U}_{\mathcal{E}}$  derived from  $\mathcal{U}_{\mathcal{E}}.sk$ , and which allows  $\mathcal{U}_{\mathcal{E}}$  to access the state of their account  $\varsigma[\mathcal{U}_{\mathcal{E}}.Addr]$ .

We denote by **SmartC** a smart-contract instance/object (i.e. deployed smart-contract with an address, Section 1.2.2), and denote by **SmartC.Addr** its address.

<sup>a</sup>The same physical person may correspond to multiple “Ethereum users” and thus control multiple accounts in the Merkle Patricia Tree.

373

### 1.2.2 Ethereum Transaction

375 We now briefly mention what **Ethereum** transactions [Woo19, Section 4.2] are, and how  
 376 they are created, signed and validated. Once more, the reader is highly encouraged to  
 377 refer to [Woo19] for a detailed presentation. Informally, a transaction object ( $tx$ ) is a  
 378 signed message originating from an **Ethereum** user  $\mathcal{U}_{\mathcal{E}}$  (the *transaction originator*, or  
 379 simply *sender*) that represents a state transition on the distributed state machine (i.e. a  
 380 change in the “World state”  $\varsigma$ ).

### 381 Raw Transaction

382 In the following, we define a raw transaction as an unsigned transaction (Table 1.2).

Field	Description	Value
$nce$	Transaction nonce	$\mathbb{N}_{\text{WORDLEN}}$
$gasP$	gasPrice	$\mathbb{N}_{\text{WORDLEN}}$
$gasL$	gasLimit	$\mathbb{N}_{\text{WORDLEN}}$
$to$	Recipient’s address	$\mathbb{B}^{\text{ADDRLEN}}$
$val$	Value of the transaction in Wei	$\mathbb{N}_{\text{WORDLEN}}$
$init / data$	Contract Creation data $init$ Message call data $data$	$\mathbb{B}^*$

Table 1.2: Structure of a *raw transaction data type* **TxRawDType**

### 383 Finalizing raw transactions

384 A raw transaction needs to be finalized to be accepted. In the context of this document,  
 385 “finalizing a raw transaction” will be a synonym of “signing a raw transaction”. The  
 386 transaction structure is represented in Table 1.3.

Field	Description	Value
$tx_{raw}$	Raw transaction object	<b>TxRawDType</b>
$v$	Field $v$ of ECDSA signature used for public key recovery	$\mathbb{B}^{\text{BYTELEN}}$
$r$	Field $r$ of ECDSA signature [Por13]	$\mathbb{F}_{\text{PSECP}}$
$s$	Field $s$ of ECDSA signature [Por13]	$\mathbb{F}_{\text{PSECP}}$

Table 1.3: Structure of a (finalized) *transaction data type* TxDType

We define the transaction generation function Fig. 1.1 as the function taking the sender's ECDSA signing key and the components of a raw transaction as arguments, and returning a signed (or finalized) transaction ( $tx_{final}$  or  $tx$  for short).

$$\begin{aligned}
tx_{final} &= \text{TxGen}(sk_{\text{ECDSA}}, nce_{in}, gasP_{in}, gasL_{in}, to_{in}, val_{in}, init_{in}, data_{in}) \\
tx_{final} &= \{ \\
&\quad \left. \begin{array}{l} nce : nce_{in}, \\ gasP : gasP_{in}, \\ gasL : gasL_{in}, \\ to : to_{in}, \\ val : val_{in}, \\ init/data : init_{in}/data_{in}, \end{array} \right\} tx_{raw} \\
&\quad \left. \begin{array}{l} v : \sigma_{\text{ECDSA}}.v, \\ r : \sigma_{\text{ECDSA}}.r, \\ s : \sigma_{\text{ECDSA}}.s \end{array} \right\} \sigma_{\text{ECDSA}} \\
&\}
\end{aligned}$$

387 To sign a transaction, the sender first computes the hash of the raw transaction  
388 using Keccak256 (Eq. (1.1)), and then uses their ECDSA signing key,  $sk_{\text{ECDSA}}$ , to sign  
389 the obtained digest (Eq. (1.2)). The signature is then appended to the raw transaction  
390 to obtain a finalized transaction (Fig. 1.1).

$$digest_{\text{ECDSA}} = \text{Keccak256}(nce_{in}, gasP_{in}, gasL_{in}, to_{in}, val_{in}, init_{in}/data_{in}) \quad (1.1)$$

$$\sigma_{\text{ECDSA}} = \text{SigSch}_{\text{ECDSA}}.\text{Sig}(sk_{\text{ECDSA}}, digest_{\text{ECDSA}}) (= (v, r, s)) \quad (1.2)$$

---

```

TxGen( $sk_{\text{ECDSA}}, nce_{in}, gasP_{in}, gasL_{in}, to_{in}, val_{in}, init_{in}, data_{in}$ )
1: if  $to_{in} == \emptyset$  do
2:    $tx_{raw} \leftarrow \{nce : nce_{in}, gasP : gasP_{in}, gasL : gasL_{in}, to : to_{in}, val : val_{in}, init : init_{in}\};$ 
3: else
4:    $tx_{raw} \leftarrow \{nce : nce_{in}, gasP : gasP_{in}, gasL : gasL_{in}, to : to_{in}, val : val_{in}, data : data_{in}\};$ 
5: endif
6:  $\sigma_{\text{ECDSA}} \leftarrow \text{SigSch}_{\text{ECDSA}}.\text{Sig}(sk_{\text{ECDSA}}, \text{Keccak256}(tx_{raw}));$ 
7:  $tx_{final} \leftarrow \{tx_{raw}, v : \sigma_{\text{ECDSA}}.v, r : \sigma_{\text{ECDSA}}.r, s : \sigma_{\text{ECDSA}}.s\};$ 
8: return  $tx_{final}$ ;

```

---

Figure 1.1: Transaction generation function TxGen

**Remark 4.** As one can see, there is no “from” attribute in a transaction. The sender’s *Ethereum* address can be recovered from the ECDSA signature. This method is defined in the *Ethereum* yellow paper as a “sender function”  $S$  [Woo19, Appendix F] which maps each transaction to its sender.

## Types of transactions

While only two types of transactions are described in [Woo19, Section 4.2]; namely those which result in message calls and those which result in the creation of new accounts with associated code, we will instead differentiate the types of transactions based on their purpose. The reader is encouraged to read [Woo19] for a formal discussion.

Informally, a transaction can be used to achieve three things: transferring Wei from an EOA to another EOA, creating a new account with associated code (i.e. “deploying a smart-contract”), and calling a function of a smart-contract. We will detail here the differences between these usages.

- **Creating a contract:** The  $tx.to$  address is set to  $\emptyset$  in the transaction. The contract creation data ( $tx.init$ ) includes the new contract’s code. The contract address is computed as the rightmost ADDRLEN bits of the Keccak256 hash of the RLP encoding [wc19] of the transaction originator’s address and account nonce [Woo19, Section 6].
- **Calling a contract function:** The  $tx.to$  address is set to the address of the contract. The message call data byte array ( $tx.data$ ) is set to the contract’s function address (or “Function Selector” [abi]) which are the first 4 bytes of the Keccak256 hash of the function signature, and the function input arguments (WORDLEN bits per input) [Woo19, Section 8].
- **Transferring Wei from an EOA to another EOA:** This corresponds to a “plain transaction” spending Wei from an address to send them to another. In that case the  $tx.to$  address corresponds to the recipient’s address while the transaction data is left empty.

### Note

In order to keep notations simple, we assume, in the rest of the document, that smart-contract functions are uniquely determined by their name. As such, we denote by  $\text{FS}(\cdot): \mathbb{B}^* \rightarrow \mathbb{B}^{4 \cdot \text{BYTELEN}}$  the function that takes a function name as input and returns its function selector.

418

### 419 Transaction validity

420 Importantly, not all finalized transactions constitute valid state transitions on the state  
421 machine [Woo19, Section 6]. We denote by `EthVerifyTx` the function that takes an  
422 **Ethereum** transaction object  $tx$  as input and return `true` (resp. `false`) if  $tx$  is valid  
423 (resp. invalid). To be deemed valid, a transaction **MUST** satisfy *all* the following condi-  
424 tions:

- 425 1. The transaction is correctly RLP encoded, with no additional trailing bytes;
- 426 2. the transaction signature  $(v, r, s)$  is valid;
- 427 3. the transaction nonce  $(tx.nce)$  is valid, i.e. it is equal to the account nonce of the  
428 transaction originator;
- 429 4. the gas limit is no smaller than the gas used by the transaction;
- 430 5. the transactor has enough funds on his account balance to cover at least the cost  
431  $tx.val + tx.gasP \cdot tx.gasL$ .

### 432 Lifecycle of a transaction, and miners' incentives

433 After the creation of an **Ethereum** transaction  $tx$  by a user from an **Ethereum** client (ma-  
434 chine running a piece of software that enables to be connected to the **Ethereum** network),  
435 the transaction is broadcasted to the network and received by a set of peers/nodes.

436 The transaction is then stored in each node's transaction pool, which is a data  
437 structure containing all transactions that should be validated (pending transactions) by  
438 the node and mined. To maximize miners' returns, the transaction pools are ordered  
439 according to the gas price of the transactions. As such, transactions with the highest  
440  $tx.gasP$  are subject to be validated and included into a block first. Once  $tx$  is selected  
441 from the transaction pool, it is validated (fed into `EthVerifyTx`), executed, and included  
442 into a block (i.e. "mined"). The block is then broadcasted to all the nodes of the network  
443 and is used as the predecessor for the next block to be mined on the network (i.e. "it is  
444 added to the chain").

### 1.2.3 Ethereum events and Bloom filters

The EVM contains the set of “LOGX” instructions enabling smart-contract functions to “emit events” (i.e. log data) when they are executed<sup>2</sup>

As such, when a block is generated by a miner or verified by the rest of the network, the address of any logging contract, and all the indexed fields from the logs generated by executing those transactions are added to a Bloom filter [Blo70], which is included in the block header [Woo19, Section 4.3]. Importantly, the actual logs *are not included in the block data* in order to save space. As such, when an application wants to find (“consume”) all the log entries from a given contract, or with specific indexed fields (or both), the node can quickly scan over the header of each block, checking the bloom filter to see if it may contain relevant logs. If it does, *the node re-executes the transactions from that block, regenerating the logs, and returning the relevant ones to the application* [Joh16].

#### Note

The ability for a smart-contract function to “emit” some pieces of data when executed, and for an application to “consume” such pieces of data, is used in Zeth in order to construct a *confidential receiver-anonymous channel* [KMO<sup>+</sup>13].

## 1.3 zk-SNARKs

In this section we introduce notions necessary to understand zero-knowledge proofs, define properties crucial for them, and introduce zkSNARKs. We refer the reader to Section 3.6 in which we describe the zkSNARK scheme used in Zeth.

### 1.3.1 Preliminary definitions

**NP class of languages.** Since the considered proof systems are designed to work with languages in NP we begin with defining this class. Intuitively, a language  $\mathbf{L}$  belongs to NP if for each element  $prim$  from the language there is a short witness  $aux$  that allows to efficiently<sup>3</sup> verify that in fact  $prim \in \mathbf{L}$ .

**Definition 4** (NP class of languages). *We say that a language  $\mathbf{L}$  belongs to a class NP if there exist a polynomial  $p$  and a Turing machine  $\mathbf{M}$  such that for every primary input  $prim \in \{0, 1\}^*$ ,  $prim \in \mathbf{L}$  iff there exists an auxiliary input  $aux$  such that  $\mathbf{M}$  accepts the pair  $(prim, aux)$  in time at most  $p(\text{length}(prim))$ .*

*The set of all pairs  $(prim, aux)$  acceptable by  $\mathbf{M}$  constitutes an NP relation  $\mathbf{R}$  corresponding to the language  $\mathbf{L}$ .*

<sup>2</sup>see: <https://ethgastable.info/>

<sup>3</sup>Informally we say that an algorithm is efficient if it runs in time polynomial in the size of its inputs.

473 **Non-interactive zero knowledge.** A non-interactive zero-knowledge proof system  
 474 NIZK for an NP language  $\mathbf{L}$  is a tuple of four algorithms  $\text{NIZK} = (\text{KGen}, \text{P}, \text{V}, \text{Sim})$ . NIZK  
 475 for a language  $\mathbf{L}$  and instance  $\text{prim} \in \mathbf{L}$  allows a party, called prover and denoted by  $\text{P}$ ,  
 476 to convince another party, called verifier and denoted by  $\text{V}$ , that  $\text{prim} \in \mathbf{L}$ .

477 Without loss of generality, we focus on zk-proof systems that are universal, that  
 478 is, are able to work with any given NP relation  $\mathbf{R}$ . To that end, we define a *relation*  
 479 *generator*  $\mathcal{R}$  that on input  $1^\lambda$  (i.e. the security parameter represented in unary) outputs  
 480 an NP relation  $\mathbf{R}$ . We assume that the security parameter  $\lambda$  can be easily deduced from  
 481  $\mathbf{R}$ .

482 We require from a NIZK to have three substantial properties:

**Completeness** that assures that an honest prover, who proves that  $\text{prim} \in \mathbf{L}$  succeeds,  
 i.e. gets his proof accepted by the verifier  $\text{V}$ . Formally we require that for any  $\lambda$ ,  
 $\mathbf{R} \leftarrow \mathcal{R}(1^\lambda)$ ,  $(\text{prim}, \text{aux}) \in \mathbf{R}$

$$\Pr \left[ \text{V}(\mathbf{R}, \text{srs}, \text{prim}, \text{P}(\mathbf{R}, \text{srs}, \text{prim}, \text{aux})) \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (\text{srs}, \text{td}) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda) \end{array} \right] = 1 .$$

**Computational soundness** which states that in case  $\text{prim} \notin \mathbf{L}$  the verifier accepts  
 the proof for  $\text{prim}$  with negligible probability only. Formally we require that for  
 any  $\lambda$ ,  $\mathbf{R} \leftarrow \mathcal{R}(1^\lambda)$ ,  $(\text{prim}, \text{aux}) \in \mathbf{R}$  and PPT adversary  $\mathcal{A}$

$$\Pr \left[ \text{V}(\mathbf{R}, \text{srs}, \text{prim}, \pi) \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (\text{srs}, \text{td}) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (\text{prim}, \pi) \leftarrow \mathcal{A}(\mathbf{R}, \text{srs}) \end{array} \right] \leq \text{negl}(\lambda) .$$

**Zero knowledge** assures that the verifier learns from a proof nothing except the ve-  
 racity of the proven statement. More precisely we require that there exist a PPT  
 algorithm  $\text{Sim}$  and negligible function  $\eta(\lambda)$  such that for every adversary  $\mathcal{A}$  and  
 security parameter  $\lambda$

$$\left| \Pr \left[ \mathcal{A}(\mathbf{R}, \pi) = 1 \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (\text{srs}, \text{td}) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (\text{prim}, \text{aux}) \leftarrow \mathcal{A}(\mathbf{R}, \text{srs}); \\ \pi \leftarrow \text{Sim}(\mathbf{R}, \text{srs}, \text{td}, \text{prim}) \end{array} \right] - \Pr \left[ \mathcal{A}(\mathbf{R}, \pi) = 1 \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (\text{srs}, \text{td}) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (\text{prim}, \text{aux}) \leftarrow \mathcal{A}(\mathbf{R}, \text{srs}); \\ \pi \leftarrow \text{P}(\mathbf{R}, \text{srs}, \text{prim}, \text{aux}) \end{array} \right] \right| \leq \eta(\lambda) .$$

483 We say that NIZK is *perfectly* zero-knowledge if  $\eta = 0$ .

484 We note that the existence of the simulator which by using the trapdoor is able to  
 485 make a proof for a false statement (i.e. for  $prim \notin \mathbf{L}$ ) makes the whole zk-proof system  
 486 vulnerable to adversaries that also know the trapdoor. More precisely, an adversary  
 487 who knows a trapdoor  $td$  can break the soundness property. This vulnerability comes  
 488 with each SRS-based NIZK (for languages in NP). Thus in the real-life deployment of a  
 489 SRS-based NIZK it has to be enforced that nobody learns the trapdoor.

490 A zkSNARK scheme, denoted  $\mathbf{ZkSnarkSch}$ , is a special type of NIZK which is equipped  
 491 with two more properties. First, zkSNARKs are arguments *of knowledge*, as such they  
 492 have to follow a stronger definition of soundness, called *knowledge soundness*.

**Knowledge soundness** assures that if a prover provided a proof  $\pi$  for a statement  
 $prim$  acceptable to a verifier, then she knows the corresponding auxiliary input  
 $aux$ . More precisely, we require that for each security parameter  $\lambda$ ,  $\mathbf{R} \leftarrow \mathcal{R}(1^\lambda)$ ,  
 and malicious PPT prover  $\mathcal{A}$  there exists a machine  $\text{Ext}_{\mathcal{A}}$ , called extractor, that  
 given access to randomness  $r$  used by  $\mathcal{A}$  and its inputs, *extracts* the auxiliary input  
 $aux$  from  $\mathcal{A}$ ; that is:

$$\Pr \left[ \begin{array}{c} \neg(\mathbf{R}(prim, aux)) \wedge \\ \mathbf{V}(\mathbf{R}, srs, prim, \pi) \end{array} \middle| \begin{array}{c} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (srs, td) \leftarrow \mathbf{KGen}(\mathbf{R}, 1^\lambda); \\ (prim, \pi) \leftarrow \mathcal{A}(\mathbf{R}, srs; r); \\ aux \leftarrow \text{Ext}_{\mathcal{A}}(\mathbf{R}, srs; r) \end{array} \right] \leq \text{negl}(\lambda).$$

493 Second, zkSNARKs are *succinct*, and so we require that proofs produced by  $\mathbf{ZkSnarkSch.P}$   
 494 are short, i.e. sublinear to the size of the primary and auxiliary inputs.

### 495 1.3.2 Computation representation — Arithmetization

496 In **Zeth** the sender shows that the transaction is correct by arguing (in zero knowledge,  
 497 i.e. hiding private inputs) about correctness of evaluation of some predefined predicate.  
 498 This predicate ensures that the soundness of the blockchain system is not violated,  
 499 i.e. the zk-proof is used to prove that a transaction follows the “rules of the system”  
 500 without disclosing its attributes. The proof system **Zeth** uses operates on an algebraic  
 501 representation of the “predicate to prove”. Informally, representing the computation, to  
 502 generate a proof of computational integrity for, as a set of algebraic constraints is called  
 503 “arithmetization”. One of such representations are Quadratic Arithmetic Programs  
 504 (QAP) [GGPR13], which, following [Gro16], is used in **Zeth**. This representation is  
 505 considered one of the most efficient for a general arithmetic circuits.

506 **QAP (R1CS).** Let  $\mathbf{C}$  be an arithmetic circuit of fan-in 2 over  $\mathbb{F}_p$ . The number of  
 507 multiplication gates in  $\mathbf{C}$  is denoted by  $constNo$ . Likewise, the number of all wires in  $\mathbf{C}$   
 508 is denoted by  $inpNo$ .

509 Before we formally introduce the QAP relation  $\mathbf{R}_{\text{QAP}}$  we provide some intuitions  
 510 behind it. First, we observe that the circuit  $\mathbf{C}$  can be represented by three matrices



511  $\vec{A}, \vec{B}, \vec{C}$  all in  $\mathbb{F}_p^{constNo \times inpNo+1}$  such that the  $i$ -th row in matrix  $\vec{A}$  (and  $\vec{B}$ ) denotes left  
 512 (and right) input to the  $i$ -th multiplication gate, which is also the  $k$ -th input to the  
 513 circuit. That is for a circuit evaluation  $z \in \mathbb{F}_p^{inpNo+1}$  the left input for the  $i$ -th gate is  
 514  $\sum_{j=0}^{inpNo} A_{ij} z_j$  and the right input is  $\sum_{j=0}^{inpNo} B_{ij} z_j$ . Furthermore, entry  $\vec{C}_{ik}$  contains the  
 515 output of  $i$ -th multiplication gate that is  $k$ -th input to the circuit.

516 Second, for the sake of efficiency we represent each matrix as a sequence of poly-  
 517 nomials. Each matrix's column is represented by a polynomial in  $\mathbb{F}_p[X]$  such that the  
 518 column's  $i$ -th input equals polynomial's evaluation at  $\omega^i$  — the  $i$ -th primitive root of  
 519 unity modulo  $p$ . More precisely, we define polynomials:

- 520 •  $u_j(X)$ , for  $j \in \{0, \dots, inpNo\}$ , such that  $u_j(\omega^i) = \vec{A}_{ij}$ ;
- 521 •  $v_j(X)$ , for  $j \in \{0, \dots, inpNo\}$ , such that  $v_j(\omega^i) = \vec{B}_{ij}$ ;
- 522 •  $w_j(X)$ , for  $j \in \{0, \dots, inpNo\}$ , such that  $w_j(\omega^i) = \vec{C}_{ij}$ .

We consider inputs from 1 to  $inpNoPrim$  public (primary input), for some  $inpNoPrim \leq inpNo$ . The rest of the inputs is considered private (auxiliary input). The QAP relation  $\mathbf{R}_{QAP}$  is defined as follows:

$$\mathbf{R}_{QAP} = \left\{ (prim, aux) \left| \begin{array}{l} a_0 = 1; prim = (a_1, \dots, a_{inpNoPrim}) \in \mathbb{F}_p^{inpNoPrim}; \\ aux = (a_{inpNoPrim+1}, \dots, a_{inpNo}) \in \mathbb{F}_p^{inpNo - inpNoPrim}; \\ \sum_{j=0}^{inpNo} a_j u_j(X) \cdot \sum_{j=0}^{inpNo} a_j v_j(X) = \sum_{j=0}^{inpNo} a_j w_j(X) \end{array} \right. \right\}.$$

#### Note

Importantly, we note that efficient computation on standard hardware may not necessarily lead to an efficient QAP representation. As such, a function can be very efficient to evaluate on a standard computer, but very slow to evaluate in QAP form.

523

## 524 1.4 Decentralized Anonymous Payment schemes (DAP)

525 **Zeth** [RZ19] is a Decentralized Anonymous Payment scheme (DAP) [BSCG<sup>+</sup>14, Section  
 526 3] defined on top of an **Ethereum** ledger  $L$ . A DAP is a tuple of polynomial-time algo-  
 527 rithms  $DAP = (\text{Setup}, \text{GenAddr}, \text{SendTx}, \text{VerifyTx}, \text{Receive})$  that manipulate (*create*,  
 528 *spend*) data objects called *Notes*. These objects are bound to a given owner, and have  
 529 a value  $v$  attribute (see: Section 2.1).

530 **System Setup** The algorithm **Setup** takes the security parameter  $\lambda$  as input and gen-  
 531 erates the public parameters  $pp$ . The algorithm **Setup** is executed by a trusted  
 532 party. The resulting public parameters  $pp$  are published and made available to all  
 533 parties.

534 **Creating Zeth addresses** The algorithm `GenAddr` takes as input the public parame-  
 535 ters  $pp$  and generates a new DAP address object  $Addr = \{pub : Addr_{pk}, priv : Addr_{sk}\}$ . More precisely,  $Addr_{pk}$  is an object referred to as the “payment ad-  
 536 dress” (Table 1.4), and  $Addr_{sk}$  is an object referred to as the “private address”  
 537 (Table 1.5) [ZCa19].  
 538

539 **Transfer notes** The algorithm `SendTx` is used to transfer some public input  $vin$  as  
 540 well as the value of a set of input (“old”)  $Notes$  into a set of output (“new”)  $Notes$   
 541 as well as some public output value  $vout$ . The inputs  $Notes$  are marked as  
 542 “consumed” (alternatively, we say that the input  $Notes$  are “spent”). `SendTx` takes  
 543 as inputs the public parameters  $pp$ , the input value and the input (“old”)  $Notes$   
 544 to be transferred, as well as the Merkle root and the Merkle authentications paths  
 545 of the commitments to the input  $Notes$ , the “spending keys” related to the input  
 546  $Notes$ , the output value to create and the “payment addresses” for the output  
 547 (“new”)  $Notes$ . If the joinsplit equation is satisfied, the algorithm returns the new  
 548  $Notes$  and the corresponding **Ethereum** transaction  $tx$ , else it returns  $\perp$ .

549 **Verifying transactions** The algorithm `VerifyTx` checks the validity of a transaction.  
 550 It takes as inputs the public parameters  $pp$ , a transaction and the current ledger  
 551  $L$  and outputs a bit equal 1 iff the transaction is valid.

552 **Receiving notes** The algorithm `Receive` scans the ledger  $L$  and retrieves unspent  $Notes$   
 553 paid to a particular user address. It takes as input the recipient address key pair  
 554  $\{pub : Addr_{pk}, priv : Addr_{sk}\}$  and the current ledger  $L$  and outputs the set of  
 555 (unspent) received  $Notes$ .

#### Note

In the rest of this document, we will refer to a *Zeth user*  $\mathcal{U}_Z$  as a person, modeled as an object, holding one **Zeth** address (object attribute), and thus holding a *private address*,  $Addr_{sk}$ . We denote by  $\mathcal{U}_Z.Addr$  the **Zeth** address of  $\mathcal{U}_Z$  derived from  $Addr_{sk}$ , and which allows  $\mathcal{U}_Z$  to be the recipient of payments via **Zeth**, and to send funds via **Zeth**. Importantly, *not all Ethereum users are Zeth users, and vice-versa!*

556

Field	Description
$apk$	The <i>paying key</i>
$pkenc$	The <i>transmission key</i>

Table 1.4: “Payment address”,  $Addr_{pk}$ , of a DAP address

Field	Description
<i>ask</i>	The <i>spending key</i>
<i>skenc</i>	The <i>receiving key</i>

Table 1.5: “Private address”,  $Addr_{sk}$ , of a DAP address

557 **Zeth** leverages zk-snarks (Section 1.3) and the possibility to deploy smart-contracts  
 558 to specify privacy-preserving state transitions altering the **Ethereum** state  $\varsigma$  (Section 1.2).  
 559 As such, **Zeth** defines a smart-contract,  $\widetilde{\mathbf{Mixer}}$ , that keeps track of the set of *ZethNotes*  
 560 (Section 2.1) in a committed form, stored in a Merkle tree; and which verifies the va-  
 561 lidity of the state transitions generated by the **Zeth** users. As such a **Zeth** DAP is  
 562 entirely determined by  $\widetilde{\mathbf{Mixer}}$ , the instance of the mixer smart-contract deployed on the  
 563 **Ethereum** ledger. State transitions are executed on-chain by calling the **Mix** function of  
 564  $\widetilde{\mathbf{Mixer}}$ , which implements the algorithm **VerifyTx** of DAP, and which modifies  $\varsigma$  iff the  
 565 transaction is deemed valid.

#### Note

We denote by  $Mix_{in}$  the inputs taken by the **Mix** function defined on  $\widetilde{\mathbf{Mixer}}$ . Let  $zdata$  be the value of the *data* field of an **Ethereum** transaction such that:

$$zdata = \text{FS}(\text{Mix}) \parallel Mix_{in}$$

Then, we define  $tx_{\text{Mix}}$  as being the **Ethereum** transaction object returned by **SendTx** such that:

$$tx_{\text{Mix}}.to = \widetilde{\mathbf{Mixer}}.Addr \wedge tx_{\text{Mix}}.data = zdata$$

Importantly, when it is clear from context, we will omit the function selector from the definition of  $zdata$ , and only assume that  $zdata = Mix_{in}$ .

566

## 567 1.5 Security Assumptions

**Definition 5** (Discrete Log (DLog) based on Bellare and Shoup [BS07]). *Let  $\mathbb{G}$  denote a group whose order  $p$  is prime and written over  $\lambda$  bits. We let  $\log_{\mathbf{g}}(h)$  denote the discrete logarithm of  $h$  in the basis  $\mathbf{g}$ . We assume  $\mathbb{G}, p$  are fixed and known to all parties. Let us denote the advantage of a PPT adversary  $\mathcal{A}$  in attacking the discrete logarithm problem as follows,*

$$\text{Adv}_{\mathbb{G}, \mathcal{A}}^{\text{dlog}} = \Pr[\mathbf{g} \leftarrow \mathbb{G}^*, x \leftarrow \mathbb{F}_p, x' \leftarrow \mathcal{A}([1], [x]) : [x'] = [x]]$$

568 We say that the DLog is hard in  $\mathbb{G}$  if and only if  $\text{Adv}_{\mathcal{A}}^{\text{dlog}}(\lambda)$  is negligible.

**Definition 6** (One More Discrete Log (om-DLog) based on Paillier and Vergnaud [PV05]). Let  $\mathbb{G}$  denote a group whose order  $p$  is prime and written over  $\lambda$  bits. We let  $\log_{\mathbf{g}}(h)$  denote the discrete logarithm of  $h$  in the basis  $\mathbf{g}$ . A PPT adversary  $\mathcal{A}$  solving the om-DLog is given  $q + 1$  random group elements as well as a limited access to a discrete logarithm oracle  $\mathcal{O}^{\text{DLog}_{\mathbf{g}}}(q)$ .  $\mathcal{A}$  is allowed to query this oracle at most  $q$  times, thus obtaining the discrete logarithm of  $q$  group elements of his choice with respect to a fixed base  $\mathbf{g}$ . Eventually,  $\mathcal{A}$  must output the  $q + 1$  discrete logarithms. Let us denote the advantage of a PPT adversary  $\mathcal{A}$  in attacking the one more discrete logarithm problem as follows,

$$\text{Adv}_{\mathcal{A}}^{\text{om-dlog}}(\lambda) = \Pr \left[ \begin{array}{l} \mathbf{g} \leftarrow_{\$} \mathbb{G}^*, \{ \llbracket r_i \rrbracket \}_{i \in [q+1]} \leftarrow_{\$} \mathbb{G}^{q+1}, \\ \{ r'_i \}_{i \in [q+1]} \leftarrow \mathcal{A}^{\mathcal{O}^{\text{DLog}_{\mathbf{g}}}(q)}(\llbracket 1 \rrbracket, \{ \llbracket r_i \rrbracket \}_{i \in [q+1]}) : \\ \forall i \leq q + 1, r'_i = \log_{\mathbf{g}}(\llbracket r_i \rrbracket) \end{array} \right]$$

569 We say that the om-DLog is hard in  $\mathbb{G}$  if and only if  $\text{Adv}_{\mathcal{A}}^{\text{om-dlog}}(\lambda)$  is negligible.

## 570 1.6 Definitions

### 571 1.6.1 Negligible function

572 **Definition 7** (Negligible function, [KL14, Definition 3.4]). A function  $f$  from  $\mathbb{N}$  to  $\mathbb{R}^+$   
573 (positive real numbers) is negligible if for every positive polynomial  $p$  there exists  $N$  such  
574 that for all integers  $n > N$  it holds that  $f(n) < \frac{1}{p(n)}$ .

### 575 1.6.2 Basic algebra notions

576 **Definition 8** (Group, see [Bou03, Section I.4]). A group is given by a tuple  $(\mathbb{G}, \otimes)$ ,  
577 where  $\mathbb{G}$  is a set and  $\otimes$  is a binary operation in  $\mathbb{G}$ , i.e.  $\otimes : \mathbb{G} \times \mathbb{G} \rightarrow \mathbb{G}$ , with the  
578 following properties:

- 579 •  $(\mathbf{g} \otimes \mathbf{h}) \otimes \mathbf{k} = \mathbf{g} \otimes (\mathbf{h} \otimes \mathbf{k})$  (associativity)
- 580 • There exists an element  $\epsilon \in \mathbb{G}$  s.t. for each  $\mathbf{g} \in \mathbb{G}$ ,  $\mathbf{g} \otimes \epsilon = \epsilon \otimes \mathbf{g} = \mathbf{g}$  (identity  
581 element).
- 582 • For each  $\mathbf{g} \in \mathbb{G}$  there exist  $\mathbf{h} \in \mathbb{G}$  s.t.  $\mathbf{g} \otimes \mathbf{h} = \mathbf{h} \otimes \mathbf{g} = \epsilon$  (inverse element).

For simplicity, we may also use the additive notation for groups:  $\otimes$  is denoted as  $+$ , the identity element as  $\mathbf{o}$  and the inverse element of  $\mathbf{g}$  as  $-\mathbf{g}$ . Given  $\mathbf{g} \in \mathbb{G}$  and  $x \in \mathbb{Z}$ , we have that:

$$x \cdot \mathbf{g} = \begin{cases} \mathbf{o} & \text{if } x = 0 \\ \mathbf{g} + \dots + \mathbf{g}, (x \text{ times}) & \text{if } x > 0. \\ -\mathbf{g} + \dots + (-\mathbf{g}), (x \text{ times}) & \text{if } x < 0 \end{cases}$$

583 **Definition 9** (Finite Cyclic Group, adapted from [KL14, Sections 7.1.3, 7.3.2]). A  
584 finite cyclic group is given by a tuple  $(q, \mathbb{G}, \mathbf{g}, \otimes)$ , called the group description, where  $\mathbb{G}$

585 represents the set of group elements,  $\mathbf{g}$  is a generator and  $q$  is the order. The generator  
 586  $\mathbf{g}$  generates the group; namely, each  $h \in \mathbb{G}$  can be expressed by the generator as  $\mathbf{h} =$   
 587  $\mathbf{g} \otimes \dots \otimes \mathbf{g}$ . Given a scalar  $x$ , we denote by  $\llbracket x \rrbracket$  the encoding of  $x$  in  $\mathbb{G}$ : i.e.  $\llbracket x \rrbracket = \mathbf{g} \otimes \dots \otimes \mathbf{g}$   
 588 ( $x$  times). As consequence,  $\llbracket 1 \rrbracket = \mathbf{g}$ .

589 For theoretical purposes, we introduce the **SetupG** algorithm that for a given security  
 590 parameter  $\lambda$  outputs a cyclic group, formally:

591 **Definition 10** (Group Setup Algorithm, taken from [KL14, Sections 7.1.3, 7.3.2]). A  
 592 group setup algorithm **SetupG** is a PPT algorithm which takes as input a security pa-  
 593 rameter  $1^\lambda$  and outputs a group description  $(q, \mathbb{G}, \mathbf{g}, \otimes)$ , where the binary representation  
 594 of  $q$  is given by  $\lambda$  bits and each group element can be represented by  $gLen(\lambda)$  bits. Note  
 595 that  $gLen$  is  $\text{poly}(\lambda)$ .<sup>4</sup>

### 596 1.6.3 Symmetric Encryption

597 **Definition 11** (Symmetric Encryption, [KL14, Definition 3.8]). A symmetric encryption  
 598 scheme **Sym** is given by a tuple of PPT algorithms  $(\text{KGen}, \text{Enc}, \text{Dec})$  where:

- 599 • **KGen**, the key generation algorithm, takes a security parameter  $1^\lambda$  and outputs a  
 600 secret key  $ek$ ; we assume, without loss of generality, that  $kLen(\lambda) = \text{length}(ek) \geq \lambda$ .  
 601 Note that  $kLen(\lambda)$  is a polynomial function in  $\lambda$ .<sup>5</sup>
- 602 • **Enc**, the encryption algorithm, takes a key  $ek$ , a plaintext  $m \in \{0, 1\}^*$  and returns  
 603 a ciphertext  $ct$ .
- 604 • **Dec**, the decryption algorithm, takes a key  $ek$  and a ciphertext  $ct$ , and returns a  
 605 message  $m$ . We assume, without loss of generality, that **Dec** is deterministic.

606 For every security parameter  $\lambda$ , key  $ek$  output by  $\text{KGen}(1^\lambda)$ , and message  $m \in \{0, 1\}^*$ ,  
 607 it holds that  $\text{Dec}(ek, \text{Enc}(ek, m)) = m$  (correctness property).

608 If  $(\text{KGen}, \text{Enc}, \text{Dec})$  is such that for key  $ek$  output by  $\text{KGen}(1^\lambda)$ , algorithm  $\text{Enc}(ek, \cdot)$   
 609 is only defined for messages  $m \in \{0, 1\}^{l(\lambda)}$ , then we say that  $(\text{KGen}, \text{Enc}, \text{Dec})$  is a *fixed-*  
 610 *length symmetric encryption scheme* with length parameter  $l(\lambda)$  ( $l$  is  $\text{poly}(\lambda)$ ). A security  
 611 notion for **Sym** follows:

**Definition 12** (IND-CPA). Let **Sym** be a symmetric encryption scheme and let  $\mathcal{A}$  be an  
 adversary. Consider the IND-CPA game described in Figure 1.2. We define the IND-CPA  
 advantage of  $\mathcal{A}$  as follows:

$$\text{Adv}_{\text{Sym}, \mathcal{A}}^{\text{ind-cpa}}(\lambda) = |2 \cdot \Pr[\text{IND-CPA}(\lambda) = 1] - 1|.$$

612 **Sym** is said to be IND-CPA secure if, for every PPT adversary  $\mathcal{A}$ , the advantage  $\text{Adv}_{\text{Sym}, \mathcal{A}}^{\text{ind-cpa}}(\lambda)$   
 613 is a negligible function.

<sup>4</sup>For simplicity we may denote  $gLen(\lambda)$  as  $gLen$ .

<sup>5</sup>For simplicity we may denote  $kLen(\lambda)$  as  $kLen$ .

### IND-CPA( $\lambda$ )

---

```

 $ek \leftarrow \text{KGen}(1^\lambda)$ 
 $(m_0, m_1, \text{state}) \leftarrow \mathcal{A}^{\text{O}^{\text{Enc}_{ek}}} \text{ with } \text{length}(m_0) = \text{length}(m_1)$ 
 $b \leftarrow_{\$} \{0, 1\}$ 
 $ct \leftarrow \text{Enc}(ek, m_b)$ 
 $\tilde{b} \leftarrow \mathcal{A}^{\text{O}^{\text{Enc}_{ek}}}(ct, \text{state})$ 
return  $\tilde{b} = b$ 

```

Figure 1.2: IND-CPA game for Sym.

## 1.6.4 Asymmetric Encryption

**Definition 13** (Asymmetric encryption, [KL14, Definition 10.1]). *An asymmetric encryption scheme  $\text{Asym}$  is given by a tuple of PPT algorithms  $(\text{KGen}, \text{Enc}, \text{Dec})$  where:*

- *$\text{KGen}$ , the key generation algorithm, takes a security parameter  $1^\lambda$  and returns a pair of keys  $(sk, pk)$ . We refer to the first of these as private key and the second as public key. We assume for convenience that  $pk$  and  $sk$  each have length at least  $\lambda$ , and that  $\lambda$  can be determined from  $pk, sk$ ;*
- *$\text{Enc}$ , the encryption algorithm, takes a public key  $pk$ , a plaintext  $m$ , from some underlying plaintext space (that may depend on  $pk$ ) and returns a ciphertext  $ct$ ;*
- *$\text{Dec}$ , the decryption algorithm, takes a private key  $sk$  and a ciphertext  $ct$ , and returns a message  $m$  or a special symbol  $\perp$  denoting the decryption failure. We assume, without loss of generality, that  $\text{Dec}$  is deterministic.*

*We require that for all  $(sk, pk)$  returned by  $\text{KGen}$ , and every message  $m$  in the appropriate underlying plaintext space, it holds that  $\text{Dec}(sk, \text{Enc}(pk, m)) = m$  (correctness property).*

Secure communication usually requires ciphertext indistinguishability (e.g. IND-CCA2 [ABR99, Definition 8]). In **Zeth**, however, the key privacy property IK-CCA [BBDP01] is also required: it ensures indistinguishability of the key under which an encryption is performed.

**Definition 14** (IK-CCA). *Let  $\text{Asym} = (\text{KGen}, \text{Enc}, \text{Dec})$  be an asymmetric encryption scheme and let  $\mathcal{A}$  be an adversary. Given the IK-CCA game described in Figure 1.3, with the condition that  $\mathcal{A}$  cannot query  $\text{O}^{\text{Dec}_{sk_0}}$  or  $\text{O}^{\text{Dec}_{sk_1}}$  on the challenge ciphertext  $ct^6$ , we define the IK-CCA advantage of  $\mathcal{A}$  as follows:*

$$\text{Adv}_{\text{Asym}, \mathcal{A}}^{\text{ik-cca}}(\lambda) = |2 \cdot \Pr[\text{IK-CCA}(\lambda) = 1] - 1|$$

*We say that  $\text{Asym}$  is IK-CCA secure if for every PPT adversary  $\mathcal{A}$  the advantage  $\text{Adv}_{\text{Asym}, \mathcal{A}}^{\text{ik-cca}}(\lambda)$  is a negligible function.*

---

<sup>6</sup>state is some state information that the adversary outputs after the choice of the message to encrypt. It can be some preprocessed information that can be helpful to win the game

---

IK-CCA( $\lambda$ )

$(sk_0, pk_0), (sk_1, pk_1) \leftarrow \text{KGen}(1^\lambda)$   
 $(m, state) \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(pk_0, pk_1)$   
 $b \leftarrow \{0, 1\}$   
 $ct \leftarrow \text{Enc}(pk_b, m)$   
 $\tilde{b} \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(ct, state)$   
**return**  $\tilde{b} = b$

Figure 1.3: IK-CCA game.

### 1.6.5 (Block-cipher-based) Compression functions

**Definition 15.** Let  $kl, il > 1$ . A block cipher is a map  $E: \{0, 1\}^{kl} \times \{0, 1\}^{il} \rightarrow \{0, 1\}^{il}$  where, for each key  $k \in \{0, 1\}^{kl}$ , the function  $E_k(\cdot) = E(k, \cdot)$  is a permutation on  $\{0, 1\}^{il}$ . If  $E$  is a block cipher then  $E^{-1}$  is its inverse, that on input  $(k, y)$  returns  $m$  such that  $E_k(m) = y$ .

Let  $\mathcal{BCK}(kl, il)$  be the set of all block ciphers  $E: \{0, 1\}^{kl} \times \{0, 1\}^{il} \rightarrow \{0, 1\}^{il}$ . In order to analyze the security properties of block-cipher based cryptographic constructions it is common to use a security model denoted *the ideal cipher model (ICM)*. Informally speaking, in ICM attackers are allowed to query an oracle simulating a random block cipher and they have no information on the internal structure. We formalize this notion in the following definition:

**Definition 16** (Ideal Cipher Model [HKT11]). *The Ideal Cipher Model (ICM), is a security model where all parties are granted access to an ideal cipher  $E: \{0, 1\}^{kl} \times \{0, 1\}^{il} \rightarrow \{0, 1\}^{il}$ , a random primitive such that the restrictions  $E(k, \cdot)$  for  $k \in \{0, 1\}^{kl}$  are  $2^{kl}$  independent random permutations.*

For fixed  $kl$  and  $il$ , each party is given access to the oracles  $\mathcal{O}^E$  and  $\mathcal{O}^{E^{-1}}$ , simulating  $E$  and  $E^{-1}$ , which can be queried for encryption and decryption a polynomial number of times. The encryption oracle takes as input a key,  $k \in \{0, 1\}^{kl}$ , and a preimage,  $m \in \{0, 1\}^{il}$ , and returns a tuple comprising the image,  $y \in \{0, 1\}^{il}$ , along with the inputs,  $k$  and  $m$ . If  $(k, m)$  is queried for the first time, the image,  $y$ , is taken uniformly at random and added to the oracle's table. Otherwise, the oracle returns the  $y$  associated with query  $(k, m)$  in its table. The decryption oracle is defined similarly with the image and key defined as inputs and the preimage chosen randomly (see: Fig. 1.4).

**Definition 17** (Block-cipher based compression function [BRS02]). *A block cipher-based compression function is a map  $f$  such that*

$$f: \mathcal{BCK}(kl, il) \times \{0, 1\}^a \times \{0, 1\}^b \rightarrow \{0, 1\}^c$$

where  $kl, il, a, b, c > 1$  and  $a + b > c$ . The function  $f$ , given  $m \in (\{0, 1\}^a \times \{0, 1\}^b)$ , computes  $f(E, m)$  using an  $E$ -oracle.

$O^E(k, m)$	$O^{E^{-1}}(k, y)$
<b>if</b> $(k, m, \cdot) \notin \text{Table}_O$	<b>if</b> $(k, \cdot, y) \notin \text{Table}_O$
$y \leftarrow_{\$} \{0, 1\}^{\text{il}}$	$m \leftarrow_{\$} \{0, 1\}^{\text{il}}$
$\text{Table}_O.\text{append}(k, m, y)$	$\text{Table}_O.\text{append}(k, m, y)$
<b>else</b> $y = \text{Table}_O(k, m)$	<b>else</b> $m = \text{Table}_O(k, y)$
<b>return</b> $(k, m, y)$	<b>return</b> $(k, m, y)$

Figure 1.4: Oracles of an ideal block cipher, with  $\text{Table}_O$  being a table of tuples (key, preimage, image) of queries already answered by the oracle.

**Remark 5.** We use the notation  $f_E$  if a compression function  $f$  is defined over a given block-cipher  $E$ , i.e.  $f_E : \{0, 1\}^a \times \{0, 1\}^b \rightarrow \{0, 1\}^c$  and  $f_E = f(E, \cdot)$ , for  $a, b, c$  as given in the definition above.

Let  $f$  be a compression function based on a block-cipher. Fix a constant  $h_0 \in \{0, 1\}^c$  and an adversary  $\mathcal{A}$ . We define the advantage in finding a collision in  $f$  as the real number

$$\text{Adv}_{f, \mathcal{A}}^{\text{coll}} = \Pr \left[ E \leftarrow_{\$} \mathcal{BK}(\text{kl}, \text{il}); ((k, m), (k', m')) \leftarrow \mathcal{A}^{O^E, O^{E^{-1}}}(f_E, h_0) : \begin{aligned} &((k, m) \neq (k', m') \wedge f_E(k, m) = f_E(k', m')) \vee f_E(k, m) = h_0 \end{aligned} \right].$$

The previous definition gives credit for finding an  $(k, m)$  such that  $f_E(k, m) = h_0$  for a fixed  $h_0 \in \{0, 1\}^c$ .

### 1.6.6 Hash functions

**Definition 18** (Hash function, [KL14, Definition 4.9]). A hash function  $\mathcal{H}$  is a pair of algorithms  $(\text{Setup}, H)$  fulfilling the following properties:

- **Setup** is a PPT algorithm which takes as input a security parameter  $1^\lambda$  and outputs a key  $hk$ . We assume that  $1^\lambda$  is included in  $hk$ .
- **H** is (deterministic) polynomial-time algorithm that takes as input a key  $hk$  and any string  $x \in \{0, 1\}^*$ , and outputs a string  $H(hk, x) = H_{hk}(x) \in \{0, 1\}^{hLen}$ , where  $hLen$  is a polynomial function in  $\lambda$ .<sup>7</sup>

If for every  $\lambda$  and  $hk$ ,  $H_{hk}$  is defined only over inputs of length  $hInpLen(\lambda)$  and  $hInpLen(\lambda) > hLen(\lambda)$ , then we say that  $\mathcal{H}$  is a fixed-length hash function with length parameter  $hInpLen$ . Note that  $hInpLen(\lambda)$  is a polynomial function in  $\lambda$ .

Informally, for a given function  $f$  we say that  $(x, y)$  is a collision if  $f(x) = f(y)$  and  $x \neq y$ . In the following, we formalize this notion for a hash function  $\mathcal{H}$ .

<sup>7</sup>For simplicity we may denote  $hLen(\lambda)$  as  $hLen$ .



**Definition 19** (Collision Resistance [KL14, Definitions 4.10]). A hash function  $\mathcal{H} = (\text{Setup}, \text{H})$  is collision resistant if for all probabilistic polynomial-time adversaries  $\mathcal{A}$  there exists a negligible function  $\text{negl}(\lambda)$  such that:

$$\text{Adv}_{\mathcal{H}, \mathcal{A}}^{\text{cr}}(\lambda) = \Pr \left[ hk \leftarrow \text{Setup}(1^\lambda), (x, y) \leftarrow \mathcal{A}(hk) : x \neq y \wedge \text{H}_{hk}(x) = \text{H}_{hk}(y) \right]$$

## 676 HDHI and HDHI2 assumptions

677 The Hash Diffie-Hellmann Independence (HDHI) assumption states that, given  $\text{H}$  in  $\mathcal{H}$   
 678 and a group description  $(p, \mathbb{G}, \mathbf{g}, \otimes)$ , for  $\llbracket u \rrbracket$  and  $\llbracket v \rrbracket$ , with  $u, v$  sampled at random, it is  
 679 hard for an attacker to distinguish  $\text{H}(\llbracket u \rrbracket \parallel \llbracket uv \rrbracket)$  from a random string of the same size.<sup>8</sup>  
 680 This is formalized in Definition 20, where an attacker can also access an oracle  $\mathcal{O}^{\text{HDHI}_v}$   
 681 that on inputs  $x \in \mathbb{G}$  returns  $\text{H}(x \parallel v \cdot x)$  (queries on  $\llbracket u \rrbracket$  are forbidden).<sup>9</sup> In other words,  
 682 the HDHI assumption measures the sense in which  $\text{H}$  is “independent” of the underlying  
 683 Diffie-Hellman problem.

**Definition 20** (HDHI, [ABR99, Definition 7]). Let  $\mathcal{H}$  be a hash function,  $\text{SetupG}$  be a group generation algorithm and  $\mathcal{A}$  be an adversary. Consider the HDHI game described in Figure 1.5. We define the advantage of  $\mathcal{A}$  in violating the HDHI assumption as follows:

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi}}(\lambda) = |2 \cdot \Pr[\text{HDHI}(\lambda) = 1] - 1|.$$

684 Note that above definition corresponds to [ABR99, Section 3.2.1, Definition 3]. In the  
 685 following, we introduce a similar notion denoted as HDHI2 (this is an adaptation of ODH2  
 686 notion in [ABN10, Section 6]) and it will be useful in the IK-CCA proof Section 3.5.4.

**Definition 21** (HDHI2). Let  $\mathcal{H}$  be a hash function,  $\text{SetupG}$  a group generation algorithm and let  $\mathcal{A}$  be an adversary. Consider the HDHI2 game described in Figure 1.6. We define the advantage of  $\mathcal{A}$  in violating the HDHI2 assumption as follows:

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi2}}(\lambda) = |2 \cdot \Pr[\text{HDHI2}(\lambda) = 1] - 1|.$$

**Lemma 1.** Let  $\mathcal{A}$  be an adversary with advantage  $\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi2}}$  in solving the HDHI2 problem. Then there exists an adversary  $\mathcal{B}$  such that

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi2}}(\lambda) \leq 2 \cdot \text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{B}}^{\text{hdhi}}(\lambda).$$

687 *Proof.* We can reuse the proof described in [ABN10, Lemma 6.1] by applying mi-  
 688 nor modifications. In fact, HDHI and HDHI2 are, respectively, slightly different from  
 689 ODH and ODH2 notions: in the related security games, if  $b = 0$  the challenges are  
 690 constructed as  $\text{H}(\llbracket u \rrbracket \parallel \llbracket uv \rrbracket)$  and  $\{\text{H}(\llbracket u \rrbracket \parallel \llbracket uv_0 \rrbracket), \text{H}(\llbracket u \rrbracket \parallel \llbracket uv_1 \rrbracket)\}$  instead of  $\text{H}(\llbracket uv \rrbracket)$  and  
 691  $\{\text{H}(\llbracket uv_0 \rrbracket), \text{H}(\llbracket uv_1 \rrbracket)\}$ . By accordingly changing the instances of  $\text{H}$  in the games  $\mathcal{G}_0, \mathcal{G}_1, \mathcal{G}_2$   
 692 of [ABN10, Lemma 6.1] our lemma follows.  $\square$

<sup>8</sup>Note that  $\text{H}$  takes as inputs bit strings, so technically we should make use of an encoding function from  $\mathbb{G}$  to  $\{0, 1\}^{gLen}$  but we may omit this step through the document to improve readability.

<sup>9</sup>In [ABR99, Section 3.2.1] this notion is denoted as adaptive HDH independence assumption. Since we only introduce the adaptive version we denote it as HDHI.

---

$\text{HDHI}(\lambda)$   
 $hk \leftarrow \mathcal{H}.\text{Setup}(1^\lambda)$   
 $(q, \mathbb{G}, \mathfrak{g}, \otimes) \leftarrow \text{SetupG}(1^\lambda)$   
 $u, v \leftarrow_{\$} [q]$   
 $w_0 \leftarrow \mathcal{H}.H_{hk}(\llbracket u \rrbracket \parallel \llbracket uv \rrbracket)$   
 $w_1 \leftarrow_{\$} \{0, 1\}^{hLen}$   
 $b \leftarrow_{\$} \{0, 1\}$   
 $\tilde{b} \leftarrow \mathcal{A}^{\text{O}^{\text{HDHI}_v}}(\llbracket u \rrbracket, \llbracket v \rrbracket, w_b)$   
**return**  $\tilde{b} = b$

Figure 1.5: HDHI game.

---

$\text{HDHI2}(\lambda)$   
 $hk \leftarrow \mathcal{H}.\text{Setup}(1^\lambda)$   
 $(q, \mathbb{G}, \mathfrak{g}, \otimes) \leftarrow \text{SetupG}(1^\lambda)$   
 $u, v_0, v_1 \leftarrow_{\$} [q]$   
 $w_{0,0} \leftarrow \mathcal{H}.H_{hk}(\llbracket u \rrbracket \parallel \llbracket uv_0 \rrbracket), w_{0,1} \leftarrow \mathcal{H}.H_{hk}(\llbracket u \rrbracket \parallel \llbracket uv_1 \rrbracket)$   
 $w_{1,0} \leftarrow_{\$} \{0, 1\}^{hLen}, w_{1,1} \leftarrow_{\$} \{0, 1\}^{hLen}$   
 $b \leftarrow_{\$} \{0, 1\}$   
 $\tilde{b} \leftarrow \mathcal{A}^{\text{O}^{\text{HDHI}_{v_0}}, \text{O}^{\text{HDHI}_{v_1}}}(\llbracket u \rrbracket, \llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket, w_{b,0}, w_{b,1})$   
**return**  $\tilde{b} = b$

Figure 1.6: HDHI2 game.

### 1.6.7 Pseudo Random Functions

Informally speaking, a pseudorandom function family  $\mathcal{PRF} = \{\text{PRF}_k : D \rightarrow C\}_{k \in \mathcal{K}}$  is a collection of functions such that for randomly chosen  $k \in \mathcal{K}$ , the function  $\text{PRF}_k$  is indistinguishable from a random function that maps  $D$  to  $C$ .

**Definition 22** (PRF Family [KL14, Definition 3.24]). *Let  $\mathcal{F} : \{0, 1\}^* \times \{0, 1\}^* \rightarrow \{0, 1\}^*$  be an efficient, length-preserving, keyed function. We say  $\mathcal{F}$  is a pseudo random function if for all probabilistic polynomial-time distinguishers  $\text{Dist}$ , there exists a negligible function  $\text{negl}$  such that:*

$$\text{Adv}_{\mathcal{F}, \text{Dist}}^{\text{prf}}(\lambda) = \left| \Pr \left[ \text{Dist}^{\mathcal{F}_k(\cdot)}(1^\lambda) = 1 \right] - \Pr \left[ \text{Dist}^{f_\lambda(\cdot)}(1^\lambda) = 1 \right] \right| \leq \text{negl}(\lambda)$$

where  $k \leftarrow_{\$} \mathcal{K} = \{0, 1\}^\lambda$  is chosen uniformly at random and  $f_\lambda$  is chosen uniformly at random from the set of functions mapping  $\lambda$ -bit strings to  $\lambda$ -bit strings.

### 1.6.8 Commitment scheme

**Definition 23** (Non-interactive commitment scheme [BCC<sup>+</sup>15, Section 2.1]). *A non-interactive commitment scheme  $\text{ComSch}$  is defined by the following algorithms:*

- **Setup**, is a PPT algorithm that takes a security parameter  $1^\lambda$  and outputs public parameters  $pp$ .
- **Com**, is a polynomial-time algorithm that takes a message  $m \in \mathbb{B}^{\text{il}}$ , a random coin  $r \in \mathbb{B}^{\text{nl}}$  and outputs a commitment  $cm \in \mathbb{B}^{\text{ol}}$ .

We assume that  $pp$  is implicitly passed to **Com**.

**Definition 24** (Computationally Hiding). *We say that a commitment scheme is computationally hiding if for all PPT adversary  $\mathcal{A}$  the following:*

$$\left| \Pr \left[ \begin{array}{l} pp \leftarrow \text{Setup}(1^\lambda), (m_0, m_1) \leftarrow \mathcal{A}(pp), b \leftarrow_{\$} \{0, 1\}, \\ r \leftarrow_{\$} \mathbb{B}^{\text{nl}}, cm \leftarrow \text{Com}(m_b; r), \tilde{b} \leftarrow \mathcal{A}(cm), b = \tilde{b} \end{array} \right] - \frac{1}{2} \right|$$

711 *is at most negligible in  $\lambda$ .*

**Definition 25** (Computationally Binding). *We say that a commitment scheme is computationally binding if for all PPT adversary  $\mathcal{A}$  the following:*

$$\Pr \left[ \begin{array}{l} pp \leftarrow \text{Setup}(1^\lambda), (m_0, r_0, m_1, r_1) \leftarrow \mathcal{A}(pp) \\ m_0 \neq m_1 \wedge \text{Com}(m_0; r_0) = \text{Com}(m_1; r_1) \end{array} \right]$$

712 *is at most negligible in  $\lambda$ .*

713 Note that the previous definitions can be made *statistical* if we consider unbounded  
714 attackers  $\mathcal{A}$ .

### 715 1.6.9 Digital Signature

716 **Definition 26** (Digital signature [KL14, Definition 12.1]). *A digital signature scheme*  
717 *SigSch is defined by the tuple of functions  $\text{SigSch} = (\text{KGen}, \text{Sig}, \text{Vf})$ ,*

- 718 •  $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$ . *Key Generation randomized algorithm takes as input the*  
719 *security parameter  $1^\lambda$  and returns a signing key  $sk$  and verifying key  $vk$ .*
- 720 •  $\sigma \leftarrow \text{Sig}(sk, m)$ . *Given a signing key  $sk$  and a message  $m$ , the Sig algorithm*  
721 *computes and outputs a signature  $\sigma$ .*
- 722 •  $\{0, 1\} \leftarrow \text{Vf}(vk, m, \sigma)$ . *Given a verification key  $vk$ , a message  $m$  and a signature*  
723  *$\sigma$ , the Vf algorithm returns 1 if  $\sigma$  is a valid signature else 0.*

724 A signature scheme must satisfy the *correctness property* (i.e  $\text{Vf}(vk, m, \text{Sig}(sk, m)) =$   
725  $\text{true}$ , where  $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$ ) and *unforgeable* (i.e. it is intractable to produce a  
726 signature, without knowing the signing key  $sk$ , on a message that has not been signed  
727 yet). In addition to these properties, certain digital signature schemes have an additional  
728 property called “one-timeness”.

729 **Definition 27** (Unforgeability (UF-CMA) [KL14, Definition 12.2]). *A digital signa-*  
730 *ture scheme SigSch is UF-CMA if the probability for any PPT adversary  $\mathcal{A}$  to win the*  
731 *UF-CMA game depicted in Fig. 1.7 is negligible.*

732 **Definition 28** (Strong Unforgeability (SUF-CMA)). *A digital signature scheme SigSch*  
733 *is SUF-CMA if the probability for any PPT adversary  $\mathcal{A}$  to win the SUF-CMA game*  
734 *depicted in Fig. 1.8 is negligible.*

735 **Definition 29** (One-Time (OT) Signature [KL14, Definition 12.6]). *A one-time signa-*  
736 *ture scheme is a digital signature scheme that uses each key-pair at most once.*

737 **Remark 6.** *It is worth noting that users may use one-time signing keys to sign multiple*  
738 *messages. In this case no security guarantees can be ensured.*

### UF-CMA( $1^\lambda, t, q$ )

---

```

1:  $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$ 
2:  $state \leftarrow \mathcal{A}^{\text{O}^{\text{Sig}_{sk}}}(vk, \cdot)$ 
3: //  $state = \{(m_i, \sigma_i)\}_{i \in [q]}$  where  $m_i$  denotes
4: // the  $i$ th query made to  $\text{O}^{\text{Sig}_{sk}}$  and
5: //  $\sigma_i$  denotes the  $i$ th oracle answers
6:  $(m^*, \sigma^*) \leftarrow \mathcal{A}(state)$ 
7: return  $\text{Vf}(vk, m^*, \sigma^*) = 1$ 
8:  $\wedge m^* \notin \{m_i\}_{i \in [q]}$ 

```

Figure 1.7: UF-CMA game

### SUF-CMA( $1^\lambda, t, q$ )

---

```

1:  $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$ 
2:  $state \leftarrow \mathcal{A}^{\text{O}^{\text{Sig}_{sk}}}(vk, \cdot)$ 
3: //  $state = \{(m_i, \sigma_i)\}_{i \in [q]}$  where  $m_i$  denotes
4: // the  $i$ th query made to  $\text{O}^{\text{Sig}_{sk}}$  and
5: //  $\sigma_i$  denotes the  $i$ th oracle answers
6:  $(m^*, \sigma^*) \leftarrow \mathcal{A}(state)$ 
7: return  $\text{Vf}(vk, m^*, \sigma^*) = 1$ 
8:  $\wedge (m^*, \sigma^*) \notin \{(m_i, \sigma_i)\}_{i \in [q]}$ 

```

Figure 1.8: SUF-CMA game

## 1.6.10 Message Authentication Code

A message authentication code is a scheme that enables users to tag data for the purpose of authenticity and integrity. Formally:

**Definition 30** (Message Authentication Code, [KL14, Definition 4.1]). *A message authentication code MAC is given by a tuple of PPT algorithms  $(\text{KGen}, \text{Tag}, \text{Vf})$  where:*

- **KGen**, the key generation algorithm, takes a security parameter  $1^\lambda$ , and returns a key  $mk \in \{0, 1\}^{mLen(\lambda)}$ .<sup>10</sup>
- **Tag**, the tag generation algorithm, takes a key  $mk$  and a message  $y \in \{0, 1\}^*$  and returns a string  $\tau \in \{0, 1\}^*$ , called tag.
- **Vf**, the tag verification algorithm, takes a key  $mk$ , a message  $y \in \{0, 1\}^*$  and a tag  $\tau \in \{0, 1\}^*$ . It returns a value in  $\{0, 1\}$  where: 0 denotes that the message was rejected (i.e. is deemed unauthentic) and 1 denotes that the message was accepted (i.e. is deemed authentic).

We require that for all  $mk \in \{0, 1\}^\lambda$  and  $y \in \{0, 1\}^*$  we have  $\text{Vf}(mk, y, \text{Tag}(mk, y)) = 1$ . If  $\text{Tag}(mk, \cdot)$  is defined only over messages of length  $l(\lambda)$  and  $\text{Vf}(mk, y, \tau)$  outputs 0 for every  $y$  that is not of length  $l(\lambda)$ , then we say that  $(\text{KGen}, \text{Tag}, \text{Vf})$  is a fixed-length MAC with length parameter  $l(\lambda)$ .

A security notion for MAC follows:

**Definition 31** (SUF-CMA, [ABR99, Section 3.2.3]). *Let  $\text{MAC} = (\text{KGen}, \text{Tag}, \text{Vf})$  be a message authentication scheme and let  $\mathcal{A}$  be an adversary. Consider the SUF-CMA game described in Figure 1.9, with the condition that  $\text{Tag}(mk, \bar{y}) \neq \bar{\tau}$ . We say adversary  $\mathcal{A}$  has forged when it outputs a pair  $(\bar{y}, \bar{\tau})$  such that  $\text{Vf}_k(\bar{y}, \bar{\tau}) = 1$  and  $(\bar{y}, \bar{\tau})$  was not previously obtained via a query to the tag oracle.*

<sup>10</sup>For simplicity we may denote  $mLen(\lambda)$  as  $mLen$ .

```

SUF-CMA ( $\lambda$ )
-----
 $mk \leftarrow \text{KGen}(1^\lambda)$ 
 $(\bar{y}, \bar{\tau}) \leftarrow \mathcal{A}^{\text{O}^{\text{Tag}_{mk}}, \text{O}^{\text{Vf}_{mk}}}$ 
return  $\text{Vf}(mk, \bar{y}, \bar{\tau}) = 1$ 

```

Figure 1.9: SUF-CMA game.

We define the SUF-CMA advantage of  $\mathcal{A}$  as follows:

$$\text{Adv}_{\text{MAC}, \mathcal{A}}^{\text{suf-cma}}(\lambda) = \Pr[\text{SUF-CMA}(\lambda) = 1]$$

762 We say that MAC is SUF-CMA secure if for every PPT adversary  $\mathcal{A}$  the advantage  
763  $\text{Adv}_{\text{MAC}, \mathcal{A}}^{\text{suf-cma}}(\lambda)$  is a negligible function.

## Chapter 2

# Zeth protocol

In this section, we detail the **Zeth** protocol and provide a set of requirements that need to be respected to guarantee the security of the protocol.

## 2.1 Zeth Data Types

We begin by describing, and giving intuition about, the data types (see: Section 1.1) used in **Zeth**. We follow some design rationale from **ZeroCash** [BSCG<sup>+</sup>14], and **Zcash** [ZCa19] in order to prevent the transaction malleability attack, and the Faerie Gold attack [ZCa19, Section 8.4]. We refer the reader to Appendix A for more details.

**ZethNoteDType** Represents a note in **Zeth**. It comprises the note’s owner public address  $apk$ , identifier  $\rho$ , randomness  $r$  and value  $v$ .

Field	Description	Value
$apk$	Note owner’s paying key	$\mathbb{B}^{\text{PRFADDRROUTLEN}}$
$r$	Note randomness	$\mathbb{B}^{\text{RTRAPLEN}}$
$v$	Note value	$\mathbb{B}^{\text{ZVALUELEN}}$
$\rho$	Note identifier	$\mathbb{B}^{\text{PRFRHOOUTLEN}}$

Table 2.1: **ZethNoteDType** data type

**JSInputDType** Denotes a joinsplit input. It comprises the opening of a commitment  $cm$  which is in the set of leaves in the Merkle tree of **Mixer** (i.e. a *ZethNote*), its address  $mkaddr$  and authentication path  $mkpath$  on the contract’s Merkle tree as well as the spending key  $ask$  of the note holder and the note nullifier  $nf$ .

Field	Description	Value
$mkpath$	Merkle authentication path to the commitment corresponding to the <i>ZethNote</i> to spend	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}^{\text{MKDEPTH}}$
$mkaddr$	Commitment address in the Merkle tree	$\mathbb{B}^{\text{MKDEPTH}}$
$znote$	Zeth note object	<b>ZethNoteDType</b>
$cm$	Zeth note commitment	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}$
$ask$	Note owner's spending key	$\mathbb{B}^{\text{ASKLEN}}$
$nf$	Note nullifier	$\mathbb{B}^{\text{PRNFOUTLEN}}$

Table 2.2: JSInputDType data type

779 **PrimInputDType** Represents the primary inputs used to generate the zk-snark  $\pi$ . *prim*  
780 is a tuple defined as the current Merkle root *mkroot* of the Merkle tree maintained  
781 by **Mixer**, the input notes nullifiers  $nfs = (nf_0, \dots, nf_{\text{JSIN}-1})$ , the output notes  
782 commitments  $cms = (cm_0, \dots, cm_{\text{JSOUT}-1})$ , the signature hash *hsig*, the message  
783 authentication tags  $htags = (h_0, \dots, h_{\text{JSIN}-1})$  and the residual bits field *rsd*, which  
784 aggregates the former's fields bits which could not be contained in a field element.

Field	Description	Value
$mkroot$	Merkle root of the Merkle tree	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}$
$nfs$	Indexed set of nullifiers of the “old” notes to spend	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}}^{\text{NFFLEN}})^{\text{JSIN}}$
$cms$	Indexed set of commitments to the newly created notes	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^{\text{JSOUT}}$
$hsig$	Signature hash (non-malleability, see: Appendix A)	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}^{\text{HSIGFLEN}}$
$htags$	Indexed set of message authentication tags (non-malleability, see: Appendix A)	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}}^{\text{HFLEN}})^{\text{JSIN}}$
$rsd$	Residual bits corresponding to unpacked bits of former fields	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}^{\text{RSDFLEN}}$

Table 2.3: PrimInputDType data type

785 **AuxInputDType** Represents the auxiliary inputs used to generate the zk-snark  $\pi$ . *aux* is  
786 a tuple defined as joinsplit inputs (i.e. “old outputs to be spent”), the new *ZethNotes*,  
787 the joinsplit's randomness  $\phi$  as well the public values *vin* and *vout*, the signature  
788 hash *hsig* and the message authentication tags  $htags = (h_0, \dots, h_{\text{JSIN}-1})$ .

Field	Description	Value
<i>jsins</i>	Indexed set of JSIN joinsplit inputs	$\text{JSInputDType}^{\text{JSIN}}$
<i>znotes</i>	Indexed set of JSOUT newly created notes	$\text{ZethNoteDType}^{\text{JSOUT}}$
$\phi$	The joinsplit randomness (non-malleability, see: Appendix A)	$\mathbb{B}^{\text{PHILEN}}$
<i>vin</i>	Public input value to the joinsplit	$\mathbb{B}^{\text{ZVALUELEN}}$
<i>vout</i>	Public output value to the joinsplit	$\mathbb{B}^{\text{ZVALUELEN}}$
<i>hsig</i>	Signature hash (non-malleability, see: Appendix A)	$\mathbb{B}^{\text{CRHHSIGOUTLEN}}$
<i>htags</i>	Indexed set of message authentication tags (non-malleability, see: Appendix A)	$(\mathbb{B}^{\text{PRFPKOUTLEN}})^{\text{JSIN}}$

Table 2.4: **AuxInputDType** data type

789 **MixInputDType** Represents the set of inputs to the Mix function of  $\widetilde{\text{Mixer}}$ . The input of  
790 the Mix function is a tuple defined as the primary inputs *prim*, the zk-proof  $\pi$ , the  
791 ciphertexts of the newly created notes *ciphers* =  $(ct_0, \dots, ct_{\text{JSOUT}-1})$ , a one-time  
792 signature  $\sigma$  and the associated verification key *vk*.

Field	Description	Value
<i>primIn</i>	Primary input object associated with the zk-proof $\pi$	<b>PrimInputDType</b>
<i>proof</i>	The zk-SNARK associated to the Zeth statement (see: Section 2.2)	<b>ZKPDType</b> (see: Section 3.6)
<i>otssig</i>	The one-time signature used to prevent transaction malleability (see: Appendix A)	<b>SigOtsDType</b> (see: Section 3.4.2)
<i>otsvk</i>	The verification key associated with the signature <i>otssig</i> used to prevent transaction malleability (see: Appendix A)	<b>VKOtsDType</b> (see: Section 3.4.2)
<i>ciphers</i>	Indexed set of ciphertexts of the newly generated notes	$(\mathbb{B}^{\text{ENCZETHNOTELEN}})^{\text{JSOUT}}$ (see: Section 3.5)

Table 2.5: **MixInputDType** data type

793 **MixEventDType** Represents the data emitted as an **Ethereum** event (Section 1.2.3) dur-  
794 ing a successful execution of the Mix function of  $\widetilde{\text{Mixer}}$ . Clients are required to  
795 read this data and use it to update their representation of  $\widetilde{\text{Mixer}}$ 's state.



Field	Description	Value
<i>mkroot</i>	New root of Merkle tree of commitments	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}$
<i>nfs</i>	Nullifiers for input notes consumed	$(\mathbb{B}^{\text{PRNFOUTLEN}})^{\text{JSIN}}$
<i>cms</i>	Commitments to the output notes	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^{\text{JSOUT}}$
<i>ciphers</i>	Ciphertexts for the output notes	$(\mathbb{B}^{\text{ENCZETHNOTELEN}})^{\text{JSOUT}}$

Table 2.6: `MixEventDType` data type

## 2.2 Zeth statement

As explained in [RZ19], the Mix function of  $\widetilde{\mathbf{Mixer}}$  verifies the validity of  $\pi$  on the given primary inputs in order to determine whether the state transition is valid. As such,  $\mathbf{Mixer}$  verifies whether for  $\pi$ , and primary input *prim*, there exists an auxiliary input *aux*, such that the tuple  $(\text{prim}, \text{aux})$  satisfies the NP-relation  $\mathbf{R}^z$ , consisting of the following constraints:

- For each  $i \in [\text{JSIN}]$ :

1.  $\text{aux.jsins}[i].\text{note.apk} = \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{addr}}(0)$
2.  $\text{aux.jsins}[i].\text{cm} = \text{ComSch.Com}(\text{aux.jsins}[i].\text{note.apk}, \text{aux.jsins}[i].\text{note}.\rho, \text{aux.jsins}[i].\text{note}.v; \text{aux.jsins}[i].\text{note}.r)$
3.  $\text{aux.jsins}[i].\text{nf} = \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{nf}}(\text{aux.jsins}[i].\text{note}.\rho)$
4.  $\text{aux.htags}[i] = \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{pk}}(i, \text{aux.hsigs})$  (malleability fix, see: Appendix A)
5.  $(\text{aux.jsins}[i].\text{note}.v) \cdot (1 - e) = 0$  is satisfied for the boolean value  $e$  set such that if  $\text{aux.jsins}[i].\text{note}.v > 0$  then  $e = 1$ .
6. The Merkle root  $\text{mkroot}'$  obtained after checking the Merkle authentication path  $\text{aux.jsins}[i].\text{mkpath}$  of commitment  $\text{aux.jsins}[i].\text{cm}$ , with MKHASH, equals to  $\text{prim.mkroot}$  if  $e = 1$ .
7.  $\text{prim.nfs}[i] = \{\text{Pack}_{\mathbb{F}_{\text{r}_{\text{BN}}}}(\text{aux.jsins}[i].\text{nf}[k \cdot \text{BNFIELD CAP}:(k+1) \cdot \text{BNFIELD CAP}])\}_{k \in [\lfloor \text{PRNFOUTLEN}/\text{BNFIELD CAP} \rfloor]}$
8.  $\text{prim.htags}[i] = \{\text{Pack}_{\mathbb{F}_{\text{r}_{\text{BN}}}}(\text{aux.htags}[i][k \cdot \text{BNFIELD CAP}:(k+1) \cdot \text{BNFIELD CAP}])\}_{k \in [\lfloor \text{PRFPKOUTLEN}/\text{BNFIELD CAP} \rfloor]}$

- For each  $j \in [\text{JSOUT}]$ :

1.  $\text{aux.znotes}[j].\rho = \text{PRF}_{\text{aux}.\phi}^{\text{rho}}(j, \text{aux.hsigs})$  (malleability fix, see: Appendix A)
2.  $\text{prim.cms}[j] = \text{ComSch.Com}(\text{aux.znotes}[j].\text{apk}, \text{aux.znotes}[j].\rho, \text{aux.znotes}[j].v; \text{aux.znotes}[j].r)$

- $\text{prim.hsigs} = \{\text{Pack}_{\mathbb{F}_{\text{r}_{\text{BN}}}}(\text{aux.hsigs}[k \cdot \text{BNFIELD CAP}:(k+1) \cdot \text{BNFIELD CAP}])\}_{k \in [\lfloor \text{CRHHSIGOUTLEN}/\text{BNFIELD CAP} \rfloor]}$

- $prim.rsd = \text{Pack}_{rsd}(\{aux.jsins[i].nf\}_{i \in [JSIN]}, aux.vin, aux.vout, aux.hsig, \{aux.htags[i]\}_{i \in [JSIN]})$
- Check that the “joinsplit is balanced”, i.e. check that the joinsplit equation holds:<sup>1</sup>

$$\begin{aligned} & \text{Pack}_{\mathbb{F}_{\text{rBN}}}(aux.vin) + \sum_{i \in [JSIN]} \text{Pack}_{\mathbb{F}_{\text{rBN}}}(aux.jsins[i].znote.v) \\ &= \sum_{j \in [JSOUT]} \text{Pack}_{\mathbb{F}_{\text{rBN}}}(aux.znotes[j].v) + \text{Pack}_{\mathbb{F}_{\text{rBN}}}(aux.vout) \end{aligned}$$

## 2.3 Generating the inputs of the Mix function ( $\text{Mix}_{in}$ )

In the following section, we assume that the system is initialized. In other words, we assume that a ledger  $L$  is available (i.e. an **Ethereum** network is operated by a set of miners), the **Mixer** contract is deployed on  $L$ . Likewise, we assume that the public parameters  $pp_{\text{ZkSnrkSch}} \leftarrow \text{ZkSnrkSch.KGen}(1^\lambda, \mathbf{R}^z)$  are available to **Mixer** and to all parties willing to call the Mix function of **Mixer**. Furthermore, we assume that there exists as set of **Ethereum** and **Zeth** users, and that the *payment address* of each **Zeth** user is easily discoverable. In the rest of this section, the set of *payment addresses* discovered by a zeth user  $\mathcal{U}_Z$  is represented as a list attribute  $\mathcal{U}_Z.keystore$  indexed by usernames.

In order for  $\mathcal{U}_Z$  to transact via **Zeth**,  $\mathcal{U}_Z$  needs to create an object  $\text{Mix}_{in}$  of type **MixInputDType** in order to call the Mix function of **Mixer**:

1. Create an object  $prim$  of type **PrimInputDType** to represent the primary input, and an object  $aux$  of type **AuxInputDType** to represent the auxiliary input, where:
  - (a)  $prim.mkroot \in \text{Roots}$ , where  $\text{Roots}$  is the set of *all* Merkle roots corresponding to one of the state of the Merkle tree on **Mixer** containing *all* the commitments to the input notes, in  $aux.jsins$ , in its set of leaves.
  - (b)  $aux.znotes[j].r \leftarrow_{\$} \mathbb{B}^{\text{RTRAPLEN}}, \forall j \in [JSOUT]$ , and  $aux.\phi \leftarrow_{\$} \mathbb{B}^{\text{PHILEN}}$
  - (c) The public values  $(aux.vin, aux.vout) \in (\mathbb{B}^{\text{ZVALUELEN}})^2$ ,  $aux.znotes[j].v$  and  $aux.znotes[j].apk \forall j \in [JSOUT]$  are all set by the sender,  $\mathcal{U}_Z$ , as desired as long as they satisfy the joinsplit equation.
  - (d) All attributes of the  $prim$  and  $aux$  objects should be derived as specified in the statement (see: Section 2.2) **alongside a signature hash ( $aux.hsig$ ) that is generated as the hash of the nullifiers and a one-time signing verification key (malleability fix, see: Appendix A), using the desired signature scheme  $\text{SigSch}_{\text{OT-SIG}}$  (see: Section 3.4):**

$$(sk_{\text{OT-SIG}}, vk_{\text{OT-SIG}}) = \text{SigSch}_{\text{OT-SIG}}.\text{KGen}(1^\lambda) \quad (2.1)$$

$$aux.hsig = \text{CRH}^{\text{hsig}}(\{aux.jsins[i].nf\}_{i \in [JSIN]}, vk_{\text{OT-SIG}}) \quad (2.2)$$

<sup>1</sup>where  $\text{Pack}_{\mathbb{F}_{\text{rBN}}}(x)$  outputs the numerical value of  $x$  in  $\mathbb{F}_{\text{rBN}}$ . We rely on the fact that  $\text{ZVALUELEN} < \text{BNFIELD CAP}$  to perform this sum.

843 (e)  $\text{Mix}_{in}.primIn \leftarrow prim$

#### Note

If one of the attributes of  $prim$  and  $aux$  is not correctly generated, then the proof of computational integrity generated in the next step will be rejected on **Mixer**, and the state of **Mixer** will not be modified.

844

845 2. Generate a zk-snark  $\pi$  to prove, in zero-knowledge, that the relation  $\mathbf{R}^z$  (Sec-  
846 tion 2.2) holds on the primary and auxiliary inputs, using the desired zk-snark  
847 scheme  $\text{ZkSnarkSch}$  (see: Section 3.6):

(a)

$$\pi \leftarrow \text{ZkSnarkSch.P}(pp_{\text{ZkSnarkSch}}, prim, aux)$$

848

(b)  $\text{Mix}_{in}.proof \leftarrow \pi$

849 3. Encrypt all the  $aux.znotes$  using the recipient's *payment address*, using the en-  
850 cryptation scheme  $\text{EncSch}$  (see: Section 3.5).

(a) For all  $j \in [\text{JSOUT}]$ , do:

$$ct_j = \text{EncSch.Enc}(aux.znotes[j], \mathcal{U}_Z.keystore[recipient_j].pub.pkenc)$$

851

(b)  $\text{Mix}_{in}.ciphers \leftarrow \{ct_j\}_{j \in [\text{JSOUT}]}$

852 4. Generate a signature  $\sigma_{\text{OT-SIG}}$  on the inputs of the  $\text{Mix}$  function, in order to prevent  
853 any malleability attacks (c.f. Appendix A), using the desired signature scheme  
854  $\text{SigSch}_{\text{OT-SIG}}$  (see: Section 3.4):

(a) The one-time signature keypair has already been generated in Eq. (2.1).

$$\begin{aligned} dataToBeSigned &= \mathcal{S}_E.Addr \parallel \text{Mix}_{in}.primIn \parallel \text{Mix}_{in}.\pi \parallel \text{Mix}_{in}.ciphers \\ \sigma_{\text{OT-SIG}} &= \text{SigSch}_{\text{OT-SIG}}.\text{Sig}(sk_{\text{OT-SIG}}, \text{CRH}^{\text{ots}}(dataToBeSigned)) \end{aligned}$$

855

(b)  $\text{Mix}_{in}.otssig \leftarrow \sigma_{\text{OT-SIG}}$

856

(c)  $\text{Mix}_{in}.otsvk \leftarrow vk_{\text{OT-SIG}}$

857 Here,  $\mathcal{S}_E.Addr$  represents the address of the Ethereum user  $\mathcal{S}_E$  who must sign the  
858 transaction (see: Section 2.4). In general, this is likely to be owned by the holder  
859  $\mathcal{U}_Z$  of the Zeth notes to be spent, but this is not a requirement.

## 860 2.4 Creating an Ethereum transaction $tx_{\text{Mix}}$ to call $\widetilde{\text{Mixer}}$

After generating a  $\text{Mix}_{in}$  object,  $\mathcal{U}_Z$  can generate an object  $tx_{raw}$  of type  $\text{TxRawDType}$ , such that:

$$tx_{raw}.to \leftarrow \widetilde{\text{Mixer}}.Addr \wedge tx_{raw}.data \leftarrow zdata$$

861 Then, an **Ethereum** user  $\mathcal{S}_E$  can ECDSA sign  $tx_{raw}$ , under  $\mathcal{S}_E.sk$  in order to transform  
862 this object of type  $\text{TxRawDType}$  into an finalized transaction, i.e. an object  $tx_{\text{Mix}}$  of type  
863  $\text{TxDType}$ .

864 Finally, the transaction  $tx_{\text{Mix}}$  is broadcasted on the **Ethereum** network and eventually  
865 gets mined.

### Note

Here, the Ethereum user  $\mathcal{S}_E$  who sends the final transaction, and the Zeth user  $\mathcal{U}_Z$  may represent the same person or entity, or not!

## 867 2.5 Processing $tx_{\text{Mix}}$

868 When a  $tx_{\text{Mix}}$  is mined (hence assuming that  $\text{EthVerifyTx}(tx_{\text{Mix}})$  returns true), the state  
869 transition specified by the **Mix** function of  $\widetilde{\text{Mixer}}$  is executed.

870 To preserve the soundness of **Zeth**, and make sure that no  $\mathcal{U}_Z$  is able to create  
871 value by double spending *ZethNotes*, various checks need to be satisfied. The function  
872 **ZethVerifyTx** is defined as the function that returns true if all the checks are satisfied,  
873 and false otherwise.

874 If  $\text{ZethVerifyTx}(tx_{\text{Mix}})$  returns true, then **Mix** modifies the “World state”  $\varsigma$  to account  
875 for the spent *ZethNotes* and the newly generated ones. However, if  $\text{ZethVerifyTx}(tx_{\text{Mix}})$   
876 returns false, then the state transition ends.

### Note

Even if  $\text{ZethVerifyTx}(tx_{\text{Mix}})$  returns false,  $\varsigma$  is modified since the **Ethereum** balances of the transaction originator is decremented by the sum of **DGAS** and the gas consumed by the **ZethVerifyTx** function, and the balance of the **Ethereum** account of the miner gets incremented by the same amount.

877 Thus, **Mix** proceeds as follows:  
878

1. Check that all the values of the primary inputs' ( $\text{Mix}_{in}.primIn$ ) entries are elements of the scalar field over which the zk-proof is generated:

$$\text{Mix}_{in}.primIn \in \mathbb{F}_{\text{rBN}}^*$$

2. Unpack the nullifiers, signature hash and public values:

$$\begin{aligned} nf_i &= \text{Unpack}_{nf}(\text{Mix}_{in}.primIn.nfs[i], \text{Mix}_{in}.primIn.rsd) \quad \forall i \in [\text{JSIN}] \\ vin &= \text{decode}_{\mathbb{N}}(\text{Unpack}_{vin}(), \text{Mix}_{in}.primIn.rsd) \\ vout &= \text{decode}_{\mathbb{N}}(\text{Unpack}_{vout}(), \text{Mix}_{in}.primIn.rsd) \\ hsig &= \text{Unpack}_{hsig}(\text{Mix}_{in}.primIn.hsig, \text{Mix}_{in}.primIn.rsd) \end{aligned}$$

879 3. Check the validity of the  $tx_{\text{Mix}}$  object ( $\text{ZethVerifyTx}$ ):

- Check that  $\text{Mix}_{in}.primIn.hsig$  is correctly computed, i.e. check that the following equation holds (to prevent transaction malleability, see: Appendix A):

$$hsig = \text{CRH}^{\text{hsig}}(\text{Mix}_{in}.primIn.nfs, \text{Mix}_{in}.otsvk)$$

- Check that  $\pi$  is a valid zk-snark for  $\text{Mix}_{in}.primIn$ , i.e. check that:

$$\text{ZkSnarkSch.V}(pp_{\text{ZkSnarkSch}}, \pi, \text{Mix}_{in}.primIn) = \text{true}$$

- Check that none of the nullifiers in  $\text{Mix}_{in}.primIn.nfs$  have already been used, i.e. check that:

$$nf_i \notin \text{Nulls}, \forall i \in [\text{JSIN}]$$

880 where  $\text{Nulls}$  is the set of all nullifiers that are “declared” on  $\widetilde{\text{Mixer}}$ .

- Check that  $\text{Mix}_{in}.otssig$  is a valid signature of the Ethereum sender’s address  $\text{Addr}$  (see: Section 2.4) and the attributes of  $\text{Mix}_{in}$ , to prevent transaction malleability (see: Appendix A), i.e. check that:

$$\begin{aligned} \text{SigSch}_{\text{OT-SIG}}.Vf(\text{Mix}_{in}.otsvk, m, \text{Mix}_{in}.otssig) &= \text{true} \\ \text{where } m &= \text{CRH}^{\text{ots}}(\text{Addr} \parallel \text{Mix}_{in}.primIn \parallel \text{Mix}_{in}.\pi \parallel \text{Mix}_{in}.ciphers) \end{aligned}$$

- Check that  $\widetilde{\text{Mix}_{in}.primIn.mkroot}$  corresponds to a valid state of the Merkle tree held on  $\widetilde{\text{Mixer}}$ , i.e. check that:

$$\text{Mix}_{in}.primIn.mkroot \in \text{Roots}'$$

881 where  $\text{Roots}'$  is the set of all Merkle roots corresponding to one of the state  
882 of the Merkle tree.

- Check that  $vin$  corresponds to the value  $val$  of the transaction object, i.e. check that:

$$vin = tx_{\text{Mix}}.val$$

883 4. If all checks above pass, i.e. if  $\text{ZethVerifyTx}(tx_{\text{Mix}})$  returns **true**, then the following  
884 additional modifications are made in  $\varsigma$ :

- 885 • Add the commitments  $\text{Mix}_{in}.primIn.cms$  to the Merkle tree held on  $\widetilde{\text{Mixer}}$ .

- 886 •  $Roots' \leftarrow Roots' \cup \{mkroot'\}$ , where  $mkroot'$  is the Merkle root of the Merkle  
887 tree after insertion of the commitments  $Mix_{in}.primIn.cms$  in the Merkle tree.
- 888 •  $Nulls \leftarrow Nulls \cup \{nf_i\}_{i \in [JSIN]}$ , i.e. the nullifiers  $nfs$  become “declared”.
- 889 • Modify the **Ethereum** balances according to the public values:
  - 890 –  $\varsigma[SE.Addr].bal = \varsigma[SE.Addr].bal - vin$
  - 891 –  $\varsigma[SE.Addr].bal = \varsigma[SE.Addr].bal + vout$
  - 892 –  $\widetilde{Mixer}.bal = \widetilde{Mixer}.bal + vin$
  - 893 –  $\widetilde{Mixer}.bal = \widetilde{Mixer}.bal - vout$
- 894 • Emit an event (Section 1.2.3)  $evMixOut$  of type **MixEventDType**, contain-  
895 ing the new root  $mkroot'$  of the Merkle tree of commitments, the nullifiers  
896  $\{nf_i\}_{i \in [JSIN]}$ , commitments to the newly created *ZethNotes*  $Mix_{in}.primIn.cms$ ,  
897 and the corresponding ciphertexts  $Mix_{in}.primIn.ciphers$ .

## 898 2.6 Receiving *ZethNotes*

899 In order to confirm the reception of *ZethNotes*,  $\mathcal{R}_Z$  must listen to the events (Sec-  
900 tion 1.2.3) of type **MixEventDType** emitted by the processing of  $tx_{Mix}$ , and try to decrypt  
901 the ciphertexts using  $\mathcal{R}_Z.priv.skenc$ . If the decryption is successful,  $\mathcal{R}_Z$  must verify that  
902 the *ZethNote* recovered is the opening of a commitment in the Merkle tree of  $\widetilde{Mixer}$ . If  
903 not,  $\mathcal{R}_Z$  rejects the payment.

904 For each event  $evMixOut \in \mathbf{MixEventDType}$  emitted by  $\widetilde{Mixer}$ , the steps to be  
905 followed are detailed below:

- 906 1. Compute the new root  $mkroot'$  of the Merkle tree of commitments, after adding the  
907 new values  $evMixOut.cms$ . If this value does not match the new root  $evMixOut.mkroot$   
908 emitted by  $\widetilde{Mixer}$ , abort.
2. Try to decrypt the ciphertexts:

$$zn_j = \text{EncSch.Dec}(evMixOut.ciphers[j], \mathcal{R}_Z.priv.skenc)$$

- 909 3. For each successful decryption, let  $j$  be the index of the decrypted ciphertext:
  - 910 (a) Check whether the recovered plaintext  $zn_j$  is a well-formed *ZethNote*. Abort  
911 if it is not well-formed.
  - (b) Check that the recovered *ZethNote*  $zn_j$  is the opening of the corresponding  
commitment  $evMixOut.cms[j]$ :

$$evMixOut.cms[j] = \text{ComSch.Com}(zn_j.apk, zn_j.\rho, zn_j.v; zn_j.r)$$

912 Abort if the note is not a valid opening.

913 (c) Additionally, if sender  $\mathcal{S}_Z$ , and recipient  $\mathcal{R}_Z$  had a contractual agreement,  
 914 then  $\mathcal{R}_Z$  needs to check that the terms of this agreement are fulfilled by all  
 915 the recovered *ZethNotes*, abort otherwise.

916 Note that steps 1 and 3b are required to ensure that data decrypted by  $\mathcal{R}_Z$  exactly  
 917 matches the data committed to in **Mixer**. In particular, step 1 requires  $\mathcal{R}_Z$  to maintain  
 918 or have access to some representation of the Merkle tree of commitments. See Sec-  
 919 tion 4.1.2 for further details.

## 920 2.7 Security requirements for the primitives

921 We list below the security requirements to instantiate the primitives of the **Zeth** protocol.

- 922 •  $\text{CRH}^{\text{hsig}}$  and  $\text{CRH}^{\text{ots}}$  MUST be collision resistant functions (see: Definition 19).
- 923 •  $\text{PRF}^{\text{addr}}$ ,  $\text{PRF}^{\text{nf}}$ ,  $\text{PRF}^{\text{rho}}$  and  $\text{PRF}^{\text{pk}}$  MUST be PRF when keyed by  $ask$  and  $\phi$  and  
 924 collision-resistant (see: Definition 19, and Section 1.6.7).
- 925 •  $\text{SigSch}_{\text{OT-SIG}}$  MUST be UF-CMA (see: Definition 27 and Appendix A.2.3).
- 926 •  $\text{ComSch}$  MUST be computationally hiding, and binding (see: Section 1.6.8).
- 927 •  $\text{MKHASH}$  MUST be collision resistant with  $h_0 = 0_{\mathbb{F}_{\text{rBN}}}$  (see: Section 1.6.5).<sup>2</sup>
- 928 •  $\text{EncSch}$  MUST be IND-CCA2 and IK-CCA (see, respectively, [ABR99, Definition 8]  
 929 and Definition 14).
- 930 •  $\text{Pack}$ ,  $\text{Pack}_{\text{rsd}}$  and  $\text{Unpack}$  MUST be bijective, i.e. the encoding (resp. decoding) of  
 931 a variable via  $\text{Pack}$  and  $\text{Pack}_{\text{rsd}}$  (resp.  $\text{Unpack}$ ) MUST be bijective.
- 932 •  $\text{encode}$  and  $\text{decode}$  MUST be bijective.

### 933 2.7.1 Additional notes

#### 934 Defining *hsig*

The signature hash *hsig* is a variable used to bind the signature keys to the primary inputs. We use the same definition of *hsig* as **Zcash** to prevent the Faerie Gold attack. As a private transaction is uniquely determined by its nullifiers  $nfs = (nf_0, \dots, nf_{\text{JSIN}-1})$ , and because of the collision resistance of  $\text{CRH}^{\text{hsig}}$ , a transaction is uniquely determined by *hsig* (with overwhelming probability). We did not use the *randomSeed* defined in **Zcash** however, since this is only necessary to achieve uniqueness of *hsig* for transactions *in transit* (i.e. not mined yet) [Hop16]. The uniqueness of *hsig* is a requirement to prevent the Fairy Gold attack.

$$hsig = \text{CRH}^{\text{hsig}}(nfs, vk)$$

<sup>2</sup>This security requirement is equivalent to the one in [ZCa19, Section 5.4.1.3] where finding a preimage of  $0^{\text{SHA256DLEN}}$  must be hard.

935 **Security Requirement.**

- 936 • The variable *hsig* MUST be derived from the nullifiers  $\{nf_i\}_{i \in [\text{JSIN}]}$  and the signing  
937 key *vk* using a collision resistant function. Doing so, makes sure that *hsig* is unique  
938 for each  $tx_{\text{Mix}}$  with overwhelming probability.

939 **Defining  $\rho$**

We define  $\rho$  like in **Zcash** in order to prevent the Faerie Gold attack. A malicious sender could reuse the same  $\rho$  for a given recipient, hence correctly generating a *ZethNote* which could become unspendable by the recipient. Making  $\rho$  the output of a collision resistant PRF with random variable  $\phi$  as key and with  $tx_{\text{Mix}}$ 's *hsig* as input ensures the uniqueness of  $\rho$  — with overwhelming probability — and prevents this attack.

$$\rho_j = \text{PRF}_{\phi}^{\text{ho}}(j, \text{hsig})$$

940 **Message authentication tags  $h_i$**

The message authentication tags are used to bind the signature hash with the input notes spending keys to show ownership of the spent notes. Each tag derived from a note owner's spending key and the signature hash MUST be unique for each note with overwhelming probability.

$$h_i = \text{PRF}_{ask_i}^{\text{pk}}(i, \text{hsig})$$



## Chapter 3

# Instantiation of the cryptographic primitives

In this chapter, we start by instantiating the cryptographic building blocks used in previous sections to describe the Zeth DAP design. Finally, we proceed by providing security proofs justifying that our instantiation complies with the security requirements listed in previous sections.

### 3.1 Instantiating the PRFs, ComSch and CRHs

The functions  $\text{CRH}^{\text{hsig}}$  and  $\text{CRH}^{\text{ots}}$  are instantiated with SHA256 [oST15] which we assume to be collision resistant. Furthermore,  $\text{ComSch}$ ,  $\text{PRF}^{\text{pk}}(x)$ ,  $\text{PRF}^{\text{rho}}(x)$ ,  $\text{PRF}^{\text{addr}}(x)$ , and  $\text{PRF}^{\text{nf}}(x)$  are all instantiated with Blake2's hash function optimized for 32-bit platforms, Blake2s, which we prove in the Weakly Ideal Cipher Model [LMN16] to be from a family of PRF and collision resistant functions. The Weakly Ideal Cipher model assumes that Blake2's underlying block cipher is ideal, has a distinguisher but no structural weaknesses (see: D.2). In addition to that, and to ensure that the functions  $\text{PRF}^{\text{pk}}(x)$ ,  $\text{PRF}^{\text{rho}}(x)$ ,  $\text{PRF}^{\text{addr}}(x)$ , and  $\text{PRF}^{\text{nf}}(x)$  compute images lying in different domains, we use different message prefixes (or “domain separators”) for the PRFs inputs. This approach ensures that the  $\text{apk}_i$ 's,  $\text{nf}_i$ 's,  $\rho_i$ 's, and  $h_i$ 's have independent distributions from a PPT adversary point of view.

#### Note

It is important to note that, for this approach to be secure, the hash function used needs to be secure against *chosen-prefix collisions attacks* [Ste15].

Furthermore, we take:

- $\text{RTRAPLEN}, \text{ASKLEN}, \text{PHILEN} = \text{BLAKE2sCLEN}$

### 3.1.1 Blake2 primitive

Blake [AHMP08] is a hash family that was presented as a candidate at the SHA3 competition. Blake2 is the next iteration of the family which has been further optimized to achieve higher throughput thanks to some optimizations and by being less conservative on its security<sup>1</sup>. Blake and Blake2 are based on the ChaCha stream cipher [Ber08a] composed with the HAIFA framework [BD07]. ChaCha defined over 20 rounds, as used in Blake2, is deemed secure and a PRF based on today’s cryptanalysis [?, CM16]. Blake2 is specified in RFC-7693 [MJS15] and licensed under CC0. Blake2s is an instantiation of Blake2 optimized for 32-bit platforms. As such, to reason about the security of Blake2s we prove the security of Blake2.

**Blake Security** Blake security has been heavily scrutinized through the SHA3 competition [VNP10, MQZ10, AMP10, AAM12, AMPŠ12, ALM12, HMRS12]. Blake2 has also been thoroughly cryptanalyzed independently [GKN<sup>+</sup>14, Hao14, EFK15, NA19]. For  $n$ -bit long digests/outputs, the hash and compression functions present  $n/2$ -bit of collision resistance and  $n$ -bit of pre-image resistance, immunity to length extension, and indistinguishability from a random oracle [ANWOW13]. They have furthermore been demonstrated secure in the Weakly Ideal Cipher Model [LMN16] (WICM, see Appendix D.1.1). More particularly, Luykx et al. show that Blake2, is indistinguishable from a random oracle in this model and is a PRF.

#### Note

We assume that the encryption scheme used in the Blake2 underlying compression function — which is derived from ChaCha20 — has no exploitable structural behavior. More precisely, that this encryption scheme behaves like a weak ideal cipher. We provide proofs in this model.

We use that result in Appendix D.2 to show the collision resistance. We also prove that, given that Blake2 is collision resistant and a PRF,  $\text{Blake2}(r\|x)$  is computationally binding and computationally hiding commitment scheme for input  $x$  and randomness  $r$ .

### 3.1.2 Commitment scheme

We define our commitment scheme as follows,

$$\begin{aligned} \text{ComSch.Setup} &: \left\{ 1^\lambda \text{ s.t. } \lambda \in \mathbb{N} \right\} \rightarrow \mathbb{B}^* \\ \text{ComSch.Com} &: (\mathbb{B}^{\text{PRFADDRROUTLEN}} \times \mathbb{B}^{\text{PRFRHOOUTLEN}} \times \mathbb{B}^{\text{ZVALUELEN}}) \times \mathbb{B}^{\text{RTRAPLEN}} \rightarrow \mathbb{F}_{\text{rBN}} \end{aligned}$$

<sup>1</sup>The authors increased the number of rounds of Blake for the SHA3 competition to be more conservative on security. They however showed afterwards that this change was not “meaningfully more secure” and thus reverted it for Blake2 (see: [ANWOW13, Section 2.1]).

We instantiate the commitment scheme with Blake2s as follows,

$$\begin{aligned} pp &= \text{ComSch.Setup}(1^\lambda) \text{ (corresponds to Blake2s's constant PB and } \mathbf{r}_{\text{BN}}) \\ cm &= \text{ComSch.Com}(m = (apk, \rho, v); r) \\ &= \text{decode}_{\mathbb{N}}(\text{Blake2s}(r \| apk \| \rho \| v)) \pmod{\mathbf{r}_{\text{BN}}} \end{aligned}$$

987 **Remark 7.** We set the commitment digest length in the parameter block PB [MJS15].

### 988 Security proof.

989 The commitment scheme defined above is computationally hiding and binding in the  
990 WICM, see Appendix D.2.4. However, because of the modulo  $\mathbf{r}_{\text{BN}}$  operation, the scheme  
991 is only  $(\text{BNFIELDLEN}/2)$ -bit binding.

### 992 3.1.3 PRFs

We show in this section how we instantiate the PRFs with Blake primitives. As a reminder, the PRFs are defined as follows,

$$\begin{aligned} \text{PRF}^{\text{addr}} &: \mathbb{B}^{\text{ASKLEN}} \times \{0\} \rightarrow \mathbb{B}^{\text{PRFADDRROUTLEN}} \\ \text{PRF}^{\text{pk}} &: (\mathbb{B}^{\text{ASKLEN}} \times [\text{JSIN}]) \times \mathbb{B}^{\text{CRHHSIGOUTLEN}} \rightarrow \mathbb{B}^{\text{PRFPKOUTLEN}} \\ \text{PRF}^{\text{nf}} &: \mathbb{B}^{\text{ASKLEN}} \times \mathbb{B}^{\text{PRFRHOOUTLEN}} \rightarrow \mathbb{B}^{\text{PRFNFOUTLEN}} \\ \text{PRF}^{\text{rho}} &: (\mathbb{B}^{\text{PHILEN}} \times [\text{JSOUT}]) \times \mathbb{B}^{\text{CRHHSIGOUTLEN}} \rightarrow \mathbb{B}^{\text{PRFRHOOUTLEN}} \end{aligned}$$

As we instantiate the PRFs with Blake2s, we have,

$$\text{PRFADDRROUTLEN}, \text{PRFNFOUTLEN}, \text{PRFPKOUTLEN}, \text{PRFRHOOUTLEN} = \text{BLAKE2sCLEN}$$

To ensure that the PRFs have independent distributions, we first introduce tagging functions  $\text{tag}^x$  which truncate and prepend with a distinct tag the PRFs key. We have,

$$\begin{aligned} \text{tag}^{\text{addr}} &: \mathbb{B}^{\text{ASKLEN}} \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \\ \text{tag}^{\text{pk}} &: \mathbb{B}^{\text{ASKLEN}} \times [\text{JSIN}] \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \\ \text{tag}^{\text{nf}} &: \mathbb{B}^{\text{ASKLEN}} \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \\ \text{tag}^{\text{rho}} &: \mathbb{B}^{\text{PHILEN}} \times [\text{JSOUT}] \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \end{aligned}$$

The tagging functions are instantiated as follows,

$$\begin{aligned}
\text{tag}^{\text{addr}}(aux.jsins[i].ask) &= \text{tag}_{ask}^{\text{addr}} \\
&= (1) \parallel (1)^{\lceil \frac{\text{JSMAX}}{2} \rceil} \parallel (0, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN}-3-\lceil \frac{\text{JSMAX}}{2} \rceil}(aux.jsins[i].ask) \\
\text{tag}^{\text{nf}}(aux.jsins[i].ask) &= \text{tag}_{ask}^{\text{nf}} \\
&= (1) \parallel (1)^{\lceil \frac{\text{JSMAX}}{2} \rceil} \parallel (1, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN}-3-\lceil \frac{\text{JSMAX}}{2} \rceil}(aux.jsins[i].ask) \\
\text{tag}^{\text{pk}}(aux.jsins[i].ask, i) &= \text{tag}_{ask, i}^{\text{pk}} \\
&= (0) \parallel \text{pad}_{\lceil \frac{\text{JSMAX}}{2} \rceil}(\text{encode}_{\mathbb{N}}(i)) \parallel (0, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN}-3-\lceil \frac{\text{JSMAX}}{2} \rceil}(aux.jsins[i].ask) \\
\text{tag}^{\text{rho}}(aux.\phi, j) &= \text{tag}_{ask, j}^{\text{rho}} \\
&= (0) \parallel \text{pad}_{\lceil \frac{\text{JSMAX}}{2} \rceil}(\text{encode}_{\mathbb{N}}(j)) \parallel (1, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN}-3-\lceil \frac{\text{JSMAX}}{2} \rceil}(aux.\phi)
\end{aligned}$$

We now present how the PRFs are instantiated,

$$\begin{aligned}
\text{PRF}_{aux.jsins[i].ask}^{\text{addr}}(0) &= aux.jsins[i].znote.apk \\
&= \text{Blake2s}(\text{tag}^{\text{addr}}(aux.jsins[i].ask) \parallel \text{pad}_{\text{BLAKE2sCLEN}}(0)) \\
\text{PRF}_{aux.jsins[i].ask}^{\text{nf}}(aux.jsins[i].\rho) &= \text{prim.nfs}[i] \\
&= \text{Blake2s}(\text{tag}^{\text{nf}}(aux.jsins[i].ask) \parallel aux.jsins[i].znote.\rho) \\
\text{PRF}_{aux.jsins[i].ask}^{\text{pk}}(i, \text{prim.hsig}) &= \text{prim.htags}[i] \\
&= \text{Blake2s}(\text{tag}^{\text{pk}}(aux.jsins[i].ask, i) \parallel \text{prim.hsig}) \\
\text{PRF}_{aux.\phi}^{\text{rho}}(j, \text{prim.hsig}) &= aux.znotes[j].\rho \\
&= \text{Blake2s}(\text{tag}^{\text{rho}}(aux.\phi, j) \parallel \text{prim.hsig})
\end{aligned}$$

**Remark 8.** We set the PRFs' output length in the Blake2s's parameter block PB.

#### Security proof.

The functions defined above are collision resistant and PRFs in the WICM, see Appendix D.2. Because of the tagging functions, the security parameter of the PRFs becomes  $\lambda = \text{BLAKE2sCLEN}/2 - \text{JSMAX}/4 - 3/2$ .

#### 3.1.4 Collision Resistant Hashes

We instantiate in this section the collision resistant hash functions  $\text{CRH}^{\text{hsig}}$  and  $\text{CRH}^{\text{ots}}$  with SHA256. As a consequence, we have,

$$\text{CRHHSIGOUTLEN} = \text{CRHOTSOUTLEN} = \text{SHA256DLEN}$$

**SHA256 Security** SHA-256 (Secure Hash Algorithm 256) is a hash function designed by the National Security Agency (NSA) in 2001. It is based on the Merkle–Damgård structure, the Davies–Meyer compression function construct [BRS02, Function f5 in Figure 3] and the classified SHACAL-2 block cipher.

Collision attacks have been thoroughly studied by the research community [SS08, MNS11]. The best attacks at this day, are second-order differential attack by Lamberger et al. [LM11] on the SHA-256 compression function reduced to 46 out of 64 rounds.

Many researchers [IS09, AGM<sup>+</sup>09] have also studied preimage attacks on SHA-256 with reduced rounds. Guo et al. [GLRW10] in particular were among the first to use the meet in the middle strategy [AS09] and achieved more efficient ones on 42-step SHA-256. Khovratovich et al. in 2012 [KRS12] have so far presented the best preimage attacks, on 45-round and 52-round SHA-256 as well as a 52-round attack on the SHA-256 compression function.

Li et al. have published in 2012 [LIS12] a noteworthy paper on converting meet in the middle preimage attack into pseudo collision attack. Using preimage attacks by bicliques, they found pseudo collisions attacks on 52 steps of SHA-256.

**Lemma 2.** *SHA256 is 128-bit collision resistant.*

## 3.2 Instantiating MKHASH

In this section, we show how we instantiate MKHASH with a compression function based on MIMC [AGR<sup>+</sup>16]. We start by showing how the compression function is constructed. Then, we instantiate MKHASH, and we finally prove that our instantiation complies with the security requirements mentioned in Section 2.7

### 3.2.1 MIMC Encryption

MIMC is a block cipher with a simple design. During a round, the message  $m$  is first mixed with the encryption key  $k$  and a randomly chosen constant  $c[i]$ . We then apply a permutation function on this result, this is done by raising it with a carefully chosen exponent  $e$  (see: Section 3.2.1). This round is repeated a specific number of time *rounds*, where *rounds* depends on the desired security level  $\lambda$ . We denote the encryption function by MIMC-Encrypt and illustrate it in Fig. 3.1.

---

```

MIMC-Encrypt( $k, m, c, e, rounds$ )
1 : foreach  $i \in [rounds]$  :
2 :    $m = (k \text{ OP } c[i] \text{ OP } m)^e$ 
3 : return ( $m \text{ OP } k$ )

```

Figure 3.1: MIMC Encryption function.

MIMC-Encrypt can be defined on both binary and prime fields. As such, the OP operation corresponds to either  $\oplus$  or  $+$  (mod  $p$ ) [AGR<sup>+</sup>16, GRR<sup>+</sup>16]. We denote by

MIMC<sub>p</sub> (respectively MIMC<sub>2<sup>n</sup></sub>) the MIMC-Encrypt function defined over  $\mathbb{F}_p$  (respectively  $\mathbb{F}_{2^n}$ ). Since block ciphers are usually defined over the product space of keys and messages, we consider the variables  $c$ ,  $rounds$  and  $e$  as fixed. The functions' signature thus become,

$$\begin{aligned} \text{MIMC}_p &: \mathbb{F}_p \times \mathbb{F}_p \rightarrow \mathbb{F}_p \\ \text{MIMC}_{2^n} &: \mathbb{F}_{2^n} \times \mathbb{F}_{2^n} \rightarrow \mathbb{F}_{2^n} \end{aligned}$$

1028 We will from now on only consider MIMC defined over prime fields as we consider  
1029 SNARKs based on arithmetic circuits defined over a finite field  $\mathbb{F}_p$ ,  $p$  prime.

### 1030 Security parameters and analysis

1031 Assume our prime  $p$  is written over at least  $\lambda$  bits. To ensure that the exponentiation  
1032 leads to a permutation in  $\mathbb{F}_p$ , we have to choose  $e$  of the form  $2^t - 1$  such that  $\gcd(e, p -$   
1033  $1) = 1$ . To achieve a security of  $\lambda$ , we furthermore define the number of rounds as  
1034 follows,  $rounds = \lceil \frac{\log_2 p}{\log_2 e} \rceil$ .<sup>2</sup>

1035 We refer to MIMC paper [AGR<sup>+</sup>16, Section 4.2 and 5.1] for more details on the  
1036 security analysis and attacks on the scheme. We can note that MIMC<sub>p</sub> does not suffer  
1037 from *inversion subfield attacks* as the only subgroup in a prime order group are itself  
1038 and the trivial group - induced by the identity element.

### 1039 3.2.2 MIMC-based compression function

1040 To construct a hash function from a block-cipher (or permutation), there exists two main  
1041 techniques: sponge functions [BDPVA07] and iterated compression functions [BRS02].  
1042 In a Merkle Tree, we need to hash two fixed size inputs into one. This is the definition of  
1043 a compression function. As such, we decided to use a compression function and settled  
1044 for Miyaguchi-Preneel compression function [BRS02,  $f_3$  function] construct for its proven  
1045 security (it is more secure than the more common Davies-Meyer construct [BRS02,  $f_5$   
1046 function]), and we do not require the flexibility of the latter [GFBR06, Section 3]).

### 1047 Miyaguchi-Preneel Compression construct

1048 Miyaguchi-Preneel (MP) [BRS02,  $f_3$  function] is a general construction that allows to  
1049 build compression functions from block ciphers (see: Section 1.6.5). Given a block cipher  
1050  $E$ , we denote the corresponding compression function by  $f_E^{\text{MP}}$  and its description is given  
1051 in Fig. 3.2. We focus on ciphers working on  $\mathbb{F}_p$ , for some prime  $p$ , while the original  
1052 construction is defined over binary fields. As such, we replace the bitwise addition  $\oplus$  by  
1053 the modular addition in  $\mathbb{F}_{\text{rBN}}$  (see: [Har19]).

1054 We denote by MIMC-MP the compression function defined by the application of the  
1055 Miyaguchi-Preneel construct over MIMC. Similarly, we define MIMC-MP<sub>p</sub>, illustrated  
1056 in Fig. 3.3, by using MIMC<sub>p</sub> as the following compression function defined over  $\mathbb{F}_p$ ,

---

<sup>2</sup>We do not consider exponents of type  $2^t + 1$  since the polynomial representing MIMC-Encrypt would be sparse and as such, the number of rounds becomes constant (see: [AGR<sup>+</sup>16, Section 5.3]).

$f_E^{\text{MP}}(k, m)$

1 :  $res = E_k(m)$   
 2 : **return**  $(res + m + k) \pmod{p}$

Figure 3.2: MP construct in  $\mathbb{F}_p$ .

$e$	$rounds$	$Constraints$
7	91	365
31	52	417
127	37	445
511	29	465
2047	24	481

$\text{MIMC-MP}_p(k, m)$

1 :  $res = \text{MIMCp}(k, m)$   
 2 : **return**  $(res + k + m) \pmod{p}$

Figure 3.3: MIMC-MP<sub>p</sub> construction.

$e$	$rounds$	$Constraints$
8191	20	481
32676	17	477
131071	15	481
524287	14	505
2097151	13	521

Table 3.1: Number of constraints for MIMC-MP for different exponents.

### 3.2.3 An efficient instantiation of MIMC primitives

In this section, we give instances of  $\text{MIMCp}$  and  $\text{MIMC-MP}_p$  that offer the best compromise between having an efficient prover and a cheap (in terms of gas consumption) on-chain compression function, and where  $p = r_{\text{BN}}$ .

We observe that a small  $e$  would reduce the number of constraints in the arithmetic circuit (see: Section 2.2) while a large one would reduce the gas necessary to carry out Merkle Tree operations on the contract (see: Section 2.5). This is due to the fact that the exponentiation is cheaper to execute on a contract than on a circuit, and that the number of rounds decreases with higher  $e$ . For instance, choosing  $e = 7$  results in 365 constraints and  $\approx 20k$  gas while  $e = 31$  corresponds to 417 constraints (+15%) and  $\approx 17k$  (−10%) in gas consumption. Repeating the same process for different exponents, we observe roughly the same order of magnitude gain on the gas consumption and loss on the number of constraints.

The number of constraints of MIMC-MP for different exponents  $e$  of type  $2^t - 1$  is given by the following formula,

$$Constraints = rounds \cdot mults + 1$$

where  $rounds = \lceil \frac{\log_2 r_{\text{BN}}}{\log_2 e} \rceil$  and  $mults = 2 \cdot t - 2$  (using the *square-and-multiply* algorithm [MVOV96, 14.06]) and the last constraint represents the final message and key addition. The number of constraints is showed in Table 3.1, with  $e$  chosen such that  $\gcd(e, r_{\text{BN}} - 1) = 1$  and the number of rounds  $rounds$  chosen to attain a security level of  $\log(r_{\text{BN}})$ . As we expect, the number of constraints increases with  $t$ .

In conclusion, we choose  $e = 7$  which is the value for the exponent  $e$  that offers the best balance between the number of constraints in the arithmetic circuit and the gas

cost of hashing on the contract. We call these “snark-efficient” instances  $\text{MIMC}[7]_{\mathbf{r}_{\text{BN}}}$  and  $\text{MIMC}[7]\text{-MP}_{\mathbf{r}_{\text{BN}}}$  and show them in Fig. 3.4 and Fig. 3.5.

$\text{MIMC}[7]_{\mathbf{r}_{\text{BN}}}(k, m)$	$c = (c[0], \dots, c[\text{rounds} - 1])$
1 : <b>foreach</b> $i \in [91]$ :	$iv = \text{Keccak256}(\text{“clearmatics\_mt\_seed”})$
2 : $m = (k + c[i] + m)^7 \pmod{\mathbf{r}_{\text{BN}}}$	$c[0], c[1] = 0, \text{Keccak256}(iv)$
3 : <b>return</b> $(m + k) \pmod{\mathbf{r}_{\text{BN}}}$	$\forall i > 1, c[i] = \text{Keccak256}(c[i - 1])$

Figure 3.4:  $\text{MIMC}[7]_{\mathbf{r}_{\text{BN}}}$  construction.

$\text{MIMC}[7]\text{-MP}_{\mathbf{r}_{\text{BN}}}(k, m)$	$\text{MKHASH}(m_0, m_1)$
<b>return</b> $\text{MIMC}[7]_{\mathbf{r}_{\text{BN}}}(k, m) + m + k \pmod{\mathbf{r}_{\text{BN}}}$	<b>return</b> $\text{MIMC}[7]\text{-MP}_{\mathbf{r}_{\text{BN}}}(m_0, m_1)$

Figure 3.5:  $\text{MIMC}[7]\text{-MP}_{\mathbf{r}_{\text{BN}}}$  construction.

Figure 3.6: MKHASH instantiation.

**Remark 9.** Note that  $\text{rounds} = 91$  targets 254 bit security level for MIMC-MP and that Keccak256 is the 256 bit digest instance of the Keccak family that won the NIST SHA-3 competition [GJMG11]. Compared to other hash functions, its use in solidity contracts is quite convenient since an EVM opcode is available (see: [W<sup>+</sup>, Appendix G]).

### 3.2.4 MKHASH instantiation

We instantiate MKHASH with  $\text{MIMC}[7]\text{-MP}_{\mathbf{r}_{\text{BN}}}$  defined in Fig. 3.6,

$$\text{MKHASH} : \mathbb{F}_{\mathbf{r}_{\text{BN}}} \times \mathbb{F}_{\mathbf{r}_{\text{BN}}} \rightarrow \mathbb{F}_{\mathbf{r}_{\text{BN}}}$$

**Remark 10.** To increase the security of the MKHASH, different round constants for each level of the Merkle tree could be used.

### 3.2.5 Security requirements satisfaction

After presenting the state of the art of MiMC cryptanalysis, we present the security proof of MIMC-MP collision resistance.

### Cryptanalysis of MiMC block cipher and primitives

MIMC’s security is increasingly being analysed for its traction in zero-knowledge and cryptocurrency communities. As of today, we do not know of any attacks breaking MIMC on prime fields on full rounds.

The first attack on MIMC was an interpolation attack [LP19] which targets a reduced-round version for a scenario in which the attacker has only limited memory. An attack on Feistel-based MIMC [Bon19] was discovered shortly after, by using generic properties of the used Feistel construction (instead of exploiting properties of the primitive itself).



1097 Additionally, [ACG<sup>+</sup>19] proposes an attack based on Gröbner basis. The authors state  
 1098 that by introducing a new intermediate variable in each round, the resulting multivariate  
 1099 system of equations is a Gröbner basis. As such, the first step of a Gröbner basis attack  
 1100 can be obtained for free. However, the following steps of the attack are so computation-  
 1101 ally demanding that the attack becomes infeasible in practice. A recent work [EGL<sup>+</sup>20]  
 1102 targets MIMC on binary fields, and achieves a full-round break of the scheme. While,  
 1103 the attack presented does not apply to prime fields, the authors note that it “can be  
 1104 generalized to include ciphers over  $\mathbb{F}_p$ ”, and that only the lack of efficient distinguishers  
 1105 over prime fields precludes this. Another attack from Beyne et al [BCD<sup>+</sup>20] uses a low  
 1106 complexity distinguisher against full MIMC permutation leading to a practical collision  
 1107 attack on reduced round sponge-based MIMC hash defined with security of 128 bits.

### 1108 Security proof of MIMC-MP collision resistance

1109 We now prove that this compression scheme satisfies all the security requirements listed  
 1110 in Section 2.7. To do so, we first assume that the round constants are pseudo-random,  
 1111 i.e. that Keccak256 is a PRF.

1112 **Lemma 3.** *Keccak256 is a PRF with  $\lambda = 128$ .*

1113 The security of MIMC-MP derives from a more general result, i.e. by modelling MIMC  
 1114 as an ideal cipher (see: Definition 16). More specifically, we show a security result for  
 1115 the MP construction on  $\mathbb{F}_{\mathbf{r}_{\text{BN}}}$  by proving that, in the Ideal Cipher Model, the collision  
 1116 resistance advantage of any adversary is bounded by  $\frac{q(q+1)}{\mathbf{r}_{\text{BN}}}$ , where  $q$  is the number of  
 1117 different queries that the attacker does to the oracle. This means that, assuming a  
 1118 maximum  $q$  number of possible encryption/decryption queries, parameter  $\mathbf{r}_{\text{BN}}$  can be  
 1119 chosen to make the advantage small as needed and  $\mathbf{f}_{\text{E}}^{\text{MP}}$  considered collision resistant.  
 1120 Similar result applies to the  $2^n$  case.

1121 The instance of MIMC we use is modelled as an ideal cipher defined on field elements,  
 1122 for this reason we consider a variant of the ICM model where the keys, inputs and outputs  
 1123 are field elements in  $\mathbb{F}_{\mathbf{r}_{\text{BN}}}$  and the block cipher scheme, with key  $k$ , correspond to a family  
 1124 of  $\mathbf{r}_{\text{BN}}$  independent random permutations  $f_k : \mathbb{F}_{\mathbf{r}_{\text{BN}}} \times \mathbb{F}_{\mathbf{r}_{\text{BN}}} \rightarrow \mathbb{F}_{\mathbf{r}_{\text{BN}}}$ .

1125 In the proof, without loss of generality, we assume the following conventions for an  
 1126 adversary  $\mathcal{A}$ :

- 1127 • the adversary asks distinct queries: i.e. if  $\mathcal{A}$  asks a query  $\text{O}^{\text{E}}(k, m)$  and this returns  
 1128  $y$ , then  $\mathcal{A}$  does not ask a subsequent query of  $\text{O}^{\text{E}}(k, m)$  or  $\text{O}^{\text{E}^{-1}}(k, y)$ , and inversely;
- 1129 • the adversary necessarily obtained the candidate collision from the oracle. This  
 1130 property follows suite from modelling MIMC as an ideal cipher.

1131 **Lemma 4.** *Let  $\mathbf{f}_{\text{E}}^{\text{MP}}$  be the MP compression function built on an ideal block-cipher E on  
 1132  $\mathbb{F}_{\mathbf{r}_{\text{BN}}}$ , the probability for an adversary  $\mathcal{A}$  to find a collision is lower than  $q(q+1)/\mathbf{r}_{\text{BN}}$   
 1133 where  $q$  is the non-null number of distinct oracle queries.*

1134 The following proof has been adapted from [BRS02, Lemma 3.3]<sup>3</sup>.

1135 *Proof.* Fix  $h_0 \in \mathbb{F}_{\mathbf{r}_{\text{BN}}}$ . Let  $\mathcal{A}$  be an adversary attacking the compression function  $f_{\text{E}}^{\text{MP}}$ .  
 1136 Assume that  $\mathcal{A}$  asks the oracles  $\text{O}^{\text{E}}$  and  $\text{O}^{\text{E}^{-1}}$  a total of *distinct*  $q$  queries. Let us  
 1137 denote the result of the  $q$  queries and output of the attacker (candidate collision) as  
 1138  $((k_1, m_1, y_1), \dots, (k_q, m_q, y_q), \text{out})$ . If  $\mathcal{A}$  is successful it means that it outputs  $(k, m)$ ,  
 1139  $(k', m')$  such that either  $(k, m) \neq (k', m')$  and  $f_{\text{E}}^{\text{MP}}(k, m) = f_{\text{E}}^{\text{MP}}(k', m')$  or  $f_{\text{E}}^{\text{MP}}(k, m) =$   
 1140  $h_0$ . By the definition of  $f_{\text{E}}^{\text{MP}}$ , we have that  $E_k(m) + m + k = E_{k'}(m') + m' + k'$  for  
 1141 the first case, or  $E_k(m) + m + k = h_0$  for the second. So either there are distinct  
 1142  $r, s \in [1, \dots, q]$  such that  $(k_r, m_r, y_r) = (k, m, E_k(m))$  and  $(k_s, m_s, y_s) = (k', m', E_{k'}(m'))$   
 1143 and  $E_{k_r}(m_r) + m_r + k_r = E_{k_s}(m_s) + m_s + k_s$  or else there is an  $r \in [1, \dots, q]$  s.t.  $(k_r, m_r, y_r) =$   
 1144  $(k, m, h_0)$  and  $E_{k_r}(m_r) + m_r + k_r = h_0$ . We show that this event is unlikely.

In fact, for each  $i \in [1, \dots, q]$ , let  $C_i$  be the event that either  $y_i + m_i + k_i = h_0$  or  
 does exist  $j \in [1, \dots, i-1]$  s.t.  $y_i + m_i + k_i = y_j + m_j + k_j$ . When carrying out the  
 simulation  $y_i$  or  $m_i$  was randomly selected from a set of at least  $\mathbf{r}_{\text{BN}} - (i-1)$  elements,  
 so  $\Pr[C_i] \leq i/(\mathbf{r}_{\text{BN}} - i)$ . This means that for the collision advantage of  $\mathcal{A}$ ,  $\text{Adv}_{f_{\text{E}}^{\text{MP}}, \mathcal{A}}^{\text{coll}}$   
 holds that  $\text{Adv}_{f_{\text{E}}^{\text{MP}}, \mathcal{A}}^{\text{coll}} \leq \Pr[C_1 \vee \dots \vee C_q] \leq \sum_{i=1}^q \Pr[C_i]$ . For  $q \leq \frac{\mathbf{r}_{\text{BN}}}{2}$  this probability is  
 bounded by  $l \cdot \frac{q(q+1)}{\mathbf{r}_{\text{BN}}}$ . However, we allow only polynomial number of queries, thus for  
 $q = \text{poly}(\lambda)$  this probability becomes,

$$\frac{\text{poly}(\lambda)}{\mathbf{r}_{\text{BN}}}$$

1145

□

1146 **Remark 11.** Note that from Lemma 4 follows that the collision resistance security of  
 1147 the **Zeth** Merkle tree is  $\log_2(\mathbf{r}_{\text{BN}}/2)$  (around 127 bits).

#### Note

Lemma 4 is applicable to our case by the strong assumption of  $\text{MIMC}[7]_{\mathbf{r}_{\text{BN}}}$  being  
 an ideal cipher. In other words, the proof does not take into account any structural  
 weakness or knowledge that an attacker is aware of. These additional information  
 makes Lemma 4 not applicable anymore to our case and as consequence can be  
 used to break the collision resistance.

1148

### 1149 3.3 Zeth statement after primitive instantiation

1150 After instantiating the various primitives and providing security proofs to justify that  
 1151 they comply with the security requirements listed in previous sections,  $\mathbf{R}^{\text{Z}}$ , now becomes:

<sup>3</sup>It states the collision resistance of a set of compression functions  $f_1, \dots, f_{12}$ , denoted as *group-1 compression functions* and showed in [BRS02, Figure 3]. As mentioned above, Miyaguchi-Preneel corresponds to  $f_3$  of that group. Since the proof of [BRS02, Lemma 3.3] shows collision resistance of  $f_1$ , we slightly modified it to work for  $f_3$ .

- 1152 • For each  $i \in [\text{JSIN}]$ :
  - 1153 1.  $\text{aux.jsins}[i].\text{znote.apk} = \text{Blake2s}(\text{tag}_{ask}^{\text{addr}} \parallel \text{pad}_{\text{BLAKE2sCLEN}}(0))$
  - 1154 with  $\text{tag}_{ask}^{\text{addr}}$  defined in Section 3.1.3
  - 1155 2.  $\text{aux.jsins}[i].\text{nf} = \text{Blake2s}(\text{tag}_{ask}^{\text{nf}} \parallel \text{aux.jsins}[i].\text{znote}.\rho)$
  - 1156 with  $\text{tag}_{ask}^{\text{nf}}$  defined in Section 3.1.3
  - 1157 3.  $\text{aux.jsins}[i].\text{cm} = \text{Blake2s}(\text{aux.jsins}[i].\text{znote}.r \parallel m)$
  - 1158 with  $m = \text{aux.jsins}[i].\text{znote.apk} \parallel \text{aux.jsins}[i].\text{znote}.\rho \parallel \text{aux.jsins}[i].\text{znote}.v$
  - 1159 4.  $\text{aux.htags}[i] = \text{Blake2s}(\text{tag}_{ask,i}^{\text{pk}} \parallel \text{prim.hsigs})$  (malleability fix, see: Appendix A)
  - 1160 with  $\text{tag}_{ask,i}^{\text{pk}}$  defined in Section 3.1.3
  - 1161 5.  $(\text{aux.jsins}[i].\text{znote}.v) \cdot (1 - e) = 0$  is satisfied for the boolean value  $e$  set such
  - 1162 that if  $\text{aux.jsins}[i].\text{znote}.v > 0$  then  $e = 1$ .
  - 1163 6. The Merkle root  $\text{mkroot}'$  obtained after checking the Merkle authentication
  - 1164 path  $\text{aux.jsins}[i].\text{mkpath}$  of commitment  $\text{aux.jsins}[i].\text{cm}$ , with  $\text{MIMC}[7]\text{-MP}_{\text{rBN}}$ ,
  - 1165 equals to  $\text{prim.mkroot}$  if  $e = 1$ .
  - 1166 7.  $\text{prim.nfs}[i]$
  - 1167  $= \{\text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.jsins}[i].\text{nf}[k \cdot \text{BNFIELD CAP}:(k+1) \cdot \text{BNFIELD CAP}])\}_{k \in [\lfloor \text{PRFNFOUTLEN}/\text{BNFIELD CAP} \rfloor]}$
  - 1168 8.  $\text{prim.htags}[i]$
  - 1169  $= \{\text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.htags}[i][k \cdot \text{BNFIELD CAP}:(k+1) \cdot \text{BNFIELD CAP}])\}_{k \in [\lfloor \text{PRFPKOUTLEN}/\text{BNFIELD CAP} \rfloor]}$
- 1170 • For each  $j \in [\text{JSOUT}]$ :
  - 1171 1.  $\text{aux.znotes}[j].\rho = \text{Blake2s}(\text{tag}_{ask,j}^{\rho} \parallel \text{prim.hsigs})$  (malleability fix, see: Appendix A)
  - 1172 with  $\text{tag}_{ask,j}^{\rho}$  defined in Section 3.1.3
  - 1173 2.  $\text{prim.cms}[j] = \text{Blake2s}(\text{aux.znotes}[j].r \parallel m)$
  - 1174 with  $m = \text{aux.znotes}[j].\text{apk} \parallel \text{aux.znotes}[j].\rho \parallel \text{aux.znotes}[j].v$
- 1175 •  $\text{prim.hsigs} = \{\text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.hsigs}[k \cdot \text{BNFIELD CAP}:(k+1) \cdot \text{BNFIELD CAP}])\}_{k \in [\lfloor \text{CRHHSIGOUTLEN}/\text{BNFIELD CAP} \rfloor]}$
- 1176 •  $\text{prim.rsd} = \text{Pack}_{\text{rsd}}(\{\text{aux.jsins}[i].\text{nf}\}_{i \in [\text{JSIN}]}, \text{aux.vin}, \text{aux.vout}, \text{aux.hsigs}, \{\text{aux.htags}[i]\}_{i \in [\text{JSIN}]})$
- Check that the “joinsplit is balanced”, i.e. check that the joinsplit equation holds:

$$\begin{aligned}
 & \text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.vin}) + \sum_{i \in [\text{JSIN}]} \text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.jsins}[i].\text{znote}.v) \\
 &= \sum_{j \in [\text{JSOUT}]} \text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.znotes}[j].v) + \text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.vout})
 \end{aligned}$$

1177 **Remark 12.** For higher security, we could use Blake2b with 32-byte output instead  
 1178 of SHA256. In fact, since a precompiled contract computing the Blake2 compression  
 1179 function [MJS15] has been added to the Istanbul release of *Ethereum* (EIP 152 [THH15]),  
 1180 it could be possible to write a small wrapper on the smart contracts, in order to hash  
 1181 with Blake2b with any parameter.

### 3.3.1 Instantiating the Packing functions

As we consider SNARKs based on arithmetic circuits based over a prime finite field, our statement’s variables are interpreted as field elements. As such, we can take advantage of the packing strategy to reduce the length of the primary inputs and diminish the cost of the on-chain verification. Indeed, the cost of Groth proof verification is linear in the number of primary inputs, each variable is the input of a costly scalar multiplication of an element in  $\mathbb{G}_1$ . Hence, while packing puts more “pressure” on the prover — by adding constraints in the circuit — it simplifies the verifier’s work.

We detail in this section how we pack the primary inputs to minimize their number and present efficient packing and unpacking functions.

The original primary inputs (see: [RZ19, Section 3.4.3]) were the input nullifiers, the output commitments, the public values **to which we added the signature hash and the authentication tags for security** (malleability fix, see: Appendix A).

$$(\{prim.nf_i\}_{i \in [JSIN]}, \{prim.cms[j]\}_{j \in [JSOUT]}, vin, vout, hsig, \{prim.htags[i]\}_{i \in [JSIN]})$$

Consider the binary variables, that is the nullifiers *nfs*, the public values *vin* and *vout*, the signature hash *hsig* and the authentication tags *htags*. For each of these variables  $x$ , let  $\alpha_x = \lceil \text{length}(x) / \text{BNFIELD CAP} \rceil$ , represent the number of field elements to encode all of the variable’s bits. Let  $\beta_x = \lfloor \text{length}(x) / \text{BNFIELD CAP} \rfloor$  represent the number of field elements which are fully “filled” and  $\gamma_x = \text{length}(x) \pmod{\text{BNFIELD CAP}}$  the remaining ones. For simplicity, as we assume  $\text{ZVALUELEN} < \text{BNFIELD CAP}$ , we only define the notation  $\gamma_v = \text{ZVALUELEN}$  for the public values *vin* and *vout*.

We denote by  $\text{RSDBLEN}$  the total number of bits which could not fully fill a field element, i.e. the weighted sum of the  $\gamma_x$ ,

$$\text{RSDBLEN} = \gamma_{hsig} + 2 \cdot \gamma_v + \text{JSIN} \cdot (\gamma_{nf} + \gamma_h)$$

A simple packing strategy would have been to encode each primary input’s field  $x$  as  $\alpha_x$  field elements. However, this leads to a potentially large waste of space when the PRFs and hash digests are slightly longer than  $\text{BNFIELD CAP}$  bits. Another strategy would have been to encode the binary string which results from the concatenation of all binary variables. This would have made the necessary unpacking (see: Section 2.5) costly. We thus decided to keep the fully filled  $\sum_x \beta_x$  field elements which are “fully” filled and aggregate the remaining  $\text{RSDBLEN}$  “residual” bits in a new variable *rsd*:

$$\begin{aligned} \text{NFFLEN} &= \lfloor \text{PRFNFOUTLEN} / \text{BNFIELD CAP} \rfloor \\ \text{HSIGFLEN} &= \lfloor \text{CRHHSIGOUTLEN} / \text{BNFIELD CAP} \rfloor \\ \text{HFLEN} &= \lfloor \text{PRFPKOUTLEN} / \text{BNFIELD CAP} \rfloor \\ \text{RSDFLEN} &= \lceil \text{RSDBLEN} / \text{BNFIELD CAP} \rceil \end{aligned}$$

To facilitate the unpacking of the primary inputs, we chose to first aggregate the primary inputs’ singular elements. More precisely, regardless of the values  $\text{JSIN}$  and

JSOUT, the public values  $vin$  and  $vout$ , and the hash signature  $hsig$  will always be at the same location in the  $rsd$  string. The residual bits  $rsd$  are thus formatted as follows,

$$vin || vout || hsig || nfs || htags.$$

To format the unpacked primary inputs into field elements, we define the following functions. The algorithm **Pack** (see: Fig. 3.7), given a bit string of length less than  $\text{BNFIELD CAP}$ , returns a field element. The algorithm  $\text{Pack}_{rsd}$  (see: Fig. 3.8) given the nullifiers, public values and authentication tags outputs the residual bits. For a given field, the algorithm **Unpack**, given the associated packed field elements and the residual bits returns the variables reassembled in binary strings. For instance, we have that  $\text{Unpack}_{nf}(\text{prim}.nfs, rsd) = \{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}$ .

$$\text{Pack} : \mathbb{B}^{\leq \text{BNFIELD CAP}} \rightarrow \mathbb{F}_{\text{rBN}}$$

$$\text{Pack}_{rsd} : (\mathbb{B}^{\text{PRFNFOUTLEN}})^{\text{JSIN}} \times \mathbb{B}^{\text{ZVALUELEN}^2} \times \mathbb{B}^{\text{CRHHSIGOUTLEN}} \times (\mathbb{B}^{\text{PRFPKOUTLEN}})^{\text{JSIN}} \rightarrow \mathbb{F}_{\text{rBN}}^{\text{RSDFLEN}}$$

$$\text{Unpack} : \mathbb{F}_{\text{rBN}}^* \times \mathbb{F}_{\text{rBN}}^{\text{RSDFLEN}} \rightarrow \mathbb{B}^*$$

More particularly, we use the function **Unpack** for the nullifiers, public values and signature hash. As such, we have,

$$\text{Unpack}_{hsig} : \mathbb{F}_{\text{rBN}}^{\text{HSIGFLEN}} \times \mathbb{F}_{\text{rBN}}^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{CRHHSIGOUTLEN}}$$

$$\text{Unpack}_{nf} : \mathbb{F}_{\text{rBN}}^{\text{NFFLEN}} \times \mathbb{F}_{\text{rBN}}^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{PRFNFOUTLEN}}$$

$$\text{Unpack}_{vin} : \mathbb{F}_{\text{rBN}}^0 \times \mathbb{F}_{\text{rBN}}^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{ZVALUELEN}}$$

$$\text{Unpack}_{vout} : \mathbb{F}_{\text{rBN}}^0 \times \mathbb{F}_{\text{rBN}}^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{ZVALUELEN}}$$

## 1199 Packing Policy Security

1200 **Proposition 1** (Packing security). *The encoding (resp. decoding) of a variable via **Pack***  
 1201 *and  $\text{Pack}_{rsd}$  (resp. **Unpack**) is bijective.*

## 1202 Packing Policy Example

In the case where  $\text{JSIN} = \text{JSOUT} = 2$ ,  $\text{BNFIELD CAP}$  and all PRFs, and  $\text{CRH}^{\text{hsig}}$  output 256 bits, the unpacked primary inputs are 2167-bit long. The packing parameters thus are,

$$\text{RSDBLEN} = 143$$

$$\text{NFFLEN} = \text{HSIGFLEN} = \text{HFLEN} = \text{RSDFLEN} = 1$$

The packed primary inputs are 2277 bits long which corresponds to a small space overhead ( $\approx 5\%$  unused bits). Moreover, as the residual bits are 143 bit long, they can be written over a single field element. As such, the primary inputs are 9 field element long. Finally, the residual bits are formatted as follows,

$$\underbrace{\text{padding}}_{113 \text{ bits}} || \underbrace{vin}_{64 \text{ bits}} || \underbrace{vout}_{64 \text{ bits}} || \underbrace{hsig}_{3 \text{ bits}} || \underbrace{nf_0}_{3 \text{ bits}} || \underbrace{nf_1}_{3 \text{ bits}} || \underbrace{h_0}_{3 \text{ bits}} || \underbrace{h_1}_{3 \text{ bits}}$$

---

 $\text{Pack}_{\mathbb{F}_{\text{rBN}}}(x)$ 


---

```

out  $\leftarrow 0_{\mathbb{F}_{\text{rBN}}}$ ;
foreach  $i \in [\text{length}(x)]$  do :
  if  $x[i] = 1$  do :
    out  $\leftarrow \text{out} +_{\mathbb{F}_{\text{rBN}}} 2^{\text{length}(x)-1-i}$ 
return out;

```

Figure 3.7: Packing algorithm in Big Endian.

---

 $\text{Pack}_{\text{rsd}}(nfs, vin, vout, hsig, htags)$ 


---

```

out  $\leftarrow []$ ;  $r \leftarrow \epsilon$ ;
 $r \leftarrow vin[ \lfloor \text{ZVALUELEN}/\text{BNFIELD CAP} \rfloor \cdot \text{BNFIELD CAP} : ]$ ;
 $r \leftarrow r || vout[ \lfloor \text{ZVALUELEN}/\text{BNFIELD CAP} \rfloor \cdot \text{BNFIELD CAP} : ]$ ;
 $r \leftarrow r || hsig[ \lfloor \text{CRHSIGOUTLEN}/\text{BNFIELD CAP} \rfloor \cdot \text{BNFIELD CAP} : ]$ ;
for  $i \in [\text{JSIN}]$  do :
   $r \leftarrow r || nfs[i][ \lfloor \text{PRFNFOUTLEN}/\text{BNFIELD CAP} \rfloor \cdot \text{BNFIELD CAP} : ]$ ;
for  $i \in [\text{JSIN}]$  do :
   $r \leftarrow r || htags[i][ \lfloor \text{PRFPKOUTLEN}/\text{BNFIELD CAP} \rfloor \cdot \text{BNFIELD CAP} : ]$ ;
for  $i \in [\lceil \text{length}(r)/\text{BNFIELD CAP} \rceil]$  do :
  out  $\leftarrow \text{Pack}_{\mathbb{F}_{\text{rBN}}}(r[i \cdot \text{BNFIELD CAP} : (i+1) \cdot \text{BNFIELD CAP}])$ ;
return out;

```

Figure 3.8: Packing residual bits algorithm.

### 3.4 Instantiate SigSch<sub>OT-SIG</sub>

In **Zeth**, we chose to use the one-time Schnorr-based signature scheme introduced by Bellare and Shoup [BS07], over **BN-254**, for its long proven security, simplicity, speed and size. Its security relies on the one-more discrete log problem (see: Definition 6) and the collision resistance of the underlying hash function CRH (see: Definition 19) that we instantiate with **SHA256**.

This one-time signature scheme (see: Definition 29) is defined by the two-tier signature scheme over a cyclic group  $(p, \mathbb{G}, \langle \mathfrak{g} \rangle, \otimes)$ . In the two-tier signature scheme, the hash function CRH only needs to be collision resistant (the random oracle model is not used). Similarly, the variable  $hk$  represents the key of the hash function (a particular instance).

To turn this two-tier signature scheme into a one-time signature scheme, one simply has to define the one-time signature key generation **KGen** as the combination of both primary and secondary key generations of the two-tier (see: [BS07, Section 6]). The one-time signing key (respectively verification key) of the one time signature scheme is defined as both the primary and secondary signing key (respectively verification key) of the two-tier scheme, Fig. 3.9

#### 3.4.1 Security requirements satisfaction

We now prove that this signature scheme satisfies all the security requirements listed in Section 2.7.

**Theorem 1.** *The One-Time Schnorr signature is strongly unforgeable under chosen-message attacks (SUF-CMA) assuming that the om-DLog problem is hard in  $\mathbb{G}$  and that the hash function is collision resistant.*

*Proof.* See: [BS07, Theorems 5.1, 5.2 and 6.1]. □

<b>KGen()</b> : $hk \leftarrow_{\$} \mathbb{B}^{kl}$ $\mathfrak{g} \leftarrow_{\$} \mathbb{G}^*$ $x \leftarrow_{\$} \mathbb{F}_p$ $pk1 = (hk, \mathfrak{g}, \llbracket x \rrbracket)$ $sk1 = (hk, \mathfrak{g}, x)$ $y \leftarrow_{\$} \mathbb{F}_p$ $pk2 = \llbracket y \rrbracket$ $sk2 = (y, \llbracket y \rrbracket)$ $pk = (pk1, pk2)$ $sk = (sk1, sk2)$	<b>Sig</b> ( $sk, m$ ) : $hk, \mathfrak{g}, x = sk.sk1$ $y, \llbracket y \rrbracket = sk.sk2$ $c = \text{CRH}(hk, \llbracket y \rrbracket \  m)$ $\sigma = y \bmod p$ $\sigma += c \cdot x \bmod p$ <b>return</b> $\sigma$	<b>Vf</b> ( $pk, m, \sigma$ ) : $hk, \mathfrak{g}, \llbracket x \rrbracket = pk.pk1$ $\llbracket y \rrbracket = pk.pk2$ $c = \text{CRH}(hk, \llbracket y \rrbracket \  m)$ <b>if</b> $\sigma \stackrel{?}{=} \llbracket y \rrbracket \otimes c \cdot \llbracket x \rrbracket$ <b>then</b> : <b>return</b> 1 <b>endif</b> <b>return</b> 0
---	--	---

Figure 3.9: One-time signature scheme from two tier Schnorr based signature scheme by Bellare and Shoup [BS07]

### 1226 3.4.2 Data types

1227 We now describe the data types associated with this signature scheme defined over  
1228 BN-254.

1229 **VK0tsDType** Denotes the verification key associated with the one-time signature scheme.

Field	Description	Value
$pk1$	Encoding of the scalar $x$ in the group	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^2$
$pk2$	Encoding of the scalar $y$ in the group	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^2$

Table 3.2: VK0tsDType data type

1230 **SK0tsDType** Denotes the signing key associated with the one-time signature scheme.

Field	Description	Value
$sk1$	Scalar element $x$	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}$
$sk21$	Scalar element $y$	$\mathbb{F}_{\mathbf{r}_{\text{BN}}}$
$sk22$	Encoding of the scalar $y$ in the group	$(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^2$

Table 3.3: SK0tsDType data type

1231 **Sig0tsDType** Denotes the signature data type associated with the one-time signature  
 1232 scheme. **Sig0tsDType** is an alias for  $(\mathbb{F}_{\mathbf{r}_{\text{BN}}})^2$ .

## 1233 3.5 Instantiate EncSch

1234 In this section we describe the instantiation of **EncSch** primitive introduced in Section 2.3.  
 1235 First, we present a general asymmetric encryption scheme called **DHAES** (Diffie-Hellman  
 1236 Asymmetric Encryption Scheme [ABR99]), which satisfies all the required security prop-  
 1237 erties for the in-band encryption scheme **EncSch** (see: Section 1.5). Then, we give details  
 1238 of the concrete algorithms used for the implementation.

### 1239 3.5.1 DHAES encryption scheme

1240 Given a symmetric encryption scheme **Sym**, a group defined by **SetupG**, a family of hash  
 1241 function  $\mathcal{H}^4$  and a message authentication scheme **MAC** as defined in Section 1.6, we  
 1242 define a **DHAES** scheme as the following public-key encryption scheme:

- 1243 • **Setup**, setup algorithm, takes as input a security parameter (in unary repre-  
 1244 sentation)  $1^\lambda$ . It runs  $\mathcal{H}.\text{Setup}$ , **SetupG** and returns public parameters  $pp =$   
 1245  $(hk, (q, \mathbb{G}, \mathbf{g}, +))$ .
- 1246 • **KGen**, key generation algorithm, takes as input public parameters  $pp$ . It samples  
 1247 at random  $v \leftarrow_{\$} [q]$  and returns a keypair  $(sk, pk) = (v, \llbracket v \rrbracket)$ .
- 1248 • **Enc**, encryption algorithm, takes as input public parameters  $pp$ , a message  $m$  and  
 1249 a public key  $pk$ . It runs **KGen** that returns an ephemeral keypair  $(esk, epk) =$   
 1250  $(u, \llbracket u \rrbracket)$ . Then, it computes a shared secret  $ss = H_{hk}(epk \parallel esk \cdot pk) = H_{hk}(epk \parallel esk \cdot$   
 1251  $pk)$ , parsed as  $ek \parallel mk^5$ . It computes  $ct_{\text{Sym}} = \text{Sym}.\text{Enc}(ek, m)$  and  $\tau = \text{MAC}.\text{Tag}(mk, ct_{\text{Sym}})$   
 1252 and finally outputs the ciphertext  $epk \parallel ct_{\text{Sym}} \parallel \tau$ .
- 1253 • **Dec**, decryption algorithm, takes as input public parameters  $pp$ , a private key  
 1254  $sk$  a ciphertext  $epk \parallel ct_{\text{Sym}} \parallel \tau$ . It computes  $ss = H_{hk}(epk \parallel sk \cdot epk)$  and parses it,  
 1255 as above, as  $ek \parallel mk$ . If **MAC** verification passes, i.e.  $\text{MAC}.\text{Vf}(mk, \tau) = 1$ , the  
 1256 algorithm returns  $\text{Sym}.\text{Dec}(ek, ct_{\text{Sym}})$  and  $\perp$  otherwise.

1257 The **DHAES** definition given above is an asymptotic adaptation of [ABR99, Section  
 1258 1.3].

### 1259 Inclusion of ephemeral key in hash input

1260 Given an ephemeral keypair  $(u_0, \llbracket u_0 \rrbracket)$ , If the group  $\langle \mathbf{g} \rangle$ , generated by **SetupG**, has com-  
 1261 posite order, then  $\llbracket u_0 \rrbracket$  is required to be part of the hash input because  $\llbracket u_0 v \rrbracket$  and  $\llbracket v \rrbracket$

<sup>4</sup>Here, we only consider fixed-length hash functions with  $h\text{InpLen}(\lambda) = 2g\text{Len}$  and  $h\text{Len}(\lambda) = k\text{Len}(\lambda) + m\text{Len}(\lambda)$  (see Section 1.6).

<sup>5</sup>Note that  $ek$  and  $mk$  must have same length.



together may not uniquely determine  $\llbracket u_0 \rrbracket$ . Equivalently, there may exist two values  $u_0$  and  $u_1$  such that  $u_0 \neq u_1$  and  $\llbracket u_0 v \rrbracket = \llbracket u_1 v \rrbracket$ . As a result, both  $u_0$  and  $u_1$  can be used to produce two different *valid* ciphertexts of the same plaintext  $m$ , under different ephemeral keys ( $\llbracket u_0 \rrbracket, \llbracket u_1 \rrbracket$ ). It is easy to show this, for example, in the multiplicative group  $\mathbb{Z}_p \setminus \{0\}$ , where  $p$  is a prime (see [ABR99, Section 3.1]). A scheme having such malleability property clearly cannot be proven IND-CCA2 secure: an attacker could easily win the related security game by altering the challenged ciphertext and query the decryption oracle that would not recognize that as a not allowed query. If the group has prime order this problem does not arise so only  $\llbracket u_0 v \rrbracket$  is required as input of the  $H$  function [ABR01, Section 3].

### 3.5.2 A DHAES instance

#### Curve25519

As cyclic group, we propose the use of a subgroup of Curve25519 described in [Ber06] and in [LHT16]. Curve25519 is a Montgomery elliptic curve [Mon87] defined by the equation  $y^2 = x^3 + 486662x^2 + x$  and coordinates on  $\mathbb{F}_p$ , where  $p$  is the prime number  $2^{255} - 19$ . It has a prime order subgroup of order  $2^{252} + 27742317777372353535851937790883648493$  and cofactor 8. Curve25519 comes with an efficient scalar multiplication denoted as X25519<sup>6</sup>. In a Diffie-Hellman-based scheme it allows to have 32-byte long public and private keys (given a point  $P = (x, y)$  only the  $x$  coordinate is actually used) and the 32-byte sequence representing 9 is specified as base point.

#### Efficiency and security of Curve25519

High-speed and timing-attack resistant implementations of X25519 are available and its security level is conjectured to be 128 bits [Ber06, Section 1]. However, combined attacks can lead to 124 bits of security (see [BL, Section “Twist Security”]). By design, Curve25519 is resistant to state-of-the-art attacks and satisfies all security criteria and principles listed in *Safecurves* [BL]<sup>7</sup>.

Interestingly, Curve25519 does not require *Public Key Validation*<sup>8</sup>, while we know that, on other curves, active attacks - consisting in sending malformed public keys - could be carried out by adversaries, to violate the confidentiality of private keys, e.g. [ABM<sup>+</sup>03]. However, Curve25519 specification mandates the *clamping* of private keys: that is, after the random sampling of 32 bytes, the user clears bits 0, 1 and 2 of the first byte, clears bit 7 and sets bit 6 of the last byte. The resulting 32 bytes are then

<sup>6</sup>X25519 is actually introduced in [LHT16] in order to avoid notation issues due to the use Curve25519 to indicate both curve and scalar multiplication as done in [Ber06]

<sup>7</sup>In this work, the authors take into account both Elliptic Curve Discrete Logarithm Problem (ECDLP) and Elliptic Curve Cryptosystems (ECC) security, that allows to have an overall evaluation of the security guarantees.

<sup>8</sup>Informally, it is a set of security checks that a user performs before using a not trusted public key (e.g. see [BCK<sup>+</sup>18])

1294 used as private key. This particular structure for private keys prevents various types of  
1295 attacks (see [Ber06, Section 3] for more details).

#### Note

Note that the *clamping* procedure is vital to ensure the security guarantees of the Curve25519 specification, and implementations **MUST** perform this exactly as described.

1296

## 1297 Chacha20

1298 ChaCha20 is a ARX-based<sup>9</sup> stream cipher introduced in [Ber08a]. It is an improved ver-  
1299 sion of Salsa20 [Ber08b] that won the *eSTREAM* challenge [est]. Compared with Salsa20,  
1300 it has been designed to improve diffusion per round, conjecturally increasing resistance  
1301 to cryptanalysis, while preserving time per round. It is considerably faster than AES in  
1302 software-only implementations and can be easily implemented to be timing-attacks resis-  
1303 tant. Several versions of the cipher can be used. The original paper presents ChaCha20  
1304 with a 128-bit key and 64-bit nonce/block count. However, the length of the key, nonce  
1305 and block count - which indicates how many chunks can be processed by using the same  
1306 key and nonce - can be modified depending on the application. In [LN18][Section 2.3], for  
1307 instance, the key is a 256-bit string, the nonce is a string of 96 bits and the block count  
1308 is encoded on a 32-bit word. This configuration allows to process around  $2^{32}$  blocks,  
1309 corresponding to roughly 256GB of data. We propose to use the same parameters in  
1310 Zeth.

$$\text{ChaCha20} : \mathbb{B}^{256} \times \mathbb{B}^{32} \times \mathbb{B}^{96} \times \mathbb{B}^* \rightarrow \mathbb{B}^*$$

## 1311 Security of Chacha

1312 Recent cryptanalysis results for ChaCha are available in [AFK<sup>+</sup>08, Ish12, SZFW12,  
1313 Mai16, CM16, CM17]: all of them makes use of advanced cryptanalysis techniques able  
1314 to perform key-recovery attacks only on reduced versions (6 and 7 rounds) of ChaCha.

#### Note

Importantly, the security properties of ChaCha rely on the fact that, for a given key, all blocks are processed with distinct values in the state words 12 to 15 (storing the counter and the nonce) [LN18, Section 2.3].

1315

## 1316 Poly1305

1317 Poly1305 [Ber05] is a high-speed message authentication code, easy to implement and  
1318 make side-channel attack resistant. It takes a 32-byte one-time key  $mk$  and a message  $m$

---

<sup>9</sup>Addition ( mod  $2^n$ ), Rotation, Xor

1319 and produces a 16-byte tag  $\tau$  that authenticates the message.  $mk$  must be unpredictable  
 1320 and it is represented as a couple  $(r, s)$ , where both components are given as a sequence  
 1321 of 16 bytes each. It can be generated by using pseudorandom algorithms: in [Ber05,  
 1322 Section 2], for example, AES and a nonce are used to generate  $s$ . The second part of  
 1323 the key,  $r$ , is expected to have a given form [Ber05, Section 2], and must be “clamped”  
 1324 as follows: top four bits of  $r[3]$ ,  $r[7]$ ,  $r[11]$ ,  $r[15]$  and bottom two bits of  $r[4]$ ,  $r[8]$ ,  $r[12]$   
 1325 are cleared (see also Section 3.5.3).

#### Note

Similarly to Curve25519, the *clamping* procedure here is essential to the security of the Poly1305 scheme. Implementations **MUST** ensure that this is performed correctly in order for all security guarantees to hold.

1326

1327 We refer to [LN18, Section 2.5, Section 3] for Tag and Vf implementations of Poly1305.

$$\begin{aligned}\text{Poly1305.Tag} &: \mathbb{B}_Y^{32} \times \mathbb{B}_Y^* \rightarrow \mathbb{B}_Y^{16} \\ \text{Poly1305.Vf} &: \mathbb{B}_Y^{32} \times \mathbb{B}_Y^{16} \times \mathbb{B}_Y^* \rightarrow \mathbb{B}\end{aligned}$$

### 1328 Security of Poly1305

1329 Citing Poly1305 [LN18, Section 4], “the Poly1305 authenticator is designed to ensure that  
 1330 forged messages are rejected with a probability of  $1 - (n/(2^{102}))$  for a  $16n$ -byte message,  
 1331 even after sending  $2^{64}$  legitimate messages, so it is SUF-CMA (strong unforgeability  
 1332 against chosen-message attacks)”.

### 1333 Blake2b-512

1334 Since we need a total of 64 bytes for the key material (32 for ChaCha20 and 32 for  
 1335 Poly1305) Blake2b512 can be used. ZCash protocol [ZCa19, Section 5.4.3], instead, makes  
 1336 use of Blake2b256 since a DHAES variant, denoted as ChaCha20-Poly1305, is adopted  
 1337 (see [LN18, Section 2.8]).

$$\text{Blake2b512} : \mathbb{B}^* \rightarrow \mathbb{B}_Y^{32}$$

### 1338 3.5.3 EncSch instantiation

In the following we instantiate EncSch as a DHAES scheme, detailing the KGen, Enc and Dec components. Before, we introduce some required constant values:

$$\begin{aligned} \text{ESKBYTELEN} &= 32 \\ \text{EPKBYTELEN} &= 32 \\ \text{NOTEBYTELEN} &= (\text{PRFADDRROUTLEN} + \text{RTRAPLEN} + \text{ZVALUELEN} + \text{PRFRHOOUTLEN})/\text{BYTELEN} \\ \text{SYMKEYBYTELEN} &= 32 \\ \text{MACKEYBYTELEN} &= 32 \\ \text{KDFDIGESTBYTELEN} &= \text{SYMKEYBYTELEN} + \text{MACKEYBYTELEN} \\ \text{CTBYTELEN} &= \text{EPKBYTELEN} + \text{NOTEBYTELEN} + \text{TAGBYTELEN} \\ \text{TAGBYTELEN} &= 16 \\ \text{CHACHANONCEVALUE} &= 0^{32} \\ \text{CHACHABLOCKCOUNTERVALUE} &= 0^{96} \end{aligned}$$

#### 1339 EncSch.KGen

1340 The keypair generation  $(sk, pk)$  is defined as:

- 1341 • Randomly sample a sequence of ESKBYTELEN bytes and assign to  $sk$ .
- Clamp  $sk$  as follows:

$$\begin{aligned} sk[0] &= sk[0] \& 0xF8 \\ sk[31] &= sk[31] \& 0x7F \\ sk[31] &= sk[31] \mid 0x40 \end{aligned}$$

1342 where  $|$  and  $\&$  denotes, respectively, OR and AND binary operators between bit  
1343 strings of same the length.<sup>10</sup>

- 1344 • Compute  $pk = \text{X25519}(sk, 0x09)$ .
- 1345 • Return  $(sk, pk) \in \mathbb{B}_{\mathbb{Y}}^{\text{ESKBYTELEN}} \times \mathbb{B}_{\mathbb{Y}}^{\text{EPKBYTELEN}}$

#### 1346 EncSch.Enc

1347 The encryption, on inputs  $(pk, m) \in \mathbb{B}_{\mathbb{Y}}^{\text{EPKBYTELEN}} \times \mathbb{B}_{\mathbb{Y}}^{\text{NOTEBYTELEN}}$ , is defined as follows:

- 1348 1. Generate an ephemeral Curve25519 keypair  $(esk, epk) \in \mathbb{B}_{\mathbb{Y}}^{\text{ESKBYTELEN}} \times \mathbb{B}_{\mathbb{Y}}^{\text{EPKBYTELEN}}$   
1349 (as above).

---

<sup>10</sup>E.g Given two bytes 0x15 and 0x03 then  $0x15|0x03 = 0x17$  and  $0x15\&0x03 = 0x01$ .

2. Compute the shared secret<sup>11</sup>  $ss \in \mathbb{B}_Y^{\text{EPKBYTELEN}}$ :

$$ss = \text{X25519}(esk, pk) \in \mathbb{B}_Y^{\text{EPKBYTELEN}}$$

3. Generates a session key:

$$\text{Blake2b512}(\text{encTag} \parallel epk \parallel ss) \in \mathbb{B}_Y^{\text{KDFDIGESTBYTELEN}}$$

where  $\text{encTag} = 0x5A \parallel 0x65 \parallel 0x74 \parallel 0x68 \parallel 0x45 \parallel 0x6E \parallel 0x63$ , that is the UTF-8 encoding of “ZethEnc” string (used for domain separation purposes). The result, then, is parsed as follows:

$$\begin{aligned} ek &= \text{Blake2b512}(\text{encTag} \parallel epk \parallel ss)[\text{SYMKEYBYTELEN} - 1] \\ mk &= \text{Blake2b512}(\text{encTag} \parallel epk \parallel ss)[\text{SYMKEYBYTELEN} : \text{SYMKEYBYTELEN} + \text{MACKEYBYTELEN} - 1]. \end{aligned}$$

4. Encrypt the confidential data:

$$ct_{\text{Sym}} = \text{ChaCha20}(ek, \text{CHACHABLOCKCOUNTERVALUE}, \text{CHACHANONCEVALUE}, m) \in \mathbb{B}^{\text{NOTEBYTELEN} * \text{BYTELEN}}$$

1350 **Remark 13.** Formally speaking we should have written  $ct_{\text{Sym}} \in \mathbb{B}^n$ , where  $n$  is  
 1351 the length of binary representation of the encrypted message  $m$ . In Zeth how-  
 1352 ever, the only data encrypted are the notes. As such the size of the plaintexts is  
 1353  $\text{NOTEBYTELEN} * \text{BYTELEN}$  bits.

1354 **Remark 14.** In the following, we omit the explicit conversion from  $\mathbb{B}^n$  to  $\mathbb{B}_Y^{\lceil n/\text{BYTELEN} \rceil}$   
 1355 when passing the output of ChaCha20 to the Poly1305 algorithms.

5. Randomly generate  $(r, s) \in \mathbb{B}_Y^{\text{MACKEYBYTELEN}/2} \times \mathbb{B}_Y^{\text{MACKEYBYTELEN}/2}$  and clamp it:

$$\begin{aligned} r[3] &= r[3] \ \& \ 0x0F \\ r[7] &= r[7] \ \& \ 0x0F \\ r[11] &= r[11] \ \& \ 0x0F \\ r[15] &= r[15] \ \& \ 0x0F \\ r[4] &= r[4] \ \& \ 0xFC \\ r[8] &= r[8] \ \& \ 0xFC \\ r[12] &= r[12] \ \& \ 0xFC \end{aligned}$$

6. Generate the related tag:

$$\tau = \text{Poly1305.Tag}(mk, ct_{\text{Sym}}) \in \mathbb{B}_Y^{\text{TAGBYTELEN}}.$$

7. Create the asymmetric ciphertext as:

$$ct = epk \parallel ct_{\text{Sym}} \parallel \tau \in \mathbb{B}_Y^{\text{CTBYTELEN}}.$$

- 1356 8. Return  $ct$ . As consequence  $\text{ENCZETHNOTELEN} = \text{CTBYTELEN} * \text{BYTELEN}$  bits.

---

<sup>11</sup>We assume here that  $esk$  has been clamped as discussed in Section 3.5.2

1357 EncSch.Dec

1358 The decryption, on inputs  $(sk, ct) \in \mathbb{B}_{\mathbb{Y}}^{\text{ESKBYTELEN}} \times \mathbb{B}_{\mathbb{Y}}^{\text{CTBYTELEN}}$ , is defined as follows:

1. Parse the ciphertext  $ct$  as:

$$\begin{aligned} epk &\leftarrow ct[: \text{EPKBYTELEN} - 1] \\ ct_{\text{Sym}} &\leftarrow ct[\text{EPKBYTELEN} : \text{EPKBYTELEN} + \text{NOTEBYTELEN} - 1] \\ \tau &\leftarrow ct[\text{EPKBYTELEN} + \text{NOTEBYTELEN} : \text{EPKBYTELEN} + \text{NOTEBYTELEN} + \text{TAGBYTELEN} - 1] \end{aligned}$$

2. Recover the shared secret

$$ss = \text{X25519}(sk, epk).$$

3. Compute the  $ek||mk$

$$\begin{aligned} ek &= \text{Blake2b512}(\text{encTag}||epk||ss)[: \text{SYMKEYBYTELEN} - 1] \\ mk &= \text{Blake2b512}(\text{encTag}||epk||ss)[\text{SYMKEYBYTELEN} : \text{SYMKEYBYTELEN} + \text{MACKEYBYTELEN} - 1]. \end{aligned}$$

4. Verify that the ciphertext has not been forged:

$$\text{Poly1305.Vf}(mk, \tau, ct_{\text{Sym}})$$

5. (If the MAC verifies) decrypt:

$$m = \text{ChaCha20.Dec}(ek, \text{CHACHABLOCKCOUNTERVALUE}, \text{CHACHANONCEVALUE}, ct_{\text{Sym}})$$

1359 6. Return  $m$ .

### 1360 3.5.4 Security requirements satisfaction

1361 DHAES has already been proved to be IND-CCA2 secure (see [ABR99, Section 3.5, The-  
1362 orem 6])<sup>12</sup> and to the best of our knowledge there is no paper showing IK-CCA security.  
1363 The only proof we have found is related to DHIES scheme [ABN10], that is a prime order  
1364 group version of DHAES. In the following, we provide a proof for IK-CCA security of  
1365 DHAES by adapting that proof to our case.

**Theorem 2** (IK-CCA of DHAES). *Let DHAES be the asymmetric encryption scheme as defined above. Let  $\mathcal{A}$  be an adversary for the IK-CCA game, then there exists a HDHI adversary  $\mathcal{B}$  of  $(\mathcal{H}, \text{SetupG})$  and a SUF-CMA adversary  $\mathcal{C}$  of MAC such that*

$$\text{Adv}_{\text{DHAES}, \mathcal{A}}^{\text{ik-cca}}(\lambda) \leq 2 \cdot \text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{B}}^{\text{hdhi}}(\lambda) + \text{Adv}_{\text{MAC}, \mathcal{C}}^{\text{suf-cma}}(\lambda).$$

1366 The adversaries  $\mathcal{B}$  and  $\mathcal{C}$  have the same running time as  $\mathcal{A}$ <sup>13</sup>.

<sup>12</sup>Specifically, it holds if Sym is IND-CPA secure, H is HDHI secure and MAC is SUF-CMA secure.

<sup>13</sup>In order to give an asymptotic version of the theorem, the number of queries  $q$  has been substituted by the fact of considering PPT adversaries.

1367 *Proof.* As already mentioned, DHAES is similar to DHIES scheme, except for the under-  
 1368 lying group and the way the symmetric keys are constructed. As consequence, IK-CCA  
 1369 property for DHAES can be showed similarly as done in [ABN10, Theorem 6.2]. More  
 1370 precisely, they show that one can construct from an attacker  $\mathcal{A}$  for the IK-CCA game  
 1371 two attackers  $\mathcal{B}$  and  $\mathcal{C}$  for the ODH and SUF-CMA games. Actually, they make use of  
 1372 a  $\bar{\mathcal{B}}$  attacker for the ODH2 game [ABN10, Figure 20] and then apply [ABN10, Lemma  
 1373 6.1] to obtain an attacker  $\mathcal{B}^{14}$  in the ODH game. We adopt a similar strategy, working  
 1374 with HDHI, HDHI2 and Lemma 1.

Let  $\mathcal{A}$  be an attacker for the IK-CCA game, and let  $\bar{\mathcal{B}}$  be an attacker for the HDHI2  
 game described in Fig. 3.10. We show that,

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \bar{\mathcal{B}}}^{\text{hdhi2}}(\lambda) = |\Pr[\text{IK-CCA}^{\mathcal{A}}(\lambda) = 1] + \Pr[\text{G}_0^{\mathcal{A}}(\lambda) = 1] - 1|$$

1375 where  $\text{G}_0$  is the security game described in Fig. 3.11.

1376 Given an HDHI2 challenge  $(\llbracket u \rrbracket, \llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket, w_{b_2,0}, w_{b_2,1})$ , an adversary  $\bar{\mathcal{B}}$  samples  $b \leftarrow_{\$} \{0, 1\}$   
 1377 and runs  $\mathcal{A}$  on  $\llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket$  (note that  $b_2$  is the random bit chosen by the  $\bar{\mathcal{B}}$  challenger in the  
 1378 HDHI2 game).  $\bar{\mathcal{B}}$  constructs oracles  $\text{O}^{\text{Dec}_{sk_i}}$  where the queries  $(\mathfrak{r} \parallel ct_{\text{Sym}} \parallel \tau)$  are processed  
 1379 as follows: if  $\mathfrak{r} \neq \llbracket u \rrbracket$ , then  $\bar{\mathcal{B}}$  queries related HDHI2 oracle to obtain  $ek \parallel mk \leftarrow \text{O}^{\text{HDHI}_{v_i}}(\mathfrak{r})$   
 1380 (see Fig. 3.10). If  $\mathfrak{r} = \llbracket u \rrbracket$ ,  $w_{b_2,i}$  is parsed as  $ek \parallel mk$ . In both cases, it checks that  
 1381  $\text{MAC.Vf}(mk, ct_{\text{Sym}}, \tau) = 1$  and, if so, returns  $m \leftarrow \text{Sym.Dec}(ek, ct_{\text{Sym}})$ . We note that  
 1382  $\mathcal{A}$  cannot query the challenged ciphertext.  $\bar{\mathcal{B}}$  returns 0 if and only if  $b = \tilde{b}$ . It easy to  
 1383 see that if  $b_2$  is equal to 0, then all symmetric encryption and MAC keys used for the  
 1384 challenge ciphertext  $(\mathfrak{r}^* \parallel ct_{\text{Sym}}^* \parallel \tau^*)$  and decryption responses are exactly as in a DHAES  
 1385 game.

1386 If  $b_2 = 1$ , then  $w_{1,0}$  and  $w_{1,1}$  are random strings and the challenge ciphertext and  
 1387 decryption responses are given as in the  $\text{G}_0$  game described in Fig. 3.11.

So we get,

$$\Pr[\text{HDHI2}^{\bar{\mathcal{B}}}(\lambda) = 1] = \frac{1}{2} \cdot \Pr[\text{IK-CCA}^{\mathcal{A}}(\lambda) = 1] + \frac{1}{2} \cdot \Pr[\text{G}_0^{\mathcal{A}}(\lambda) = 1]$$

1388 And from the definition of HDHI2 advantage, we have,

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \bar{\mathcal{B}}}^{\text{hdhi2}}(\lambda) = |\Pr[\text{IK-CCA}^{\mathcal{A}}(\lambda) = 1] + \Pr[\text{G}_0^{\mathcal{A}}(\lambda) = 1] - 1|$$

1389 At this point, we can conclude as done in [ABN10, Theorem 6.2], with the only  
 1390 difference of applying Lemma 1 instead of [ABN10, Lemma 6.1] and by defining a game  
 1391  $\text{G}_1$  that is *identical until bad*<sup>15</sup>  $\text{G}_0$  defined in Fig. 3.11.  $\square$

<sup>14</sup>Note that in [ABN10] the IK-CCA game is a particular case of the AI-CCA game that requires two  
 input messages in the LR query. In order to reason only about the key-privacy, the two messages  $m_0$   
 and  $m_1$  are constrained to be equal.

<sup>15</sup>Games  $\text{G}_i$  and  $\text{G}_j$  are said to be *identical until bad* if they differ only in statements that follow the  
 setting of the **bad** variable to *True*. **bad** is initialized with *False*

Adversary $\bar{\mathcal{B}}(\llbracket u \rrbracket, \llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket, w_{b_2,0}, w_{b_2,1})$	$\bar{\mathcal{B}}$ simulation of $\mathcal{O}^{\text{Dec}_{sk_i}}(\mathfrak{r} \parallel ct_{\text{Sym}} \parallel \tau)$
$b \leftarrow_{\$} \{0, 1\}$	<b>if</b> $\mathfrak{r} \neq \llbracket u \rrbracket$
$(m, state) \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(\llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket)$	$ek \parallel mk \leftarrow \mathcal{O}^{\text{HDH}_{v_i}}(\mathfrak{r})$
$ek \parallel mk \leftarrow w_{b_2,b}$	<b>else</b>
$\mathfrak{r}^* \leftarrow u$	$ek \parallel mk \leftarrow w_{b_2,i}$
$ct_{\text{Sym}}^* \leftarrow \text{Sym.Enc}(ek, m)$	<b>fi</b>
$\tau^* \leftarrow \text{MAC.Tag}(mk, ct_{\text{Sym}}^*)$	<b>if</b> $\text{MAC.Vf}(mk, ct_{\text{Sym}}, \tau) = 1$
$\tilde{b} \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(\mathfrak{r}^* \parallel ct_{\text{Sym}}^* \parallel \tau^*, state)$	<b>return</b> $\text{Sym.Dec}(ek, ct_{\text{Sym}})$
<b>return</b> $\tilde{b} = b$	<b>else</b>
	<b>return</b> $\perp$
	<b>fi</b>

Figure 3.10: Description of the adversary  $\bar{\mathcal{B}}$  for HDH12, simulating DHAES game for  $\mathcal{A}$ .

$\mathcal{G}_0(\lambda)$	Oracle $\mathcal{O}^{\overline{\text{Dec}_{sk_i}}}(\mathfrak{r} \parallel ct_{\text{Sym}} \parallel \tau)$
$(q, \mathbb{G}, \mathfrak{g}, +) \leftarrow \text{SetupG}(1^\lambda)$	<b>if</b> $\mathfrak{r} = \mathfrak{r}^*$
$(sk_0, pk_0), (sk_1, pk_1) \leftarrow_{\$} \text{KGen}(1^\lambda)$	$m \leftarrow \perp$
$\mathfrak{r}^* \leftarrow_{\$} \mathbb{G}$	<b>if</b> $\text{MAC.Vf}(mk^*, ct_{\text{Sym}}, \tau) = 1$
$ek^* \leftarrow_{\$} \{0, 1\}^{kLen}$	<b>bad</b> $\leftarrow \text{true}$
$mk^* \leftarrow_{\$} \{0, 1\}^{mLen}$	$m \leftarrow \text{Sym.Dec}(ek^*, ct_{\text{Sym}})$
$(m, state) \leftarrow \mathcal{A}^{\mathcal{O}^{\overline{\text{Dec}_{sk_0}}}, \mathcal{O}^{\overline{\text{Dec}_{sk_1}}}}(pk_0, pk_1)$	<b>fi</b>
$b \leftarrow_{\$} \{0, 1\}$	<b>else</b>
$ct_{\text{Sym}}^* \leftarrow \text{Sym.Enc}(ek^*, m)$	$m \leftarrow \text{Dec}(sk_i, \mathfrak{r} \parallel ct_{\text{Sym}} \parallel \tau)$
$\tau^* \leftarrow \text{MAC.Tag}(mk^*, ct_{\text{Sym}}^*)$	<b>fi</b>
$\tilde{b} \leftarrow \mathcal{A}^{\mathcal{O}^{\overline{\text{Dec}_{sk_0}}}, \mathcal{O}^{\overline{\text{Dec}_{sk_1}}}}(\mathfrak{r}^* \parallel ct_{\text{Sym}}^* \parallel \tau^*, state)$	<b>return</b> $m$
<b>return</b> $\tilde{b} = b$	

Figure 3.11:  $\mathcal{G}_0$  game and related decryption oracles for Theorem 2.



### 1392 3.5.5 Final notes and observations

1393 In this section we list some notes regarding Zcash approach (see [ZCa19, Section 8.7])  
 1394 and other observations:

- 1395 • *Key derivation parameters:* in DHAES construction, the only required input vari-  
 1396 ables are the shared secret  $ss$  and  $epk$ . In Sprout, additional parameters were  
 1397 added (i.e.  $h_{sig}$ ,  $pk_{enc}$  and a counter  $i$ ) (see [ZCa19, 5.4.4.2]): they state that  $h_{sig}$   
 1398 was used in order to get a different random extractor for each joinsplit transfer  
 1399 in order to limit the degradation of the security and weaken assumption on the  
 1400 hash. The authors believed, about the use of long-standing public key  $pk_{enc}$ , that  
 1401 it might be necessary for IND-CCA2 security and for post-quantum privacy (in the  
 1402 case where the quantum attacker does not have the public key) [zca]. None of  
 1403 these additional components are used any longer starting from the Sapling release  
 1404 (see [ZCa19, 5.4.4.4]). To the best of our knowledge there is no formal reason  
 1405 to use note counter  $i$  into the KDF: an explanation could be to avoid that same  
 1406 session key is reused, but this should be fine as soon as a different nonce or block  
 1407 counter is used for the symmetric cipher (actually this is already mandated in the  
 1408 case  $epk$ , reuse as described below).
- 1409 • *Reuse of ephemeral keys  $epk$ :* ZCash reuses the same ephemeral keys  $epk$  (and  
 1410 different nonces) for two ciphertexts in a joinsplit description, claiming that this  
 1411 does not affect the security of the scheme as soon as the HDHI assumption of the  
 1412 DHAES security proof is adapted. Note that the proof they refer is related to the  
 1413 IND-CCA2 notion.
- 1414 • Note that in Zcash Sprout and Sapling being able to break the Elliptic Curve  
 1415 Diffie-Hellman Problem on Curve25519 or Jubjub would not help to decrypt the  
 1416 transmitted notes ciphertext unless the receiver  $pk_{enc}$  is known or guessed. On the  
 1417 other hand, having  $pk_{enc}$  into the hash (as used in Sprout) may violate in principle  
 1418 the key-privacy of the encryption scheme. For these reasons, we underline that  
 1419 the protocol should enforce a mechanism that does not reveal users public keys to  
 1420 increase the security.
- 1421 • In [ABN10], the concept of *robustness* for an asymmetric encryption scheme is  
 1422 introduced: it formalizes the infeasibility to produce a ciphertext valid under two  
 1423 different public encryption keys. We note that this is particularly useful for Zeth  
 1424 since only the intended receiver will be able to decrypt the encrypted note. The  
 1425 definition, actually, is more general, since it covers also the case when a decryption  
 1426 is successful but returns an incorrect plaintext. This prevent situations where  
 1427 a user, scanning the Mixer logs for incoming transactions, gets a false positive  
 1428 decryption and stores garbage notes.

## Note

We note however, that the “false-positive” situation above can be prevented by relying on a weaker notion of robustness called *collision-freeness* [Moh10]. In fact, as described in Section 2.6, the procedure to receive a *ZethNote* requires to decrypt the ciphertext emitted by the **Mixer**, and then to verify that the recovered plaintext is the opening of a commitment in the Merkle tree. As such, since the *collision-freeness* of the encryption ensures that plaintexts recovered under different keys are different (i.e. “do not produce a collision”), then we know that plaintexts recovered by parties who are not the intended recipient will fail the “commitment opening verification”, leading the payment to be rejected, and solving the false-positive issue aforementioned.

In [ABN10, Section 6], the authors prove that DHIES can be made strongly robust. The proof can be easily adapted to work with DHAES.

- *No public key validation for X25519*: community of cryptographers have been discussing about the absence of any mandated public key validation or checks on the result of X25519. For example, in [LHT16, Section 6.1], an optional zero check is introduced in order to assure that the result of X25519 is not 0: this avoids that one of the two party can force the result of the key-exchange by using a small order point as public key. This property is generally defined as *contributory behavior*, that is, none of the parties is able to force the output of a key exchange. However, protocols do not have all the same security requirements and adding default checks in the Curve25519 specifications would be superfluous in most cases and would add complexities that Bernstein has deliberately chosen to avoid (*simple implementation principle*). More importantly, Diffie-Hellman does not require *contributory behavior* property [Per17]: modern view is that the only requirements are key indistinguishability and, in case of an active attacker, that the output of the key exchange should not produce a low-entropy function of the honest party’s private key (e.g. small-subgroup and invalid-curve attacks). Since these two properties are considered satisfied by Curve25519, there is no need to add extra checks to the Curve25519 specification. We conclude by observing that in Sprout release, Zcash protocol does not specify any point validation and makes use only of the private key clamping to keep Diffie-Hellman key exchange secure.

## 3.6 Instantiate ZkSnarkSch

The Groth’s zkSNARK Groth16 [Gro16] is the most efficient (in terms of the proof size and proof and verification cost) known zkSNARK for QAPs, thus one of the most efficient zkSNARK for proving statements on arithmetic circuits. Below we present Groth16’s key generation, prover, verifier, and simulator algorithms, adjusted as described in [BGM17]

1456 to further reduce the size of  $srs$  and proofs, and to make the **KGen** algorithm more  
 1457 amenable to implementation as a multi-party computation.

1458 In what follows, let the number  $constNo$  of constraints in the relation  $\mathbf{R}$  be fixed.  
 1459 Without loss of generality we consider  $constNo$  to be an *upper bound* on the number of  
 1460 constraints in the  $\mathbf{R}$  parameter, and assume that there exists some  $constNo$ -th root of  
 1461 unity  $\omega \in \mathbb{F}_{\mathbf{r}_{\text{BN}}}$ . Define  $\ell_i(X)$  to be the  $i$ -th Lagrange polynomial of degree  $(constNo - 1)$   
 1462 over the set  $\{\omega^i\}_{i \in [constNo]}$ , and let  $\ell(X)$  be the unique non-zero polynomial of degree  
 1463  $constNo$  that satisfies  $\ell(\omega^i) = 0$  for all  $i \in [constNo]$ .

1464 We note that the requirement that there exists a  $constNo$ -th root of unity  $\omega$  imposes  
 1465 a restriction on the maximum number of constraints in  $\mathbf{R}$  that the scheme can support.  
 1466 In particular case of  $\omega \in \mathbb{F}_{\mathbf{r}_{\text{BN}}}$ , the restriction becomes  $constNo \leq 2^{28}$ .

1467 **KGen**( $\mathbf{R}, 1^\lambda$ ):

- 1468 i. Pick trapdoor  $td = (\tau, \alpha, \beta, \delta) \leftarrow (\mathbb{Z}_p^* \setminus \{\omega^{i-1}\}_{i=1}^{constNo}) \times (\mathbb{Z}_p^*)^3$ ;  
 ii. For  $j \in \{1, \dots, inpNo\}$ , let

$$\begin{aligned} u_j(\tau) &= \sum_{i=1}^{constNo} U_{ij} \ell_i(\tau), \\ v_j(\tau) &= \sum_{i=1}^{constNo} V_{ij} \ell_i(\tau), \\ w_j(\tau) &= \sum_{i=1}^{constNo} W_{ij} \ell_i(\tau); \end{aligned}$$

- iii. Set

$$\begin{aligned} srs_{\mathbf{P}} &\leftarrow \left( \llbracket \alpha \rrbracket_1, \llbracket \beta \rrbracket, \llbracket \delta \rrbracket, \{\llbracket u_j(\tau) \rrbracket_1\}_{j=1}^{inpNo}, \{\llbracket v_j(\tau) \rrbracket\}_{j=0}^{inpNo}, \right. \\ &\quad \left. \{\llbracket (u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau))/\delta \rrbracket_1\}_{j=inpNoPrim+1}^{inpNo}, \right. \\ &\quad \left. \{\llbracket \tau^i \ell(\tau)/\delta \rrbracket_1\}_{i=0}^{constNo-2} \right) \\ srs_{\mathbf{V}} &\leftarrow \left( \llbracket \alpha \rrbracket_1, \llbracket \beta \rrbracket_2, \llbracket \delta \rrbracket_2, \{\llbracket \beta u_j(\tau) + \alpha v_j(\tau) + w_j \rrbracket_1\}_{j=0}^{inpNoPrim} \right) \\ srs &\leftarrow (srs_{\mathbf{P}}, srs_{\mathbf{V}}) \end{aligned}$$

1469 **return**  $srs, td$

1470 **P**( $\mathbf{R}, srs_{\mathbf{P}}, prim = (inp_j)_{j=1}^{inpNoPrim}, aux = (inp_j)_{j=inpNoPrim+1}^{inpNo}$ ):

- i. Define

$$a^\dagger(X) = \sum_{j=1}^{inpNo} inp_j u_j(X), \quad b^\dagger(X) = \sum_{j=1}^{inpNo} inp_j v_j(X), \quad c^\dagger(X) = \sum_{j=1}^{inpNo} inp_j w_j(X);$$

- 1471 ii. Define the polynomial  $h(X) = (a^\dagger(X)b^\dagger(X) - c^\dagger(X))/\ell(X)$  and compute the  
 1472 coefficients  $\{h_i\}_{i=0}^{constNo-2}$  of  $h$ , such that  $h(X) = \sum_{i=0}^{constNo-2} h_i X^i$ .

1473     iii.  $r_a \leftarrow_{\$} \mathbb{Z}_p$ ;  
 1474     iv.  $r_b \leftarrow_{\$} \mathbb{Z}_p$ ;  
        v. Compute proof elements:

$$\begin{aligned}
 \mathbf{a} &\leftarrow \sum_{j=1}^{inpNo} inp_j \llbracket u_j(\tau) \rrbracket_1 + \llbracket \alpha \rrbracket_1 + r_a \llbracket \delta \rrbracket_1 \\
 \mathbf{b} &\leftarrow \sum_{j=1}^{inpNo} inp_j \llbracket v_j(\tau) \rrbracket_2 + \llbracket \beta \rrbracket_2 + r_b \llbracket \delta \rrbracket_2 \\
 \mathbf{c} &\leftarrow r_b \mathbf{a} + r_a \left( \sum_{j=1}^{inpNo} inp_j \llbracket v_j(\tau) \rrbracket_1 + \llbracket \beta \rrbracket_1 \right) + \\
 &\quad \sum_{j=inpNoPrim+1}^{inpNo} inp_j \left\llbracket \frac{u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau)}{\delta} \right\rrbracket_1 + \\
 &\quad \sum_{i=0}^{constNo-2} h_i \llbracket \tau^i \ell(\tau) / \delta \rrbracket_1
 \end{aligned}$$

1475     **return**  $\pi \leftarrow (\mathbf{a}, \mathbf{b}, \mathbf{c})$ ;

1476      $\mathbf{V}(\mathbf{R}, srs_V, prim = (inp_j)_{j=1}^{inpNoPrim}, \pi)$ :

       i. Check that:

$$\begin{aligned}
 \mathbf{a} \bullet \mathbf{b} &= \mathbf{c} \bullet \llbracket \delta \rrbracket_2 \\
 &+ \left( \sum_{j=1}^{inpNoPrim} inp_j \llbracket u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau) \rrbracket_1 \right) \bullet \llbracket 1 \rrbracket_2 \\
 &+ \llbracket \alpha \rrbracket_1 \bullet \llbracket \beta \rrbracket_2
 \end{aligned}$$

1477     Note that  $\llbracket \alpha \rrbracket_1$  and  $\llbracket \beta \rrbracket_2$  are stored individually and used by the prover to re-  
 1478     compute  $\llbracket \alpha\beta \rrbracket_T$  seemingly redundantly. This is required in order to leverage the  
 1479     pairing check functionality built in to **Ethereum**, which accepts a sequence of tuples  
 1480     in  $\mathbb{G}_1 \times \mathbb{G}_2$  and returns true if and only if the product of the resulting pairings  
 1481     equals  $\llbracket 1 \rrbracket_T$ .

1482      $\mathbf{Sim}(\mathbf{R}, srs, td, prim)$ :

1483     i. Sample  $\mathbf{a}^* \leftarrow_{\$} \mathbb{Z}_p$ ;  $\mathbf{b}^* \leftarrow_{\$} \mathbb{Z}_p$ ;

ii. Compute proof elements:

$$\begin{aligned}
\mathbf{a} &\leftarrow \llbracket \mathbf{a}^* \rrbracket_1 + \llbracket \alpha \rrbracket_1 \\
\mathbf{b} &\leftarrow \llbracket \mathbf{b}^* \rrbracket_1 + \llbracket \beta \rrbracket_2 \\
\mathbf{c} &\leftarrow \frac{1}{\delta} \cdot \left[ \mathbf{a}^* \mathbf{b}^* \llbracket 1 \rrbracket_1 + \mathbf{a}^* \llbracket \beta \rrbracket_1 + \mathbf{b}^* \llbracket \alpha \rrbracket_1 \right. \\
&\quad \left. - \sum_{j=1}^{inpNoPrim} inp_j \llbracket u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau) \rrbracket_1 \right]
\end{aligned}$$

1484

**return**  $\pi \leftarrow (\mathbf{a}, \mathbf{b}, \mathbf{c});$

## Chapter 4

# Implementation considerations and optimizations

### 4.1 Client Security Considerations

In this section we consider some details of client *wallet* software that manages user's private and public keys, **Zeth** notes, and interacts with the **Mixer** contract.

Due to the processing and storage requirements involved, we consider it impractical for all **Zeth** client implementations to include dedicated Ethereum nodes (miners or archivers) to be run on the same host as the wallet. Therefore, in order to interact with the Ethereum network, wallet software must communicate with external Ethereum P2P nodes via their RPC channel, and must assume that these nodes are completely of the wallet's control. From a security standpoint, connected Ethereum nodes should therefore be considered untrusted, and in particular the details of all RPC calls and responses should be considered publicly visible. Note that even if the connected Ethereum node itself is not malicious, 3rd parties able to see network traffic may also be able to gain an insight into the RPC communication of a specific **Zeth** client.

#### Note

Note that there are several possible models besides the fully untrusted Ethereum node. Organizations or individuals could host one or more “trusted” Ethereum nodes, which clients can securely connect to (if they trust the host). This centralization would represent a security trade-off. From the point of view of clients it would create a single point of trust, and for potential malicious observers or attackers it would represent a valuable target.

In what follows we focus on preventing data leaks through network traffic. We do not consider adversaries with physical access to the machine running the wallet (see: Appendix C).

#### 1505 4.1.1 Syncing and Waiting

1506 **Zeth** clients must periodically synchronize with the latest state of the blockchain. This  
1507 is necessary to keep track of the data held by the **Mixer** contract, and to detect notes  
1508 received by the user of the wallet, storing them for future transactions.

1509 Clients should synchronize with Ethereum nodes in such a way that information is  
1510 not leaked. As such:

- 1511 1. Clients **MUST** use consensus evidence and block headers to verify all data they  
1512 receive from Ethereum nodes.
- 1513 2. Clients **MUST** locally store all **Mixer** state they require in order to function.
- 1514 3. Clients **MUST** obtain all such information by “synchronizing” with the Ethereum  
1515 blockchain and parsing relevant events emitted by **Mixer**. Clients **MUST NOT**  
1516 query the **Mixer** state via RPC.
- 1517 4. Clients **SHOULD** take steps to avoid being identified while synchronizing (see: Ap-  
1518 pendix C.2. For example, clients **SHOULD** vary the set of Ethereum nodes that they  
1519 connect to, and **SHOULD NOT** always sync from the block following the last one that  
1520 they processed.
- 1521 5. Clients **SHOULD NOT** re-request blocks or transaction receipts that are of particular  
1522 interest to them.
- 1523 6. Clients **SHOULD NOT** make any RPC calls or change their externally visible behavior  
1524 in response to blocks or transaction receipts that are of interest to them.

#### 1525 Use of Contract Queries

1526 We suggest that clients **SHOULD NOT** directly query the contract state, for the reasons  
1527 discussed in Appendix C.2 and Appendix C.3 (and consequently, Section 4.2 suggests  
1528 that the **Mixer** contract should, as far as possible, not expose public methods). The  
1529 operation of the protocol does not require such queries, and they introduce a risk that  
1530 some client implementations will leak information by using them.

1531 Implementors may have application-specific requirements for public **Mixer** methods,  
1532 but should be aware of possible risks associated with introducing them. See: Section 4.1.4  
1533 for an possible example of this.

#### 1534 4.1.2 Note Management

1535 **Mix** calls on the **Mixer** contract emit log events containing new commitment values,  
1536 nullifiers, the new Merkle root and the secret data for new notes (encrypted using a key  
1537 derived from the recipients public key). As clients synchronize with the latest state of  
1538 the blockchain, they **MUST** read these events and correctly process the data within.

- 1539 1. Clients **MUST** process the `MixEventDType` event for every Mix transaction, in the  
1540 order in which they are appear in the blockchain.
- 1541 2. Clients implementing spending functionality **MUST** use commitment values in events  
1542 to track the state of the Merkle tree. The Merkle tree state will be used to generate  
1543 Merkle paths for future transactions, and **MUST** be made available to the client  
1544 without the need to query the contract. (Note that not all commitments must  
1545 necessarily be persisted - see Section 4.3).
- 1546 3. Clients that can receive notes **MUST** attempt to decrypt the ciphertexts for every  
1547 transaction (see Item 2).
- 1548 4. Clients **MUST NOT** perform any network-related action, including closing the RPC  
1549 connection, dependent on successful / unsuccessful decryption of ciphertexts (see  
1550 Appendix C.3).
- 1551 5. Clients that can receive notes **MUST** attempt to parse any successfully decrypted  
1552 plaintext (that is, ensure it is well-formed as in Item 3a).
- 1553 6. Clients **MUST NOT** perform any network-related action, including closing the con-  
1554 nection, dependent on successful / unsuccessful parsing (see Appendix C.4).
- 1555 7. Clients that can receive notes **MUST** verify that successfully parsed plaintext data  
1556 is the opening of the corresponding commitment in the transaction (see Item 3b).
- 1557 8. Clients **MUST NOT** perform any network-related action, including closing the con-  
1558 nection, dependent on whether the parsed note data is the opening of the corre-  
1559 sponding commitment (see Appendix C.4).
- 1560 9. Clients **MUST** confirm that, after adding the new commitments, the local repre-  
1561 sentation of the Merkle tree of commitments has a root consistent with the event  
1562 data.
- 1563 10. Clients **SHOULD** record data related to valid decrypted notes. This will be required  
1564 in order to spend the notes in a future transaction.
- 1565 11. Clients implementing spending functionality **SHOULD** process all nullifiers in Mix  
1566 transaction events, checking for any corresponding notes previously recorded. Any  
1567 such note should be marked as spent.

#### 1568 4.1.3 Prepare Arguments for Mix Transaction

1569 Clients **MUST NOT** query Ethereum nodes while generating any arguments to a Mix call.  
1570 In particular, Merkle paths **MUST** be calculated using the clients representation of the  
1571 Merkle tree of commitments that was constructed by parsing events.

1572 Where the zero-knowledge proof is generated by some external process, clients **MUST**  
1573 put in place sufficient security schemes to ensure that:



- 1574 • they are communicating with an authentic proof generation process (not a man-  
1575 in-the-middle), and
- 1576 • data sent to and from the proving process cannot be observed in transit by a third  
1577 party, and
- 1578 • the proof has been generated for the correct witness.

1579 Without these safe-guards, the operation of the system and the secret data required  
1580 to spend the input notes may be compromised. See: Appendix C.6.

#### 1581 4.1.4 Wallet Backup and Recovery

1582 Wallets are expected to rely on locally stored data such as **Zeth** note secret values, key  
1583 data, information about the **Mixer**state, etc. If such data is lost, it can be reconstructed  
1584 by replaying all **Mixer** events from the creation of **Mixer**, *if the user's Zeth secret*  
1585 *address is available*. As well as note data (recovered by decrypting the ciphertext from  
1586 event logs), nullifiers (also emitted to the event logs) reveal which of the users notes has  
1587 been spent.

1588 Thereby, given only the user's Zeth secret address the full wallet state should be  
1589 recoverable, although this operation may be computationally expensive and time con-  
1590 suming. To support this operation, implementors may consider making an exception  
1591 to the guidance in Section 4.1.1 by introducing a public method on **Mixer** to return a  
1592 snapshot of the *entire* contract state. The intention being to avoid the need to parse the  
1593 entire **Mixer** history. Such a query would not be for specific commitments or nullifiers,  
1594 and so the amount of information leaked is limited and may be considered acceptable.  
1595 However, the utility of this “snapshot” approach may be limited in most realistic situa-  
1596 tions.

1597 Firstly, **Mixer** may not trivially support a query returning the full set of nullifiers,  
1598 without some significant storage overhead. This is because nullifiers are most naturally  
1599 stored as a Solidity map, which cannot be enumerated efficiently. Clients must also  
1600 iterate through the full contract history attempting decrypt every ciphertext in the event  
1601 logs, in order to recover the user's **Zeth** notes they may contain. Thereby, while some  
1602 simple clients (e.g. payment verifiers) may not require historical nullifier or note data,  
1603 iteration through the **Mixer** history seems unavoidable for recovery in most realistic  
1604 cases.

1605 Although wallets **SHOULD** support recovery from key data alone, it is preferable to  
1606 avoid such expensive operations. Wallets **SHOULD** support backing up sufficient data  
1607 to allow recovery in a “reasonable” amount of time. Naturally, the definition of “rea-  
1608 sonable” here and the nature of the costs involved will depend heavily on the specific  
1609 use-cases being targeted by the implementation.

## 4.2 Contract Security Considerations

Section 4.1 mentions several considerations for client implementations, concerning how they interact with the contract. These must be taken into account when authoring the contract code, to ensure that clients can securely retrieve the information needed, and to discourage implementors from using insecure operations.

1. **Mixer** MUST validate inputs, the contract needs to ensure that the primary inputs are elements of the scalar field  $\mathbb{F}_{\mathbf{r}_{\text{BN}}}$  ( $< \mathbf{r}_{\text{BN}}$ ).
2. **Mixer** MUST output events for valid Mix calls, including:
  - (a) Commitment for each new note.
  - (b) Nullifier for each spent note.
  - (c) Value of new Merkle root of commitments.
  - (d) Ciphertexts for each new note.
  - (e) Implementation-specific data (such as the one-time sender public key specified in Section 3.5, required to decrypt the ciphertexts).
3. Except the Mix method, **Mixer** SHOULD NOT expose any public methods unless strictly required by the client.
4. The Mix function MUST be payable, to support non-zero *vin*.

Note that Item 4 requires  $tx_{\text{Mix}}.val$  to be set to a non-zero value, even if  $vin = 0$ . To work around this, **Mixer** must be authored so that it refunds the sender of the Ethereum transaction with the value  $tx_{\text{Mix}}.val$  in the case that  $vin = 0$ . In this way, clients calling Mix with  $vin = 0$  can set  $tx_{\text{Mix}}.val = 1\text{Wei}$ , in order to satisfy this requirement without losing any Ether.

## 4.3 Efficiency and Scalability

### 4.3.1 Importance of Performance

Poor performance and scalability has several impacts on the viability of the system.

Efficiency and performance are arguably most important for the **Mixer** contract, where gas usage directly affects the monetary cost of using Zeth to transfer value. That is, high gas costs could make transactions very expensive, and therefore not practical for many use-cases, undermining the utility and viability.

High storage or compute requirements on the client would severely restrict the set of devices on which Zeth client software can run, and long delays when sending or receiving transactions can adversely affect the user-experience, discouraging some users and undermining the privacy promises of the system.

1643 Although the proof-of-concept implementation of Zeth is not intended to be used in a  
1644 production environment, one of its aims is to demonstrate the practicality of the protocol  
1645 in terms of transaction costs. Therefore, some of the techniques described here have been  
1646 included in the proof-of-concept implementation, while in some cases implementors of  
1647 production software may wish to make different trade-offs.

### 1648 4.3.2 Cost Centers

1649 One important performance factor, primarily affecting client performance, is the cost of  
1650 zero-knowledge proof generation. This is directly related to the number of constraints  
1651 used to represent the statement in Section 2.2, which in turn depends on the specific  
1652 cryptographic primitives used (see Chapter 3).

1653 Note that cryptographic primitives which are “snark-friendly” (i.e. can be imple-  
1654 mented using fewer gates in an arithmetic circuit) may not necessarily run efficiently on  
1655 the EVM or on standard hardware. As such, trade-offs must be made between proof  
1656 generation cost and the gas costs of state transitions. An example of this is the hash  
1657 function used in the Merkle tree of commitments. This is not only used in the statement  
1658 of Section 2.2 (to verify Merkle proofs, see Section 2.2), but also on the client (to create  
1659 Merkle proofs, see Section 2.3) and in the **Mixer** contract (to compute the Merkle root,  
1660 see Section 2.5).

1661 Aside from the specific hash function used, implementors have some freedom in the  
1662 data structures and algorithms used to maintain the Merkle tree and generate proofs.  
1663 Because of this freedom, and the importance of the chosen algorithms on performance  
1664 across all components of the system, the majority of this section focuses on the details  
1665 of the Merkle tree.

1666 As described in Chapter 2, Zeth notes are maintained and secured by the Merkle tree,  
1667 whose depth MKDEPTH must be fixed when the contract is deployed. Therefore, MKDEPTH  
1668 determines the maximum number of notes ( $2^{\text{MKDEPTH}}$ ) that may be created over the lifetime  
1669 of the deployment. To ensure the utility of Zeth, MKDEPTH must be sufficiently large,  
1670 and therefore the following includes discussion of *scalability* with respect to MKDEPTH.

1671 Also, due to the fact that MKDEPTH is fixed, we assume that Merkle proofs are com-  
1672 puted as MKDEPTH-tuples, no matter how many leaves have been populated. Unpopulated  
1673 leaves are assumed to take some default value (usually a string of zero bits).

### 1674 4.3.3 Client Performance

#### 1675 Commitment Merkle Tree

1676 The simplest possible implementation, which stores only the data items at the leaves  
1677 of the tree, requires  $2^{\text{MKDEPTH}} - 1$  hash invocations to compute the Merkle root or to  
1678 generate a Merkle proof. The cost of this is too high to be practical for non-trivial  
1679 values of MKDEPTH.

1680 An immediate improvement in compute costs can be achieved by simply storing  
1681 all nodes (or all nodes whose value is not the default value) and updating only those

necessary as new commitments are added. This reduces the cost of updating the tree to a fixed number (linear in  $\log_2(\text{JSOUT})$  and  $\text{MKDEPTH}$ ), but doubles the storage cost.

In the case of the client, the Merkle tree will only be used to generate proofs for notes owned by the user of the client. Thereby clients need only store nodes of the Merkle tree that are required for this purpose, and may discard all others. In particular, any full sub-tree need only contain nodes that are part of Merkle paths associated with the user's notes. Implementations that discard unnecessary nodes achieve vast savings in storage space.

## Zero-Knowledge Proof Generation

As well as keeping the number of constraints as low as possible, it is important to ensure that the prover implementation is optimal to ensure proving times are not unnecessarily long. Proof generation should also exploit any available parallelism, to help reduce the time taken. This may require specific programming languages or frameworks to be used, necessitating that proof generation be performed by some external process (as is the case in the proof-of-concept implementation).

The proof generation process can also be very memory intensive (in part due to the FFT calculations required), and so ensuring that enough RAM is present in the system is important to avoid long proof times.

See Appendix C.6 for a discussion of related security concerns.

## 4.3.4 Contract Performance

For most components of the contract, the set of operations to be performed is strictly defined and the set of possible algorithmic optimizations that can be made is limited. In these cases, it is important to ensure that code is benchmarked and optimized to a reasonable degree, to minimize gas costs. We note that apart from the number and type of compute instructions executed, store and lookup operations have a significant impact on the gas used. In particular, storing new values is more expensive than over-writing existing values, and a gas rebate is made when contracts release stored values. See [Woo19, Appendix H.1] for further details.

The primary component in which algorithmic optimizations can be made is the Merkle tree of note commitments. The **Mixer** contract must compute (and store) the new Merkle root after adding the **JSOUT** new commitments as leaves.

As in Section 4.3.3, the simplest possible implementation which stores only the data items at the leaves of the tree, requires the full root to be recomputed, involving  $2^{\text{MKDEPTH}} - 1$  hash invocations. This quickly becomes impractical for non-trivial values of  $\text{MKDEPTH}$ .

The first-pass optimization (also described in Section 4.3.3) can be used to ensure that the cost of updating the Merkle tree (number of hash computations, stores and loads) is bounded by a constant that is linear in the Merkle tree depth. This is the strategy used in the proof-of-concept implementation of Mixer.

It may be possible to gain further improvements in gas costs by discarding nodes from the Merkle Tree that are not required. Unlike clients, **Mixer** is only required to

1722 compute the new Merkle root, and does not need to create or validate Merkle proofs  
 1723 (as these are checked as part of the zero-knowledge proof). Consequently, *all* nodes in a  
 1724 sub-tree can be discarded when the sub-tree is full, and the optimization is much simpler  
 1725 to implement than on the client.

1726 Another possible strategy for decreasing the gas costs associated with Merkle trees  
 1727 is *Merkle Shrubs*, described in [Lab19, Section 2.2]. Under this scheme, the contract  
 1728 maintains a “frontier” of roots of sub-trees and Merkle proofs provided by clients (as  
 1729 auxiliary inputs to the  $\mathbf{R}^Z$  circuit) contain a path from the leaf to one of the nodes in the  
 1730 frontier. The gas savings in this scheme are due to the fact that, for new commitments,  
 1731 the contract need only recompute the value of nodes from the leaf to the “frontier” (not  
 1732 all the way to the root of the tree). However this comes at the cost of complexity in the  
 1733 arithmetic circuit, which must verify a Merkle path to one of several frontier nodes.

1734 When choosing cryptographic primitives to be used on the EVM (and considering  
 1735 the trade-off with other platforms, described in Section 4.3.1) it may be valuable to note  
 1736 that the EVM supports so-called “pre-compiled contracts”. These behave like built-  
 1737 in functions providing very gas-efficient access to certain algorithms, such as Keccak.  
 1738 However, pre-compiled contracts exist only for a limited set of algorithms. Others must  
 1739 be implemented using EVM instructions.

#### 1740 4.3.5 Optimizing Blake2’s circuit.

1741 After presenting Blake2s circuit and its components working on little endian variables,  
 1742 we show a few optimizations.

#### 1743 Helper circuits

1744 We first define the following helper circuits needed in the Blake2s routine working on  
 1745  $w$ -bit long words.

1746 **XOR circuits** The following XOR circuits on  $w$ -bit long variables have been imple-  
 1747 mented, **we assume the inputs are boolean (this is not checked)**,

- 1748 • “Classic” XOR circuit, which xors 2 variables,  
 1749  $a \oplus b = c$ ;
- 1750 • XOR with constant, which xors two variables and a constant,  
 1751  $a \oplus b \oplus c = d$ , with  $c$  constant;
- 1752 • XOR with rotation, which xors two variables and rotates the result.  
 1753  $a \oplus b \ggg r = c$ , with  $r$  constant, and  $\ggg$  the rightward rotation [MJS15, Section  
 1754 2.3]; i.e. for and constant  $r < w$  we have  $\forall 0 < i < w, a_i \oplus b_i = c_{i+r \pmod w}$

1755 Each of these circuits presents  $w$  constraints. Assuming that the inputs are boolean,  
 1756 the output is automatically boolean. To ascertain that both inputs are boolean ( $a$  and  
 1757  $b$ ), we would need  $2 \cdot w$  more gates per circuit.

1758 **Modular addition** We present here two circuits to verify modular arithmetic.

1759 **Double modular addition:  $a + b = c \pmod{2^w}$ .** This circuit checks that the  
 1760 sum of two  $w$ -bit long variables in little endian format modulo  $2^w$  is equal to a  $w$ -bit  
 1761 long variable. More precisely, it checks the equality of the modular addition of  $a + b$   
 1762  $\pmod{2^w}$  and  $c$  and the booleanness of the later. **We assume the inputs are boolean (this**  
 1763 **is not checked).**

1764 As the addition of two  $w$ -bit long integers results in at most an  $w + 1$ -bit integer  
 1765 however, we consider  $c$  to be  $w + 1$  bit-long. We do not care about the last bit value,  
 1766  $c_w$ , however but have to ensure its booleanness.

1767 The circuit presents the following  $w + 2$  constraints, for  $a$  and  $b$  of size  $w$ , where  
 1768  $w = 32$  in practice, and variable  $c$  of size  $w + 1$ , that:

$$\sum_{i=0}^{w-1} (a_i + b_i) \cdot 2^i = \sum_{j=0}^w c_j \cdot 2^j \quad (4.1)$$

$$\forall j \in [0, w], (c_j - 0) \cdot (c_j - 1) = 0 \quad (4.2)$$

1769 **Triple modular addition:  $a + b + c = d \pmod{2^w}$ .** This circuit checks the  
 1770 equality of a  $w$ -bit long variable  $d$  with the sum of three  $w$ -bit long variables in little  
 1771 endian format modulo  $2^w$ . More precisely, it checks the equality of the modular addition  
 1772 of  $a + b + c \pmod{2^w}$  and  $d$  and the booleanness of the later. **We assume the inputs are**  
 1773 **boolean (this is not checked).**

1774 As the addition of three  $w$ -bit long integers results in at most an  $w + 2$ -bit integer  
 1775 however, we consider  $d$  to be  $w + 2$  bit-long. We do not care about the last two bits'  
 1776 values however ( $d_w$  and  $d_{w+1}$ ), but have to ensure their booleanness.

The circuit presents the following  $w + 3$  constraints, for  $a$ ,  $b$  and  $c$  of size  $w$ , where  
 $w = 32$  in practice, and variable  $d$  of size  $w + 2$ , that:

$$\sum_{i=0}^{w-1} (a_i + b_i + c_i) \cdot 2^i = \sum_{j=0}^{w+1} d_j \cdot 2^j \quad (4.3)$$

$$\forall j \in [0, w + 1], (d_j - 0) \cdot (d_j - 1) = 0 \quad (4.4)$$

## 1777 Blake2s routine circuit

1778 We define in this section the circuit of Blake2 routine (see: [MJS15, Section 3.1] and Fig. 4.1),  
 1779 also known as “G function” [ANWOW13, Section 2.4]. G is based on ChaCha encryption  
 1780 [Ber08a], works on  $w$ -bit long words, and presents  $8 \cdot w + 10$  constraints. The function  
 1781 mixes a state ( $a$ ,  $b$ ,  $c$ , and  $d$ ) with the inputs ( $x$ , and  $y$ ) and returns the updated state.

1782 **This circuit does not check the booleanness of the inputs or state.** However, given that  
 1783 the state is boolean, the output is automatically boolean due to the use of the modular  
 1784 addition circuits.

1785 For Blake2s, we have  $w = 32$ ,  $r_1 = 16$ ,  $r_2 = 12$ ,  $r_3 = 3$  and  $r_4 = 7$ .

$G(a, b, c, d; x, y) \rightarrow (a_2, b_2, c_2, d_2)$	getSigma()
1 : $a_1 = a + b + x \pmod{2^w}$	1 : $\Sigma \in (\mathbb{N}^{16})^{10}$
2 : $d_1 = d \oplus a_1 \ggg r_1$	2 : $\Sigma[0] = (0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15)$
3 : $c_1 = c + d_1 \pmod{2^w}$	3 : $\Sigma[1] = (14, 10, 4, 8, 9, 15, 13, 6, 1, 12, 0, 2, 11, 7, 5, 3)$
4 : $b_1 = b \oplus c_1 \ggg r_2$	4 : $\Sigma[2] = (11, 8, 12, 0, 5, 2, 15, 13, 10, 14, 3, 6, 7, 1, 9, 4)$
5 : $a_2 = a_1 + b_1 + y \pmod{2^w}$	5 : $\Sigma[3] = (7, 9, 3, 1, 13, 12, 11, 14, 2, 6, 5, 10, 4, 0, 15, 8)$
6 : $d_2 = d_1 \oplus a_2 \ggg r_3$	6 : $\Sigma[4] = (9, 0, 5, 7, 2, 4, 10, 15, 14, 1, 11, 12, 6, 8, 3, 13)$
7 : $c_2 = c_1 + d_2 \pmod{2^w}$	7 : $\Sigma[5] = (2, 12, 6, 10, 0, 11, 8, 3, 4, 13, 7, 5, 15, 14, 1, 9)$
8 : $b_2 = b_1 \oplus c_2 \ggg r_4$	8 : $\Sigma[6] = (12, 5, 1, 15, 14, 13, 4, 10, 0, 7, 6, 3, 9, 2, 8, 11)$
9 : <b>return</b> $a_2, b_2, c_2, d_2$	9 : $\Sigma[7] = (13, 11, 7, 14, 12, 1, 3, 9, 5, 0, 15, 4, 8, 6, 2, 10)$
Figure 4.1: <b>G</b> primitive [MJS15, Section 3.1]	10 : $\Sigma[8] = (6, 15, 14, 9, 11, 3, 0, 8, 12, 2, 13, 7, 1, 4, 10, 5)$
	11 : $\Sigma[9] = (10, 2, 8, 4, 7, 6, 1, 5, 15, 11, 9, 14, 3, 12, 13, 0)$
	12 : <b>return</b> $\Sigma$

Figure 4.2: Blake2 permutation table [MJS15, Section 2.7]

#### 1786 Blake2s compression function circuit

The compression function is defined as follows, for more details see: Fig. 4.3,

$$\text{Blake2sC} : \mathbb{B}^n \times \mathbb{B}^{2n} \times \mathbb{B}^{n/4} \times \mathbb{B}^{n/4} \rightarrow \mathbb{B}^n$$

1787 Blake2C takes as input a state  $h \in \mathbb{B}^n$  which is used as chaining value when hashing,  
 1788 the message to compress  $x \in \mathbb{B}^{2n}$ , a message length written in binary  $t \in \mathbb{B}^{n/4}$  which is  
 1789 incremented when hashing and a binary flag  $f \in \mathbb{B}^{n/4}$  to tell whether the current block  
 1790 is the last to be compressed to prevent length extension attacks.

1791 Blake2C uses **G** function iteratively over **rounds** rounds on a state and message. We  
 1792 hardcoded the constant initialization vector **IV** and the permutation table  $\Sigma$ . **Blake2sC**  
 1793 works in little endian (see: [MJS15, Section 2.4]) on  $n = 256$ -bit long variables and  $w =$   
 1794 32-bit long words and the rotation constants specified in Section 4.3.5 (see: [MJS15, Sec-  
 1795 tion 2.1]). We have the following constants (see: specifications [ANWOW13] and [MJS15,  
 1796 Section 2.2]),

- 1797 • **IV** is the  $8 \cdot w$ -bit long Initialisation Vector; it corresponds to the first  $w$  bits of the  
 1798 fractional parts of the square roots of the first eight prime numbers  $(2, 3, 5, 7, \dots)$   
 1799 (see: [MJS15, Section 2.6]);
- 1800 •  $\Sigma$  are the  $10 \cdot 16$  permutation constants of Blake2 (see: Fig. 4.2 and [MJS15, Section  
 1801 2.7]);
- 1802 • **rounds**, the number of rounds: 10 for Blake2sC, 12 for Blake2bC().

1803 We have the following variables (see: specifications [ANWOW13] and [MJS15, Section  
 1804 2.2]),

- 1805 •  $H$  is the  $8 \cdot w$ -bit long initial state while  $v$  is the  $16 \cdot w$ -bit long final state;
- 1806 •  $T[i]$  are two  $w$ -bit long counters encoding the block length;
- 1807 •  $F[i]$  are two  $w$ -bit long finalization flags. We set the first one  $F[0]$  to  $2^w - 1$  to state
- 1808 when the input block is the last one to be hashed. The second is unused,  $F[1] = 0$ ,
- 1809 as it is only used for tree hashing mode which is not our case.

1810 We introduce the following functions to write Blake2C (see: specifications [ANWOW13]  
1811 and [MJS15, Section 2.6]):

- 1812 • The function **prime** takes a positive integer  $i$  as input and outputs the  $i^{th}$  prime
- 1813 number;
- 1814 • The function **dec** takes a real number  $x$  as input outputs its positive decimal part.

1815 This circuit presents  $(64 \cdot \mathbf{rounds} + 8) \cdot w + 8 \cdot \mathbf{rounds} + 10$  constraints. For Blake2sC,  
1816 as  $w = 32$  and  $\mathbf{rounds} = 10$ , we have 21536 constraints.

1817 **We do not check the input block booleanness in this circuit.** Given that the initial  
1818 state is boolean, the output is automatically boolean. This can be proved iteratively by  
1819 the booleanness of  $G$  primitive's output.

1820 **Security requirement** To ensure the correct use of Blake2sC, Blake2sC's inputs **MUST** be  
1821 boolean.

## 1822 Blake2s hash function

The hash function is defined as follows, for more details see: Fig. 4.3,

$$\mathbf{Blake2s} : \mathbb{B}^{\leq 2n} \times \mathbb{B}^* \rightarrow \mathbb{B}^n$$

1823 Blake2 takes as input a hash key  $k \in \mathbb{B}^n$  and the message to hash  $x \in \mathbb{B}^{2n}$ .

1824 Blake2 uses Blake2C function iteratively over each  $2n$ -bit long chunk of the padded  
1825 message. If the key is non null, it is used as the first block to be hashed. We hardcoded  
1826 the constant initialization vector  $IV$  and part of the parameter block  $PB$ . We have the  
1827 following constants (see: specifications [ANWOW13] and [MJS15, Section 2.2]),

- 1828 •  $IV$  is the  $8 \cdot w$ -bit long Initialisation Vector; it corresponds to the first  $w$  bits of the
- 1829 fractional parts of the square roots of the first eight prime numbers  $(2, 3, 5, 7, \dots)$
- 1830 (see: [MJS15, Section 2.6]).

1831 We have the following variables (see: specifications [ANWOW13] and [MJS15, Section  
1832 2.2]),

- 1833 •  $PB$  is the  $16 \cdot w$ -bit long parameter block used to initialize the state (see: [MJS15,  
1834 Section 2.5]). In big endian, the first byte would correspond to the digest length  
1835 (fixed to 32 bytes), the second byte to the key length, the third and fourth bytes  
1836 correspond to the use of the serial mode;



### Blake2C( $h, m, t, f$ )

---

```

1 :  $\mathbf{T}, \mathbf{F}, \mathbf{H}, \mathbf{IV}, v \in (\mathbb{B}^w)^2 \times (\mathbb{B}^w)^2 \times (\mathbb{B}^w)^8 \times (\mathbb{B}^w)^8 \times (\mathbb{B}^w)^{16}$ 
2 :  $\{\mathbf{IV}[i]\}_{i \in [8]} \leftarrow \{ \lfloor 2^w \cdot \text{dec}(\sqrt{\text{prime}(i+1)}) \rfloor \}_{i \in [8]}$ 
3 :  $\Sigma \leftarrow \text{getSigma}()$ 
4 :  $\{\mathbf{H}[i]\}_{i \in [8]} \leftarrow \{h[i \cdot w:(i+1) \cdot w]\}_{i \in [8]}$ 
5 :  $\{m[i]\}_{i \in [8]} \leftarrow \{x[i \cdot w:(i+1) \cdot w]\}_{i \in [8]}$ 
6 :  $\mathbf{T}[0], \mathbf{T}[1] \leftarrow t[w:2w], t[0:w]$ 
7 :  $\mathbf{F}[0], \mathbf{F}[1] \leftarrow f[w:2w], f[0:w]$ 
8 :  $\{v[i]\}_{i \in [8]} \leftarrow \{\mathbf{H}[i]\}_{i \in [8]}$ 
9 :  $\{v[i+8]\}_{i \in [8]} \leftarrow \{\mathbf{IV}[i]\}_{i \in [8]}$ 
10 :  $v[12], v[13] \leftarrow v[12] \oplus \mathbf{T}[0], v[13] \oplus \mathbf{T}[1]$ 
11 :  $v[14], v[15] \leftarrow v[14] \oplus \mathbf{F}[0], v[15] \oplus \mathbf{F}[1]$ 
12 : foreach  $r \in [\text{rounds}]$  do
13 :    $\tau \leftarrow \Sigma[r \pmod{15}]$ 
14 :    $v[0], v[4], v[8], v[12] \leftarrow \mathbf{G}(v[0], v[4], v[8], v[12], m[\tau[0]], m[\tau[1]])$ 
15 :    $v[1], v[5], v[9], v[13] \leftarrow \mathbf{G}(v[1], v[5], v[9], v[13], m[\tau[2]], m[\tau[3]])$ 
16 :    $v[2], v[6], v[10], v[14] \leftarrow \mathbf{G}(v[2], v[6], v[10], v[14], m[\tau[4]], m[\tau[5]])$ 
17 :    $v[3], v[7], v[11], v[15] \leftarrow \mathbf{G}(v[3], v[7], v[11], v[15], m[\tau[6]], m[\tau[7]])$ 
18 :    $v[0], v[5], v[10], v[15] \leftarrow \mathbf{G}(v[0], v[5], v[10], v[15], m[\tau[8]], m[\tau[9]])$ 
19 :    $v[1], v[6], v[11], v[12] \leftarrow \mathbf{G}(v[1], v[6], v[11], v[12], m[\tau[10]], m[\tau[11]])$ 
20 :    $v[2], v[7], v[8], v[13] \leftarrow \mathbf{G}(v[2], v[7], v[8], v[13], m[\tau[12]], m[\tau[13]])$ 
21 :    $v[3], v[4], v[9], v[14] \leftarrow \mathbf{G}(v[3], v[4], v[9], v[14], m[\tau[14]], m[\tau[15]])$ 
22 : return  $\|_{i=0}^8 \mathbf{H}[i] \oplus v[i] \oplus v[i+8]$ 

```

Figure 4.3: Blake2 compression function [MJS15, Section 3.2]. Set  $n$ ,  $w$  and  $\mathbf{G}$ 's constants to obtain Blake2sC.

1837 •  $\mathbf{H} \in \mathbb{B}^{\text{BLAKE2sCLEN}}$ , the chaining value.

1838 **We do not check the input block booleanness in this circuit.** Given that the initial  
1839 state is boolean, the output is automatically boolean. This can be proved iteratively by  
1840 the booleanness of Blake2C primitive's output.

1841 **Security requirement** To ensure the correct use of Blake2s, Blake2s's inputs **MUST** be  
1842 boolean.

### 1843 Optimizing the circuits

1844 We presented beforehand helper circuits to build Blake2s compression function. We show  
1845 here two exclusive methods to optimize these circuits.

## Blake2( $k, x$ )

---

```

1 :  H, IV, PB  $\in \mathbb{B}^{8w} \times \mathbb{B}^{8w} \times \mathbb{B}^{8w}$ 
2 :  PB  $\leftarrow \text{pad}_{8 \cdot w}(\text{encode}_{\mathbb{N}}(0x0101)) \parallel \text{pad}_w(\text{encode}_{\mathbb{N}}(\lceil \text{length}(k)/\text{BYTELEN} \rceil)) \parallel \text{encode}_{\mathbb{N}}(0x20)$ 
3 :  IV  $\leftarrow \parallel_{i=0}^8 \left[ 2^w \cdot \text{dec}(\sqrt{\text{prime}(i+1)}) \right]$ 
4 :  H  $\leftarrow \text{PB} \oplus \text{IV}$ 
5 :  y  $\leftarrow x$ 
6 :  if length( $k$ )  $\neq 0$  do
7 :    y  $\leftarrow \text{pad}_{2n}(k) \parallel y$ 
8 :    z  $\leftarrow \text{pad}_{2n \cdot \lceil \text{length}(y)/2n \rceil}(y)$ 
9 :    for  $i \in [\lceil \text{length}(z)/2n \rceil]$  do
10 :      if  $i = \lceil \text{length}(z)/2n \rceil - 1$  do
11 :        H  $\leftarrow \text{Blake2C}(H, z[i \cdot 2n:(i+1) \cdot 2n], \text{pad}_{2w}(\text{encode}_{\mathbb{N}}(\lceil \text{length}(y)/\text{BYTELEN} \rceil)), \text{pad}_{2w}(\text{encode}_{\mathbb{N}}(2^w - 1)))$ 
12 :      else
13 :        H  $\leftarrow \text{Blake2C}(H, z[i \cdot 2n:(i+1) \cdot 2n], \text{pad}_{2w}(\text{encode}_{\mathbb{N}}((i+1) \cdot 2n/\text{BYTELEN})), \text{pad}_{2w}(0))$ 
14 :    return H

```

Figure 4.4: Blake2 hash function [MJS15, Section 3.3]. Set  $n = 16w$  and G's constants accordingly to obtain Blake2s.

## 1846 Optimizing the Modular additions

**Double modular addition:  $a + b = c \pmod{2^w}$ .** We present here an optimization on the circuit to save one constraint by merging the modular constraint with a boolean constraint. The optimized circuit presents the following constraints:

$$\left( \sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i \right) \cdot \left( \sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i - 2^w \right) = 0 \quad (4.5)$$

$$\forall j \in [0, w-1], (c_j - 0) \cdot (c_j - 1) = 0 \quad (4.6)$$

1847 with  $\sum_{i=0}^{w-1} x_i \cdot 2^i$  a binary encoding of  $x$  ( $x_i$  is the  $i^{\text{th}}$  bit of  $x$ ).

1848 These equations describe  $w + 1$  constraints to prove the bit equality  $a + b = c$  (note  
1849 that an additional  $2 \cdot w$  constraints would be required to prove the booleanness of input  
1850 variables  $a$  and  $b$ ). We now explain how we obtained them.

1851 *Proof.* The most straightforward way to prove that  $a + b = c \pmod{2^w}$  and  $c$  booleanness  
1852 is with the set of constraints illustrated in Eq. (4.1) and in Eq. (4.2).

As we perform arithmetic modulo  $2^w$ , we do not care about the value of  $c_w$  but would like to ensure its booleanness. As one may notice, the summing constraint Eq. (4.1) is an equality of two linear combinations with no multiplication by a variable. Hence, we can combine it with the boolean constraint of  $c_w$  to remove any reference to  $c_w$  and still

have a bilinear gate. To do so, we first rewrite Eq. (4.1) as an equality check over  $c_w \cdot 2^w$  and multiply Eq. (4.2) for  $j = n$  by  $2^{2 \cdot w}$ .

$$\sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i = c_w \cdot 2^w \quad (4.7)$$

$$2^w \cdot (c_w - 0) \cdot 2^w \cdot (c_w - 1) = 0 \quad (4.8)$$

We finally replace  $c_w \cdot 2^w$  in Eq. (4.8) by the value from Eq. (4.7).

$$\begin{aligned} 0 &= 2^w \cdot (c_w - 0) \cdot 2^w \cdot (c_w - 1) = 2^w \cdot c_w \cdot (2^w \cdot c_w - 2^w) \\ &= \left( \sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i \right) \cdot \left[ \left( \sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i \right) - 2^w \right] \end{aligned}$$

1853 This results in Eq. (4.5) and Eq. (4.6). All references to  $c_w$  disappeared and, with a  
1854 single multiplication by a variable, we still have bilinear gates.  $\square$

1855 **Triple modular addition:  $a + b + c = d \pmod{2^w}$ .** To optimize, we will use  
1856 the former circuit twice. We will define a temporary variable  $d'$  such that  $a + b = d'$   
1857  $\pmod{2^w}$ . As such, we have  $c + d' = d \pmod{2^w}$ . As  $d'$  is the addition of two  $w$ -bit long  
1858 variables, it is  $w + 1$  bit long. However as we do the sum modulo  $2^w$ , we discard the last  
1859 bit of  $d'$ . We do similarly for  $d$ . To ensure that  $d$  is boolean, we check the booleanness  
1860 of the  $w + 1$  bits of  $d$  as well as the booleanness of the last bit of  $d'$  (to account for  $d$ 's  
1861  $w + 2^{th}$  bit in the original expression  $(a + b + c = d \pmod{2^w})$ ).

We thus obtain the following circuit with  $w + 2$  constraints,

$$\begin{aligned} &\left( \sum_{i=0}^{w-1} (a_i + b_i - d'_i) \cdot 2^i \right) \cdot \left( \sum_{i=0}^{w-1} (a_i + b_i - d'_i) \cdot 2^i - 2^w \right) = 0 \\ &\left( \sum_{i=0}^{w-1} (c_i + d'_i - d_i) \cdot 2^i \right) \cdot \left( \sum_{i=0}^{w-1} (c_i + d'_i - d_i) \cdot 2^i - 2^w \right) = 0 \\ &\forall j \in [0, w - 1], (d_j - 0) \cdot (d_j - 1) = 0 \end{aligned}$$

1862 These optimizations leads to a gain of 320 constraints ( $= 4 \cdot 8 \cdot rounds$ ).

1863 **Optimizing Blake2s routine's circuit** As seen in Fig. 4.1, our routine presents 2  
1864 double and 2 triple modular additions. Each of these circuits comprises at least one  
1865 modular constraint which pack several  $w$ -bit long variables. The circuit is however  
1866 processed in  $\mathbb{F}_{\text{rBN}}$ , that is to say most integers can be written over BNFIELDCAP bits. We  
1867 can thus batch the modular constraints. As the G primitive performs 2 double modular  
1868 and 2 triple modular, we have in total 6 modular checks per iteration. We can batch  
1869 up to BNFIELDCAP/ $w$  constraints together. For  $w = 32$  and BNFIELDCAP = 253, we can

1870 encode up to 7 words per field element, that is to say we can include all the modular  
1871 constraints into a single one.

1872 This optimization leads to a gain of 274 constraints ( $= 4 \cdot 8 \cdot 10 - \lceil \frac{4 \cdot 8 \cdot 10}{7} \rceil$ ).

1873 **Optimization conclusion** Using the more efficient optimization on the modular ad-  
1874 ditions, the Blake2s compression function comprises 21216 constraints.

#### 1875 Increasing the PRF security with Blake

1876 As Blake2 comprises a personalization tag in its parameter block PB, one could ensure  
1877 the independence of the PRFs by writing different tags for each of them. We would be  
1878 able to consider up to  $2^{30}$  inputs and outputs and have a security parameter of  $\lambda$ . We  
1879 did not choose to write this enhancement in the instantiation to keep a general tagging  
1880 method in case of a change of hash function.

## 1881 4.4 Encryption of the notes

1882 In this section we give some remarks concerning the implementation of the **Zeth** en-  
1883 cryption scheme, described in Section 3.5. As noted, there are several details in the  
1884 specification of the underlying primitives which can impact security if not carefully im-  
1885 plemented. The following list is by no means exhaustive but includes several details  
1886 noted during development of the proof-of-concept system.

- 1887 • Private keys for Curve25519 **MUST** be randomly generated as 32 bytes where the  
1888 first byte is a multiple of 8, and the last byte takes a value between 64 and 127  
1889 (inclusive). Further details are given in [Ber06], including an example algorithm  
1890 for generation. Implementations **MUST** take care to ensure that their code, or any  
1891 external libraries they rely upon, correctly perform this step.
- 1892 • A similar observation holds for Poly1305 in which the  $r$  component of the MAC  
1893 key  $(r, s)$  **MUST** be *clamped* in a specific way (see Section 3.5.3). This step is also  
1894 essential and **MUST** be performed.
- 1895 • In the implementation of the ChaCha stream cipher, correct use of the *key*, *counter*  
1896 and *nonce* **MUST** be ensured in order to adhere to the standard and guarantee the  
1897 appropriate security properties.

1898 During the proof-of-concept implementation it was not obvious that the cryptogra-  
1899 phy library<sup>1</sup> adhered to the specifications with respect to the above points. In particular,  
1900 it was not clear whether key clamping was performed at generation time and/or when  
1901 performing operations, and the interface to the ChaCha cipher accepted a different set  
1902 of input parameters (namely *key* and *nonce* with no *counter*). This left some ambiguity  
1903 about the responsibility for clamping, and whether the ChaCha block data would be

---

<sup>1</sup><https://cryptography.io/en/latest/>

1904 updated as described in the specification. Details of how this was resolved are given in  
1905 the proof-of-concept encryption code, which may prove a useful reference for implemen-  
1906 tors<sup>2</sup>.

---

<sup>2</sup>see: <https://github.com/clearmatics/zeth/blob/v0.4/client/zeth/encryption.py>

# Appendix A

## Transaction Non Malleability

The transaction malleability problem for a DAP (Section 1.4) is characterized by an experiment TR-NM, which involves a polynomial-size adversary  $\mathcal{A}$  attempting to break DAP.

**Definition 32.** Let DAP be a (candidate) Decentralized Anonymous Payment scheme.

$$\text{DAP} = (\text{Setup}, \text{GenAddr}, \text{SendTx}, \text{VerifyTx}, \text{Receive})$$

We say that DAP is TR-NM secure if, for every  $\text{poly}(\lambda)$ -size adversary  $\mathcal{A}$  and sufficiently large  $\lambda$ ,

$$\text{Adv}_{\text{DAP}, \mathcal{A}}^{\text{tr-nm}}(\lambda) < \text{negl}(\lambda),$$

where  $\text{Adv}_{\text{DAP}, \mathcal{A}}^{\text{tr-nm}}(\lambda) = \Pr[\text{TR-NM}(\text{DAP}, \mathcal{A}, \lambda) = 1]$  is  $\mathcal{A}$ 's advantage in the TR-NM experiment.

Below, we adapt [BSCG<sup>+</sup>14, Appendix C.2] to our specific DAP, **Zeth**.

We now describe the TR-NM experiment mentioned above for **Zeth**. Given a (candidate) **Zeth** DAP, adversary  $\mathcal{A}$ , and security parameter  $\lambda$ , the (probabilistic) experiment TR-NM(DAP,  $\mathcal{A}$ ,  $\lambda$ ) consists of an interaction between  $\mathcal{A}$  and a challenger  $\mathcal{C}$ , terminating with a binary output by  $\mathcal{C}$ .

At the beginning of the experiment,  $\mathcal{C}$  samples  $pp \leftarrow \text{Setup}(\lambda)$  and sends  $pp$  to  $\mathcal{A}$ . Next,  $\mathcal{C}$  initializes a DAP oracle  $\text{O}^{\text{DAP}}$  with  $pp$  and allows  $\mathcal{A}$  to issue queries to it [RZ19, Appendix B].

At the end of the experiment,  $\mathcal{A}$  sends to  $\mathcal{C}$  a **Mixer** contract call transaction  $tx_{\text{Mix}}^*$ , and  $\mathcal{C}$  outputs 1 iff the following conditions hold. Letting  $T$  be the set of transactions generated by  $\text{O}^{\text{DAP}}$  in response to **SendTx** queries, there exists  $tx_{\text{Mix}} \in T$  such that:

1.  $tx_{\text{Mix}}$  was not inserted in  $L$  by  $\mathcal{A}$ ;
2.  $tx_{\text{Mix}}^*.data \neq tx_{\text{Mix}}.data$ ;
3.  $\text{VerifyTx}(pp, tx_{\text{Mix}}^*, L') = 1$  where  $L'$  is the portion of the ledger  $L$  preceding  $tx_{\text{Mix}}$ ;
4. a serial number revealed in  $tx_{\text{Mix}}^*$  is also revealed in  $tx_{\text{Mix}}$ .

## A.1 Transaction malleability attack on Zeth

In this section we present the threat related to the transaction malleability attack on Zeth and expose the solutions by ZeroCash [BSCG<sup>+</sup>14] and Zcash [ZCa19] that we adapted.

First, we start by assuming that none of the checks related to transaction malleability attack have been added in the protocol Chapter 2. As such, we assume that *hsig* and *htags* are not attributes of `PrimInputDType`,  $\phi$  is not an attribute of `AuxInputDType`, and *otssig* and *otsvk* are not attributes of the `MixInputDType` data type anymore. As a consequence, all checks related to these attributes are removed from the protocol. Moreover, if *zn* is an object of type `ZethNoteDType`, then *zn*. $\rho$  is chosen at random. Finally, the NP-relation used in Zeth, now denoted  $\mathbf{R}^{\text{mal}}$ , becomes the following:

- For each  $i \in [\text{JSIN}]$ :

1.  $\text{aux.jsins}[i].\text{znote.apk} = \text{Blake2s}(\text{tag}_{\text{ask}}^{\text{addr}} \parallel \text{pad}_{\text{BLAKE2sCLEN}}(0))$   
with  $\text{tag}_{\text{ask}}^{\text{addr}}$  defined in Section 3.1.3
2.  $\text{aux.jsins}[i].\text{nf} = \text{Blake2s}(\text{tag}_{\text{ask}}^{\text{nf}} \parallel \text{aux.jsins}[i].\text{znote}.\rho)$   
with  $\text{tag}_{\text{ask}}^{\text{nf}}$  defined in Section 3.1.3
3.  $\text{aux.jsins}[i].\text{cm} = \text{Blake2s}(\text{aux.jsins}[i].\text{znote}.r \parallel m)$   
with  $m = \text{aux.jsins}[i].\text{znote.apk} \parallel \text{aux.jsins}[i].\text{znote}.\rho \parallel \text{aux.jsins}[i].\text{znote}.v$
4.  $(\text{aux.jsins}[i].\text{znote}.v) \cdot (1 - e) = 0$  is satisfied for the boolean value  $e$  set such that if  $\text{aux.jsins}[i].\text{znote}.v > 0$  then  $e = 1$ .
5. The Merkle root  $\text{mkroot}'$  obtained after checking the Merkle authentication path  $\text{aux.jsins}[i].\text{mkpath}$  of commitment  $\text{aux.jsins}[i].\text{cm}$ , with  $\text{MIMC}[7]\text{-MP}_{\text{rBN}}$ , equals to  $\text{prim.mkroot}$  if  $e = 1$ .
6.  $\text{prim.nfs}[i]$   
 $= \{\text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.jsins}[i].\text{nf}[k \cdot \text{BNFIELD CAP}:(k + 1) \cdot \text{BNFIELD CAP}])\}_{k \in [\lfloor \text{PRNFOUTLEN} / \text{BNFIELD CAP} \rfloor]}$

- For each  $j \in [\text{JSOUT}]$ :

1.  $\text{prim.cms}[j] = \text{Blake2s}(\text{aux.znotes}[j].r \parallel m)$   
with  $m = \text{aux.znotes}[j].\text{apk} \parallel \text{aux.znotes}[j].\rho \parallel \text{aux.znotes}[j].v$

- $\text{prim.rsd} = \text{Pack}_{\text{rsd}}(\{\text{aux.jsins}[i].\text{nf}\}_{i \in [\text{JSIN}]}, \text{aux.vin}, \text{aux.vout})$
- Check that the “joinsplit is balanced”, i.e. check that the joinsplit equation holds:

$$\begin{aligned} & \text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.vin}) + \sum_{i \in [\text{JSIN}]} \text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.jsins}[i].\text{znote}.v) \\ &= \sum_{j \in [\text{JSOUT}]} \text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.znotes}[j].v) + \text{Pack}_{\mathbb{F}_{\text{rBN}}}(\text{aux.vout}) \end{aligned}$$

### A.1.1 The transaction malleability attack

In order to win the game TR-NM on the weak Zeth DAP above, an adversary  $\mathcal{A}$  intercepts a target transaction  $tx_{\text{Mix}}$  by passively listening to the network (remember that transactions are broadcasted to the Ethereum network in order to be mined, see: Section 1.2.2), extracts the zk-proof and primary inputs from  $tx_{\text{Mix}}.data$  and uses these extracted pieces of information in order to create a malicious transaction  $tx_{\text{Mix}}'$ , where the ciphertexts are replaced by arbitrary data. The adversary can then broadcast  $tx_{\text{Mix}}'$  to the network in order for it to be mined. If the malicious transaction gets mined before the legitimate one, the input notes become spent and the ciphertexts are undecipherable making the new notes unredeemable (by any Zeth user!), since all attempts to decrypt the ciphertexts will fail (see: Section 2.6).

---

```

TxMalGen( $sk'_{\text{ECDSA}}, nce_{in}, tx_{\text{Mix}}$ )
1:  $p \leftarrow tx_{\text{Mix}}.gasP + 1$ 
2:  $l \leftarrow tx_{\text{Mix}}.gasL + 1$ 
3:  $zdata' \leftarrow tx_{\text{Mix}}.data$ 
4:  $zdata'.ciphers \leftarrow \mathbb{B}^*$ 
5:  $tx_{raw} \leftarrow \{nce : nce_{in}, gasP : p, gasL : l, to : tx_{\text{Mix}}.to, val : tx_{\text{Mix}}.val, data : zdata'\};$ 
6:  $\sigma_{\text{ECDSA}} \leftarrow \text{SigSch}_{\text{ECDSA}}.\text{Sig}(sk'_{\text{ECDSA}}, \text{Keccak256}(tx_{raw}));$ 
7:  $tx_{final} \leftarrow \{tx_{raw}, v : \sigma_{\text{ECDSA}}.v', r : \sigma_{\text{ECDSA}}.r', s : \sigma_{\text{ECDSA}}.s'\};$ 
8: return  $tx_{final}$ ;

```

---

Figure A.1: Transaction malleability attack function TxMalGen

As shown on Fig. A.1, during the attack, the adversary extracts the proof and primary inputs from the honest transaction, and replaces the ciphertexts by some arbitrary information. The attacker then formats this data into a transaction that calls the Mix function of **Mixer**, and submits it to the network. While the data fields ( $tx_{\text{Mix}}.data$  and  $tx_{\text{Mix}}'.data$ ) are different, the nullifiers revealed by each transactions are the same (i.e.  $tx_{\text{Mix}}.data.proof = tx_{\text{Mix}}'.data.proof$ , and  $tx_{\text{Mix}}.data.prim = tx_{\text{Mix}}'.data.prim$ ). As a consequence, if the adversary makes sure that  $tx_{\text{Mix}}'$  satisfies all the checks of EthVerifyTx (Section 1.2.2), he can be ensured that ZethVerifyTx( $tx_{\text{Mix}}'$ ) will return the same value as ZethVerifyTx( $tx_{\text{Mix}}$ ). Furthermore, if  $tx_{\text{Mix}}'.gasP > tx_{\text{Mix}}.gasP$ , then the adversary maximizes his chances to see his transaction mined first (Section 1.2.2), and so maximizes the chances for the malleability attack to be successful; leading to lost funds on **Mixer**.

**Remark 15.** Note that, although not directly contained within the data field of a **Mixer** call transaction, the Ethereum address  $\mathcal{S}_{\mathcal{E}}.Addr$  of the transaction sender is also used by the **Mixer** call (this is either the calling contract's address, or the transaction signer's address recovered as described in Remark 4). In particular, the balance of this Ethereum address is incremented by the value  $vout$  by successful Mix calls. If we again assume the



1986 *absence of the malleability checks, an attacker could re-sign any  $\widetilde{\text{Mixer}}$  call transaction*  
 1987 *with a key under his control, rebroadcast it as described above, and (with some reasonable*  
 1988 *probability) become the recipient of any public vout output value.*

## 1989 **A.2 Solutions to address the transaction malleability at-** 1990 **tack**

### 1991 **A.2.1 ZeroCash solution**

1992 The idea of the solution [BSCG<sup>+</sup>14] is to use a one-time, SUF-CMA digital signature  
 1993 and bind its verification key with the zk-proof primary inputs to *prevent an adversary*  
 1994 *from corrupting part of a transaction's data.*

1995 Specifically, to transact via **Zeth**, the user first samples a key pair  $(sk, vk)$  for a  
 1996 one-time signature scheme. He then computes the hash  $hsig = \text{CRH}(vk)$  (CRH being a  
 1997 collision resistant hash function, see: [BSCG<sup>+</sup>14]) and derives a value  $h_i = \text{PRF}_{ask_i}^{\text{pk}}(hsig)$ ,  
 1998 for each input note (i.e.  $i \in [\text{JSIN}]$ ), which acts as message authentication code (MAC)  
 1999 binding  $hsig$  to the address spending key of a note ( $ask_i$ ).

2000 The user then generates the zk-proof with the additional statement that the values  
 2001  $\{h_i\}_{i \in [\text{JSIN}]}$  are computed correctly. He finally uses  $sk$  to sign every value associated with  
 2002 the operation, thus obtaining a signature, which is included, along with the signature  
 2003 verification key  $vk$ , in the transaction. To verify a transaction on the DAP, it is nec-  
 2004 essary to verify that the primary inputs are correctly formatted, that the Merkle root  
 2005 corresponds to one of the previous states of the Merkle tree, that the nullifiers have not  
 2006 been declared in a previous transaction, that the  $hsig$  is correctly computed from  $vk$  and  
 2007 that both the zk-proof and the one-time signature verifications pass successfully.

2008 Now, an adversary trying to carry out the attack aforementioned will have to either  
 2009 change the ciphertexts or the encryption key. Nevertheless, doing so should lead to  
 2010 the one-time signature verification to fail or, should yield an attack that breaks the  
 2011 UF-CMA property of the one-time signature (as this corresponds to creating a forgery  
 2012 on a different message (not changing the signature)). Thereby, the adversary also has to  
 2013 modify the signature, however he does not know the one-time signing key used by the  
 2014 creator of the targeted transaction. As such, the adversary needs to use another signing  
 2015 key pair, however this leads to the check, verifying that  $hsig$  is correctly computed,  
 2016 to fail. If the adversary attempts to change  $hsig$ , the zk-proof verification fails as the  
 2017 NP-statement has changed. Hence, any attempt to carry out a malleability attack  
 2018 results in the violation of at least one check in the verification of the transaction on  
 2019 the DAP. The solution presented effectively solves the transaction-malleability attack  
 2020 initially described.

2021 **Remark 16.** *The one-timeness property of the signature scheme was required in ZeroCash to*  
 2022 *retain anonymity. It also makes analysing non-adaptive adversary sufficient. As Ethereum*  
 2023 *transaction senders need to pay the gas cost associated with their transactions, they*  
 2024 *are (the senders) not anonymous. This said, making sure that Zeth is designed with*

2025 *anonymity in mind is worth the effort in order to minimize information leakages and be*  
 2026 *ready if/when **Ethereum** incorporates protocol changes that enable anonymous transac-*  
 2027 *tions.*

## 2028 **A.2.2 Zcash’s solution**

2029 In addition to the changes aforementioned, **Zcash**’s solution [ZCa19] also consists in:

- Redefining the variable *hsig* as,

$$hsig = \text{CRH}(\text{randomSeed}, \{nf_i\}_{i \in [\text{JSIN}]}, vk)$$

2030 for some random seed *randomSeed*.

- Defining a new random variable  $\phi$  and using it with *hsig*, as key and input of a PRF respectively, to compute the identifier of each output notes  $\rho_j$  ( $j \in [\text{JSOUT}]$ ) and ensure their uniqueness (with overwhelming probability).

2034 These changes were made to prevent the Faerie Gold attack [ZCa19, Section 8.4], as well  
 2035 as to prevent linkability: if *hsig* were repeated in two transactions, the circuit would  
 2036 leak, via  $\{h_i\}_{i \in [\text{JSIN}]}$ , the fact that the input notes in both transactions were spent with  
 2037 the same *ask<sub>i</sub>* (if that were the case).

2038 More particularly, using the input notes’ nullifiers to derive *hsig* ensures that *hsig* is  
 2039 unique with overwhelming probability for all *accepted* transaction. Furthermore, us-  
 2040 ing *randomSeed* ensures the uniqueness of *hsig* for transactions *in transit* (as before  
 2041 validation there may be several in transit transactions with the same set of nullifiers).

## 2042 **A.2.3 Solution on Ethereum**

2043 As described in the Ethereum yellow paper [Woo19, Appendix F], Ethereum transactions  
 2044 are ECDSA signed. Further, as described in Section 2.3, the one-time signature used to  
 2045 sign the Mix data also signs the Ethereum address used to sign the transaction. As such,  
 2046 any modification to the transaction object will result in a new transaction hash, and  
 2047 any attempt to sign the transaction with a different ECDSA key will be rejected by the  
 2048 **Mixer** contract (see: Section 2.5). We thereby conclude that the one-time signature  
 2049 used to sign the transaction data, does not need to be SUF-CMA, but *only needs to*  
 2050 *achieve* UF-CMA.

2051 Specifically, carrying out any change on the one-time signature will change the  
 2052 **Ethereum** transaction data and result in a failure to verify the ECDSA signature of  
 2053 the **Ethereum** transaction. To obtain a new valid signature on this transaction, the  
 2054 adversary needs to break the UF-CMA property of the ECDSA signature scheme or use  
 2055 another ECDSA keypair to sign the transaction. In the last case, the one-time signature  
 2056 will no longer be valid.

2057 Note that including the Ethereum transaction sender in the data to be signed by the  
 2058 one-time signature scheme also addresses the possible attack described in Remark 15.

2059 An attacker trying to resign the same Ethereum transaction with a different key will  
 2060 cause **Mixer** to reject the transaction when the one-time signature is checked.

**Remark 17.** *We note that the transaction malleability issue can also be addressed in another way. In fact, one could use the ECDSA signatures on **Ethereum** transactions to fix all inputs and ciphertexts, and then tie the sender of the **Ethereum** transaction to the zk-snark by putting the sender address  $\mathcal{S}_{\mathcal{E}}.\text{Addr}$  in  $h\text{sig}$ . In other words, it is also possible to define  $h\text{sig}$  as:*

$$h\text{sig} = \text{CRH}(\{\text{nf}_i\}_{i \in [\text{JSIN}]}, \mathcal{S}_{\mathcal{E}}.\text{Addr})$$

2061 *As such, if an attacker extracts the ciphertexts, in a  $\text{tx}_{\text{Mix}}$  transaction, to craft another*  
 2062 *malicious transaction  $\text{tx}_{\text{Mix}}$ , the key-pair used to sign  $\text{tx}_{\text{Mix}}$  differs from the one used*  
 2063 *to sign  $\text{tx}_{\text{Mix}}$ , which changes the transaction sender address recovered on **Mixer**. As a*  
 2064 *consequence, the check on  $h\text{sig}$  would fail on the **Mixer**, invalidating the transaction,*  
 2065 *and preventing the attack.*

2066 *While such solution would avoid the need to generate one-time signing keys and*  
 2067 *could avoid a signature check in the **Mixer**, it would also require every **Zeth** user to*  
 2068 *have an Ethereum account. Doing so, would be a major hindrance toward the design of*  
 2069 *mechanisms aiming to provide anonymity to **Zeth** transactions initiators. In fact, the*  
 2070 *addressing scheme used in **Zeth** along with the solution to the malleability introduced in*  
 2071 ***Zcash** makes it possible to generate raw **Zeth** transactions without having an Ethereum*  
 2072 *account. These raw transactions could then be broadcasted — to a set of Ethereum user*  
 2073 *nodes — on an anonymous p2p network, before being finalized and submitted to the*  
 2074 *Ethereum network by Ethereum users who would be rewarded according to an incentive*  
 2075 *structure. While such protocol is outside of the scope of this document, it shows that*  
 2076 *defining  $h\text{sig}$  using the senders address alters the flexibility of **Zeth**; hence why this*  
 2077 *solution has not been favored.*

## Appendix B

# Double Spend Attack on Equivalent class

The primary inputs of our zk-snarks are elements of  $\mathbb{F}_{r_{BN}}$  and they can be written over  $\text{BNFIELDLEN}$  bits. Because  $r_{BN}$  is an odd prime number and we have that  $\text{encode}_{\mathbb{F}_{r_{BN}}}(\mathbb{F}_{r_{BN}}) \subset \mathbb{B}^{\text{BNFIELDLEN}}$ , the projection of  $\mathbb{B}^{\text{BNFIELDLEN}}$  in  $\mathbb{F}_{r_{BN}}$  is surjective.

When we pass the primary inputs to the Mixer contracts, they are interpreted as elements of  $\mathbb{B}^{\text{WORDLEN}}$ ,  $\mathbb{B}^{\text{BNFIELDLEN}} \subset \mathbb{B}^{\text{WORDLEN}}$ . As previously noted, this means that there exists elements of  $\mathbb{B}^{\text{WORDLEN}}$  with the same projection in  $\mathbb{F}_{r_{BN}}$ . An adversary could make use of this to perform a double spend attack.

Indeed, to check that a coin is not double spent, the contract stores the nullifiers of spent coins (as elements of  $\mathbb{B}^{\text{WORDLEN}}$ ) and verifies that the nullifier of the coin to be spent is not stored. The adversary could thus modify the nullifier to a different value with the same projection. As the snark verification operates in  $\mathbb{F}_{r_{BN}}$ , the proof would still be valid. However, the value stored for this nullifier would be different from the adversarial one. Hence, the nullifier would be validated, the transaction would succeed and the coin would be double spent. In practice, the adversary can perform the attack by simply adding  $r_{BN}$  to one of the elements representing the nullifier.

To prevent this attack, the contract checks that all primary inputs are elements of  $\mathbb{F}_{r_{BN}}$ , that is to say that they are smaller than  $r_{BN}$ . As one may see, the attack describe above is not due to the necessary packing of hash digests into field elements but to the contract storage of field elements.

## 2100 Appendix C

# 2101 Side-Channel Attacks and 2102 Information Leaks

2103 The following subsections describe several side-channel attacks and possible weaknesses  
2104 that implementors should be aware of and attempt to mitigate.

2105 We consider cases in which the attacker is able to observe the RPC communications  
2106 between **Zeth** client software, and Ethereum P2P nodes. This situation may occur if an  
2107 observer is able to monitor the network traffic between the Ethereum node and the **Zeth**  
2108 client software, or if the Ethereum node itself is compromised.

### Note

In this discussion, we do not consider adversaries with physical access to the machine running the client software. Such adversaries could make precise measurements of timing, power-consumption or other physical quantities that could reveal fine-grained details of the operations being carried out by the software, or the data it is operating on. Protecting against attacks of this kind often involves implementation techniques such as: avoiding branches based on private data, using only predictable memory access patterns, and adding “noise” to the information leaked on certain side-channels. We leave consideration of these attacks and prevention methods for a future discussion.

2109

## 2110 C.1 Counterfeit Data

2111 Malicious Ethereum nodes or attackers able to compromise the network, have the op-  
2112 portunity to send invalid data to RPC clients. This could be used to inject invalid data  
2113 into the client’s record of state, which could prevent it from generating valid Mix calls  
2114 or allow it to be identified in the future. In general, data from any remote host should  
2115 be treated as malicious, unless accompanied by evidence that convinces the client of its  
2116 authenticity.

2117 In the case of Ethereum event logs (the main source of data used to track the on-  
2118 chain state - see: Section 4.1.1 for details), clients **MUST** leverage the consensus evidence  
2119 and block headers to verify that log data it receives is genuine and has been committed  
2120 to the blockchain. See: Section 1.2.3 for further information about how such data is  
2121 secured.

## 2122 C.2 Data Leaked during Synchronization

2123 In order to receive private payments and keep client data up-to-date, **Zeth** client software  
2124 must scan the blockchain and process *all* the event data emitted by **Mixer** during  
2125 Mix calls (as described in Section 4.1.1). There are several issues to consider when  
2126 determining exactly how and when this “synchronization” takes place.

2127 Client implementations that only connect to the RPC endpoint in response to user  
2128 input, or in preparation for performing a Mix call, may leak information. Observers may  
2129 deduce that such client are likely to be the recipient of a recent or upcoming transaction,  
2130 or that they may be about to perform a Mix call.

2131 Similarly, payment provider software that only listens for events when awaiting a  
2132 transaction, and remains disconnected otherwise, may reveal that it is the recipient of  
2133 an upcoming transaction, and possibly *which* transaction or block it was paid by (based  
2134 on when it stops listening).

2135 Further, consider wallet software that performs RPC operations to explicitly wait for  
2136 the Ethereum transaction corresponding to a specific Mix call. (This most likely to be  
2137 for transactions it has sent, to inform the user and update the wallet state once payment  
2138 is complete, but could possibly happen on the receiver if he somehow knows the ID of  
2139 the transaction of interest - e.g. via off-chain communication with the sender). If such  
2140 a *wait* procedure is implemented by querying the state of a specific transaction by its  
2141 ID, or by listening for blocks *until* the transaction of interest is received, the connected  
2142 Ethereum node may infer that this client is interested in the transaction, and likely to  
2143 be the sender or recipient.

2144 Consider a client which periodically connects to some Ethereum node and requests  
2145 all relevant data from the last block it saw, up to the latest block available. Each client  
2146 will have information up to some block  $n$  (where  $n$  varies per client), and  $n$  is known to  
2147 the Ethereum node that served the client. The client could then potentially be identified  
2148 by  $n$  (even if it hides its IP for each connection) since a client that connects and queries  
2149 **Zeth** transactions from block  $n + 1$  reveals that it is one of the clients who synced up to  
2150 block  $n$  when it last connected.

2151 Note that, if the client always broadcasts the mix transaction via this same Ethereum  
2152 node, then the Ethereum node may already deduce that the client is the sender. However,  
2153 implementations may wish to use techniques (such as sending transactions from other  
2154 nodes or hiding their IP address in other way) to obfuscate any relationship between  
2155 transactions and the clients that originated them.

## C.3 Queries on Successful Decryption

The event data emitted by Mixer contains the note data for new commitments, encrypted using a key derived from the recipients public key. As described in Section 2.6, clients scan the blockchain for these events and attempt to decrypt the ciphertext using their secret decryption keys. If they are successful, they are the recipient of the note and can try to parse the plaintext to extract the secret note data.

When decryption is successful and the note data has been extracted from the plaintext (we discuss parsing failure in Appendix C.4), clients **MUST** check that this note data does indeed open the commitment for the note.

A naive implementation of this check could query the state of Mixer via RPC to check the relevant entry in the set of commitments. However, this would reveal to an observer that the client had successfully decrypted and parsed the corresponding ciphertext, and was therefore the recipient of that note.

For this reason, the protocol specifies that Mixer **MUST** emit events informing clients of new commitment values and locations in the Merkle tree. Clients **MUST** consume *all* such data to maintain their view of contract state (as described in Appendix C.2). Further, clients **MUST** attempt to decrypt *all* ciphertexts and, for successful decryptions, **MUST** verify that the plaintext opens the note’s commitment. This avoids the need for any extra RPC queries that would reveal which ciphertexts were successfully decrypted.

## C.4 Invalid Ciphertext

The attack described in [TBP20, Section 4.2.1] illustrates the importance of correctly handling invalid data in client software. A so-called “REJECT Attack” is described whereby an attacker creates a Mix call with specially crafted ciphertext. The ciphertext can be successfully decrypted by the correct recipient, (that is, the plaintext note is encrypted with an encryption key derived from the recipients public key), but the corresponding plaintext is invalid and cannot be parsed correctly by the recipient.

### Note

Note that the above is possible because the plaintext is not verified either by the  $\mathbf{R}^Z$  circuit, or by the contract (which is unable to decrypt it). Hence, Zeth allows such transactions with malicious ciphertexts to be accepted by the Mixer contract, and clients must handle this case with care.

In the case described in [TBP20], there is no distinction between wallet and the underlying P2P nodes. Before a fix was applied (see: [?]), nodes explicitly rejected transactions of the above form, proving to their peers that they were able to decrypt the ciphertext and were therefore the intended recipient.

In Zeth, P2P nodes and wallet software are separated, so there will be no explicit rejection of the transaction. However, careless error handling (such as exceptions which causes the RPC connection to be closed) could potentially be detected by the connected

2190 Ethereum node. As in the “REJECT Attack”, this reveals that the connected RPC  
2191 client is the intended recipient of a transaction, and the owner of the corresponding  
2192 encryption key.

## 2193 C.5 Using (and Retrieving) Nullifiers

2194 Any non-trivial wallet implementation will need to track which of the user’s **Zeth** notes  
2195 have been spent, and which are still available. Naturally, the wallet software could mark  
2196 the notes as it broadcasts transactions that spend them. However, this approach is  
2197 subject to several problems.

2198 Firstly, for each note spent, the client software must record the ID of the spending  
2199 transaction, in order to track it and confirm that it is accepted into a block. Once each  
2200 spending transaction is accepted the client can finally mark the appropriate **Zeth** notes  
2201 as “spent”. This requires significant complexity in order to asynchronously mark the  
2202 notes, and to deal with the issues described in Appendix C.2.

2203 Secondly, this approach does not support multiple wallets using the same key, or  
2204 wallets being restored from **Zeth** addresses. A user that wishes to rebuild his wallet  
2205 (see the discussion in Section 4.1.4), or check for any spending activity by other wallets,  
2206 would not be able to do so by simply scanning the blockchain.

2207 By using the nullifiers passed to **Mix** calls, clients can determine the availability of  
2208 notes in a more robust way. That is, to determine whether a note is spent or available,  
2209 the client can compute the nullifier and check whether that nullifier has been seen by  
2210 the **Mixer** contract.

2211 In a similar way to Appendix C.3, queries to **Mixer** for specific nullifiers reveals  
2212 to observers that the client was the sender of any previous or future transaction that  
2213 generates such a nullifier. To mitigate this, **Mixer** MUST include nullifier values in the  
2214 event data it emits, and clients SHOULD use this to track which of their notes are spent.  
2215 This MUST happen as part of the regular sync operation, so that no extra RPC traffic is  
2216 generated and observers cannot distinguish between clients that do and do not recognize  
2217 any given nullifier. Note that this approach also supports tracking spent notes from  
2218 multiple wallets, and rebuilding wallets by re-syncing the blockchain.

## 2219 C.6 Proof Generation

2220 Generation of the zero-knowledge proofs, required for valid **Mix** calls, is a very compu-  
2221 tationally intensive process. In itself, the proof generation does not require any com-  
2222 munication with external parties, and so may not directly leak information about the  
2223 client, but implementors should consider some indirect ways in which information may  
2224 be leaked.

2225 Implementors may also wish to consider the possible indirect impact of proof gen-  
2226 eration on the RPC channel. For example, a client that “waits” for proof generation  
2227 without servicing the RPC connection may fail to respond to, or take significantly longer



2228 to respond to, new log events, allowing the connected Ethereum node to deduce that it  
2229 is generating a proof and therefore likely to be the sender of an upcoming transaction.

#### Note

As stated in the introduction to this chapter, this discussion does not consider general timing attacks. We mention this extreme case, of a client that completely stalls during proof generation, only to illustrate how a poor implementation may leak information to its RPC peer.

2230

2231 In the case where proof generation is carried out on some external host, or by an  
2232 external process on the same host, there may be a risk of network traffic or other IPC  
2233 traffic being observed. An observer able to detect that a given client is communicating  
2234 with a prover process can reliably deduce that it will be the sender of an upcoming  
2235 transaction.

2236 An observer able to see the content of the communication between the wallet and  
2237 prover process will also gain knowledge of the auxiliary inputs to the proof (including  
2238 the data required to spend the input notes and secret attributes of the output notes).  
2239 It is therefore important to secure any such connection, protect any prover process from  
2240 being maliciously modified or observed, and to ensure that wallets only communicate  
2241 with trusted processes.

## 2242 C.7 Simple Mixer Calls

2243 The public parameters to a Mix call can reveal information about the nature of a trans-  
2244 action, even though they do not reveal recipient details or note amounts. For example,  
2245 a Mix call in which  $\text{Mix}_{in}.primIn.vout == 0$  and  $\text{Mix}_{in}.primIn.vin \neq 0$  may indicate  
2246 a simple “deposit” of funds into the mixer. Similarly, if both  $\text{Mix}_{in}.primIn.vout$  and  
2247  $\text{Mix}_{in}.primIn.vin$  are zero, the transaction must be spending only notes already within  
2248 **Mixer**, into new notes. Finally, if  $\text{Mix}_{in}.primIn.vin == 0$  and  $\text{Mix}_{in}.primIn.vout \neq 0$ ,  
2249 the sender may be performing a simple “withdrawal” of funds from some existing notes.

2250 A Mix call can combine all of the above logical operations in a single transaction.  
2251 That is, it can deposit value into the mixer, spend existing notes, create new notes, and  
2252 withdraw value from **Mixer** at the same time. Combining logical operations in this way  
2253 makes it much more difficult for observer to attribute a specific purpose to the Mix call.

2254 Clients can also perform Mix calls in which  $vin = vout = 0$  and 0-valued notes are  
2255 created from other 0-valued notes. Such “dummy” self-payments can further obfuscate  
2256 the activity of a wallet, by adding “noise” to the system. Note, however, that the gas  
2257 cost for such transactions must still be paid.

2258 Wallet implementations **SHOULD** encourage the use of these complex calls where pos-  
2259 sible, either via the user interface or by automatically adding complexity to transactions,

2260 and **SHOULD** support features such as adding “noise”<sup>1</sup> if the user wishes to pay for extra  
2261 protection of this kind.

### 2262 C.7.1 Small Anonymity Sets

2263 Until there is a large number of commitments and users of the mixer, it may be easy  
2264 for an observer to infer some of the private data that is intended to be hidden by mixer  
2265 calls.

2266 In the simple case, if there are very few commitments in  $\widetilde{\text{Mixer}}$ ’s Merkle tree, an  
2267 attacker has a small list of candidate commitments that are being spent by subsequent  
2268 Mix calls. Similarly, if the number of distinct Ethereum addresses that have been used  
2269 to call  $\widetilde{\text{Mixer}}$  is very small, observers can trace the original source of funds subsequently  
2270 withdrawn, to a small set of original depositors.

2271 Client software may wish to track metrics about the  $\widetilde{\text{Mixer}}$  state, and either prevent  
2272 certain actions or design the user interface to discourage users<sup>2</sup> from creating transactions  
2273 whose features can be identified with high probability.

- 2274 • **Number of Commitments.** If there is a low absolute number of commitments,  
2275 clearly any non-zero output must spend one of these. (Although we note that only  
2276 *vout* can be publicly known to be non-zero).
- 2277 • **Number of Unspent Commitments.** If  $\#Comms - \#Nulls$  is small, and a  
2278 new commitment is created and then spent, observers can deduce that there is a  
2279 high chance that the spend operation targeted the new commitment.
- 2280 • **Number of Ethereum Addresses.** While very few distinct addresses (or groups  
2281 of addresses that are not associated) have used the contract, observers can deduce  
2282 that subsequent Mix calls are likely to spend commitments created by clients as-  
2283 sociated with one of this small set of addresses.

2284 The set of Ethereum addresses that have interacted with the contract can leak data  
2285 in other ways. An Ethereum address that withdraws value from the contract, but has not  
2286 previously been used to make a Mix call (or a Mix call that deposits value into  $\widetilde{\text{Mixer}}$ ),  
2287 must have been the recipient of zeth notes created by a previous depositor. The details  
2288 may not be directly available to an observer, but this is another example of information  
2289 which could be combined with other leaked data to infer connections between entities  
2290 and transactions.

---

<sup>1</sup>By randomly scheduling dummy payments, for instance

<sup>2</sup>By, for example, displaying warning messages and/or asking the user for confirmation

## Appendix D

# Security proofs of Blake2

This appendix proves the collision resistance, PRF-ness, binding and hiding properties of the Blake2 hash function in the Weakly Ideal Cipher model (WICM, see [LMN16]). The proofs use definitions and results of Luykx et al. [LMN16], regarding the indistinguishability of Blake2 and a random oracle in the Weakly Ideal Cipher Model (WICM). Assuming Blake2s is as secure as Blake2, these results may be applied to Blake2s.

### D.1 Security model of Blake2

The security analysis treats Blake2 as hash function built on top of a block-cipher-based compression function in the WICM (which derives from the Ideal Cipher Model). In this section, we present the WICM and prove that Blake2 is a collision resistant PRF, and thus a commitment scheme.

#### D.1.1 Weakly Ideal Cipher Model

The research community believes that Blake’s underlying block cipher has no known weaknesses and could reasonably be modeled as an ideal cipher [LMN16, Section 2.1]. However, Blake2 admits weak keys with a specific structure [LMN16, Section 2.1]. Blake2 is therefore more appropriately analysed in the WICM, which is an extension of the Ideal Cipher Model that represents a block cipher as a set of independent random permutations [HKT11]. The WICM may also be viewed as a specialization for Blake2 of the Weak Cipher Model [MP15], which aims to be realistic by modeling particular characteristics, invariants or properties a block cipher may have.

A number of definitions in what follows are quoted directly from Luykx et al. [LMN16].

**The Weakly Ideal Cipher Model.** Let  $\mathcal{W}$  and  $\mathcal{S}$  be the following partition of  $\mathbb{B}^{2 \cdot \text{ol}}$  into weak and strong sets, where  $w$  is the word length ( $16 \cdot w = 2 \cdot \text{ol}$ ):

$$\begin{aligned}\mathcal{W} &= \left\{aaaabbbbccccdddd \in \mathbb{B}^{2 \cdot \text{ol}} \mid a, b, c, d \in \mathbb{B}^w\right\} \\ \mathcal{S} &= \mathbb{B}^{2 \cdot \text{ol}} \setminus \mathcal{W}\end{aligned}$$

Let  $\mathcal{BLC}(2 \cdot \text{ol}, 2 \cdot \text{ol})$  denote the set of all block ciphers  $E : \mathbb{B}^{2 \cdot \text{ol}} \times \mathbb{B}^{2 \cdot \text{ol}} \rightarrow \mathbb{B}^{2 \cdot \text{ol}}$ . Define  $\mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  as the set of all block ciphers  $E \in \mathcal{BLC}(2 \cdot \text{ol}, 2 \cdot \text{ol})$  with the additional restriction that  $E(k_w, \cdot)$  is  $\mathcal{W}$ - and  $\mathcal{S}$ -subspace invariant for all keys  $k_w \in \mathcal{K}_{\text{weak}}$ . That is, inputs in  $\mathcal{W}$  map to  $\mathcal{W}$ , and likewise for  $\mathcal{S}$ . Here,  $\mathcal{K}_{\text{weak}}$  is the set of weak keys, defined as

$$\mathcal{K}_{\text{weak}} = \left\{ k = \text{kkkkkkkkkkkkkkkk} \in \mathbb{B}^{2 \cdot \text{ol}} \mid k \in \mathbb{B}^w \right\}.$$

2313 A random  $E \leftarrow \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  can now be modeled as follows:

- 2314 • on input of  $(k, x) \in \mathcal{K}_{\text{weak}} \times \mathcal{W}$ ,  $E$  generates its response  $y$  randomly from  $\mathcal{W}$  up  
2315 to repetition;
- 2316 • on input of  $(k, x) \in \mathcal{K}_{\text{weak}} \times \mathcal{S}$ ,  $E$  generates its response  $y$  randomly from  $\mathcal{S}$  up to  
2317 repetition.

2318 For key values  $k \in \mathbb{B}^{2 \cdot \text{ol}} \setminus \mathcal{K}_{\text{weak}}$ ,  $E$  behaves like an ideal cipher: it either outputs a  
2319 new random value or if the key-message-image tuple has already been queried the tuple's  
2320 image. The case of inverse queries is analogous.

Blake2C is defined over the following domains and codomain:

$$\text{Blake2C} : \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol}) \times \mathbb{B}^{\text{ol}} \times \mathbb{B}^{2 \cdot \text{ol}} \times \mathbb{B}^{\text{ol}/4} \times \mathbb{B}^{\text{ol}/4} \rightarrow \mathbb{B}^{\text{ol}}$$

2321 We write  $\text{Blake2C}_E(h, m, t, f)$  for the output of the Blake2 compression function, defined  
2322 over encryption scheme  $E$  on inputs  $h$ ,  $m$ ,  $t$  and  $f$ . The compression function, in par-  
2323 ticular, computes the state  $x = (h \parallel \text{pad}_{\text{ol}/2}(0) \parallel t \parallel f) \oplus (\text{pad}_{\text{ol}}(0) \parallel \text{IV})$  for some  $\text{IV}$ . It then  
2324 encrypts  $x$  under  $m$  (where  $m$  is treated as a key for the encryption) and splits  $E(m, x)$   
2325 in two same size variables, the left part  $l_E$  and right part  $r_E$ . It finally outputs  $l_E \oplus r_E \oplus h$ .

2327 Zethuses the Blake2 compression function with a fixed encryption scheme  $E^*$  based on  
2328 ChaCha stream cipher [Ber08a]. Accordingly we write  $\text{Blake2C}(h, m, t, f) = \text{Blake2C}_{E^*}(h, m, t, f)$ .

Blake2 is defined over the following domain and codomain:

$$\text{Blake2} : \mathbb{B}^{\leq 2 \cdot \text{ol}} \times \mathbb{B}^* \rightarrow \mathbb{B}^{\text{ol}}$$

2329 **Indifferentiability.** One way to measure the extent to which a certain cryptographic  
2330 function behaves like a random function is via the indistinguishability framework, where  
2331 a distinguisher is given oracle access to either the cryptographic function or the random  
2332 function, with the goal of determining which one it has access to.

**Definition 33.** Let  $\mathcal{C}$  be a construction with oracle access to an ideal primitive  $\mathcal{P}$ . Let  $\mathcal{R}$  be an ideal primitive with the same domain and codomain as  $\mathcal{C}$ . Let  $\text{Sim}$  be a simulator with the same domain and codomain as  $\mathcal{P}$  with oracle access to  $\mathcal{R}$ , and let  $\text{Dist}$  be a PPT distinguisher. The indifferentiability advantage of  $\text{Dist}$  is defined as:

$$\text{Indiff}_{\mathcal{C}^{\mathcal{P}}, \text{Sim}}(\text{Dist}) = \left| \Pr \left[ \text{Dist}^{\mathcal{C}^{\mathcal{P}}, \mathcal{P}} = 1 \right] - \Pr \left[ \text{Dist}^{\mathcal{R}, \text{Sim}^{\mathcal{R}}} = 1 \right] \right|$$

2333 The distinguisher  $\text{Dist}$  can query both its left oracle (either  $\mathcal{C}$  or  $\mathcal{R}$ ) and its right  
 2334 oracle (either  $\mathcal{P}$  or  $\text{Sim}$ ). We refer to  $\mathcal{C}^{\mathcal{P}}$ ,  $\mathcal{P}$  as the real world, and to  $\mathcal{R}$ ,  $\text{Sim}^{\mathcal{R}}$  as the  
 2335 simulated world; the distinguisher  $\text{Dist}$  converses with either of these worlds and its goal  
 2336 is to tell both worlds apart.

**Theorem 3** (Indifferentiability of Blake2 [LMN16]). *Let an encryption scheme  $E \leftarrow_{\$} \mathcal{BCK}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  be a weakly ideal cipher, and consider the hash function  $\text{Blake2}_E$  that internally uses  $E$ . There exists a simulator  $\text{Sim}$  such that for any distinguisher  $\text{Dist}$  with total complexity  $q$ , we have:*

$$\text{Indiff}_{\text{Blake2}_E, \text{Sim}}(\text{Dist}) \leq \frac{\binom{q}{2}}{2^{2\text{ol}}} + \frac{2\binom{q}{2}}{2^{\text{ol}}} + \frac{q}{2^{\text{ol}/2}}$$

2337 where  $\text{Sim}$  makes at most  $O(q^3)$  queries to a random function  $\mathcal{R}$ .

2338 *Proof.* See: [LMN16, Corollary 1]. □

2339 For asymptotic security, we assume the distinguisher to be PPT and that the number  
 2340 of queries done is polynomial  $q \leq \text{poly}(\text{ol})$ .

2341 **Additional remarks.** Luykx et al. [LMN16] remark that, by resorting to the WICM,  
 2342 they do not make stronger assumptions than those used in previous results (ICM), and,  
 2343 despite the fact that they give distinguishers more power (by weakening the cipher),  
 2344 they are able to get similar results.

## 2345 D.2 Security proofs

### 2346 D.2.1 Proof of Blake2 function's PRFness

2347 Luykx et al. already prove the PRFness of Blake2 *keyed* hash function in the multi-key  
 2348 setting.

**Definition 34** (PRF in multi-key setting [ML15]). *Let  $\mu \geq 1$  and  $k \leftarrow_{\$} \mathcal{K}^{\mu}$ . Let  $\mathcal{C}$  be a keyed construction with key space  $\mathcal{K}$  and with oracle access to an ideal primitive  $\mathcal{P}$ . Let  $\mathcal{R}_1, \dots, \mathcal{R}_{\mu}$  be random functions with the same domains and ranges as  $\mathcal{C}_{k_1}, \dots, \mathcal{C}_{k_{\mu}}$ . Let  $D$  be a distinguisher. The PRF distinguishing advantage of  $D$  is defined as,*

$$\text{PRF}_{\mathcal{C}^{\mathcal{P}}}(\mathcal{D}) = \left| \Pr \left[ \text{Dist}_{k_1, \dots, k_{\mu}, \mathcal{P}}^{\mathcal{C}^{\mathcal{P}}} = 1 \right] - \Pr \left[ \text{Dist}_{\mathcal{R}_1, \dots, \mathcal{R}_{\mu}, \mathcal{P}} = 1 \right] \right|$$

Blake2 supports keyed hashing by simply prepending the key to the message:

$$\text{Blake2}_{E,k}(m) = \text{Blake2}_E(k \| 0^{2\text{ol}-\text{kl}} \| m)$$

2349 where  $\text{kl} \leq 2\text{ol}$  denotes the key size. In other words, the key gets processed as other data  
 2350 and the HAIFA counter and flags are designated to the key in a similar fashion as if they  
 2351 were for normal data blocks.

**Theorem 4** (PRF-Security of Blake2 keyed mode [LMN16]). *Let  $\mu \geq 1$  and let  $k \leftarrow_{\$} (\mathbb{B}^{kl})^\mu$ . Let an encryption scheme  $E \leftarrow_{\$} \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  be a weakly ideal cipher, and consider the keyed hash function  $\text{Blake2}_{E,k}$  that internally uses  $\text{Blake2C}_E$  that internally uses  $E$ . For any distinguisher  $\text{Dist}$  with total complexity  $q$ :*

$$\text{PRF}_{\text{Blake2}_{E,k}}(\text{Dist}) \leq \frac{\binom{q}{2}}{2^{2\text{ol}}} + \frac{2\binom{q}{2}}{2^{\text{ol}}} + \frac{q}{2^{\text{ol}/2}} + \frac{\mu q}{2^{kl}} + \frac{\binom{\mu}{2}}{2^{kl}}$$

2352 *Proof.* See: [LMN16, Corollary 3]. □

2353 **Remark 18.** *We can note that in the case of keyed hashing, the key is padded only to be*  
 2354 *processed in a single block to differentiate the key from the message. The security proof*  
 2355 *of Theorem 4 does not rely on this padding and as such also works with no padding.*

**Theorem 5** (PRF-Security of Blake2 with key [LMN16]). *Let  $k \leftarrow_{\$} \mathbb{B}^{kl}$ . Let an encryption scheme  $E \leftarrow_{\$} \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$  be a weakly ideal cipher, and consider the keyed hash function  $\text{Blake2}_E(k, \cdot) = \text{Blake2}_E(k \parallel \cdot)$  that internally uses  $\text{Blake2C}_E$  that internally uses  $E$ . For any distinguisher  $\text{Dist}$  with total complexity  $q$ :*

$$\text{PRF}_{\text{Blake2}_E}(\text{Dist}) \leq \frac{\binom{q}{2}}{2^{2\text{ol}}} + \frac{2\binom{q}{2}}{2^{\text{ol}}} + \frac{q}{2^{\text{ol}/2}} + \frac{q}{2^{kl}}$$

2356 *Proof.* See: Remark 18 and Theorem 4 with  $\mu = 1$ . □

## 2357 D.2.2 Proof of Blake2 collision resistance

2358 We want to prove here the collision resistance of Blake2. To do so, we are going to prove  
 2359 by contradiction that if Blake2 is not collision resistant, it is not indifferentiable.

2360 **Theorem 6.** *Blake2 is collision resistant.*

2361 *Informal proof.* Let us assume that there exists a PPT adversary  $\mathcal{B}$  which breaks the  
 2362 collision resistance of Blake2. We build an adversary  $\mathcal{A}$  that uses this adversary to  
 2363 differentiate between the real and simulated worlds. More particularly,  $\mathcal{A}$  gets left and  
 2364 right oracles (see [LMN16, Figure 3]), which are either an oracle for a hash function and  
 2365 for a weakly ideal block cipher or a random oracle and an encryption simulator with  
 2366 oracle access to the random oracle.

2367 On each  $\mathcal{B}$ 's query  $m_i$ ,  $i \in \{1, \dots, q\}$ ,  $\mathcal{A}$  passes them to his left oracle and returns  
 2368 the answer  $h_i$  to  $\mathcal{B}$ . Eventually, if  $\mathcal{B}$  finds a collision, that is a pair  $(m_i, m_j)$  such that  
 2369  $i \neq j$  and  $h_i = h_j$ ,  $\mathcal{A}$  guesses that his oracles were real; else  $\mathcal{A}$  returns a random guess.  
 2370 Otherwise  $\mathcal{A}$  guesses his oracles were simulated. If the left oracle was a random oracle,  
 2371 the probability of finding a collision would be negligible for  $q \leq \text{poly}(\lambda)^1$ .

---

<sup>1</sup>The probability would be  $\frac{q^2}{2^{\text{ol}}}$  which is negligible for a polynomial number of queries  $q$ . This is the sum of the probabilities of finding a collision when doing the  $i^{\text{th}}$  query. Indeed, let us suppose the adversary has done  $i - 1$ ,  $i > 2$ , queries without finding a collision, i.e. he knows  $i - 1$  distinct tuples of input-output. When receiving the  $i^{\text{th}}$  value, the adversary has thus  $i - 1$  chance to find a collision. The probability for the new output to be equal to any of the previous outputs is thus  $(i - 1) \cdot \frac{1}{2^{\text{ol}}}$  (as we are in the random oracle model). Summing this probability over all queries, we find the probability of finding a collision after doing  $q$  queries.

2372 On the other hand, as assumed  $\mathcal{B}$  finds a collision with non-negligible probability  
 2373 if the oracles were real. Hence,  $\mathcal{A}$  wins the indistinguishability game with non-negligible  
 2374 advantage what is a contradiction.  $\square$

### 2375 D.2.3 Blake2 as a commitment scheme

2376 We prove here that Blake2 is a commitment scheme, i.e. is binding and hiding. To do so  
 2377 we rely on the previous results that Blake2 is collision resistant and a PRF.

2378 **Theorem 7.** *Let  $E \leftarrow \$_{\mathcal{B}\mathcal{L}\mathcal{K}}(2\text{ol}, 2\text{ol})$  and for a message  $x \in \mathbb{B}^*$  and randomness  $r \in \mathbb{B}^{\text{rl}}$   
 2379 commitment to  $x$  using  $r$  be  $\text{ComSch.Com}(x; r) = \text{Blake2}_E(r \| x)$ . Then ComSch is hiding  
 2380 and binding.*

2381 *Informal proof. Hiding.* A commitment scheme ComSch is computationally hiding if,  
 2382 knowing two potential openings, a PPT adversary cannot distinguish which was com-  
 2383 mitted. Let us assume that there exists a PPT adversary  $\mathcal{B}$  which breaks the hiding  
 2384 property of Blake2 with a non-negligible advantage  $\eta$ . We build an adversary  $\mathcal{A}$  that  
 2385 uses  $\mathcal{B}$  to break the PRF property of Blake2 with advantage  $\eta/2$ .

2386 First, the PRF game is initiated, that is, the challenger chooses a random encryption  
 2387 scheme  $E$  and key  $k \in \mathbb{B}^{\text{rl}}$  and instantiates two oracles  $O^{\text{Blake2}_k} = \text{Blake2}_E(k, \cdot)$  and  $O^R$   
 2388 a random function. The challenger picks an oracle randomly and gives  $\mathcal{A}$  access to it.  
 2389  $\mathcal{B}$  sends  $q$  oracle queries to  $\mathcal{A}$  who passes them to his left oracle and returns the oracle  
 2390 answers to him.  $\mathcal{B}$  then outputs two challenge messages  $(m_0, m_1)$  and sends them to  $\mathcal{A}$   
 2391 who randomly message selects  $m_b$ , sends it to his left oracle, gets back  $y_b$  as an answer  
 2392 and sends  $y_b$  to  $\mathcal{B}$ .  $\mathcal{B}$  returns the decision bit  $\tilde{b}$  to  $\mathcal{A}$ . If  $b = \tilde{b}$ ,  $\mathcal{A}$  answers to the challenger  
 2393 that the oracle was instantiating the PRF. Otherwise,  $\mathcal{A}$  answers with a random guess.  
 2394 The advantage of  $\mathcal{A}$  equals advantage of  $\mathcal{B}$  if it interacts with a real hash function. The  
 2395 advantage of  $\mathcal{A}$  equals half the advantage of  $\mathcal{B}$  when interacting with a random oracle  
 2396 and simulator.

2397 *Binding.* A commitment scheme ComSch is said to be computationally binding if  
 2398 it is infeasible to find  $x, x'$  and  $r, r'$  such that  $x \neq x'$  and  $\text{Com}(x; r) = \text{Com}(x'; r')$ .  
 2399 This is implied by collision resistance of Blake2. Thus if  $\mathcal{B}$  is an algorithm that breaks  
 2400 the binding property with advantage  $\eta$ , there is another algorithm  $\mathcal{A}$  that breaks Blake2  
 2401 collision resistance with the same advantage.  $\square$

2402 Assuming that Blake2s is as secure as Blake2, a commitment scheme based on a  
 2403 Blake2s, i.e.  $\text{Com}(x; r) = \text{Blake2s}_E(r \| x)$  is hiding and binding.

## 2404 D.2.4 Proof of commitment scheme security

To prove the binding and hiding property of  $\text{ComSch}$  (see: Section 3.1.2), we introduce the following commitment scheme  $\text{ComSch}^*$ ,

$$\begin{aligned} \text{ComSch}^*.\text{Setup} : \{1^\lambda \text{ s.t. } \lambda \in \mathbb{N}\} &\rightarrow \mathbb{B}^* \\ \text{ComSch}^*.\text{Com} : \mathcal{BLC}^*(2 \cdot \text{BLAKE2sCLEN}, 2 \cdot \text{BLAKE2sCLEN}) \times \mathbb{B}^{2 \cdot \text{BLAKE2sCLEN}} \\ &\times (\mathbb{B}^{\text{PRFADDRROUTLEN}} \times \mathbb{B}^{\text{PRFRHOOUTLEN}} \times \mathbb{B}^{\text{ZVALUELEN}}) \times \mathbb{B}^{\text{RTRAPLEN}} \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \end{aligned}$$

The commitment scheme is defined as follows,

$$\begin{aligned} \text{ComSch}^*.\text{Setup}(1^\lambda) &= pp^* = \epsilon \\ \text{ComSch}^*.\text{Com}(m = (apk, \rho, v); r) &= cm \\ &= \text{Blake2E}^*(r \| apk \| \rho \| v) \end{aligned}$$

Given a commitment scheme  $\text{ComSch}^*$ , the bijective function  $\text{decode}_{\mathbb{N}}(\cdot)$  and  $p_\lambda \in \mathbb{N}$ , a prime which can be represented using  $\lambda$  bits, we define the commitment scheme  $\text{ComSch}'$  as follows:

$$\begin{aligned} \text{ComSch}'.\text{Setup}(1^\lambda) &= (\text{ComSch}^*.\text{Setup}(1^\lambda), p_\lambda) \\ \text{ComSch}'.\text{Com}(m; r) &= \text{decode}_{\mathbb{N}}(\text{ComSch}^*.\text{Com}(m; r)) \pmod{p_\lambda} \text{ for } m = (apk \| \rho \| v) \end{aligned}$$

2405 Note that  $\text{ComSch}$  is a particular instantiation of  $\text{ComSch}'$  where  $\text{E}^*$  is set as ChaCha  
2406 encryption scheme [Ber08a],  $k^*$  is a random key, and  $p_\lambda$  is  $r_{\text{BN}}$ .

2407 **Theorem 8** (Hiding). *If  $\text{ComSch}^*$  is hiding then  $\text{ComSch}'$  is hiding.*

2408 *Proof.* We prove it by contradiction i.e. we assume that there exists an adversary  $\mathcal{B}$   
2409 that breaks  $\text{ComSch}'$ 's hiding property and show an adversary  $\mathcal{A}$  that uses  $\mathcal{B}$  to break  
2410  $\text{ComSch}^*$ 's hiding property with non-negligible probability.

2411 Let  $\mathcal{C}$  be a challenger that sets up the hiding game for  $\text{ComSch}^*$  and  $\mathcal{A}$ . The adversary  
2412  $\mathcal{A}$ , given public parameters  $pp^*$  of  $\text{ComSch}^*$  and access to an oracle that runs the  $\text{Com}$   
2413 algorithm of  $\text{ComSch}^*$  scheme, simulates a hiding game for  $\text{ComSch}'$  for  $\mathcal{B}$ . The adversary  
2414  $\mathcal{A}$  starts by setting public parameters  $pp'$  for  $\text{ComSch}'$  using public parameters  $pp^*$   
2415 given by  $\mathcal{C}$ . Parameters  $pp'$  are passed to  $\mathcal{B}$  who outputs a pair of messages  $m_0, m_1$ .  
2416 The adversary  $\mathcal{A}$  forwards them to the challenger who samples a bit  $b$  at random and  
2417 generates  $cm^* = \text{ComSch}^*.\text{Com}(m_b; r)$  for some randomness  $r$ . The result is returned to  
2418  $\mathcal{A}$  (see: Definition 24). Then  $\mathcal{A}$  passes  $cm = \text{decode}_{\mathbb{N}}(cm^*) \pmod{p_\lambda}$  to  $\mathcal{B}$  who returns  
2419 his guess  $b'$ . The adversary  $\mathcal{A}$  returns the same  $b'$  to the challenger.

2420 By construction, it is clear that  $\mathcal{A}$  wins the hiding game with the same probability  
2421 that  $\mathcal{B}$  wins the simulated hiding game. Since  $\mathcal{B}$  advantage is non-negligible, this means  
2422 that  $\mathcal{A}$  adversary wins the  $\text{ComSch}^*$  hiding game with non-negligible probability as  
2423 well.  $\square$



2424 **Theorem 9** (Binding). *Let  $\text{ComSch}^*$  be a computationally binding commitment scheme*  
 2425 *and  $\text{ComSch}^*.\text{Com}$  indifferntiable from a random oracle. Then  $\text{ComSch}'$  is also compu-*  
 2426 *tationally binding if  $l = \lceil 2^\lambda/p_\lambda \rceil$  is at most  $\text{poly}(\lambda)$ .*

2427 *Proof.* Assume that  $\mathcal{A}$  asks the  $\text{ComSch}'$  commit and open oracles a total of distinct  
 2428  $q_\lambda$  queries. Let us denote the result of the  $q_\lambda$  queries and output of the attacker  
 2429 (candidate collision) as  $((m_1, r_1, y_1), \dots, (m_{q_\lambda}, r_{q_\lambda}, y_{q_\lambda}), \text{out})$ . If  $\mathcal{A}$  is successful it means  
 2430 that it outputs  $(m, r)$ ,  $(m', r')$  such that  $(m, r) \neq (m', r')$  and  $\text{ComSch}'.\text{Com}(m; r) =$   
 2431  $\text{ComSch}'.\text{Com}(m'; r')$ .

By the definition of  $\text{ComSch}'$ , we have that,

$$\text{ComSch}'.\text{Com}(m; r) = \text{decode}_\mathbb{N}(\text{ComSch}^*.\text{Com}(m; r)) \pmod{p_\lambda}$$

Hence, we have a collision in  $\text{ComSch}'$  if there exists  $k \in [l]$ ,  $l$  being the ratio of the  
 codomains of  $\text{ComSch}^*.\text{Com}$  and  $\text{ComSch}'.\text{Com}$ , such that,

$$|\text{decode}_\mathbb{N}(\text{ComSch}^*.\text{Com}(m; r)) - \text{decode}_\mathbb{N}(\text{ComSch}^*.\text{Com}(m'; r'))| = k \cdot p_\lambda.$$

2432 We show that this event is unlikely.

2433 In fact, for each  $i \in [q_\lambda]$ , let  $C_i$  be the event that the adversary wins at the  $i$ -th  
 2434 query. That is, the last commitment  $y_i$  is a  $\text{ComSch}'$  collision with one of the previous  
 2435  $y_j$ . More precisely there exists  $j \leq i$  and  $k < l$  such that  $y_i = y_j + k \cdot p_\lambda$ .

2436 Since  $\text{ComSch}^*$  is a random oracle,  $y_i$  is randomly selected from a set of at least  $p_\lambda$   
 2437 elements. As such, we have  $\Pr[C_i] \leq i \cdot l/p_\lambda$ .

Thus the probability of finding a collision after  $q_\lambda$  queries is  $\Pr[C_1 \vee \dots \vee C_{q_\lambda}] \leq$   
 $\sum_{i=1}^{q_\lambda} \Pr[C_i] = l/p_\lambda \cdot \sum_{i=1}^{q_\lambda} i$ . This probability is bounded by  $l \cdot \frac{q_\lambda(q_\lambda+1)}{p_\lambda}$ . However,  
 we allow only polynomial number of queries. Thus for  $q_\lambda = \text{poly}(\lambda)$  this probability  
 becomes,

$$\frac{2^\lambda \cdot \text{poly}(\lambda)}{p_\lambda^2},$$

2438 what is negligible for  $2^\lambda/p_\lambda \leq \text{poly}(\lambda)$ . □

2439 **Remark 19.** *We observe that in **Zeth**'s commitment scheme, we set  $p_\lambda = \mathbf{r}_{\text{BN}}$  and*  
 2440  *$2^\lambda = 2^{\text{BLAKE2sCLEN}}$ . We thus have  $l = 6$  and the probability of an attacker breaking*  
 2441 *the binding property is because of the reduction modulo  $\mathbf{r}_{\text{BN}}$  increases by a factor of*  
 2442 *approximately 6 and remains negligible.*

2443 **Corollary 1.** *Assume that Blake2 is indifferntiable from a random oracle and a PRF,*  
 2444 *then  $\text{ComSch}^*$  is computationally binding and computationally hiding. Furthermore,*  
 2445 *the reduction is tight. That is, the advantage of any PPT adversary against binding*  
 2446 *(resp. hiding) property is the same as the advantage of an adversary against collision*  
 2447 *resistance and binding (resp. hiding).*

2448

2449

# Glossary

- 2450 **joinsplit** Set of JSIN input *ZethNotes*, and JSOUT output *ZethNotes* as well as the  
2451 public values *vin* and *vout* used in a  $tx_{\text{Mix}}$  transaction. 29–31, 33, 50, 85, 103
- 2452 **joinsplit equation** Equation that checks that the sum of the values of the SendTx  
2453 algorithm of DAP is equal to the sum of the values of its outputs. This equations  
2454 checks that the joinsplit is “balanced” and thus, that no value is created while  
2455 creating new *ZethNotes*. 17, 33, 50, 85, 103

## 2456 **Acronyms**

2457 **DOS** Denial of Service (Attack). 7, 103

2458 **EOA** Externally Owned Account. 8, 9, 11, 103

2459 **EVM** Ethereum Virtual Machine. 7, 8, 12, 74, 76, 103

2460 **MAC** Message Authentication Code. 87, 103

2461 **PoC** Proof of Concept. 103

2462 **RLP** Recursive Length Prefix. 11, 12, 103

# Bibliography

- 2464 [AAM12] Imad Fakhri Alshaikhli, Mohammad A Alahmad, and Khanssaa Munthir.  
2465 Comparison and analysis study of sha-3 finalists. In *2012 International*  
2466 *Conference on Advanced Computer Science Applications and Technologies*  
2467 *(ACSAT)*, pages 366–371. IEEE, 2012.
- 2468 [abi] Contract abi specification, section "function selector". [https://solidity.readthedocs.io/en/develop/abi-spec.html#](https://solidity.readthedocs.io/en/develop/abi-spec.html#function-selector)  
2469 [function-selector](https://solidity.readthedocs.io/en/develop/abi-spec.html#function-selector).  
2470
- 2471 [ABM<sup>+</sup>03] Adrian Antipa, Daniel Brown, Alfred Menezes, René Struik, and Scott  
2472 Vanstone. Validation of elliptic curve public keys. In *International Work-*  
2473 *shop on Public Key Cryptography*, pages 211–223. Springer, 2003.
- 2474 [ABN10] Michel Abdalla, Mihir Bellare, and Gregory Neven. Robust encryption.  
2475 In *Theory of Cryptography Conference*, pages 480–497. Springer, 2010.  
2476 <https://eprint.iacr.org/2008/440.pdf>.
- 2477 [ABR99] Michel Abdalla, Mihir Bellare, and Phillip Rogaway. Dhaes:  
2478 An encryption scheme based on the diffie-hellman prob-  
2479 lem. 1999. [https://pdfs.semanticscholar.org/95f4/](https://pdfs.semanticscholar.org/95f4/63d097086fba325086a4cf88706648dafd09.pdf)  
2480 [63d097086fba325086a4cf88706648dafd09.pdf](https://pdfs.semanticscholar.org/95f4/63d097086fba325086a4cf88706648dafd09.pdf).
- 2481 [ABR01] Michel Abdalla, Mihir Bellare, and Phillip Rogaway. Dhies: An en-  
2482 cryption scheme based on the diffie-hellman problem., 2001. <https://web.cs.ucdavis.edu/~rogaway/papers/dhies.pdf>.  
2483
- 2484 [ACG<sup>+</sup>19] Martin R Albrecht, Carlos Cid, Lorenzo Grassi, Dmitry Khovratovich,  
2485 Reinhard Lüftenegger, Christian Rechberger, and Markus Schofnegger.  
2486 Algebraic cryptanalysis of stark-friendly designs: application to marvel-  
2487 lous and mimc. In *International Conference on the Theory and Appli-*  
2488 *cation of Cryptology and Information Security*, pages 371–397. Springer,  
2489 2019.
- 2490 [AFK<sup>+</sup>08] Jean-Philippe Aumasson, Simon Fischer, Shahram Khazaei, Willi Meier,  
2491 and Christian Rechberger. New features of latin dances: analysis of salsa,  
2492 chacha, and rumba. In *International Workshop on Fast Software Encryp-*  
2493 *tion*, pages 470–488. Springer, 2008.

- 2494 [AG18] Andreas M. Antonopoulos and Wood Gavin. *Mastering Ethereum*.  
2495 O'Reilly Media, 2018.
- 2496 [AGM<sup>+</sup>09] Kazumaro Aoki, Jian Guo, Krystian Matusiewicz, Yu Sasaki, and Lei  
2497 Wang. Preimages for step-reduced sha-2. In *International Conference*  
2498 *on the Theory and Application of Cryptology and Information Security*,  
2499 pages 578–597. Springer, 2009. [https://link.springer.com/content/  
2500 pdf/10.1007/978-3-642-10366-7\\_34.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-10366-7_34.pdf).
- 2501 [AGR<sup>+</sup>16] Martin Albrecht, Lorenzo Grassi, Christian Rechberger, Arnab Roy, and  
2502 Tyge Tiessen. Mimc: Efficient encryption and cryptographic hashing  
2503 with minimal multiplicative complexity. In *International Conference on*  
2504 *the Theory and Application of Cryptology and Information Security*, pages  
2505 191–219. Springer, 2016.
- 2506 [AHMP08] Jean-Philippe Aumasson, Luca Henzen, Willi Meier, and Raphael C-W  
2507 Phan. Sha-3 proposal blake. *Submission to NIST*, 229:230, 2008.
- 2508 [ALM12] Elena Andreeva, Atul Luykx, and Bart Mennink. Provable security of  
2509 blake with non-ideal compression function. In *International Conference*  
2510 *on Selected Areas in Cryptography*, pages 321–338. Springer, 2012.
- 2511 [AMP10] Elena Andreeva, Bart Mennink, and Bart Preneel. Security reductions  
2512 of the second round sha-3 candidates. In *International Conference on*  
2513 *Information Security*, pages 39–53. Springer, 2010.
- 2514 [AMPŠ12] Elena Andreeva, Bart Mennink, Bart Preneel, and Marjan Škrobot. Se-  
2515 curity analysis and comparison of the sha-3 finalists blake, grøstl, jh,  
2516 keccak, and skein. In *International Conference on Cryptology in Africa*,  
2517 pages 287–305. Springer, 2012.
- 2518 [ANWOW13] Jean-Philippe Aumasson, Samuel Neves, Zooko Wilcox-O’Hearn, and  
2519 Christian Winnerlein. Blake2: simpler, smaller, fast as md5. In *Interna-*  
2520 *tional Conference on Applied Cryptography and Network Security*, pages  
2521 119–135. Springer, 2013. <https://eprint.iacr.org/2013/322.pdf>.
- 2522 [AS09] Kazumaro Aoki and Yu Sasaki. Meet-in-the-middle preimage attacks  
2523 against reduced sha-0 and sha-1. In *Annual International Cryptology Con-*  
2524 *ference*, pages 70–89. Springer, 2009.
- 2525 [BBDP01] Mihir Bellare, Alexandra Boldyreva, Anand Desai, and David Pointcheval.  
2526 Key-privacy in public-key encryption. In *International Conference on the*  
2527 *Theory and Application of Cryptology and Information Security*, pages  
2528 566–582. Springer, 2001. [https://iacr.org/archive/asiacrypt2001/  
2529 22480568.pdf](https://iacr.org/archive/asiacrypt2001/22480568.pdf).

- 2530 [BCC<sup>+</sup>15] Jonathan Bootle, Andrea Cerulli, Pyrros Chaidos, Essam Ghadafi, Jens  
2531 Groth, and Christophe Petit. Short accountable ring signatures based on  
2532 ddh. In *European Symposium on Research in Computer Security*, pages  
2533 243–265. Springer, 2015.
- 2534 [BCD<sup>+</sup>20] Tim Beyne, Anne Canteaut, Itai Dinur, Maria Eichlseder, Gregor Le-  
2535 ander, Gaëtan Leurent, María Naya-Plasencia, Léo Perrin, Yu Sasaki,  
2536 Yosuke Todo, and Friedrich Wiemer. Out of oddity – new cryptana-  
2537 lytic techniques against symmetric primitives optimized for integrity proof  
2538 systems. Cryptology ePrint Archive, Report 2020/188, 2020. <https://eprint.iacr.org/2020/188>.  
2539
- 2540 [BCK<sup>+</sup>18] Elaine Barker, Lily Chen, Sharon Keller, Allen Roginsky, Apostol  
2541 Vassilev, and Richard Davis. Recommendation for pair-wise key-  
2542 establishment schemes using discrete logarithm cryptography. Technical  
2543 report, National Institute of Standards and Technology, 2018. [Online;  
2544 last accessed 10-January-2020].
- 2545 [BD07] Eli Biham and Orr Dunkelman. A framework for iterative hash  
2546 functions—haifa. Technical report, Computer Science Department, Techn-  
2547 nion, 2007. <https://eprint.iacr.org/2007/278.pdf>.
- 2548 [BDPVA07] Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche.  
2549 Sponge functions. In *ECRYPT hash workshop*, volume 2007. Citeseer,  
2550 2007.
- 2551 [Ber05] Daniel J Bernstein. The poly1305-aes message-authentication code.  
2552 In *International Workshop on Fast Software Encryption*, pages 32–49.  
2553 Springer, 2005. <https://cr.yp.to/mac/poly1305-20050329.pdf>.
- 2554 [Ber06] Daniel J Bernstein. Curve25519: new diffie-hellman speed records. In  
2555 *International Workshop on Public Key Cryptography*, pages 207–228.  
2556 Springer, 2006. <https://cr.yp.to/ecdh/curve25519-20060209.pdf>.
- 2557 [Ber08a] Daniel J Bernstein. Chacha, a variant of salsa20. In *Workshop Record*  
2558 *of SASC*, volume 8, pages 3–5, 2008. [https://cr.yp.to/chacha/](https://cr.yp.to/chacha/chacha-20080120.pdf)  
2559 [chacha-20080120.pdf](https://cr.yp.to/chacha/chacha-20080120.pdf).
- 2560 [Ber08b] Daniel J. Bernstein. New stream cipher designs. chapter The Salsa20 Fam-  
2561 ily of Stream Ciphers, pages 84–97. Springer-Verlag, Berlin, Heidelberg,  
2562 2008.
- 2563 [BGM17] Sean Bowe, Ariel Gabizon, and Ian Miers. Scalable multi-party compu-  
2564 tation for zk-snark parameters in the random beacon model. Cryptology  
2565 ePrint Archive, Report 2017/1050, 2017. [https://eprint.iacr.org/](https://eprint.iacr.org/2017/1050)  
2566 [2017/1050](https://eprint.iacr.org/2017/1050).

2567 [BL] Daniel J. Bernstein and Tanja Lange. Safecurves: choosing safe curves  
2568 for elliptic-curve cryptography. <https://safecurves.cr.yp.to>. [Online;  
2569 last accessed 09-December-2019].

2570 [Blo70] Burton H. Bloom. Space/time trade-offs in hash coding with allowable  
2571 errors. *Commun. ACM*, 13(7):422–426, 1970.

2572 [Bon19] Xavier Bonnetain. Collisions on feistel-mimc and univariate gmimc. 2019.

2573 [Bou03] Nicolas Bourbaki. *Elements of mathematics: Algebra*. Springer, 2003.

2574 [Bra97] S. Bradner. Key words for use in rfcs to indicate requirement levels. RFC  
2575 2119, RFC Editor, March 1997.

2576 [BRS02] John Black, Phillip Rogaway, and Thomas Shrimpton. Black-box  
2577 analysis of the block-cipher-based hash-function constructions from  
2578 pgv. In *Annual International Cryptology Conference*, pages 320–335.  
2579 Springer, 2002. [https://link.springer.com/content/pdf/10.1007/  
2580 3-540-45708-9\\_21.pdf](https://link.springer.com/content/pdf/10.1007/3-540-45708-9_21.pdf).

2581 [BS07] Mihir Bellare and Sarah Shoup. Two-tier signatures, strongly unforge-  
2582 able signatures, and fiat-shamir without random oracles. In *International  
2583 Workshop on Public Key Cryptography*, pages 201–216. Springer, 2007.

2584 [BSCG<sup>+</sup>14] Eli Ben Sasson, Alessandro Chiesa, Christina Garman, Matthew Green,  
2585 Ian Miers, Eran Tromer, and Madars Virza. Zerocash: Decentralized  
2586 anonymous payments from bitcoin. In *2014 IEEE Symposium on Security  
2587 and Privacy*, pages 459–474. IEEE, 2014.

2588 [Cle19] Clearmatics. Zeth Sproken release. [https://github.com/clearmatics/  
2589 zeth/releases/tag/v0.2](https://github.com/clearmatics/zeth/releases/tag/v0.2), 2019. [Online; released 04-April-2019].

2590 [CM16] Arka Rai Choudhuri and Subhamoy Maitra. Differential cryptanalysis of  
2591 salsa and chacha-an evaluation with a hybrid model. *IACR Cryptology  
2592 ePrint Archive*, 2016:377, 2016. [https://eprint.iacr.org/2016/377.  
2593 pdf](https://eprint.iacr.org/2016/377.pdf).

2594 [CM17] Arka Rai Choudhuri and Subhamoy Maitra. Significantly improved multi-  
2595 bit differentials for reduced round salsa and chacha. *IACR Transactions  
2596 on Symmetric Cryptology*, 2016(2):261–287, Feb. 2017.

2597 [EFK15] Thomas Espitau, Pierre-Alain Fouque, and Pierre Karpman. Higher-order  
2598 differential meet-in-the-middle preimage attacks on sha-1 and blake. In  
2599 *Annual Cryptology Conference*, pages 683–701. Springer, 2015. [https:  
2600 //eprint.iacr.org/2015/515](https://eprint.iacr.org/2015/515).

- 2601 [EGL<sup>+</sup>20] Maria Eichlseder, Lorenzo Grassi, Reinhard Lüftenegger, Morten Øygar-  
 2602 den, Christian Rechberger, Markus Schofnegger, and Qingju Wang. An  
 2603 algebraic attack on ciphers with low-degree round functions: Applica-  
 2604 tion to full mimc. *Cryptology ePrint Archive*, Report 2020/182, 2020.  
 2605 <https://eprint.iacr.org/2020/182>.
- 2606 [est] Estream project. <https://en.wikipedia.org/wiki/ESTREAM>.
- 2607 [GFBR06] Decio Gazzoni Filho, Paulo SLM Barreto, and Vincent Rijmen. The  
 2608 maelstrom-0 hash function. In *Brazilian Symposium on Information and*  
 2609 *Computer System Security*. , 2006.
- 2610 [GGPR13] Rosario Gennaro, Craig Gentry, Bryan Parno, and Mariana Raykova.  
 2611 Quadratic span programs and succinct nizks without pcps. In Thomas  
 2612 Johansson and Phong Q. Nguyen, editors, *Advances in Cryptology - EU-*  
 2613 *ROCRYPT 2013, 32nd Annual International Conference on the Theory*  
 2614 *and Applications of Cryptographic Techniques, Athens, Greece, May 26-*  
 2615 *30, 2013. Proceedings*, volume 7881 of *Lecture Notes in Computer Science*,  
 2616 pages 626–645. Springer, 2013.
- 2617 [GJMG11] B Guido, D Joan, P Michaël, and VA Gilles. The keccak sha-3 submission.  
 2618 2011.
- 2619 [GKN<sup>+</sup>14] Jian Guo, Pierre Karpman, Ivica Nikolić, Lei Wang, and Shuang Wu.  
 2620 Analysis of blake2. In *Cryptographers’ Track at the RSA Conference*,  
 2621 pages 402–423. Springer, 2014. [https://eprint.iacr.org/2013/467](https://eprint.iacr.org/2013/467.pdf).  
 2622 pdf.
- 2623 [GLRW10] Jian Guo, San Ling, Christian Rechberger, and Huaxiong Wang. Ad-  
 2624 vanced meet-in-the-middle preimage attacks: First results on full tiger,  
 2625 and improved results on md4 and sha-2. In *International Conference on*  
 2626 *the Theory and Application of Cryptology and Information Security*, pages  
 2627 56–75. Springer, 2010. <https://eprint.iacr.org/2010/016.pdf>.
- 2628 [Gro16] Jens Groth. On the size of pairing-based non-interactive arguments. In  
 2629 *Annual International Conference on the Theory and Applications of Cryp-*  
 2630 *tographic Techniques*, pages 305–326. Springer, 2016.
- 2631 [GRR<sup>+</sup>16] Lorenzo Grassi, Christian Rechberger, Dragos Rotaru, Peter Scholl, and  
 2632 Nigel P Smart. Mpc-friendly symmetric key primitives. In *Proceedings of*  
 2633 *the 2016 ACM SIGSAC Conference on Computer and Communications*  
 2634 *Security*, pages 430–443. ACM, 2016. [https://eprint.iacr.org/2016/](https://eprint.iacr.org/2016/542)  
 2635 542.
- 2636 [Hao14] Yonglin Hao. The boomerang attacks on blake and blake2. In *Interna-*  
 2637 *tional Conference on Information Security and Cryptology*, pages 286–310.  
 2638 Springer, 2014. <https://eprint.iacr.org/2014/1012.pdf>.



- [Har19] HarryR. Conversation about Miyaguchi-Preneel security. <https://github.com/HarryR/ethsnarks/issues/119>, 2019. Online; accessed June-2019.
- [HKT11] Thomas Holenstein, Robin Künzler, and Stefano Tessaro. The equivalence of the random oracle model and the ideal cipher model, revisited. In *Proceedings of the forty-third annual ACM symposium on Theory of computing*, pages 89–98. ACM, 2011.
- [HMRS12] Ekawat Homsirikamol, Paweł Morawiecki, Marcin Rogawski, and Marian Srebrny. Security margin evaluation of sha-3 contest finalists through sat-based attacks. In *IFIP International Conference on Computer Information Systems and Industrial Management*, pages 56–67. Springer, 2012.
- [Hop16] Daira Hopwood. Daira’s comment on: ”ensure spec retains distinctness assumption in hsig”, 2016.
- [IS09] Takanori Isobe and Kyoji Shibutani. Preimage attacks on reduced tiger and sha-2. In *International Workshop on Fast Software Encryption*, pages 139–155. Springer, 2009. [https://link.springer.com/content/pdf/10.1007/978-3-642-03317-9\\_9.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-03317-9_9.pdf).
- [Ish12] Tsukasa Ishiguro. Modified version of” latin dances revisited: New analytic results of salsa20 and chacha”. 2012. <https://eprint.iacr.org/2012/065.pdf>.
- [JMV01] Don Johnson, Alfred Menezes, and Scott Vanstone. The elliptic curve digital signature algorithm (ecdsa). *International journal of information security*, 1(1):36–63, 2001.
- [Joh16] Nick Johnson. Response to ”how does ethereum make use of bloom filters?”. <https://ethereum.stackexchange.com/questions/3418/how-does-ethereum-make-use-of-bloom-filters>, 2016. [Online; last accessed 10-January-2020].
- [KL14] Jonathan Katz and Yehuda Lindell. *Introduction to modern cryptography*. Chapman and Hall/CRC, 2014. <https://repo.zenk-security.com/Cryptographie%20.%20Algorithmes%20.%20Steganographie/Introduction%20to%20Modern%20Cryptography.pdf>.
- [KMO<sup>+</sup>13] Markulf Kohlweiss, Ueli Maurer, Cristina Onete, Björn Tackmann, and Daniele Venturi. Anonymity-preserving public-key encryption: A constructive approach. In Emiliano De Cristofaro and Matthew K. Wright, editors, *Privacy Enhancing Technologies - 13th International Symposium, PETS 2013, Bloomington, IN, USA, July 10-12, 2013. Proceedings*, volume 7981 of *Lecture Notes in Computer Science*, pages 19–39. Springer, 2013.

- 2677 [KRS12] Dmitry Khovratovich, Christian Rechberger, and Alexandra Savelieva.  
2678 Bicliques for preimages: attacks on skein-512 and the sha-2 family. In  
2679 *International Workshop on Fast Software Encryption*, pages 244–263.  
2680 Springer, 2012. [https://link.springer.com/content/pdf/10.1007/](https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5_15.pdf)  
2681 [978-3-642-34047-5\\_15.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5_15.pdf).
- 2682 [Lab19] Matter Labs. Merkle shrubs. [https://github.com/matter-labs/](https://github.com/matter-labs/MerkleShrubs)  
2683 [MerkleShrubs](https://github.com/matter-labs/MerkleShrubs), 2019.
- 2684 [LHT16] Adam Langley, Mike Hamburg, and Sean Turner. Elliptic curves for se-  
2685 curity. RFC 7748, <https://tools.ietf.org/pdf/rfc7748.pdf>, 2016.
- 2686 [LIS12] Ji Li, Takanori Isobe, and Kyoji Shibutani. Converting meet-in-the-  
2687 middle preimage attack into pseudo collision attack: Application to sha-2.  
2688 In *International Workshop on Fast Software Encryption*, pages 264–286.  
2689 Springer, 2012. [https://link.springer.com/content/pdf/10.1007/](https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5_16.pdf)  
2690 [978-3-642-34047-5\\_16.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5_16.pdf).
- 2691 [LM11] Mario Lamberger and Florian Mendel. Higher-order differential attack  
2692 on reduced sha-256. *IACR Cryptology ePrint Archive*, 2011:37, 2011.  
2693 <https://eprint.iacr.org/2011/037.pdf>.
- 2694 [LMN16] Atul Luykx, Bart Mennink, and Samuel Neves. Security analysis of  
2695 blake2’s modes of operation. *IACR Transactions on Symmetric Cryp-*  
2696 *tology*, pages 158–176, 2016. [https://www.esat.kuleuven.be/cosic/](https://www.esat.kuleuven.be/cosic/publications/article-2705.pdf)  
2697 [publications/article-2705.pdf](https://www.esat.kuleuven.be/cosic/publications/article-2705.pdf).
- 2698 [LN18] Adam Langley and Yoav Nir. Chacha20 and poly1305 for ietf protocols.  
2699 *RFC 8439*, 2018. <https://tools.ietf.org/html/rfc8439>.
- 2700 [LP19] Chaoyun Li and Bart Preneel. Improved interpolation attacks on cryp-  
2701 tographic primitives of low algebraic degree. In *International Conference*  
2702 *on Selected Areas in Cryptography*, pages 171–193. Springer, 2019.
- 2703 [Mai16] Subhamoy Maitra. Chosen iv cryptanalysis on reduced round chacha and  
2704 salsa. *Discrete Applied Mathematics*, 208:88–97, 2016.
- 2705 [MJS15] Ed. M-J. Saarinen. Blake Compression Function F. [https://](https://tools.ietf.org/html/rfc7693#section-3.2)  
2706 [tools.ietf.org/html/rfc7693#section-3.2](https://tools.ietf.org/html/rfc7693#section-3.2), 2015. [Online; accessed  
2707 November-2019].
- 2708 [ML15] Nicky Mouha and Atul Luykx. Multi-key security: The even-mansour  
2709 construction revisited. In *Annual Cryptology Conference*, pages 209–223.  
2710 Springer, 2015. <https://hal.inria.fr/hal-01240988/document>.
- 2711 [MNS11] Florian Mendel, Tomislav Nad, and Martin Schl  ffer. Finding sha-2 char-  
2712 acteristics: searching through a minefield of contradictions. In *Inter-*  
2713 *national Conference on the Theory and Application of Cryptology and*

- 2714 *Information Security*, pages 288–307. Springer, 2011. [https://link.springer.com/content/pdf/10.1007/978-3-642-25385-0\\_16.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-25385-0_16.pdf).  
2715
- 2716 [Moh10] Payman Mohassel. A closer look at anonymity and robustness in encryption schemes. In *Advances in Cryptology - ASIACRYPT 2010 - 16th International Conference on the Theory and Application of Cryptology and Information Security*, volume 6477 of *Lecture Notes in Computer Science*, pages 501–518. Springer, 2010.  
2717  
2718  
2719  
2720
- 2721 [Mon87] Peter L. Montgomery. Speeding the Pollard and elliptic curve methods of factorization. *Math. Comp.*, 48(177):243–264, 1987.  
2722
- 2723 [MP15] Bart Mennink and Bart Preneel. On the impact of known-key attacks on hash functions. In *International Conference on the Theory and Application of Cryptology and Information Security*, pages 59–84. Springer, 2015. <https://eprint.iacr.org/2015/909.pdf>.  
2724  
2725  
2726
- 2727 [MQZ10] Mao Ming, He Qiang, and Shaokun Zeng. Security analysis of blake-32 based on differential properties. In *2010 International Conference on Computational and Information Sciences*, pages 783–786. IEEE, 2010.  
2728  
2729
- 2730 [MVOV96] Alfred J Menezes, Paul C Van Oorschot, and Scott A Vanstone. Handbook of applied cryptography. <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.99.2838&rep=rep1&type=pdf>, 1996.  
2731  
2732
- 2733 [NA19] Samuel Neves and Filipe Araujo. An observation on norx, blake2, and chacha. *Information Processing Letters*, 149:1–5, 2019.  
2734
- 2735 [oST15] National Institute of Standards and Technology. Secure Hash Standard (SHS). <https://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.180-4.pdf>, 2015.  
2736  
2737
- 2738 [Per17] Trevor Perrin. X25519 and zero outputs. <https://moderncrypto.org/mail-archive/curves/2017/000896.html>, 2017. [Online; last accessed 08-January-2020].  
2739  
2740
- 2741 [Por13] Thomas Pornin. Deterministic Usage of the Digital Signature Algorithm (DSA) and Elliptic Curve Digital Signature Algorithm (ECDSA). <https://tools.ietf.org/html/rfc6979>, 2013.  
2742  
2743
- 2744 [PV05] Pascal Paillier and Damien Vergnaud. Discrete-log-based signatures may not be equivalent to discrete log. In *International Conference on the Theory and Application of Cryptology and Information Security*, pages 1–20. Springer, 2005.  
2745  
2746  
2747
- 2748 [Qu99] Minghua Qu. Sec 2: Recommended elliptic curve domain parameters. *Certicom Res., Mississauga, ON, Canada, Tech. Rep. SEC2-Ver-0.6*, 1999.  
2749

2750 [RZ19] Antoine Rondelet and Michal Zajac. ZETH: On Integrating Zerocash on  
2751 Ethereum. [Online; released April-2019], 2019.

2752 [SS08] Somitra Kumar Sanadhya and Palash Sarkar. New collision attacks  
2753 against up to 24-step sha-2. In *International conference on cryptology in*  
2754 *India*, pages 91–103. Springer, 2008. [https://eprint.iacr.org/2008/](https://eprint.iacr.org/2008/270.pdf)  
2755 [270.pdf](https://eprint.iacr.org/2008/270.pdf).

2756 [Ste15] Stevens, Marc. On collisions for md5. [https://www.win.tue.](https://www.win.tue.nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.%20Stevens.pdf)  
2757 [nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.](https://www.win.tue.nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.%20Stevens.pdf)  
2758 [%20Stevens.pdf](https://www.win.tue.nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.%20Stevens.pdf), 2015.

2759 [SZFW12] Zhenqing Shi, Bin Zhang, Dengguo Feng, and Wenling Wu. Improved  
2760 key recovery attacks on reduced-round salsa20 and chacha. In *Interna-*  
2761 *tional Conference on Information Security and Cryptology*, pages 337–351.  
2762 Springer, 2012.

2763 [TBP20] Florian Tramèr, Dan Boneh, and Kenneth G. Paterson. Remote side-  
2764 channel attacks on anonymous transactions. Cryptology ePrint Archive,  
2765 Report 2020/220, 2020. <https://eprint.iacr.org/2020/220>.

2766 [THH15] Piotr Dyrąga Tjaden Hess, Matt Luongo and James Hancock. EIP 152:  
2767 Add BLAKE2 compression function ‘F’. [https://eips.ethereum.org/](https://eips.ethereum.org/EIPS/eip-152)  
2768 [EIPS/eip-152](https://eips.ethereum.org/EIPS/eip-152), 2015. [Online; accessed November-2019].

2769 [VNP10] Janoš Vidali, Peter Nose, and Enes Pašalić. Collisions for variants of the  
2770 blake hash function. *Information processing letters*, 110(14-15):585–590,  
2771 2010.

2772 [W<sup>+</sup>] Gavin Wood et al. Ethereum: A secure decentralised generalised trans-  
2773 action ledger.

2774 [wc] Ethereum wiki contributors. Patricia tree. [https://github.com/](https://github.com/ethereum/wiki/wiki/Patricia-Tree)  
2775 [ethereum/wiki/wiki/Patricia-Tree](https://github.com/ethereum/wiki/wiki/Patricia-Tree).

2776 [wc19] Ethereum wiki contributors. RLP. [https://github.com/ethereum/](https://github.com/ethereum/wiki/wiki/RLP)  
2777 [wiki/wiki/RLP](https://github.com/ethereum/wiki/wiki/RLP), 2019. [Online; accessed December-2019].

2778 [Woo19] Dr Gavin Wood. ETHEREUM: A Secure Decentralised Gener-  
2779 alised Transaction Ledger Byzantium. [https://ethereum.github.io/](https://ethereum.github.io/yellowpaper/paper.pdf)  
2780 [yellowpaper/paper.pdf](https://ethereum.github.io/yellowpaper/paper.pdf), 2019. [VERSION 7e819ec - 2019-10-20].

2781 [zca] On the security of sprout/sapling in-band en-  
2782 cryption. [https://forum.zcashcommunity.com/t/](https://forum.zcashcommunity.com/t/on-the-security-of-sprout-sapling-in-band-encryption/34986)  
2783 [on-the-security-of-sprout-sapling-in-band-encryption/34986](https://forum.zcashcommunity.com/t/on-the-security-of-sprout-sapling-in-band-encryption/34986).

2784 [ZCa19] ZCash. ZCash protocol specification. [https://github.com/zcash/zips/](https://github.com/zcash/zips/blob/master/protocol/protocol.pdf)  
2785 [blob/master/protocol/protocol.pdf](https://github.com/zcash/zips/blob/master/protocol/protocol.pdf), 2019. [Online; initially released  
2786 14-December-2015].