

1

Zeth Protocol Specification

2

Clearmatics Cryptography R&D

3

April 30, 2021

4

Abstract

5 This document specifies the Zeth protocol with various security fixes and performance
6 improvements from the initial design [RZ19].

7 **Keywords**— Ethereum, Zerocash, Zcash, financial-privacy, zero-knowledge proofs,
8 Zeth, privacy-preserving state transitions

9 Contents

10	1 Preliminaries	11
11	1.1 Data structures and representation	11
12	1.1.1 Structured data	11
13	1.1.2 Representations	14
14	1.2 Ethereum	15
15	1.2.1 Ethereum account	16
16	1.2.2 Ethereum transaction	17
17	1.2.3 Ethereum events and Bloom filters	20
18	1.3 zk-SNARKs	21
19	1.3.1 Preliminary definitions	21
20	1.3.2 Computation representation – arithmetization	23
21	1.4 Decentralized Anonymous Payment schemes (DAP)	24
22	1.5 Definitions	26
23	1.5.1 Negligible function	26
24	1.5.2 Basic algebra notions	27
25	1.5.3 Security assumptions	27
26	1.5.4 Symmetric encryption	28
27	1.5.5 Asymmetric encryption	29
28	1.5.6 Block cipher-based compression functions	30
29	1.5.7 Hash functions	31
30	1.5.8 Pseudo Random Functions	33
31	1.5.9 Commitment scheme	33
32	1.5.10 Digital Signature	34
33	1.5.11 Message Authentication Code	35
34	2 Zeth protocol	37
35	2.1 Zeth Data Types	37
36	2.2 Zeth statement	40
37	2.3 Generating the inputs of the Mix function ($\widetilde{\text{Mix}}_{in}$)	41
38	2.4 Creating an Ethereum transaction tx_{Mix} to call Mixer	43
39	2.5 Processing tx_{Mix}	43
40	2.6 Receiving <i>ZethNotes</i>	46
41	2.7 Security requirements for the primitives	47

42	2.7.1	Additional notes	48
43	3	Instantiation of the cryptographic primitives	49
44	3.1	Instantiating the PRFs, ComSch and CRHs	50
45	3.1.1	Blake2 primitive	50
46	3.1.2	Commitment scheme	51
47	3.1.3	PRFs	51
48	3.1.4	Collision resistant hashes	53
49	3.2	Instantiating MKHASH	53
50	3.2.1	MIMC Encryption	53
51	3.2.2	MIMC-based compression function	55
52	3.2.3	An efficient instantiation of MIMC primitives	55
53	3.2.4	Security requirements satisfaction	58
54	3.3	Zeth statement after primitive instantiation	60
55	3.3.1	Instantiating the packing functions	61
56	3.4	Instantiate SigSch _{OT-SIG}	65
57	3.4.1	Security requirements satisfaction	66
58	3.4.2	Data types	66
59	3.5	Instantiate EncSch	67
60	3.5.1	DHAES encryption scheme	67
61	3.5.2	A DHAES instance	68
62	3.5.3	EncSch instantiation	71
63	3.5.4	Security requirements satisfaction	74
64	3.5.5	Final notes and observations	75
65	3.6	ZkSnarkSch instantiation	78
66	4	Implementation considerations and optimizations	81
67	4.1	Client security considerations	81
68	4.1.1	Syncing and waiting	82
69	4.1.2	Note management	83
70	4.1.3	Prepare arguments for Mix transaction	84
71	4.1.4	Wallet backup and recovery	84
72	4.2	Contract security considerations	85
73	4.3	Efficiency and scalability	85
74	4.3.1	Importance of performance	85
75	4.3.2	Cost centers	86
76	4.3.3	Client performance	86
77	4.3.4	Zero-knowledge proof verification (on-chain)	87
78	4.3.5	Merkle tree updates (on-chain)	88
79	4.3.6	Optimizing Blake2's circuit.	89
80	4.4	Encryption of the notes	96

81	A Transaction non malleability	99
82	A.1 Transaction malleability attack on Zeth	100
83	A.1.1 The attack	101
84	A.2 Solutions to address the transaction malleability attack	102
85	A.2.1 ZeroCash solution	102
86	A.2.2 Zcash’s solution	103
87	A.2.3 Solution on Ethereum	103
88	B Double spend attack on equivalent class	105
89	C Side-channel attacks and information leaks	106
90	C.1 Counterfeit data	106
91	C.2 Data leaked during synchronization	107
92	C.3 Queries on successful decryption	108
93	C.4 Invalid ciphertext	108
94	C.5 Using (and retrieving) nullifiers	109
95	C.6 Proof generation	110
96	C.7 Simple mixer calls	110
97	C.7.1 Small anonymity sets	111
98	D Security proofs of Blake2	113
99	D.1 Security model of Blake2	113
100	D.1.1 Weakly Ideal Cipher Model	113
101	D.2 Security proofs	115
102	D.2.1 Blake2 is a PRF	115
103	D.2.2 Proof of Blake2 collision resistance	116
104	D.2.3 Blake2 as a commitment scheme	117
105	D.2.4 Proof of commitment scheme security	118
106	E Fuzzy message detection	120
107	F Extended discussion on the security of MIMC in different settings	122
108		

109 Notation

110 Basic mathematical notation

111 \emptyset The empty set, i.e. $\emptyset = \{\}$

112 $\#S$ The number of elements in the finite set S (also referred to as “cardinality of the
113 set S ”). By convention, $\#\emptyset = 0$

114 $x \in S$ Represents that x is an element of S . If x is a variable such that $x \in S$, we will
115 say that “ x has type S ”, i.e. the unordered collection of objects S represents all
116 the values that x can take

117 $S \setminus T$ Set difference of sets S and T , i.e. $S \setminus T = \{x \in S : x \notin T\}$ (voiced “the set of
118 elements x in S such that x is not in T ”)

119 $S \subseteq T$ S is a subset of T , i.e. $x \in S \Rightarrow x \in T$

120 $S \subset T$ S is a *proper* (or “strict”) subset of T , i.e. $x \in S \Rightarrow x \in T \wedge \exists y \in T, y \notin S$

121 $S = T$ $S \subseteq T \wedge T \subseteq S$

122 $S \cup T$ Union of set S and set T , i.e. $\{x : x \in S \vee x \in T\}$

123 $S \cap T$ Intersection of set S with set T , i.e. $\{x : x \in S \wedge x \in T\}$

124 $f: S \rightarrow T$ Function f that maps elements of the non-empty set S , the “domain”, to the
125 non-empty set T , the “codomain”

126 \mathbb{N} Set of natural numbers. \mathbb{N}^+ represents $\mathbb{N} \setminus \{0\} = \{1, 2, \dots\}$, where $\{n, \dots\}$ rep-
127 represents the application of the successor operator $\text{Succ}(n) = n + 1$, defined by the
128 Peano axioms, infinitely many times

129 \mathbb{Z} Set of integers, i.e. $\{\dots, -2, -1, 0, 1, 2, \dots\}$, where $\{\dots, n\}$ represents the appli-
130 cation of the predecessor operator $\text{Pred}(n) = n - 1$ infinitely many times

131 \mathbb{Q}, \mathbb{R} Set of rational, real numbers

132 $[n]$ Set $\{0, \dots, n - 1\}$, where $n \in \mathbb{N}$

133 $\{a, \dots, b\}$ Set of integers from a through b inclusive, where $a \leq b$

134	$(a_0, a_1, \dots, a_{n-1})$	n -tuple, i.e. ordered collection of items of length n . If $n = 1$, we call
135		it a “singleton”, if $n = 2$, we call the tuple a “pair”. Finally, if $n = 3$, we call it
136		a “triple”. We use the terms “tuples” and “lists” interchangeably.
137	$S \times T$	Cartesian product of sets S and T , i.e. set of all ordered pairs $\{(x, y) : x \in S \wedge y \in T\}$
138	S^n	n -fold Cartesian product of S with itself, i.e. $S^n = \{(x_0, \dots, x_{n-1}) : x_i \in S \forall i \in$
139		$[n]\}$, where $n \in \mathbb{N}$
140	Λ	General notation for an alphabet, i.e. a <i>non-empty finite set</i> such that every string
141		(ordered collection of symbols, or letters, all in Λ) has a unique decomposition.
142		The number of symbols in a string is denoted the “length” of the string
143	ε	The empty string. ε is a string over any alphabet.
144	Λ^n	Set of all strings, defined over alphabet Λ , containing n symbols (i.e. “of length
145		n ”)
146	Λ^*	The Kleene star of Λ represents the set of all strings of finite length, defined over
147		alphabet Λ , including the empty string ε , i.e. $\Lambda^* = \bigcup_{n \in \mathbb{N}} \Lambda^n$
148	$\text{length}(x)$	$\text{length} : \Lambda^* \rightarrow \mathbb{N}$ computes the length of a string x defined over Λ , i.e. $\text{length}(x)$
149		returns the number of symbols composing the string x . By convention, $\text{length}(\varepsilon) =$
150		0
151	$x y$	Infix notation for the concatenation function, $: \Lambda^* \times \Lambda^* \rightarrow \Lambda^*$. If $\text{length}(x) =$
152		$n, \text{length}(y) = m$ and $(n, m) \in \mathbb{N}^2$, then for $z = x y$ holds $\text{length}(z) = n + m$
153	$\text{trunc}_x(k)$	$\text{trunc} : \Lambda^* \rightarrow \Lambda^k$ is the truncation function that returns the sequence formed
154		from the first k elements of x , where $x \in \Lambda^*$. If $k > \text{length}(x)$, then $\text{trunc}_x(k) = x$
155	$x[a:b]$	$[:] : \Lambda^n \times \mathbb{N} \times \mathbb{N} \rightarrow \Lambda^{\leq b-a}$ is the slice function that, if $b \geq a$, returns the string
156		starting at index $\min(n, a)$ of x and finishing at index $\min(n, b)$. The function
157		additionally interprets $x[:b]$ as $x[0:b]$ and $x[a:]$ as $x[a:n]$
158	$\text{pad}_n(x)$	$\text{pad} : \Lambda^{\leq n} \rightarrow \Lambda^n$ is the padding function which pads x by 0’s to reach a size of
159		n . The padding depends on the variable type and endianness.
160	$\text{append}(l, x)$	$\text{append} : D^n \times D \rightarrow D^{n+1}$ is the algorithm that appends x to the list of n
161		element(s) l , if all x and l share the same data type D
162	\mathbb{B}	Alphabet of binary symbols, i.e. $\{0, 1\}$
163	$\langle \mathbf{g}_1, \dots, \mathbf{g}_l \rangle$	Cyclic group generated by $\{\mathbf{g}_1, \dots, \mathbf{g}_l\}$
164	$(q, \mathbb{G}, \mathbf{g}, \otimes)$	Description of the cyclic group $\mathbb{G} = \langle \mathbf{g} \rangle$ of order q , with operation \otimes
165	\mathbb{G}_{CUR}	Safe subgroup of the cyclic group induced by the set of points on the elliptic curve
166		Curve (i.e. elliptic curve subgroup suited for cryptographic use, in which hardness
167		assumptions hold)

168 $\mathbb{Z}/r\mathbb{Z}$ Quotient group defined as the set of equivalence classes modulo r . $\mathbb{Z}/r\mathbb{Z}$, also writ-
169 ten \mathbb{Z}_r , is an additive group. If $r = p$ a prime number, then $\mathbb{Z}_p = \{0, \dots, p - 1\} =$
170 $\mathbb{Z}/p\mathbb{Z}$ is a finite field of elements modulo prime p , also denoted \mathbb{F}_p , where $0_{\mathbb{F}_p}$ and
171 $1_{\mathbb{F}_p}$ respectively represent the additive and multiplicative identity

172 \mathbb{F}_q Finite field of cardinality $q = p^m$, where p is prime, and $m \in \mathbb{N}$

173 $\llbracket x \rrbracket$ Represents the encoding of the scalar x in a group \mathbb{G} described as $(p, \mathbb{G}, \langle \mathbf{g} \rangle, \otimes)$,
174 i.e. $\llbracket x \rrbracket = x \cdot \llbracket 1 \rrbracket = \mathbf{g} \otimes \dots \otimes \mathbf{g}$ (x times). Thus, by convention, $\llbracket 1 \rrbracket = \mathbf{g}$

175 • Represents an inline operator for bilinear pairing. That is for a bilinear pairing
176 from $\mathbb{G}_1 \times \mathbb{G}_2$ to \mathbb{G}_T and elements $\llbracket a \rrbracket_1, \llbracket b \rrbracket_2$ we write $\llbracket ab \rrbracket_t = \llbracket a \rrbracket_1 \bullet \llbracket b \rrbracket_2$

177 $\lceil x \rceil$ Round $x \in \mathbb{R}$ to the next integer

178 $\lfloor x \rfloor$ Round $x \in \mathbb{R}$ to the previous integer

179 $\log_b(x)$ Logarithm with respect to base b , i.e. $x = b^y, \log_b(x) = y$

180 **Algorithmic notation**

181 $x \leftarrow \$ \mathcal{X}$ Element chosen uniformly at random from set \mathcal{X}

182 $x \leftarrow y$ The value y is assigned to the variable x (i.e. “ x receives the value y ”)

183 PPT Probabilistic polynomial time. A polynomial time algorithm A is one for which
184 there exists a polynomial f such that the running time of A on input $x \in \{0, 1\}^*$
185 is $f(|x|)$. A probabilistic algorithm has the ability to “flip” random coins and use
186 the result of these coin tosses in its computation

187 NUPPT Non-uniform probabilistic polynomial time

188 $\mathcal{O}(f)$ Big-O notation

189 il, kl, nl, rl, ol The input il , key kl , nonce nl , randomness rl and output ol length

190 **Cryptography notation**

191 $\mathcal{O}^X(n)$ Public oracle for algorithm X which can be accessed at most n times; \mathcal{O}^X is an
192 unrestricted oracle for algorithm X

193 λ Security parameter ($\lambda \in \mathbb{N}$)

194 negl Negligible function. In this document, negligible will usually mean $\mathcal{O}(2^{-\lambda})$

195 poly Polynomial function

196 \mathcal{A} Adversary algorithm

197	$\text{Adv}_{F,\mathcal{A}}^{\text{prop}}(\lambda)$	Advantage of the adversary \mathcal{A} with regard to the attack game prop on F	
198		(e.g. F can be a function, a family of functions or a group on which a given	
199		property represented by the game prop is supposed to hold)	
200	$\text{prop}^{\mathcal{A}}$	Adversary \mathcal{A} running a security game prop	
201	Zeth notation		
202	π	Output of the proving algorithm of a zk-SNARK scheme. π is also informally	
203		referred to as a “zk-SNARK proof”, “zk-proof”, or simply “proof”	
204	$\mathcal{P}_{\mathcal{Z}}$	Standard notation for a Zeth user	
205	$\widetilde{\text{Mixer}}$	The mixer smart-contract instance	
206	EncSch	In-band encryption scheme used to share Zeth notes	
207	Ethereum notation		
208	Account	Standard notation for an Ethereum account object	
209	$\widetilde{\text{Ctrct}}$	Standard notation for an Ethereum smart-contract instance	
210	$\mathcal{P}_{\mathcal{E}}$	Standard notation for an Ethereum user	
211	ς	Mapping representing the Ethereum state (i.e. “World state”)	
212	$\varsigma[a]$	Account object stored at address a in ς if it exists, \perp is returned otherwise	
213	Constants		
214	ADDRLEN	The bit-length of an Ethereum address	160 <i>bits</i>
215	BLAKE2sCLEN	Output size of Blake2s compression function [ANWOW13]	256 <i>bits</i>
216	FIELD_{BLS}CAP	Field capacity of $\mathbb{F}_{\mathbf{r}_{\text{BLS}}}$.	$\lfloor \log_2 \mathbf{r}_{\text{BLS}} \rfloor = 252$ bits
217	FIELD_{BLS}LEN	Bit-length of a field element $x \in \mathbb{F}_{\mathbf{r}_{\text{BLS}}}$	$\lfloor \log_2 \mathbf{r}_{\text{BLS}} \rfloor = 253$ bits
218	FIELD_{BN}CAP	Field capacity of $\mathbb{F}_{\mathbf{r}_{\text{BN}}}$.	$\lfloor \log_2 \mathbf{r}_{\text{BN}} \rfloor = 253$ bits
219	FIELD_{BN}LEN	Bit-length of a field element $x \in \mathbb{F}_{\mathbf{r}_{\text{BN}}}$	$\lfloor \log_2 \mathbf{r}_{\text{BN}} \rfloor = 254$ bits
220	BYTELEN	Bit-length of a byte	8 <i>bits</i>
221	ENCZETHNOTELEN	Size of an encrypted note (see Section 3.5.3)	$\text{CTBYTELEN} * \text{BYTELEN}$ <i>bits</i>
222	ETHWORDLEN	Width of a storage cell on the Ethereum Virtual Machine stack, i.e. size of	
223		a word on the EVM	256 <i>bits</i>

224	FIELD CAP	Field capacity of $\mathbb{F}_{r_{\text{CUR}}}$, defined as the maximum bit length l such that all	
225		numbers x encoded on l bits are elements of $\mathbb{F}_{r_{\text{CUR}}}$. In other words, $\text{FIELD CAP} =$	
226		$\max_{x \in \mathbb{F}_{r_{\text{CUR}}}} \{\lceil \log_2 x \rceil\}$ s.t. $\sum_{i \in [\text{FIELD CAP}]} 2^i \in \mathbb{F}_{r_{\text{CUR}}}$	
227	FIELD LEN	Bit-length of elements in field element $x \in \mathbb{F}_{r_{\text{CUR}}}$	$\lceil \log_2 r_{\text{CUR}} \rceil$ bits
228	JSIN, JSOUT, JS MAX	The number of inputs and outputs of a joinsplit and $\text{JS MAX} = \max \{\text{JSIN}, \text{JSOUT}\}$	
229	KEK256 DLEN	Message digest size of Keccak256 [GJMG11]	256 bits
230	MKDEPTH	The depth of the Merkle tree used to store commitments	
231	p_{BLS}	Characteristic of the prime (base) finite field over which curve BLS12-377 is de-	
232		defined, $p_{\text{BLS}} = 2586644260129690940106527336948935335363935127549146605398$	
233		$84262666720468348340822774968888139573360124440321458177$ [BCG ⁺ 20]	
234	p_{BN}	Characteristic of the prime (base) finite field over which curve BN-254 is defined,	
235		$p_{\text{BN}} = 218882428718392752224640574525727508869631115729782366268903789$	
236		4645226208583 [Rk19]	
237	p_{SECP}	Characteristic of the prime (base) finite field over which curve secp256k1 is de-	
238		defined, $p_{\text{SECP}} = 115792089237316195423570985008687907853269984665640564039$	
239		457584007908834671663 [wik]	
240	r_{BLS}	Characteristic of the scalar field of BLS12-377, $r_{\text{BLS}} = 84444617494283704242488$	
241		$24938781546531375899335154063827935233455917409239041$ [BCG ⁺ 20]	
242	r_{BN}	Characteristic of the scalar field of BN-254, $r_{\text{BN}} = 2188824287183927522246405$	
243		$745257275088548364400416034343698204186575808495617$ [Rk19]	
244	r_{CUR}	Characteristic of the scalar field of some chosen curve Curve	
245	r_{SECP}	Characteristic of the scalar field of secp256k1, $r_{\text{SECP}} = 1157920892373161954235$	
246		$70985008687907852837564279074904382605163141518161494337$ [wik]	
247	SHA256 BLEN	Block size of SHA256 [oST15, Figure 1]	512 bits
248	SHA256 DLEN	Message digest size of SHA256 [oST15, Figure 1]	256 bits
249	SHA256 MLEN	Message size of SHA256 [oST15, Figure 1]	$< 2^{64}$ bits
250	DGAS	The default/intrinsic cost of an Ethereum transaction	21000 gas
251	ZVALUE LEN	Size in bits of the transferable maximal value	64 bits

252

Change log

253

- **Version:** 0.0, **Date:** 04/12/2019, **Contributor:** Antoine Rondelet, **Description:** Creation of the document. Established initial table of content and started to populate sections with bullet lists to develop in further versions of the document.

254

255

256

- **Version:** 0.1, **Date:** 20/12/2019, **Contributor:** Antoine Rondelet, **Description:** Refactored the structure of the document. Finalized the table of content, wrote sections on notations and preliminaries, and introduced the content related to the malleability fix.

257

258

259

260

- **Version:** 0.2, **Date:** 24/02/2020, **Contributor:** Clearmatics Cryptography R&D, **Description:**

261

262

- **Date:** 26/02/2020, **Contributor:** Duncan Tebbs, **Description:** Wrote section on wallet implementation and side-channel attacks considerations.

263

264

- **Date:** 02/03/2020, **Contributor:** Giuseppe Giffone, **Description:** Changed Merkle tree hash function to MiMC compression function.

265

266

- **Date:** 04/03/2020, **Contributor:** Duncan Tebbs, Michal Zajac, **Description:** Added background on Groth16 SNARK and SNARK scheme instantiation in the protocol.

267

268

269

- **Date:** 04/03/2020, **Contributor:** Raphael Toledo, **Description:** Wrote section on the packing policy and corresponding attack.

270

271

- **Date:** 04/03/2020, **Contributor:** Duncan Tebbs, **Description:** Refactored the data structures preliminary section.

272

273

- **Date:** 24/03/2020, **Contributor:** Raphael Toledo, **Description:** Changed the PRF and commitment instantiation with Blake2s compression function.

274

275

- **Date:** 17/04/2020, **Contributor:** Giuseppe Giffone, **Description:** Added DHAES encryption scheme.

276

277

- **Version:** 0.3, **Date:** 09/06/2020, **Contributor:** Antoine Rondelet, **Description:** Fixed various inconsistencies throughout the document (notational mistakes in document body and in proofs, latex macros, and typos).

278

279

280

- **Version:** 1.0, **Date:** 30/06/2020, **Contributor:** Antoine Rondelet, Duncan Tebbs, Michal Zajac, **Description:** Global review of the document: fixed inconsistencies in definitions and notations, corrected grammatical mistakes and typos,

281

282

283 added examples, figures and merged sections 5 and 6 of Chapter 1 for clarity, added
284 missing references.

285 • **Version:** 1.1, **Date:** 06/11/2020, **Contributor:** Duncan Tebbs, **Description:**
286 Specification in terms of a generic curve, with constants provided for BN-254 and
287 BLS12-377. Some clarification and grammar fixes.

288 • **Version:** 1.2, **Date:** 27/01/2021, **Contributor:** Antoine Rondelet, Michal Zajac,
289 **Description:** Fix erroneous prime fields used in data types, added elliptic curve
290 group notations, fixed erroneous Groth16 formulas and notation inconsistencies
291 between group and field elements.

292 • **Version:** 1.3, **Date:** 08/02/2021, **Contributor:** Duncan Tebbs, **Description:**
293 Explanatory remarks about hashing public zk-proof data to single primary input,
294 and delegating proof verification to external protocols.

295 • **Version:** 1.4, **Date:** 16/04/2021, **Contributor:** Antoine Rondelet, **Descrip-**
296 **tion:** Added appendix about fuzzy message detection.

297 • **Version:** 1.5, **Date:** 19/04/2021, **Contributor:** Antoine Rondelet, Duncan
298 Tebbs, Michal Zajac, **Description:** Added more context with regard to the secu-
299 rity of MiMC in the different settings (round functions of different degrees, lower
300 bounds on the number of rounds etc.). Special thanks to Lorenzo Grassi for very
301 informative discussions.

302 Chapter 1

303 Preliminaries

304 **Zeth** is a protocol which enables private transactions on **Ethereum** [Woo19]. It is a modi-
305 fication of the Decentralized Anonymous Payment (DAP) system **ZeroCash** [BSCG⁺14].
306 The design described in [RZ19] presents the mechanisms by which **ZeroCash** can be used
307 on **Ethereum**, and argues that the information leakages of the solution are well defined
308 and controlled. This document, however, serves as a specification of the protocol and
309 provides security fixes and optimizations from the first proof of concept release of the
310 protocol [Cle19].

311 This document assumes familiarity with blockchain and **Ethereum** in particular. It
312 does not, in any way, aim at replacing the Ethereum yellow paper [Woo19]. The reader is
313 strongly advised to read about **Ethereum** before delving into this specification document.

314 The key words **MUST**, **MUST NOT**, **SHOULD**, **SHOULD NOT**, **MAY**, and **RECOMMENDED** in this
315 document are to be interpreted as described in [Bra97] when they appear in **ALL CAPS**.
316 These words may also appear in this document in lower case as plain English words,
317 absent their normative meanings.

318 1.1 Data structures and representation

319 1.1.1 Structured data

320 When describing the operations to be performed and the data to be manipulated as
321 part of the protocol, we commonly employ tuples of related data where each element
322 of the tuple has some associated semantic meaning and which must often satisfy some
323 conditions. In this section, we develop a framework to reason about such *structured*
324 data, where a single datum may consist of one or more logical parts (called *fields*). The
325 framework is built on top of simple mathematical concepts such as sets, and mappings
326 between them, ensuring that we can always reason about structured data in a rigorous
327 way. We also define notation to aid the specification of structured data, and to refer to
328 specific components of a datum. This will be used extensively in the specification of the
329 protocol.

330 As a simple motivating example, consider a protocol that processes data relating to
 331 individual people. This fictional system may send and receive data such as *name*, *age*
 332 and *address* for a single person, grouping this data into a logical unit. Further, each
 333 piece of data must satisfy specific conditions (*name* must be a series of characters from
 334 some alphabet, *age* must be a positive integer, etc.) We shall make use of this example
 335 several times during the formulation below.

336 In what follows, let $\text{STR} = \{a, b, \dots, y, z\}^*$ (the Kleene star of the *Roman alphabet*).
 337 In our formulation, field names f_i will be elements in this set.

338 **Remark 1.1.1.** Note that a similar formulation could be made using an arbitrary set,
 339 such as the same alphabet augmented with specific symbols, or the alphabet of a different
 340 language. Our choice of STR here is for simplicity.

341 We begin by defining a data type as a set of values called “fields”, each with a “name”
 342 from STR . Abstract sets are used to constrain the values of each field.

Definition 1.1.2 (Structured Data Type). Let f_0, \dots, f_{n-1} be n distinct elements of
 STR and let V_0, \dots, V_{n-1} be sets, for some $n \in \mathbb{N}$. We define *the structured data type \mathbf{T}*
with fields $\{(f_i, V_i)\}_{i \in [n]}$ to be a set of values:

$$\mathbf{T} = V_0 \times \dots \times V_{n-1}$$

with associated post-fix “dot” operators $.f_i : \mathbf{T} \rightarrow V_i$ for $i = 0, \dots, n-1$, acting on values
 $\mathbf{x} \in \mathbf{T}$ to extract the individual elements:

$$\mathbf{x}.f_i = v_i, \text{ where } \mathbf{x} = (v_0, \dots, v_{n-1}) \in \mathbf{T}$$

343 Here, we say that the i -th field has *field name* f_i , with *value set* V_i . Each “dot”
 344 operator $.f_i$ *extracts* the i -th component, or the *value with field name* f_i .

Example 1.1.3. Consider our example protocol that processes information about peo-
 ple. A potentially useful structured data type \mathbf{Person} may be defined with fields:

$$\{(name, \text{STR}), (age, \mathbb{N}), (height, \mathbb{R}^+)\}$$

345 Values \mathbf{p} in \mathbf{Person} are simply tuples in $\text{STR} \times \mathbb{N} \times \mathbb{R}^+$, with semantic meaning (name,
 346 age, height) assigned to each component of \mathbf{p} .

Examples of valid elements in \mathbf{Person} include $\mathbf{a} = (alice, 28, 1.65)$ and $\mathbf{b} = (bob, 31, 1.74)$,
 where the following equalities hold:

$$\begin{aligned} \mathbf{a}.name &= alice, \\ \mathbf{b}.age &= 31, \\ \mathbf{b}.height &= 1.74; \end{aligned}$$

347 For clarity, structured data types may be specified using tables of names, descriptions
 348 and value sets, rather than sets of the form $\{(f_i, V_i)\}_{i \in [n]}$. Similarly, it is frequently
 349 convenient to include the *field names* alongside values when specifying structured data
 350 values.

351 **Example 1.1.4.** Person from Example 1.1.3 might be described in table-form as follows:

Field	Description	Data type
<i>name</i>	Name of the person	STR
<i>age</i>	Age in years	\mathbb{N}
<i>height</i>	Height in meters	\mathbb{R}^+

Example 1.1.5. The values \mathbf{a} and \mathbf{b} in Example 1.1.3 might be written as follows:

$$\mathbf{a} = \{name : alice, age : 28, height : 1.65\}$$

$$\mathbf{b} = \{name : bob, age : 31, height : 1.74\}$$

352 **Remark 1.1.6** (“dot” operators in assignment). The “dot” operators may be used in
 353 algorithm descriptions to indicate *assignment to a specific component*. For example
 354 $\mathbf{a}.age \leftarrow 29$ means that the value of the *age* field of \mathbf{a} is replaced by the value 29.

Formally, for a structured data type \mathbf{T} with fields $\{(f_i, V_i)\}_{i \in [n]}$ where $\mathbf{x} = (v_0, \dots, v_{n-1}) \in \mathbf{T}$ and $v'_i \in V_i$:

$$\mathbf{x}.f_i \leftarrow v'_i$$

is equivalent to:

$$\mathbf{x} \leftarrow (v_0, \dots, v_{i-1}, v'_i, v_{i+1}, \dots, v_{n-1})$$

355 We define one further operator and related assignment notation, convenient in cases
 356 where $V_i = X^m$ for sets X and $m \in \mathbb{N}$.

Definition 1.1.7 (Square bracket operator). For $m \in \mathbb{N}$ and set X , define the operator $[\] : X^m \times [m] \rightarrow X$ as:

$$\mathbf{x}[i] = x_i \text{ where } \mathbf{x} = (x_0, \dots, x_m)$$

For the set X^* , the operator takes the form $[\] : X^* \times \mathbb{N} \rightarrow X$, defined as:

$$\mathbf{x}[i] = \begin{cases} x_i & \text{if } \text{length}(\mathbf{x}) > i \text{ where } \mathbf{x} = (x_0, \dots) \\ \perp & \text{otherwise} \end{cases}$$

Remark 1.1.8 (Square bracket operators in assignment). Similarly to Remark 1.1.6, we develop assignment notation for the square bracket operator $[\]$. Let $\mathbf{x} = (x_0, \dots, x_{m-1})$ be a member of X^m , and x'_i be some element in X . The statement:

$$\mathbf{x}[i] \leftarrow x'_i$$

is equivalent to:

$$\mathbf{x} \leftarrow (x_0, \dots, x_{i-1}, x'_i, x_{i+1}, \dots, x_{m-1})$$

357 Informally, this can be interpreted as replacing the i -th component of \mathbf{x} with x'_i .

358 **Remark 1.1.9** (Deep structures and chained “dot” operators). Consider the case of
 359 structured data \mathbf{T} with fields $\{(f_i, V_i)\}_{i \in [n]}$ for $n \in \mathbb{N}$. Let \mathbf{T}' be another structured data
 360 type with fields $\{(f'_i, V'_i)\}_{i \in [n']}$ for $n' \in \mathbb{N}$, and assume that $V_j = \mathbf{T}'$ for some $j \in [n]$.
 361 Informally, the values of the j -th field of elements of \mathbf{T} are themselves structured data
 362 of type \mathbf{T}' .

363 In this case, “dot” operators may be *chained*, so that $\mathbf{x}.f_j.f'_k$ refers to the k -th field
 364 v'_k of the j -th field v_j of $\mathbf{x} \in \mathbf{T}$.

365 **Example 1.1.10.** Define a structured data type **Address** with fields $(country, \text{STR}), (zipcode, \text{STR})$.
 366 We redefine the structured data type **Person** from Example 1.1.3, with an extra field
 367 *address* of type **Address**. That is, **Person** is the structured data type with fields:

Field	Description	Data type
<i>name</i>	Name of the person	STR
<i>age</i>	Age in years	\mathbb{N}
<i>height</i>	Height in meters	\mathbb{R}^+
<i>address</i>	Address of the person	Address

An example element \mathbf{a} in **Person** is:

$$\mathbf{a} = \left\{ \begin{array}{l} name : \textit{alice}, \\ age : 28, \\ height : 1.65, \\ address : (country : UK, zipcode : SW1A) \end{array} \right\}$$

In this case, the following equalities using the dot and square bracket operators all hold:

$$\begin{aligned} \mathbf{a}.name &= \textit{alice} \\ \mathbf{a}.height &= 1.65 \\ \mathbf{a}.address.country &= UK \\ \mathbf{a}.address.zipcode &= SW1A \\ \mathbf{a}.address.country[1] &= K \end{aligned}$$

368 1.1.2 Representations

369 The binary alphabet $\{0, 1\}$, denoted \mathbb{B} , is used to represent the presence or absence of an
 370 electrical signal in a computer. In fact, every piece of information in a computer is rep-
 371 resented as a string of bits. We assume the existence of an efficient binary representation
 372 for some set of primitive datatypes (such as the natural numbers \mathbb{N} , or alphanumeric

373 characters). Structured data types built up from primitive types (as described above)
 374 can then recursively be assigned similarly efficient representations. This is used to define
 375 the following functions to *encode* data to its bit-string representation, and to *decode* such
 376 bit-strings back to elements of the original type.

Definition 1.1.11. For a set X of values which are to be represented as bit strings, we define functions:

$$\begin{aligned} \text{encode}_X &: X \rightarrow \mathbb{B}^* \\ \text{decode}_X &: \mathbb{B}^* \rightarrow X \cup \perp \end{aligned}$$

satisfying

$$\text{decode}_X(\text{encode}_X(x)) = x \quad \forall x \in X$$

377 to be the functions which encode (resp. decode) elements of X into (resp. from) the
 378 bit-string representations chosen above. We note that decode_X may return \perp in the case
 379 that the input bit-string is malformed.

380 Without ambiguity, we overload the functions `encode` and `decode` to mean encode_X
 381 and decode_X where the set X is clear from context.

In the following sections, we assume that elements of \mathbb{N} are encoded as big-endian binary numbers in the natural way. We denote by \mathbb{N}_b the set of natural numbers that can be uniquely encoded in this way using b bits (possibly with padding). In other words,

$$\mathbb{N}_b = \left\{ x \in \mathbb{N} \text{ s.t. } \text{encode}_{\mathbb{N}}(x) \in \mathbb{B}^b \right\}$$

382 1.2 Ethereum

383 In a nutshell, **Ethereum** is a distributed deterministic state machine, consisting of a glob-
 384 ally accessible singleton state (“the World state”) and a virtual machine that applies
 385 changes to that state [AG18]. State transitions in the state machine are represented
 386 by transactions on the system. As such, each transaction represents a change in the
 387 global state represented as a Merkle Patricia Tree [wc] whose nodes are objects called
 388 “accounts” (Section 1.2.1). The Ethereum Virtual Machine (EVM) allows state transi-
 389 tions to be specified by creating a type of accounts which are associated with a piece
 390 of code (smart-contracts). The code of such accounts, and so, the corresponding state
 391 transitions, can be executed to transition to another state in the automata, by creating
 392 a transaction that calls the given piece of code (Section 1.2.2).

393 To prevent unbounded state transitions in the state machine, each instruction exe-
 394 cuted by the EVM is associated with a cost in **Wei**, referred to as “the gas necessary to
 395 run the operation”. The “gas cost” of a transaction needs to be paid by the transaction
 396 originator (deduced from their account balance), and is awarded to the miner (added
 397 to their account balance) who successfully mines the block containing the transaction.
 398 In addition to the cost of every instruction executed as part of a state transition, every

399 transaction has an intrinsic cost of DGAS gas [Woo19, Appendix G]. Bounding modifica-
 400 tions to the Ethereum state by the amount of Wei held in the transaction originator’s
 401 account allows the system to avoid the Halting problem¹ and protects against a range
 402 of Denial of Service (DOS) attacks.

403 1.2.1 Ethereum account

404 An Ethereum account [Woo19, Section 4.1] is an object containing 4 attributes, as rep-
 405 resented Table 1.1. We distinguish two types of accounts:

- 406 • “Externally Owned Accounts” (EOA), that are created by derivation of an ECDSA se-
 407 cret key; and
- 408 • Smart-contract accounts, that are derived from EVM code specifying a state tran-
 409 sition on the state machine.

410 Each account object is accessible in the Merkle Patricia Tree representing the “World
 411 state” by a unique ADDRLEN-bit long identifier called the address. In the context of EOA,
 412 the address is obtained by generating a new ECDSA [JMV01] key pair (sk, vk) over curve
 413 secp256k1 [Qu99] and taking the rightmost ADDRLEN bits of the Keccak256 hash of the
 414 verification key vk .

Field	Description	Data type
<i>nce</i>	The nonce of an account is a scalar value representing the number of transactions that have originated from the account, starting at 0.	$\mathbb{N}_{\text{ETHWORDLEN}}$
<i>bal</i>	The balance of an account is a scalar value representing the amount of Wei in the account.	$\mathbb{N}_{\text{ETHWORDLEN}}$
<i>sRoot</i>	The storage root is the Keccak256 hash representing the storage of the account.	$\mathbb{B}^{\text{KEK256DLEN}}$
<i>codeh</i>	The code hash is the hash of the EVM code governing the account. If this field is the Keccak256 hash of the empty string, then the account is said to be an “Externally owned Account” (EOA), and is controlled by the corresponding ECDSA private key. If, however, this field is not the Keccak256 hash of the empty string, the account represents a smart contract whose interactions are governed by its EVM code.	$\mathbb{B}^{\text{KEK256DLEN}}$

Table 1.1: Ethereum Account structure

¹https://en.wikipedia.org/wiki/Halting_problem

Note

In the rest of this document, we will refer to an *Ethereum user* $\mathcal{U}_{\mathcal{E}}$ as a person, modeled as an object, holding *one*^a secret key, sk (object attribute), associated with an existing EOA in the “World state”. We denote by $\mathcal{U}_{\mathcal{E}}.Addr$ the **Ethereum** address of $\mathcal{U}_{\mathcal{E}}$ derived from $\mathcal{U}_{\mathcal{E}}.sk$, and which allows $\mathcal{U}_{\mathcal{E}}$ to access the state of their account $\varsigma[\mathcal{U}_{\mathcal{E}}.Addr]$.

We denote by $\widetilde{\mathbf{SmartC}}$ a smart-contract instance/object (i.e. deployed smart-contract with an address, Section 1.2.2), and denote by $\widetilde{\mathbf{SmartC}}.Addr$ its address.

^aThe same physical person may correspond to multiple “Ethereum users” and thus control multiple accounts in the Merkle Patricia Tree.

415

416 1.2.2 Ethereum transaction

417 We now briefly mention what **Ethereum** transactions [Woo19, Section 4.2] are, and how
418 they are created, signed and validated. Once more, the reader is highly encouraged to
419 refer to [Woo19] for a detailed presentation. Informally, a transaction object (tx) is a
420 signed message originating from an **Ethereum** user $\mathcal{U}_{\mathcal{E}}$ (the *transaction originator*, or
421 simply *sender*) that represents a state transition on the distributed state machine (i.e. a
422 change in the “World state” ς).

423 Raw transaction

424 In the following, we define a raw transaction as an unsigned transaction (Table 1.2).

Field	Description	Data type
nce	Transaction nonce	$\mathbb{N}_{\text{ETHWORDLEN}}$
$gasP$	gasPrice	$\mathbb{N}_{\text{ETHWORDLEN}}$
$gasL$	gasLimit	$\mathbb{N}_{\text{ETHWORDLEN}}$
to	Recipient’s address	$\mathbb{B}^{\text{ADDRLEN}}$
val	Value of the transaction in Wei	$\mathbb{N}_{\text{ETHWORDLEN}}$
$init / data$	Contract Creation data $init$ Message call data $data$	\mathbb{B}^*

Table 1.2: Structure of a *raw transaction data type* TxRawDType

425 Finalizing raw transactions

426 A raw transaction needs to be finalized to be accepted. In the context of this document,
427 “finalizing a raw transaction” will be a synonym of “signing a raw transaction”. The
428 transaction structure is represented in Table 1.3.

Field	Description	Data type
tx_{raw}	Raw transaction object	TxRawDType
v	Field v of ECDSA signature used for public key recovery	$\mathbb{B}^{\text{BYTELEN}}$
r	Field r of ECDSA signature [Por13]	$\mathbb{F}_{\text{rSECP}}$
s	Field s of ECDSA signature [Por13]	$\mathbb{F}_{\text{rSECP}}$

Table 1.3: Structure of a (finalized) *transaction data type* TxDType

We define the transaction generation function, cf. Fig. 1.1, as the function taking the sender’s ECDSA signing key and the components of a raw transaction as arguments, and returning a signed (or finalized) transaction (tx_{final} or tx for short).

$$tx_{final} = \text{TxGen}(sk_{\text{ECDSA}}, nce_{in}, gasP_{in}, gasL_{in}, to_{in}, val_{in}, init_{in}, data_{in})$$

$$tx_{final} = \left\{ \begin{array}{l} nce : nce_{in}, \\ gasP : gasP_{in}, \\ gasL : gasL_{in}, \\ to : to_{in}, \\ val : val_{in}, \\ init/data : init_{in}/data_{in}, \end{array} \right\} tx_{raw}$$

$$\left. \begin{array}{l} v : \sigma_{\text{ECDSA}}.v, \\ r : \sigma_{\text{ECDSA}}.r, \\ s : \sigma_{\text{ECDSA}}.s \end{array} \right\} \sigma_{\text{ECDSA}}$$

429 To sign a transaction, the sender first computes the hash of the raw transaction using
430 Keccak256, cf. Eq. (1.1), and then uses their ECDSA signing key, sk_{ECDSA} , to sign the
431 obtained digest. cf. Eq. (1.2). The signature is then appended to the raw transaction to
432 obtain a finalized transaction, cf. Fig. 1.1.

$$digest_{\text{ECDSA}} = \text{Keccak256}(nce_{in}, gasP_{in}, gasL_{in}, to_{in}, val_{in}, init_{in}/data_{in}) \quad (1.1)$$

$$\sigma_{\text{ECDSA}} = \text{SigSch}_{\text{ECDSA}}.\text{Sig}(sk_{\text{ECDSA}}, digest_{\text{ECDSA}}) (= (v, r, s)) \quad (1.2)$$

```

TxGen( $sk_{\text{ECDSA}}, nce_{in}, gasP_{in}, gasL_{in}, to_{in}, val_{in}, init_{in}, data_{in}$ )


---


1 : if  $to_{in} = \emptyset$  do
2 :    $tx_{raw} \leftarrow \{nce : nce_{in}, gasP : gasP_{in}, gasL : gasL_{in}, to : to_{in}, val : val_{in}, init : init_{in}\};$ 
3 : else
4 :    $tx_{raw} \leftarrow \{nce : nce_{in}, gasP : gasP_{in}, gasL : gasL_{in}, to : to_{in}, val : val_{in}, data : data_{in}\};$ 
5 : endif
6 :  $\sigma_{\text{ECDSA}} \leftarrow \text{SigSch}_{\text{ECDSA}}.\text{Sig}(sk_{\text{ECDSA}}, \text{Keccak256}(tx_{raw}));$ 
7 :  $tx_{final} \leftarrow \{tx_{raw}, v : \sigma_{\text{ECDSA}}.v, r : \sigma_{\text{ECDSA}}.r, s : \sigma_{\text{ECDSA}}.s\};$ 
8 : return  $tx_{final}$ ;

```

Figure 1.1: Transaction generation function TxGen

433 **Remark 1.2.1.** As one can see, there is no “from” attribute in a transaction. The
434 sender’s Ethereum address can be recovered from the ECDSA signature. This method
435 is defined in the Ethereum yellow paper as a “sender function” S [Woo19, Appendix F]
436 which maps each transaction to its sender.

437 Types of transactions

438 While only two types of transactions are described in [Woo19, Section 4.2]; namely those
439 which result in message calls and those which result in the creation of new accounts with
440 associated code, we will instead differentiate the types of transactions based on their
441 purpose. The reader is encouraged to read [Woo19] for a formal discussion.

442 Informally, a transaction can be used to achieve three things: transferring Wei from
443 an EOA to another EOA, creating a new account with associated code (i.e. “deploying
444 a smart-contract”), and calling a function of a smart-contract. We will detail here the
445 differences between these usages.

446 **Creating a contract** The $tx.to$ address is set to \emptyset in the transaction. The contract
447 creation data ($tx.init$) includes the new contract’s code. The contract address is
448 computed as the rightmost ADDRLEN bits of the Keccak256 hash of the RLP en-
449 coding [wc19] of the transaction originator’s address and account nonce [Woo19,
450 Section 6].

451 **Calling a contract function** The $tx.to$ address is set to the address of the contract.
452 The message call data byte array ($tx.data$) is set to the contract’s function address
453 (or “Function Selector” [abi]) which are the first 4 bytes of the Keccak256 hash
454 of the function signature, and the function input arguments (ETHWORDLEN bits per
455 input) [Woo19, Section 8].

456 **Transferring Wei from an EOA to another EOA** This corresponds to a “plain trans-
457 action” spending Wei from an address to send them to another. In that case the
458 $tx.to$ address corresponds to the recipient’s address while the transaction data is
459 left empty.

Note

In order to keep notations simple, we assume, in the rest of the document, that smart-contract functions are uniquely determined by their name. As such, we denote by $\text{FS}(\cdot): \mathbb{B}^* \rightarrow \mathbb{B}^{4 \cdot \text{BYTELEN}}$ the function that takes a function name as input and returns its function selector.

460

461 Transaction validity

462 Importantly, not all finalized transactions constitute valid state transitions on the state
463 machine [Woo19, Section 6]. We denote by `EthVerifyTx` the function that takes an
464 `Ethereum` transaction object tx as input and return `true` (resp. `false`) if tx is valid (resp. in-
465 valid). To be deemed valid, a transaction **MUST** satisfy *all* the following conditions:

- 466 1. The transaction is correctly RLP encoded, with no additional trailing bytes;
- 467 2. the transaction signature (v, r, s) is valid;
- 468 3. the transaction nonce $(tx.nce)$ is valid, i.e. it is equal to the account nonce of the
469 transaction originator;
- 470 4. the gas limit is no smaller than the gas used by the transaction;
- 471 5. the transactor has enough funds on his account balance to cover at least the cost
472 $tx.val + tx.gasP \cdot tx.gasL$.

473 Lifecycle of a transaction, and miners' incentives

474 After the creation of an `Ethereum` transaction tx by a user from an `Ethereum` client (ma-
475 chine running a piece of software that enables to be connected to the `Ethereum` network),
476 the transaction is broadcasted to the network and received by a set of peers/nodes.

477 The transaction is then stored in each node's transaction pool, which is a data
478 structure containing all transactions that should be validated (pending transactions) by
479 the node and mined. To maximize miners' returns, the transaction pools are ordered
480 according to the gas price of the transactions. As such, transactions with the highest
481 $tx.gasP$ are subject to be validated and included into a block first. Once tx is selected
482 from the transaction pool, it is validated (fed into `EthVerifyTx`), executed, and included
483 into a block (i.e. "mined"). The block is then broadcasted to all the nodes of the network
484 and is used as the predecessor for the next block to be mined on the network (i.e. "it is
485 added to the chain").

486 1.2.3 Ethereum events and Bloom filters

487 The EVM contains the set of "LOGX" instructions enabling smart-contract functions
488 to "emit events" (i.e. log data) when they are executed²

²see <https://ethgastable.info/>

489 As such, when a block is generated by a miner or verified by the rest of the network,
490 the address of any logging contract, and all the indexed fields from the logs generated by
491 executing those transactions are added to a Bloom filter [Blo70], which is included in the
492 block header [Woo19, Section 4.3]. Importantly, the actual logs *are not included in the*
493 *block data* in order to save space. As such, when an application wants to find (“consume”)
494 all the log entries from a given contract, or with specific indexed fields (or both), the
495 node can quickly scan over the header of each block, checking the Bloom filter to see if
496 it may contain relevant logs. If it does, *the node re-executes the transactions from that*
497 *block, regenerating the logs, and returning the relevant ones to the application* [Joh16].

Note

The ability for a smart-contract function to “emit” some pieces of data when executed, and for an application to “consume” such pieces of data, is used in Zeth in order to construct a *confidential receiver-anonymous channel* [KMO⁺13].

498

499 1.3 zk-SNARKs

500 In this section we introduce notions necessary to understand zero-knowledge proofs,
501 define properties crucial for them, and introduce zk-SNARKs. We refer the reader
502 to Section 3.6 in which we describe the zk-SNARK scheme used in Zeth.

503 1.3.1 Preliminary definitions

504 **NP class of languages.** Since the considered proof systems are designed to work with
505 languages in NP we begin with defining this class. Intuitively, a language \mathbf{L} belongs to
506 NP if for each element $prim$ from the language there is a short witness aux that allows
507 to efficiently³ verify that in fact $prim \in \mathbf{L}$.

508 **Definition 1.3.1** (NP class of languages, cf. [Gol01]). We say that a language \mathbf{L} belongs
509 to a class NP if there exist a polynomial p and a Turing machine M such that for every
510 primary input $prim \in \{0, 1\}^*$, $prim \in \mathbf{L}$ iff there exists an auxiliary input aux such that
511 M accepts the pair $(prim, aux)$ in time at most $p(\text{length}(prim))$.

512 The set of all pairs $(prim, aux)$ acceptable by M constitutes an NP *relation* \mathbf{R} corre-
513 sponding to the language \mathbf{L} .

514 **Non-interactive zero knowledge.** A non-interactive zero-knowledge proof system
515 NIZK for an NP language \mathbf{L} is a tuple of four algorithms $\text{NIZK} = (\text{KGen}, \text{P}, \text{V}, \text{Sim})$. NIZK
516 for a language \mathbf{L} and instance $prim \in \mathbf{L}$ allows a party, called prover and denoted by P ,
517 to convince another party, called verifier and denoted by V , that $prim \in \mathbf{L}$ and nothing
518 else.

519 Without loss of generality, we focus on zk-proof systems that are universal, that
520 is, are able to work with any given NP relation \mathbf{R} . To that end, we define a *relation*

³Informally we say that an algorithm is efficient if it runs in time polynomial in the size of its inputs.

521 generator \mathcal{R} that on input 1^λ (i.e. the security parameter represented in unary) outputs
 522 an NP relation \mathbf{R} . We assume that the security parameter λ can be easily deduced from
 523 \mathbf{R} .

524 We require from a NIZK to have three substantial properties, cf. [Gro06]:

Completeness that assures that an honest prover, who proves that $prim \in \mathbf{L}$ succeeds, i.e. gets his proof accepted by the verifier \mathbf{V} . Formally we require that for any λ , $\mathbf{R} \leftarrow \mathcal{R}(1^\lambda)$, $(prim, aux) \in \mathbf{R}$

$$\Pr \left[\mathbf{V}(\mathbf{R}, crs, prim, P(\mathbf{R}, crs, prim, aux)) \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (crs, td) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda) \end{array} \right] = 1 .$$

Computational soundness which states that in case $prim \notin \mathbf{L}$ the verifier accepts the proof for $prim$ with negligible probability only. Formally we require that for any $\mathbf{R} \leftarrow \mathcal{R}(1^\lambda)$ and PPT adversary \mathcal{A}

$$\Pr \left[\mathbf{V}(\mathbf{R}, crs, prim, \pi) \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (crs, td) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (prim, \pi) \leftarrow \mathcal{A}(\mathbf{R}, crs); \\ prim \notin \mathbf{L} \end{array} \right] \leq \text{negl}(\lambda).$$

Zero knowledge assures that the verifier learns from a proof nothing except the veracity of the proven statement. More precisely we require that there exist a PPT algorithm Sim and negligible function $\eta(\lambda)$ such that for every adversary \mathcal{A} and security parameter λ

$$\left| \Pr \left[\mathcal{A}(\mathbf{R}, crs, \pi) = 1 \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (crs, td) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (prim, aux) \leftarrow \mathcal{A}(\mathbf{R}, crs); \\ (prim, aux) \in \mathbf{R}; \\ \pi \leftarrow \text{Sim}(\mathbf{R}, crs, td, prim) \end{array} \right] - \Pr \left[\mathcal{A}(\mathbf{R}, crs, \pi) = 1 \mid \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (crs, td) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (prim, aux) \leftarrow \mathcal{A}(\mathbf{R}, crs); \\ (prim, aux) \in \mathbf{R}; \\ \pi \leftarrow P(\mathbf{R}, crs, prim, aux) \end{array} \right] \right| \leq \eta(\lambda).$$

525 We say that NIZK is *perfectly* zero-knowledge if $\eta = 0$.

526 We note that the existence of the simulator which by using the trapdoor is able to
 527 make a proof for a false statement (i.e. for $prim \notin \mathbf{L}$) makes the whole zk-proof system

528 vulnerable to adversaries that also know the trapdoor. More precisely, an adversary
 529 who knows a trapdoor td can break the soundness property. This vulnerability comes
 530 with each CRS-based NIZK (for languages in NP). Thus in the real-life deployment of a
 531 CRS-based NIZK it has to be enforced that nobody learns the trapdoor.

532 A zk-SNARK scheme, denoted ZkSnarkSch , is a special type of NIZK which is equipped
 533 with two more properties. First, zk-SNARKs are arguments *of knowledge*, as such they
 534 have to follow a stronger definition of soundness, called *knowledge soundness*.

Knowledge soundness assures that if a prover provided a proof π for a statement
 $prim$ acceptable to a verifier, then she knows the corresponding auxiliary input aux .
 More precisely, we require that for each $\mathbf{R} \leftarrow \mathcal{R}(1^\lambda)$, and malicious PPT prover \mathcal{A}
 there exists a machine $\text{Ext}_{\mathcal{A}}$, called extractor, that given access to randomness r
 used by \mathcal{A} and its inputs, *extracts* the auxiliary input aux from \mathcal{A} ; that is:

$$\Pr \left[\begin{array}{l} \neg(\mathbf{R}(prim, aux)) \wedge \\ \mathbf{V}(\mathbf{R}, crs, prim, \pi) \end{array} \middle| \begin{array}{l} \mathbf{R} \leftarrow \mathcal{R}(1^\lambda); \\ (crs, td) \leftarrow \text{KGen}(\mathbf{R}, 1^\lambda); \\ (prim, \pi) \leftarrow \mathcal{A}(\mathbf{R}, crs; r); \\ aux \leftarrow \text{Ext}_{\mathcal{A}}(\mathbf{R}, crs; r) \end{array} \right] \leq \text{negl}(\lambda).$$

535 Second, zk-SNARKs are *succinct*, and so we require that proofs produced by ZkSnarkSch.P
 536 are short, i.e. sublinear to the size of the primary and auxiliary inputs. Importantly, in
 537 many modern zk-SNARKs, like [Gro16, MBKM19, Gab19, GWC19, CHM⁺20] the proof
 538 size is constant regardless the size of the input.

539 1.3.2 Computation representation – arithmetization

540 In **Zeth** the sender shows that the transaction is correct by arguing (in zero knowledge,
 541 i.e. hiding private inputs) about correctness of evaluation of some predefined predicate.
 542 This predicate ensures that the soundness of the blockchain system is not violated, i.e. the
 543 zk-proof is used to prove that a transaction follows the “rules of the system” without
 544 disclosing its attributes. The proof system that **Zeth** uses operates on an algebraic
 545 representation of the “predicate to prove”. Informally, representing the computation as
 546 a set of algebraic constraints is called *arithmetization*. One of such representations is
 547 Quadratic Arithmetic Programs (QAP) [GGPR13], which, following [Gro16], is used in
 548 **Zeth**. This representation is considered one of the most efficient for general arithmetic
 549 circuits.

550 **Remark 1.3.2.** Preprocessing SNARKs such as [Gro16] rely on common reference
 551 strings with a specific structure. As such, we may use crs and srs (*structured refer-*
 552 *ence string*) interchangeably in the rest of this document.

553 **QAP (R1CS).** Let C be an arithmetic circuit of fan-in 2 over \mathbb{F}_p . The number of
 554 multiplication gates in C is denoted by $constNo$. Likewise, the number of all wires in C
 555 is denoted by $inpNo$.

556 Before we formally introduce the QAP relation \mathbf{R}_{QAP} we provide some intuitions
 557 behind it. First, we observe that the circuit \mathbf{C} can be represented by three matrices
 558 $\vec{A}, \vec{B}, \vec{C}$ all in $\mathbb{F}_p^{\text{constNo} \times \text{inpNo} + 1}$ such that the i -th row in matrix \vec{A} (and \vec{B}) denotes left
 559 (and right) input to the i -th multiplication gate, which is also the k -th input to the
 560 circuit. That is for a circuit evaluation $z \in \mathbb{F}_p^{\text{inpNo} + 1}$ the left input for the i -th gate is
 561 $\sum_{j=0}^{\text{inpNo}} A_{ij} z_j$ and the right input is $\sum_{j=0}^{\text{inpNo}} B_{ij} z_j$. Furthermore, entry \vec{C}_{ik} contains the
 562 output of i -th multiplication gate that is k -th input to the circuit.

563 Second, for the sake of efficiency we represent each matrix as a sequence of poly-
 564 nomials. Each matrix's column is represented by a polynomial in $\mathbb{F}_p[X]$ such that the
 565 column's i -th input equals polynomial's evaluation at ω^i – the i -th primitive root of
 566 unity modulo p . More precisely, we define polynomials:

- 567 • $u_j(X)$, for $j \in \{0, \dots, \text{inpNo}\}$, such that $u_j(\omega^i) = \vec{A}_{ij}$;
- 568 • $v_j(X)$, for $j \in \{0, \dots, \text{inpNo}\}$, such that $v_j(\omega^i) = \vec{B}_{ij}$;
- 569 • $w_j(X)$, for $j \in \{0, \dots, \text{inpNo}\}$, such that $w_j(\omega^i) = \vec{C}_{ij}$.

We consider inputs from 1 to inpNoPrim public (primary input), for some $\text{inpNoPrim} \leq \text{inpNo}$. The rest of the inputs is considered private (auxiliary input). The QAP relation \mathbf{R}_{QAP} is defined as follows:

$$\mathbf{R}_{\text{QAP}} = \left\{ (prim, aux) \left| \begin{array}{l} a_0 = 1; prim = (a_1, \dots, a_{\text{inpNoPrim}}) \in \mathbb{F}_p^{\text{inpNoPrim}}; \\ aux = (a_{\text{inpNoPrim}+1}, \dots, a_{\text{inpNo}}) \in \mathbb{F}_p^{\text{inpNo} - \text{inpNoPrim}}; \\ \sum_{j=0}^{\text{inpNo}} a_j u_j(X) \cdot \sum_{j=0}^{\text{inpNo}} a_j v_j(X) = \sum_{j=0}^{\text{inpNo}} a_j w_j(X) \end{array} \right. \right\}.$$

Note

Importantly, we note that efficient computation on standard hardware may not necessarily lead to an efficient QAP representation. As such, a function can be very efficient to evaluate on a standard computer, but very slow to evaluate in QAP form.

570

571 1.4 Decentralized Anonymous Payment schemes (DAP)

572 **Zeth** [RZ19] is a Decentralized Anonymous Payment scheme (DAP) [BSCG⁺14, Section
 573 3] defined on top of an Ethereum ledger L . A DAP is a tuple of polynomial-time algo-
 574 rithms $\text{DAP} = (\text{Setup}, \text{GenAddr}, \text{SendTx}, \text{VerifyTx}, \text{Receive})$ that manipulate (*create*,
 575 *spend*) data objects called *Notes*. These objects are bound to a given owner and have a
 576 value v attribute (see Section 2.1).

577 **System Setup** The algorithm **Setup** takes the security parameter λ as input and gen-
 578 erates the public parameters pp . The algorithm **Setup** is executed by a trusted

579 party. The resulting public parameters pp are published and made available to all
 580 parties.

581 **Creating Zeth addresses** The algorithm `GenAddr` takes as input the public parame-
 582 ters pp and generates a new DAP address object $Addr = \{pub : Addr_{pk}, priv :$
 583 $Addr_{sk}\}$. More precisely, $Addr_{pk}$ is an object referred to as the “payment ad-
 584 dress” (Table 1.4), and $Addr_{sk}$ is an object referred to as the “private address”
 585 (Table 1.5) [ZCa19].

586 **Transfer notes** The algorithm `SendTx` is used to transfer some public input vin as
 587 well as the value of a set of input (“old”) $Notes$ into a set of output (“new”)
 588 $Notes$ as well as some public output value $vout$. The inputs $Notes$ are marked as
 589 “consumed” (alternatively, we say that the input $Notes$ are “spent”). `SendTx` takes
 590 as inputs the public parameters pp , the input value and the input (“old”) $Notes$
 591 to be transferred, as well as the Merkle root and the Merkle authentications paths
 592 of the commitments to the input $Notes$, the “spending keys” related to the input
 593 $Notes$, the output value to create and the “payment addresses” for the output
 594 (“new”) $Notes$. If the joinsplit equation is satisfied, the algorithm returns the new
 595 $Notes$ and the corresponding `Ethereum` transaction tx , else it returns \perp .

596 **Verifying transactions** The algorithm `VerifyTx` checks the validity of a transaction.
 597 It takes as inputs the public parameters pp , a transaction and the current ledger
 598 L and outputs a bit equal 1 iff the transaction is valid, 0 otherwise.

599 **Receiving notes** The algorithm `Receive` scans the ledger L and retrieves unspent $Notes$
 600 paid to a particular user address. It takes as input the recipient address key pair
 601 $\{pub : Addr_{pk}, priv : Addr_{sk}\}$ and the current ledger L and outputs the set of
 602 (unspent) received $Notes$.

Note

In the rest of this document, we will refer to a *Zeth user* U_Z as a person, modeled as an object, holding one `Zeth` address (object attribute), and thus holding a *private address*, $Addr_{sk}$. We denote by $U_Z.Addr$ the `Zeth` address of U_Z derived from $Addr_{sk}$, and which allows U_Z to be the recipient of payments via `Zeth`, and to send funds via `Zeth`. Importantly, *not all Ethereum users are Zeth users, and vice-versa!*

603

Field	Description
apk	The <i>paying key</i>
$pkenc$	The <i>transmission key</i>

Table 1.4: “Payment address”, $Addr_{pk}$, of a DAP address

Field	Description
<i>ask</i>	The <i>spending key</i>
<i>skenc</i>	The <i>receiving key</i>

Table 1.5: “Private address”, $Addr_{sk}$, of a DAP address

604 **Zeth** leverages zk-SNARKs (Section 1.3) and the possibility to deploy smart-contracts
605 to specify privacy-preserving state transitions altering the **Ethereum** state ς (Section 1.2).
606 As such, **Zeth** defines a smart-contract, $\widetilde{\mathbf{Mixer}}$, that keeps track of the set of *ZethNotes*
607 (Section 2.1) in a committed form, stored in a Merkle tree; and which verifies the va-
608 lidity of the state transitions generated by the **Zeth** users. As such a **Zeth** DAP is
609 entirely determined by $\widetilde{\mathbf{Mixer}}$, the instance of the mixer smart-contract deployed on the
610 **Ethereum** ledger. State transitions are executed on-chain by calling the `Mix` function of
611 $\widetilde{\mathbf{Mixer}}$, which implements the algorithm `VerifyTx` of DAP, and which modifies ς iff the
612 transaction is deemed valid.

Note

We denote by Mix_{in} the inputs taken by the `Mix` function defined on $\widetilde{\mathbf{Mixer}}$. Let $zdata$ be the value of the *data* field of an **Ethereum** transaction such that:

$$zdata = FS(\text{Mix}) \parallel Mix_{in}$$

Then, we define tx_{Mix} as being the **Ethereum** transaction object returned by `SendTx` such that:

$$tx_{\text{Mix}}.to = \widetilde{\mathbf{Mixer}}.Addr \wedge tx_{\text{Mix}}.data = zdata$$

Importantly, when it is clear from context, we will omit the function selector from the definition of $zdata$, and only assume that $zdata = Mix_{in}$.

613

614 1.5 Definitions

615 1.5.1 Negligible function

616 **Definition 1.5.1** (Negligible function, [KL14, Definition 3.4]). A function f from \mathbb{N} to
617 \mathbb{R}^+ (positive real numbers) is negligible if for every positive polynomial p there exists N
618 such that for all integers $n > N$ it holds that $f(n) < \frac{1}{p(n)}$.

619 **1.5.2 Basic algebra notions**

620 **Definition 1.5.2** (Group, see [Bou03, Section I.4]). A group is given by a tuple (\mathbb{G}, \otimes) ,
 621 where \mathbb{G} is a set and \otimes is a binary operation in \mathbb{G} , i.e. $\otimes : \mathbb{G} \times \mathbb{G} \rightarrow \mathbb{G}$, with the following
 622 properties:

- 623 • $(\mathbf{g} \otimes \mathbf{h}) \otimes \mathbf{k} = \mathbf{g} \otimes (\mathbf{h} \otimes \mathbf{k})$ (associativity)
- 624 • There exists an element $\epsilon \in \mathbb{G}$ s.t. for each $\mathbf{g} \in \mathbb{G}$, $\mathbf{g} \otimes \epsilon = \epsilon \otimes \mathbf{g} = \mathbf{g}$ (identity
 625 element).
- 626 • For each $\mathbf{g} \in \mathbb{G}$ there exist $\mathbf{h} \in \mathbb{G}$ s.t. $\mathbf{g} \otimes \mathbf{h} = \mathbf{h} \otimes \mathbf{g} = \epsilon$ (inverse element).

For simplicity, we may also use the additive notation for groups: \otimes is denoted as $+$, the identity element as \mathbf{o} and the inverse element of \mathbf{g} as $-\mathbf{g}$. Given $\mathbf{g} \in \mathbb{G}$ and $x \in \mathbb{Z}$, we have that:

$$x \cdot \mathbf{g} = \begin{cases} \mathbf{o} & \text{if } x = 0 \\ \mathbf{g} + \dots + \mathbf{g}, (x \text{ times}) & \text{if } x > 0. \\ -\mathbf{g} + \dots + (-\mathbf{g}), (x \text{ times}) & \text{if } x < 0 \end{cases}$$

627 **Definition 1.5.3** (Finite Cyclic Group, adapted from [KL14, Sections 7.1.3, 7.3.2]). A
 628 finite cyclic group is given by a tuple $(q, \mathbb{G}, \mathbf{g}, \otimes)$, called the *group description*, where \mathbb{G}
 629 represents the set of group elements, \mathbf{g} is a generator and q is the order. The generator
 630 \mathbf{g} generates the group; namely, each $\mathbf{h} \in \mathbb{G}$ can be expressed by the generator as $\mathbf{h} =$
 631 $\mathbf{g} \otimes \dots \otimes \mathbf{g}$. Given a scalar x , we denote by $\llbracket x \rrbracket$ the *encoding* of x in \mathbb{G} : i.e. $\llbracket x \rrbracket = \mathbf{g} \otimes \dots \otimes \mathbf{g}$
 632 (x times). As consequence, $\llbracket 1 \rrbracket = \mathbf{g}$.

633 For theoretical purposes, we introduce the **SetupG** algorithm that for a given security
 634 parameter λ outputs a cyclic group, formally:

635 **Definition 1.5.4** (Group Setup Algorithm, taken from [KL14, Sections 7.1.3, 7.3.2]).
 636 A group setup algorithm **SetupG** is a PPT algorithm which takes as input a security pa-
 637 rameter 1^λ and outputs a group description $(q, \mathbb{G}, \mathbf{g}, \otimes)$, where the binary representation
 638 of q is given by λ bits and each group element can be represented by $gLen(\lambda)$ bits. Note
 639 that $gLen$ is $\text{poly}(\lambda)$.⁴

640 **1.5.3 Security assumptions**

Definition 1.5.5 (Discrete Log Problem(DLog), cf. [BS07]). Let \mathbb{G} denote a group
 (Section 1.5.2) whose order p is prime and written over λ bits. We let $\log_{\mathbf{g}}(h)$ denote
 the discrete logarithm of h in the basis \mathbf{g} . We assume \mathbb{G}, p are fixed and known to
 all parties. We denote the advantage of a PPT adversary \mathcal{A} in attacking the discrete
 logarithm problem as

$$\text{Adv}_{\mathbb{G}, \mathcal{A}}^{\text{dlog}} = \Pr[\mathbf{g} \leftarrow \mathbb{G}^*, x \leftarrow \mathbb{F}_p, x' \leftarrow \mathcal{A}(\llbracket 1 \rrbracket, \llbracket x \rrbracket) : \llbracket x' \rrbracket = \llbracket x \rrbracket]$$

⁴For simplicity, we may denote $gLen(\lambda)$ as $gLen$.

641 We say that the DLog is hard in \mathbb{G} if and only if $\text{Adv}_{\mathbb{G}, \mathcal{A}}^{\text{dlog}}(\lambda)$ is negligible for any PPT
 642 adversary \mathcal{A} .

Definition 1.5.6 (One More Discrete Log Problem (om-DLog), cf. [PV05]). Let \mathbb{G} denote a group whose order p is prime and written over λ bits. We let $\log_{\mathfrak{g}}(h)$ denote the discrete logarithm of h in the basis \mathfrak{g} . A PPT adversary \mathcal{A} solving the om-DLog is given $q + 1$ random group elements as well as limited access to a discrete logarithm oracle $\text{O}^{\text{DLog}_{\mathfrak{g}}}(q)$. \mathcal{A} is allowed to query this oracle at most q times, thus obtaining the discrete logarithm of q group elements of his choice with respect to a fixed base \mathfrak{g} . Eventually, \mathcal{A} must output the $q + 1$ discrete logarithms. We denote the advantage of a PPT adversary \mathcal{A} in attacking the one more discrete logarithm problem as

$$\text{Adv}_{\mathbb{G}, \mathcal{A}}^{\text{om-dlog}}(\lambda) = \Pr \left[\begin{array}{l} \mathfrak{g} \leftarrow \mathbb{G}^*, \{[r_i]\}_{i \in [q+1]} \leftarrow \mathbb{G}^{q+1}, \\ \{r'_i\}_{i \in [q+1]} \leftarrow \mathcal{A}^{\text{O}^{\text{DLog}_{\mathfrak{g}}}(q)}(\llbracket 1 \rrbracket, \{[r_i]\}_{i \in [q+1]}) : \\ \forall i \in [q+1], r'_i = \log_{\mathfrak{g}}([r_i]) \end{array} \right]$$

643 We say that the om-DLog is hard in \mathbb{G} if and only if $\text{Adv}_{\mathbb{G}, \mathcal{A}}^{\text{om-dlog}}(\lambda)$ is negligible for any
 644 PPT adversary \mathcal{A} .

645 1.5.4 Symmetric encryption

646 **Definition 1.5.7** (Symmetric Encryption, [KL14, Definition 3.8]). A symmetric encryption
 647 scheme Sym is given by a tuple of PPT algorithms $(\text{KGen}, \text{Enc}, \text{Dec})$ where:

- 648 • KGen , the key generation algorithm, takes a security parameter 1^λ and outputs a
 649 secret key ek ; we assume, without loss of generality, that $kLen(\lambda) = \text{length}(ek) \geq \lambda$.
 650 Note that $kLen(\lambda)$ is a polynomial function in λ .⁵
- 651 • Enc , the encryption algorithm, takes a key ek , a plaintext $m \in \{0, 1\}^*$ and returns
 652 a ciphertext ct .
- 653 • Dec , the decryption algorithm, takes a key ek and a ciphertext ct , and returns a
 654 message m . We assume, without loss of generality, that Dec is deterministic.

655 For every security parameter λ , key ek output by $\text{KGen}(1^\lambda)$, and message $m \in \{0, 1\}^*$,
 656 it holds that $\text{Dec}(ek, \text{Enc}(ek, m)) = m$ (*correctness property*).

657 Let $(\text{KGen}, \text{Enc}, \text{Dec})$ be a symmetric encryption scheme. If there exists a polynomial
 658 l such that, for all $\lambda > 0$ and key ek output by $\text{KGen}(1^\lambda)$, $\text{Enc}(ek, \cdot)$ is only defined for
 659 messages $m \in \{0, 1\}^{l(\lambda)}$, then we say that $(\text{KGen}, \text{Enc}, \text{Dec})$ is a *fixed-length symmetric*
 660 *encryption scheme* with *length parameter* $l(\lambda)$. A security notion for Sym follows:

Definition 1.5.8 (IND-CPA). Let Sym be a symmetric encryption scheme and let \mathcal{A} be an adversary. Consider the IND-CPA game described in Figure 1.2. We define the IND-CPA advantage of \mathcal{A} as follows:

$$\text{Adv}_{\text{Sym}, \mathcal{A}}^{\text{ind-cpa}}(\lambda) = |2 \cdot \Pr[\text{IND-CPA}(\lambda) = 1] - 1|.$$

⁵For simplicity, we may denote $kLen(\lambda)$ as $kLen$.

```

IND-CPA( $\lambda$ )
-----
 $ek \leftarrow \text{KGen}(1^\lambda)$ 
 $(m_0, m_1, \text{state}) \leftarrow \mathcal{A}^{\text{O}^{\text{Enc}_{ek}}}$  with  $\text{length}(m_0) = \text{length}(m_1)$ 
 $b \leftarrow_{\$} \{0, 1\}$ 
 $ct \leftarrow \text{Enc}(ek, m_b)$ 
 $\tilde{b} \leftarrow \mathcal{A}^{\text{O}^{\text{Enc}_{ek}}}(ct, \text{state})$ 
return  $\tilde{b} = b$ 

```

Figure 1.2: IND-CPA game for Sym.

661 Sym is said to be IND-CPA secure if, for every PPT adversary \mathcal{A} , the advantage $\text{Adv}_{\text{Sym}, \mathcal{A}}^{\text{ind-cpa}}(\lambda)$
662 is a negligible function.

663 1.5.5 Asymmetric encryption

664 **Definition 1.5.9** (Asymmetric encryption, [KL14, Definition 10.1]). An *asymmetric*
665 *encryption scheme* Asym is given by a tuple of PPT algorithms $(\text{KGen}, \text{Enc}, \text{Dec})$ where:

- 666 • KGen , the key generation algorithm, takes a security parameter 1^λ and returns a
667 pair of keys (sk, pk) . We refer to the first of these as the *private key* and the second
668 as the *public key*. We assume for convenience that pk and sk each have length at
669 least λ , and that λ can be determined from pk, sk ;
- 670 • Enc , the encryption algorithm, takes a public key pk , a plaintext m , from some
671 underlying plaintext space (that may depend on pk) and returns a ciphertext ct ;
- 672 • Dec , the decryption algorithm, takes a private key sk and a ciphertext ct , and
673 returns a message m or a special symbol \perp denoting decryption failure. We assume,
674 without loss of generality, that Dec is deterministic.

675 We require that for all (sk, pk) returned by KGen , and every message m in the appropriate
676 underlying plaintext space, it holds that $\text{Dec}(sk, \text{Enc}(pk, m)) = m$ (*correctness property*).

677 Secure communication usually requires ciphertext indistinguishability (e.g. IND-CCA2
678 [ABR99, Definition 8]). In **Zeth**, however, the key privacy property IK-CCA [BBDP01]
679 is also required – it ensures indistinguishability of the key under which an encryption is
680 performed.

Definition 1.5.10 (IK-CCA). Let $\text{Asym} = (\text{KGen}, \text{Enc}, \text{Dec})$ be an asymmetric encryp-
tion scheme and let \mathcal{A} be an adversary. Given the IK-CCA game described in Figure 1.3,
with the condition that \mathcal{A} cannot query $\text{O}^{\text{Dec}_{sk_0}}$ or $\text{O}^{\text{Dec}_{sk_1}}$ on the challenge ciphertext

```

IK-CCA( $\lambda$ )
-----
( $sk_0, pk_0$ ), ( $sk_1, pk_1$ )  $\leftarrow$  KGen( $1^\lambda$ )
( $m, state$ )  $\leftarrow$   $\mathcal{A}^{\mathcal{O}^{Dec_{sk_0}}, \mathcal{O}^{Dec_{sk_1}}}$ ( $pk_0, pk_1$ )
 $b \leftarrow_{\$} \{0, 1\}$ 
 $ct \leftarrow$  Enc( $pk_b, m$ )
 $\tilde{b} \leftarrow$   $\mathcal{A}^{\mathcal{O}^{Dec_{sk_0}}, \mathcal{O}^{Dec_{sk_1}}}$ ( $ct, state$ )
return  $\tilde{b} = b$ 

```

Figure 1.3: IK-CCA game.

ct^6 , we define the IK-CCA advantage of \mathcal{A} as follows:

$$\text{Adv}_{\text{Asym}, \mathcal{A}}^{\text{ik-cca}}(\lambda) = |2 \cdot \Pr[\text{IK-CCA}(\lambda) = 1] - 1|$$

681 We say that Asym is IK-CCA secure if for every PPT adversary \mathcal{A} the advantage $\text{Adv}_{\text{Asym}, \mathcal{A}}^{\text{ik-cca}}(\lambda)$
682 is a negligible function.

683 1.5.6 Block cipher-based compression functions

684 **Definition 1.5.11.** Let $kl, il > 1$. A *block cipher* is a map $E: \{0, 1\}^{kl} \times \{0, 1\}^{il} \rightarrow \{0, 1\}^{il}$
685 where, for each key $k \in \{0, 1\}^{kl}$, the function $E_k(\cdot) = E(k, \cdot)$ is a permutation on $\{0, 1\}^{il}$.
686 If E is a block cipher then E^{-1} is its inverse, that on input (k, y) returns m such that
687 $E_k(m) = y$.

688 Let $\mathcal{BCK}(kl, il)$ be the set of all block ciphers $E: \{0, 1\}^{kl} \times \{0, 1\}^{il} \rightarrow \{0, 1\}^{il}$. In order
689 to analyse the security properties of block cipher-based cryptographic constructions it
690 is common to use a security model denoted *the ideal cipher model (ICM)*. Informally
691 speaking, in ICM attackers are allowed to query an oracle simulating a random block
692 cipher, but have no information about the oracle's internal structure. We formalize this
693 notion in the following definition:

694 **Definition 1.5.12** (Ideal Cipher Model [HKT11]). The Ideal Cipher Model (ICM),
695 is a security model where all parties are granted access to an ideal cipher $E: \{0, 1\}^{kl} \times$
696 $\{0, 1\}^{il} \rightarrow \{0, 1\}^{il}$, a random primitive such that $E(k, \cdot)$ for $k \in \{0, 1\}^{kl}$ are 2^{kl} independent
697 random permutations.

698 For fixed kl and il , each party is given access to the oracles \mathcal{O}^E and $\mathcal{O}^{E^{-1}}$, simulating
699 E and E^{-1} , which can be queried for encryption and decryption a polynomial number
700 of times. The encryption oracle takes as input a key, $k \in \{0, 1\}^{kl}$, and a preimage,
701 $m \in \{0, 1\}^{il}$, and returns a tuple comprising the image, $y \in \{0, 1\}^{il}$, along with the
702 inputs, k and m . If (k, m) is queried for the first time, the image y is taken uniformly

⁶*state* is some state information that the adversary outputs after the choice of the message to encrypt. It can be some preprocessed information that can be helpful to win the game

$O^E(k, m)$	$O^{E^{-1}}(k, y)$
if $(k, m, \cdot) \notin \text{Table}_O$ $y \leftarrow \{0, 1\}^{\text{il}}$ $\text{Table}_O.\text{append}(k, m, y)$	if $(k, \cdot, y) \notin \text{Table}_O$ $m \leftarrow \{0, 1\}^{\text{il}}$ $\text{Table}_O.\text{append}(k, m, y)$
else $y \leftarrow \text{Table}_O(k, m)$	else $m \leftarrow \text{Table}_O(k, y)$
return (k, m, y)	return (k, m, y)

Figure 1.4: Oracles of an ideal block cipher, with Table_O being a table of tuples (key, preimage, image) of queries already answered by the oracle.

703 at random from $\{0, 1\}^{\text{il}}$ and added to the oracle's table. Otherwise, the oracle returns
704 y associated with query (k, m) in its table. The decryption oracle is defined similarly
705 with the image and key defined as inputs and the preimage chosen randomly, for details
706 see Fig. 1.4.

Definition 1.5.13 (Block cipher-based compression function [BRS02]). A *block cipher-based compression function* is a map f such that

$$f: \mathcal{BK}(kl, \text{il}) \times \{0, 1\}^a \times \{0, 1\}^b \rightarrow \{0, 1\}^c$$

707 where $kl, \text{il}, a, b, c > 1$ and $a + b > c$. The function f , given $m \in \{0, 1\}^a \times \{0, 1\}^b$,
708 computes $f(E, m)$ using an E -oracle.

709 **Remark 1.5.14.** We use f_E to denote a block cipher-based compression function f
710 restricted to a given block cipher E , i.e. $f_E: \{0, 1\}^a \times \{0, 1\}^b \rightarrow \{0, 1\}^c$ and $f_E = f(E, \cdot)$,
711 for a, b, c as given in the definition above.

Let f be a compression function based on a block cipher. Fix a constant $h_0 \in \{0, 1\}^c$ and an adversary \mathcal{A} . We define the advantage in finding a collision in f as

$$\text{Adv}_{f, \mathcal{A}}^{\text{coll}} = \Pr \left[E \leftarrow \mathcal{BK}(kl, \text{il}); ((k, m), (k', m')) \leftarrow \mathcal{A}^{O^E, O^{E^{-1}}}(f_E, h_0) : \begin{array}{l} ((k, m) \neq (k', m') \wedge f_E(k, m) = f_E(k', m')) \vee f_E(k, m) = h_0 \end{array} \right].$$

712 The previous definition gives credit for finding an (k, m) such that $f_E(k, m) = h_0$ for
713 a fixed $h_0 \in \{0, 1\}^c$.

714 1.5.7 Hash functions

715 **Definition 1.5.15** (Hash function, [KL14, Definition 4.9]). A hash function \mathcal{H} is a pair
716 of algorithms (Setup, H) fulfilling the following properties:

- 717 • **Setup** is a PPT algorithm which takes as input a security parameter 1^λ and outputs
718 a key hk . We assume that 1^λ is included in hk .

719 • H is (deterministic) polynomial-time algorithm that takes as input a key hk and
 720 any string $x \in \{0, 1\}^*$, and outputs a string $H(hk, x) = H_{hk}(x) \in \{0, 1\}^{hLen}$, where
 721 $hLen$ is a polynomial in λ .⁷

722 If for every λ and hk , H_{hk} is defined only over inputs of length $hInpLen(\lambda)$ and $hLen(\lambda) >$
 723 $hLen(\lambda)$, then we say that \mathcal{H} is a *fixed-length hash function* with length parameter
 724 $hInpLen$. Note that $hInpLen(\lambda)$ is a polynomial in λ .

725 Informally, for a given function f we say that (x, y) is a *collision* if $f(x) = f(y)$ and
 726 $x \neq y$. In the following, we formalize this notion for a hash function \mathcal{H} .

Definition 1.5.16 (Collision Resistance [KL14, Definitions 4.10]). A hash function $\mathcal{H} = (\text{Setup}, H)$ is collision resistant if for all PPT adversaries \mathcal{A} there exists a negligible function $\text{negl}(\lambda)$ such that:

$$\text{Adv}_{\mathcal{H}, \mathcal{A}}^{\text{cr}}(\lambda) = \Pr \left[hk \leftarrow \text{Setup}(1^\lambda), (x, y) \leftarrow \mathcal{A}(hk) : x \neq y \wedge H_{hk}(x) = H_{hk}(y) \right] \leq \text{negl}(\lambda).$$

727 HDHI and HDHI2 assumptions

728 The Hash Diffie-Hellman Independence (HDHI) assumption states that, given H in \mathcal{H}
 729 and a group description $(p, \mathbb{G}, \mathbf{g}, \otimes)$, for $\llbracket u \rrbracket$ and $\llbracket v \rrbracket$, with u, v sampled at random, it
 730 is hard for an attacker to distinguish $H(\llbracket u \rrbracket \parallel \llbracket uv \rrbracket)$ from a random string of the same
 731 size.⁸ This is formalized in Definition 1.5.17, where an attacker can also access an oracle
 732 $\mathcal{O}^{\text{HDHI}_v}$ that on input $\mathfrak{r} \in \mathbb{G}$ returns $H(\mathfrak{r} \parallel v \cdot \mathfrak{r})$ (queries on $\llbracket u \rrbracket$ are forbidden).⁹ In other
 733 words, the HDHI assumption measures the sense in which H is “independent” of the
 734 underlying Diffie-Hellman problem.

Definition 1.5.17 (HDHI, [ABR99, Definition 7]). Let \mathcal{H} be a hash function, SetupG be a group generation algorithm and \mathcal{A} be an adversary. Consider the HDHI game described in Figure 1.5. We define the advantage of \mathcal{A} in violating the HDHI assumption as:

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi}}(\lambda) = |2 \cdot \Pr[\text{HDHI}(\lambda) = 1] - 1|.$$

735 Note that the above definition corresponds to [ABR99, Section 3.2.1, Definition 3].
 736 In the following, we introduce a similar notion denoted as HDHI2 (this is an adapta-
 737 tion of the ODH2 notion in [ABN10, Section 6]) which will be useful in the IK-CCA
 738 proof Section 3.5.4.

Definition 1.5.18 (HDHI2). Let \mathcal{H} be a hash function, SetupG a group generation algorithm and let \mathcal{A} be an adversary. Consider the HDHI2 game described in Figure 1.6. We define the advantage of \mathcal{A} in violating the HDHI2 assumption as:

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi2}}(\lambda) = |2 \cdot \Pr[\text{HDHI2}(\lambda) = 1] - 1|.$$

⁷For simplicity, we may denote $hLen(\lambda)$ as $hLen$.

⁸Note that H takes as inputs bit strings, so technically we should make use of an encoding function from \mathbb{G} to $\{0, 1\}^{gLen}$ but we may omit this step through the document to improve readability.

⁹In [ABR99, Section 3.2.1] this notion is denoted as adaptive HDH independence assumption. Since we only introduce the adaptive version we denote it as HDHI.

HDHI(λ)
 $hk \leftarrow \mathcal{H}.\text{Setup}(1^\lambda)$
 $(q, \mathbb{G}, \mathfrak{g}, \otimes) \leftarrow \text{SetupG}(1^\lambda)$
 $u, v \leftarrow_{\$} [q]$
 $w_0 \leftarrow \mathcal{H}.\text{H}_{hk}(\llbracket u \rrbracket \parallel \llbracket uv \rrbracket)$
 $w_1 \leftarrow_{\$} \{0, 1\}^{hLen}$
 $b \leftarrow_{\$} \{0, 1\}$
 $\tilde{b} \leftarrow \mathcal{A}^{\text{HDHI}_v}(\llbracket u \rrbracket, \llbracket v \rrbracket, w_b)$
return $\tilde{b} = b$

Figure 1.5: HDHI game.

HDHI2(λ)
 $hk \leftarrow \mathcal{H}.\text{Setup}(1^\lambda)$
 $(q, \mathbb{G}, \mathfrak{g}, \otimes) \leftarrow \text{SetupG}(1^\lambda)$
 $u, v_0, v_1 \leftarrow_{\$} [q]$
 $w_{0,0} \leftarrow \mathcal{H}.\text{H}_{hk}(\llbracket u \rrbracket \parallel \llbracket uv_0 \rrbracket), w_{0,1} \leftarrow \mathcal{H}.\text{H}_{hk}(\llbracket u \rrbracket \parallel \llbracket uv_1 \rrbracket)$
 $w_{1,0} \leftarrow_{\$} \{0, 1\}^{hLen}, w_{1,1} \leftarrow_{\$} \{0, 1\}^{hLen}$
 $b \leftarrow_{\$} \{0, 1\}$
 $\tilde{b} \leftarrow \mathcal{A}^{\text{HDHI}_{v_0}, \text{HDHI}_{v_1}}(\llbracket u \rrbracket, \llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket, w_{b,0}, w_{b,1})$
return $\tilde{b} = b$

Figure 1.6: HDHI2 game.

Lemma 1.5.1. *Let \mathcal{A} be an adversary with advantage $\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi2}}$ in solving the HDHI2 problem. Then there exists an adversary \mathcal{B} such that*

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{A}}^{\text{hdhi2}}(\lambda) \leq 2 \cdot \text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{B}}^{\text{hdhi}}(\lambda).$$

739 *Proof.* We reuse the proof described in [ABN10, Lemma 6.1] by applying minor mod-
740 ifications. In fact, HDHI and HDHI2 are, respectively, slightly different from ODH and
741 ODH2 notions: in the related security games, if $b = 0$ the challenges are constructed as
742 $\text{H}(\llbracket u \rrbracket \parallel \llbracket uv \rrbracket)$ and $\{\text{H}(\llbracket u \rrbracket \parallel \llbracket uv_0 \rrbracket), \text{H}(\llbracket u \rrbracket \parallel \llbracket uv_1 \rrbracket)\}$ instead of $\text{H}(\llbracket uv \rrbracket)$ and $\{\text{H}(\llbracket uv_0 \rrbracket), \text{H}(\llbracket uv_1 \rrbracket)\}$.
743 By accordingly changing the instances of H in the games $\text{G}_0, \text{G}_1, \text{G}_2$ of [ABN10, Lemma
744 6.1] our lemma follows. \square

745 1.5.8 Pseudo Random Functions

746 Informally, a pseudorandom function family $\mathcal{PRF} = \{\text{PRF}_k : D \rightarrow C\}_{k \in \mathcal{K}}$ is a collection
747 of functions such that for a randomly chosen $k \in \mathcal{K}$, the function PRF_k is indistinguish-
748 able from a random function that maps D to C .

Definition 1.5.19 (PRF Family [KL14, Definition 3.23]). Let $\mathcal{F} : \{0, 1\}^* \times \{0, 1\}^* \rightarrow \{0, 1\}^*$ be an efficient, length-preserving, keyed function. We say \mathcal{F} is a pseudo random function if for all probabilistic polynomial-time distinguishers Dist , there exists a negligible function negl such that:

$$\text{Adv}_{\mathcal{F}, \text{Dist}}^{\text{prf}}(\lambda) = \left| \Pr \left[\text{Dist}^{\mathcal{F}_k(\cdot)}(1^\lambda) = 1 \right] - \Pr \left[\text{Dist}^{f_\lambda(\cdot)}(1^\lambda) = 1 \right] \right| \leq \text{negl}(\lambda),$$

749 where $k \leftarrow_{\$} \mathcal{K} = \{0, 1\}^\lambda$ is chosen uniformly at random and f_λ is chosen uniformly at
750 random from the set of functions mapping λ -bit strings to λ -bit strings.

751 1.5.9 Commitment scheme

752 **Definition 1.5.20** (Non-interactive commitment scheme [BCC⁺15, Section 2.1]). A
753 non-interactive commitment scheme ComSch is defined by the following algorithms:

- 754 • **Setup**, is a PPT algorithm that takes a security parameter 1^λ and outputs public
755 parameters pp .
- 756 • **Com**, is a polynomial-time algorithm that takes a message $m \in \mathbb{B}^{\text{nl}}$, a random coin
757 $r \in \mathbb{B}^{\text{nl}}$ and outputs a commitment $cm \in \mathbb{B}^{\text{ol}}$.

758 We assume that pp is implicitly passed to **Com**.

Definition 1.5.21 (Computational hiding). We say that a commitment scheme is computationally hiding if for all PPT adversary \mathcal{A} , the advantage:

$$\left| \Pr \left[\begin{array}{l} pp \leftarrow \text{Setup}(1^\lambda), (m_0, m_1) \leftarrow \mathcal{A}(pp), b \leftarrow_{\$} \{0, 1\}, \\ r \leftarrow_{\$} \mathbb{B}^{\text{nl}}, cm \leftarrow \text{Com}(m_b; r), \tilde{b} \leftarrow \mathcal{A}(cm), b = \tilde{b} \end{array} \right] - \frac{1}{2} \right|$$

759 is at most negligible in λ .

Definition 1.5.22 (Computational binding). We say that a commitment scheme is computationally binding if for all PPT adversary \mathcal{A} , the advantage:

$$\Pr \left[\begin{array}{l} pp \leftarrow \text{Setup}(1^\lambda), (m_0, r_0, m_1, r_1) \leftarrow \mathcal{A}(pp) \\ m_0 \neq m_1 \wedge \text{Com}(m_0; r_0) = \text{Com}(m_1; r_1) \end{array} \right]$$

760 is at most negligible in λ .

761 Note that the previous definitions can be made *statistical* if we consider unbounded
762 attackers \mathcal{A} .

763 1.5.10 Digital Signature

764 **Definition 1.5.23** (Digital signature [KL14, Definition 12.1]). A digital signature scheme
765 **SigSch** is defined by the tuple of functions $\text{SigSch} = (\text{KGen}, \text{Sig}, \text{Vf})$,

- 766 • $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$. Key Generation randomized algorithm takes as input the
767 security parameter 1^λ and returns a signing key sk and verifying key vk .
- 768 • $\sigma \leftarrow \text{Sig}(sk, m)$. Given a signing key sk and a message m , the **Sig** algorithm
769 computes and outputs a signature σ .
- 770 • $\{0, 1\} \leftarrow \text{Vf}(vk, m, \sigma)$. Given a verification key vk , a message m and a signature
771 σ , the **Vf** algorithm returns 1 if σ is a valid signature else 0.

772 A signature scheme must satisfy the *correctness property* (i.e $\text{Vf}(vk, m, \text{Sig}(sk, m)) =$
773 true , where $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$) and be *unforgeable* (i.e. it is intractable to produce a
774 signature, without knowing the signing key sk , on a message that has not been signed
775 yet). In addition to these properties, certain digital signature schemes have an additional
776 property called *one-timeness*, also defined below.

UF-CMA($1^\lambda, t, q$)

1 : $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$
2 : $state \leftarrow \mathcal{A}^{\text{OSig}_{sk}}(vk, \cdot)$
3 : // $state = \{(m_i, \sigma_i)\}_{i \in [q]}$ where m_i denotes
4 : // the i th query made to OSig_{sk} and
5 : // σ_i denotes the i th oracle answers
6 : $(m^*, \sigma^*) \leftarrow \mathcal{A}(state)$
7 : **return** $\text{Vf}(vk, m^*, \sigma^*) = 1$
8 : $\wedge m^* \notin \{m_i\}_{i \in [q]}$

Figure 1.7: UF-CMA game

SUF-CMA($1^\lambda, t, q$)

1 : $(sk, vk) \leftarrow \text{KGen}(1^\lambda)$
2 : $state \leftarrow \mathcal{A}^{\text{OSig}_{sk}}(vk, \cdot)$
3 : // $state = \{(m_i, \sigma_i)\}_{i \in [q]}$ where m_i denotes
4 : // the i th query made to OSig_{sk} and
5 : // σ_i denotes the i th oracle answers
6 : $(m^*, \sigma^*) \leftarrow \mathcal{A}(state)$
7 : **return** $\text{Vf}(vk, m^*, \sigma^*) = 1$
8 : $\wedge (m^*, \sigma^*) \notin \{(m_i, \sigma_i)\}_{i \in [q]}$

Figure 1.8: SUF-CMA game

777 **Definition 1.5.24** (Unforgeability (UF-CMA) [KL14, Definition 12.2]). A digital signa-
778 ture scheme SigSch is UF-CMA if for any PPT adversary \mathcal{A} , the probability that \mathcal{A} wins
779 the UF-CMA game, depicted in Fig. 1.7, is negligible.

780 **Definition 1.5.25** (Strong Unforgeability (SUF-CMA)). A digital signature scheme
781 SigSch is SUF-CMA if the probability that any PPT adversary \mathcal{A} wins the SUF-CMA game,
782 depicted in Fig. 1.8, is negligible.

783 **Definition 1.5.26** (One-Time (OT) Signature [KL14, Definition 12.6]). A *one-time*
784 signature scheme is a digital signature scheme that uses each key-pair at most once.

785 **Remark 1.5.27.** It is worth noting that users may use one-time signing keys to sign
786 multiple messages. In this case no security claims can be made.

787 1.5.11 Message Authentication Code

788 A message authentication code is a scheme that enables users to tag data for the purpose
789 of authenticity and integrity. Formally:

790 **Definition 1.5.28** (Message Authentication Code, [KL14, Definition 4.1]). A message
791 authentication code MAC is given by a tuple of PPT algorithms $(\text{KGen}, \text{Tag}, \text{Vf})$ where:

- 792 • KGen , the key generation algorithm, takes a security parameter 1^λ , and returns a
793 key $mk \in \{0, 1\}^{mLen(\lambda)}$.¹⁰
- 794 • Tag , the tag generation algorithm, takes a key mk and a message $y \in \{0, 1\}^*$ and
795 returns a string $\tau \in \{0, 1\}^*$, called *tag*.
- 796 • Vf , the tag verification algorithm, takes a key mk , a message $y \in \{0, 1\}^*$ and a tag
797 $\tau \in \{0, 1\}^*$. It returns a value in $\{0, 1\}$ where: 0 denotes that the message was
798 rejected (i.e. deemed unauthentic) and 1 denotes that the message was accepted
799 (i.e. deemed authentic).

¹⁰For simplicity, we may denote $mLen(\lambda)$ as $mLen$.

```

SUF-CMA ( $\lambda$ )
-----
 $mk \leftarrow \text{KGen}(1^\lambda)$ 
 $(\bar{y}, \bar{\tau}) \leftarrow \mathcal{A}^{\text{O}^{\text{Tag}_{mk}}, \text{O}^{\text{Vf}_{mk}}}$ 
return  $\text{Vf}(mk, \bar{y}, \bar{\tau}) = 1$ 

```

Figure 1.9: SUF-CMA game.

800 We require that for all $mk \in \{0, 1\}^\lambda$ and $y \in \{0, 1\}^*$ we have $\text{Vf}(mk, y, \text{Tag}(mk, y)) = 1$.
801 If $\text{Tag}(mk, \cdot)$ is defined only over messages of length $l(\lambda)$ and $\text{Vf}(mk, y, \tau)$ outputs 0 for
802 every y that is not of length $l(\lambda)$, then we say that $(\text{KGen}, \text{Tag}, \text{Vf})$ is a *fixed-length MAC*
803 with length parameter $l(\lambda)$.

804 A security notion for MAC follows:

805 **Definition 1.5.29** (SUF-CMA, [ABR99, Section 3.2.3]). Let $\text{MAC} = (\text{KGen}, \text{Tag}, \text{Vf})$ be
806 a message authentication scheme and let \mathcal{A} be an adversary. Consider the SUF-CMA
807 game described in Figure 1.9, with the condition that $\text{Tag}(mk, \bar{y}) \neq \bar{\tau}$. We say that an
808 adversary \mathcal{A} has *forged* a tag when it outputs a pair $(\bar{y}, \bar{\tau})$ such that $\text{Vf}_k(\bar{y}, \bar{\tau}) = 1$, where
809 $(\bar{y}, \bar{\tau})$ was not previously obtained via a query to the tag oracle.

We define the SUF-CMA advantage of \mathcal{A} as follows:

$$\text{Adv}_{\text{MAC}, \mathcal{A}}^{\text{suf-cma}}(\lambda) = \Pr[\text{SUF-CMA}(\lambda) = 1]$$

810 We say that MAC is SUF-CMA secure if for every PPT adversary \mathcal{A} the advantage
811 $\text{Adv}_{\text{MAC}, \mathcal{A}}^{\text{suf-cma}}(\lambda)$ is a negligible function.

812 Chapter 2

813 Zeth protocol

814 In this section, we detail the **Zeth** protocol and provide a set of requirements that need
815 to be respected to guarantee the security of the protocol.

816 2.1 Zeth Data Types

817 We begin by describing, and giving intuition about, the data types (see Section 1.1) used
818 in **Zeth**. We follow some design rationale from **ZeroCash** [BSCG⁺14], and **Zcash** [ZCa19]
819 in order to prevent the transaction malleability attack, and the Faerie Gold attack [ZCa19,
820 Section 8.4]. We refer the reader to Appendix A for more details.

821 In what follows **Curve** represents a curve with scalar field $\mathbb{F}_{r_{\text{CUR}}}$, satisfying the re-
822 quirements of Section 3.6. The specification is described in terms of this generic curve,
823 with examples and notes relating to specific instances of interest (namely BN-254 and
824 BLS12-377, see Chapter 3).

825 **ZethNoteDType** Represents a note in **Zeth**. This data type consists of the note owner's
826 public address apk , identifier ρ , randomness r and value v .

Field	Description	Data type
apk	Note owner's paying key	$\mathbb{B}^{\text{PRFADDRROUTLEN}}$
r	Note randomness	$\mathbb{B}^{\text{RTRAPLEN}}$
v	Note value	$\mathbb{B}^{\text{ZVALUELEN}}$
ρ	Note identifier	$\mathbb{B}^{\text{PRFRHOOUTLEN}}$

Table 2.1: **ZethNoteDType** data type

827 **JSInputDType** Denotes a joinsplit input. It comprises the opening of a commitment
828 cm which is in the set of leaves in the Merkle tree of **Mixer** (i.e. a *ZethNote*), its
829 address $mkaddr$ and authentication path $mkpath$ on the contract's Merkle tree as
830 well as the spending key ask of the note holder and the note nullifier nf .

Field	Description	Data type
<i>mkpath</i>	Merkle authentication path to the commitment corresponding to the <i>ZethNote</i> to spend	$(\mathbb{F}_{\text{rCUR}})^{\text{MKDEPTH}}$
<i>mkaddr</i>	Commitment address in the Merkle tree	$\mathbb{B}^{\text{MKDEPTH}}$
<i>znote</i>	Zeth note object	ZethNoteDType
<i>cm</i>	Zeth note commitment	\mathbb{F}_{rCUR}
<i>ask</i>	Note owner's spending key	$\mathbb{B}^{\text{ASKLEN}}$
<i>nf</i>	Note nullifier	$\mathbb{B}^{\text{PRNFOUTLEN}}$

Table 2.2: JSInputDType data type

831 **PrimInputDType** Represents the primary inputs used to generate the zk-SNARK proof
832 π . *prim* is a tuple defined as the current Merkle root *mkroot* of the Merkle tree
833 maintained by **Mixer**, the input notes nullifiers $nfs = (nf_0, \dots, nf_{\text{JSIN}-1})$, the
834 output notes commitments $cms = (cm_0, \dots, cm_{\text{JSOUT}-1})$, the signature hash *hsig*,
835 the message authentication tags $htags = (h_0, \dots, h_{\text{JSIN}-1})$ and the residual bits
836 field *rsd*, which aggregates the former's fields bits which could not be contained in
837 a field element.

Field	Description	Data type
<i>mkroot</i>	Merkle root of the Merkle tree	\mathbb{F}_{rCUR}
<i>nfs</i>	Indexed set of nullifiers of the “old” notes to spend (see Section 3.3.1 for definition of NFFLEN)	$((\mathbb{F}_{\text{rCUR}})^{\text{NFFLEN}})^{\text{JSIN}}$
<i>cms</i>	Indexed set of commitments to the newly created notes	$(\mathbb{F}_{\text{rCUR}})^{\text{JSOUT}}$
<i>hsig</i>	Signature hash (non-malleability, see Appendix A and Section 3.3.1 for definition of HSIGFLEN)	$(\mathbb{F}_{\text{rCUR}})^{\text{HSIGFLEN}}$
<i>htags</i>	Indexed set of message authentication tags (non-malleability, see Appendix A and Section 3.3.1 for definition of HFLEN)	$((\mathbb{F}_{\text{rCUR}})^{\text{HFLEN}})^{\text{JSIN}}$
<i>rsd</i>	Residual bits corresponding to unpacked bits of former fields (see Section 3.3.1 for definition of RSDFLEN)	$(\mathbb{F}_{\text{rCUR}})^{\text{RSDFLEN}}$

Table 2.3: PrimInputDType data type

838 **AuxInputDType** Represents the auxiliary inputs used to generate the zk-SNARK proof
839 π . *aux* is a tuple defined as joinsplit inputs (i.e. “old outputs to be spent”), the new
840 *ZethNotes*, the joinsplit’s randomness ϕ as well the public values *vin* and *vout*, the
841 signature hash *hsig* and the message authentication tags *htags* = $(h_0, \dots, h_{\text{JSIN}-1})$.

Field	Description	Data type
<i>jsins</i>	Indexed set of JSIN joinsplit inputs	$\text{JSInputDType}^{\text{JSIN}}$
<i>znotes</i>	Indexed set of JSOUT newly created notes	$\text{ZethNoteDType}^{\text{JSOUT}}$
ϕ	The joinsplit randomness (non-malleability, see Appendix A)	$\mathbb{B}^{\text{PHILEN}}$
<i>vin</i>	Public input value to the joinsplit	$\mathbb{B}^{\text{ZVALUELEN}}$
<i>vout</i>	Public output value to the joinsplit	$\mathbb{B}^{\text{ZVALUELEN}}$
<i>hsig</i>	Signature hash (non-malleability, see Appendix A)	$\mathbb{B}^{\text{CRHHSIGOUTLEN}}$
<i>htags</i>	Indexed set of message authentication tags (non-malleability, see Appendix A)	$(\mathbb{B}^{\text{PRFPKOUTLEN}})^{\text{JSIN}}$

Table 2.4: **AuxInputDType** data type

842 **MixInputDType** Represents the set of inputs to the Mix function of $\widetilde{\text{Mixer}}$. The input of
843 the Mix function is a tuple defined as the primary inputs *prim*, the zk-proof π , the
844 ciphertexts of the newly created notes *ciphers* = $(ct_0, \dots, ct_{\text{JSOUT}-1})$, a one-time
845 signature σ and the associated verification key *vk*.

Field	Description	Data type
<i>primIn</i>	Primary input object associated with the zk-proof π	PrimInputDType
<i>proof</i>	The zk-SNARK associated to the <i>Zeth</i> statement (see Section 2.2)	ZKPDType (see Section 3.6)
<i>otssig</i>	The one-time signature used to prevent transaction malleability (see Appendix A)	SigOtsDType (see Section 3.4.2)
<i>otsvk</i>	The verification key associated with the signature <i>otssig</i> used to prevent transaction malleability (see Appendix A)	VKOtsDType (see Section 3.4.2)
<i>ciphers</i>	Indexed set of ciphertexts of the newly generated notes	$(\mathbb{B}^{\text{ENCZETHNOTELEN}})^{\text{JSOUT}}$ (see Section 3.5)

Table 2.5: **MixInputDType** data type

846 **MixEventDType** Represents the data emitted as an **Ethereum** event (Section 1.2.3) dur-
 847 ing a successful execution of the **Mix** function of **Mixer**. Clients are required to
 848 read this data and use it to update their representation of **Mixer**'s state.

Field	Description	Data type
<i>mkroot</i>	New root of Merkle tree of commitments	\mathbb{F}_{rCUR}
<i>nfs</i>	Nullifiers for input notes consumed	$(\mathbb{B}^{\text{PRFNFOUTLEN}})^{\text{JSIN}}$
<i>cms</i>	Commitments to the output notes	$(\mathbb{F}_{\text{rCUR}})^{\text{JSOUT}}$
<i>ciphers</i>	Ciphertexts for the output notes	$(\mathbb{B}^{\text{ENCZETHNOTELEN}})^{\text{JSOUT}}$

Table 2.6: MixEventDType data type

849 2.2 Zeth statement

850 As explained in [RZ19], the **Mix** function of **Mixer** verifies the validity of π on the
 851 given primary inputs in order to determine whether the state transition is valid. As
 852 such, **Mixer** verifies whether for π , and primary input *prim*, there exists an auxiliary
 853 input *aux*, such that the tuple $(\text{prim}, \text{aux})$ satisfies the NP-relation \mathbf{R}^z , consisting of the
 854 following constraints:

- 855 • For each $i \in [\text{JSIN}]$:
 - 856 1. $\text{aux.jsins}[i].\text{znote.apk} = \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{addr}}(0)$
 - 857 2. $\text{aux.jsins}[i].\text{cm} = \text{ComSch.Com}(\text{aux.jsins}[i].\text{znote.apk}, \text{aux.jsins}[i].\text{znote}.\rho, \text{aux.jsins}[i].\text{znote}.v;$
 858 $\text{aux.jsins}[i].\text{znote}.r)$
 - 859 3. $\text{aux.jsins}[i].\text{nf} = \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{nf}}(\text{aux.jsins}[i].\text{znote}.\rho)$
 - 860 4. $\text{aux.htags}[i] = \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{pk}}(i, \text{aux.hsig})$ (non-malleability, see Appendix A)
 - 861 5. $(\text{aux.jsins}[i].\text{znote}.v) \cdot (1 - e) = 0$ is satisfied for the boolean value e set such
 862 that if $\text{aux.jsins}[i].\text{znote}.v > 0$ then $e = 1$.
 - 863 6. The Merkle root mkroot' obtained after checking the Merkle authentica-
 864 tion path $\text{aux.jsins}[i].\text{mkpath}$ of commitment $\text{aux.jsins}[i].\text{cm}$, with MKHASH,
 865 equals to prim.mkroot if $e = 1$.
 - 866 7. $\text{prim.nfs}[i]$
 867 $= \{\text{Pack}_{\mathbb{F}_{\text{rCUR}}}(\text{aux.jsins}[i].\text{nf}[k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}])\}_{k \in [[\text{PRFNFOUTLEN}/\text{FIELD CAP}]}$
 868 (see Section 3.3.1 for definition of Pack)
 - 869 8. $\text{prim.htags}[i]$
 870 $= \{\text{Pack}_{\mathbb{F}_{\text{rCUR}}}(\text{aux.htags}[i][k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}])\}_{k \in [[\text{PRFPKOUTLEN}/\text{FIELD CAP}]}$
 871 (see Section 3.3.1 for definition of Pack)

- 872 • For each $j \in [\text{JSOUT}]$:
- 873 1. $aux.znotes[j].\rho = \text{PRF}_{aux.\phi}^{\text{rho}}(j, aux.hsigs)$ (non-malleability, see Appendix A)
- 874 2. $prim.cms[j] = \text{ComSch.Com}(aux.znotes[j].apk, aux.znotes[j].\rho, aux.znotes[j].v;$
- 875 $aux.znotes[j].r)$
- 876 • $prim.hsigs = \{\text{Pack}_{\mathbb{F}_{\text{rCUR}}}(aux.hsigs[k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}])\}_{k \in [\lfloor \text{CRHHSIGOUTLEN}/\text{FIELD CAP} \rfloor]}$
- 877 (see Section 3.3.1 for definition of Pack)
- 878 • $prim.rsd = \text{Pack}_{rsd}(\{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}, aux.vin, aux.vout, aux.hsigs, \{aux.htags[i]\}_{i \in [\text{JSIN}]})$
- 879 (see Section 3.3.1 for definition of Pack_{rsd})
- Check that the “joinsplit is balanced”, i.e. check that the joinsplit equation holds:¹

$$\begin{aligned}
& \text{Pack}_{\mathbb{F}_{\text{rCUR}}}(aux.vin) + \sum_{i \in [\text{JSIN}]} \text{Pack}_{\mathbb{F}_{\text{rCUR}}}(aux.jsins[i].note.v) \\
&= \sum_{j \in [\text{JSOUT}]} \text{Pack}_{\mathbb{F}_{\text{rCUR}}}(aux.znotes[j].v) + \text{Pack}_{\mathbb{F}_{\text{rCUR}}}(aux.vout)
\end{aligned}$$

880 2.3 Generating the inputs of the Mix function (Mix_{in})

881 In the following section, we assume that the system is initialized. In other words, we
882 assume that a ledger L is available (i.e. an **Ethereum** network is operated by a set of
883 miners), the **Mixer** contract is deployed on L . Likewise, we assume that the public
884 parameters $pp_{\text{ZkSnarkSch}} \leftarrow \text{ZkSnarkSch.KGen}(1^\lambda, \mathbf{R}^z)$ are available to **Mixer** and to all
885 parties willing to call the Mix function of **Mixer**. Furthermore, we assume that there
886 exists a set of **Ethereum** and **Zeth** users, and that the *payment address* of each **Zeth** user
887 is easily discoverable. In the rest of this section, the set of *payment addresses* discovered
888 by a zeth user \mathcal{U}_Z is represented as a list attribute $\mathcal{U}_Z.\text{keystore}$ indexed by usernames.

889 In order for \mathcal{U}_Z to transact via **Zeth**, \mathcal{U}_Z needs to create an object Mix_{in} of type
890 **MixInputDType** to pass to the Mix function of **Mixer**:

- 891 1. Create an object *prim* of type **PrimInputDType** to represent the primary input,
- 892 and an object *aux* of type **AuxInputDType** to represent the auxiliary input, where:
 - 893 (a) $prim.mkroot \in \text{Roots}$, where *Roots* is the set of *all* Merkle roots corresponding
 - 894 to one of the state of the Merkle tree on **Mixer** containing *all* the commit-
 - 895 ments to the input notes, in $aux.jsins$, in its set of leaves.
 - 896 (b) $aux.znotes[j].r \leftarrow_{\$} \mathbb{B}^{\text{RTRAPLEN}}, \forall j \in [\text{JSOUT}]$, and $aux.\phi \leftarrow_{\$} \mathbb{B}^{\text{PHILEN}}$

¹where $\text{Pack}_{\mathbb{F}_{\text{rCUR}}}(x)$ outputs the numerical value of x in \mathbb{F}_{rCUR} . We rely on the fact that $\text{ZVALUELEN} < \text{FIELD CAP}$ to perform this sum.

897 (c) The public values $(aux.vin, aux.vout) \in (\mathbb{B}^{\text{ZVALUELEN}})^2$, $aux.znotes[j].v$ and
 898 $aux.znotes[j].apk \forall j \in [\text{JSOUT}]$ are all set by the sender, $\mathcal{U}_{\mathcal{Z}}$, as desired as
 899 long as they satisfy the joinsplit equation.

(d) All attributes of the *prim* and *aux* objects should be derived as specified in the statement (see Section 2.2), alongside a signature hash ($aux.hsigs$) that is generated as the hash of the nullifiers and a one-time signing verification key (non-malleability, see Appendix A), using the desired signature scheme $\text{SigSch}_{\text{OT-SIG}}$ (see Section 3.4):

$$(sk_{\text{OT-SIG}}, vk_{\text{OT-SIG}}) = \text{SigSch}_{\text{OT-SIG}}.\text{KGen}(1^\lambda) \quad (2.1)$$

$$aux.hsigs = \text{CRH}^{\text{hsig}}(\{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}, vk_{\text{OT-SIG}}) \quad (2.2)$$

900 (e) $\text{Mix}_{in}.primIn \leftarrow prim$

Note

If one of the attributes of *prim* and *aux* is not correctly generated, then the proof of computational integrity generated in the next step will be rejected on **Mixer**, and the state of **Mixer** will not be modified.

901

902 2. Generate a zk-SNARK proof π to prove, in zero-knowledge, that the relation \mathbf{R}^z
 903 (Section 2.2) holds on the primary and auxiliary inputs, using the desired zk-
 904 SNARK scheme ZkSnarkSch (see Section 3.6):

905 (a) $\pi \leftarrow \text{ZkSnarkSch}.\text{P}(pp_{\text{ZkSnarkSch}}, prim, aux)$

906 (b) $\text{Mix}_{in}.proof \leftarrow \pi$

907 3. Encrypt all the *aux.znotes* using the recipient's *payment address*, using the en-
 908 cryptation scheme EncSch (see Section 3.5).

(a) For all $j \in [\text{JSOUT}]$, do:

$$ct_j \leftarrow \text{EncSch}.\text{Enc}(aux.znotes[j], \mathcal{U}_{\mathcal{Z}}.keystore[recipient_j].pub.pkenc)$$

909 (b) $\text{Mix}_{in}.ciphers \leftarrow \{ct_j\}_{j \in [\text{JSOUT}]}$

910 4. Generate a signature $\sigma_{\text{OT-SIG}}$ on the inputs of the *Mix* function, in order to prevent
 911 any malleability attacks (c.f. Appendix A), using the desired signature scheme
 912 $\text{SigSch}_{\text{OT-SIG}}$ (see Section 3.4):

(a) Using the one-time signature keypair generated in Eq. (2.1), do:

$$\begin{aligned} dataToBeSigned &= \mathcal{S}_{\mathcal{E}}.Addr \parallel \text{Mix}_{in}.primIn \parallel \text{Mix}_{in}.\pi \parallel \text{Mix}_{in}.ciphers \\ \sigma_{\text{OT-SIG}} &= \text{SigSch}_{\text{OT-SIG}}.\text{Sig}(sk_{\text{OT-SIG}}, \text{CRH}^{\text{ots}}(dataToBeSigned)) \end{aligned}$$

913 (b) $\text{Mix}_{in}.otssig \leftarrow \sigma_{\text{OT-SIG}}$

914 (c) $\text{Mix}_{in}.otsvk \leftarrow vk_{\text{OT-SIG}}$

915 Here, $\mathcal{S}_{\mathcal{E}}.Addr$ represents the address of the **Ethereum** user $\mathcal{S}_{\mathcal{E}}$ who must sign the
916 transaction (see Section 2.4). In general, this is likely to be owned by the holder
917 $\mathcal{U}_{\mathcal{Z}}$ of the **Zeth** notes to be spent, but this is not a requirement.

Remark 2.3.1. Remark 2.5.1 describes a situation in which the proof data $\text{Mix}_{in}.\pi$ is not available (having been verified by some external mechanism). In such cases, *dataToBeSigned* in Item 4a MAY be replaced with:

$$dataToBeSigned = \mathcal{S}_{\mathcal{E}}.Addr \parallel \text{Mix}_{in}.primIn \parallel \text{Mix}_{in}.ciphers$$

918 This modification, mentioned here for completeness, **MUST NOT** be used except as de-
919 scribed in Remark 2.5.1.

920 2.4 Creating an Ethereum transaction tx_{Mix} to call $\widetilde{\text{Mixer}}$

After generating a Mix_{in} object, $\mathcal{U}_{\mathcal{Z}}$ can generate an object tx_{raw} of type `TxRawDType`, such that:

$$tx_{raw}.to = \widetilde{\text{Mixer}}.Addr \wedge tx_{raw}.data = zdata$$

921 Then, an **Ethereum** user $\mathcal{S}_{\mathcal{E}}$ can ECDSA sign tx_{raw} , under $\mathcal{S}_{\mathcal{E}}.sk$ in order to transform
922 this object of type `TxRawDType` into an finalized transaction, i.e. an object tx_{Mix} of type
923 `TxDType`.

924 Finally, the transaction tx_{Mix} is broadcasted on the **Ethereum** network and eventually
925 gets mined.

Note

Here, the **Ethereum** user $\mathcal{S}_{\mathcal{E}}$ who sends the final transaction, and the **Zeth** user $\mathcal{U}_{\mathcal{Z}}$ may represent the same person or entity, but this is not necessarily the case. It is perfectly feasible (and in some cases may be desirable) for a **Zeth** user $\mathcal{U}_{\mathcal{Z}}$ to create a **Zeth** transaction which is later signed by a distinct party $\mathcal{S}_{\mathcal{E}}$. In particular, the only identifying information that appears in plaintext on the ledger will be that of $\mathcal{S}_{\mathcal{E}}$.

926

927 2.5 Processing tx_{Mix}

928 When a tx_{Mix} is mined (hence assuming that $\text{EthVerifyTx}(tx_{\text{Mix}})$ returns `true`), the state
929 transition specified by the `Mix` function of $\widetilde{\text{Mixer}}$ is executed.

930 To preserve the soundness of **Zeth**, and make sure that no $\mathcal{U}_{\mathcal{Z}}$ is able to create
931 value by double spending *ZethNotes*, various checks need to be satisfied. The function
932 `ZethVerifyTx` is defined as the function that returns `true` if all the checks are satisfied,
933 and `false` otherwise.

934 If $\text{ZethVerifyTx}(tx_{\text{Mix}})$ returns `true`, then `Mix` modifies the “World state” ς to account
 935 for the spent *ZethNotes* and the newly generated ones. However, if $\text{ZethVerifyTx}(tx_{\text{Mix}})$
 936 returns `false`, then the state transition ends.

Note

Even if $\text{ZethVerifyTx}(tx_{\text{Mix}})$ returns `false`, ς is modified since the `Ethereum` balances of the transaction originator is decremented by the sum of `DGAS` and the gas consumed by the ZethVerifyTx function, and the balance of the `Ethereum` account of the miner gets incremented by the same amount.

937

938 Thus, `Mix` proceeds as follows:

1. Check that all the values of the primary inputs’ ($\text{Mix}_{in}.primIn$) entries are elements of the scalar field over which the zk-proof is generated:

$$\text{Mix}_{in}.primIn \in \mathbb{F}_{\text{rCUR}}^*$$

2. Unpack the nullifiers, signature hash and public values (see Section 3.3.1 for the definitions of the `Unpack` functions):

$$\begin{aligned} nf_i &= \text{Unpack}_{nf}(\text{Mix}_{in}.primIn.nfs[i], \text{Mix}_{in}.primIn.rsd) \quad \forall i \in [\text{JSIN}] \\ v_{in} &= \text{decode}_{\mathbb{N}}(\text{Unpack}_{v_{in}}(), \text{Mix}_{in}.primIn.rsd) \\ v_{out} &= \text{decode}_{\mathbb{N}}(\text{Unpack}_{v_{out}}(), \text{Mix}_{in}.primIn.rsd) \\ hsig &= \text{Unpack}_{hsig}(\text{Mix}_{in}.primIn.hsig, \text{Mix}_{in}.primIn.rsd) \end{aligned}$$

939

3. Check the validity of the tx_{Mix} object (ZethVerifyTx):

- (a) Check that $\text{Mix}_{in}.primIn.hsig$ is correctly computed, i.e. check that the following equation holds (to prevent transaction malleability, see Appendix A):

$$hsig = \text{CRH}^{\text{hsig}}(\text{Mix}_{in}.primIn.nfs, \text{Mix}_{in}.otsvk)$$

- (b) Check that π is a valid zk-SNARK proof for $\text{Mix}_{in}.primIn$, i.e. check that:

$$\text{ZkSnarkSch.V}(pp_{\text{ZkSnarkSch}}, \pi, \text{Mix}_{in}.primIn) = \text{true}$$

- (c) Check that none of the nullifiers in $\text{Mix}_{in}.primIn.nfs$ have already been used, i.e. check that:

$$nf_i \notin \text{Nulls}, \forall i \in [\text{JSIN}]$$

940

where *Nulls* is the set of all nullifiers that are “declared” on $\widetilde{\text{Mixer}}$.

- (d) Check that $\text{Mix}_{in}.otssig$ is a valid signature of the `Ethereum` sender’s address *Addr* (see Section 2.4) and the attributes of Mix_{in} , to prevent transaction malleability (see Appendix A), i.e. check that:

$$\text{SigSch}_{\text{OT-SIG}}.\text{Vf}(\text{Mix}_{in}.otsvk, m, \text{Mix}_{in}.otssig) = \text{true}$$

where $dataToBeSigned = Addr \parallel \widetilde{Mix_{in}.primIn} \parallel Mix_{in}.\pi \parallel Mix_{in}.ciphers$,
and $m = CRH^{ots}(dataToBeSigned)$

- (e) Check that $\widetilde{Mix_{in}.primIn.mkroot}$ corresponds to a valid state of the Merkle tree held on **Mixer**, i.e. check that:

$$\widetilde{Mix_{in}.primIn.mkroot} \in Roots'$$

941 where $Roots'$ is the set of all Merkle roots corresponding to one of the states
942 of the Merkle tree.

- (f) Check that vin corresponds to the value val of the transaction object, i.e. check that:

$$vin = tx_{Mix}.val$$

943 4. If all checks above pass, i.e. if $ZethVerifyTx(tx_{Mix})$ returns true, then the following
944 additional modifications are made in ς :

- 945 (a) Add the commitments $Mix_{in}.primIn.cms$ to the Merkle tree held on $\widetilde{\mathbf{Mixer}}$.
946 (b) $Roots' \leftarrow Roots' \cup \{mkroot'\}$, where $mkroot'$ is the Merkle root of the Merkle
947 tree after insertion of the commitments $Mix_{in}.primIn.cms$ in the Merkle tree.
948 (c) $Nulls \leftarrow Nulls \cup \{nf_i\}_{i \in [JSIN]}$, i.e. the nullifiers nfs become “declared”.
949 (d) Modify the **Ethereum** balances according to the public values:
950 • $\varsigma[\mathcal{S}_{\mathcal{E}}.Addr].bal = \varsigma[\mathcal{S}_{\mathcal{E}}.Addr].bal - vin$
951 • $\varsigma[\mathcal{S}_{\mathcal{E}}.Addr].bal = \varsigma[\mathcal{S}_{\mathcal{E}}.Addr].bal + vout$
952 • $\widetilde{\mathbf{Mixer}}.bal = \widetilde{\mathbf{Mixer}}.bal + vin$
953 • $\widetilde{\mathbf{Mixer}}.bal = \widetilde{\mathbf{Mixer}}.bal - vout$
954 (e) Emit an event (Section 1.2.3) $evMixOut$ of type $MixEventDType$, contain-
955 ing the new root $mkroot'$ of the Merkle tree of commitments, the nullifiers
956 $\{nf_i\}_{i \in [JSIN]}$, commitments to the newly created $ZethNotes$ $Mix_{in}.primIn.cms$,
957 and the corresponding ciphertexts $Mix_{in}.primIn.ciphers$.

958 **Remark 2.5.1.** In some deployments, verification of the zk-SNARK proof π may be
959 delegated to an external mechanism (in such a way that integrity of the system can
960 still be guaranteed), and π may not appear as public data on-chain. For example, where
961 multiple **Zeth** transactions are aggregated by a system such as that described in [Ron20],
962 the original zk-SNARK proofs become auxiliary inputs to a “wrapping” SNARK, which
963 checks their validity via a single proof verification. A modified version of the Mix function
964 receives Mix parameters from a specific contract (known to behave correctly with respect
965 to the delegation protocol) without $Mix_{in}.\pi$.

In this case, the value of $dataToBeSigned$ in Item 3d may be replaced by:

$$dataToBeSigned = Addr \parallel \widetilde{Mix_{in}.primIn} \parallel Mix_{in}.ciphers,$$

966 and the equivalent change must be made when generating the `Mix` parameters, as de-
 967 scribed in Remark 2.3.1. The transaction, as presented to an aggregator, is malleable
 968 since the `Groth16` proofs can be modified in a way that preserves the validity. How-
 969 ever, once aggregated, the transaction proof data for the transaction does not appear on
 970 chain (it is an auxiliary input to some externally generated “wrapping” poof). Hence,
 971 the transaction can only be identified by the remaining public data, which is protected
 972 by the one-time signature $\sigma_{\text{OT-SIG}}$, and hence non-malleable.

973 The external scheme used to verify the zk-SNARK proof must specify the exact re-
 974 quirements of the contract and how it should be modified, including any further checks
 975 that must be carried out. Thus, this modification is not part of the core `Zeth` pro-
 976 tocol described in this document, but an augmentation forming part of an external
 977 protocol. However, for completeness, we briefly describe a *dispatch* entry point in the
 978 proof-of-concept `Mixer` implementation which supports delegation of proof verification,
 979 as described above. The *dispatch* entry point performs the following checks:

- 980 • Check that the `Mixer` has been deployed with the (immutable) address of a trusted
 981 contract, permitted to call this entry point. Otherwise, abort.
- 982 • Check that the caller *msg.sender* matches the permitted caller set at deployment
 983 time, otherwise abort.
- 984 • Perform all checks related to the `Mix` parameters, except Item 3b, with the modi-
 985 fication to Item 3d described in this remark.

986 After these checks, the state-transition continues as normal.

987 The `Zeth` client implementation is also augmented to include a flag to enable the
 988 corresponding change described in Remark 2.3.1 (to generate a signature on the modified
 989 *dataToBeSigned*). Naturally, the *dispatch* entry point can only be used with parameters
 990 generated using this flag (otherwise the signature check will fail).

991 Such modifications **MUST NOT** be implemented except as described by the secure ex-
 992 ternal scheme for delegating proof verification.

993 2.6 Receiving *ZethNotes*

994 In order to confirm the reception of *ZethNotes*, \mathcal{R}_Z must listen to the events (Sec-
 995 tion 1.2.3) of type `MixEventDType` emitted by the processing of tx_{MIX} , and try to decrypt
 996 the ciphertexts using $\mathcal{R}_Z.\text{priv.skenc}$ to see if he is the recipient of a `Zeth` payment. If
 997 the decryption is successful (\mathcal{R}_Z is the recipient of a payment), \mathcal{R}_Z must verify that the
 998 *ZethNote* recovered is the opening of a commitment in the Merkle tree of `Mixer`. If not,
 999 \mathcal{R}_Z rejects the (invalid) payment.

1000 We describe below the steps that \mathcal{R}_Z needs to carry out for all events *evMixOut* \in
 1001 `MixEventDType` emitted by `Mixer`, in order to receive payments:

1002 1. Compute the new root $mkroot'$ of the Merkle tree of commitments, after adding the
 1003 new values $ev\widetilde{MixOut}.cms$. If this value does not match the new root $evMixOut.mkroot$
 1004 emitted by **Mixer**, abort.

2. Try to decrypt the ciphertexts:

$$zn_j = \text{EncSch.Dec}(\mathcal{R}_Z.\text{priv.skenc}, evMixOut.ciphers[j])$$

1005 3. For each successful decryption, let j be the index of the decrypted ciphertext:

1006 (a) Check whether the recovered plaintext zn_j is a well-formed *ZethNote*. Abort
 1007 if it is not well-formed.

(b) Check that the recovered *ZethNote* zn_j is the opening of the corresponding
 commitment $evMixOut.cms[j]$:

$$evMixOut.cms[j] = \text{ComSch.Com}(zn_j.apk, zn_j.\rho, zn_j.v; zn_j.r)$$

1008 Abort if the note is not a valid opening.

1009 (c) Additionally, if sender \mathcal{S}_Z , and recipient \mathcal{R}_Z had a contractual agreement,
 1010 then \mathcal{R}_Z needs to check that the terms of this agreement are fulfilled by all
 1011 the recovered *ZethNotes*, abort otherwise.

1012 Note that Steps 1 and 3b are required to ensure that data decrypted by \mathcal{R}_Z ex-
 1013 actly matches the data committed to in $\widetilde{\text{Mixer}}$. In particular, Step 1 requires \mathcal{R}_Z
 1014 to maintain or have access to some representation of the Merkle tree of commitments.
 1015 See Section 4.1.2 for further details.

1016 2.7 Security requirements for the primitives

1017 We list below the security requirements to instantiate the primitives of the **Zeth** protocol.

- 1018 • CRH^{hsig} and CRH^{ots} MUST be collision resistant functions (see Definition 1.5.16).
- 1019 • PRF^{addr} , PRF^{nf} , PRF^{rho} and PRF^{pk} MUST be PRF when keyed by ask and ϕ , and be
 1020 collision resistant (see Definition 1.5.16, and Section 1.5.8).
- 1021 • $\text{SigSch}_{\text{OT-SIG}}$ MUST be UF-CMA (see Definition 1.5.24 and Appendix A.2.3).
- 1022 • ComSch MUST be computationally hiding and binding (see Section 1.5.9).
- 1023 • MKHASH MUST be collision resistant with $h_0 = 0_{\mathbb{F}_{\text{rCUR}}}$ (see Section 1.5.6).²
- 1024 • EncSch MUST be IND-CCA2 and IK-CCA (see, respectively, [ABR99, Definition 8]
 1025 and Definition 1.5.10).
- 1026 • $\text{Unpack}(\text{Pack}(X)) = X$ and $\text{Unpack}(\text{Pack}_{\text{rsa}}(X)) = X$ MUST hold.
- 1027 • $\text{decode}(\text{encode}(X)) = X$ MUST hold.

²This security requirement is equivalent to the one in [ZCa19, Section 5.4.1.3] where finding a preimage of $0^{\text{SHA256DLEN}}$ must be hard.

1028 **2.7.1 Additional notes**

1029 **Defining $hsig$**

The signature hash $hsig$ is a variable used to bind the signature keys to the primary inputs. We use the same definition of $hsig$ as **Zcash** to prevent the Faerie Gold attack and thus

$$hsig = \text{CRH}^{\text{hsig}}(nfs, vk).$$

1030 As a private transaction is uniquely determined by its nullifiers $nfs = (nf_0, \dots, nf_{\text{JSIN}-1})$,
1031 and because of the collision resistance of CRH^{hsig} , a transaction is uniquely determined
1032 by $hsig$ (with overwhelming probability). We did not use the *randomSeed* defined in
1033 **Zcash** however, since this is only necessary to achieve uniqueness of $hsig$ for transactions
1034 *in transit* (i.e. not mined yet) [Hop16]. The uniqueness of $hsig$ is a requirement to
1035 prevent the Fairy Gold attack.

1036 **Security Requirement.**

- 1037 • The variable $hsig$ **MUST** be derived from the nullifiers $\{nf_i\}_{i \in [\text{JSIN}]}$ and the signing
1038 key vk using a collision resistant function. Doing so, makes sure that $hsig$ is unique
1039 for each tx_{Mix} with overwhelming probability.

1040 **Defining ρ**

We define ρ like in **Zcash** in order to prevent the Faerie Gold attack. A malicious sender could reuse the same ρ for a given recipient, hence correctly generating a *ZethNote* which could become unspendable by the recipient. Making ρ the output of a collision resistant PRF with random variable ϕ as key and with tx_{Mix} 's $hsig$ as input ensures, with overwhelming probability, the uniqueness of ρ and prevents this attack. Thus,

$$\rho_j = \text{PRF}_{\phi}^{\text{rho}}(j, hsig).$$

1041 **Message authentication tags h_i**

The message authentication tags are used to bind the signature hash to the input notes spending keys, to show ownership of the spent notes. Each tag derived from a note owner's spending key and the signature hash **MUST** be unique for each note with overwhelming probability. We define

$$h_i = \text{PRF}_{ask_i}^{\text{pk}}(i, hsig).$$

1042 Chapter 3

1043 Instantiation of the cryptographic 1044 primitives

1045 In this chapter, we start by instantiating the cryptographic building blocks used in
1046 previous sections to describe the `Zeth` DAP design. Finally, we proceed by providing
1047 security proofs justifying that our instantiation complies with the security requirements
1048 listed in previous sections.

1049 Note that, in several cases, it is necessary to specify details in terms of concrete
1050 properties of the curve `Curve` and associated scalar field \mathbb{F}_{CUR} . In these cases, we focus
1051 on two curves of interest: `BN-254` and `BLS12-377`. We note, however, that other suitable
1052 curves could be used.

1053 `BN-254` [Rk19] has several properties that make it implementation-friendly. Ele-
1054 ments of both the base field and scalar field can be represented in `ETHWORDLEN` bits (the
1055 native word size of the EVM), allowing efficient encoding and manipulation of such ele-
1056 ments. Moreover, a subset of operations on `BN-254` are supported by the EVM through
1057 precompiled contracts. These precompiled contracts enable verification of signatures
1058 (Section 3.4) and zero-knowledge proofs (Section 3.6), required by this protocol, with
1059 minimal gas overhead.

1060 `BLS12-377` [BCG⁺20], like `BN-254`, has the advantage that scalar field elements can
1061 be represented within `ETHWORDLEN`-bit words (although the same is not true of base
1062 field elements). However, the `EVM` provides no native support for `BLS12-377`, which
1063 increases the complexity of the `Mixer` implementation (see Section 2.5 for details of the
1064 operations to be performed). An advantage that `BLS12-377` does provide, is that it is
1065 the “inner” curve of a one-layer chain (as described in [BCG⁺20, HG20]). Therefore
1066 zero-knowledge proofs using `BLS12-377` can be efficiently verified by statements in other
1067 zero-knowledge proofs using an appropriate “outer” pairing. Support for `BLS12-377`
1068 in `Zeth` therefore admits several applications (no explicitly covered by this document),
1069 such as aggregation of proofs over multiple `Zeth` transactions (e.g. [Ron20]).

1070 Further details related to implementation and optimization are given in Chapter 4.

3.1 Instantiating the PRFs, ComSch and CRHs

The functions CRH^{hsig} and CRH^{ots} are instantiated with SHA256 [oST15] which we assume to be collision resistant. Furthermore, ComSch , $\text{PRF}^{\text{pk}}(x)$, $\text{PRF}^{\text{rho}}(x)$, $\text{PRF}^{\text{addr}}(x)$, and $\text{PRF}^{\text{nf}}(x)$ are all instantiated with Blake2’s hash function optimized for 32-bit platforms, Blake2s, which we prove in the Weakly Ideal Cipher Model [LMN16] to be from a family of PRF and collision resistant functions. The Weakly Ideal Cipher model assumes that Blake2’s underlying block cipher is ideal and has no structural weaknesses (see Appendix D.2). In addition to that, and to ensure that the functions $\text{PRF}^{\text{pk}}(x)$, $\text{PRF}^{\text{rho}}(x)$, $\text{PRF}^{\text{addr}}(x)$, and $\text{PRF}^{\text{nf}}(x)$ compute images lying in different domains, we use different message prefixes (or “domain separators”) for the PRFs inputs. This approach ensures that the apk_i ’s, nf_i ’s, ρ_i ’s, and h_i ’s have independent distributions from a PPT adversary point of view.

Note

It is important to note that, for this approach to be secure, the hash function used needs to be secure against *chosen-prefix collision attacks* [Ste15].

Furthermore, we take:

- $\text{RTRAPLEN}, \text{ASKLEN}, \text{PHILEN} = \text{BLAKE2sCLEN}$

3.1.1 Blake2 primitive

Blake [AHMP08] is a hash family that was presented as a candidate at the SHA3 competition. Blake2 is the next iteration of the family which has been further optimized to achieve higher throughput thanks to some optimizations and by being less conservative on its security¹. Blake and Blake2 are based on the ChaCha stream cipher [Ber08a] composed with the HAIFA framework [BD07]. ChaCha defined over 20 rounds, as used in Blake2, is deemed secure and a PRF based on today’s cryptanalysis [Pro14, CM16]. Blake2 is specified in RFC-7693 [MJS15] and licensed under CC0. Blake2s is an instantiation of Blake2 optimized for 32-bit platforms. As such, to reason about the security of Blake2s we prove the security of Blake2.

Blake security Blake security has been heavily scrutinized through the SHA3 competition [VNP10, MQZ10, AMP10, AAM12, AMPŠ12, ALM12, HMRS12]. Blake2 has also been thoroughly cryptanalyzed independently [GKN⁺14, Hao14, EFK15, NA19]. For n -bit long digests/outputs, the hash and compression functions present $n/2$ -bit of collision resistance and n -bit of preimage resistance, immunity to length extension, and indistinguishability from a random oracle [ANWOW13]. They have furthermore been demonstrated secure in the Weakly Ideal Cipher Model [LMN16] (WICM, see Appendix D.1.1). More

¹The authors increased the number of rounds of Blake for the SHA3 competition to be more conservative on security. They however showed afterwards that this change was not “meaningfully more secure” and thus reverted it for Blake2 (see [ANWOW13, Section 2.1]).

1103 particularly, Luykx et al. show that Blake2 is indifferentiable from a random oracle in
 1104 this model and is a PRF. We use this result in Appendix D.2 to show the collision
 1105 resistance of Blake2. We also prove that, given that Blake2 is collision resistant and a
 1106 PRF, $\text{Blake2}(r||x)$ is a computationally binding and computationally hiding commitment
 1107 scheme for input x and randomness r .

Note

We assume that the encryption scheme used in the Blake2 underlying compression function – which is derived from ChaCha20 – has no exploitable structural behaviour. More precisely, that this encryption scheme behaves like a weak ideal cipher. We provide proofs in this model.

1108

1109 3.1.2 Commitment scheme

We define our commitment scheme as follows,

$$\begin{aligned} \text{ComSch.Setup} & : \{1^\lambda \text{ s.t. } \lambda \in \mathbb{N}\} \rightarrow \mathbb{B}^* \\ \text{ComSch.Com} & : (\mathbb{B}^{\text{PRFADDRROUTLEN}} \times \mathbb{B}^{\text{PRFRHOOUTLEN}} \times \mathbb{B}^{\text{ZVALUELEN}}) \times \mathbb{B}^{\text{RTRAPLEN}} \rightarrow \mathbb{F}_{\mathbf{r}_{\text{CUR}}} \end{aligned}$$

We instantiate the commitment scheme with Blake2s as follows,

$$\begin{aligned} pp & = \text{ComSch.Setup}(1^\lambda) \text{ (corresponds to Blake2s's constant PB and } \mathbf{r}_{\text{CUR}}) \\ cm & = \text{ComSch.Com}(m = (apk, \rho, v); r) \\ & = \text{decode}_{\mathbb{N}}(\text{Blake2s}(r||apk||\rho||v)) \pmod{\mathbf{r}_{\text{CUR}}} \end{aligned}$$

1110 **Remark 3.1.1.** We set the commitment digest length in the parameter block PB [MJS15].

1111 Security proof

1112 The commitment scheme defined above is computationally hiding and binding in the
 1113 WICM, see Appendix D.2.4. However, because of the modulo \mathbf{r}_{CUR} operation, the scheme
 1114 is only $(\text{FIELDLEN}/2)$ -bit binding.

1115 3.1.3 PRFs

We show in this section how we instantiate the PRFs with Blake primitives. As a reminder the PRFs are defined as follows,

$$\begin{aligned} \text{PRF}^{\text{addr}} & : \mathbb{B}^{\text{ASKLEN}} \times \{0\} \rightarrow \mathbb{B}^{\text{PRFADDRROUTLEN}} \\ \text{PRF}^{\text{pk}} & : (\mathbb{B}^{\text{ASKLEN}} \times [\text{JSIN}]) \times \mathbb{B}^{\text{CRHHSIGOUTLEN}} \rightarrow \mathbb{B}^{\text{PRFPKOUTLEN}} \\ \text{PRF}^{\text{nf}} & : \mathbb{B}^{\text{ASKLEN}} \times \mathbb{B}^{\text{PRFRHOOUTLEN}} \rightarrow \mathbb{B}^{\text{PRFNFOUTLEN}} \\ \text{PRF}^{\text{rho}} & : (\mathbb{B}^{\text{PHILEN}} \times [\text{JSOUT}]) \times \mathbb{B}^{\text{CRHHSIGOUTLEN}} \rightarrow \mathbb{B}^{\text{PRFRHOOUTLEN}} \end{aligned}$$

As we instantiate the PRFs with Blake2s we have,

$$\text{PRFADDRROUTLEN, PRFNFOUTLEN, PRFPKOUTLEN, PRFRHOOUTLEN} = \text{BLAKE2sCLEN}$$

To ensure that the PRFs have independent distributions, we first introduce tagging functions tag^x which truncate and prepend with a distinct tag the PRFs key. We have,

$$\begin{aligned} \text{tag}^{\text{addr}} &: \mathbb{B}^{\text{ASKLEN}} \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \\ \text{tag}^{\text{pk}} &: \mathbb{B}^{\text{ASKLEN}} \times [\text{JSIN}] \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \\ \text{tag}^{\text{nf}} &: \mathbb{B}^{\text{ASKLEN}} \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \\ \text{tag}^{\text{rho}} &: \mathbb{B}^{\text{PHILEN}} \times [\text{JSOUT}] \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \end{aligned}$$

The tagging functions are instantiated as follows,

$$\begin{aligned} \text{tag}^{\text{addr}}(\text{aux.jsins}[i].\text{ask}) &= \text{tag}_{\text{ask}}^{\text{addr}} \\ &= (1) \parallel (1) \left\lceil \frac{\text{JSMAX}}{2} \right\rceil \parallel (0, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN} - 3 - \left\lceil \frac{\text{JSMAX}}{2} \right\rceil}(\text{aux.jsins}[i].\text{ask}) \\ \text{tag}^{\text{nf}}(\text{aux.jsins}[i].\text{ask}) &= \text{tag}_{\text{ask}}^{\text{nf}} \\ &= (1) \parallel (1) \left\lceil \frac{\text{JSMAX}}{2} \right\rceil \parallel (1, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN} - 3 - \left\lceil \frac{\text{JSMAX}}{2} \right\rceil}(\text{aux.jsins}[i].\text{ask}) \\ \text{tag}^{\text{pk}}(\text{aux.jsins}[i].\text{ask}, i) &= \text{tag}_{\text{ask}, i}^{\text{pk}} \\ &= (0) \parallel \text{pad}_{\left\lceil \frac{\text{JSMAX}}{2} \right\rceil}(\text{encode}_{\mathbb{N}}(i)) \parallel (0, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN} - 3 - \left\lceil \frac{\text{JSMAX}}{2} \right\rceil}(\text{aux.jsins}[i].\text{ask}) \\ \text{tag}^{\text{rho}}(\text{aux}.\phi, j) &= \text{tag}_{\text{ask}, j}^{\text{rho}} \\ &= (0) \parallel \text{pad}_{\left\lceil \frac{\text{JSMAX}}{2} \right\rceil}(\text{encode}_{\mathbb{N}}(j)) \parallel (1, 0) \parallel \text{trunc}_{\text{BLAKE2sCLEN} - 3 - \left\lceil \frac{\text{JSMAX}}{2} \right\rceil}(\text{aux}.\phi) \end{aligned}$$

1116 where $\text{pad}_{\left\lceil \frac{\text{JSMAX}}{2} \right\rceil}(\text{encode}_{\mathbb{N}}(i))$ is the function that pads the binary representation of i by
 1117 adding 0's before the most significant bit (e.g. assuming big endian encoding, $\text{pad}_2(\text{encode}_{\mathbb{N}}(1)) =$
 1118 01).

We now present how the PRFs are instantiated,

$$\begin{aligned} \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{addr}}(0) &= \text{aux.jsins}[i].\text{znote.apk} \\ &= \text{Blake2s}(\text{tag}^{\text{addr}}(\text{aux.jsins}[i].\text{ask}) \parallel \text{pad}_{\text{BLAKE2sCLEN}}(0)) \\ \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{nf}}(\text{aux.jsins}[i].\rho) &= \text{prim.nfs}[i] \\ &= \text{Blake2s}(\text{tag}^{\text{nf}}(\text{aux.jsins}[i].\text{ask}) \parallel \text{aux.jsins}[i].\text{znote}.\rho) \\ \text{PRF}_{\text{aux.jsins}[i].\text{ask}}^{\text{pk}}(i, \text{prim.hsig}) &= \text{prim.htags}[i] \\ &= \text{Blake2s}(\text{tag}^{\text{pk}}(\text{aux.jsins}[i].\text{ask}, i) \parallel \text{prim.hsig}) \\ \text{PRF}_{\text{aux}.\phi}^{\text{rho}}(j, \text{prim.hsig}) &= \text{aux.znotes}[j].\rho \\ &= \text{Blake2s}(\text{tag}^{\text{rho}}(\text{aux}.\phi, j) \parallel \text{prim.hsig}) \end{aligned}$$

1119 **Remark 3.1.2.** We set the PRFs' output length in the Blake2s's parameter block PB.

1120 Security proof

1121 The functions defined above are collision resistant and PRFs in the WICM, see Ap-
1122 pendix D.2. Because of the tagging functions, the security parameter of the PRFs be-
1123 comes $\lambda = \text{BLAKE2sCLEN}/2 - \text{JSMAX}/4 - 3/2$.

1124 3.1.4 Collision resistant hashes

We instantiate in this section the collision resistant hash functions CRH^{hsig} and CRH^{ots} with SHA256. As a consequence, we have,

$$\text{CRHHSIGOUTLEN} = \text{CRHOTSOUTLEN} = \text{SHA256DLEN}$$

1125 **SHA256 Security** SHA-256 (Secure Hash Algorithm 256) is a hash function designed
1126 by the National Security Agency (NSA) in 2001. It is based on the Merkle–Damgård
1127 structure, the Davies–Meyer compression function construct [BRS02, Function f_5 in
1128 Figure 3] and the classified SHACAL-2 block cipher.

1129 Collision attacks have been thoroughly studied by the research community [SS08,
1130 MNS11]. The best attacks at this day, are second-order differential attack by Lamberger
1131 et al. [LM11] on the SHA-256 compression function reduced to 46 out of 64 rounds.

1132 Many researchers [IS09, AGM⁺09] have also studied preimage attacks on SHA-256
1133 with reduced rounds. Guo et al. [GLRW10] in particular were among the first to use
1134 the meet in the middle strategy [AS09] and achieved more efficient ones on 42-step
1135 SHA-256. Khovratovich et al. in 2012 [KRS12] have so far presented the best preimage
1136 attacks, on 45-round and 52-round SHA-256 as well as a 52-round attack on the SHA-256
1137 compression function.

1138 Li et al. have published in 2012 [LIS12] a noteworthy paper on converting meet in
1139 the middle preimage attack into pseudo collision attack. Using preimage attacks by
1140 bicliques, they found pseudo collisions attacks on 52 steps of SHA-256.

1141 **Claim 1.** SHA256 is 128-bit collision resistant.

1142 3.2 Instantiating MKHASH

1143 In this section we describe the instantiation of MKHASH with a compression function
1144 based on MIMC [AGR⁺16]. We firstly show how the compression function is constructed,
1145 and prove that this instantiation complies with the security requirements mentioned
1146 in Section 2.7

1147 3.2.1 MIMC Encryption

1148 MIMC is a block cipher with a simple design, consisting of a number of rounds (denoted
1149 *rounds*). During the i -th round, the message m is mixed with the encryption key k and a
1150 randomly chosen constant $c[i]$, and a permutation function is applied to generate a new
1151 value of m . The permutation function consists of exponentiation with a carefully chosen

1152 exponent e (see Section 3.2.1). Note that *rounds* depends on the desired security level
 1153 λ . We denote the encryption function by **MIMC-Encrypt** and illustrate it in Fig. 3.1.

```

MIMC-Encrypt( $k, m, c, e, rounds$ )
  1 : foreach  $i \in [rounds]$  :
  2 :    $m \leftarrow (k \text{ OP } c[i] \text{ OP } m)^e$ 
  3 : return ( $m \text{ OP } k$ )
  
```

Figure 3.1: MIMC Encryption function.

1154 **MIMC-Encrypt** can be defined on both binary and prime fields, and as such the **OP**
 1155 operation corresponds to either \oplus or $+$ (mod p) [AGR⁺16, GRR⁺16]. For general prime
 1156 p (resp. positive integer n), we denote by **MIMC p** (resp. **MIMC $_{2^n}$**) the **MIMC-Encrypt**
 1157 function defined over \mathbb{F}_p (resp. \mathbb{F}_{2^n}).

1158 Security parameters and analysis

1159 In this document, we only consider MIMC defined over prime fields (in particular, the
 1160 field $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$ over which **ZkSnarkSch** operates).

Since block ciphers are usually defined over the product space of keys and messages,
 we consider the variables c , *rounds* and e as fixed. We thereby consider an instantiation
 of MIMC with signature

$$\mathbf{MIMC}_{\mathbf{r}_{\text{CUR}}} : \mathbb{F}_{\mathbf{r}_{\text{CUR}}} \times \mathbb{F}_{\mathbf{r}_{\text{CUR}}} \rightarrow \mathbb{F}_{\mathbf{r}_{\text{CUR}}}$$

In the sections below, and as in [AGR⁺16], we will consider exponents of the form
 $e = 2^t - 1$ and $e = 2^t + 1$ where $\gcd(e, \mathbf{r}_{\text{CUR}} - 1) = 1$. We note that the term cancellation
 happening with exponents of the form $e = 2^t + 1$ does not immediately translate to
 the context where MIMC is carried out over prime fields of large odd characteristic. In
 fact, in the case of $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$, where $\mathbf{r}_{\text{CUR}} > \binom{e}{\lfloor e/2 \rfloor}$, polynomials $(x + y)^e$ are not sparse. This
 comes from the *Binomial Theorem*

$$(x + y)^e = \sum_{i=0}^e \binom{e}{i} x^i y^{e-i}$$

1161 and the observation that if $\binom{e}{i} < \mathbf{r}_{\text{CUR}}$ then $\binom{e}{i} \bmod \mathbf{r}_{\text{CUR}} = \binom{e}{i}$, hence ensuring that all
 1162 the polynomial coefficients are greater than 0, and that the polynomial is dense.

1163 To achieve a security of λ , we require that *rounds* $\geq \lambda \log_e(2)$. Importantly, since we
 1164 use MIMC over prime fields $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$ which are large (where \mathbf{r}_{CUR} is the prime characteristic
 1165 of the scalar field of an elliptic curve group, such that $\lceil \log_2(\mathbf{r}_{\text{CUR}}) \rceil > \lambda^2$), then, picking
 1166 *rounds* = $\left\lceil \frac{\log_2 \mathbf{r}_{\text{CUR}}}{\log_2 e} \right\rceil > \lambda \log_e(2)$ provides a margin of safety on the number of rounds
 1167 selected to instantiate MIMC with desired security level λ .

²Longer elements are needed in Elliptic Curve Cryptography (ECC) to resist algebraic attacks such
 as Number Field Sieve (NFS)-based attacks on discrete logs [Gor93] for instance

1168 We refer to the MIMC paper [AGR⁺16, Section 4.2 and 5.1] and to Appendix F for
 1169 more details on the security analysis and attacks on the scheme in the different settings.
 1170 Note that $\text{MIMC}_{\mathbf{r}_{\text{CUR}}}$ does not suffer from *inversion subfield attacks* as there are no proper
 1171 subfields of $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$.

1172 3.2.2 MIMC-based compression function

1173 There exist two main techniques to construct a hash function from a block-cipher (or
 1174 permutation): sponge functions [BDPVA07] and iterated compression functions [BRS02].

1175 A Merkle tree is a binary tree of values of fixed size, where the values in each “layer”
 1176 are generated by hashing pairs of values from the previous “layer”. That is, we require a
 1177 compression function MKHASH, which we construct via the Miyaguchi-Preneel scheme.
 1178 (Miyaguchi-Preneel is more secure [BRS02, f_5 function] than the more flexible Davies-
 1179 Meyer construct [GFBR06, Section 3], but this flexibility is not required in our case).

1180 Miyaguchi-Preneel compression construct

1181 Miyaguchi-Preneel (MP) [BRS02, f_3 function] is a general scheme for constructing com-
 1182 pression functions from block ciphers (see Section 1.5.6). Given a block cipher E , the
 1183 corresponding compression function by f_E^{MP} is given in Fig. 3.2. The original construc-
 1184 tion is defined over binary fields, however Z_{eth} operates over prime fields. Hence, in the
 1185 general discussion here we replace the bitwise addition operator \oplus by modular addition
 1186 in $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$ (see [Har19]).

1187 We denote by MIMC-MP the compression function defined by the application of
 1188 the Miyaguchi-Preneel construct over MIMC. Similarly, for general prime p we denote
 1189 by MIMC-MP_p (see Fig. 3.3) the compression function defined by application of the
 1190 Miyaguchi-Preneel construct over MIMC_p .

$$\begin{array}{l} \hline f_E^{\text{MP}}(k, m) \\ \hline 1 : \text{res} \leftarrow E_k(m) \\ 2 : \text{return } (\text{res} + m + k) \pmod{\mathbf{r}_{\text{CUR}}} \end{array}$$

Figure 3.2: MP construct in $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$.

$$\begin{array}{l} \hline \text{MIMC-MP}_{\mathbf{r}_{\text{CUR}}}(k, m) \\ \hline 1 : \text{res} \leftarrow \text{MIMC}_{\mathbf{r}_{\text{CUR}}}(k, m) \\ 2 : \text{return } (\text{res} + k + m) \pmod{\mathbf{r}_{\text{CUR}}} \end{array}$$

Figure 3.3: $\text{MIMC-MP}_{\mathbf{r}_{\text{CUR}}}$ construction.

1191 3.2.3 An efficient instantiation of MIMC primitives

1192 To select appropriate instances of $\text{MIMC}_{\mathbf{r}_{\text{CUR}}}$ and $\text{MIMC-MP}_{\mathbf{r}_{\text{CUR}}}$, we consider the cost
 1193 (in terms of gas consumption and prover efficiency). For given e and rounds , the final
 1194 definition of $\text{MIMC-MP}_{\mathbf{r}_{\text{CUR}}}$ is given in Fig. 3.4 and Fig. 3.5.

1195 **Remark 3.2.1.** Note that Keccak256 is the 256-bit digest instance of the Keccak family
 1196 that won the NIST SHA-3 competition [GJMG11]. It is supported by the EVM via an
 1197 opcode (see [W⁺, Appendix G]), making it convenient for use in smart contracts.

MIMC $r_{\text{CUR}}(k, m)$	InitRoundConstants ()
1 : $c \leftarrow \text{InitRoundConstants}()$ 2 : foreach $i \in [\text{rounds}]$: 3 : $m \leftarrow (k + c[i] + m)^e \pmod{r_{\text{CUR}}}$ 4 : return $(m + k) \pmod{r_{\text{CUR}}}$	$iv \leftarrow \text{Keccak256}(\text{"clearmatics_mt_seed"})$ $c[0] \leftarrow 0$ $c[1] \leftarrow \text{Keccak256}(iv)$ foreach $i \in \{2, \dots, \text{rounds}\}$: $c[i] \leftarrow \text{Keccak256}(c[i - 1])$ return $c = (c[0], \dots, c[\text{rounds} - 1])$

Figure 3.4: **MIMC** r_{CUR} full construction

MIMC-MP $r_{\text{CUR}}(k, m)$
return MIMC $r_{\text{CUR}}(k, m) + m + k \pmod{r_{\text{CUR}}}$

Figure 3.5: **MIMC-MP** r_{CUR} full construction

1198 **Remark 3.2.2.** To increase the security of the MKHASH, different round constants for
1199 each level of the Merkle tree could be used.

We define MKHASH to be **MIMC-MP** over $\mathbb{F}_{r_{\text{CUR}}}$. Thereby, for input values m_0 and m_1 , $\text{MKHASH} : \mathbb{F}_{r_{\text{CUR}}} \times \mathbb{F}_{r_{\text{CUR}}} \rightarrow \mathbb{F}_{r_{\text{CUR}}}$ is defined by

$$\text{MKHASH}(m_0, m_1) = \text{MIMC-MP}_{r_{\text{CUR}}}(m_0, m_1) \quad (3.1)$$

1200 For specific values of r_{CUR} (such as r_{BN} for BN-254 or r_{BLS} for BLS12-377), it remains
1201 to select concrete values of e and rounds , where $\text{rounds} = \lceil \frac{\log_2 r_{\text{CUR}}}{\log_2 e} \rceil$. These values
1202 influence the number of constraints in the arithmetic circuit (see Section 2.2 for details of
1203 the statement) and the gas cost of Merkle tree operations on the contract (see Section 2.5
1204 for details of the specific operations).

1205 In the arithmetic circuit, an invocation of **MIMC-MP** requires $\text{rounds} \cdot \text{mults}$ con-
1206 straints, where mults is the number of multiplications required for exponentiation. For
1207 exponents of the form $e = 2^t - 1$, we have $\text{mults} = 2 \cdot t - 2$, (using the *square-and-multiply*
1208 algorithm [MVOV96]), and for $e = 2^t + 1$ we have $\text{mults} = t + 1$. Thus we expect that
1209 exponents of the latter form are more optimal. The implementation in the contract
1210 performs a very similar set of arithmetic operations (exponentiation in the field through
1211 a series of multiplications and modulo reductions), and so the cost is dominated by the
1212 same number $\text{rounds} \cdot \text{mults}$ as for the circuit. Hence, values of e and rounds that are
1213 optimal for the circuit will also result in gas-efficient implementations in the contract.

1214 For several concrete values of e , the number of rounds required to attain the desired
1215 security level, along with the number of constraints, are shown in Table 3.1.

1216 For the case of BN-254 we set $e = 17$ with $\text{rounds} = 65$, to achieve a 254-bit security
1217 level. For BLS12-377 we set $e = 17$ with $\text{rounds} = 62$, to achieve 253-bit security. These
1218 values are chosen such that they satisfy the requirement that $\text{gcd}(e, r_{\text{CUR}} - 1) = 1$, and
1219 give a balance between the number of constraints in the arithmetic circuit and the gas
1220 cost of hashing on the contract.

e	BN-254		BLS12-377	
	<i>rounds</i>	<i>constraints</i>	<i>rounds</i>	<i>constraints</i>
5	110	331		
7	91	365		
17	65	316	62	311
31	52	417	51	409
127	37	445	37	445
257	32	289	32	289
511	29	465		
2047	24	481	23	461
8191	20	481	20	481
32676	17	477		
65537	16	273	16	273
131071	15	481	15	481
524287	14	505	14	505
1048577	13	274	13	274
2097151	13	521		

Table 3.1: Arithmetic constraints required to represent MIMC-MP as an R1CS program, for different exponents e and curves. Grey (resp. white) lines represent exponents of shape $2^t + 1$ (resp. $2^t - 1$). Missing entries where $\gcd(e, r_{\text{CUR}} - 1) \neq 1$

1221 3.2.4 Security requirements satisfaction

1222 After presenting the state of the art of MiMC cryptanalysis, we present the security
1223 proof of MIMC-MP collision resistance.

1224 Cryptanalysis of MIMC block cipher and primitives

1225 MIMC’s security is increasingly being analysed since the primitive has gained traction
1226 in zero-knowledge and cryptocurrency communities for its succinct algebraic constraint
1227 representation. As of today, we do not know of any attacks breaking MIMC on prime
1228 fields on full rounds.

1229 The first attack on MIMC was an interpolation attack [LP19] which targets a reduced-
1230 round version for a scenario in which the attacker has only limited memory. An attack
1231 on Feistel-based MIMC [Bon19] was discovered shortly after, by using generic properties
1232 of the used Feistel construction (instead of exploiting properties of the primitive itself).
1233 Additionally, [ACG⁺19] proposes an attack based on Gröbner basis. The authors state
1234 that by introducing a new intermediate variable in each round, the resulting multivariate
1235 system of equations is a Gröbner basis. As such, the first step of a Gröbner basis attack
1236 can be obtained for free. However, the following steps of the attack are so computation-
1237 ally demanding that the attack becomes infeasible in practice. A recent work [EGL⁺20]
1238 targets MIMC on binary fields, and achieves a full-round break of the scheme. While,
1239 the attack presented does not apply to prime fields, the authors note that it “can be
1240 generalized to include ciphers over \mathbb{F}_p ”, and that only the lack of efficient distinguishers
1241 over prime fields precludes this. Another attack from Beyne et al [BCD⁺20] uses a low
1242 complexity distinguisher against full MIMC permutation leading to a practical collision
1243 attack on reduced round sponge-based MIMC hash defined with security of 128 bits.

1244 Security proof of MIMC-MP collision resistance

1245 We now prove that this compression scheme satisfies all the security requirements listed
1246 in Section 2.7. To do so, we first assume that the round constants are pseudo-random,
1247 i.e. that Keccak256 is a PRF.

1248 **Lemma 3.2.1.** *Keccak256 is a PRF with $\lambda = 128$.*

1249 The security of MIMC-MP derives from a more general result, i.e. from modelling
1250 MIMC as an ideal cipher (see Definition 1.5.12). More specifically, we show a security
1251 result for the MP construction on $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$ by proving that, in the Ideal Cipher Model,
1252 the collision resistance advantage of any adversary is bounded by $\frac{q(q+1)}{\mathbf{r}_{\text{CUR}}}$, where q is the
1253 number of different queries that the attacker makes to the oracle. This means that,
1254 assuming a maximum q number of possible encryption/decryption queries, parameter
1255 \mathbf{r}_{CUR} can be chosen to make the advantage small as needed and $\mathbf{f}_{\text{E}}^{\text{MP}}$ considered collision
1256 resistant. Similar result applies to the 2^n case.

1257 The instance of MIMC we use is modelled as an ideal cipher defined on field elements,
1258 for this reason we consider a variant of the ICM model where the keys, inputs and outputs

1259 are field elements in $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$ and the block cipher scheme, with key k , correspond to a family
 1260 of \mathbf{r}_{CUR} independent random permutations $f_k : \mathbb{F}_{\mathbf{r}_{\text{CUR}}} \times \mathbb{F}_{\mathbf{r}_{\text{CUR}}} \rightarrow \mathbb{F}_{\mathbf{r}_{\text{CUR}}}$.

1261 In the proof, without loss of generality, we assume the following conventions for an
 1262 adversary \mathcal{A} :

- 1263 • the adversary asks distinct queries: i.e. if \mathcal{A} asks a query $\mathcal{O}^E(k, m)$ and this returns
 1264 y , then \mathcal{A} does not ask a subsequent query of $\mathcal{O}^E(k, m)$ or $\mathcal{O}^{E^{-1}}(k, y)$, and inversely;
- 1265 • the adversary necessarily obtained the candidate collision from the oracle. This
 1266 property follows suite from modelling MIMC as an ideal cipher.

1267 **Lemma 3.2.2.** *Let f_E^{MP} be the MP compression function built on an ideal block-cipher
 1268 E on $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$, the probability for an adversary \mathcal{A} to find a collision is not greater than
 1269 $q(q+1)/\mathbf{r}_{\text{CUR}}$ where q is a (positive) number of distinct oracle queries.*

1270 The following proof has been adapted from [BRS02, Lemma 3.3]³.

1271 *Proof.* Fix $h_0 \in \mathbb{F}_{\mathbf{r}_{\text{CUR}}}$. Let \mathcal{A} be an adversary attacking the compression function f_E^{MP} .
 1272 Assume that \mathcal{A} asks the oracles \mathcal{O}^E and $\mathcal{O}^{E^{-1}}$ a total of *distinct* q queries. Let us
 1273 denote the result of the q queries and output of the attacker (candidate collision) as
 1274 $((k_1, m_1, y_1), \dots, (k_q, m_q, y_q), \text{out})$. If \mathcal{A} is successful it means that it outputs (k, m) ,
 1275 (k', m') such that either $(k, m) \neq (k', m')$ and $f_E^{\text{MP}}(k, m) = f_E^{\text{MP}}(k', m')$ or $f_E^{\text{MP}}(k, m) =$
 1276 h_0 . By the definition of f_E^{MP} , we have that $E_k(m) + m + k = E_{k'}(m') + m' + k'$ for
 1277 the first case, or $E_k(m) + m + k = h_0$ for the second. So either there are distinct
 1278 $r, s \in [1, \dots, q]$ such that $(k_r, m_r, y_r) = (k, m, E_k(m))$ and $(k_s, m_s, y_s) = (k', m', E_{k'}(m'))$
 1279 and $E_{k_r}(m_r) + m_r + k_r = E_{k_s}(m_s) + m_s + k_s$ or else there is an $r \in [1, \dots, q]$ s.t. $(k_r, m_r, y_r) =$
 1280 (k, m, h_0) and $E_{k_r}(m_r) + m_r + k_r = h_0$. We show that this event is unlikely.

1281 In fact, for each $i \in [1, \dots, q]$, let C_i be the event that either $y_i + m_i + k_i = h_0$ or
 1282 does exist $j \in [1, \dots, i-1]$ s.t. $y_i + m_i + k_i = y_j + m_j + k_j$. When carrying out the
 1283 simulation y_i or m_i was randomly selected from a set of at least $\mathbf{r}_{\text{CUR}} - (i-1)$ elements,
 1284 so $\Pr[C_i] \leq i/(\mathbf{r}_{\text{CUR}} - i)$. This means that for the collision advantage of \mathcal{A} , $\text{Adv}_{f_E^{\text{MP}}, \mathcal{A}}^{\text{coll}}$ it
 1285 holds that $\text{Adv}_{f_E^{\text{MP}}, \mathcal{A}}^{\text{coll}} \leq \Pr[C_1 \vee \dots \vee C_q] \leq \sum_{i=1}^q \Pr[C_i]$. For $q \leq \frac{\mathbf{r}_{\text{CUR}}}{2}$ this probability is
 1286 bounded by $l \cdot \frac{q(q+1)}{\mathbf{r}_{\text{CUR}}}$. However, we allow only a polynomial number of queries, thus for
 1287 $q = \text{poly}(\lambda)$ this probability becomes $\frac{\text{poly}(\lambda)}{\mathbf{r}_{\text{CUR}}}$, where $\mathbf{r}_{\text{CUR}} \approx 2^\lambda$. \square

³It states the collision resistance of a set of compression functions f_1, \dots, f_{12} , denoted as *group-1 compression functions* and showed in [BRS02, Figure 3]. As mentioned above, Miyaguchi-Preneel corresponds to f_3 of that group. Since the proof of [BRS02, Lemma 3.3] shows collision resistance of f_1 , we slightly modified it to work for f_3 .

Note

Lemma 3.2.2 is applicable to our case by the strong assumption of $\text{MIMC}_{\mathbf{r}_{\text{CUR}}}$ being an ideal cipher. In other words, the proof does not take into account any structural weakness or knowledge that an attacker is aware of. Any such additional information could make Lemma 3.2.2 invalid, and consequently could be used to break the collision resistance.

1288

1289 **Remark 3.2.3.** Note that from Lemma 3.2.2 follows that the collision resistance security
1290 of the Zeth Merkle tree is $\log_2(\mathbf{r}_{\text{CUR}}/2)$ (around 127 bits for $\mathbf{r}_{\text{CUR}} = \mathbf{r}_{\text{BN}}$ or \mathbf{r}_{BLS}).

Note

MIMC has *not* received as much cryptanalytic scrutiny as other “older” and more established hash functions. This is important to note since, for these type of primitives which are not provably secure, the amount of attacks received by a scheme is a great indicator of its security and robustness. A natural alternative to MIMC here consists in using Pedersen hash which is provably collision resistant under the discrete-logarithm assumption.

1291

3.3 Zeth statement after primitive instantiation

1292

1293 After instantiating the various primitives and providing security proofs to justify that
1294 they comply with the security requirements listed in previous sections, \mathbf{R}^z now becomes:

1295

- For each $i \in [\text{JSIN}]$:

1296

1. $aux.jsins[i].znote.apk = \text{Blake2s}(tag_{ask}^{addr} \parallel \text{pad}_{\text{BLAKE2sCLEN}}(0))$
with tag_{ask}^{addr} defined in Section 3.1.3

1297

1298

2. $aux.jsins[i].nf = \text{Blake2s}(tag_{ask}^{nf} \parallel aux.jsins[i].znote.\rho)$
with tag_{ask}^{nf} defined in Section 3.1.3

1299

1300

3. $aux.jsins[i].cm = \text{Blake2s}(aux.jsins[i].znote.r \parallel m)$

1301

with $m = aux.jsins[i].znote.apk \parallel aux.jsins[i].znote.\rho \parallel aux.jsins[i].znote.v$

1302

4. $aux.htags[i] = \text{Blake2s}(tag_{ask,i}^{pk} \parallel prim.hsigs)$ (malleability fix, see Appendix A)
with $tag_{ask,i}^{pk}$ defined in Section 3.1.3

1303

1304

5. $(aux.jsins[i].znote.v) \cdot (1 - e) = 0$ is satisfied for the boolean value e set such that if $aux.jsins[i].znote.v > 0$ then $e = 1$.

1305

1306

6. The Merkle root $mkroot'$ used to check the Merkle authentication path $aux.jsins[i].mkpath$ of commitment $aux.jsins[i].cm$, with $\text{MIMC-MP}_{\mathbf{r}_{\text{CUR}}}$, equals $prim.mkroot$ if $e = 1$.

1307

1308

1309

7. $prim.nfs[i]$

1310

$$= \left\{ \text{Pack}_{\mathbb{F}_{\mathbf{r}_{\text{CUR}}}}(aux.jsins[i].nf[k \cdot \text{FIELD CAP} : (k + 1) \cdot \text{FIELD CAP}]) \right\}_{k \in \llbracket \text{PRFNFOUTLEN} / \text{FIELD CAP} \rrbracket}$$

- 1311 8. $prim.htags[i]$
1312 $= \{ \text{Pack}_{\mathbb{F}_{r_{\text{CUR}}}}(aux.htags[i][k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}]) \}_{k \in [\lfloor \text{PRFPKOUTLEN}/\text{FIELD CAP} \rfloor]}$
- 1313 • For each $j \in [\text{JSOUT}]$:
- 1314 1. $aux.znotes[j].\rho = \text{Blake2s}(tag_{ask,j}^\rho || prim.hsigs)$ (malleability fix, see Appendix A)
1315 with $tag_{ask,j}^\rho$ defined in Section 3.1.3
- 1316 2. $prim.cms[j] = \text{Blake2s}(aux.znotes[j].r || m)$
1317 with $m = aux.znotes[j].apk || aux.znotes[j].\rho || aux.znotes[j].v$
- 1318 • $prim.hsigs = \{ \text{Pack}_{\mathbb{F}_{r_{\text{CUR}}}}(aux.hsigs[k \cdot \text{FIELD CAP}:(k+1) \cdot \text{FIELD CAP}]) \}_{k \in [\lfloor \text{CRHHSIGOUTLEN}/\text{FIELD CAP} \rfloor]}$
- 1319 • $prim.rsd = \text{Pack}_{rsd}(\{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}, aux.vin, aux.vout, aux.hsigs, \{aux.htags[i]\}_{i \in [\text{JSIN}]})$
- Check that the “joinsplit is balanced”, i.e. check that the joinsplit equation holds:

$$\begin{aligned} & \text{Pack}_{\mathbb{F}_{r_{\text{CUR}}}}(aux.vin) + \sum_{i \in [\text{JSIN}]} \text{Pack}_{\mathbb{F}_{r_{\text{CUR}}}}(aux.jsins[i].znote.v) \\ &= \sum_{j \in [\text{JSOUT}]} \text{Pack}_{\mathbb{F}_{r_{\text{CUR}}}}(aux.znotes[j].v) + \text{Pack}_{\mathbb{F}_{r_{\text{CUR}}}}(aux.vout) \end{aligned}$$

1320 **Remark 3.3.1.** For higher security, we could use Blake2b with 32-byte output instead
1321 of SHA256. In fact, since a precompiled contract computing the Blake2 compression
1322 function [MJS15] has been added to the Istanbul release of **Ethereum** (EIP 152 [THH15]),
1323 it could be possible to write a small wrapper on the smart contracts, in order to hash
1324 with Blake2b with any parameter.

1325 3.3.1 Instantiating the packing functions

1326 As we consider SNARKs based on arithmetic circuits defined over a prime field, all
1327 variables in the constraint system are interpreted as field elements. Nevertheless, as
1328 illustrated in Section 2.2, part of the statement consists of functions whose co-domains
1329 are sets of binary strings (which may be longer than the bit representation of elements of
1330 the finite field). While a bit (i.e. $\{0, 1\}$) is an element of \mathbb{F}_p (p prime), it is important to
1331 minimize the number of gates in the arithmetic circuit (for proof generation efficiency),
1332 and to minimize the number of input wires (to improve verification time). This can be
1333 done by representing fragments of binary strings as the base 2 decomposition of field
1334 elements, thereby “packing” binary strings into multiple elements. Converting binary
1335 strings into field elements requires the addition of some arithmetic gates (extending
1336 the statement to be proven), but reduces the number of primary inputs (reducing the
1337 complexity of the SNARK verification carried out on-chain). The cost of Groth16 zk-
1338 SNARK [Gro16] proof verification is linear in the number of primary inputs, since each
1339 input acts as a scalar in a costly scalar multiplication of a curve point in \mathbb{G}_1 . Hence,
1340 while packing slightly increases the prover cost – by adding constraints to the circuit –
1341 it simplifies the verifier’s work.

1342 In this section, we detail the method by which we encode (resp. decode) a set of
 1343 binary strings to (resp. from) sets of field elements. In the rest of this section, the notion
 1344 of *packing policy* refers to the set of *packing* and *unpacking* functions.

The set of primary inputs is composed of the input nullifiers, the output commitments, the public values (see [RZ19, Section 3.4.3]) along with the signature hash and the authentication tags for security (malleability fix, see Appendix A). The complete description of the public inputs is represented in Eq. (3.2).

$$(\{prim.nf_i\}_{i \in [JSIN]}, \{prim.cms[j]\}_{j \in [JSOUT]}, vin, vout, hsig, \{prim.htags[i]\}_{i \in [JSIN]}) \quad (3.2)$$

1345 The primary inputs that consist of binary strings are: the nullifiers *nfs*, the public
 1346 values *vin* and *vout*, the signature hash *hsig* and the authentication tags *htags*.

1347 For a binary string x , let $\alpha_x = \lceil \text{length}(x) / \text{FIELD CAP} \rceil$ be the number of field elements
 1348 required to completely encode x and let $\beta_x = \lfloor \text{length}(x) / \text{FIELD CAP} \rfloor$ be the number of
 1349 field elements whose capacity is fully used. Let $\gamma_x = \text{length}(x) \pmod{\text{FIELD CAP}}$ be the
 1350 number of “residual” bits remaining after fully using β_x field elements.

1351 **Example 3.3.2.** Consider binary strings $A \in \{0, 1\}^7$ of length 7, to be encoded over the
 1352 field \mathbb{F}_{41} . This field has a capacity of 5 bits, and therefore $\alpha_A = 2$, $\beta_A = 1$, and $\gamma_A = 2$.
 1353 That is, A can be represented as 2 field elements, or as 1 field element with 2 “residual”
 1354 bits.

1355 Consider $A = (1111011)$. Fig. 3.6 illustrates how A can be packed as field elements.
 1356 Note that the 2 residual bits are taken from the “beginning” of the bit string, that is,
 1357 the highest order bits.

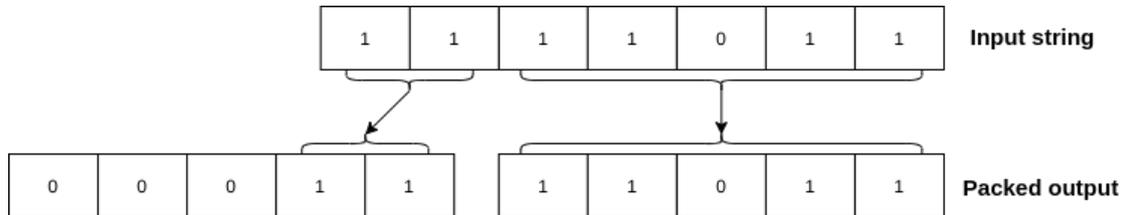


Figure 3.6: Packing of string A (see Example 3.3.2)

1358 We now consider strategies to pack all primary inputs that are binary strings. A naive
 1359 approach is to encode each binary string x as α_x field elements. In general, this results
 1360 in significant waste (and consequently more field elements than necessary), especially
 1361 when the number of residual bits is small compared to **FIELD CAP** (see Fig. 3.7). An
 1362 alternative strategy could be to concatenate all binary strings into a single string y and
 1363 pack this string into α_y field elements. While this approach minimizes the set of unused
 1364 bits, each unpack operation would require different shift and mask operations over 2 or
 1365 3 field elements. This significantly increases the complexity of the unpacking operation
 1366 that must be performed on-chain, resulting in a higher gas cost (due to extra logic) or
 1367 more contract code (if each unpack operation is hard-coded).

The `Zeth` protocol requires that each binary string variable x is packed into β_x field elements, and the residual bits from all binary strings, along with the public values vin and $vout$, are aggregated into a variable rsd . Let `RSDBLEN` be the total number of residual bits, and `RSDFLEN` be the number of field elements required to represent rsd . We assume that `ZVALUELEN` $<$ `FIELD CAP`, and define the notation $\gamma_v = \text{ZVALUELEN}$ for the bit lengths of public values vin and $vout$. Thus `RSDBLEN` is given by

$$\text{RSDBLEN} = \gamma_{hsig} + 2 \cdot \gamma_v + \text{JSIN} \cdot (\gamma_{nf} + \gamma_h)$$

and the lengths, in field elements, of each of the corresponding public inputs are

$$\begin{aligned} \text{NFFLEN} &= \beta_{nf} \\ \text{HSIGFLEN} &= \beta_{hsig} \\ \text{HFLEN} &= \beta_h \\ \text{RSDFLEN} &= \lceil \text{RSDBLEN} / \text{FIELD CAP} \rceil \end{aligned}$$

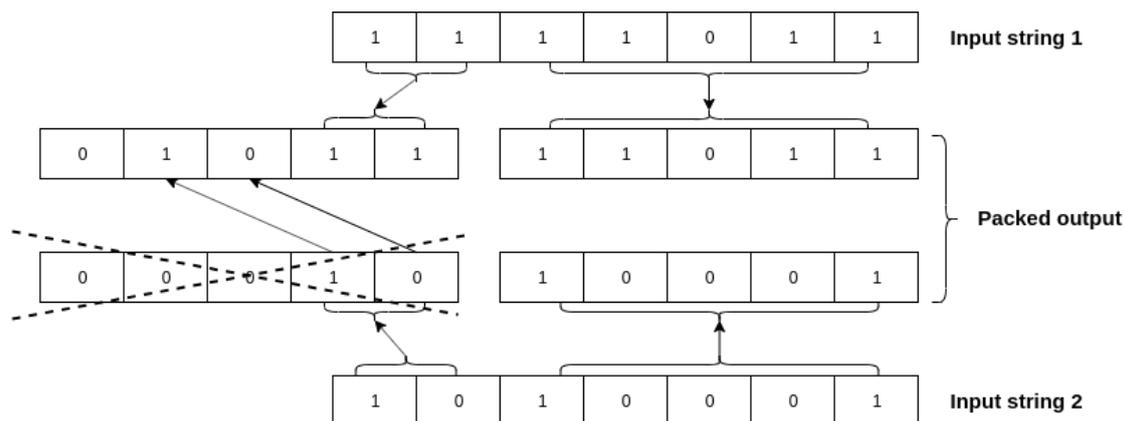


Figure 3.7: Packing of multiple strings. Observe that, by carefully arranging the bits of the input strings, it is possible to output fewer field elements

The residual bits rsd are formatted as follows:

$$\widetilde{hsig} \parallel \widetilde{nfs} \parallel \widetilde{htags} \parallel vin \parallel vout$$

1368 where \widetilde{hsig} , \widetilde{nfs} , \widetilde{htags} are, respectively, the γ_{hsig} , γ_{nf} , γ_h bits.

1369 Note that the public values are packed into the “last”, or lowest order, $2 \cdot \gamma_v$ bits of
 1370 the resulting field element(s). In this way, their unpack functions are independent of the
 1371 values `JSIN` and `JSOUT` and of the number of residual bits required for each bit string
 1372 (and consequently, independent of the finite field used).

To format the unpacked primary inputs into field elements, we define the following functions. Given a bit string of length less than `FIELD CAP`, the algorithm `Pack` (see Fig. 3.8) returns a field element. Given the nullifiers, public values and authentication tags, the algorithm `Packrsd` (see Fig. 3.9) outputs the residual bits. Given a set of

$\text{Pack}_{\mathbb{F}_{\text{rCUR}}}(x)$	$\text{Pack}_{rsd}(nfs, vin, vout, hsig, htags)$
$out \leftarrow 0_{\mathbb{F}_{\text{rCUR}}};$ for $i \in [\text{length}(x)]$ do : if $x[i] = 1$ do : $out \leftarrow out +_{\mathbb{F}_{\text{rCUR}}} 2^{\text{length}(x)-1-i}$ return out ;	$out \leftarrow []; r \leftarrow \epsilon;$ $r \leftarrow vout;$ $r \leftarrow vin r;$ for $i \in [\text{JSIN}]$ do : $r \leftarrow htags[i][\beta_{htags[i]} \cdot \text{FIELD CAP}] r;$ for $i \in [\text{JSIN}]$ do : $r \leftarrow nfs[i][\beta_{nfs[i]} \cdot \text{FIELD CAP}] r;$ $r \leftarrow hsig[\beta_{hsig} \cdot \text{FIELD CAP}] r;$ for $i \in [[\text{length}(r)/\text{FIELD CAP}]]$ do : $out[i] \leftarrow \text{Pack}_{\mathbb{F}_{\text{rBN}}}(r[i \cdot \text{FIELD CAP}:(i+1) \cdot \text{FIELD CAP}]);$ return out ;

Figure 3.8: Algorithm to pack bits into a field element.

Figure 3.9: Algorithm to pack residual bits.

packed field elements and the residual bits, the algorithm `Unpack` returns the variables reassembled as binary strings. In particular, we have that $\text{Unpack}_{nf}(\text{prim}.nfs, rsd) = \{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}$.

$$\begin{aligned}
\text{Pack} &: \mathbb{B}^{\leq \text{FIELD CAP}} \rightarrow \mathbb{F}_{\text{rCUR}} \\
\text{Pack}_{rsd} &: (\mathbb{B}^{\text{PRFNFOUTLEN}})^{\text{JSIN}} \times (\mathbb{B}^{\text{ZVALUELEN}})^2 \times \mathbb{B}^{\text{CRHHSIGOUTLEN}} \times (\mathbb{B}^{\text{PRFPKOUTLEN}})^{\text{JSIN}} \rightarrow (\mathbb{F}_{\text{rCUR}})^{\text{RSDFLEN}} \\
\text{Unpack} &: \mathbb{F}_{\text{rCUR}}^* \times (\mathbb{F}_{\text{rCUR}})^{\text{RSDFLEN}} \rightarrow \mathbb{B}^*
\end{aligned}$$

The `Unpack` functions for nullifiers, public values and signature hash are represented as follows.

$$\begin{aligned}
\text{Unpack}_{hsig} &: (\mathbb{F}_{\text{rCUR}})^{\text{HSIGFLEN}} \times (\mathbb{F}_{\text{rCUR}})^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{CRHHSIGOUTLEN}} \\
\text{Unpack}_{nf} &: (\mathbb{F}_{\text{rCUR}})^{\text{NFLEN}} \times (\mathbb{F}_{\text{rCUR}})^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{PRFNFOUTLEN}} \\
\text{Unpack}_{vin} &: \mathbb{F}_{\text{rCUR}}^0 \times (\mathbb{F}_{\text{rCUR}})^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{ZVALUELEN}} \\
\text{Unpack}_{vout} &: \mathbb{F}_{\text{rCUR}}^0 \times (\mathbb{F}_{\text{rCUR}})^{\text{RSDFLEN}} \rightarrow \mathbb{B}^{\text{ZVALUELEN}}
\end{aligned}$$

1373 Packing Policy Security

1374 **Proposition 3.3.1** (Packing security). *For a binary string x , it holds that $\text{Unpack}(\text{Pack}(x)) =$*
1375 *x and $\text{Unpack}(\text{Pack}_{rsd}(x)) = x$.*

1376 Packing Policy Example

In the case where $\text{JSIN} = \text{JSOUT} = 2$, the BN-254 is being used (in which field elements hold $\text{FIELD CAP}_{\text{BN}}$ bits) and all PRFs and CRH^{hsig} output bit-strings of length 256, the

unpacked primary inputs are 2167-bit long. The packing parameters are therefore:

$$\begin{aligned} \text{RSDBLEN} &= 5 \times 3 + 64 + 64 = 143 \\ \text{NFFLEN} &= \text{HSIGFLEN} = \text{HFLEN} = \text{RSDFLLEN} = 1 \end{aligned}$$

The packed primary inputs are 2277 bits long, corresponding to a small space overhead of $\approx 5\%$ unused bits. Moreover, the 143-bit residual bits can be packed into a single field element. As such, the primary inputs are encoded as 9 field elements. Finally, the residual bits are formatted as follows,

$$\underbrace{\text{padding}}_{113 \text{ bits}} \parallel \underbrace{\text{hsig}}_{3 \text{ bits}} \parallel \underbrace{\text{nf}_1}_{3 \text{ bits}} \parallel \underbrace{\text{nf}_0}_{3 \text{ bits}} \parallel \underbrace{h_1}_{3 \text{ bits}} \parallel \underbrace{h_0}_{3 \text{ bits}} \parallel \underbrace{\text{vin}}_{64 \text{ bits}} \parallel \underbrace{\text{vout}}_{64 \text{ bits}}$$

For the analogous case using BLS12-377 (in which field elements hold $\text{FIELD CAP}_{\text{BLS}}$ bits), the packing parameters are:

$$\begin{aligned} \text{RSDBLEN} &= 5 \times 4 + 64 + 64 = 148 \\ \text{NFFLEN} &= \text{HSIGFLEN} = \text{HFLEN} = \text{RSDFLLEN} = 1 \end{aligned}$$

The residual bits can be packed into a single field element of the form

$$\underbrace{\text{padding}}_{108 \text{ bits}} \parallel \underbrace{\text{hsig}}_{4 \text{ bits}} \parallel \underbrace{\text{nf}_0}_{4 \text{ bits}} \parallel \underbrace{\text{nf}_1}_{4 \text{ bits}} \parallel \underbrace{h_0}_{4 \text{ bits}} \parallel \underbrace{h_1}_{4 \text{ bits}} \parallel \underbrace{\text{vin}}_{64 \text{ bits}} \parallel \underbrace{\text{vout}}_{64 \text{ bits}}$$

1377 and the primary inputs are again encoded as 9 field elements.

1378 3.4 Instantiate SigSch_{OT-SIG}

1379 **Zeth** uses the one-time Schnorr-based signature scheme introduced by Bellare and Shoup [BS07]
 1380 for its long proven security, simplicity, speed and size. Its security relies on the one-more
 1381 discrete log problem (see Definition 1.5.6) and the collision resistance of the underlying
 1382 hash function CRH (see Definition 1.5.16) that we instantiate with SHA256.

1383 Note that no signature operations or data are used in the arithmetic circuit describing
 1384 the **Zeth** statement. Hence the curve used for the signature scheme can be chosen
 1385 independently of **Curve** (the scalar field of which is used for the arithmetic circuit, and
 1386 consequently for commitments and bit string encodings described in Section 3.1 and
 1387 Section 3.2). BN-254 is used since it is supported by the EVM, in the form of precompiled
 1388 contracts. This allows a gas-efficient implementation in the **Mixer** contract.

1389 This one-time signature scheme (see Definition 1.5.26) is defined by the two-tier
 1390 signature scheme over a cyclic group $(p, \mathbb{G}, \langle \mathbf{g} \rangle, \otimes)$. In the two-tier signature scheme, the
 1391 hash function CRH only needs to be collision resistant (the random oracle model is not
 1392 used). Similarly, the variable hk represents the key of the hash function (a particular
 1393 instance).

1394 To turn this two-tier signature scheme into a one-time signature scheme, one simply
 1395 has to define the one-time signature key generation KGen as the combination of both

1396 primary and secondary key generations of the two-tier (see [BS07, Section 6]). The
 1397 one-time signing key (respectively verification key) of the one time signature scheme is
 1398 defined as both the primary and secondary signing key (respectively verification key) of
 1399 the two-tier scheme, Fig. 3.10

$\underline{\text{KGen}(1^\lambda)} :$ $hk \leftarrow_{\$} \mathbb{B}^{kl}$ $\mathbf{g} \leftarrow_{\$} \mathbb{G}^*$ $x \leftarrow_{\$} \mathbb{F}_p$ $ppk = (hk, \mathbf{g}, \llbracket x \rrbracket)$ $psk = (hk, \mathbf{g}, x)$ $y \leftarrow_{\$} \mathbb{F}_p$ $spk = \llbracket y \rrbracket$ $ssk = (y, \llbracket y \rrbracket)$ $pk = (ppk, spk)$ $sk = (psk, ssk)$	$\underline{\text{Sig}(sk, m)} :$ $hk, \mathbf{g}, x = sk.psk$ $y, \llbracket y \rrbracket = sk.ssk$ $c = \text{CRH}(hk, \llbracket y \rrbracket m)$ $\sigma = y \bmod p$ $\sigma += c \cdot x \bmod p$ $\mathbf{return} \sigma$	$\underline{\text{Vf}(pk, m, \sigma)} :$ $hk, \mathbf{g}, \llbracket x \rrbracket = pk.ppk$ $\llbracket y \rrbracket = pk.spk$ $c = \text{CRH}(hk, \llbracket y \rrbracket m)$ $\mathbf{if} \sigma = \llbracket y \rrbracket \otimes c \cdot \llbracket x \rrbracket \mathbf{then}$ $\quad \mathbf{return} 1$ \mathbf{else} $\quad \mathbf{return} 0$ \mathbf{endif}
--	---	---

Figure 3.10: One-time signature scheme from two tier Schnorr based signature scheme by Bellare and Shoup [BS07]

1400 3.4.1 Security requirements satisfaction

1401 We now prove that this signature scheme satisfies all the security requirements listed
 1402 in Section 2.7.

1403 **Theorem 3.4.1.** *The One-Time Schnorr signature is strongly unforgeable under chosen-*
 1404 *message attacks (SUF-CMA) assuming that the om-DLog problem is hard in \mathbb{G} and that*
 1405 *the hash function CRH is collision resistant.*

1406 *Proof.* See [BS07, Theorems 5.1, 5.2 and 6.1]. □

1407 3.4.2 Data types

1408 We now describe the data types and operations associated with this signature scheme.

1409 **VK0tsDType** Denotes the verification key associated with the one-time signature scheme.

Field	Description	Data type
ppk	Encoding of the scalar x in the group	\mathbb{G}_{BN}
spk	Encoding of the scalar y in the group	\mathbb{G}_{BN}

Table 3.2: VK0tsDType data type

1410 **SK0tsDType** Denotes the signing key associated with the one-time signature scheme.

Field	Description	Data type
psk	Scalar element x	\mathbb{F}_{rBN}
ssk	Scalar element y and its encoding in the group	$\mathbb{F}_{\text{rBN}} \times \mathbb{G}_{\text{BN}}$

Table 3.3: SK0tsDType data type

1411 **Sig0tsDType** Denotes the signature data type associated with the one-time signature
1412 scheme. **Sig0tsDType** is an alias for \mathbb{F}_{rBN} .

1413 3.5 Instantiate EncSch

1414 In this section we describe the instantiation of **EncSch** primitive introduced in Section 2.3.
1415 First, we present a general asymmetric encryption scheme called **DHAES** (Diffie-Hellman
1416 Asymmetric Encryption Scheme [ABR99]), which satisfies all the required security prop-
1417 erties for the in-band encryption scheme **EncSch** (see Section 1.5.3). Then, we give details
1418 of the concrete algorithms used for the implementation.

1419 3.5.1 DHAES encryption scheme

1420 Given a symmetric encryption scheme **Sym**, a group defined by **SetupG**, a family of hash
1421 function \mathcal{H}^4 and a message authentication scheme **MAC** as defined in Section 1.5, we
1422 define a **DHAES** scheme as the following public-key encryption scheme:

- 1423 • **Setup**, setup algorithm, takes as input a security parameter 1^λ . It runs $\mathcal{H}.\text{Setup}$,
1424 SetupG and returns public parameters $pp = (hk, (q, \mathbb{G}, \mathfrak{g}, +))$.
- 1425 • **KGen**, key generation algorithm, takes as input public parameters pp . It samples
1426 at random $v \leftarrow_{\$} [q]$ and returns a keypair $(sk, pk) = (v, \llbracket v \rrbracket)$.
- 1427 • **Enc**, encryption algorithm, takes as input public parameters pp , a message m and
1428 a public key pk . It runs **KGen** that returns an ephemeral keypair $(esk, epk) =$
1429 $(u, \llbracket u \rrbracket)$. Then, it computes a shared secret $ss = \text{H}_{hk}(epk \parallel esk \cdot pk) = \text{H}_{hk}(epk \parallel sk \cdot$

⁴Here, we only consider fixed-length hash functions with $h\text{InpLen}(\lambda) = 2g\text{Len}$ and $h\text{Len}(\lambda) = k\text{Len}(\lambda) + m\text{Len}(\lambda)$ (see Section 1.5).

1430 epk), parsed as $ek\|mk$ ⁵. It computes $ct_{\text{Sym}} = \text{Sym.Enc}(ek, m)$ and $\tau = \text{MAC.Tag}(mk, ct_{\text{Sym}})$
1431 and finally outputs the ciphertext $epk\|ct_{\text{Sym}}\|\tau$.

1432 • Dec, decryption algorithm, takes as input public parameters pp , a private key sk
1433 a ciphertext $epk\|ct_{\text{Sym}}\|\tau$. It computes $ss = H_{hk}(epk\|sk \cdot epk)$ and parses it, as
1434 above, as $ek\|mk$. If MAC verification passes, i.e. $\text{MAC.Vf}(mk, \tau) = 1$, the algorithm
1435 returns $\text{Sym.Dec}(ek, ct_{\text{Sym}})$ and \perp otherwise.

1436 The DHAES definition given above is an asymptotic adaptation of [ABR99, Section
1437 1.3].

1438 Inclusion of ephemeral key in hash input

1439 Given an ephemeral keypair $(u_0, \llbracket u_0 \rrbracket)$, If the group $\langle \mathfrak{g} \rangle$, generated by SetupG , has com-
1440 posite order, then $\llbracket u_0 \rrbracket$ is required to be part of the hash input because $\llbracket u_0 v \rrbracket$ and $\llbracket v \rrbracket$
1441 together may not uniquely determine $\llbracket u_0 \rrbracket$. Equivalently, there may exist two values
1442 u_0 and u_1 such that $u_0 \neq u_1$ and $\llbracket u_0 v \rrbracket = \llbracket u_1 v \rrbracket$. As a result, both u_0 and u_1 can be
1443 used to produce two different *valid* ciphertexts of the same plaintext m , under different
1444 ephemeral keys $(\llbracket u_0 \rrbracket, \llbracket u_1 \rrbracket)$. It is easy to show this, for example, in the multiplicative
1445 group $\mathbb{Z}_p \setminus \{0\}$, where p is a prime (see [ABR99, Section 3.1]). A scheme having such
1446 malleability property clearly cannot be proven IND-CCA2 secure: an attacker could eas-
1447 ily win the related security game by altering the challenged ciphertext and query the
1448 decryption oracle that would not recognize that as a not allowed query. If the group
1449 has prime order this problem does not arise so only $\llbracket u_0 v \rrbracket$ is required as input of the H
1450 function [ABR01, Section 3].

1451 3.5.2 A DHAES instance

1452 Curve25519

1453 For a cyclic group we propose the use of a subgroup of Curve25519 described in [Ber06]
1454 and in [LHT16]. Curve25519 is a Montgomery elliptic curve [Mon87] defined by the
1455 equation $y^2 = x^3 + 486662x^2 + x$ and coordinates on \mathbb{F}_p , where p is the prime number
1456 $2^{255} - 19$. It has a prime order subgroup of order $2^{252} + 27742317777372353535851937$
1457 790883648493 and cofactor 8. Curve25519 comes with an efficient scalar multiplication
1458 denoted as X25519⁶. In a Diffie-Hellman-based scheme it allows to have 32-byte long
1459 public and private keys (given a point $P = (x, y)$ only the x coordinate is actually used)
1460 and the 32-byte sequence representing 9 is specified as base point.

1461 Efficiency and security of Curve25519

1462 High-speed and timing-attack resistant implementations of X25519 are available and
1463 its security level is conjectured to be 128 bits [Ber06, Section 1]. However, combined

⁵Note that ek and mk must have the same length.

⁶X25519 is actually introduced in [LHT16] in order to avoid notation issues due to the use Curve25519 to indicate both curve and scalar multiplication as done in [Ber06]

1464 attacks can lead to 124 bits of security (see [BL, Section “Twist Security”]). By design,
1465 Curve25519 is resistant to state-of-the-art attacks and satisfies all security criteria and
1466 principles listed in *Safecurves* [BL]⁷.

1467 Interestingly, Curve25519 does not require *public key validation*⁸, while we know that,
1468 on other curves, active attacks – consisting of sending malformed public keys – could be
1469 carried out by adversaries, to violate the confidentiality of private keys, e.g. [ABM⁺03].
1470 However, Curve25519 specification mandates the *clamping* of private keys: that is, after
1471 the random sampling of 32 bytes, the user clears bits 0, 1 and 2 of the first byte, clears
1472 bit 7 and sets bit 6 of the last byte. The resulting 32 bytes are then used as private key.
1473 This particular structure for private keys prevents various types of attacks (see [Ber06,
1474 Section 3] for more details).

Note

Note that the *clamping* procedure is vital to ensure the security guarantees of the Curve25519 specification, and implementations **MUST** perform this exactly as described.

1475

1476 Chacha20

1477 ChaCha20 is an ARX-based⁹ stream cipher introduced in [Ber08a]. It is an improved
1478 version of Salsa20 [Ber08b] that won the *eSTREAM* challenge [est]. Compared with
1479 Salsa20, it has been designed to improve diffusion per round, conjecturally increasing
1480 resistance to cryptanalysis, while preserving time efficiency per round. It is considerably
1481 faster than AES in software-only implementations and can be easily implemented to be
1482 timing-attacks resistant. Several versions of the cipher can be used. The original paper
1483 presents ChaCha20 with a 128-bit key and 64-bit nonce/block count. However, the length
1484 of the key, nonce and block count – which indicates how many chunks can be processed
1485 by using the same key and nonce – can be modified depending on the application.
1486 In [LN18][Section 2.3], for instance, the key is a 256-bit string, the nonce is a string of
1487 96 bits and the block count is encoded on a 32-bit word. This configuration allows to
1488 process around 2^{32} blocks, corresponding to roughly 256 GB of data. We propose to use
1489 the same parameters in *Zeth*.

$$\text{ChaCha20} : \mathbb{B}^{256} \times \mathbb{B}^{32} \times \mathbb{B}^{96} \times \mathbb{B}^* \rightarrow \mathbb{B}^*$$

⁷In this work, the authors take into account both Elliptic Curve Discrete Logarithm Problem (ECDLP) and Elliptic Curve Cryptosystems (ECC) security, that allows to have an overall evaluation of the security guarantees.

⁸Informally, it is a set of security checks that a user performs before using a not trusted public key (e.g. see [BCK⁺18])

⁹Addition-Rotation-XOR

1490 **Security of Chacha**

1491 Recent cryptanalysis results for ChaCha are available in [AFK⁺08, Ish12, SZFW12,
1492 Mai16, CM16, CM17]: all of them make use of advanced cryptanalysis techniques able
1493 to perform key-recovery attacks only on reduced versions (6 and 7 rounds) of ChaCha.

Note

Importantly, the security properties of ChaCha rely on the fact that, for a given key, all blocks are processed with distinct values in the state words 12 to 15 (storing the counter and the nonce) [LN18, Section 2.3].

1494

1495 **Poly1305**

1496 Poly1305 [Ber05] is a high-speed message authentication code, easy to implement and
1497 make side-channel attack resistant. It takes a 32-byte one-time key mk and a message m
1498 and produces a 16-byte tag τ that authenticates the message. mk must be unpredictable
1499 and it is represented as a couple (r, s) , where both components are given as a sequence
1500 of 16 bytes each. It can be generated by using pseudorandom algorithms: in [Ber05,
1501 Section 2], for example, AES and a nonce are used to generate s . The second part of
1502 the key, r , is expected to have a given form [Ber05, Section 2], and must be “clamped”
1503 as follows: top four bits of $r[3]$, $r[7]$, $r[11]$, $r[15]$ and bottom two bits of $r[4]$, $r[8]$, $r[12]$
1504 are cleared (see also Section 3.5.3).

Note

Similarly to Curve25519, the *clamping* procedure here is essential to the security of the Poly1305 scheme. Implementations MUST ensure that this is performed correctly in order for all security guarantees to hold.

1505

1506 We refer to [LN18, Section 2.5, Section 3] for Tag and Vf implementations of Poly1305.

$$\text{Poly1305.Tag} : \mathbb{B}_Y^{32} \times \mathbb{B}_Y^* \rightarrow \mathbb{B}_Y^{16}$$

$$\text{Poly1305.Vf} : \mathbb{B}_Y^{32} \times \mathbb{B}_Y^{16} \times \mathbb{B}_Y^* \rightarrow \mathbb{B}$$

1507 **Security of Poly1305**

1508 Citing Poly1305 [LN18, Section 4], “the Poly1305 authenticator is designed to ensure that
1509 forged messages are rejected with a probability of $1 - (n/(2^{102}))$ for a $16n$ -byte message,
1510 even after sending 2^{64} legitimate messages, so it is SUF-CMA (strong unforgeability
1511 against chosen-message attacks)”.

1512 **Blake2b-512**

1513 Since we need a total of 64 bytes for the key material (32 for ChaCha20 and 32 for
 1514 Poly1305) Blake2b512 can be used. ZCash protocol [ZCa19, Section 5.4.3], instead, makes
 1515 use of Blake2b256 since a DHAES variant, denoted as ChaCha20-Poly1305, is adopted
 1516 (see [LN18, Section 2.8]).

$$\text{Blake2b512} : \mathbb{B}^* \rightarrow \mathbb{B}_{\mathbb{Y}}^{32}$$

1517 **3.5.3 EncSch instantiation**

In the following we instantiate EncSch as a DHAES scheme, detailing the KGen, Enc and Dec components. First, we introduce some required constant values:

$$\begin{aligned} \text{ESKBYTELEN} &= 32 \\ \text{EPKBYTELEN} &= 32 \\ \text{NOTEBYTELEN} &= (\text{PRFADDRROUTLEN} + \text{RTRAPLEN} + \text{ZVALUELEN} + \text{PRFRHOOUTLEN})/\text{BYTELEN} \\ \text{SYMKEYBYTELEN} &= 32 \\ \text{MACKEYBYTELEN} &= 32 \\ \text{KDFDIGESTBYTELEN} &= \text{SYMKEYBYTELEN} + \text{MACKEYBYTELEN} \\ \text{CTBYTELEN} &= \text{EPKBYTELEN} + \text{NOTEBYTELEN} + \text{TAGBYTELEN} \\ \text{TAGBYTELEN} &= 16 \\ \text{CHACHANONCEVALUE} &= 0^{32} \\ \text{CHACHABLOCKCOUNTERVALUE} &= 0^{96} \end{aligned}$$

1518 **EncSch.KGen**

1519 The keypair (sk, pk) generation is defined as:

- 1520
- Randomly sample a sequence of ESKBYTELEN bytes and assign to sk .
 - Clamp sk as follows:

$$\begin{aligned} sk[0] &\leftarrow sk[0] \& 0xF8 \\ sk[31] &\leftarrow sk[31] \& 0x7F \\ sk[31] &\leftarrow sk[31] | 0x40 \end{aligned}$$

1521 where $|$ and $\&$ denotes, respectively, OR and AND binary operators between bit
 1522 strings of same the length.¹⁰

- 1523
- Compute $pk = \text{X25519}(sk, 0x09)$.
 - Return $(sk, pk) \in \mathbb{B}_{\mathbb{Y}}^{\text{ESKBYTELEN}} \times \mathbb{B}_{\mathbb{Y}}^{\text{EPKBYTELEN}}$
- 1524

¹⁰E.g Given two bytes 0x15 and 0x03 then 0x15|0x03 = 0x17 and 0x15&0x03 = 0x01.

1525 EncSch.Enc

1526 The encryption, on inputs $(pk, m) \in \mathbb{B}_{\mathbb{Y}}^{\text{EPKBYTELEN}} \times \mathbb{B}_{\mathbb{Y}}^{\text{NOTEBYTELEN}}$, is defined as follows:

1527 1. Generate an ephemeral Curve25519 keypair $(esk, epk) \in \mathbb{B}_{\mathbb{Y}}^{\text{ESKBYTELEN}} \times \mathbb{B}_{\mathbb{Y}}^{\text{EPKBYTELEN}}$
1528 (as above).

2. Compute the shared secret¹¹ $ss \in \mathbb{B}_{\mathbb{Y}}^{\text{EPKBYTELEN}}$:

$$ss = \text{X25519}(esk, pk) \in \mathbb{B}_{\mathbb{Y}}^{\text{EPKBYTELEN}}$$

3. Generate a session key:

$$\text{Blake2b512}(\text{encTag} \| epk \| ss) \in \mathbb{B}_{\mathbb{Y}}^{\text{KDFDIGESTBYTELEN}}$$

where $\text{encTag} = 0x5A \| 0x65 \| 0x74 \| 0x68 \| 0x45 \| 0x6E \| 0x63$, that is the UTF-8 encoding of “ZethEnc” string (used for domain separation purposes). The result, then, is parsed as follows:

$$ek = \text{Blake2b512}(\text{encTag} \| epk \| ss)[\text{SYMKEYBYTELEN} - 1]$$

$$mk = \text{Blake2b512}(\text{encTag} \| epk \| ss)[\text{SYMKEYBYTELEN} : \text{SYMKEYBYTELEN} + \text{MACKEYBYTELEN} - 1].$$

4. Encrypt the confidential data:

$$ct_{\text{sym}} = \text{ChaCha20}(ek, \text{CHACHABLOCKCOUNTERVALUE}, \text{CHACHANONCEVALUE}, m) \in \mathbb{B}^{\text{NOTEBYTELEN} * \text{BYTELEN}}$$

1529 **Remark 3.5.1.** Formally speaking we should have written $ct_{\text{sym}} \in \mathbb{B}^n$, where
1530 n is the length of binary representation of the encrypted message m . In Zeth
1531 however, the only data encrypted are the notes. As such, the size of the plaintexts
1532 is $\text{NOTEBYTELEN} * \text{BYTELEN}$ bits.

1533 **Remark 3.5.2.** In the following, we omit the explicit conversion from \mathbb{B}^n to
1534 $\mathbb{B}_{\mathbb{Y}}^{\lceil n/\text{BYTELEN} \rceil}$ when passing the output of ChaCha20 to the Poly1305 algorithms.

5. Randomly generate $(r, s) \in \mathbb{B}_{\mathbb{Y}}^{\text{MACKEYBYTELEN}/2} \times \mathbb{B}_{\mathbb{Y}}^{\text{MACKEYBYTELEN}/2}$ and clamp it:

$$\begin{aligned} r[3] &\leftarrow r[3] \ \& \ 0x0F \\ r[7] &\leftarrow r[7] \ \& \ 0x0F \\ r[11] &\leftarrow r[11] \ \& \ 0x0F \\ r[15] &\leftarrow r[15] \ \& \ 0x0F \\ r[4] &\leftarrow r[4] \ \& \ 0xFC \\ r[8] &\leftarrow r[8] \ \& \ 0xFC \\ r[12] &\leftarrow r[12] \ \& \ 0xFC \end{aligned}$$

¹¹We assume here that esk has been clamped as discussed in Section 3.5.2

6. Generate the related tag:

$$\tau = \text{Poly1305.Tag}(mk, ct_{\text{Sym}}) \in \mathbb{B}_Y^{\text{TAGBYTELEN}}.$$

7. Create the asymmetric ciphertext as:

$$ct = epk \| ct_{\text{Sym}} \| \tau \in \mathbb{B}_Y^{\text{CTBYTELEN}}.$$

1535 8. Return ct . As consequence $\text{ENCZETHNOTELEN} = \text{CTBYTELEN} * \text{BYTELEN}$ bits.

1536 **EncSch.Dec**

1537 The decryption, on inputs $(sk, ct) \in \mathbb{B}_Y^{\text{ESKBYTELEN}} \times \mathbb{B}_Y^{\text{CTBYTELEN}}$, is defined as follows:

1. Parse the ciphertext ct as:

$$\begin{aligned} epk &\leftarrow ct[: \text{EPKBYTELEN} - 1] \\ ct_{\text{Sym}} &\leftarrow ct[\text{EPKBYTELEN} : \text{EPKBYTELEN} + \text{NOTEBYTELEN} - 1] \\ \tau &\leftarrow ct[\text{EPKBYTELEN} + \text{NOTEBYTELEN} : \text{EPKBYTELEN} + \text{NOTEBYTELEN} + \text{TAGBYTELEN} - 1] \end{aligned}$$

2. Recover the shared secret

$$ss = \text{X25519}(sk, epk).$$

3. Compute the $ek \| mk$

$$\begin{aligned} ek &= \text{Blake2b512}(\text{encTag} \| epk \| ss)[: \text{SYMKEYBYTELEN} - 1] \\ mk &= \text{Blake2b512}(\text{encTag} \| epk \| ss)[\text{SYMKEYBYTELEN} : \text{SYMKEYBYTELEN} + \text{MACKEYBYTELEN} - 1]. \end{aligned}$$

4. Verify that the ciphertext has not been forged:

$$\text{Poly1305.Vf}(mk, \tau, ct_{\text{Sym}})$$

5. (If the MAC verifies) decrypt:

$$m = \text{ChaCha20.Dec}(ek, \text{CHACHABLOCKCOUNTERVALUE}, \text{CHACHANONCEVALUE}, ct_{\text{Sym}})$$

1538 6. Return m .

1539 3.5.4 Security requirements satisfaction

1540 DHAES has already been proved to be IND-CCA2 secure (see [ABR99, Section 3.5, The-
 1541 orem 6])¹² and to the best of our knowledge there is no paper showing IK-CCA security.
 1542 The only proof we have found is related to DHIES scheme [ABN10], that is a prime order
 1543 group version of DHAES. In the following, we provide a proof for IK-CCA security of
 1544 DHAES by adapting that proof to our case.

Theorem 3.5.1 (IK-CCA of DHAES). *Let DHAES be the asymmetric encryption scheme as defined above. Let \mathcal{A} be an adversary for the IK-CCA game, then there exists a HDHI adversary \mathcal{B} of $(\mathcal{H}, \text{SetupG})$ and a SUF-CMA adversary \mathcal{C} of MAC such that*

$$\text{Adv}_{\text{DHAES}, \mathcal{A}}^{\text{ik-cca}}(\lambda) \leq 2 \cdot \text{Adv}_{\mathcal{H}, \text{SetupG}, \mathcal{B}}^{\text{hdhi}}(\lambda) + \text{Adv}_{\text{MAC}, \mathcal{C}}^{\text{suf-cma}}(\lambda).$$

1545 *The adversaries \mathcal{B} and \mathcal{C} have the same running time as \mathcal{A} ¹³.*

1546 *Informal proof.* As already mentioned, DHAES is similar to DHIES scheme, except for
 1547 the underlying group and the way the symmetric keys are constructed. As consequence,
 1548 IK-CCA property for DHAES can be shown similarly to the approach in [ABN10, Theorem
 1549 6.2]. More precisely, they show that one can construct from an attacker \mathcal{A} for the IK-CCA
 1550 game two attackers \mathcal{B} and \mathcal{C} for the ODH and SUF-CMA games. Actually, they make
 1551 use of a $\bar{\mathcal{B}}$ attacker for the ODH2 game [ABN10, Figure 20] and then apply [ABN10,
 1552 Lemma 6.1] to obtain an attacker \mathcal{B} ¹⁴ in the ODH game. We adopt a similar strategy,
 1553 working with HDHI, HDHI2 and Lemma 1.5.1.

Let \mathcal{A} be an attacker for the IK-CCA game, and let $\bar{\mathcal{B}}$ be an attacker for the HDHI2 game described in Fig. 3.11. We show that,

$$\text{Adv}_{\mathcal{H}, \text{SetupG}, \bar{\mathcal{B}}}^{\text{hdhi2}}(\lambda) = |\Pr[\text{IK-CCA}^{\mathcal{A}}(\lambda) = 1] + \Pr[\text{G}_0^{\mathcal{A}}(\lambda) = 1] - 1|$$

1554 where G_0 is the security game described in Fig. 3.12.

1555 Given an HDHI2 challenge $(\llbracket u \rrbracket, \llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket, w_{b_2,0}, w_{b_2,1})$, an adversary $\bar{\mathcal{B}}$ samples $b \leftarrow_{\$} \{0, 1\}$
 1556 and runs \mathcal{A} on $\llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket$ (note that b_2 is the random bit chosen by the $\bar{\mathcal{B}}$ challenger in the
 1557 HDHI2 game). $\bar{\mathcal{B}}$ constructs oracles $\text{O}^{\text{Dec}_{sk_i}}$ where the queries $(\tau \parallel ct_{\text{Sym}} \parallel \tau)$ are processed
 1558 as follows: if $\tau \neq \llbracket u \rrbracket$, then $\bar{\mathcal{B}}$ queries related HDHI2 oracle to obtain $ek \parallel mk \leftarrow \text{O}^{\text{HDHI}_{v_i}}(\tau)$
 1559 (see Fig. 3.11). If $\tau = \llbracket u \rrbracket$, $w_{b_2,i}$ is parsed as $ek \parallel mk$. In both cases, it checks that
 1560 $\text{MAC.Vf}(mk, ct_{\text{Sym}}, \tau) = 1$ and, if so, returns $m \leftarrow \text{Sym.Dec}(ek, ct_{\text{Sym}})$. We note that
 1561 \mathcal{A} cannot query the challenged ciphertext. $\bar{\mathcal{B}}$ returns 0 if and only if $b = \tilde{b}$. It easy to
 1562 see that if b_2 is equal to 0, then all symmetric encryption and MAC keys used for the
 1563 challenge ciphertext $(\tau^* \parallel ct_{\text{Sym}}^* \parallel \tau^*)$ and decryption responses are exactly as in a DHAES
 1564 game.

¹²Specifically, if Sym is IND-CPA secure, it holds that H is HDHI secure and MAC is SUF-CMA secure.

¹³In order to give an asymptotic version of the theorem, the number of queries q has been substituted by the fact of considering PPT adversaries.

¹⁴Note that in [ABN10] the IK-CCA game is a particular case of the AI-CCA game that requires two input messages in the LR query. In order to reason only about the key-privacy, the two messages m_0 and m_1 are constrained to be equal.

Adversary $\bar{\mathcal{B}}(\llbracket u \rrbracket, \llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket, w_{b_2,0}, w_{b_2,1})$	$\bar{\mathcal{B}}$ simulation of $\mathcal{O}^{\text{Dec}_{sk_i}}(\tau \parallel ct_{\text{Sym}} \parallel \tau)$
$b \leftarrow \$_\{0,1\}$	if $\tau \neq \llbracket u \rrbracket$
$(m, state) \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(\llbracket v_0 \rrbracket, \llbracket v_1 \rrbracket)$	$ek \parallel mk \leftarrow \mathcal{O}^{\text{HDH}_{v_i}}(\tau)$
$ek \parallel mk \leftarrow w_{b_2,b}$	else
$\tau^* \leftarrow u$	$ek \parallel mk \leftarrow w_{b_2,i}$
$ct_{\text{Sym}}^* \leftarrow \text{Sym.Enc}(ek, m)$	fi
$\tau^* \leftarrow \text{MAC.Tag}(mk, ct_{\text{Sym}}^*)$	if $\text{MAC.Vf}(mk, ct_{\text{Sym}}, \tau) = 1$
$\tilde{b} \leftarrow \mathcal{A}^{\mathcal{O}^{\text{Dec}_{sk_0}}, \mathcal{O}^{\text{Dec}_{sk_1}}}(\tau^* \parallel ct_{\text{Sym}}^* \parallel \tau^*, state)$	return $\text{Sym.Dec}(ek, ct_{\text{Sym}})$
return $\tilde{b} = b$	else
	return \perp
	fi

Figure 3.11: Description of the adversary $\bar{\mathcal{B}}$ for HDHI2, simulating DHAES game for \mathcal{A} .

If $b_2 = 1$, then $w_{1,0}$ and $w_{1,1}$ are random strings and the challenge ciphertext and decryption responses are given as in the G_0 game described in Fig. 3.12. So we get,

$$\Pr[\text{HDHI2}^{\bar{\mathcal{B}}}(\lambda) = 1] = \frac{1}{2} \cdot \Pr[\text{IK-CCA}^{\mathcal{A}}(\lambda) = 1] + \frac{1}{2} \cdot \Pr[\mathsf{G}_0^{\mathcal{A}}(\lambda) = 1].$$

1565 And from the definition of HDHI2 advantage we have

$$\text{Adv}_{\mathcal{H}, \text{Setup}_{\mathsf{G}, \bar{\mathcal{B}}}}^{\text{hdhi2}}(\lambda) = |\Pr[\text{IK-CCA}^{\mathcal{A}}(\lambda) = 1] + \Pr[\mathsf{G}_0^{\mathcal{A}}(\lambda) = 1] - 1|.$$

1566 At this point, we can conclude as in [ABN10, Theorem 6.2], with the only difference
 1567 of applying Lemma 1.5.1 instead of [ABN10, Lemma 6.1] and by defining a game G_1
 1568 that is *identical until bad*¹⁵ G_0 defined in Fig. 3.12. \square

1569 3.5.5 Final notes and observations

1570 In this section we list some notes regarding the approach taken in Zcash (see [ZCa19,
 1571 Section 8.7]), and other observations:

- 1572 • *Key derivation parameters:* in DHAES construction, the only required input vari-
 1573 ables are the shared secret ss and epk . In the Sprout release of Zcash, additional
 1574 parameters were added (i.e. h_{sig} , pk_{enc} and a counter i) (see [ZCa19, 5.4.4.2]):
 1575 they state that h_{sig} was used in order to get a different randomness extractor for
 1576 each joinsplit transfer in order to limit the degradation of the security and weaken
 1577 assumption on the hash. The authors believed, about the use of long-standing
 1578 public key pk_{enc} , that it might be necessary for IND-CCA2 security and for post-
 1579 quantum privacy (in the case where the quantum attacker does not have the public

¹⁵Games G_i and G_j are said to be *identical until bad* if they differ only in statements that follow the setting of the **bad** variable to *True*. **bad** is initialized with *False*

$\begin{aligned} & \mathbf{G}_0(\lambda) \\ \hline & (q, \mathbb{G}, \mathbf{g}, +) \leftarrow \text{SetupG}(1^\lambda) \\ & (sk_0, pk_0), (sk_1, pk_1) \leftarrow \text{KGen}(1^\lambda) \\ & \tau^* \leftarrow \text{\$ } \mathbb{G} \\ & ek^* \leftarrow \text{\$ } \{0, 1\}^{kLen} \\ & mk^* \leftarrow \text{\$ } \{0, 1\}^{mLen} \\ & (m, state) \leftarrow \mathcal{A}^{\mathcal{O}^{\overline{\text{Dec}}_{sk_0}}, \mathcal{O}^{\overline{\text{Dec}}_{sk_1}}}(pk_0, pk_1) \\ & b \leftarrow \text{\$ } \{0, 1\} \\ & ct_{\text{Sym}}^* \leftarrow \text{Sym.Enc}(ek^*, m) \\ & \tau^* \leftarrow \text{MAC.Tag}(mk^*, ct_{\text{Sym}}^*) \\ & \tilde{b} \leftarrow \mathcal{A}^{\mathcal{O}^{\overline{\text{Dec}}_{sk_0}}, \mathcal{O}^{\overline{\text{Dec}}_{sk_1}}}(\tau^* \ ct_{\text{Sym}}^* \ \tau^*, state) \\ & \mathbf{return } \tilde{b} = b \end{aligned}$	$\begin{aligned} & \text{Oracle } \mathcal{O}^{\overline{\text{Dec}}_{sk_i}}(\tau \ ct_{\text{Sym}} \ \tau) \\ \hline & \mathbf{if } \tau = \tau^* \\ & \quad m \leftarrow \perp \\ & \quad \mathbf{if } \text{MAC.Vf}(mk^*, ct_{\text{Sym}}, \tau) = 1 \\ & \quad \quad \mathbf{bad} \leftarrow \mathbf{true} \\ & \quad \quad m \leftarrow \text{Sym.Dec}(ek^*, ct_{\text{Sym}}) \\ & \quad \mathbf{fi} \\ & \mathbf{else} \\ & \quad m \leftarrow \text{Dec}(sk_i, \tau \ ct_{\text{Sym}} \ \tau) \\ & \mathbf{fi} \\ & \mathbf{return } m \end{aligned}$
---	--

Figure 3.12: \mathbf{G}_0 game and related decryption oracles for Theorem 3.5.1.

1580 key) [zcaa]. None of these additional components are used any longer starting from
1581 the Sapling release (see [ZCa19, 5.4.4.4]). To the best of our knowledge there is
1582 no formal reason to use the note counter i as an input to the KDF: an explana-
1583 tion could be to avoid the same session key being reused for multiple notes, but
1584 this should not be a problem since a different nonce or block counter is used for
1585 the symmetric cipher (actually this is already mandated in the case where epk is
1586 reused, as described below).

1587 • *Reuse of ephemeral keys epk* : **Zcash** reuses the same ephemeral keys epk (and
1588 different nonces) for two ciphertexts in a joinsplit description, claiming that this
1589 does not affect the security of the scheme as soon as the HDHI assumption of the
1590 DHAES security proof is adapted. Note that the proof they refer to is related to
1591 the IND-CCA2 notion.

1592 • Note that in **Zcash Sprout** and **Sapling**, being able to break the Elliptic Curve
1593 Diffie-Hellman Problem on **Curve25519** or **Jubjub** would not help to decrypt the
1594 transmitted notes ciphertext unless the receiver pk_{enc} is known or guessed. On the
1595 other hand, having pk_{enc} into the hash (as used in **Sprout**) may violate in principle
1596 the key-privacy of the encryption scheme. For these reasons, we underline that
1597 the protocol should enforce a mechanism that does not reveal users public keys to
1598 increase the security.

1599 • In [ABN10], the concept of *robustness* for an asymmetric encryption scheme is
1600 introduced: it formalizes the infeasibility of producing a ciphertext valid under two
1601 different public encryption keys. We note that this is particularly useful for **Zeth**
1602 since only the intended receiver will be able to decrypt the encrypted note. In fact,
1603 the definition is more general since it also covers the case in which a decryption

1604 is successful but returns an incorrect plaintext. This prevents situations where
1605 a user, scanning the **Mixer** logs for incoming transactions, gets a false positive
1606 decryption and stores garbage notes.

Note

We note however, that the “false-positive” situation above can be prevented by relying on a weaker notion of robustness called *collision-freeness* [Moh10]. In fact, as described in Section 2.6, the procedure to receive a *ZethNote* requires to decrypt the ciphertext emitted by the **Mixer**, and then to verify that the recovered plaintext is the opening of a commitment in the Merkle tree. As such, since the *collision-freeness* of the encryption ensures that plaintexts recovered under different keys are different (i.e. “do not produce a collision”), then we know that plaintexts recovered by parties who are not the intended recipient will fail the “commitment opening verification”, leading the payment to be rejected, and solving the aforementioned false-positive issue.

1607

1608 In [ABN10, Section 6], the authors prove that DHIES can be made strongly robust.
1609 The proof can be easily adapted to work with DHAES.

1610 • *No public key validation for X25519*: cryptographers have been discussing the ab-
1611 sence of any mandated public key validation or checks on the result of X25519.
1612 For example, in [LHT16, Section 6.1], an optional zero check is introduced in order
1613 to assure that the result of X25519 is not 0: this avoids a situation in which one
1614 of the two parties can force the result of the key-exchange by using a small order
1615 point as public key. This property is generally defined as *contributory behaviour*,
1616 that is, none of the parties is able to force the output of a key exchange. However,
1617 protocols do not have all the same security requirements and adding default checks
1618 in the Curve25519 specifications would be superfluous in most cases and would add
1619 complexities that Bernstein has deliberately chosen to avoid (*simple implementa-*
1620 *tion principle*). More importantly, Diffie-Hellman does not require *contributory*
1621 *behaviour* property [Per17]: modern view is that the only requirements are key
1622 indistinguishability and, in case of an active attacker, that the output of the key
1623 exchange should not produce a low-entropy function of the honest party’s private
1624 key (e.g. small-subgroup and invalid-curve attacks). Since these two properties are
1625 considered satisfied by Curve25519, there is no need to add extra checks to the
1626 Curve25519 specification. We conclude by observing that in the Sprout release, the
1627 Zcash protocol does not specify any point validation and makes use only of the
1628 private key clamping to keep Diffie-Hellman key exchange secure.

1629 3.6 ZkSnarkSch instantiation

1630 Groth’s proof system Groth16 [Gro16] is the most efficient known zk-SNARK (in terms
 1631 of the proof size and proof and verification cost) for QAPs, and thus one of the most
 1632 efficient NIZK for proving statements on arithmetic circuits (consisting of addition and
 1633 multiplication gates over a finite field \mathbb{F}). Below we present Groth16’s key generation,
 1634 prover, verifier, and simulator algorithms, adjusted as described in [BGM17] to further
 1635 reduce the size of *srs* and proofs, and to make the KGen algorithm more amenable to
 1636 implementation as a multi-party computation.

1637 In what follows, let the number *constNo* of constraints in the relation \mathbf{R} be fixed.
 1638 Without loss of generality we consider *constNo* to be an *upper bound* on the number of
 1639 constraints in the \mathbf{R} parameter, and assume that there exists some *constNo*-th root of
 1640 unity $\omega \in \mathbb{F}_{\text{rCUR}}$. Define $\ell_i(X)$ to be the *i*-th Lagrange polynomial of degree $(\text{constNo} - 1)$
 1641 over the set $\{\omega^i\}_{i \in [\text{constNo}]}$, and let $\ell(X)$ be the unique non-zero polynomial of degree
 1642 *constNo* that satisfies $\ell(\omega^i) = 0$ for all $i \in [\text{constNo}]$.

1643 We note that the requirement that there exists a *constNo*-th root of unity ω imposes
 1644 a restriction on the maximum number of constraints in \mathbf{R} that the scheme can support.
 1645 In the particular case of $\omega \in \mathbb{F}_{\text{rBN}}$, the restriction becomes $\text{constNo} \leq 2^{28}$. For \mathbb{F}_{rBLS} this
 1646 becomes $\text{constNo} \leq 2^{47}$.

1647 Furthermore, we denote by $\text{inp} \in \mathbb{F}^{\text{inpNo}+1}$ the tuple of variables (i.e. “circuit wires”)
 1648 in the algebraic representation of the relation \mathbf{R} , such that:

- 1649 • $\text{inp}_0 = 1_{\mathbb{F}}$ (the multiplicative identity in \mathbb{F}),
- 1650 • $(\text{inp}_1, \dots, \text{inp}_{\text{inpNoPrim}})$ represent variables in the statement,
- 1651 • $(\text{inp}_{\text{inpNoPrim}+1}, \dots, \text{inp}_{\text{inpNo}})$ represent variables in the witness (so-called “auxil-
 1652 iary input”).

1653 KGen($\mathbf{R}, 1^\lambda$):

- 1654 i. Pick trapdoor $td = (\tau, \alpha, \beta, \delta) \leftarrow \$ (\mathbb{Z}_p^* \setminus \{\omega^{i-1}\}_{i=1}^{\text{constNo}}) \times (\mathbb{Z}_p^*)^3$;
- ii. For $j \in \{0, \dots, \text{inpNo}\}$, let

$$u_j(\tau) = \sum_{i=1}^{\text{constNo}} U_{ij} \ell_i(\tau),$$

$$v_j(\tau) = \sum_{i=1}^{\text{constNo}} V_{ij} \ell_i(\tau),$$

$$w_j(\tau) = \sum_{i=1}^{\text{constNo}} W_{ij} \ell_i(\tau);$$

iii. Set

$$\begin{aligned}
srs_{\mathbf{P}} &\leftarrow \left(\begin{array}{l} \llbracket \alpha \rrbracket_1, \llbracket \beta \rrbracket, \llbracket \delta \rrbracket, \{ \llbracket u_j(\tau) \rrbracket_1 \}_{j=0}^{inpNo}, \{ \llbracket v_j(\tau) \rrbracket \}_{j=0}^{inpNo}, \\ \{ \llbracket (u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau))/\delta \rrbracket_1 \}_{j=inpNoPrim+1}^{inpNo}, \\ \{ \llbracket \tau^i \ell(\tau)/\delta \rrbracket_1 \}_{i=0}^{constNo-2} \end{array} \right) \\
srs_{\mathbf{V}} &\leftarrow \left(\llbracket \alpha \rrbracket_1, \llbracket \beta \rrbracket_2, \llbracket \delta \rrbracket_2, \{ \llbracket \beta u_j(\tau) + \alpha v_j(\tau) + w_j \rrbracket_1 \}_{j=0}^{inpNoPrim} \right) \\
srs &\leftarrow (srs_{\mathbf{P}}, srs_{\mathbf{V}})
\end{aligned}$$

1655 **return** srs, td

1656 $\mathbf{P}(\mathbf{R}, srs_{\mathbf{P}}, prim = (inp_j)_{j=1}^{inpNoPrim}, aux = (inp_j)_{j=inpNoPrim+1}^{inpNo})$:

i. Define

$$a^\dagger(X) = \sum_{j=0}^{inpNo} inp_j u_j(X), \quad b^\dagger(X) = \sum_{j=0}^{inpNo} inp_j v_j(X), \quad c^\dagger(X) = \sum_{j=0}^{inpNo} inp_j w_j(X);$$

1657 ii. Define the polynomial $h(X) = (a^\dagger(X)b^\dagger(X) - c^\dagger(X))/\ell(X)$ and compute the
1658 coefficients $\{h_i\}_{i=0}^{constNo-2}$ of h , such that $h(X) = \sum_{i=0}^{constNo-2} h_i X^i$.

1659 iii. $r_a \leftarrow \mathbb{Z}_p$;

1660 iv. $r_b \leftarrow \mathbb{Z}_p$;

v. Compute proof elements:

$$\begin{aligned}
\mathbf{a} &\leftarrow \sum_{j=0}^{inpNo} inp_j \llbracket u_j(\tau) \rrbracket_1 + \llbracket \alpha \rrbracket_1 + r_a \llbracket \delta \rrbracket_1 \\
\mathbf{b} &\leftarrow \sum_{j=0}^{inpNo} inp_j \llbracket v_j(\tau) \rrbracket_2 + \llbracket \beta \rrbracket_2 + r_b \llbracket \delta \rrbracket_2 \\
\mathbf{c} &\leftarrow \sum_{j=inpNoPrim+1}^{inpNo} inp_j \left\llbracket \frac{u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau)}{\delta} \right\rrbracket_1 + \sum_{i=0}^{constNo-2} h_i \llbracket \tau^i \ell(\tau)/\delta \rrbracket_1 + \\
&\quad r_b \mathbf{a} + r_a \left(\sum_{j=0}^{inpNo} inp_j \llbracket v_j(\tau) \rrbracket_1 + \llbracket \beta \rrbracket_1 + r_b \llbracket \delta \rrbracket_1 \right) - r_a r_b \llbracket \delta \rrbracket_1
\end{aligned}$$

1661 **return** $\pi \leftarrow (\mathbf{a}, \mathbf{b}, \mathbf{c})$;

1662 $\mathbf{V}(\mathbf{R}, srs_{\mathbf{V}}, prim = (inp_j)_{j=1}^{inpNoPrim}, \pi)$:

i. Check that:

$$\begin{aligned}
\mathbf{a} \bullet \mathbf{b} &= \mathbf{c} \bullet \llbracket \delta \rrbracket_2 \\
&\quad + \left(\sum_{j=0}^{inpNoPrim} inp_j \llbracket u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau) \rrbracket_1 \right) \bullet \llbracket 1 \rrbracket_2 \\
&\quad + \llbracket \alpha \rrbracket_1 \bullet \llbracket \beta \rrbracket_2
\end{aligned}$$

1663 Note that $\llbracket \alpha \rrbracket_1$ and $\llbracket \beta \rrbracket_2$ are stored individually and used by the prover to re-
 1664 compute $\llbracket \alpha\beta \rrbracket_T$ seemingly redundantly. This is required in order to leverage the
 1665 pairing check functionality built in to **Ethereum**, which accepts a sequence of tuples
 1666 in $\mathbb{G}_1 \times \mathbb{G}_2$ and returns **true** if and only if the product of the resulting pairings
 1667 equals $\llbracket 1 \rrbracket_T$.

1668 **Sim**($\mathbf{R}, srs, td, prim$):

- 1669 i. Sample $a \leftarrow \$\mathbb{Z}_p$; $b \leftarrow \$\mathbb{Z}_p$;
 ii. Compute proof elements:

$$\begin{aligned} \mathbf{a} &\leftarrow \llbracket a \rrbracket_1 \\ \mathbf{b} &\leftarrow \llbracket b \rrbracket_2 \\ \mathbf{c} &\leftarrow \frac{1}{\delta} \cdot \left(ab - \alpha\beta - \sum_{j=0}^{inpNoPrim} inp_j(u_j(\tau)\beta + v_j(\tau)\alpha + w_j(\tau)) \right) \llbracket 1 \rrbracket_1 \end{aligned}$$

1670 **return** $\pi \leftarrow (\mathbf{a}, \mathbf{b}, \mathbf{c})$;

1671 Chapter 4

1672 Implementation considerations 1673 and optimizations

1674 4.1 Client security considerations

1675 In this section we consider some details of client *wallet* software that manages user's
1676 private and public keys, **Zeth** notes, and interacts with the **Mixer** contract.

1677 Due to the processing and storage requirements involved, we consider it impractical
1678 for all **Zeth** client implementations to assume that a dedicated **Ethereum** node (miner
1679 node or archive node) is run on the same host as the wallet. Therefore, in order to
1680 interact with the **Ethereum** network, wallet software must communicate with external
1681 **Ethereum** P2P nodes via their RPC channel, and must assume that these nodes are
1682 completely outside the wallet's control. *From a security standpoint, connected **Ethereum***
1683 *nodes should therefore be considered untrusted, and in particular the details of all RPC*
1684 *calls and responses should be considered publicly visible.* Note that even if the connected
1685 **Ethereum** node itself is not malicious, 3rd parties able to see network traffic may also be
1686 able to gain an insight into the RPC communication of a specific **Zeth** client.

Note

Note that there are several possible models besides the fully untrusted **Ethereum** node. Organizations or individuals could host one or more “trusted” **Ethereum** nodes, which clients can securely connect to (if they trust the host). This centralization would represent a security trade-off. From the point of view of clients it would create a single point of trust, and for potential malicious observers or attackers it would represent a valuable target.

1687

1688 In what follows we focus on preventing data leaks through network traffic. We do
1689 not consider adversaries with physical access to the machine running the wallet (see Ap-
1690 pendix C).

Note

Importantly, we focus here on information leakages intrinsic to network communication patterns of the **Zeth** protocol. However, in order to protect against sophisticated adversaries, it is necessary to use network-level anonymity solutions to protect the source of messages emitted on the network. While this is outside of the scope of the **Zeth** protocol, we highly encourage implementers to establish threat models and consider using technologies like *mixnets* to protect against network analysis (see e.g. [PHE⁺17, DG09]).

1691

1692 4.1.1 Syncing and waiting

1693 **Zeth** clients must periodically synchronize with the latest state of the blockchain. This
1694 is necessary to keep track of the data held by the **Mixer** contract, and to detect notes
1695 received by the user of the wallet, storing them for future transactions.

1696 Clients should synchronize with **Ethereum** nodes in such a way that information is
1697 not leaked. As such:

- 1698 1. Clients **MUST** use consensus evidence and block headers to verify all data they
1699 receive from **Ethereum** nodes.
- 1700 2. Clients **MUST** locally store all parts of the **Mixer** state they require in order to
1701 function.
- 1702 3. Clients **MUST** obtain all such information by “synchronizing” with the **Ethereum**
1703 blockchain and parsing relevant events emitted by **Mixer**. Clients **MUST NOT** query
1704 the **Mixer** state via RPC.
- 1705 4. Clients **SHOULD** take steps to avoid being identified while synchronizing (see Ap-
1706 pendix C.2. For example, clients **SHOULD** vary the set of **Ethereum** nodes that they
1707 connect to, and **SHOULD NOT** always sync from the block following the last one that
1708 they processed.
- 1709 5. Clients **SHOULD NOT** re-request blocks or transaction receipts that are of particular
1710 interest to them. They **SHOULD** process all events, emitted by **Mixer**, in the same
1711 way.
- 1712 6. Clients **SHOULD NOT** make any RPC calls or change their externally visible behaviour
1713 in response to blocks or transaction receipts that are of interest to them.

1714 Use of contract queries

1715 We suggest that clients **SHOULD NOT** directly query the contract state, for the reasons
1716 discussed in Appendix C.2 and Appendix C.3 (and consequently, Section 4.2 suggests
1717 that the **Mixer** contract should, as far as possible, not expose public methods). The

1718 Zeth protocol prohibits direct queries of the state of $\widetilde{\text{Mixer}}$ (via *public* smart-contract
1719 functions) because they introduce a risk that client implementations will leak information
1720 by using them.

1721 If implementers choose to add public methods to the $\widetilde{\text{Mixer}}$ contract (for application-
1722 specific reasons), they should consider carefully the security issues raised in Appendix C.
1723 This specification assumes that `Mix` is the only public method of the $\widetilde{\text{Mixer}}$ contract.

1724 4.1.2 Note management

1725 `Mix` calls on the $\widetilde{\text{Mixer}}$ contract emit log events containing new commitment values,
1726 nullifiers, the new Merkle root and the secret data for new notes (encrypted using a key
1727 derived from the recipients public key). As clients synchronize with the latest state of
1728 the blockchain, they **MUST** read these events and correctly process the data they contain.

- 1729 1. Clients **MUST** process the `MixEventDType` event for every `Mix` transaction, in the
1730 order in which they appear in the blockchain.
- 1731 2. Clients implementing spending functionality **MUST** use the commitment values in
1732 events to track the state of the Merkle tree. The Merkle tree state will be used
1733 to generate Merkle paths for future transactions, and **MUST** be made available to
1734 the client without the need to query the contract. (Note that not all commitments
1735 must necessarily be persisted – see Section 4.3).
- 1736 3. Clients that can receive notes **MUST** attempt to decrypt the ciphertexts for every
1737 transaction (see Item 2 in Section 2.6).
- 1738 4. Clients **MUST NOT** perform any network-related action, including closing the RPC
1739 connection, dependent on successful/unsuccessful decryption of ciphertexts (see
1740 Appendix C.3).
- 1741 5. Clients that can receive notes **MUST** attempt to parse any successfully decrypted
1742 plaintext (that is, ensure it is well-formed as in Item 3a in Section 2.6).
- 1743 6. Clients **MUST NOT** perform any network-related action, including closing the con-
1744 nection, dependent on successful / unsuccessful parsing (see Appendix C.4).
- 1745 7. Clients that can receive notes **MUST** verify that successfully parsed plaintext data
1746 is the opening of the corresponding commitment in the transaction (see Item 3b
1747 in Section 2.6).
- 1748 8. Clients **MUST NOT** perform any network-related action, including closing the con-
1749 nection, dependent on whether the parsed note data is the opening of the corre-
1750 sponding commitment (see Appendix C.4).
- 1751 9. Clients **MUST** confirm that, after adding the new commitments, the local repre-
1752 sentation of the Merkle tree of commitments has a root consistent with the event
1753 data.

- 1754 10. Clients **SHOULD** keep a *local* record of the data related to valid decrypted notes.
1755 This will be required in order to spend the notes in a future transaction.
- 1756 11. Clients implementing spending functionality **SHOULD** process all nullifiers in **Mix**
1757 transaction events, checking for any corresponding notes previously recorded. Any
1758 such note should be marked as spent in the client’s local record.

1759 4.1.3 Prepare arguments for **Mix** transaction

1760 Clients **MUST NOT** query **Ethereum** nodes while generating any arguments to a **Mix** call.
1761 In particular, Merkle paths **MUST** be calculated using the client’s local representation of
1762 the Merkle tree of commitments that was constructed by parsing events.

1763 Where the zero-knowledge proof is generated by some external process, clients **MUST**
1764 put in place sufficient security schemes to ensure that:

- 1765 • they are communicating with an authentic proof generation process (not a man-
1766 in-the-middle), and
- 1767 • data sent to and from the proving process cannot be observed in transit and tam-
1768 pered with by a third party, and
- 1769 • the proof has been generated for the correct instance–witness pair¹

1770 Without these safe-guards, the operation of the system and the secret data required
1771 to spend the input notes may be compromised. See Appendix C.6.

1772 4.1.4 Wallet backup and recovery

1773 Given the restrictions placed on clients and their interaction with the **Mixer** contract,
1774 it follows that clients must store all data required to spend notes owned by their users’
1775 addresses, and to verify the validity of incoming notes. If this local data is lost, it must
1776 be reconstructed before client operations can resume.

1777 **Zeth** private keys (see Table 1.5) can be used to fully restore client state. In this
1778 case, clients **MUST** retrieve all events from the beginning of the **Mixer** contract’s his-
1779 tory, decrypting notes and tracking nullifiers, as described in the previous sections, to
1780 reconstruct the set of unspent notes that they own.

1781 Without a backup of the private keys it is not possible to restore wallet state. As
1782 such, private keys are the minimal set of data that must be securely stored and backed
1783 up, and clients **SHOULD** provide support for this mode of recovery. However, to avoid the
1784 need to scan all events emitted by **Mixer** (a very expensive operation) implementations
1785 **SHOULD** also support back ups of further state data (such as the representation of the
1786 Merkle tree of note commitments, the set of unspent notes, etc) to allow more efficient
1787 modes of recovery.

¹Although given an acceptable zk-proof π for an instance *prim* it is infeasible to check which witness has been used – which comes directly from the zero-knowledge property – we need to assure security measures that prevents any third party from mauling and tampering with the proof generation process.

1788 4.2 Contract security considerations

1789 Section 4.1 mentions several considerations for client implementations, concerning how
1790 they interact with the contract. These must be taken into account when authoring the
1791 contract code, to ensure that clients can securely retrieve the information they need to
1792 operate without encouraging them to perform insecure operations.

- 1793 1. **Mixer** MUST validate inputs, the contract needs to ensure that the primary inputs
1794 are elements of the scalar field $\mathbb{F}_{\mathbf{r}_{\text{CUR}}}$ (that is, they are in the range $\{0, \dots, \mathbf{r}_{\text{CUR}} - 1\}$).
- 1795 2. **Mixer** MUST output events for valid Mix calls, including:
 - 1796 (a) commitment for each new note;
 - 1797 (b) nullifier for each spent note;
 - 1798 (c) value of new Merkle root of commitments;
 - 1799 (d) ciphertexts for each new note;
 - 1800 (e) implementation-specific data (such as the one-time sender public key specified
1801 in Section 3.5, required to decrypt the ciphertexts).
- 1802 3. The Mix function MUST be *payable*², to support non-zero *vin*.
- 1803 4. **Mixer** MUST NOT expose any public methods except for Mix.

1804 4.3 Efficiency and scalability

1805 4.3.1 Importance of performance

1806 Poor performance and scalability has several impacts on the viability of the system.

1807 Efficiency and performance are arguably most important for the **Mixer** contract,
1808 where gas usage directly affects the monetary cost of using **Zeth** to transfer value. That
1809 is, high gas costs could make transactions very expensive, and therefore not practical for
1810 many use-cases, undermining the utility and viability.

1811 High storage or compute requirements on the client would severely restrict the set
1812 of devices on which **Zeth** client software can run, and long delays when sending or
1813 receiving transactions can adversely affect the user-experience, discouraging some users
1814 and undermining the privacy promises of the system.

1815 Although the proof-of-concept implementation of **Zeth** is not intended to be used in a
1816 production environment, one of its aims is to demonstrate the practicality of the protocol
1817 in terms of transaction costs. Therefore, some of the techniques described here have been
1818 included in the proof-of-concept implementation, while in some cases implementers of
1819 production software may wish to make different trade-offs.

²see <https://solidity.readthedocs.io/en/v0.6.2/types.html?highlight=payable#function-types>

1820 4.3.2 Cost centers

1821 One important factor, primarily affecting client performance, is the cost of zero-knowledge
1822 proof generation. This is directly related to the number of constraints used to repre-
1823 sent the statement in Section 2.2, which in turn depends on the specific cryptographic
1824 primitives used (see Chapter 3).

1825 Note that cryptographic primitives which are “snark-friendly” (i.e. can be imple-
1826 mented using fewer gates in an arithmetic circuit) may not necessarily run efficiently on
1827 the EVM or on standard hardware. As such, trade-offs must be made between proof
1828 generation cost and the gas costs of state transitions. An example of this is the hash
1829 function used in the Merkle tree of commitments. This is not only used in the statement
1830 of Section 2.2 (to verify Merkle proofs, see Section 2.2), but also on the client (to create
1831 Merkle proofs, see Section 2.3) and in the $\widetilde{\text{Mixer}}$ contract (to compute the Merkle root,
1832 see Section 2.5).

1833 Aside from the specific hash function used, implementers have some freedom in the
1834 data structures and algorithms used to maintain the Merkle tree and generate proofs.
1835 Because of this freedom, and the importance of the chosen algorithms on performance
1836 across all components of the system, the majority of this section focuses on the details
1837 of the Merkle tree.

1838 As described in Chapter 2, **Zeth** notes are maintained and secured by the Merkle tree,
1839 whose depth MKDEPTH must be fixed when the contract is deployed. Therefore, MKDEPTH
1840 determines the maximum number of notes (2^{MKDEPTH}) that may be created over the lifetime
1841 of the deployment. To ensure the utility of **Zeth**, MKDEPTH must be sufficiently large,
1842 and therefore the following includes a discussion of *scalability* with respect to MKDEPTH .

1843 Also, due to the fact that MKDEPTH is fixed, we assume that Merkle proofs are com-
1844 puted as MKDEPTH -tuples, no matter how many leaves have been populated. Unpopulated
1845 leaves are assumed to take some default value (usually a string of zero bits).

1846 4.3.3 Client performance

1847 Commitment Merkle tree

1848 The simplest possible implementation, which stores only the data items at the leaves
1849 of the tree, requires $2^{\text{MKDEPTH}} - 1$ hash invocations to compute the Merkle root or to
1850 generate a Merkle proof. The cost of this is too high to be practical for non-trivial
1851 values of MKDEPTH .

1852 An immediate improvement in compute costs can be achieved by simply storing
1853 all nodes (or all nodes whose value is not the default value) and updating only those
1854 necessary as new commitments are added. When adding JSOUT consecutive leaves to
1855 the tree, after $\mathcal{O}(\log_2(\text{JSOUT}))$ layers (requiring $\mathcal{O}(\text{JSOUT})$ hashes) we reach the common
1856 ancestor of all new leaves and can update the Merkle tree by proceeding along a single
1857 branch (of approximately $\text{MKDEPTH} - \log_2(\text{JSOUT})$ layers). Thus, the cost of updating
1858 the Merkle tree for a single transaction has a fixed bound which is linear in JSOUT and

1859 MKDEPTH. However, this doubles the storage cost of the tree since non-leaf nodes must
1860 also be persisted.

1861 In the case of the client, the Merkle tree will only be used to generate proofs for notes
1862 owned by the user of the client. Thereby Zeth clients need only store nodes of the Merkle
1863 tree that are required for this purpose, and may discard all others. In particular, any
1864 full sub-tree need only contain nodes that are part of Merkle paths associated with the
1865 user’s notes. Implementations that discard unnecessary nodes can achieve vast savings
1866 in storage space.

1867 Zero-knowledge proof generation

1868 As well as keeping the number of constraints as low as possible, it is important to ensure
1869 that the prover implementation is optimal and thereby that proving times are as short as
1870 possible. Proof generation should also exploit any available parallelism, to help reduce
1871 the time taken. This may require specific programming languages or frameworks to be
1872 used, necessitating that proof generation be performed by some external process (as is
1873 the case in the proof-of-concept implementation).

1874 The proof generation process can also be very memory intensive (in part due to the
1875 FFT calculations required), and so ensuring that enough RAM is present in the system
1876 is important to avoid long proof times.

1877 See Appendix C.6 for a discussion of related security concerns.

1878 4.3.4 Zero-knowledge proof verification (on-chain)

1879 Verification of the joinsplit statement via a zero-knowledge proof represents a significant
1880 computation, which must be carried out on-chain (by the Mixer contract) for each valid
1881 Zeth transaction. As described in Section 3.6, this verification cost increases linearly
1882 with the number of primary inputs to the statement – a scalar multiplication of a group
1883 element and a group addition operation must be performed for each primary input. A
1884 technique presented in [GGPR13, Section 4.5.1] can be applied to reduce this linear cost.

Given a relation \mathbf{R} , the corresponding language \mathbf{L} , and a collision resistant hash
function $H : \mathbf{L} \rightarrow \mathbb{F}_{\mathbf{r}_{\text{CUR}}}$, let

$$\mathbf{R}' = \{(prim', aux') \mid prim' = H(prim), aux' = (prim, aux), \text{ for } (prim, aux) \in \mathbf{R}\}$$

1885 be a new relation, with corresponding language $\mathbf{L}' \subset \mathbb{F}_{\mathbf{r}_{\text{CUR}}}$. To (probabilistically) verify
1886 that $prim \in \mathbf{L}$, a verifier can compute $H(prim)$ and check that $H(prim) \in \mathbf{L}'$. (By
1887 construction, if $H(prim) \in \mathbf{L}'$, there exists $(prim_0, aux) \in \mathbf{R}$, i.e. $prim_0 \in \mathbf{L}$ with
1888 $H(prim_0) = H(prim)$. By the collision-resistance of H we have $prim_0 = prim$ with
1889 overwhelming probability.)

1890 Informally, the original circuit is transformed as follows:

- 1891 • all *primary* inputs $prim$ become *auxiliary* inputs,
- 1892 • a single primary input h is added, and

1893 • the statement is extended such that h is the digest of the original primary inputs.

1894 This slightly increases the complexity of the statement to be proven, adding to the
1895 cost of generating proofs π' for the augmented statement, but minimizes the linear
1896 component of the verification cost (since the verifier must now only process a single
1897 primary input). Note that this technique does not require any change to the initial
1898 statement itself (in this case the joinsplit statement described in Section 2.2), or the
1899 data upon which it operates. The Mixer contract must perform this hash step before
1900 the zk-SNARK verification, although we note that the parameters are also unchanged.

1901 In the proof-of-concept implementation of Zeth, this technique is employed using a
1902 snark-friendly hash function constructed as follows.

1903 The Merkle-Damgård construction (see [MVOV96, Chapter 9]) can be applied to a
1904 collision-resistant compression function to yield a collision-resistant hash function, ac-
1905 cepting an arbitrary length input. We apply this to the compression function described
1906 in Section 3.2, which is chosen to be collision resistant over domain \mathbb{F}_{rCUR} , and efficiently
1907 implementable as arithmetic constraints. Thereby, the resulting hash function, in com-
1908 mon with the underlying compression function, can also be efficiently implemented to
1909 hash lists of elements in \mathbb{F}_{rCUR} (and this is exactly the form of the original primary inputs).

1910 4.3.5 Merkle tree updates (on-chain)

1911 For most components of the contract, the set of operations to be performed is strictly
1912 defined and the set of possible algorithmic optimizations that can be made is limited.
1913 In these cases, it is important to ensure that code is benchmarked and optimized to
1914 a reasonable degree, to minimize gas costs. We note that apart from the number and
1915 type of compute instructions executed, store and lookup operations have a significant
1916 impact on the gas used. In particular, storing new values is more expensive than over-
1917 writing existing values, and a gas rebate is made when contracts release stored values.
1918 See [Woo19, Appendix H.1] for further details.

1919 The primary component in which algorithmic optimizations can be made is the
1920 Merkle tree of note commitments. The Mixer contract must compute (and store) the
1921 new Merkle root after adding the JSOUT new commitments as leaves.

1922 As in Section 4.3.3, the simplest possible implementation which stores only the data
1923 items at the leaves of the tree, requires the full root to be recomputed, involving $2^{\text{MKDEPTH}} -$
1924 1 hash invocations. This quickly becomes impractical for non-trivial values of MKDEPTH.

1925 The first-pass optimization (also described in Section 4.3.3) can be used to ensure
1926 that the cost of updating the Merkle tree (number of hash computations, stores and
1927 loads) is bounded by a constant that is linear in the Merkle tree depth. This is the
1928 strategy used in the proof-of-concept implementation of Mixer.

1929 It may be possible to gain further improvements in gas costs by discarding nodes
1930 from the Merkle tree that are not required. Unlike clients, Mixer is only required to
1931 compute the new Merkle root, and does not need to create or validate Merkle proofs
1932 (as these are checked as part of the zero-knowledge proof). Consequently, *all* nodes in a

1933 sub-tree can be discarded when the sub-tree is full, and the optimization is much simpler
1934 to implement than on the client.

1935 Another possible strategy for decreasing the gas costs associated with Merkle trees
1936 is *Merkle Shrubs*, described in [Lab19, Section 2.2]. Under this scheme, the contract
1937 maintains a “frontier” of roots of sub-trees and Merkle proofs provided by clients (as
1938 auxiliary inputs to the \mathbf{R}^z circuit) contain a path from the leaf to one of the nodes in the
1939 frontier. The gas savings in this scheme are due to the fact that, for new commitments,
1940 the contract need only recompute the value of nodes from the leaf to the “frontier” (not
1941 all the way to the root of the tree). However this comes at the cost of complexity in the
1942 arithmetic circuit, which must verify a Merkle path to one of several frontier nodes.

1943 When choosing cryptographic primitives to be used on the EVM (and considering
1944 the trade-off with other platforms, described in Section 4.3.1) it may be valuable to note
1945 that the EVM supports so-called “pre-compiled contracts”. These behave like built-
1946 in functions providing very gas-efficient access to certain algorithms, such as Keccak.
1947 However, pre-compiled contracts exist only for a limited set of algorithms. Others must
1948 be implemented using EVM instructions.

1949 **4.3.6 Optimizing Blake2’s circuit.**

1950 After presenting Blake2s circuit and its components working on little endian variables,
1951 we show a few optimizations.

1952 **Helper circuits**

1953 We first define the following helper circuits needed in the Blake2s routine, operating on
1954 w -bit long words.

1955 **XOR circuits** The following XOR circuits on w -bit long variables have been imple-
1956 mented, we assume the inputs are boolean (this is not checked in these circuits),

- 1957 • “Classic” XOR circuit, which xors 2 variables,
1958 $a \oplus b = c$;
- 1959 • XOR with constant, which xors two variables and a constant,
1960 $a \oplus b \oplus c = d$, with c constant;
- 1961 • XOR with rotation, which xors two variables and rotates the result.
1962 $a \oplus b \ggg r = c$, with r constant, and \ggg the rightward rotation [MJS15, Section
1963 2.3]; i.e. for and constant $r < w$ we have $a_i \oplus b_i = c_{i+r \pmod w}$, for $i = 0, \dots, w$,

1964 Each of these circuits presents w constraints. Assuming that the inputs are boolean,
1965 the output is automatically boolean. To ascertain that both inputs are boolean (a and
1966 b), we would need $2 \cdot w$ more gates per circuit.³

³Making sure that no gates are duplicated in the circuit is very important to keep the proving time as small as possible. One challenge of writing RICS programs is to make sure that the statement is correctly represented, without redundancy, in order to keep the constraint system as small as possible.

1967 **Modular addition** We present here two circuits to verify modular arithmetic.

1968 **Double modular addition: $a + b = c \pmod{2^w}$.** This circuit checks that the
 1969 sum of two w -bit long variables in little endian format modulo 2^w is equal to a w -bit
 1970 long variable. More precisely, it checks the equality of the modular addition of $a + b$
 1971 ($\pmod{2^w}$) and c and the booleanness of the latter. We assume the inputs are boolean (this
 1972 is not checked in this circuit).

1973 As the addition of two w -bit long integers results in at most an $(w + 1)$ -bit integer,
 1974 we consider c to be $(w + 1)$ -bit long. We do not care about the last bit value, c_w , but
 1975 have to ensure its booleanness.

1976 The circuit presents the following $w + 2$ constraints, for a and b of size w , where
 1977 $w = 32$ in practice, and variable c of size $w + 1$, that:

$$\sum_{i=0}^{w-1} (a_i + b_i) \cdot 2^i = \sum_{j=0}^w c_j \cdot 2^j \quad (4.1)$$

$$\forall j \in \{0, \dots, w\}, (c_j - 0) \cdot (c_j - 1) = 0 \quad (4.2)$$

1978 **Triple modular addition: $a + b + c = d \pmod{2^w}$.** This circuit checks the
 1979 equality of a w -bit long variable d with the sum of three w -bit long variables in little
 1980 endian format modulo 2^w . More precisely, it checks the equality of the modular addition
 1981 of $a + b + c \pmod{2^w}$ and d and the booleanness of the latter. We assume the inputs are
 1982 boolean (this is not checked in this circuit).

1983 As the addition of three w -bit long integers results in at most an $(w + 2)$ -bit integer,
 1984 we consider d to be $(w + 2)$ -bit long. We do not care about the values of the last two
 1985 bits (d_w and d_{w+1}), but have to ensure their booleanness.

The circuit presents the following $w + 3$ constraints, for a , b and c of size w , where
 $w = 32$ in practice, and variable d of size $w + 2$, that:

$$\sum_{i=0}^{w-1} (a_i + b_i + c_i) \cdot 2^i = \sum_{j=0}^{w+1} d_j \cdot 2^j \quad (4.3)$$

$$\forall j \in \{0, \dots, w + 1\}, (d_j - 0) \cdot (d_j - 1) = 0 \quad (4.4)$$

1986 Blake2s routine circuit

1987 We define in this section the circuit of the Blake2 routine (see [MJS15, Section 3.1]
 1988 and Fig. 4.1) known as “G function” [ANWOW13, Section 2.4]. G is based on ChaCha
 1989 encryption [Ber08a]. It works on w -bit long words, and presents $8 \cdot w + 10$ constraints.
 1990 The function mixes a state (a , b , c and d) with the inputs (x and y) and returns the
 1991 updated state.

1992 This circuit does not check the booleanness of the inputs or state. However, given that
 1993 the state is boolean, the output is automatically boolean due to the use of the modular
 1994 addition circuits.

$G(a, b, c, d; x, y) \mapsto (a_2, b_2, c_2, d_2)$	<code>getSigma()</code>
1 : $a_1 \leftarrow a + b + x \pmod{2^w}$	1 : $\Sigma \in (\mathbb{N}^{16})^{10}$
2 : $d_1 \leftarrow d \oplus a_1 \ggg r_1$	2 : $\Sigma[0] \leftarrow (0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15)$
3 : $c_1 \leftarrow c + d_1 \pmod{2^w}$	3 : $\Sigma[1] \leftarrow (14, 10, 4, 8, 9, 15, 13, 6, 1, 12, 0, 2, 11, 7, 5, 3)$
4 : $b_1 \leftarrow b \oplus c_1 \ggg r_2$	4 : $\Sigma[2] \leftarrow (11, 8, 12, 0, 5, 2, 15, 13, 10, 14, 3, 6, 7, 1, 9, 4)$
5 : $a_2 \leftarrow a_1 + b_1 + y \pmod{2^w}$	5 : $\Sigma[3] \leftarrow (7, 9, 3, 1, 13, 12, 11, 14, 2, 6, 5, 10, 4, 0, 15, 8)$
6 : $d_2 \leftarrow d_1 \oplus a_2 \ggg r_3$	6 : $\Sigma[4] \leftarrow (9, 0, 5, 7, 2, 4, 10, 15, 14, 1, 11, 12, 6, 8, 3, 13)$
7 : $c_2 \leftarrow c_1 + d_2 \pmod{2^w}$	7 : $\Sigma[5] \leftarrow (2, 12, 6, 10, 0, 11, 8, 3, 4, 13, 7, 5, 15, 14, 1, 9)$
8 : $b_2 \leftarrow b_1 \oplus c_2 \ggg r_4$	8 : $\Sigma[6] \leftarrow (12, 5, 1, 15, 14, 13, 4, 10, 0, 7, 6, 3, 9, 2, 8, 11)$
9 : return a_2, b_2, c_2, d_2	9 : $\Sigma[7] \leftarrow (13, 11, 7, 14, 12, 1, 3, 9, 5, 0, 15, 4, 8, 6, 2, 10)$
	10 : $\Sigma[8] \leftarrow (6, 15, 14, 9, 11, 3, 0, 8, 12, 2, 13, 7, 1, 4, 10, 5)$
	11 : $\Sigma[9] \leftarrow (10, 2, 8, 4, 7, 6, 1, 5, 15, 11, 9, 14, 3, 12, 13, 0)$
	12 : return Σ

Figure 4.1: G primitive [MJS15, Section 3.1]

Figure 4.2: Blake2 permutation table [MJS15, Section 2.7]

1995 For Blake2s, we have $w = 32$, $r_1 = 16$, $r_2 = 12$, $r_3 = 3$ and $r_4 = 7$.

1996 Blake2s compression function circuit

The compression function is defined as follows, for more details see Fig. 4.3,

$$\text{Blake2sC} : \mathbb{B}^n \times \mathbb{B}^{2n} \times \mathbb{B}^{n/4} \times \mathbb{B}^{n/4} \rightarrow \mathbb{B}^n.$$

1997 Blake2C takes as input a state $h \in \mathbb{B}^n$ which is used as chaining value when hashing,
 1998 a message to compress $x \in \mathbb{B}^{2n}$, a message length written in binary $t \in \mathbb{B}^{n/4}$ which is
 1999 incremented when hashing and a binary flag $f \in \mathbb{B}^{n/4}$ to tell whether the current block
 2000 is the last to be compressed to prevent length extension attacks.

2001 Blake2C uses the G function iteratively over **rounds** number of rounds on a state
 2002 and message. The constant initialization vector IV and the permutation table Σ are
 2003 hard-coded. Blake2sC works in little endian (see [MJS15, Section 2.4]) on n -bit long
 2004 variables ($n = 256$), w -bit long words ($w = 32$), and the rotation constants specified
 2005 in Section 4.3.6 (see [MJS15, Section 2.1]). We have the following constants (see speci-
 2006 fications [ANWOW13] and [MJS15, Section 2.2]),

- 2007 • IV is the $(8 \cdot w)$ -bit long initialization vector; it corresponds to the first w bits of the
 2008 fractional parts of the square roots of the first eight prime numbers $(2, 3, 5, 7, \dots)$
 2009 (see [MJS15, Section 2.6]);
- 2010 • Σ are the $10 \cdot 16$ permutation constants of Blake2 (see Fig. 4.2 and [MJS15, Section
 2011 2.7]);
- 2012 • **rounds**, the number of rounds: 10 for Blake2sC, 12 for Blake2bC.

2013 We have the following variables (see specifications [ANWOW13] and [MJS15, Section
2014 2.2]),

- 2015 • H is the $(8 \cdot w)$ -bit long initial state while v is the $(16 \cdot w)$ -bit long final state;
- 2016 • $T[i]$ are two w -bit long counters encoding the block length;
- 2017 • $F[i]$ are two w -bit long finalization flags. We set the first one $F[0]$ to $2^w - 1$ to state
2018 when the input block is the last one to be hashed. The second, $F[1] = 0$ is only set
2019 for tree hashing mode (which is not our case) and is therefore unused.

2020 We introduce the following functions to write Blake2C (see specifications [ANWOW13]
2021 and [MJS15, Section 2.6]):

- 2022 • The function `prime` takes a positive integer i as input and outputs the i -th prime
2023 number;
- 2024 • The function `dec` takes a real number x as input outputs its positive decimal part.

2025 This circuit presents $((64 \cdot \text{rounds} + 8) \cdot w + 8 \cdot \text{rounds} + 10)$ constraints. For Blake2sC,
2026 as $w = 32$ and $\text{rounds} = 10$, we have 21536 constraints.

2027 We do not check the input block booleanness in this circuit. Given that the initial
2028 state is boolean, the output is automatically boolean. This can be proved iteratively by
2029 the booleanness of G primitive's output.

2030 **Security requirement.** The inputs to Blake2sC MUST be boolean.

2031 Blake2s hash function

The hash function is defined as follows, for more details see Fig. 4.3,

$$\text{Blake2s} : \mathbb{B}^{\leq 2n} \times \mathbb{B}^* \rightarrow \mathbb{B}^n$$

2032 Blake2 takes as input a hash key $k \in \mathbb{B}^n$ and the message to hash $x \in \mathbb{B}^{2n}$. Blake2
2033 uses the Blake2C function iteratively over each $2n$ -bit long chunk of the padded message.
2034 If the key is non null, it is used as the first block to be hashed. The constant initialization
2035 vector IV and part of the parameter block PB are hard-coded. We have the following
2036 constants (see specifications [ANWOW13] and [MJS15, Section 2.2]),

- 2037 • IV is the $(8 \cdot w)$ -bit long Initialization Vector; it corresponds to the first w bits of the
2038 fractional parts of the square roots of the first eight prime numbers $(2, 3, 5, 7, \dots)$
2039 (see [MJS15, Section 2.6]).

2040 We have the following variables (see specifications [ANWOW13] and [MJS15, Section
2041 2.2]),

Blake2C(h, m, t, f)

```

1 :  $T, F, H, IV, v \in (\mathbb{B}^w)^2 \times (\mathbb{B}^w)^2 \times (\mathbb{B}^w)^8 \times (\mathbb{B}^w)^8 \times (\mathbb{B}^w)^{16}$ 
2 :  $\{IV[i]\}_{i \in [8]} \leftarrow \left\{ \left[ 2^w \cdot \text{dec}(\sqrt{\text{prime}(i+1)}) \right] \right\}_{i \in [8]}$ 
3 :  $\Sigma \leftarrow \text{getSigma}()$ 
4 :  $\{H[i]\}_{i \in [8]} \leftarrow \{h[i \cdot w : (i+1) \cdot w]\}_{i \in [8]}$ 
5 :  $\{m[i]\}_{i \in [8]} \leftarrow \{x[i \cdot w : (i+1) \cdot w]\}_{i \in [8]}$ 
6 :  $T[0], T[1] \leftarrow t[w:2w], t[0:w]$ 
7 :  $F[0], F[1] \leftarrow f[w:2w], f[0:w]$ 
8 :  $\{v[i]\}_{i \in [8]} \leftarrow \{H[i]\}_{i \in [8]}$ 
9 :  $\{v[i+8]\}_{i \in [8]} \leftarrow \{IV[i]\}_{i \in [8]}$ 
10 :  $v[12], v[13] \leftarrow v[12] \oplus T[0], v[13] \oplus T[1]$ 
11 :  $v[14], v[15] \leftarrow v[14] \oplus F[0], v[15] \oplus F[1]$ 
12 : foreach  $r \in [\text{rounds}]$  do
13 :    $\tau \leftarrow \Sigma[r \pmod{15}]$ 
14 :    $v[0], v[4], v[8], v[12] \leftarrow G(v[0], v[4], v[8], v[12], m[\tau[0]], m[\tau[1]])$ 
15 :    $v[1], v[5], v[9], v[13] \leftarrow G(v[1], v[5], v[9], v[13], m[\tau[2]], m[\tau[3]])$ 
16 :    $v[2], v[6], v[10], v[14] \leftarrow G(v[2], v[6], v[10], v[14], m[\tau[4]], m[\tau[5]])$ 
17 :    $v[3], v[7], v[11], v[15] \leftarrow G(v[3], v[7], v[11], v[15], m[\tau[6]], m[\tau[7]])$ 
18 :    $v[0], v[5], v[10], v[15] \leftarrow G(v[0], v[5], v[10], v[15], m[\tau[8]], m[\tau[9]])$ 
19 :    $v[1], v[6], v[11], v[12] \leftarrow G(v[1], v[6], v[11], v[12], m[\tau[10]], m[\tau[11]])$ 
20 :    $v[2], v[7], v[8], v[13] \leftarrow G(v[2], v[7], v[8], v[13], m[\tau[12]], m[\tau[13]])$ 
21 :    $v[3], v[4], v[9], v[14] \leftarrow G(v[3], v[4], v[9], v[14], m[\tau[14]], m[\tau[15]])$ 
22 : return  $\|_{i=0}^8 H[i] \oplus v[i] \oplus v[i+8]$ 

```

Figure 4.3: Blake2 compression function [MJS15, Section 3.2]. Set n , w and G 's constants to obtain Blake2sC.

- 2042 • PB is the $(16 \cdot w)$ -bit long parameter block used to initialize the state (see [MJS15,
2043 Section 2.5]). In big endian encoding, the first byte corresponds to the digest
2044 length (fixed to 32 bytes), the second byte to the key length, the third and fourth
2045 bytes correspond to the use of the serial mode;
- 2046 • $H \in \mathbb{B}^{\text{BLAKE2sCLEN}}$, the chaining value.

2047 We do not check the input block booleaness in this circuit. Given that the initial
2048 state is boolean, the output is automatically boolean. This can be proved iteratively by
2049 the booleaness of Blake2C primitive's output.

2050 **Security requirement** To ensure the correct use of Blake2s, Blake2s's inputs MUST be
2051 boolean.

Blake2(k, x)

```

1 : H, IV, PB ∈ ℤ8w × ℤ8w × ℤ8w
2 : PB ← pad8·w(encodeℕ(0x0101)) || padw(encodeℕ(⌈length(k)/BYTELEN⌉)) || encodeℕ(0x20)
3 : IV ← ||i=08 [ 2w · dec(√prime(i+1)) ]
4 : H ← PB ⊕ IV
5 : y ← x
6 : if length(k) ≠ 0 do
7 :   y ← pad2n(k) || y
8 :   z ← pad2n·⌈length(y)/2n⌉(y)
9 :   for i ∈ [⌈length(z)/2n⌉] do
10 :    if i = ⌈length(z)/2n⌉ - 1 do
11 :      H ← Blake2C(H, z[i · 2n:(i+1) · 2n], pad2w(encodeℕ(⌈length(y)/BYTELEN⌉)), pad2w(encodeℕ(2w - 1)))
12 :    else
13 :      H ← Blake2C(H, z[i · 2n:(i+1) · 2n], pad2w(encodeℕ((i+1) · 2n/BYTELEN)), pad2w(0))
14 :   return H

```

Figure 4.4: Blake2 hash function [MJS15, Section 3.3]. Set $n = 16w$ and G's constants accordingly to obtain Blake2s.

2052 Optimizing the circuits

2053 The above helper circuits form the building blocks of the Blake2s compression function.
 2054 We show here two exclusive methods to optimize these circuits.

2055 Optimizing the Modular additions

Double modular addition: $a + b = c \pmod{2^w}$. We present here an optimization on the circuit to save one constraint by merging the modular constraint with a boolean constraint. The optimized circuit presents the following constraints:

$$\left(\sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i \right) \cdot \left(\sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i - 2^w \right) = 0 \quad (4.5)$$

$$\forall j \in \{0, \dots, w-1\}, (c_j - 0) \cdot (c_j - 1) = 0 \quad (4.6)$$

2056 with $\sum_{i=0}^{w-1} x_i \cdot 2^i$ a binary encoding of x (x_i is the i^{th} bit of x).

2057 These equations describe $w + 1$ constraints to prove the bit equality $a + b = c$ (note
 2058 that an additional $2 \cdot w$ constraints would be required to prove the booleanness of input
 2059 variables a and b). We now explain how we obtained them.

2060 *Proof.* The most straightforward way to prove that $a + b = c \pmod{2^w}$ and c booleanness
 2061 is with the set of constraints illustrated in Eq. (4.1) and in Eq. (4.2).

As we perform arithmetic modulo 2^w , we do not care about the value of c_w but would like to ensure its booleanness. As one may notice, the summing constraint Eq. (4.1) is an equality of two linear combinations with no multiplication by a variable. Hence, we can combine it with the boolean constraint of c_w to remove any reference to c_w and still have a bilinear gate. To do so, we first rewrite Eq. (4.1) as an equality check over $c_w \cdot 2^w$ and multiply Eq. (4.2) for $j = n$ by $2^{2 \cdot w}$.

$$\sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i = c_w \cdot 2^w \quad (4.7)$$

$$2^w \cdot (c_w - 0) \cdot 2^w \cdot (c_w - 1) = 0 \quad (4.8)$$

We finally replace $c_w \cdot 2^w$ in Eq. (4.8) by the value from Eq. (4.7).

$$\begin{aligned} 0 &= 2^w \cdot (c_w - 0) \cdot 2^w \cdot (c_w - 1) = 2^w \cdot c_w \cdot (2^w \cdot c_w - 2^w) \\ &= \left(\sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i \right) \cdot \left(\left(\sum_{i=0}^{w-1} (a_i + b_i - c_i) \cdot 2^i \right) - 2^w \right) \end{aligned}$$

2062 This results in Eq. (4.5) and Eq. (4.6). All references to c_w have disappeared and, with
2063 a single multiplication by a variable, we still have bilinear gates. \square

2064 **Triple modular addition: $a + b + c = d \pmod{2^w}$.** To optimize, we use the
2065 above circuit twice. We define a temporary variable d' such that $a + b = d' \pmod{2^w}$.
2066 As such, we have $c + d' = d \pmod{2^w}$. As d' is the addition of two w -bit long variables,
2067 it is $(w + 1)$ -bit long. However as we evaluate the sum modulo 2^w , we discard the last bit
2068 of d' . We proceed similarly for d . To ensure that d is boolean, we check the booleanness
2069 of the $w + 1$ bits of d as well as the booleanness of the last bit of d' (to account for d 's
2070 $w + 2^{\text{th}}$ bit in the original expression ($a + b + c = d \pmod{2^w}$)).

We thus obtain the following circuit with $w + 2$ constraints,

$$\begin{aligned} \left(\sum_{i=0}^{w-1} (a_i + b_i - d'_i) \cdot 2^i \right) \cdot \left(\sum_{i=0}^{w-1} (a_i + b_i - d'_i) \cdot 2^i - 2^w \right) &= 0 \\ \left(\sum_{i=0}^{w-1} (c_i + d'_i - d_i) \cdot 2^i \right) \cdot \left(\sum_{i=0}^{w-1} (c_i + d'_i - d_i) \cdot 2^i - 2^w \right) &= 0 \\ \forall j \in \{0, \dots, w - 1\}, (d_j - 0) \cdot (d_j - 1) &= 0 \end{aligned}$$

2071 These optimizations lead to a gain of 320 constraints ($= 4 \cdot 8 \cdot \text{rounds}$).

2072 **Optimizing Blake2s routine's circuit** As seen in Fig. 4.1, our routine presents 2
2073 double and 2 triple modular additions. Each of these circuits comprises at least one
2074 modular constraint which pack several w -bit long variables. The circuit is however
2075 processed in \mathbb{F}_{CUR} , that is to say most integers can be written over FIELDCAP bits. We

2076 can thus batch the modular constraints. As the G primitive performs 2 double modular
2077 and 2 triple modular, we have in total 6 modular checks per iteration. We can batch up
2078 to $\text{FIELD CAP}/w$ constraints together. For $w = 32$ and $\text{FIELD CAP} \geq 224$ (which holds for
2079 BN-254 and BLS12-377), we can encode up to 7 words per field element, that is to say
2080 we can include all the modular constraints into a single one.

2081 This optimization leads to a gain of 274 constraints ($= 4 \cdot 8 \cdot 10 - \lceil \frac{4 \cdot 8 \cdot 10}{7} \rceil$).

2082 **Optimization conclusion** Using the more efficient optimization on the modular ad-
2083 ditions, the Blake2s compression function comprises 21216 constraints.

2084 **Increasing the PRF security with Blake**

2085 As Blake2 comprises a personalization tag in its parameter block PB , one could ensure the
2086 independence of the PRFs by writing different tags for each of them (we would be able to
2087 consider up to 2^{30} inputs and outputs). We did not choose to write this enhancement in
2088 the instantiation to keep a general tagging method in case of a change of hash function.

2089 **4.4 Encryption of the notes**

2090 In this section we give some remarks concerning the implementation of the Zeth en-
2091 crypton scheme, described in Section 3.5. As noted, there are several details in the
2092 specification of the underlying primitives which can impact security if not carefully im-
2093 plemented. The following list is by no means exhaustive but includes several details
2094 noted during development of the proof-of-concept system.

- 2095 • Private keys for Curve25519 **MUST** be randomly generated as 32 bytes where the
2096 first byte is a multiple of 8, and the last byte takes a value between 64 and 127
2097 (inclusive). Further details are given in [Ber06], including an example algorithm
2098 for generation. Implementations **MUST** take care to ensure that their code, or any
2099 external libraries they rely upon, correctly perform this step.
- 2100 • A similar observation holds for Poly1305 in which the r component of the MAC
2101 key (r, s) **MUST** be *clamped* in a specific way (see Section 3.5.3). This step is also
2102 essential and **MUST** be performed.
- 2103 • In the implementation of the ChaCha stream cipher, correct use of the *key*, *counter*
2104 and *nonce* **MUST** be ensured in order to adhere to the standard and guarantee the
2105 appropriate security properties.

2106 During the proof-of-concept implementation it was not obvious that the cryptogra-
2107 phy library⁴ adhered to the specifications with respect to the above points. In particular,
2108 it was not clear whether key clamping was performed at generation time and/or when

⁴<https://cryptography.io/en/latest/>

2109 performing operations. Moreover, the interface to the ChaCha cipher accepted a differ-
2110 ent set of input parameters (namely *key* and *nonce* with no *counter*). This left some
2111 ambiguity about the responsibility for clamping, and whether the ChaCha block data
2112 would be updated as described in the specification. Details of how this was resolved are
2113 given in the proof-of-concept encryption code, which may prove a useful reference for
2114 implementers⁵.

⁵see <https://github.com/clearmatics/zeth/blob/v0.4/client/zeth/encryption.py>

²¹¹⁵ Acknowledgements

²¹¹⁶ We thank Lorenzo Grassi for insightful comments on the security of MiMC.

2117 Appendix A

2118 Transaction non malleability

2119 The transaction malleability problem for a DAP (Section 1.4) is characterized by a game
2120 TR-NM involving a polynomial-time adversary \mathcal{A} as described below.

Definition A.0.1. Let DAP be a (candidate) Decentralized Anonymous Payment scheme.

$$\text{DAP} = (\text{Setup}, \text{GenAddr}, \text{SendTx}, \text{VerifyTx}, \text{Receive})$$

We say that DAP is TR-NM secure if, for every $\text{poly}(\lambda)$ -time adversary \mathcal{A}

$$\text{Adv}_{\text{DAP}, \mathcal{A}}^{\text{tr-nm}}(\lambda) < \text{negl}(\lambda),$$

2121 where $\text{Adv}_{\text{DAP}, \mathcal{A}}^{\text{tr-nm}}(\lambda) = \Pr[\text{TR-NM}(\text{DAP}, \mathcal{A}, \lambda) = 1]$ is \mathcal{A} 's advantage in the TR-NM exper-
2122 iment.

2123 Below, we adapt [BSCG⁺14, Appendix C.2] to our specific DAP—Zeth.

2124 We start by describing the TR-NM experiment. Given a (candidate) Zeth DAP,
2125 adversary \mathcal{A} , and security parameter λ , the (probabilistic) game TR-NM(DAP, \mathcal{A} , λ)
2126 consists of an interaction between \mathcal{A} and a challenger \mathcal{C} , terminating with a binary
2127 output by \mathcal{C} .

2128 At the beginning of the game, \mathcal{C} samples $pp \leftarrow \text{Setup}(\lambda)$ and sends pp to \mathcal{A} . Next, \mathcal{C}
2129 initializes a DAP oracle O^{DAP} with pp and allows \mathcal{A} to issue queries to it [RZ19, Appendix
2130 B].

2131 At the end of the experiment, \mathcal{A} sends to \mathcal{C} a **Mixer** contract call transaction tx_{Mixer}^* ,
2132 and \mathcal{C} outputs 1 iff the following conditions hold. Letting T be the set of transactions
2133 generated by O^{DAP} in response to **SendTx** queries, there exists $tx_{\text{Mixer}} \in T$ such that:

- 2134 1. tx_{Mixer} was not inserted in L by \mathcal{A} ;
- 2135 2. $tx_{\text{Mixer}}^*.data \neq tx_{\text{Mixer}}.data$;
- 2136 3. $\text{VerifyTx}(pp, tx_{\text{Mixer}}^*, L') = 1$ where L' is the portion of the ledger L preceding tx_{Mixer} ;
- 2137 4. a serial number revealed in tx_{Mixer}^* is also revealed in tx_{Mixer} .

2138 A.1 Transaction malleability attack on Zeth

2139 In this section we present the threat related to the transaction malleability attack on
 2140 Zeth and expose the solutions by ZeroCash [BSCG⁺14] and Zcash [ZCa19] that we
 2141 adapted.

2142 First, we start by assuming that none of the checks related to transaction malleability
 2143 attack have been added in the protocol Chapter 2. As such, we assume that *hsig* and
 2144 *htags* are not attributes of `PrimInputDType`, ϕ is not an attribute of `AuxInputDType`,
 2145 and *otssig* and *otsvk* are not attributes of the `MixInputDType` data type anymore. As
 2146 a consequence, all checks related to these attributes are removed from the protocol.
 2147 Moreover, if *zn* is an object of type `ZethNoteDType`, then *zn.ρ* is chosen at random.
 2148 Finally, the NP-relation used in Zeth, now denoted \mathbf{R}^{mal} , becomes the following:

- 2149 • For each $i \in [\text{JSIN}]$:
 - 2150 1. $aux.jsins[i].znote.apk = \text{Blake2s}(tag_{ask}^{addr} \parallel \text{pad}_{\text{BLAKE2sCLEN}}(0))$
 2151 with tag_{ask}^{addr} defined in Section 3.1.3
 - 2152 2. $aux.jsins[i].nf = \text{Blake2s}(tag_{ask}^{nf} \parallel aux.jsins[i].znote.\rho)$
 2153 with tag_{ask}^{nf} defined in Section 3.1.3
 - 2154 3. $aux.jsins[i].cm = \text{Blake2s}(aux.jsins[i].znote.r \parallel m)$
 2155 with $m = aux.jsins[i].znote.apk \parallel aux.jsins[i].znote.\rho \parallel aux.jsins[i].znote.v$
 - 2156 4. $(aux.jsins[i].znote.v) \cdot (1 - e) = 0$ is satisfied for the boolean value e set such
 2157 that if $aux.jsins[i].znote.v > 0$ then $e = 1$.
 - 2158 5. The Merkle root $mkroot'$ obtained after checking the Merkle authentication
 2159 path $aux.jsins[i].mkpath$ of commitment $aux.jsins[i].cm$, with MKHASH, is
 2160 equal to $prim.mkroot$ if $e = 1$.
 - 2161 6. $prim.nfs[i]$
 2162 $= \left\{ \text{Pack}_{\mathbb{F}_{\text{rCUR}}} (aux.jsins[i].nf[k \cdot \text{FIELD CAP} : (k + 1) \cdot \text{FIELD CAP}]) \right\}_{k \in [\lfloor \text{PRFNFOULEN} / \text{FIELD CAP} \rfloor]}$
- 2163 • For each $j \in [\text{JSOUT}]$:
 - 2164 1. $prim.cms[j] = \text{Blake2s}(aux.znotes[j].r \parallel m)$
 2165 with $m = aux.znotes[j].apk \parallel aux.znotes[j].\rho \parallel aux.znotes[j].v$
- 2166 • $prim.rsd = \text{Pack}_{rsd}(\{aux.jsins[i].nf\}_{i \in [\text{JSIN}]}, aux.vin, aux.vout)$
- Check that the “joinsplit is balanced”, i.e. check that the joinsplit equation holds:

$$\begin{aligned}
 & \text{Pack}_{\mathbb{F}_{\text{rCUR}}}(aux.vin) + \sum_{i \in [\text{JSIN}]} \text{Pack}_{\mathbb{F}_{\text{rCUR}}}(aux.jsins[i].znote.v) \\
 &= \sum_{j \in [\text{JSOUT}]} \text{Pack}_{\mathbb{F}_{\text{rCUR}}}(aux.znotes[j].v) + \text{Pack}_{\mathbb{F}_{\text{rCUR}}}(aux.vout)
 \end{aligned}$$

2167 A.1.1 The attack

2168 In order to win the game TR-NM on the weak Zeth DAP above, an adversary \mathcal{A} inter-
 2169 cepts a target transaction tx_{Mix} by passively listening to the network (remember that
 2170 transactions are broadcasted to the Ethereum network in order to be mined, see Sec-
 2171 tion 1.2.2), extracts the zk-proof and primary inputs from $tx_{\text{Mix}}.data$ and uses these
 2172 extracted pieces of information in order to create a malicious transaction tx_{Mix}' , where
 2173 the ciphertexts are replaced by arbitrary data. The adversary can then broadcast tx_{Mix}'
 2174 to the network in order for it to be mined. If the malicious transaction gets mined before
 2175 the legitimate one, the input notes become spent and the ciphertexts are undecryptable
 2176 making the new notes unredeemable (by any Zeth user!), since all attempts to decrypt
 2177 the ciphertexts will fail (see Section 2.6).

```

TxMalGen( $sk'_{\text{ECDSA}}, nce_{in}, tx_{\text{Mix}}$ )
-----
1 :  $p \leftarrow tx_{\text{Mix}}.gasP + 1$ 
2 :  $l \leftarrow tx_{\text{Mix}}.gasL + 1$ 
3 :  $zdata' \leftarrow tx_{\text{Mix}}.data$ 
4 :  $zdata'.ciphers \leftarrow \mathbb{B}^*$ 
5 :  $tx_{raw} \leftarrow \{nce : nce_{in}, gasP : p, gasL : l, to : tx_{\text{Mix}}.to, val : tx_{\text{Mix}}.val, data : zdata'\};$ 
6 :  $\sigma_{\text{ECDSA}} \leftarrow \text{SigSch}_{\text{ECDSA}}.\text{Sig}(sk'_{\text{ECDSA}}, \text{Keccak256}(tx_{raw}));$ 
7 :  $tx_{final} \leftarrow \{tx_{raw}, v : \sigma_{\text{ECDSA}}.v', r : \sigma_{\text{ECDSA}}.r', s : \sigma_{\text{ECDSA}}.s'\};$ 
8 : return  $tx_{final}$ ;

```

Figure A.1: Transaction malleability attack function TxMalGen

2178 As shown on Fig. A.1, during the attack, the adversary extracts the proof and pri-
 2179 mary inputs from the honest transaction, and replaces the ciphertexts by some arbitrary
 2180 information. The attacker then formats this data into a transaction that calls the Mix
 2181 function of **Mixer**, and submits it to the network. While the data fields ($tx_{\text{Mix}}.data$
 2182 and $tx_{\text{Mix}}'.data$) are different, the nullifiers revealed by both transactions are the same
 2183 (i.e. $tx_{\text{Mix}}.data.proof = tx_{\text{Mix}}'.data.proof$, and $tx_{\text{Mix}}.data.prim = tx_{\text{Mix}}'.data.prim$).
 2184 As a consequence, if the adversary makes sure that tx_{Mix}' satisfies all the checks of
 2185 EthVerifyTx (Section 1.2.2), he can ensure that ZethVerifyTx(tx_{Mix}') will return the same
 2186 value as ZethVerifyTx(tx_{Mix}). Furthermore, if $tx_{\text{Mix}}'.gasP > tx_{\text{Mix}}.gasP$, then the adver-
 2187 sary maximizes his chances of having his transaction mined first (Section 1.2.2), and so
 2188 maximizes the chances for the malleability attack to be successful; leading to lost funds
 2189 on **Mixer**.

2190 **Remark A.1.1.** Note that, although not directly contained within the *data* field of a
 2191 **Mixer** call transaction, the Ethereum address $\mathcal{S}_{\mathcal{E}}.Addr$ of the transaction sender is also
 2192 used by the **Mixer** call (this is either the calling contract's address, or the transaction
 2193 signer's address recovered as described in Remark 1.2.1). In particular, the balance
 2194 of this Ethereum address is incremented by the value *vout* by successful Mix calls. If

2195 we again assume the absence of the malleability checks, an attacker could re-sign any
2196 **Mixer** call transaction with a key under his control, rebroadcast it as described above,
2197 and (with some reasonable probability) become the recipient of any public output value
2198 $vout$.

2199 **Remark A.1.2.** We note that the attack described above cannot be prevented by
2200 merely substituting a malleable Groth16 zk-SNARK by a simulation-extractable one
2201 like e.g. [GM17]. This comes since the attack does not utilise malleability of the proof
2202 system, but malleability of data that are broadcasted along with the zk-proof.

2203 A.2 Solutions to address the transaction malleability at- 2204 tack

2205 A.2.1 ZeroCash solution

2206 The idea of the solution presented in [BSCG⁺14] is to use a one-time SUF-CMA digital
2207 signature and bind its verification key with the zk-proof primary inputs to prevent an
2208 adversary from corrupting part of a transaction's data.

2209 Specifically, to transact via **Zeth**, the user first samples a key pair (sk, vk) for a one-
2210 time signature scheme. He then computes the hash $hsig = CRH(vk)$, where CRH is a
2211 collision resistant hash function, see [BSCG⁺14], and derives a value $h_i = PRF_{ask_i}^{pk}(hsig)$,
2212 for each input note (i.e. $i \in [JSIN]$), which acts as a MAC binding $hsig$ to the address
2213 spending key of a note (ask_i) .

2214 The user then generates the zk-proof with the additional statement that the values
2215 $\{h_i\}_{i \in [JSIN]}$ are computed correctly. He finally uses sk to sign every value associated with
2216 the operation, thus obtaining a signature, which is included, along with the signature
2217 verification key vk , in the transaction. To verify a transaction on the DAP, it is necessary
2218 to verify that

- 2219 • the primary inputs are correctly formatted,
- 2220 • the Merkle root corresponds to one of the previous states of the Merkle tree,
- 2221 • the nullifiers have not been declared in a previous transaction,
- 2222 • the $hsig$ is correctly computed from vk , and
- 2223 • both the zk-proof and the one-time signature verifications pass successfully.

2224 Now, an adversary trying to carry out the aforementioned attack has to either change
2225 the ciphertexts or the encryption key. Nevertheless, doing so should lead to the one-time
2226 signature verification to fail or should yield an attack that breaks the UF-CMA property
2227 of the one-time signature (as this corresponds to creating a forgery on a different message,
2228 not changing the signature). Thereby, the adversary also has to modify the signature,
2229 however he does not know the one-time signing key used by the creator of the targeted
2230 transaction. As such, the adversary needs to use another signing key pair, however

2231 this leads to the check verifying that $hsig$ is correctly computed to fail. If the adversary
2232 attempts to change $hsig$, the zk-proof verification fails as the NP-statement has changed.
2233 Hence, any attempt to carry out a malleability attack results in the violation of at least
2234 one check in the verification of the transaction on the DAP. The solution presented
2235 effectively solves the transaction-malleability attack initially described.

2236 **Remark A.2.1.** The one-timeness property of the signature scheme was required in
2237 **ZeroCash** to retain anonymity. It also makes analysing non-adaptive adversary sufficient.
2238 As **Ethereum** transaction senders need to pay the gas cost associated with their trans-
2239 actions, the senders are not anonymous. This said, making sure that **Zeth** is designed
2240 with anonymity in mind is worth the effort in order to minimize information leakages
2241 and be ready if/when **Ethereum** incorporates protocol changes that enable anonymous
2242 transactions.

2243 A.2.2 Zcash’s solution

2244 In addition to the changes aforementioned, **Zcash**’s solution [ZCa19] also consists of:

- Redefining the variable $hsig$ as,

$$hsig = \text{CRH}(\text{randomSeed}, \{nf_i\}_{i \in [\text{JSIN}]}, vk)$$

2245 for some random seed randomSeed .

- Defining a new random variable ϕ and using it with $hsig$, as key and input of a
2246 PRF respectively, to compute the identifier of each output notes ρ_j ($j \in [\text{JSOUT}]$)
2247 and ensure their uniqueness (with overwhelming probability).
2248

2249 These changes were made to prevent the Faerie Gold attack [ZCa19, Section 8.4], as well
2250 as to prevent linkability: if $hsig$ were repeated in two transactions, the circuit would
2251 leak, via $\{h_i\}_{i \in [\text{JSIN}]}$, the fact that the input notes in both transactions were spent with
2252 the same ask_i (if that were the case).

2253 More particularly, using the input notes’ nullifiers to derive $hsig$ ensures that $hsig$ is
2254 unique with overwhelming probability for all *accepted* transaction. Furthermore, us-
2255 ing randomSeed ensures the uniqueness of $hsig$ for transactions *in transit* (as before
2256 validation there may be several in transit transactions with the same set of nullifiers).

2257 A.2.3 Solution on Ethereum

2258 As described in the **Ethereum** yellow paper [Woo19, Appendix F], **Ethereum** transactions
2259 are ECDSA signed. Further, as described in Section 2.3, the one-time signature used to
2260 sign the Mix data also signs the **Ethereum** address used to sign the transaction. As such,
2261 any modification to the transaction object will result in a new transaction hash, and
2262 any attempt to sign the transaction with a different ECDSA key will be rejected by the
2263 **Mixer** contract (see Section 2.5). We thereby conclude that the one-time signature used

2264 to sign the transaction data does not need to be SUF-CMA, but *only needs to achieve*
2265 UF-CMA.

2266 Specifically, carrying out any change on the one-time signature will change the
2267 **Ethereum** transaction data and result in a failure to verify the ECDSA signature of
2268 the **Ethereum** transaction. To obtain a new valid signature on this transaction, the
2269 adversary needs to break the UF-CMA property of the ECDSA signature scheme or use
2270 another ECDSA keypair to sign the transaction. In the last case, the one-time signature
2271 will no longer be valid.

2272 Note that including the **Ethereum** transaction sender in the data to be signed by the
2273 one-time signature scheme also addresses the possible attack described in Remark A.1.1.
2274 An attacker trying to resign the same **Ethereum** transaction with a different key will
2275 cause **Mixer** to reject the transaction when the one-time signature is checked.

Remark A.2.2. We note that the transaction malleability issue can also be addressed
in another way. In fact, one could use the ECDSA signatures on **Ethereum** transactions
to fix all inputs and ciphertexts, and then tie the sender of the **Ethereum** transaction to
the zk-snark by putting the sender address $\mathcal{S}_{\mathcal{E}}.Addr$ in $hsig$. In other words, it is also
possible to define $hsig$ as:

$$hsig = \text{CRH}(\{nf_i\}_{i \in [\text{JSIN}]}, \mathcal{S}_{\mathcal{E}}.Addr)$$

2276 As such, if an attacker extracts the ciphertexts of a tx_{Mix} transaction in order to craft
2277 another malicious transaction $tx_{\text{Mix}'}$, the key-pair used to sign $tx_{\text{Mix}'}$ differs from the one
2278 used to sign tx_{Mix} , which changes the transaction sender address recovered on **Mixer**. As
2279 a consequence, the check on $hsig$ would fail on the **Mixer**, invalidating the transaction,
2280 and preventing the attack.

2281 While such a solution would avoid the need to generate one-time signing keys and
2282 could avoid a signature check in the **Mixer**, it would also require every **Zeth** user to
2283 have an **Ethereum** account. Doing so, would be a major hindrance toward the design of
2284 mechanisms aiming to provide anonymity to **Zeth** transactions initiators. In fact, the
2285 addressing scheme used in **Zeth** along with the solution to the malleability introduced in
2286 **Zcash** makes it possible to generate raw **Zeth** transactions without having an **Ethereum**
2287 account. These raw transactions could then be broadcasted – to a set of **Ethereum** user
2288 nodes – on an anonymous p2p network, before being finalized and submitted to the
2289 **Ethereum** network by **Ethereum** users who would be rewarded according to an incentive
2290 structure. While such a protocol is outside of the scope of this document, it shows that
2291 defining $hsig$ using the senders address alters the flexibility of **Zeth**; hence this solution
2292 has not been favoured.

2293 Appendix B

2294 Double spend attack on 2295 equivalent class

2296 The primary inputs of our zk-SNARK are elements of $\mathbb{F}_{r_{\text{CUR}}}$ and they can be written over
2297 **FIELDLEN** bits. Note that the projection of $\mathbb{B}^{\text{FIELDLEN}}$ onto $\mathbb{F}_{r_{\text{CUR}}}$ formed by interpreting
2298 elements in $\mathbb{B}^{\text{FIELDLEN}}$ as **FIELDLEN**-bit numbers and reducing modulo r_{CUR} , is surjective.

2299 When we pass the primary inputs to the **Mixer** contract, they are interpreted as
2300 elements of $\mathbb{B}^{\text{ETHWORDLEN}}$, and $\mathbb{B}^{\text{FIELDLEN}} \subset \mathbb{B}^{\text{ETHWORDLEN}}$. As previously noted, this means
2301 that there exist pairs of elements in $\mathbb{B}^{\text{ETHWORDLEN}}$ with the same projection in $\mathbb{F}_{r_{\text{CUR}}}$. An
2302 adversary could make use of this to perform a double spend attack.

2303 Indeed, to check that a note is not double spent, the contract stores the nullifiers of
2304 spent notes (as elements of $\mathbb{B}^{\text{ETHWORDLEN}}$) and verifies that the nullifier of the note to be
2305 spent is not stored. The adversary could thus modify the nullifier to a different value
2306 with the same projection. As the SNARK verification operates in $\mathbb{F}_{r_{\text{CUR}}}$, the proof would
2307 still be valid. However, the value stored for this nullifier would be different from the
2308 adversarial one. Hence, the nullifier would be validated, the transaction would succeed
2309 and the note would be double spent. In practice, the adversary can perform the attack
2310 by simply adding r_{CUR} to one of the elements representing the nullifier.

2311 To prevent this attack, the contract checks that all primary inputs are elements of
2312 $\mathbb{F}_{r_{\text{CUR}}}$, that is to say that they are smaller than r_{CUR} . As one may see, the attack described
2313 above is not due to the packing of hash digests into field elements but to the contract
2314 storage of field elements as **Ethereum** words.

2315 Appendix C

2316 Side-channel attacks and 2317 information leaks

2318 The following subsections describe several side-channel attacks and possible weaknesses
2319 that implementers should be aware of and attempt to mitigate.

2320 We consider cases in which the attacker is able to observe the RPC communications
2321 between Zeth client software, and Ethereum P2P nodes. This situation may occur if an
2322 observer is able to monitor the network traffic between the Ethereum node and the Zeth
2323 client software, or if the Ethereum node itself is compromised.

Note

In this discussion, we do not consider adversaries with physical access to the machine running the client software. Such adversaries could make precise measurements of timing, power-consumption or other physical quantities that could reveal fine-grained details of the operations being carried out by the software, or the data it is operating on. Protecting against attacks of this kind often involves implementation techniques such as: avoiding branches based on private data, being careful with memory access patterns, and making all operations constant time, to only name a few. We leave consideration of these attacks and prevention methods for a future discussion.

2324

2325 C.1 Counterfeit data

2326 Malicious Ethereum nodes or attackers able to compromise the network have the oppor-
2327 tunity to send invalid data to RPC clients. This could be used to inject invalid data
2328 into the client's record of state, which could prevent it from generating valid Mix calls
2329 or allow it to be identified in the future. In general, data from any remote host should
2330 be treated as malicious, unless accompanied by evidence that convinces the client of its
2331 authenticity.

2332 In the case of Ethereum event logs (the main source of data used to track the on-
2333 chain state – see Section 4.1.1 for details), clients **MUST** leverage the consensus evidence
2334 and block headers to verify that log data is genuine and has been committed to the
2335 blockchain. See Section 1.2.3 for further information about how such data is secured.

2336 C.2 Data leaked during synchronization

2337 In order to receive private payments and keep their local data up-to-date, **Zeth client**
2338 software **MUST** scan the blockchain and process *all* the event data emitted by **Mixer**
2339 during Mix calls (as described in Section 4.1.1). There are several issues to consider
2340 when determining exactly how and when this “synchronization” takes place.

2341 Client implementations that only connect to the RPC endpoint in response to user
2342 input, or in preparation for performing a Mix call, may leak information. Observers may
2343 deduce that such client are likely to be the recipient of a recent or upcoming transaction,
2344 or that they may be about to perform a Mix call.

2345 Similarly, payment provider software that only listens for events when awaiting a
2346 transaction, and remains disconnected otherwise, may reveal that it is the recipient of
2347 an upcoming transaction, and possibly *which* transaction or block it was paid by (based
2348 on when it stops listening).

2349 Further, consider wallet software that performs RPC operations to explicitly wait
2350 for the Ethereum transaction corresponding to a specific Mix call. This would most
2351 likely be for transactions emitted by the **Zeth client**, in order to inform the user and
2352 update the wallet state once the payment is complete (but could possibly happen on
2353 the receiver side, if he somehow knows the ID of the transaction of interest – e.g. via
2354 off-chain communication with the sender). If such a *wait* procedure is implemented by
2355 querying the status of a specific transaction by its ID, or by listening for blocks *until*
2356 the transaction of interest is received, the connected Ethereum node may infer that this
2357 client is interested in the transaction, and likely to be the sender or recipient.

2358 Consider a client which periodically connects to some Ethereum node and requests
2359 all relevant data from the last block it saw, up to the latest block available. Each client
2360 will have information up to some block n (where n varies per client), and n is known to
2361 the Ethereum node that served the client. The client could then potentially be identified
2362 by n (even if it hides its IP for each connection) since a client that connects and queries
2363 **Zeth** transactions from block $n + 1$ reveals that it is one of the clients who synced up to
2364 block n when it last connected.

2365 Note that, if the client always broadcasts the Mix transaction via this same Ethereum
2366 node, then the Ethereum node may already deduce that the client is the sender. However,
2367 implementations may wish to use techniques (such as sending transactions from other
2368 nodes or hiding their IP address in other way) to obfuscate any relationship between
2369 transactions and the clients that originated them.

2370 C.3 Queries on successful decryption

2371 The event data emitted by $\widetilde{\text{Mixer}}$ contains the note data for new commitments, en-
2372 crypted using a key derived from the recipients' public key. As described in Section 2.6,
2373 clients scan the blockchain for these events and attempt to decrypt the ciphertext using
2374 their secret decryption keys. If they are successful, they are the recipient of the note
2375 and can try to parse the plaintext to extract the secret note data.

2376 When decryption is successful and the note data has been extracted from the plain-
2377 text (we discuss parsing failure in Appendix C.4), clients **MUST** check that this note data
2378 does indeed open the commitment for the note.

2379 A naive implementation of this check could query the state of $\widetilde{\text{Mixer}}$ via RPC to
2380 check the relevant entry in the set of commitments. However, this would reveal to
2381 an observer that the client had successfully decrypted and parsed the corresponding
2382 ciphertext, and was therefore the recipient of that note.

2383 For this reason, the protocol specifies that $\widetilde{\text{Mixer}}$ **MUST** emit events informing clients
2384 of new commitment values and locations in the Merkle tree. Clients **MUST** consume
2385 *all* such data to maintain their view of contract state (as described in Appendix C.2).
2386 Further, clients **MUST** attempt to decrypt *all* ciphertexts and, for successful decryptions,
2387 **MUST** verify that the plaintext opens the note's commitment. This avoids the need for
2388 any extra RPC queries that would reveal which ciphertexts were successfully decrypted.

Note

Emitting events containing all data necessary to carry out the local checks imple-
mented in the wallet is a way to enforce that all wallets behave exactly the same
to the eyes of network (passive) adversaries (regardless whether the user is the
recipient of a note or not).

2389

2390 C.4 Invalid ciphertext

2391 The attack described in [TBP20, Section 4.2.1] illustrates the importance of correctly
2392 handling invalid data in client software. A so-called "REJECT Attack" is described
2393 whereby an attacker creates a Mix call with specially crafted ciphertext. The cipher-
2394 text can be successfully decrypted by the correct recipient – that is, the plaintext note
2395 is encrypted with an encryption key derived from the recipients public key – but the
2396 corresponding plaintext is invalid and cannot be parsed correctly by the recipient.

Note

Note that the above is possible because the plaintext is neither verified by the
circuit encoding \mathbf{R}^Z , nor by the contract (which is unable to decrypt it). Hence,
 $\widetilde{\text{Zeth}}$ allows such transactions with malicious ciphertexts to be accepted by the
 $\widetilde{\text{Mixer}}$ contract, and clients must handle this case with care.

2397

2398 In the case described in [TBP20], there is no distinction between “client” or “wallet”
2399 software, and the underlying P2P nodes. Before a fix was applied (see [zcab]), nodes
2400 explicitly rejected transactions of the above form, proving to their peers that they were
2401 able to decrypt the ciphertext and were therefore the intended recipient.

2402 In **Zeth**, P2P nodes and wallet software are separated, so there will be no explicit
2403 rejection of the transaction. However, careless error handling (such as exceptions which
2404 causes the RPC connection to be closed) could potentially be detected by the connected
2405 Ethereum node. As in the “REJECT Attack”, this reveals that the connected RPC
2406 client is the intended recipient of a transaction, and the owner of the corresponding
2407 encryption key.

2408 C.5 Using (and retrieving) nullifiers

2409 Any non-trivial wallet implementation will need to track which of the user’s **Zeth** notes
2410 have been spent, and which are still available. Naturally, the wallet software could mark
2411 the notes as it broadcasts transactions that spend them. However, this approach is
2412 subject to several problems.

2413 Firstly, for each note spent, the client software must record the ID of the spending
2414 transaction, in order to track it and confirm that it is accepted into a block. Once each
2415 spending transaction is accepted the client can finally mark the appropriate **Zeth** notes
2416 as “spent”. This requires significant complexity in order to asynchronously mark the
2417 notes, and to deal with the issues described in Appendix C.2.

2418 Secondly, this approach does not support multiple wallets using the same key, or
2419 wallets being restored from **Zeth** addresses. A user that wishes to rebuild his wallet
2420 (see the discussion in Section 4.1.4), or check for any spending activity by other wallets,
2421 would not be able to do so by simply scanning the blockchain.

2422 By using the nullifiers passed to **Mix** calls, clients can determine the availability of
2423 notes in a more robust way. That is, to determine whether a note is spent or available,
2424 the client can compute the nullifier and check whether that nullifier has been seen by
2425 the **Mixer** contract.

2426 In a similar way to Appendix C.3, queries to **Mixer** for specific nullifiers reveals
2427 to observers that the client was the sender of any previous or future transaction that
2428 generates such a nullifier. To mitigate this, **Mixer** **MUST** include nullifier values in the
2429 event data it emits, and clients **SHOULD** use this to track which of their notes are spent.
2430 This **MUST** happen as part of the regular sync operation, so that no extra RPC traffic is
2431 generated and observers cannot distinguish between clients that do and do not recognize
2432 any given nullifier. Note that this approach also supports tracking spent notes from
2433 multiple wallets, and rebuilding wallets by re-syncing the blockchain.

2434 C.6 Proof generation

2435 Generation of the zero-knowledge proofs, required for valid Mix calls, is a very computa-
2436 tionally intensive process. The proof generation itself does not require any communica-
2437 tion with external parties, and so may not directly leak information about the client, but
2438 implementers should consider some indirect ways in which information may be leaked.

2439 Implementers may also wish to consider the possible indirect impact of proof genera-
2440 tion on the RPC channel. For example, a client that “waits” for proof generation without
2441 servicing the RPC connection may fail to respond to (or take significantly longer to re-
2442 spond to) new log events. The connected Ethereum node might then deduce that its peer
2443 is generating a proof and therefore likely to be the sender of an upcoming transaction.

Note

As stated in the introduction to this chapter, this discussion does not consider general timing attacks. We mention this extreme case of a client that completely stalls during proof generation only to illustrate how a poor implementation may leak information to its RPC peer.

2444

2445 In the case where proof generation is carried out on some external host, or by an
2446 external process on the same host, there may be a risk of network traffic or other IPC
2447 traffic being observed. If an observer can detect that a given client is communicating
2448 with a prover process, it can reliably deduce that the client will be the sender of an
2449 upcoming transaction.

2450 An observer able to see the content of the communication between the wallet and
2451 prover process will also gain knowledge of the auxiliary inputs to the proof (including
2452 the data required to spend the input notes and secret attributes of the output notes).
2453 It is therefore important to secure any such connection, protect any prover process from
2454 being maliciously modified or observed, and to ensure that wallets only communicate
2455 with trusted processes.

2456 C.7 Simple mixer calls

2457 The public parameters to a Mix call can reveal information about the nature of a trans-
2458 action, even though they do not reveal recipient details or note amounts. For example,
2459 a Mix call in which $\text{Mix}_{in}.primIn.vout = 0$ and $\text{Mix}_{in}.primIn.vin \neq 0$ may indicate
2460 a simple “deposit” of funds into the mixer. Similarly, if both $\text{Mix}_{in}.primIn.vout$ and
2461 $\text{Mix}_{in}.primIn.vin$ are zero, the transaction must be spending only notes already within
2462 **Mixer**, into new notes. Finally, if $\text{Mix}_{in}.primIn.vin = 0$ and $\text{Mix}_{in}.primIn.vout \neq 0$, the
2463 sender may be performing a simple “withdrawal” of funds from some existing notes.

2464 A Mix call can combine all of the above logical operations in a single transaction.
2465 That is, it can deposit value into the mixer, spend existing notes, create new notes, and
2466 withdraw value from **Mixer** *at the same time*. Combining logical operations in this way

2467 makes it much more difficult for an observer to attribute a specific purpose to the Mix
2468 call.

2469 Clients can also perform Mix calls in which $vin = vout = 0$ and 0-valued notes are
2470 created from other 0-valued notes. Such “dummy” self-payments can further obfuscate
2471 the activity of a wallet, by adding “noise” to the system. Note, however, that the gas
2472 cost for such transactions must still be paid.

2473 Wallet implementations **SHOULD** encourage the use of these complex calls where pos-
2474 sible, either via the user interface or by automatically adding complexity to transactions,
2475 and **SHOULD** support features such as adding “noise”¹ if the user wishes to pay for extra
2476 protection of this kind.

2477 C.7.1 Small anonymity sets

2478 Until there is a large number of commitments and users of the mixer, it may be easy
2479 for an observer to infer some of the private data that is intended to be hidden by mixer
2480 calls.

2481 In the simple case, if there are very few commitments in the $\widetilde{\text{Mixer}}$ ’s Merkle tree, an
2482 attacker has a small list of candidate commitments that are being spent by subsequent
2483 Mix calls. Similarly, if the number of distinct Ethereum addresses that have been used
2484 to call $\widetilde{\text{Mixer}}$ is very small, observers can trace the original source of funds subsequently
2485 withdrawn to a small set of original depositors.

2486 Client software may wish to track metrics about the $\widetilde{\text{Mixer}}$ state, and either prevent
2487 certain actions or design the user interface to discourage users² from creating trans-
2488 actions whose features can be identified with high probability. We provide below a
2489 non-exhaustive list of metrics of interest:

- 2490 • **Number of commitments.** If there is a low absolute number of commitments,
2491 clearly any non-zero output must spend one of these (although we note that only
2492 $vout$ can be publicly known to be non-zero).
- 2493 • **Number of unspent commitments.** If $\#Comms - \#Nulls$ is small and a new
2494 commitment is created and then spent, observers can deduce that there is a high
2495 chance that the spend operation targeted the new commitment.
- 2496 • **Number of Ethereum addresses.** While very few distinct addresses (or groups
2497 of addresses that are not associated) have used the contract, observers can de-
2498 deduce that subsequent Mix calls are likely to spend commitments created by clients
2499 associated with one of this small set of addresses.

2500 The set of Ethereum addresses that have interacted with the contract can leak data
2501 in other ways. An Ethereum address that withdraws value from the contract, but has not
2502 previously been used to make a Mix call (or a Mix call that deposits value into $\widetilde{\text{Mixer}}$),

¹By randomly scheduling dummy payments, for instance

²By, for example, displaying warning messages and/or asking the user for confirmation

2503 must have been the recipient of zeth notes created by a previous depositor. The details
2504 may not be directly available to an observer, but this is another example of information
2505 which could be combined with other leaked data to infer connections between entities
2506 and transactions.

2507 Appendix D

2508 Security proofs of Blake2

2509 This appendix proves the collision resistance, PRF-ness, binding and hiding properties
2510 of the Blake2 hash function in the Weakly Ideal Cipher model (WICM, see [LMN16]).
2511 The proofs use definitions and results of Luykx et al. [LMN16], regarding the indifferen-
2512 tiability of Blake2 and a random oracle in the Weakly Ideal Cipher Model (WICM). In
2513 the following, we assume that the optimization of Blake2 for 8- to 32-bit platforms is as
2514 secure as Blake2 as described in [LMN16].

2515 D.1 Security model of Blake2

2516 The security analysis treats Blake2 as hash function built on top of a block-cipher-based
2517 compression function in the WICM (which derives from the Ideal Cipher Model). In
2518 this section, we present the WICM and prove that Blake2 is a collision resistant PRF,
2519 and thus a commitment scheme.

2520 D.1.1 Weakly Ideal Cipher Model

2521 The research community believes that Blake’s underlying block cipher has no known
2522 weaknesses and could reasonably be modeled as an ideal cipher [LMN16, Section 2.1].
2523 However, Blake2 admits weak keys with a specific structure [LMN16, Section 2.1]. Blake2
2524 is therefore more appropriately analysed in the WICM, which is an extension of the Ideal
2525 Cipher Model that represents a block cipher as a set of independent random permuta-
2526 tions [HKT11]. The WICM may also be viewed as a specialization for Blake2 of the Weak
2527 Cipher Model [MP15], which aims to be realistic by modeling particular characteristics,
2528 invariants or properties a block cipher may have.

2529 A number of definitions in what follows are quoted directly from Luykx et al. [LMN16].

The Weakly Ideal Cipher Model. Let \mathcal{W} and \mathcal{S} be the following partition of $\mathbb{B}^{2 \cdot \text{ol}}$
into weak and strong sets, where w is the word length ($16 \cdot w = 2 \cdot \text{ol}$):

$$\mathcal{W} = \left\{ \text{aaaabbbbccccddd} \in \mathbb{B}^{2 \cdot \text{ol}} \mid a, b, c, d \in \mathbb{B}^w \right\}$$

$$\mathcal{S} = \mathbb{B}^{2 \cdot \text{ol}} \setminus \mathcal{W}$$

Let $\mathcal{BLC}(2 \cdot \text{ol}, 2 \cdot \text{ol})$ denote the set of all block ciphers $E : \mathbb{B}^{2 \cdot \text{ol}} \times \mathbb{B}^{2 \cdot \text{ol}} \rightarrow \mathbb{B}^{2 \cdot \text{ol}}$. Define $\mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$ as the set of all block ciphers $E \in \mathcal{BLC}(2 \cdot \text{ol}, 2 \cdot \text{ol})$ with the additional restriction that $E(k_w, \cdot)$ is \mathcal{W} - and \mathcal{S} -subspace invariant for all keys $k_w \in \mathcal{K}_{\text{weak}}$. That is, inputs in \mathcal{W} map to \mathcal{W} , and likewise for \mathcal{S} . Here, $\mathcal{K}_{\text{weak}}$ is the set of weak keys, defined as

$$\mathcal{K}_{\text{weak}} = \left\{ k = \text{kkkkkkkkkkkkkkkkkk} \in \mathbb{B}^{2 \cdot \text{ol}} \mid k \in \mathbb{B}^w \right\}.$$

2530 A random $E \leftarrow \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$ can now be modeled as follows:

2531 • on input of $(k, x) \in \mathcal{K}_{\text{weak}} \times \mathcal{W}$, E generates its response y randomly from \mathcal{W} up
2532 to repetition;

2533 • on input of $(k, x) \in \mathcal{K}_{\text{weak}} \times \mathcal{S}$, E generates its response y randomly from \mathcal{S} up to
2534 repetition.

2535 For key values $k \in \mathbb{B}^{2 \cdot \text{ol}} \setminus \mathcal{K}_{\text{weak}}$, E behaves like an ideal cipher: it either outputs a
2536 new random value or if the key-message-image tuple has already been queried the tuple's
2537 image. The case of inverse queries is analogous.

Blake2C is defined over the following domains and codomain:

$$\text{Blake2C} : \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol}) \times \mathbb{B}^{\text{ol}} \times \mathbb{B}^{2 \cdot \text{ol}} \times \mathbb{B}^{\text{ol}/4} \times \mathbb{B}^{\text{ol}/4} \rightarrow \mathbb{B}^{\text{ol}}$$

2538 We write $\text{Blake2C}_E(h, m, t, f)$ for the output of the Blake2 compression function, defined
2539 over encryption scheme E on inputs h , m , t and f . The compression function, in par-
2540 ticular, computes the state $x = (h \parallel \text{pad}_{\text{ol}/2}(0) \parallel t \parallel f) \oplus (\text{pad}_{\text{ol}}(0) \parallel \text{IV})$ for some IV . It then
2541 encrypts x under m (where m is treated as a key for the encryption) and splits $E(m, x)$
2542 in two same size variables, the left part l_E and right part r_E . It finally outputs $l_E \oplus r_E \oplus h$.

2543

2544 Zeth uses the Blake2 compression function with a fixed encryption scheme E^* based on
2545 ChaCha stream cipher [Ber08a]. Thus, we write $\text{Blake2C}(h, m, t, f) = \text{Blake2C}_{E^*}(h, m, t, f)$.

2546 **Indifferentiability.** One way to measure the extent to which a certain cryptographic
2547 function behaves like a random function is via the indistinguishability framework where
2548 a distinguisher is given oracle access to either the cryptographic function or the random
2549 function with the goal of determining which one it has access to.

Definition D.1.1. Let \mathcal{C} be a construction with oracle access to an ideal primitive \mathcal{P} . Let \mathcal{R} be an ideal primitive with the same domain and codomain as \mathcal{C} . Let Sim be a simulator with the same domain and codomain as \mathcal{P} with oracle access to \mathcal{R} , and let Dist be a PPT distinguisher. The indifferentiability advantage of Dist is defined as:

$$\text{Indiff}_{\mathcal{C}^{\mathcal{P}}, \text{Sim}}(\text{Dist}) = \left| \Pr \left[\text{Dist}^{\mathcal{C}^{\mathcal{P}}, \mathcal{P}} = 1 \right] - \Pr \left[\text{Dist}^{\mathcal{R}, \text{Sim}^{\mathcal{R}}} = 1 \right] \right|$$

2550 The distinguisher Dist can query both its left oracle (either \mathcal{C} or \mathcal{R}) and its right
 2551 oracle (either \mathcal{P} or Sim). We refer to $\mathcal{C}^{\mathcal{P}}$, \mathcal{P} as the real world, and to \mathcal{R} , $\text{Sim}^{\mathcal{R}}$ as the
 2552 simulated world; the distinguisher Dist converses with either of these worlds and its goal
 2553 is to tell both worlds apart.

Theorem D.1.1 (Indifferentiability of Blake2 [LMN16]). *Let an encryption scheme $E \leftarrow \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$ be a weakly ideal cipher, and consider the hash function Blake2_E that internally uses E . There exists a simulator Sim such that for any distinguisher Dist with total complexity q , we have:*

$$\text{Indiff}_{\text{Blake2}_E, \text{Sim}}(\text{Dist}) \leq \frac{\binom{q}{2}}{2^{2\text{ol}}} + \frac{2\binom{q}{2}}{2^{\text{ol}}} + \frac{q}{2^{\text{ol}/2}}$$

2554 where Sim makes at most $O(q^3)$ queries to a random function \mathcal{R} .

2555 *Proof.* See [LMN16, Corollary 1]. □

2556 For asymptotic security, we assume the distinguisher to be PPT and that the number
 2557 of queries made is polynomial $q \leq \text{poly}(\text{ol})$.

2558 **Additional remarks.** Luykx et al. [LMN16] remark that, by resorting to the WICM,
 2559 they do not make stronger assumptions than those used in previous results (ICM), and,
 2560 despite the fact that they give distinguishers more power (by weakening the cipher),
 2561 they are able to get similar results.

2562 D.2 Security proofs

2563 D.2.1 Blake2 is a PRF

2564 Luykx et al. already prove the PRFness of Blake2 *keyed* hash function in the multi-key
 2565 setting.

Definition D.2.1 (PRF in multi-key setting [ML15]). Let $\mu \geq 1$ and $k \leftarrow \mathcal{K}^\mu$. Let \mathcal{C} be a keyed construction with key space \mathcal{K} and with oracle access to an ideal primitive \mathcal{P} . Let $\mathcal{R}_1, \dots, \mathcal{R}_\mu$ be random functions with the same domains and ranges as $\mathcal{C}_{k_1}, \dots, \mathcal{C}_{k_\mu}$. Let D be a distinguisher. The PRF distinguishing advantage of D is defined as,

$$\text{PRF}_{\mathcal{C}^{\mathcal{P}}}(\text{D}) = \left| \Pr[\text{Dist}_{\mathcal{C}_{k_1}^{\mathcal{P}}, \dots, \mathcal{C}_{k_\mu}^{\mathcal{P}}, \mathcal{P}} = 1] - \Pr[\text{Dist}_{\mathcal{R}_1, \dots, \mathcal{R}_\mu, \mathcal{P}} = 1] \right|$$

Blake2 supports keyed hashing by simply prepending the key to the message:

$$\text{Blake2}_{E,k}(m) = \text{Blake2}_E(k \| 0^{2\text{ol}-\text{kl}} \| m)$$

2566 where $\text{kl} \leq 2\text{ol}$ denotes the key size. In other words, the key gets processed as other data
 2567 and the HAIFA counter and flags are designated to the key in a similar fashion as if they
 2568 were for normal data blocks.

Theorem D.2.1 (PRF-security of Blake2 keyed mode [LMN16]). *Let $\mu \geq 1$ and let $k \leftarrow_{\$} (\mathbb{B}^{\text{kl}})^{\mu}$. Let an encryption scheme $E \leftarrow_{\$} \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$ be a weakly ideal cipher, and consider the keyed hash function $\text{Blake2}_{E,k}$ that internally uses Blake2C_E that internally uses E . For any distinguisher Dist with total complexity q :*

$$\text{PRF}_{\text{Blake2}_{E,k}}(\text{Dist}) \leq \frac{\binom{q}{2}}{2^{2\text{ol}}} + \frac{2\binom{q}{2}}{2^{\text{ol}}} + \frac{q}{2^{\text{ol}/2}} + \frac{\mu q}{2^{\text{kl}}} + \frac{\binom{\mu}{2}}{2^{\text{kl}}}$$

2569 *Proof.* See [LMN16, Corollary 3]. □

2570 **Remark D.2.2.** We can note that in the case of keyed hashing, the key is padded only
 2571 to be processed in a single block to differentiate the key from the message. The security
 2572 proof of Theorem D.2.1 does not rely on this padding and as such also works with no
 2573 padding.

Theorem D.2.2 (PRF-security of Blake2 with a single key [LMN16]). *Let $k \leftarrow_{\$} \mathbb{B}^{\text{kl}}$. Let an encryption scheme $E \leftarrow_{\$} \mathcal{BLC}^*(2 \cdot \text{ol}, 2 \cdot \text{ol})$ be a weakly ideal cipher, and consider the keyed hash function $\text{Blake2}_E(k, \cdot) = \text{Blake2}_E(k \parallel \cdot)$ that internally uses Blake2C_E that internally uses E . For any distinguisher Dist with total complexity q :*

$$\text{PRF}_{\text{Blake2}_E}(\text{Dist}) \leq \frac{\binom{q}{2}}{2^{2\text{ol}}} + \frac{2\binom{q}{2}}{2^{\text{ol}}} + \frac{q}{2^{\text{ol}/2}} + \frac{q}{2^{\text{kl}}}$$

2574 *Proof.* See Remark D.2.2 and Theorem D.2.1 with $\mu = 1$. □

2575 **Remark D.2.3.** Since we analyse the security of Blake2 asymptotically, we assume that
 2576 for a security parameter λ holds $\text{ol} = \mathcal{O}(\lambda)$, $\text{kl} = \mathcal{O}(\lambda)$, and $q = \text{poly}(\lambda)$.

2577 D.2.2 Proof of Blake2 collision resistance

2578 We want to prove here the collision resistance of Blake2. To do so, we are going to
 2579 prove by contradiction that if Blake2 is not collision resistant, it is not indifferentiable
 2580 according to Definition D.1.1.

2581 **Theorem D.2.3.** *Blake2 is collision resistant.*

2582 *Informal proof.* Let us assume that there exists a PPT adversary \mathcal{B} which breaks the
 2583 collision resistance of Blake2. We build an adversary \mathcal{A} that uses this adversary to
 2584 differentiate between the real and simulated worlds. More particularly, \mathcal{A} gets left and
 2585 right oracles (see [LMN16, Figure 3]), which are either an oracle for a hash function and
 2586 for a weakly ideal block cipher or a random oracle and an encryption simulator with
 2587 oracle access to the random oracle.

2588 On each \mathcal{B} 's query m_i , $i \in \{1, \dots, q\}$, \mathcal{A} passes them to his left oracle and returns
 2589 the answer h_i to \mathcal{B} . Eventually, if \mathcal{B} finds a collision, that is a pair (m_i, m_j) such that
 2590 $m_i \neq m_j$ and $h_i = h_j$, \mathcal{A} guesses that his oracles were real; else \mathcal{A} returns a random

2591 guess. Otherwise \mathcal{A} guesses his oracles were simulated – if the left oracle was a random
 2592 oracle, the probability of finding a collision would be negligible for $q \leq \text{poly}(\lambda)^1$.

2593 On the other hand, \mathcal{B} finds a collision with non-negligible probability if the oracles
 2594 were real. Hence, \mathcal{A} wins the indistinguishability game with non-negligible advantage,
 2595 which is a contradiction. \square

2596 D.2.3 Blake2 as a commitment scheme

2597 We prove here that Blake2 is a commitment scheme, i.e. is binding and hiding. To do so
 2598 we rely on the previous results that Blake2 is collision resistant and a PRF.

2599 **Theorem D.2.4.** *Let $E \leftarrow \text{\$BLK}(2\text{ol}, 2\text{ol})$ and for a message $x \in \mathbb{B}^*$ and randomness
 2600 $r \in \mathbb{B}^l$ commitment to x using r be $\text{ComSch.Com}(x; r) = \text{Blake2}_E(r \| x)$. Then ComSch
 2601 is hiding and binding.*

2602 *Informal proof. Hiding.* A commitment scheme ComSch is computationally hiding if,
 2603 knowing two potential openings, a PPT adversary cannot distinguish which was com-
 2604 mitted. Let us assume that there exists a PPT adversary \mathcal{B} which breaks the hiding
 2605 property of Blake2 with a non-negligible advantage η . We build an adversary \mathcal{A} that
 2606 uses \mathcal{B} to break the PRF property of Blake2 with advantage $\eta/2$.

2607 First, the PRF game is initiated, that is, the challenger chooses a random encryption
 2608 scheme E and key $k \in \mathbb{B}^l$ and instantiates two oracles $O^{\text{Blake2}_k} = \text{Blake2}_E(k, \cdot)$ and O^R
 2609 a random function. The challenger picks an oracle randomly and gives \mathcal{A} access to it.
 2610 \mathcal{B} sends q oracle queries m_1, \dots, m_q to \mathcal{A} (adaptively) who extends them with random
 2611 r_1, \dots, r_q and sends $r_i \| m_i$ to his left oracle. Given the answer from the oracle, \mathcal{A} returns
 2612 them to \mathcal{B} . Eventually, \mathcal{B} then outputs two challenge messages $(\tilde{m}_0, \tilde{m}_1)$ and sends them
 2613 to \mathcal{A} who randomly selects message \tilde{m}_b , extends it with r and sends $r \| \tilde{m}_b$ to his left oracle.
 2614 The oracle answers with y_b which is also sent to \mathcal{B} . Finally, \mathcal{B} returns the decision bit \tilde{b}
 2615 to \mathcal{A} . If $b = \tilde{b}$, \mathcal{A} answers to the challenger that the oracle was instantiating the PRF.
 2616 Otherwise, \mathcal{A} answers with a random guess. The advantage of \mathcal{A} equals advantage of \mathcal{B}
 2617 if it interacts with a real hash function. The advantage of \mathcal{A} equals half the advantage
 2618 of \mathcal{B} when interacting with a random oracle and simulator.

2619 *Binding.* A commitment scheme ComSch is said to be computationally binding if
 2620 it is infeasible to find x, x' and r, r' such that $x \neq x'$ and $\text{Com}(x; r) = \text{Com}(x'; r')$.
 2621 This is implied by collision resistance of Blake2. Thus if \mathcal{B} is an algorithm that breaks
 2622 the hiding property with advantage η , there is another algorithm \mathcal{A} that breaks Blake2
 2623 collision resistance with the same advantage. \square

¹The probability would be $\frac{q^2}{2^{\text{ol}}}$ which is negligible for a polynomial number of queries q . This is the sum of the probabilities of finding a collision when doing the i^{th} query. Indeed, let us suppose the adversary has done $i - 1$, $i > 2$, queries without finding a collision, i.e. he knows $i - 1$ distinct tuples of input-output. When receiving the i^{th} value, the adversary has thus $i - 1$ chance to find a collision. The probability for the new output to be equal to any of the previous outputs is thus $(i - 1) \cdot \frac{1}{2^{\text{ol}}}$ (as we are in the random oracle model). Summing this probability over all queries, we find the probability of finding a collision after doing q queries.

2624 Assuming that Blake2s is as secure as Blake2, a commitment scheme based on a
 2625 Blake2s, i.e. $\text{Com}(x; r) = \text{Blake2s}_E(r\|x)$ is hiding and binding.

2626 D.2.4 Proof of commitment scheme security

To prove the binding and hiding property of ComSch (see Section 3.1.2), we introduce the following commitment scheme ComSch^* ,

$$\begin{aligned} \text{ComSch}^*.\text{Setup} &: \{1^\lambda \text{ s.t. } \lambda \in \mathbb{N}\} \rightarrow \mathbb{B}^* \\ \text{ComSch}^*.\text{Com} &: \mathcal{B}\mathcal{L}\mathcal{K}^*(2 \cdot \text{BLAKE2sCLEN}, 2 \cdot \text{BLAKE2sCLEN}) \times \mathbb{B}^{2 \cdot \text{BLAKE2sCLEN}} \\ &\quad \times (\mathbb{B}^{\text{PRFADDRROUTLEN}} \times \mathbb{B}^{\text{PRFRHOOUTLEN}} \times \mathbb{B}^{\text{ZVALUELEN}}) \times \mathbb{B}^{\text{RTRAPLEN}} \rightarrow \mathbb{B}^{\text{BLAKE2sCLEN}} \end{aligned}$$

The commitment scheme is defined as follows,

$$\begin{aligned} \text{ComSch}^*.\text{Setup}(1^\lambda) &= pp^* = \epsilon \\ \text{ComSch}^*.\text{Com}(m = (apk, \rho, v); r) &= cm \\ &= \text{Blake2}_{E^*}(r\|apk\|\rho\|v) \end{aligned}$$

Given a commitment scheme ComSch^* , the bijective function $\text{decode}_{\mathbb{N}}(\cdot)$ and $p_\lambda \in \mathbb{N}$, a prime which can be represented using λ bits, we define the commitment scheme ComSch' as follows:

$$\begin{aligned} \text{ComSch}'.\text{Setup}(1^\lambda) &= (\text{ComSch}^*.\text{Setup}(1^\lambda), p_\lambda) \\ \text{ComSch}'.\text{Com}(m; r) &= \text{decode}_{\mathbb{N}}(\text{ComSch}^*.\text{Com}(m; r)) \pmod{p_\lambda} \text{ for } m = (apk\|\rho\|v) \end{aligned}$$

2627 Note that ComSch (see Section 3.1.2) is a particular instantiation of ComSch' where E^*
 2628 is set as ChaCha encryption scheme [Ber08a], k^* is a random key, and p_λ is r_{CUR} .

2629 **Theorem D.2.5** (Hiding). *If ComSch^* is hiding then ComSch' is hiding.*

2630 *Proof.* We prove the theorem by contradiction i.e. we assume that there exists an adver-
 2631 sary \mathcal{B} that breaks ComSch' 's hiding property and construct an adversary \mathcal{A} that uses
 2632 \mathcal{B} to break ComSch^* 's hiding property with non-negligible probability.

2633 Let \mathcal{C} be a challenger that sets up the hiding game for ComSch^* and \mathcal{A} . The adversary
 2634 \mathcal{A} , given public parameters pp^* of ComSch^* and access to an oracle that runs the Com
 2635 algorithm of ComSch^* scheme, simulates a hiding game for ComSch' for \mathcal{B} . The adversary
 2636 \mathcal{A} starts by setting public parameters pp' for ComSch' using public parameters pp^*
 2637 given by \mathcal{C} . Parameters pp' are passed to \mathcal{B} who outputs a pair of messages m_0, m_1 .
 2638 The adversary \mathcal{A} forwards them to the challenger who samples a bit b at random and
 2639 generates $cm^* = \text{ComSch}^*.\text{Com}(m_b; r)$ for some randomness r . The result is returned
 2640 to \mathcal{A} (see Definition 1.5.21). Then \mathcal{A} passes $cm = \text{decode}_{\mathbb{N}}(cm^*) \pmod{p_\lambda}$ to \mathcal{B} who
 2641 returns his guess b' . The adversary \mathcal{A} returns the same b' to the challenger.

2642 By construction, it is clear that \mathcal{A} wins the hiding game with the same probability
 2643 that \mathcal{B} wins the simulated hiding game. Since \mathcal{B} 's advantage is non-negligible, this means
 2644 that \mathcal{A} wins the ComSch^* hiding game with non-negligible probability as well. \square

2645 **Theorem D.2.6** (Binding). *Let ComSch^* be a computationally binding commitment*
 2646 *scheme and $\text{ComSch}^*. \text{Com}$ indiffereniable from a random oracle. Then ComSch' is also*
 2647 *computationally binding if $l = \lceil 2^\lambda / p_\lambda \rceil$ is at most $\text{poly}(\lambda)$.*

2648 *Proof.* Assume that \mathcal{A} asks the ComSch' commit and open oracles a total of q_λ distinct
 2649 queries. Let us denote the result of the q_λ queries and output of the attacker (the
 2650 candidate collision) as $((m_1, r_1, y_1), \dots, (m_{q_\lambda}, r_{q_\lambda}, y_{q_\lambda}), \text{out})$. If \mathcal{A} is successful it means
 2651 that it outputs $(m, r), (m', r')$ such that $(m, r) \neq (m', r')$ and $\text{ComSch}'. \text{Com}(m; r) =$
 2652 $\text{ComSch}'. \text{Com}(m'; r')$.

By the definition of ComSch' , we have that,

$$\text{ComSch}'. \text{Com}(m; r) = \text{decode}_{\mathbb{N}}(\text{ComSch}^*. \text{Com}(m; r)) \pmod{p_\lambda}$$

Hence, we have a collision in ComSch' if there exists $k \in [l]$, l being the ratio of the
 codomains of $\text{ComSch}^*. \text{Com}$ and $\text{ComSch}'. \text{Com}$, such that,

$$|\text{decode}_{\mathbb{N}}(\text{ComSch}^*. \text{Com}(m; r)) - \text{decode}_{\mathbb{N}}(\text{ComSch}^*. \text{Com}(m'; r'))| = k \cdot p_\lambda.$$

2653 We show that this event is unlikely.

2654 In fact, for each $i \in [q_\lambda]$, let C_i be the event that the adversary wins at the i -th
 2655 query. That is, the last commitment y_i is a ComSch' collision with one of the previous
 2656 y_j . More precisely there exists $j \leq i$ and $k < l$ such that $y_i = y_j + k \cdot p_\lambda$.

2657 Since ComSch^* is a random oracle, y_i is randomly selected from a set of at least p_λ
 2658 elements. As such, we have $\Pr[C_i] \leq i \cdot l / p_\lambda$.

Thus the probability of finding a collision after q_λ queries is $\Pr[C_1 \vee \dots \vee C_{q_\lambda}] \leq$
 $\sum_{i=1}^{q_\lambda} \Pr[C_i] = l / p_\lambda \cdot \sum_{i=1}^{q_\lambda} i$. This probability is bounded by $l \cdot \frac{q_\lambda(q_\lambda+1)}{p_\lambda}$. However,
 we allow only polynomial number of queries. Thus for $q_\lambda = \text{poly}(\lambda)$ this probability
 becomes,

$$\frac{2^\lambda \cdot \text{poly}(\lambda)}{p_\lambda^2},$$

2659 what is negligible for $2^\lambda / p_\lambda \leq \text{poly}(\lambda)$. □

2660 **Remark D.2.4.** Note that in **Zeth**'s commitment scheme, we set $p_\lambda = \mathbf{r}_{\text{CUR}}$ and $2^\lambda =$
 2661 $2^{\text{BLAKE2sCLEN}}$. Thus, for **BN-254** and **BLS12-377** have $l = 6$ and $l = 14$, respectively.
 2662 Therefore, the probability of an attacker breaking the binding property due to reduction
 2663 modulo \mathbf{r}_{CUR} increases approximately by these factors. This is still negligible.

2664 **Corollary.** *Assume that Blake2 is indiffereniable from a random oracle and a PRF,*
 2665 *then ComSch^* is computationally binding and computationally hiding. Furthermore, the*
 2666 *reduction is tight. That is, the advantage of any PPT adversary against the binding*
 2667 *(resp. hiding) property is the same as the advantage of an adversary against collision*
 2668 *resistance and binding (resp. hiding).*

2669 Appendix E

2670 Fuzzy message detection

2671 As explained in 2.6 and 4.1.1, in order to receive *ZethNotes*, a **Zeth** user must listen
2672 on a broadcast channel, and try to decrypt all encrypted events emitted by the **Mixer**
2673 contract. While providing the best potential for indistinguishability (all users scan the
2674 chain data and expose the same behavior), such routine is particularly expensive to
2675 carry out, especially for computationally restricted users (i.e. users with computationally
2676 limited devices).

2677 As a way to trade-off the users' anonymity and the cost of the message detection
2678 routine in privacy-preserving protocols, Beck et al. [BLMG21] introduced the notion
2679 of *fuzzy message detection schemes*. These protocols allow the delegation of message
2680 detection to untrustworthy servers, without revealing precisely which messages belong to
2681 the receiver, by allowing receivers to enforce false-positive detection rates. Such schemes
2682 provide a promising avenue for reconciling recipient anonymity (via *key ambiguity* and
2683 *message detection ambiguity*) and the performance of the *ZethNotes* receiving algorithm
2684 that currently needs to run on a machine belonging to (or trusted by) the recipient.

2685 Nevertheless, the selection of the fuzzy detection parameters for **Zeth** is a chal-
2686 lenge, especially the selection of the false-positive rate. Under the scheme presented
2687 in [BLMG21], not only is this parameter public (an additional “leakage” of informa-
2688 tion¹, including to potentially adversarial nodes), but this parameter is likely to be set
2689 to different values by different users, based on the number of payments they receive
2690 through **Zeth**. This, coupled with the existing gas-related leakages, will increase the set
2691 of information leakages in the protocol, the consequences of which are hard to properly
2692 estimate. Furthermore, letting such parameters be set by users raises other challenges
2693 for wallet developers, user experience (UX) engineers and documentation engineers. In
2694 fact, any degree of liberty given to the user increases the potential for “deviation” from
2695 the “expected/indistinguishable” behavior. Hence, UX/documentation/wallet engineers
2696 must be able to suggest sensible default values for such parameters, must extensively
2697 document the purpose of these parameters and must extensively educate the end-users
2698 to maximize the chances of adequate parameter selections. While feasible, such tasks

¹limited to one server (in the best case), or to the whole network (in the worst case — if the adversary broadcasts all its known information)

2699 largely rely on modeling efforts², which simplify real-world systems and can only be used
2700 to simulate a limited set of situations. Moreover, not being able to easily (i.e. without
2701 distributing new keys) update the false-positive rate over time is problematic in the
2702 context of **Zeth** as it does not allow users to have adaptable false-positive probabilities
2703 to account for potential spikes in the number of payments they receive (e.g. a merchant
2704 during sales).

2705 On the other hand, and as mentioned above, being able to use *fuzzy message detection*
2706 *schemes* in **Zeth** would also widen the user base of the protocol, which, as a consequence,
2707 would widen the anonymity set.

²See e.g. <https://git.openprivacy.ca/openprivacy/fuzzytags-sim>

2708 Appendix F

2709 Extended discussion on the 2710 security of MIMC in different 2711 settings

2712 In the original design proposed in the MIMC paper [AGR⁺16], the round function is
2713 represented as a “shifted” permutation via a cubic map (i.e. the round input message is
2714 added to the key and round constant - the “shift” -, and a map x^e permutes the element
2715 in the underlying field \mathbb{F}_{2^n} , $n \in \mathbb{N}$, where $\gcd(e, 2^n - 1) = 1$). This function (which is
2716 a permutation and therefore invertible) acts as a substitution box (S-box) and brings
2717 non-linearity to the scheme, as usually required for security.

2718 In other sections of the paper, however, the MIMC authors proposed generalizations
2719 to the initial design. These allow MIMC to be used:

- 2720 • over prime fields of odd characteristic (i.e. \mathbb{F}_p , p odd prime),
- 2721 • with different permutation polynomials (i.e. using different exponents in the round
2722 function)

2723 Understanding the relationships between these various settings is required in order to
2724 use MIMC to operate over prime fields of odd characteristic with a non-cubic permutation
2725 monomial.

2726 Overall, for MIMC to be considered secure, it is important that no attack in the
2727 literature (that may provide a significant speedup compared to “exhaustive key search”)
2728 can successfully be mounted by a PPT adversary. Two main families of attacks exist:
2729 statistical attacks and algebraic attacks.

2730 **Statistical attacks.** In the “Security Analysis” section [AGR⁺16, Section 4.2], the au-
 2731 thors explain that since the cubic function is an Almost Perfect Nonlinear map (APN)¹,
 2732 linear attacks pose no threat to MIMC.

2733 We observe that this claim aligns with [HRS99, Theorem 2]. In fact if MIMC is
 2734 operated over \mathbb{F}_{2^n} , it is easy to see that the degree of the cubic map can be expressed as
 2735 $3 = 2^t + 1$, where $t = 1$, and where n and 1 are trivially coprime. This case is covered
 2736 by [HRS99, Theorem 2] which confirms that the cubic function $S(x) = x^3$ over \mathbb{F}_{2^n} is an
 2737 APN/2-uniform mapping, as desired for differential and linear cryptanalysis resistance.

2738 Likewise, in [AGR⁺16, Section 5.1] the authors claim that, provided that the cubic
 2739 map is a permutation over the prime field of interest, MIMC can be used to operate over
 2740 prime fields of odd characteristic. In this case too, $S(x) = x^3$ is an APN, provided $p \neq 3$
 2741 (as reported by [HRS99, Theorem 3, item 3]).

2742 In [AGR⁺16, Section 5.3] the choice of the map degree is relaxed to be of the general
 2743 form $2^t \pm 1$. Unfortunately, the authors showed that the case $e = 2^t + 1$ is not as good as
 2744 it initially seems in \mathbb{F}_{2^n} , due to term cancellation in fields of characteristic 2 that renders
 2745 the resulting polynomial sparse². More precisely the degree of the polynomial will be
 2746 bounded by 3^r , r being the number of rounds, which does not constitute an improvement
 2747 on the case of $e = 3$. For this very reason, exponents of the form $2^t + 1, t > 1$, may not
 2748 be of interest (sparse polynomial and more expensive arithmetic in the round function).

2749 Likewise, if the map degree e is chosen to be of the form $2^t + 1$, with $\gcd(e, 2^n - 1) = 1$
 2750 in the context of MIMC over \mathbb{F}_{2^n} , it is necessary to bear in mind that, without the extra
 2751 requirement that t needs to be coprime with n , then this case is not covered by [HRS99,
 2752 Theorem 2], and $S(x) = x^e$ does not have differential 2-uniformity anymore - violating
 2753 the claim made in [AGR⁺16, Section 4.2] paragraph “Linear Attacks” about optimal
 2754 resistance against linear and differential cryptanalysis. (In fact, depending on the value
 2755 of $g = \gcd(n, t)$, the map $S(x) = x^{2^t+1}$ would be differentially 2^g -uniform [Nyb93] -
 2756 contrasting with the setting considered in paragraph “Linear attacks” where $e = 3$).
 2757 The case $e = 2^t - 1$ does not yield an APN over \mathbb{F}_{2^n} either (except in the case $t = 2$
 2758 which reduces to the case $2^{t'} + 1$ for $t' = 1$). Similar observations show that picking
 2759 round function degrees of the form $2^t \pm 1$ in the context where MIMC is defined over
 2760 prime fields \mathbb{F}_p , p odd prime, does not yield APNs.

Overall, when studying the resistance of MIMC against statistical attacks, we are
 interested in the probability that an input difference (d) is mapped into an output
 difference (D). That is, we are studying the probability of the following event:

$$F(x + d) - F(x) = D$$

2761 where F is a function that may either represent a single round, a set of rounds, or the
 2762 full cipher.

¹A function $f(x) = x^e$ over \mathbb{F}_{p^n} is said to be differentially k -uniform if k is the maximum number of solutions $x \in \mathbb{F}_{p^n}$ of $f(x + d) - f(x) = D$ where $d, D \in \mathbb{F}_{p^n}$ and $d \neq 0$. If f is a 2-uniform mapping, we say that f is almost perfect nonlinear. See e.g. [HRS99] for more information.

²Since the round function is a polynomial, the whole scheme can be seen as a polynomial with overall degree and “sparsity” that depends on the underlying field characteristic, degree of the round function and number of rounds.

2763 Over a single round of MIMC (i.e. $F(x) = S(x)$), this probability is bounded by $(e -$
 2764 $1)/p$ provided that the exponent e is “small” (i.e. small compared to the size of the field).
 2765 By assuming that the different rounds of the scheme are independent, the probability
 2766 that an input difference gets mapped to an output difference, when F represents the full
 2767 cipher, becomes bounded as $\Pr[F(x + d) - F(x) = D] \leq ((e - 1)/p)^{\text{rounds}}$.

For security, we want $((e - 1)/p)^{\text{rounds}}$ to be bounded by $2^{-\lambda}$, where λ is the security level (e.g. 128). Hence, we need

$$\left(\frac{e - 1}{p}\right)^{\text{rounds}} \leq 2^{-\lambda}$$

that is, we want

$$\text{rounds} \geq \frac{\lambda}{\log_2\left(\frac{p}{e-1}\right)}$$

2768 As such, if the exponent is much smaller than the size of the field, few rounds are
 2769 sufficient to prevent the differential attacks.

2770 **Algebraic attacks.** While permutation monomials of degree $e = 2^t \pm 1$ may not
 2771 constitute APNs in the various MIMC settings, it is important to note (as highlighted by
 2772 Grassi in [GR21]) that when working over finite fields \mathbb{F}_p of large prime characteristic p
 2773 or extension fields \mathbb{F}_{2^n} of large extension degrees n , the algebraic attacks (exploiting the
 2774 low-degree of the cipher) are much more efficient than the statistical attacks (i.e. they
 2775 can break a much higher number of rounds).

2776 In fact, when considering security against algebraic attacks, we want (roughly speak-
 2777 ing) the polynomial that defines the cipher to be of maximum degree and full (or at
 2778 least, dense). That is we want the degree of the polynomial to be higher than 2^λ . Since
 2779 in MIMC the S-box is defined as $S(x) = x^e$, then after rounds rounds the degree of the
 2780 polynomial describing the cipher is e^{rounds} . Hence, we need

$$e^{\text{rounds}} \geq 2^\lambda$$

that is, we want

$$\text{rounds} \geq \lambda \log_e(2)$$

2781 **Remark F.0.1.** More rounds may be required as advised in [EGL⁺20] to prevent some
 2782 algebraic attacks that can be mounted when MIMC is used over binary fields \mathbb{F}_{2^n} .

2783 It is important to note that the security analysis related to algebraic attacks relies
 2784 on the fact that the polynomial describing the cipher is dense/full. If this assumption
 2785 is violated, a more granular security analysis needs to be carried out for the setting of
 2786 interest.

Note

For *small exponents and large prime fields* (e.g. for $\lambda = 128$, $p = 2^{128}$ and $e = 3$), we see that the lower bound on the number of rounds is (much) smaller in the context of statistical attacks than in the context of algebraic attacks. As such, we see that in such settings algebraic attacks are much more powerful than statistical attacks. Hence, when instantiating MIMC with a *small* exponent of the form $2^t \pm 1$, it is crucial to make sure that, even if the resulting map is not an APN, the polynomial describing the cipher remains full/dense. Importantly, if the setting is changed (e.g. to use exponents that are “big” w.r.t. the field size) the security analysis proposed in [AGR⁺16] must be changed.

2787

2788

2789 Glossary

2790 **joinsplit** Set of JSIN input *ZethNotes*, and JSOUT output *ZethNotes* as well as the public
2791 values *vin* and *vout* used in a tx_{Mix} transaction. 37, 39, 41, 61, 100, 125, 126

2792 **joinsplit equation** Equation that checks that the sum of the values of the SendTx
2793 algorithm of DAP is equal to the sum of the values of its outputs. This equations
2794 checks that the joinsplit is “balanced” and thus, that no value is created while
2795 creating new *ZethNotes*. 25, 41, 42, 61, 100, 125

2796 **Acronyms**

2797 **APN** Almost Perfect Nonlinear (function). 123, 125

2798 **DOS** Denial of Service (Attack). 16, 125

2799 **ECC** Elliptic Curve Cryptography. 54, 125

2800 **EOA** Externally Owned Account. 16, 17, 19, 125

2801 **EVM** Ethereum Virtual Machine. 15, 16, 20, 49, 55, 65, 86, 89, 125

2802 **FFT** Fast Fourier Transform. 87, 125

2803 **MAC** Message Authentication Code. 102, 125

2804 **NFS** Number Field Sieve. 54, 125

2805 **PoC** Proof of Concept. 125

2806 **RAM** Random-access Memory. 87, 125

2807 **RLP** Recursive Length Prefix. 19, 20, 125

2808 Bibliography

- 2809 [AAM12] Imad Fakhri Alshaikhli, Mohammad A Alahmad, and Khanssaa Munthir.
2810 Comparison and analysis study of sha-3 finalists. In *2012 International*
2811 *Conference on Advanced Computer Science Applications and Technologies*
2812 *(ACSAT)*, pages 366–371. IEEE, 2012.
- 2813 [abi] Contract abi specification, section "function selector". [https://solidity.readthedocs.io/en/develop/abi-spec.html#](https://solidity.readthedocs.io/en/develop/abi-spec.html#function-selector)
2814 [function-selector](https://solidity.readthedocs.io/en/develop/abi-spec.html#function-selector).
2815
- 2816 [ABM⁺03] Adrian Antipa, Daniel Brown, Alfred Menezes, René Struik, and Scott
2817 Vanstone. Validation of elliptic curve public keys. In *International Work-*
2818 *shop on Public Key Cryptography*, pages 211–223. Springer, 2003.
- 2819 [ABN10] Michel Abdalla, Mihir Bellare, and Gregory Neven. Robust encryption.
2820 In *Theory of Cryptography Conference*, pages 480–497. Springer, 2010.
2821 <https://eprint.iacr.org/2008/440.pdf>.
- 2822 [ABR99] Michel Abdalla, Mihir Bellare, and Phillip Rogaway. Dhaes:
2823 An encryption scheme based on the diffie-hellman prob-
2824 lem. 1999. [https://pdfs.semanticscholar.org/95f4/](https://pdfs.semanticscholar.org/95f4/63d097086fba325086a4cf88706648dafd09.pdf)
2825 [63d097086fba325086a4cf88706648dafd09.pdf](https://pdfs.semanticscholar.org/95f4/63d097086fba325086a4cf88706648dafd09.pdf).
- 2826 [ABR01] Michel Abdalla, Mihir Bellare, and Phillip Rogaway. Dhies: An en-
2827 cryptation scheme based on the diffie-hellman problem., 2001. [https://](https://web.cs.ucdavis.edu/~rogaway/papers/dhies.pdf)
2828 web.cs.ucdavis.edu/~rogaway/papers/dhies.pdf.
- 2829 [ACG⁺19] Martin R Albrecht, Carlos Cid, Lorenzo Grassi, Dmitry Khovratovich,
2830 Reinhard Lüftenegger, Christian Rechberger, and Markus Schofnegger.
2831 Algebraic cryptanalysis of stark-friendly designs: application to marvel-
2832 lous and mimc. In *International Conference on the Theory and Appli-*
2833 *cation of Cryptology and Information Security*, pages 371–397. Springer,
2834 2019.
- 2835 [AFK⁺08] Jean-Philippe Aumasson, Simon Fischer, Shahram Khazaei, Willi Meier,
2836 and Christian Rechberger. New features of latin dances: analysis of salsa,
2837 chacha, and rumba. In *International Workshop on Fast Software Encryp-*
2838 *tion*, pages 470–488. Springer, 2008.

- 2839 [AG18] Andreas M. Antonopoulos and Wood Gavin. *Mastering Ethereum*.
2840 O’Reilly Media, 2018.
- 2841 [AGM⁺09] Kazumaro Aoki, Jian Guo, Krystian Matusiewicz, Yu Sasaki, and Lei
2842 Wang. Preimages for step-reduced sha-2. In *International Conference*
2843 *on the Theory and Application of Cryptology and Information Security*,
2844 pages 578–597. Springer, 2009. [https://link.springer.com/content/
2845 pdf/10.1007/978-3-642-10366-7_34.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-10366-7_34.pdf).
- 2846 [AGR⁺16] Martin Albrecht, Lorenzo Grassi, Christian Rechberger, Arnab Roy, and
2847 Tyge Tiessen. Mimc: Efficient encryption and cryptographic hashing
2848 with minimal multiplicative complexity. In *International Conference on*
2849 *the Theory and Application of Cryptology and Information Security*, pages
2850 191–219. Springer, 2016.
- 2851 [AHMP08] Jean-Philippe Aumasson, Luca Henzen, Willi Meier, and Raphael C-W
2852 Phan. Sha-3 proposal blake. *Submission to NIST*, 229:230, 2008.
- 2853 [ALM12] Elena Andreeva, Atul Luykx, and Bart Mennink. Provable security of
2854 blake with non-ideal compression function. In *International Conference*
2855 *on Selected Areas in Cryptography*, pages 321–338. Springer, 2012.
- 2856 [AMP10] Elena Andreeva, Bart Mennink, and Bart Preneel. Security reductions
2857 of the second round sha-3 candidates. In *International Conference on*
2858 *Information Security*, pages 39–53. Springer, 2010.
- 2859 [AMPŠ12] Elena Andreeva, Bart Mennink, Bart Preneel, and Marjan Škrobot. Se-
2860 curity analysis and comparison of the sha-3 finalists blake, grøstl, jh,
2861 keccak, and skein. In *International Conference on Cryptology in Africa*,
2862 pages 287–305. Springer, 2012.
- 2863 [ANWOW13] Jean-Philippe Aumasson, Samuel Neves, Zooko Wilcox-O’Hearn, and
2864 Christian Winnerlein. Blake2: simpler, smaller, fast as md5. In *Interna-*
2865 *tional Conference on Applied Cryptography and Network Security*, pages
2866 119–135. Springer, 2013. <https://eprint.iacr.org/2013/322.pdf>.
- 2867 [AS09] Kazumaro Aoki and Yu Sasaki. Meet-in-the-middle preimage attacks
2868 against reduced sha-0 and sha-1. In *Annual International Cryptology Con-*
2869 *ference*, pages 70–89. Springer, 2009.
- 2870 [BBDP01] Mihir Bellare, Alexandra Boldyreva, Anand Desai, and David Pointcheval.
2871 Key-privacy in public-key encryption. In *International Conference on the*
2872 *Theory and Application of Cryptology and Information Security*, pages
2873 566–582. Springer, 2001. [https://iacr.org/archive/asiacrypt2001/
2874 22480568.pdf](https://iacr.org/archive/asiacrypt2001/22480568.pdf).

- 2875 [BCC⁺15] Jonathan Bootle, Andrea Cerulli, Pyrros Chaidos, Essam Ghadafi, Jens
2876 Groth, and Christophe Petit. Short accountable ring signatures based on
2877 ddh. In *European Symposium on Research in Computer Security*, pages
2878 243–265. Springer, 2015.
- 2879 [BCD⁺20] Tim Beyne, Anne Canteaut, Itai Dinur, Maria Eichlseder, Gregor Le-
2880 ander, Gaëtan Leurent, María Naya-Plasencia, Léo Perrin, Yu Sasaki,
2881 Yosuke Todo, and Friedrich Wiemer. Out of oddity – new cryptana-
2882 lytic techniques against symmetric primitives optimized for integrity proof
2883 systems. Cryptology ePrint Archive, Report 2020/188, 2020. <https://eprint.iacr.org/2020/188>.
2884
- 2885 [BCG⁺20] Sean Bowe, Alessandro Chiesa, Matthew Green, Ian Miers, Pratyush
2886 Mishra, and Howard Wu. ZEXE: enabling decentralized private com-
2887 putation. In *2020 IEEE Symposium on Security and Privacy, SP 2020,*
2888 *San Francisco, CA, USA, May 18-21, 2020*, pages 947–964. IEEE, 2020.
- 2889 [BCK⁺18] Elaine Barker, Lily Chen, Sharon Keller, Allen Roginsky, Apostol
2890 Vassilev, and Richard Davis. Recommendation for pair-wise key-
2891 establishment schemes using discrete logarithm cryptography. Technical
2892 report, National Institute of Standards and Technology, 2018. [Online;
2893 last accessed 10-January-2020].
- 2894 [BD07] Eli Biham and Orr Dunkelman. A framework for iterative hash
2895 functions—haifa. Technical report, Computer Science Department, Tech-
2896 nion, 2007. <https://eprint.iacr.org/2007/278.pdf>.
- 2897 [BDPVA07] Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche.
2898 Sponge functions. In *ECRYPT hash workshop*, volume 2007. Citeseer,
2899 2007.
- 2900 [Ber05] Daniel J Bernstein. The poly1305-aes message-authentication code.
2901 In *International Workshop on Fast Software Encryption*, pages 32–49.
2902 Springer, 2005. <https://cr.yp.to/mac/poly1305-20050329.pdf>.
- 2903 [Ber06] Daniel J Bernstein. Curve25519: new diffie-hellman speed records. In
2904 *International Workshop on Public Key Cryptography*, pages 207–228.
2905 Springer, 2006. <https://cr.yp.to/ecdh/curve25519-20060209.pdf>.
- 2906 [Ber08a] Daniel J Bernstein. Chacha, a variant of salsa20. In *Workshop Record*
2907 *of SASC*, volume 8, pages 3–5, 2008. [https://cr.yp.to/chacha/
2908 chacha-20080120.pdf](https://cr.yp.to/chacha/chacha-20080120.pdf).
- 2909 [Ber08b] Daniel J. Bernstein. New stream cipher designs. chapter The Salsa20 Fam-
2910 ily of Stream Ciphers, pages 84–97. Springer-Verlag, Berlin, Heidelberg,
2911 2008.

- 2912 [BGM17] Sean Bowe, Ariel Gabizon, and Ian Miers. Scalable multi-party compu-
 2913 tation for zk-snark parameters in the random beacon model. *Cryptology*
 2914 ePrint Archive, Report 2017/1050, 2017. [https://eprint.iacr.org/](https://eprint.iacr.org/2017/1050)
 2915 2017/1050.
- 2916 [BL] Daniel J. Bernstein and Tanja Lange. Safecurves: choosing safe curves
 2917 for elliptic-curve cryptography. <https://safecurves.cr.jp.to>. [Online;
 2918 last accessed 09-December-2019].
- 2919 [BLMG21] Gabrielle Beck, Julia Len, Ian Miers, and Matthew Green. Fuzzy message
 2920 detection. *IACR Cryptol. ePrint Arch.*, 2021:89, 2021.
- 2921 [Blo70] Burton H. Bloom. Space/time trade-offs in hash coding with allowable
 2922 errors. *Commun. ACM*, 13(7):422–426, 1970.
- 2923 [Bon19] Xavier Bonnetain. Collisions on feistel-mimc and univariate gmimc. 2019.
- 2924 [Bou03] Nicolas Bourbaki. *Elements of mathematics: Algebra*. Springer, 2003.
- 2925 [Bra97] S. Bradner. Key words for use in rfc's to indicate requirement levels. RFC
 2926 2119, RFC Editor, March 1997.
- 2927 [BRS02] John Black, Phillip Rogaway, and Thomas Shrimpton. Black-box
 2928 analysis of the block-cipher-based hash-function constructions from
 2929 pgv. In *Annual International Cryptology Conference*, pages 320–335.
 2930 Springer, 2002. [https://link.springer.com/content/pdf/10.1007/](https://link.springer.com/content/pdf/10.1007/3-540-45708-9_21.pdf)
 2931 3-540-45708-9_21.pdf.
- 2932 [BS07] Mihir Bellare and Sarah Shoup. Two-tier signatures, strongly unforge-
 2933 able signatures, and fiat-shamir without random oracles. In *International*
 2934 *Workshop on Public Key Cryptography*, pages 201–216. Springer, 2007.
- 2935 [BSCG⁺14] Eli Ben Sasson, Alessandro Chiesa, Christina Garman, Matthew Green,
 2936 Ian Miers, Eran Tromer, and Madars Virza. Zerocash: Decentralized
 2937 anonymous payments from bitcoin. In *2014 IEEE Symposium on Security*
 2938 *and Privacy*, pages 459–474. IEEE, 2014.
- 2939 [CHM⁺20] Alessandro Chiesa, Yuncong Hu, Mary Maller, Pratyush Mishra, Noah
 2940 Vesely, and Nicholas P. Ward. Marlin: Preprocessing zkSNARKs with
 2941 universal and updatable SRS. pages 738–768, 2020.
- 2942 [Cle19] Clearmatics. Zeth Spoken release. [https://github.com/clearmatics/](https://github.com/clearmatics/zeth/releases/tag/v0.2)
 2943 [zeth/releases/tag/v0.2](https://github.com/clearmatics/zeth/releases/tag/v0.2), 2019. [Online; released 04-April-2019].
- 2944 [CM16] Arka Rai Choudhuri and Subhamoy Maitra. Differential cryptanalysis of
 2945 salsa and chacha-an evaluation with a hybrid model. *IACR Cryptology*
 2946 *ePrint Archive*, 2016:377, 2016. [https://eprint.iacr.org/2016/377.](https://eprint.iacr.org/2016/377.pdf)
 2947 pdf.

- 2948 [CM17] Arka Rai Choudhuri and Subhamoy Maitra. Significantly improved multi-
2949 bit differentials for reduced round salsa and chacha. *IACR Transactions*
2950 *on Symmetric Cryptology*, 2016(2):261–287, Feb. 2017.
- 2951 [DG09] George Danezis and Ian Goldberg. Sphinx: A compact and provably
2952 secure mix format. In *30th IEEE Symposium on Security and Privacy*
2953 *(S&P 2009), 17-20 May 2009, Oakland, California, USA*, pages 269–282.
2954 IEEE Computer Society, 2009.
- 2955 [EFK15] Thomas Espitau, Pierre-Alain Fouque, and Pierre Karpman. Higher-order
2956 differential meet-in-the-middle preimage attacks on sha-1 and blake. In
2957 *Annual Cryptology Conference*, pages 683–701. Springer, 2015. <https://eprint.iacr.org/2015/515>.
2958
- 2959 [EGL⁺20] Maria Eichlseder, Lorenzo Grassi, Reinhard Lüftenegger, Morten Øygarden,
2960 Christian Rechberger, Markus Schafneggler, and Qingju Wang. An
2961 algebraic attack on ciphers with low-degree round functions: Application
2962 to full mimc. In Shiho Moriai and Huaxiong Wang, editors, *Advances in*
2963 *Cryptology - ASIACRYPT 2020 - 26th International Conference on the*
2964 *Theory and Application of Cryptology and Information Security, Daejeon,*
2965 *South Korea, December 7-11, 2020, Proceedings, Part I*, volume 12491 of
2966 *Lecture Notes in Computer Science*, pages 477–506. Springer, 2020.
- 2967 [est] Estream project. <https://en.wikipedia.org/wiki/ESTREAM>.
- 2968 [Gab19] Ariel Gabizon. AuroraLight: Improved prover efficiency and SRS size in
2969 a sonic-like system. Cryptology ePrint Archive, Report 2019/601, 2019.
2970 <https://eprint.iacr.org/2019/601>.
- 2971 [GFBR06] Decio Gazzoni Filho, Paulo SLM Barreto, and Vincent Rijmen. The
2972 maelstrom-0 hash function. In *Brazilian Symposium on Information and*
2973 *Computer System Security*. , 2006.
- 2974 [GGPR13] Rosario Gennaro, Craig Gentry, Bryan Parno, and Mariana Raykova.
2975 Quadratic span programs and succinct nizks without pcps. In Thomas
2976 Johansson and Phong Q. Nguyen, editors, *Advances in Cryptology - EU-*
2977 *ROCRYPT 2013, 32nd Annual International Conference on the Theory*
2978 *and Applications of Cryptographic Techniques, Athens, Greece, May 26-*
2979 *30, 2013. Proceedings*, volume 7881 of *Lecture Notes in Computer Science*,
2980 pages 626–645. Springer, 2013.
- 2981 [GJMG11] B Guido, D Joan, P Michaël, and VA Gilles. The keccak sha-3 submission.
2982 2011.
- 2983 [GKN⁺14] Jian Guo, Pierre Karpman, Ivica Nikolić, Lei Wang, and Shuang Wu.
2984 Analysis of blake2. In *Cryptographers’ Track at the RSA Conference*,

- 2985 pages 402–423. Springer, 2014. <https://eprint.iacr.org/2013/467.pdf>.
2986
- 2987 [GLRW10] Jian Guo, San Ling, Christian Rechberger, and Huaxiong Wang. Ad-
2988 vanced meet-in-the-middle preimage attacks: First results on full tiger,
2989 and improved results on md4 and sha-2. In *International Conference on*
2990 *the Theory and Application of Cryptology and Information Security*, pages
2991 56–75. Springer, 2010. <https://eprint.iacr.org/2010/016.pdf>.
- 2992 [GM17] Jens Groth and Mary Maller. Snarky signatures: Minimal signatures
2993 of knowledge from simulation-extractable snarks. In Jonathan Katz and
2994 Hovav Shacham, editors, *Advances in Cryptology - CRYPTO 2017 - 37th*
2995 *Annual International Cryptology Conference, Santa Barbara, CA, USA,*
2996 *August 20-24, 2017, Proceedings, Part II*, volume 10402 of *Lecture Notes*
2997 *in Computer Science*, pages 581–612. Springer, 2017.
- 2998 [Gol01] Oded Goldreich. *Foundations of Cryptography: Basic Tools*, volume 1.
2999 Cambridge University Press, Cambridge, UK, 2001.
- 3000 [Gor93] Daniel M. Gordon. Discrete logarithms in $GF(P)$ using the number field
3001 sieve. *SIAM J. Discret. Math.*, 6(1):124–138, 1993.
- 3002 [GR21] Lorenzo Grassi and Antoine Rondelet. Private communications. 2021.
- 3003 [Gro06] Jens Groth. Simulation-sound NIZK proofs for a practical language and
3004 constant size group signatures. pages 444–459, 2006.
- 3005 [Gro16] Jens Groth. On the size of pairing-based non-interactive arguments. In
3006 *Annual International Conference on the Theory and Applications of Cryptographic*
3007 *Techniques*, pages 305–326. Springer, 2016.
- 3008 [GRR⁺16] Lorenzo Grassi, Christian Rechberger, Dragos Rotaru, Peter Scholl, and
3009 Nigel P Smart. Mpc-friendly symmetric key primitives. In *Proceedings of*
3010 *the 2016 ACM SIGSAC Conference on Computer and Communications*
3011 *Security*, pages 430–443. ACM, 2016. [https://eprint.iacr.org/2016/](https://eprint.iacr.org/2016/542)
3012 [542](https://eprint.iacr.org/2016/542).
- 3013 [GWC19] Ariel Gabizon, Zachary J. Williamson, and Oana Ciobotaru. PLONK:
3014 Permutations over lagrange-bases for oecumenical noninteractive argu-
3015 ments of knowledge. Cryptology ePrint Archive, Report 2019/953, 2019.
3016 <https://eprint.iacr.org/2019/953>.
- 3017 [Hao14] Yonglin Hao. The boomerang attacks on blake and blake2. In *International*
3018 *Conference on Information Security and Cryptology*, pages 286–310.
3019 Springer, 2014. <https://eprint.iacr.org/2014/1012.pdf>.

- 3020 [Har19] HarryR. Conversation about Miyaguchi-Preneel security. <https://github.com/HarryR/ethsnarks/issues/119>, 2019. Online; accessed June-2019.
- 3021
3022
- 3023 [HG20] Youssef El Housni and Aurore Guillevic. Optimized and secure pairing-friendly elliptic curves suitable for one layer proof composition. Cryptology ePrint Archive, Report 2020/351, 2020. <https://eprint.iacr.org/2020/351>.
- 3024
3025
3026
- 3027 [HKT11] Thomas Holenstein, Robin Künzler, and Stefano Tessaro. The equivalence of the random oracle model and the ideal cipher model, revisited. In *Proceedings of the forty-third annual ACM symposium on Theory of computing*, pages 89–98. ACM, 2011.
- 3028
3029
3030
- 3031 [HMRS12] Ekawat Homsirikamol, Paweł Morawiecki, Marcin Rogawski, and Marian Srebrny. Security margin evaluation of sha-3 contest finalists through sat-based attacks. In *IFIP International Conference on Computer Information Systems and Industrial Management*, pages 56–67. Springer, 2012.
- 3032
3033
3034
- 3035 [Hop16] Daira Hopwood. Daira’s comment on: ”ensure spec retains distinctness assumption in hsig”, 2016.
- 3036
- 3037 [HRS99] Tor Helleseth, Chunming Rong, and Daniel Sandberg. New families of almost perfect nonlinear power mappings. *IEEE Trans. Inf. Theory*, 45(2):475–485, 1999.
- 3038
3039
- 3040 [IS09] Takatori Isobe and Kyoji Shibutani. Preimage attacks on reduced tiger and sha-2. In *International Workshop on Fast Software Encryption*, pages 139–155. Springer, 2009. https://link.springer.com/content/pdf/10.1007/978-3-642-03317-9_9.pdf.
- 3041
3042
3043
- 3044 [Ish12] Tsukasa Ishiguro. Modified version of” latin dances revisited: New analytic results of salsa20 and chacha”. 2012. <https://eprint.iacr.org/2012/065.pdf>.
- 3045
3046
- 3047 [JMV01] Don Johnson, Alfred Menezes, and Scott Vanstone. The elliptic curve digital signature algorithm (ecdsa). *International journal of information security*, 1(1):36–63, 2001.
- 3048
3049
- 3050 [Joh16] Nick Johnson. Response to ”how does ethereum make use of bloom filters?”. <https://ethereum.stackexchange.com/questions/3418/how-does-ethereum-make-use-of-bloom-filters>, 2016. [Online; last accessed 10-January-2020].
- 3051
3052
3053
- 3054 [KL14] Jonathan Katz and Yehuda Lindell. *Introduction to modern cryptography*. Chapman and Hall/CRC, 2014. <https://repo.zenk-security.com/Cryptographie%20.%20Algorithmes%20.%20Steganographie/Introduction%20to%20Modern%20Cryptography.pdf>.
- 3055
3056
3057

- 3058 [KMO⁺13] Markulf Kohlweiss, Ueli Maurer, Cristina Onete, Björn Tackmann, and
3059 Daniele Venturi. Anonymity-preserving public-key encryption: A con-
3060 structive approach. In Emiliano De Cristofaro and Matthew K. Wright,
3061 editors, *Privacy Enhancing Technologies - 13th International Symposium,*
3062 *PETS 2013, Bloomington, IN, USA, July 10-12, 2013. Proceedings*, vol-
3063 ume 7981 of *Lecture Notes in Computer Science*, pages 19–39. Springer,
3064 2013.
- 3065 [KRS12] Dmitry Khovratovich, Christian Rechberger, and Alexandra Savelieva.
3066 Bicliques for preimages: attacks on skein-512 and the sha-2 family. In
3067 *International Workshop on Fast Software Encryption*, pages 244–263.
3068 Springer, 2012. [https://link.springer.com/content/pdf/10.1007/
3069 978-3-642-34047-5_15.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5_15.pdf).
- 3070 [Lab19] Matter Labs. Merkle shrubs. [https://github.com/matter-labs/
3071 MerkleShrubs](https://github.com/matter-labs/MerkleShrubs), 2019.
- 3072 [LHT16] Adam Langley, Mike Hamburg, and Sean Turner. Elliptic curves for se-
3073 curity. RFC 7748, <https://tools.ietf.org/pdf/rfc7748.pdf>, 2016.
- 3074 [LIS12] Ji Li, Takanori Isobe, and Kyoji Shibutani. Converting meet-in-the-
3075 middle preimage attack into pseudo collision attack: Application to sha-2.
3076 In *International Workshop on Fast Software Encryption*, pages 264–286.
3077 Springer, 2012. [https://link.springer.com/content/pdf/10.1007/
3078 978-3-642-34047-5_16.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-34047-5_16.pdf).
- 3079 [LM11] Mario Lamberger and Florian Mendel. Higher-order differential attack
3080 on reduced sha-256. *IACR Cryptology ePrint Archive*, 2011:37, 2011.
3081 <https://eprint.iacr.org/2011/037.pdf>.
- 3082 [LMN16] Atul Luykx, Bart Mennink, and Samuel Neves. Security analysis of
3083 blake2’s modes of operation. *IACR Transactions on Symmetric Cryp-*
3084 *tology*, pages 158–176, 2016. [https://www.esat.kuleuven.be/cosic/
3085 publications/article-2705.pdf](https://www.esat.kuleuven.be/cosic/publications/article-2705.pdf).
- 3086 [LN18] Adam Langley and Yoav Nir. Chacha20 and poly1305 for ietf protocols.
3087 *RFC 8439*, 2018. <https://tools.ietf.org/html/rfc8439>.
- 3088 [LP19] Chaoyun Li and Bart Preneel. Improved interpolation attacks on cryp-
3089 tographic primitives of low algebraic degree. In *International Conference*
3090 *on Selected Areas in Cryptography*, pages 171–193. Springer, 2019.
- 3091 [Mai16] Subhamoy Maitra. Chosen iv cryptanalysis on reduced round chacha and
3092 salsa. *Discrete Applied Mathematics*, 208:88–97, 2016.
- 3093 [MBKM19] Mary Maller, Sean Bowe, Markulf Kohlweiss, and Sarah Meiklejohn.
3094 Sonic: Zero-knowledge SNARKs from linear-size universal and updatable
3095 structured reference strings. pages 2111–2128, 2019.

- 3096 [MJS15] Ed. M.-J. Saarinen. Blake Compression Function F. <https://tools.ietf.org/html/rfc7693#section-3.2>, 2015. [Online; accessed
3097 November-2019].
3098
- 3099 [ML15] Nicky Mouha and Atul Luykx. Multi-key security: The even-mansour
3100 construction revisited. In *Annual Cryptology Conference*, pages 209–223.
3101 Springer, 2015. <https://hal.inria.fr/hal-01240988/document>.
- 3102 [MNS11] Florian Mendel, Tomislav Nad, and Martin Schl affer. Finding sha-2 char-
3103 acteristics: searching through a minefield of contradictions. In *Inter-
3104 national Conference on the Theory and Application of Cryptology and
3105 Information Security*, pages 288–307. Springer, 2011. [https://link.
3106 springer.com/content/pdf/10.1007/978-3-642-25385-0_16.pdf](https://link.springer.com/content/pdf/10.1007/978-3-642-25385-0_16.pdf).
- 3107 [Moh10] Payman Mohassel. A closer look at anonymity and robustness in encryp-
3108 tion schemes. In *Advances in Cryptology - ASIACRYPT 2010 - 16th In-
3109 ternational Conference on the Theory and Application of Cryptology and
3110 Information Security*, volume 6477 of *Lecture Notes in Computer Science*,
3111 pages 501–518. Springer, 2010.
- 3112 [Mon87] Peter L. Montgomery. Speeding the Pollard and elliptic curve methods of
3113 factorization. *Math. Comp.*, 48(177):243–264, 1987.
- 3114 [MP15] Bart Mennink and Bart Preneel. On the impact of known-key attacks on
3115 hash functions. In *International Conference on the Theory and Applica-
3116 tion of Cryptology and Information Security*, pages 59–84. Springer, 2015.
3117 <https://eprint.iacr.org/2015/909.pdf>.
- 3118 [MQZ10] Mao Ming, He Qiang, and Shaokun Zeng. Security analysis of blake-
3119 32 based on differential properties. In *2010 International Conference on
3120 Computational and Information Sciences*, pages 783–786. IEEE, 2010.
- 3121 [MVOV96] Alfred J Menezes, Paul C Van Oorschot, and Scott A Vanstone. Handbook
3122 of applied cryptography. [http://citeseerx.ist.psu.edu/viewdoc/
3123 download?doi=10.1.1.99.2838&rep=rep1&type=pdf](http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.99.2838&rep=rep1&type=pdf), 1996.
- 3124 [NA19] Samuel Neves and Filipe Araujo. An observation on norx, blake2, and
3125 chacha. *Information Processing Letters*, 149:1–5, 2019.
- 3126 [Nyb93] Kaisa Nyberg. Differentially uniform mappings for cryptography. In Tor
3127 Helleseth, editor, *Advances in Cryptology - EUROCRYPT '93, Workshop
3128 on the Theory and Application of of Cryptographic Techniques, Lofthus,
3129 Norway, May 23-27, 1993, Proceedings*, volume 765 of *Lecture Notes in
3130 Computer Science*, pages 55–64. Springer, 1993.
- 3131 [oST15] National Institute of Standards and Technology. Secure Hash Stan-
3132 dard (SHS). [https://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.
3133 180-4.pdf](https://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.180-4.pdf), 2015.

- 3134 [Per17] Trevor Perrin. X25519 and zero outputs. [https://moderncrypto.org/
3135 mail-archive/curves/2017/000896.html](https://moderncrypto.org/mail-archive/curves/2017/000896.html), 2017. [Online; last accessed
3136 08-January-2020].
- 3137 [PHE⁺17] Ania M. Piotrowska, Jamie Hayes, Tariq Elahi, Sebastian Meiser, and
3138 George Danezis. The loopix anonymity system. In Engin Kirda and
3139 Thomas Ristenpart, editors, *26th USENIX Security Symposium, USENIX
3140 Security 2017, Vancouver, BC, Canada, August 16-18, 2017*, pages 1199–
3141 1216. USENIX Association, 2017.
- 3142 [Por13] Thomas Pornin. Deterministic Usage of the Digital Signature Algorithm
3143 (DSA) and Elliptic Curve Digital Signature Algorithm (ECDSA). [https://
3144 tools.ietf.org/html/rfc6979](https://tools.ietf.org/html/rfc6979), 2013.
- 3145 [Pro14] Gordon Procter. A security analysis of the composition of chacha20 and
3146 poly1305. *IACR Cryptology ePrint Archive*, 2014:613, 2014.
- 3147 [PV05] Pascal Paillier and Damien Vergnaud. Discrete-log-based signatures may
3148 not be equivalent to discrete log. In *International Conference on the
3149 Theory and Application of Cryptology and Information Security*, pages
3150 1–20. Springer, 2005.
- 3151 [Qu99] Minghua Qu. Sec 2: Recommended elliptic curve domain parameters. *Cer-
3152 ticom Res., Mississauga, ON, Canada, Tech. Rep. SEC2-Ver-0.6*, 1999.
- 3153 [Rk19] Antoine Rondelet and @karalabe. Go-ethereum BN256 package.
3154 [https://github.com/ethereum/go-ethereum/blob/master/crypto/
3155 bn256/cloudflare/constants.go](https://github.com/ethereum/go-ethereum/blob/master/crypto/bn256/cloudflare/constants.go), 2019. [Online; released 28-May-2019].
- 3156 [Ron20] Antoine Rondelet. Zecale: Reconciling privacy and scalability on
3157 ethereum, 2020.
- 3158 [RZ19] Antoine Rondelet and Michal Zajac. ZETH: On Integrating Zerocash on
3159 Ethereum. [Online; released April-2019], 2019.
- 3160 [SS08] Somitra Kumar Sanadhya and Palash Sarkar. New collision attacks
3161 against up to 24-step sha-2. In *International conference on cryptology in
3162 India*, pages 91–103. Springer, 2008. [https://eprint.iacr.org/2008/
3163 270.pdf](https://eprint.iacr.org/2008/270.pdf).
- 3164 [Ste15] Stevens, Marc. On collisions for md5. [https://www.win.tue.
3165 nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.
3166 %20Stevens.pdf](https://www.win.tue.nl/hashclash/On%20Collisions%20for%20MD5%20-%20M.M.J.%20Stevens.pdf), 2015.
- 3167 [SZFW12] Zhenqing Shi, Bin Zhang, Dengguo Feng, and Wenling Wu. Improved
3168 key recovery attacks on reduced-round salsa20 and chacha. In *Interna-
3169 tional Conference on Information Security and Cryptology*, pages 337–351.
3170 Springer, 2012.

- 3171 [TBP20] Florian Tramèr, Dan Boneh, and Kenneth G. Paterson. Remote side-
3172 channel attacks on anonymous transactions. Cryptology ePrint Archive,
3173 Report 2020/220, 2020. <https://eprint.iacr.org/2020/220>.
- 3174 [THH15] Piotr Dyrąga Tjaden Hess, Matt Luongo and James Hancock. EIP 152:
3175 Add BLAKE2 compression function ‘F’. [https://eips.ethereum.org/](https://eips.ethereum.org/EIPS/eip-152)
3176 [EIPS/eip-152](https://eips.ethereum.org/EIPS/eip-152), 2015. [Online; accessed November-2019].
- 3177 [VNP10] Janoš Vidali, Peter Nose, and Enes Pašalić. Collisions for variants of the
3178 blake hash function. *Information processing letters*, 110(14-15):585–590,
3179 2010.
- 3180 [W⁺] Gavin Wood et al. Ethereum: A secure decentralised generalised trans-
3181 action ledger.
- 3182 [wc] Ethereum wiki contributors. Patricia tree. [https://github.com/](https://github.com/ethereum/wiki/wiki/Patricia-Tree)
3183 [ethereum/wiki/wiki/Patricia-Tree](https://github.com/ethereum/wiki/wiki/Patricia-Tree).
- 3184 [wc19] Ethereum wiki contributors. RLP. [https://github.com/ethereum/](https://github.com/ethereum/wiki/wiki/RLP)
3185 [wiki/wiki/RLP](https://github.com/ethereum/wiki/wiki/RLP), 2019. [Online; accessed December-2019].
- 3186 [wik] Bitcoin wiki. Secp256k1. <https://en.bitcoin.it/wiki/Secp256k1>.
3187 [Online; last accessed 04-January-2020].
- 3188 [Woo19] Dr Gavin Wood. ETHEREUM: A Secure Decentralised Gener-
3189 alised Transaction Ledger Byzantium. [https://ethereum.github.io/](https://ethereum.github.io/yellowpaper/paper.pdf)
3190 [yellowpaper/paper.pdf](https://ethereum.github.io/yellowpaper/paper.pdf), 2019. [VERSION 7e819ec - 2019-10-20].
- 3191 [zcaa] On the security of sprout/sapling in-band en-
3192 cryption. [https://forum.zcashcommunity.com/t/](https://forum.zcashcommunity.com/t/on-the-security-of-sprout-sapling-in-band-encryption/34986)
3193 [on-the-security-of-sprout-sapling-in-band-encryption/34986](https://forum.zcashcommunity.com/t/on-the-security-of-sprout-sapling-in-band-encryption/34986).
- 3194 [zcab] "zcash alerts - security announcement 2019-09-24".
3195 [https://z.cash/support/security/announcements/](https://z.cash/support/security/announcements/security-announcement-2019-09-24/a)
3196 [security-announcement-2019-09-24/a](https://z.cash/support/security/announcements/security-announcement-2019-09-24/a).
- 3197 [ZCa19] ZCash. ZCash protocol specification. [https://github.com/zcash/zips/](https://github.com/zcash/zips/blob/master/protocol/protocol.pdf)
3198 [blob/master/protocol/protocol.pdf](https://github.com/zcash/zips/blob/master/protocol/protocol.pdf), 2019. [Online; initially released
3199 14-December-2015].